


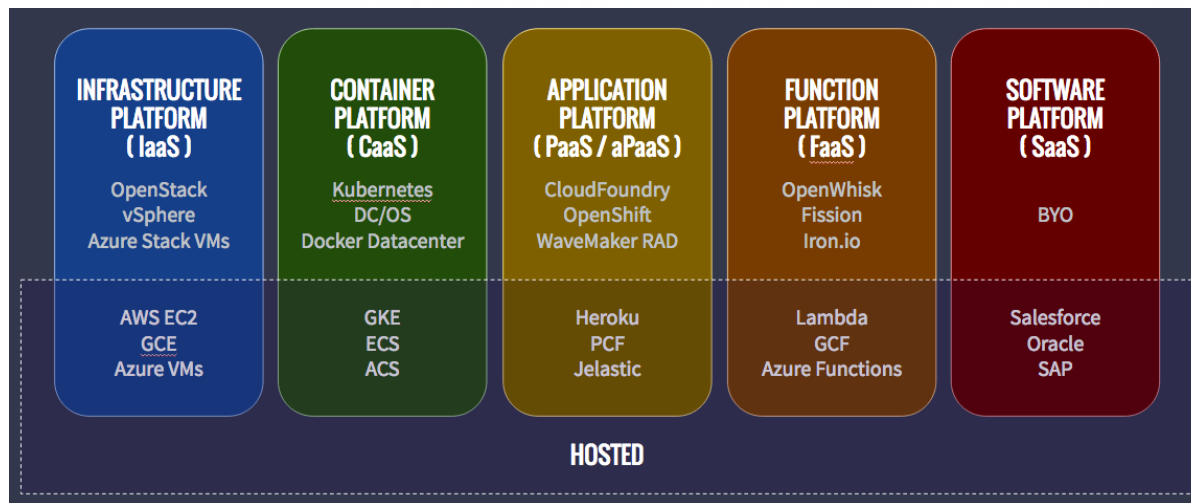
Tackling Cold Start in Serverless Computing with Container Runtime Reusing

Kun Suo*, Yong Shi*, Xiaohua Xu*
Dazhao Cheng⁺, Wei Chen[^]

*Kennesaw State University,

⁺University of North Carolina at Charlotte, [^]Nvidia

Serverless & Its Workflow



Serverless at the Edge




idle



Event
Trigger



Run ML/
DL code



Size

Plate

Speed

...

Serverless Pro vs Con

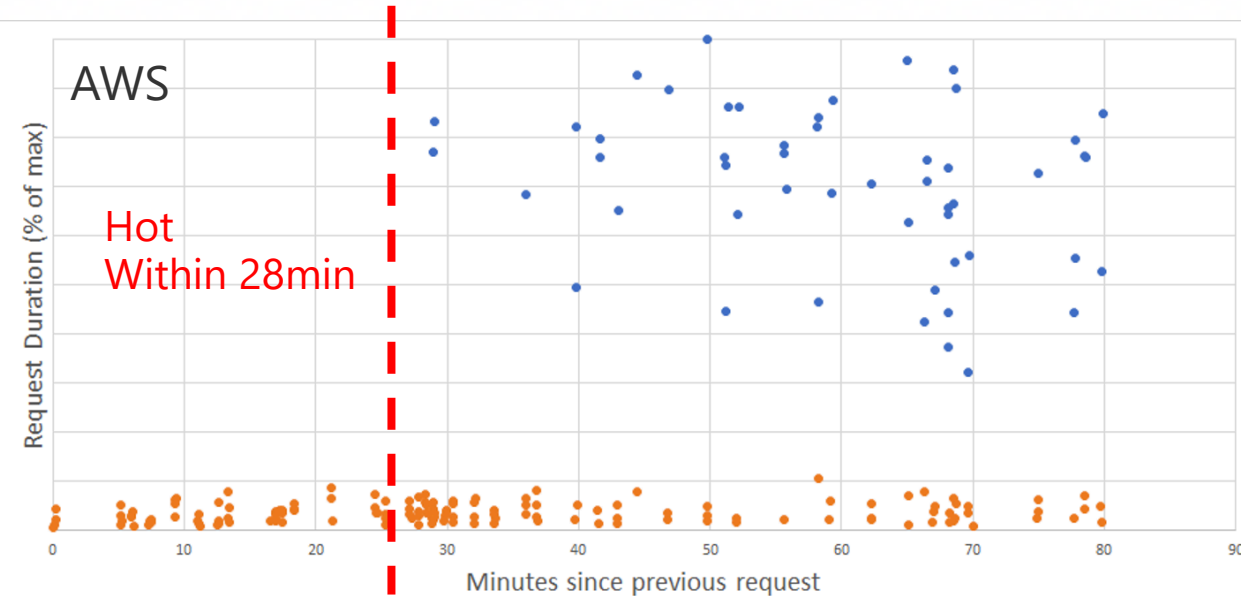
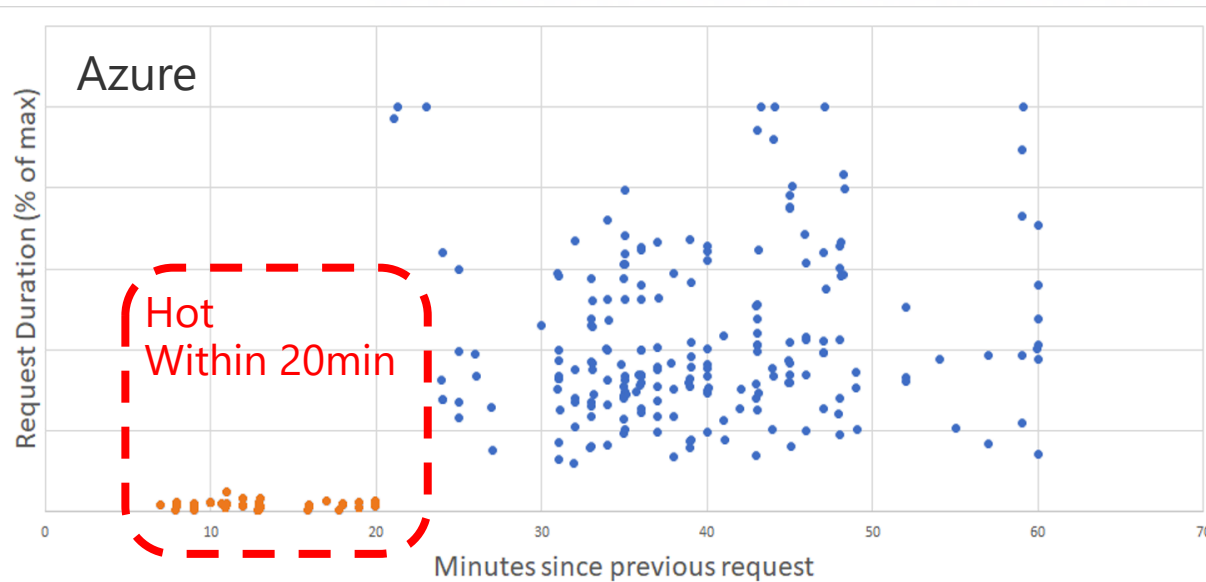
Benefits:

- Drop cost
- Fast Dev/deployment
- More security
- Microservice support
- Auto scale
- ...

Drawbacks:

- Not good for long time apps
- Deep dependency on platform
- Cold start
- Lack of debugging or monitoring
- ...

Cold Start Impact



AWS Lambda Cold Start Demo

Python:
http request return
random number

Measure the
request latency

The screenshot shows the AWS Lambda Management Console interface. The function name is 'test'. The 'Test' button is highlighted. The results section shows the following JSON output:

```
{
  "statusCode": 200,
  "number": 92
}
```

Below the output is a 'Summary' section with the following details:

Code SHA-256	Request ID
sGK910fBuTtqRilDvyMmBh5WRbACHBR3lwm4fy/EVEQ=	85bd5671-6a88-4efb-ba3a-4d677a1007df

The 'Duration' is highlighted with a red box and shows '0.38 ms'. Other metrics include 'Billed duration: 100 ms', 'Resources configured: 128 MB', and 'Max memory used: 23 MB'.

The 'Log output' section shows the following log entries:

```
START RequestId: 85bd5671-6a88-4efb-ba3a-4d677a1007df Version: $LATEST
END RequestId: 85bd5671-6a88-4efb-ba3a-4d677a1007df
REPORT RequestId: 85bd5671-6a88-4efb-ba3a-4d677a1007df Duration: 0.38 ms Billed
Duration: 100 ms Memory Size: 128 MB Max Memory Used: 23 MB
```

At the bottom of the console, there is a 'Feedback' button, a language selector set to 'English (US)', a copyright notice '© 2008 - 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved.', a timer showing '00:01', and a 'Finish' button.

Containerized Application Characteristics

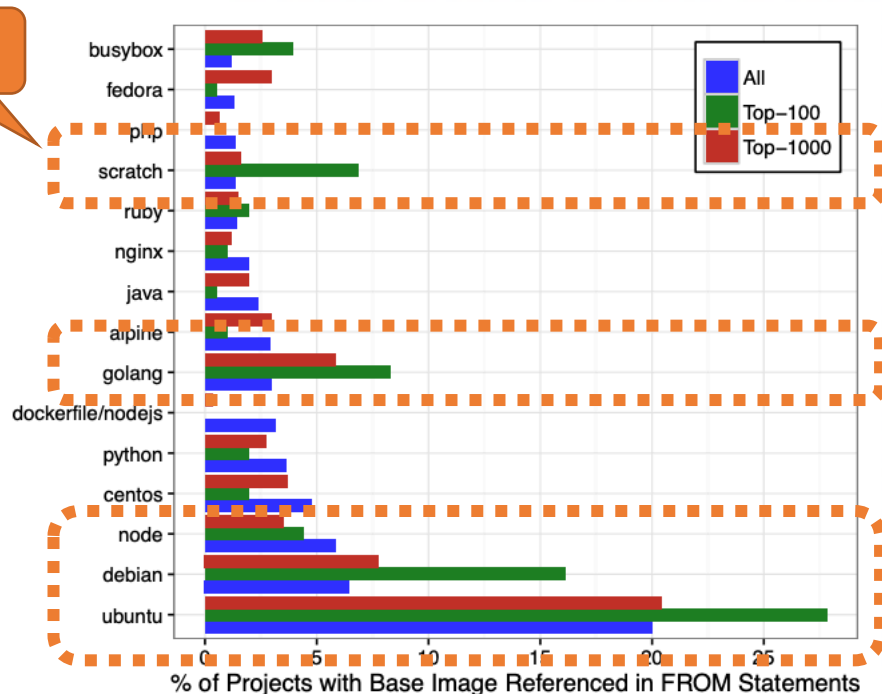


Fig. 5. Percentage of usage of 15 most commonly used base images




Fig. 4. Distribution of top 15 languages in our dataset

<https://peerj.com/preprints/2905.pdf>

Our Idea: HotC

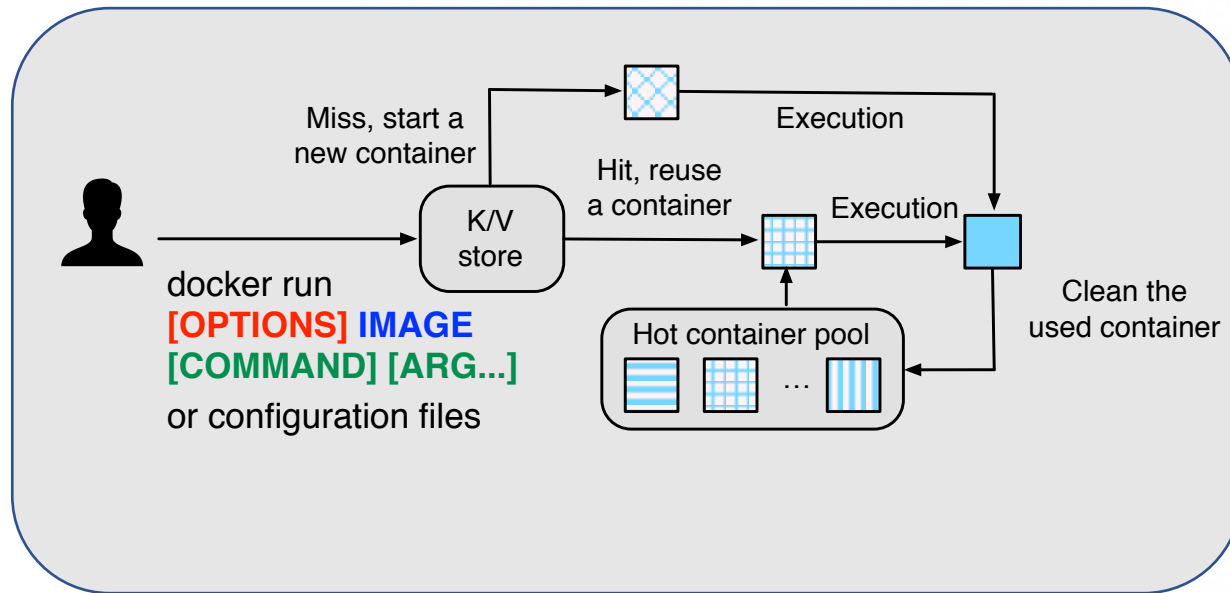



FIGURE 3: SERVERLESS ARCHITECTURE OVERVIEW



Overhead of Keeping Live Containers



The CPU usage is almost 0% when no application executes.

The memory usage increased by 0.7MB for each individual live container.

- Alpine Linux container only takes hundreds of KBs

Most of the resource consumption comes from the applications instead of the container itself


Overhead cannot be limited by precisely controlling the number of live containers



Evaluation


DELL PowerEdge T430 server:

- Dual ten-core Intel Xeon E5-2640 2.6GHz processors, 64GB memory, Gigabit Network and a 2TB 7200RPM hard drive
- Ubuntu 16.04 and Linux kernel version 4.4.20



Raspberry Pi 3:



- Quad Core 1.2GHz Broadcom BCM2837 64bit CPU, 1GB memory and 32GB storage
- Linux Raspberrypi 4.14



Docker 1.17, OpenFaaS 0.8.5



Preliminary Result



Conclusion

Serverless computing is widely used but facing challenges such as long latency due to the container cold start.

HotC, a container-based runtime management framework which leverages the lightweight containers to mitigate the cold start and improve network performance.

Our evaluation results show HotC introduces negligible overhead and can efficiently improve the performance of various applications in both cloud servers and edge devices.

Q & A

Email: ksuo@kennesaw.edu

Website: <https://kevinsuo.github.io/>

Alibaba cloud optimization for container cold start

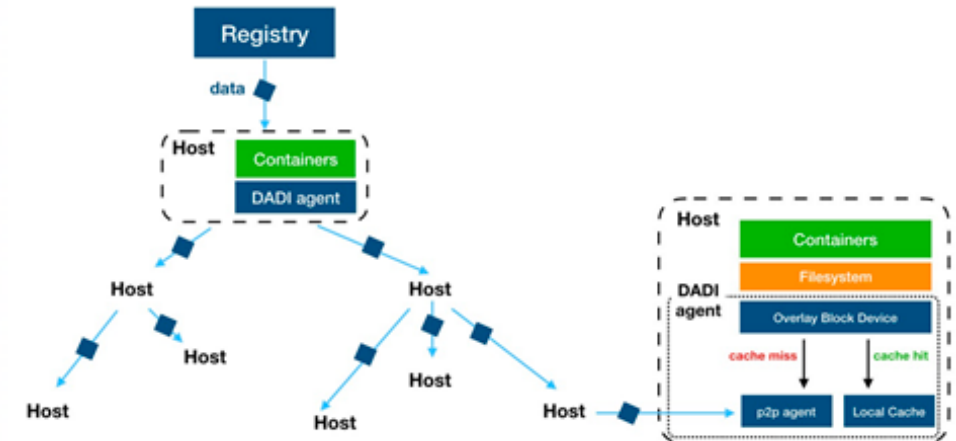
1, New image format

No need to download and unzip the full image

2, P2P network pair for data distribution

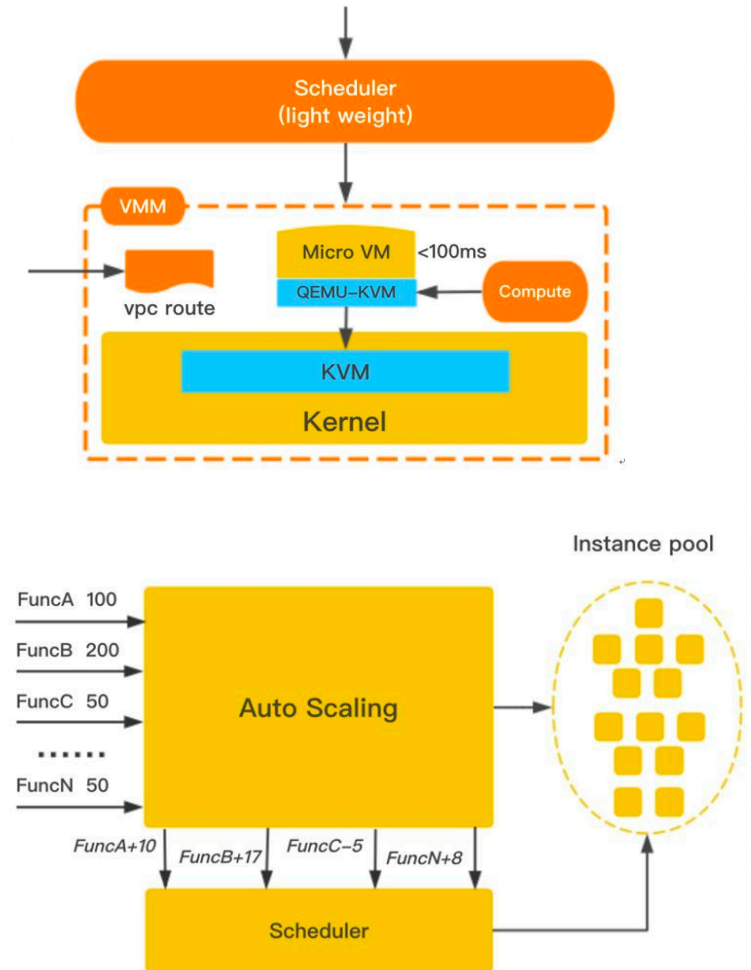
Reduce the overall load on a single point Registry

3, Efficient decompression algorithm



Tencent optimization for container cold start

- 1, Deploys their functions based on lightweight virtualization named MicroVm
- 2, New scheduling policies for active functions
real-time auto scale system which can expand or contract in second-level based on the system metrics and monitoring data

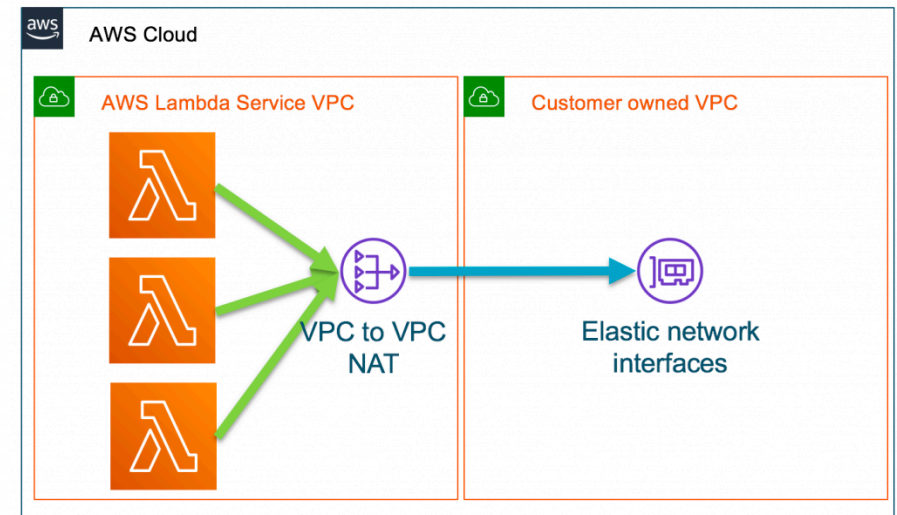


AWS optimization for container cold start

1, Keep-alive policy that retains the resources in memory for 10x minutes after a function execution

P.S. OpenWhisk also uses a 10-minute period

2, AWS Hyperplane create interfaces in advance for each serverless application instead of creating Virtual Private Cloud (VPC) for each function during execution



Azure optimization for container cold start

1, Keep-alive value for each user's workload, according to its actual invocation frequency and pattern

2, Enables the provider in many cases to pre-warm a function execution just before its invocation happens

