

# CS 3502

# Operating Systems

## System Call

Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

# Outline

---

- What is system call?
  - Kernel space vs user space
  - System call vs library call
  - What service can system call provide?
  - System call naming, input, output
- How to design a system call
  - Example
  - Project 1



# User space vs. Kernel space

- **Kernel space** is strictly reserved for running a privileged operating system kernel, kernel extensions, and most device drivers.
- **User space** is the area where application software and some drivers execute.

User mode	<b>User applications</b>	For example, <a href="#">bash</a> , <a href="#">LibreOffice</a> , <a href="#">GIMP</a> , <a href="#">Blender</a> , <a href="#">0 A.D.</a> , <a href="#">Mozilla Firefox</a> , etc.				<b>Graphics:</b> <a href="#">Mesa</a> , <a href="#">AMD Catalyst</a> , ...			
	Low-level system components:	<b>System daemons:</b> <a href="#">systemd</a> , <a href="#">runit</a> , <a href="#">logind</a> , <a href="#">networkd</a> , <a href="#">PulseAudio</a> , ...	<b>Windowing system:</b> <a href="#">X11</a> , <a href="#">Wayland</a> , <a href="#">SurfaceFlinger</a> (Android)	<b>Other libraries:</b> <a href="#">GTK+</a> , <a href="#">Qt</a> , <a href="#">EFL</a> , <a href="#">SDL</a> , <a href="#">SFML</a> , <a href="#">FLTK</a> , <a href="#">GNUstep</a> , etc.					
	<b>C standard library</b>	<a href="#">open()</a> , <a href="#">exec()</a> , <a href="#">sbrk()</a> , <a href="#">socket()</a> , <a href="#">fopen()</a> , <a href="#">calloc()</a> , ... (up to 2000 subroutines) <a href="#">glibc</a> aims to be <a href="#">POSIX/SUS</a> -compatible, <a href="#">musl</a> and <a href="#">uClibc</a> target embedded systems, <a href="#">bionic</a> written for <a href="#">Android</a> , etc.							
Kernel mode	<a href="#">Linux kernel</a>	<a href="#">stat</a> , <a href="#">splice</a> , <a href="#">dup</a> , <a href="#">read</a> , <a href="#">open</a> , <a href="#">ioctl</a> , <a href="#">write</a> , <a href="#">mmap</a> , <a href="#">close</a> , <a href="#">exit</a> , etc. (about 380 system calls) The Linux kernel <a href="#">System Call Interface</a> (SCI, aims to be <a href="#">POSIX/SUS</a> -compatible)							
		Process scheduling subsystem	IPC subsystem	Memory management subsystem	Virtual files subsystem	Network subsystem			
Other components: <a href="#">ALSA</a> , <a href="#">DRI</a> , <a href="#">evdev</a> , <a href="#">LVM</a> , <a href="#">device mapper</a> , <a href="#">Linux Network Scheduler</a> , <a href="#">Netfilter</a> <a href="#">Linux Security Modules</a> : <a href="#">SELinux</a> , <a href="#">TOMOYO</a> , <a href="#">AppArmor</a> , <a href="#">Smack</a>									
<b>Hardware</b> ( <a href="#">CPU</a> , <a href="#">main memory</a> , <a href="#">data storage devices</a> , etc.)									



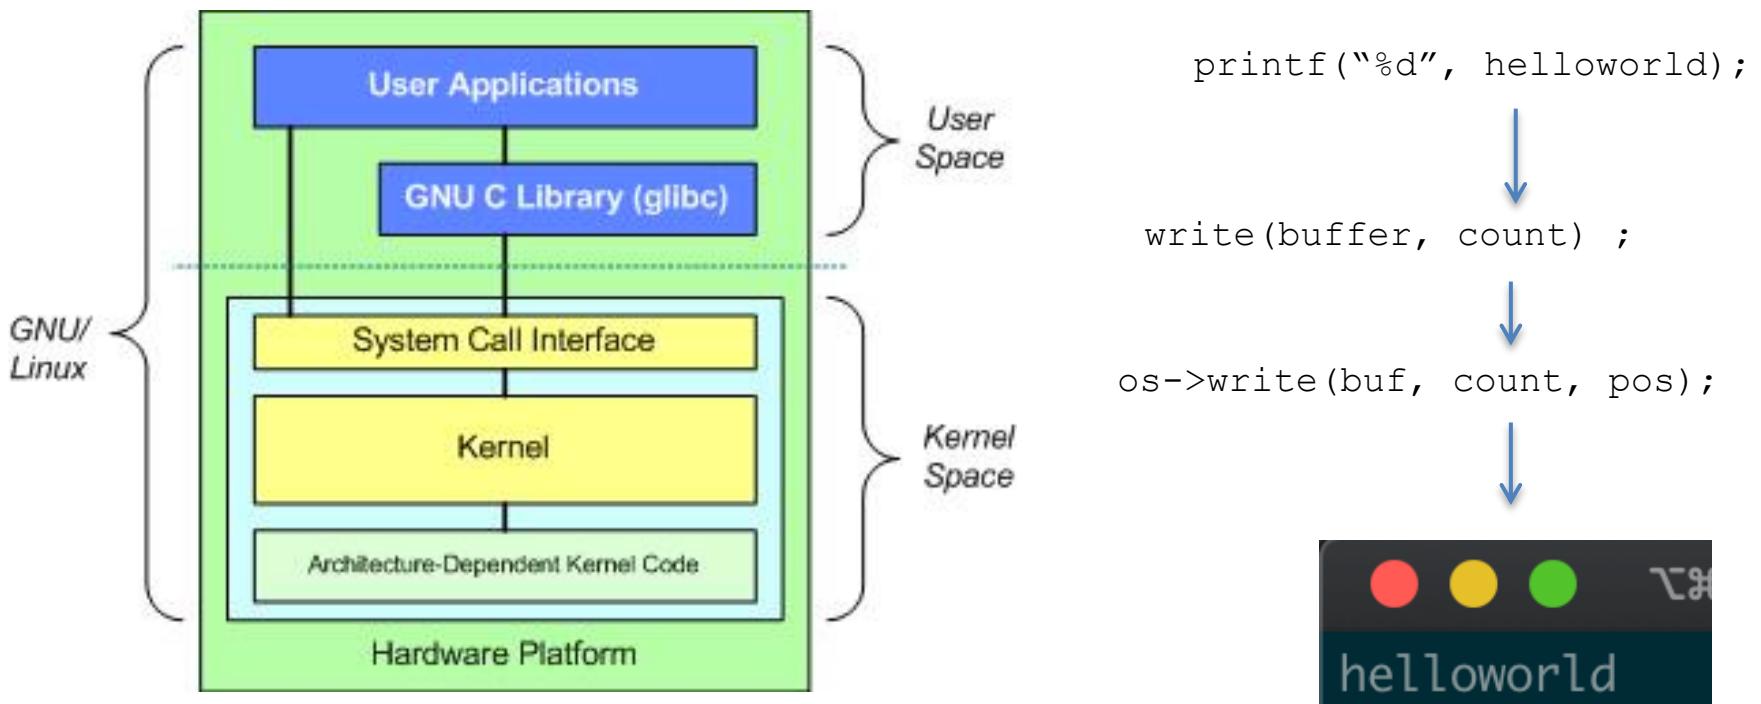
# User mode vs. Kernel mode

---

- The difference between kernel and user mode?
  - The CPU can execute **every instruction** in its instruction set and use **every feature** of the hardware when executing in kernel mode.
  - However, it can execute only **a subset of instructions** and use only **subset of features** when executing in the user mode.
- The purpose of having these two modes
  - Purpose: **protection** – to protect critical resources (e.g., privileged instructions, memory, I/O devices) from being misused by user programs.

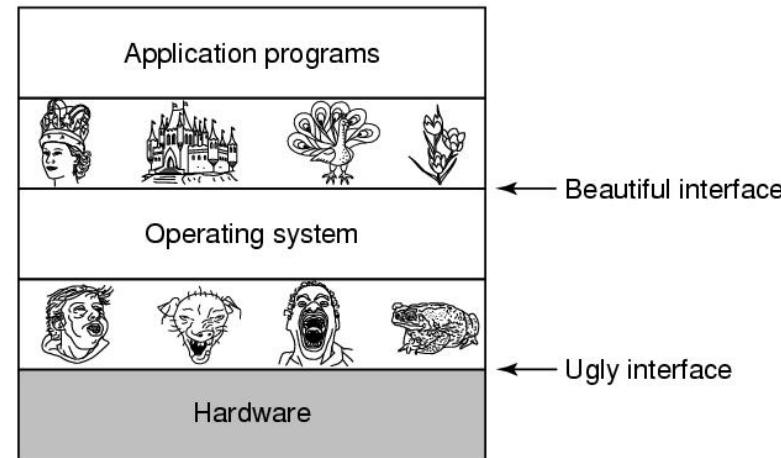
# Interact between user space and kernel space

- For applications, in order to perform privileged operations, it must transit into OS through well defined interfaces
  - System call



# System calls

- A type of special “protected procedure calls” allowing user-level processes request services from the kernel.
- System calls provide:
  - An **abstraction layer** between processes and hardware, allowing the kernel to provide access control
  - A **virtualization** of the underlying system
  - A well-defined **interface** for system services



# System calls vs. Library functions

---

- What are the similarities and differences between *system calls* and *library functions* (e.g., libc functions)?

libc functions

[https://www.gnu.org/software/libc/manual/html\\_node/Function-Index.html](https://www.gnu.org/software/libc/manual/html_node/Function-Index.html)

system calls

<https://filippo.io/linux-syscall-table/>



# System calls vs. Library functions

---

- Similarity
  - Both appear to be APIs that can be called by programs to obtain a given service
    - ▶ E.g., open,
    - ▶ <https://elixir.bootlin.com/linux/latest/source/tools/include/nolibc/nolibc.h#L2038>
  - ▶ E.g., strlen
    - ▶ [https://www.gnu.org/software/libc/manual/html\\_node/String-Length.html#index-strlen](https://www.gnu.org/software/libc/manual/html_node/String-Length.html#index-strlen)



# System calls vs. Library functions

```
1 /* strlen example */
2 #include <stdio.h>
3 #include <string.h>
4
5 int main ()
6 {
7     char szInput[256];
8     printf ("Enter a sentence: ");
9     gets (szInput);
10    printf ("The sentence entered is %u characters long.\n", (unsigned)strlen(szInput))
11    return 0;
12 }
```

Output:

```
Enter sentence: just testing
The sentence entered is 12 characters long.
```

libc functions:

<string.h> - - -> strlen() : all in user space

# System calls vs. Library functions

```
// C program to illustrate
// open system call
#include<stdio.h>
#include<fcntl.h>
#include<errno.h>
extern int errno;
int main()
{
    // if file does not have in directory
    // then file foo.txt is created.
    int fd = open("foo.txt", O_RDONLY | O_CREAT);

    printf("fd = %d/n", fd);

    if (fd == -1)
    {
        // print which type of error have in a code
        printf("Error Number % d\n", errno);

        // print program detail "Success or failure"
        perror("Program");
    }
    return 0;
}
```

System calls:

<fcntl.h> - - -> open()

- - -> do\_sys\_open() // wrapper system call

[https://elixir.bootlin.com/linux/latest/source  
/fs/open.c#L1074](https://elixir.bootlin.com/linux/latest/source/fs/open.c#L1074)

```
long do_sys_open(int dfd, const char __user *filename, int flags, umode_t mode)
{
    struct open_flags op;
    int fd = build_open_flags(flags, mode, &op);
    struct filename *tmp;

    if (fd)
        return fd;

    tmp = getname(filename);
    if (IS_ERR(tmp))
        return PTR_ERR(tmp);

    fd = get_unused_fd_flags(flags);
    if (fd >= 0) {
        struct file *f = do_filp_open(dfd, tmp, &op);
        if (IS_ERR(f)) {
            put_unused_fd(fd);
            fd = PTR_ERR(f);
        } else {
            fsnotify_open(f);
            fd_install(fd, f);
        }
    }
    putname(tmp);
    return fd;
}

SYSCALL_DEFINE3(open, const char __user *, filename, int, flags, umode_t, mode)
{
    if (force_o_largefile())
        flags |= O_LARGEFILE;

    return do_sys_open(AT_FDCWD, filename, flags, mode);
}
```

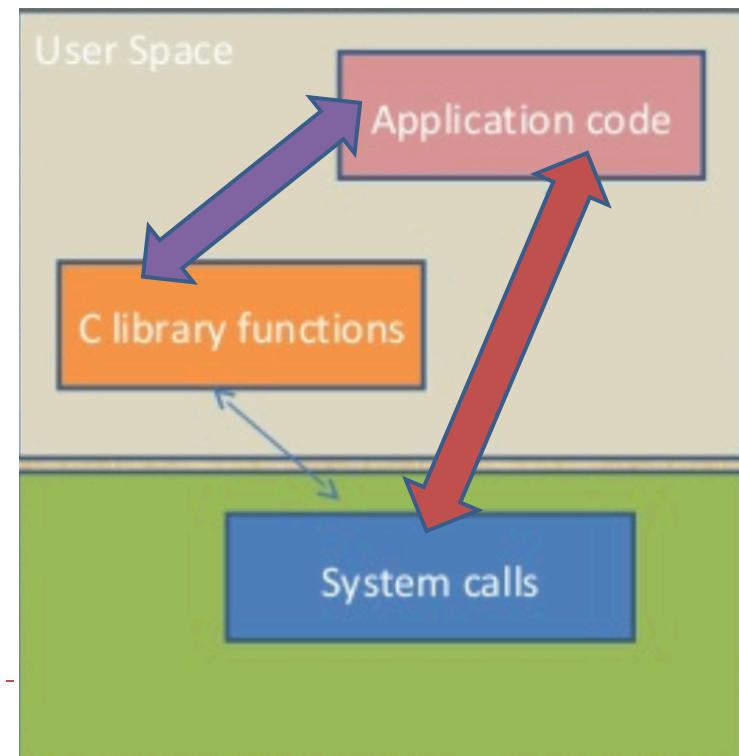


# System calls vs. Library functions

- Difference
  - Library functions execute in the user space
  - System calls execute in the kernel space

`strlen() (<string.h>)` ? → all in user space

`open() (<fcntl.h>)`? → `do_sys_open()`

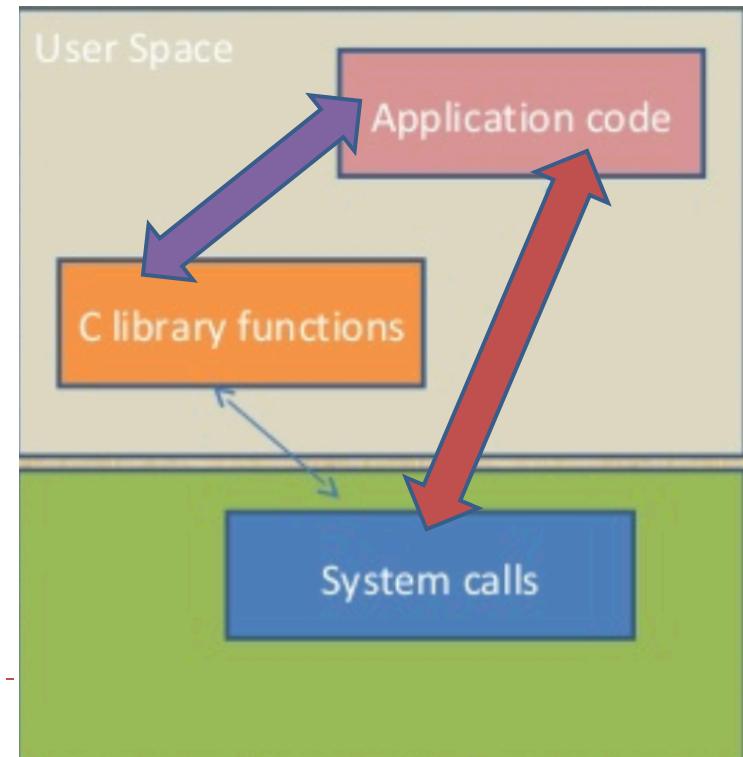


# System calls vs. Library functions

- Difference
  - Fast, no context switch
  - Slow, high cost, kernel/user context switch

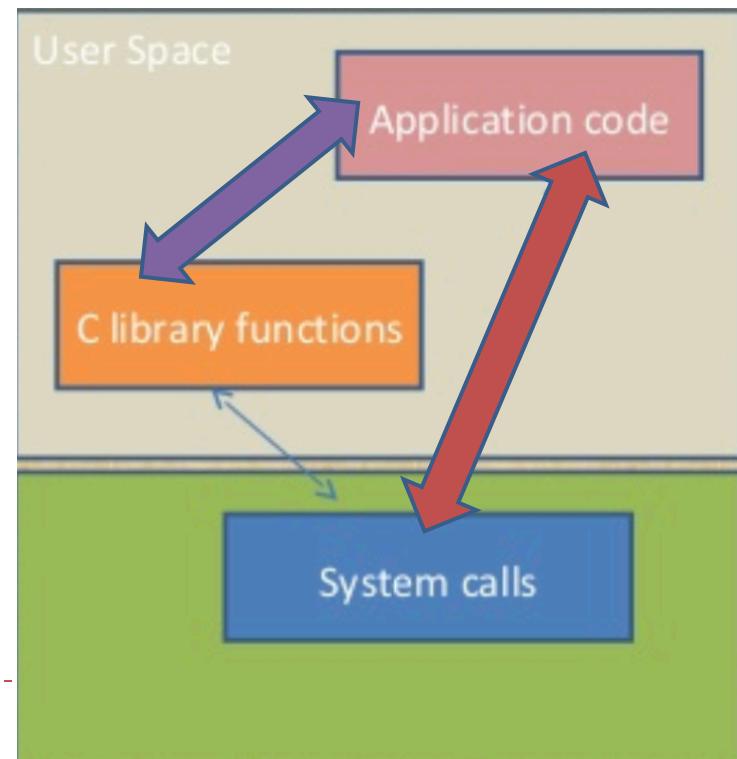
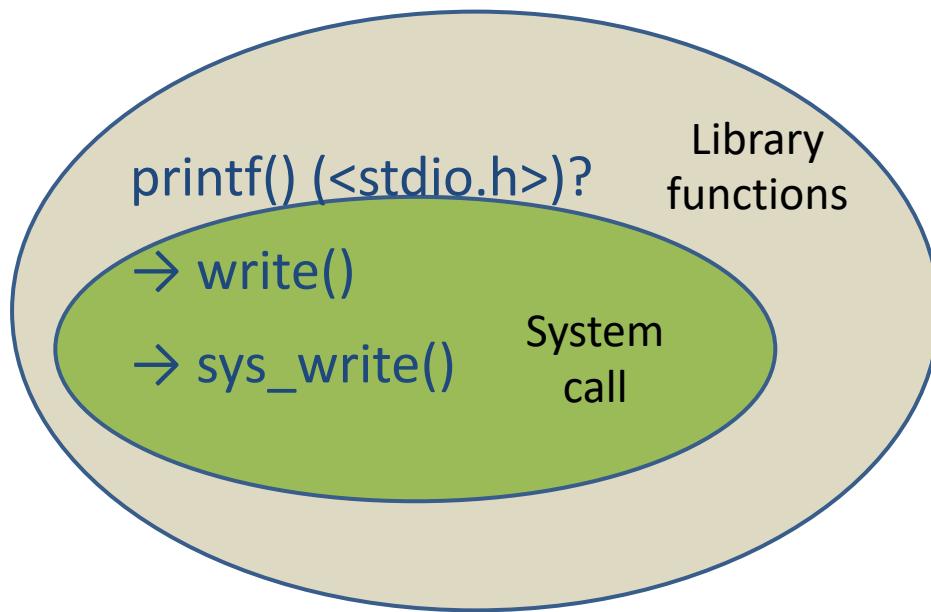
`strlen() (<string.h>)` ? → all in user space

`open() (<fcntl.h>)`? → `do_sys_open()`



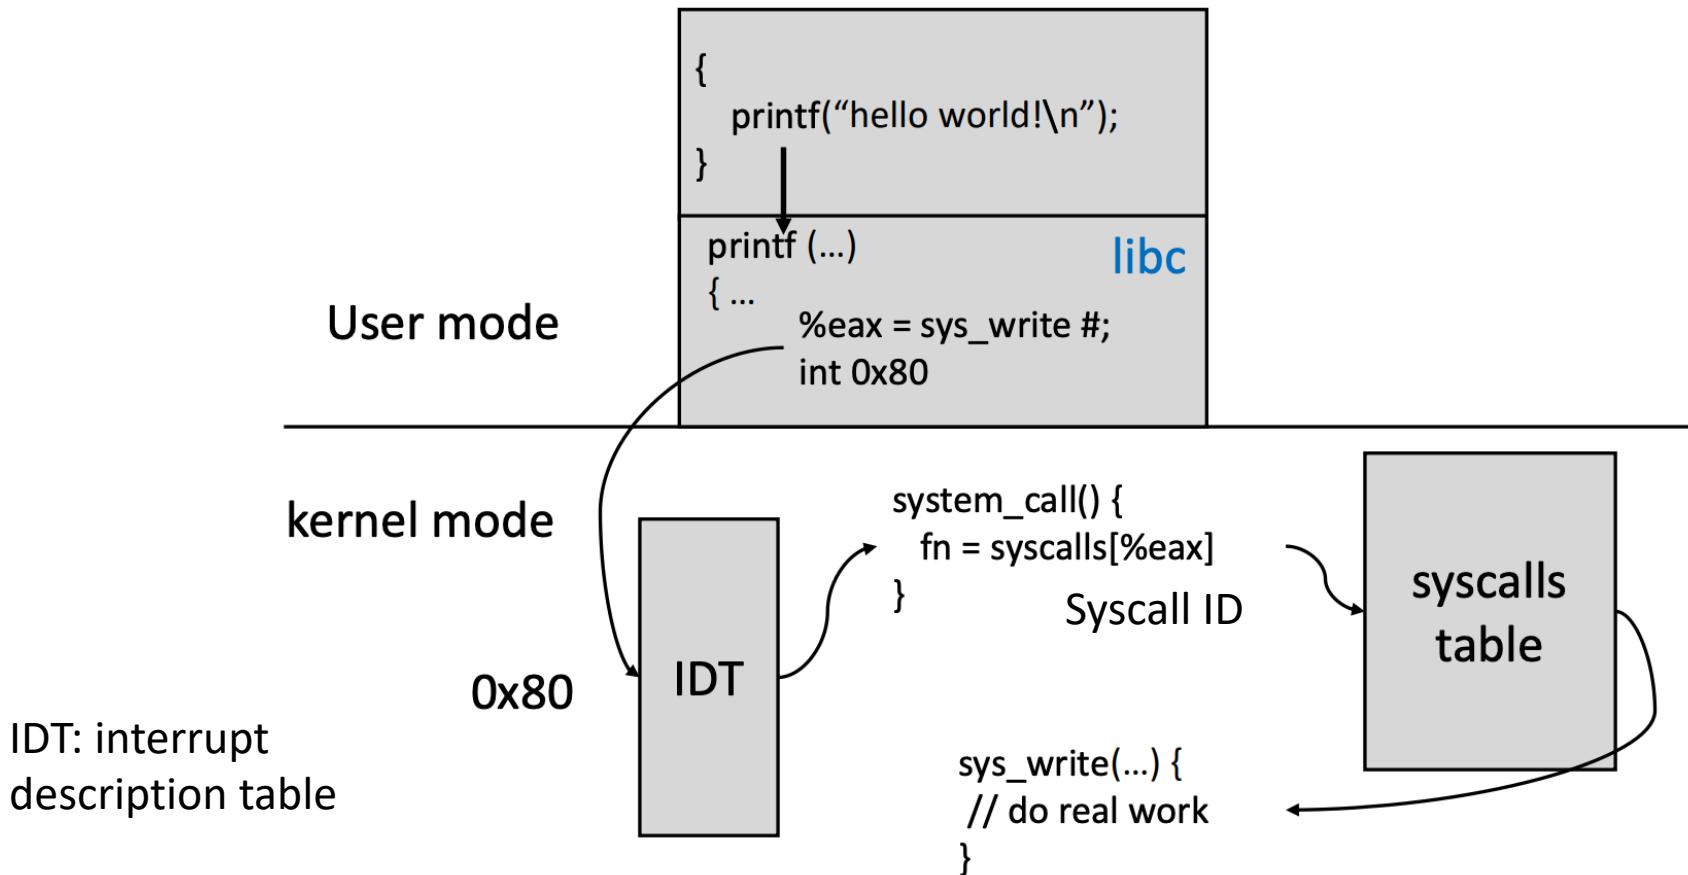
# System calls vs. Library functions

- Many system calls have a corresponding standard C library wrapper routines, which hide the details of system call entry/exit.



# System calls vs. Library functions

## Syscall Wrapper Macros



# System calls vs. Library functions

	<b>System call</b>	<b>Library function</b>
position	The functions which are a part of Kernel.	The functions which are a part of standard C library.
space	Get executed in kernel mode.	Get executed in user mode.
privilege	Runs in the supervisory mode so have every privileged.	Doesn't have much privileged.
performance	System calls are slow as there is context switch involved.	Faster execution as it doesn't involve context switch.

# Services Provided by System Calls

---

- Process creation and management
- Main memory management
- File Access, Directory and File system management
- Device handling(I/O)
- Protection, e.g., encrypt
- Networking, etc.



# Services Provided by System Calls

---

- Process related
  - end, abort
  - load, execute. E.g., exec
  - create process, terminate process. E.g., fork
  - get process attributes, set process attributes
  - wait for time. E.g., wait
  - wait event, signal event
  - allocate and free memory

<https://filippo.io/linux-syscall-table/>



# fork

---

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main()
{
    // make two process which run same
    // program after this instruction
    fork();

    printf("Hello world!\n");
    return 0;
}
```

Output:

```
Hello world!
Hello world!
```



# Services Provided by System Calls

---

- File related
  - create file, delete file
  - open, close file
  - read, write, reposition
  - get and set file attributes

<https://filippo.io/linux-syscall-table/>



# open

---

```
#include <unistd.h>
#include <fcntl.h>

int main()
{
    int filedesc = open("testfile.txt", O_WRONLY | O_APPEND);
    if(filedesc < 0)
        return 1;

    if(write(filedesc,"This will be output to testfile.txt\n", 36) != 36)
    {
        write(2,"There was an error writing to testfile.txt\n");      // str
        return 1;
    }

    return 0;
}
```



# Services Provided by System Calls

---

- Device related
  - register device, release device
  - read, write
  - get device attributes, set device attributes
  - logically attach or detach devices

<https://filippo.io/linux-syscall-table/>



# ioctl

```
int main(void) {  
    int fd;  
    int i;  
    int iomask;  
  
    if ((fd = open("/dev/gpiog", O_RDWR))<0) {  
        printf("Open error on /dev/gpiog\n");  
        exit(0);  
    }  
  
    iomask=1<<25;  
  
    for (i=0;i<10;i++) {  
        printf("Led ON\n");  
        ioctl(fd,_IO(ETRAXGPIO_IODEVICE,IO_SETBITS),iomask);  
        sleep(1);  
  
        printf("Led OFF\n");  
        ioctl(fd,_IO(ETRAXGPIO_IODEVICE,IO_CLRBITS),iomask);  
        sleep(1);  
    }  
    close(fd);  
    exit(0);  
}
```

Configure the device fd



# Services Provided by System Calls

---

- Information related
  - Get time or date, set time or date
  - Number of current users
  - Amount of free memory or disk space

<https://filippo.io/linux-syscall-table/>



# getpid

```
#include<stdio.h>
#include<dos.h>
int main()
{
    struct date dt;

    getdate(&dt);

    printf("Operating system's current date is %d-%d-%d\n"
           ,dt.da_day,dt.da_mon,dt.da_year);

    return 0;
}
```

## OUTPUT:

```
Operating system's current date is 12-01-2012
```

```
#include <stdio.h>
#include <time.h>
int main ()
{
    time_t seconds;
    seconds = time (NULL);

    printf ("Number of hours since 1970 Jan 1st "
           "is %ld \n", seconds/3600);
    return 0;
}
```

## OUTPUT:

```
Number of hours since 1970 Jan 1st is 374528
```



# Services Provided by System Calls

---

- Communication related
  - create, delete communication connection
  - send, receive messages
  - transfer status information
  - attach and detach remote devices

<https://filippo.io/linux-syscall-table/>



# pipe

---

```
int main()
{
    // We use two pipes
    // First pipe to send input string from parent to child
    // Second pipe to send concatenated string from child to parent

    int fd1[2];    // Used to store two ends of first pipe
    int fd2[2];    // Used to store two ends of second pipe

    char fixed_str[] = "forgeeks.org";
    char input_str[100];
    pid_t p;

    if (pipe(fd1)==-1)
    {
        fprintf(stderr, "Pipe Failed" );
        return 1;
    }
    if (pipe(fd2)==-1)
    {
        fprintf(stderr, "Pipe Failed" );
        return 1;
    }

    scanf("%s", input_str);
    p = fork();

    if (p < 0)
    {
        fprintf(stderr, "fork Failed" );
        return 1;
    }
}

// child process
else
{
    close(fd1[1]);    // Close writing end of first pipe

    // Read a string using first pipe
    char concat_str[100];
    read(fd1[0], concat_str, 100);

    // Concatenate a fixed string with it
    int k = strlen(concat_str);
    int i;
    for (i=0; i<strlen(fixed_str); i++)
        concat_str[k++] = fixed_str[i];

    concat_str[k] = '\0';    // string ends with null character

    // Close both reading ends
    close(fd1[0]);
    close(fd2[0]);

    // Write concatenated string and close writing end of second pipe
    write(fd2[1], concat_str, strlen(concat_str));
    close(fd2[1]);

    exit(0);
}
```

# chown

---

```
> root:bash — Konsole
root@kali:~# ls -l file1.txt
-rw-r--r-- 1 root root 12 Feb 4 12:04 file1.txt
root@kali:~# chown master file1.txt
root@kali:~# ls -l file1.txt
-rw-r--r-- 1 master root 12 Feb 4 12:04 file1.txt
root@kali:~# █
```



# Examples of System Calls

Process Control	File Manipulation	Device Manipulation	Information Maintenance	Communication	Protection
fork() exit() wait()	open() read() write() close()	ioctl() read() write()	getpid() alarm() time()	pipe() send() recv() mmap()	chmod() umask() chown()

<http://man7.org/linux/man-pages/man2/syscalls.2.html>

<https://filippo.io/linux-syscall-table/>



# Syscall interface

---

- Important to keep interface **small, stable** (for binary and backward compatibility). Every syscall does **one thing**.
- Early UNIXs had about **60** system calls, Linux 2.6 has about **300**; Solaris more, Window more still
- Aside: Windows does **not publicly** document syscalls and only documents library wrapper routines (unlike UNIX/Linux)
- Syscall numbers **cannot be reused** (!)

Modern Linux system calls:

[https://elixir.bootlin.com/linux/v5.0/source/arch/x86/entry/syscalls/syscall\\_64.tbl](https://elixir.bootlin.com/linux/v5.0/source/arch/x86/entry/syscalls/syscall_64.tbl)



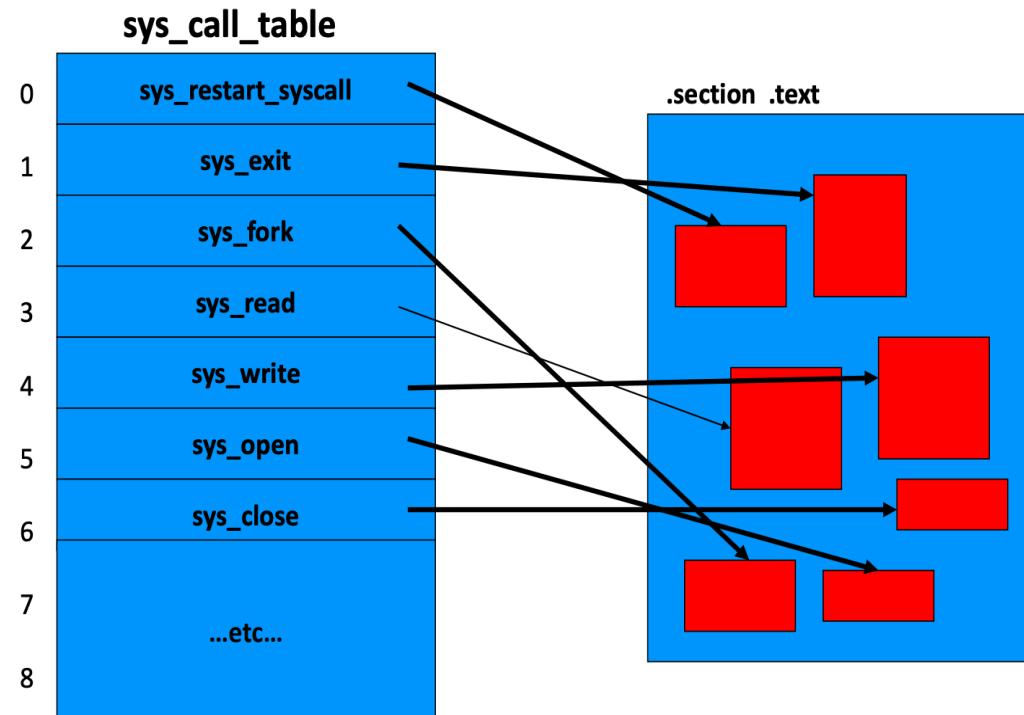
# Syscall interface APIs

---

- Three most common APIs in OSes are
  - Win32 API for Windows,
  - POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X),
  - Java API for the Java virtual machine (JVM)

# System call table

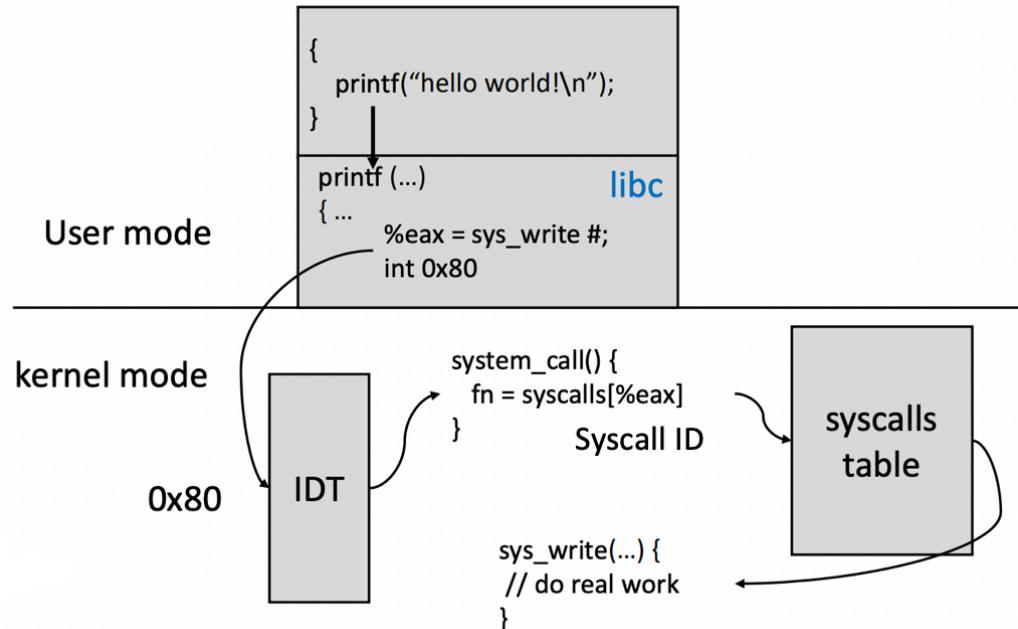
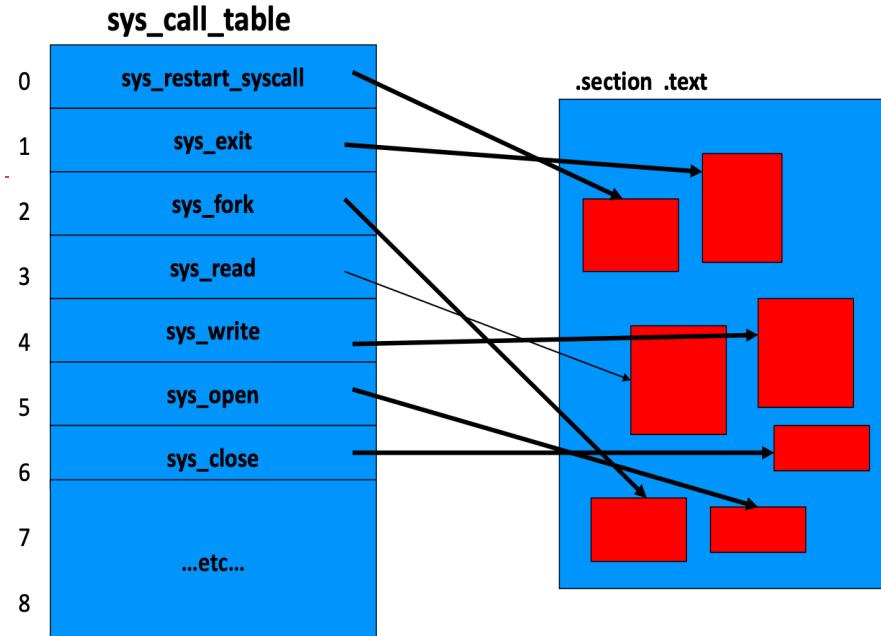
- There are approximately 300 system-calls in Linux 2.6.
- An **array** of function-pointers (identified by the ID number)
- This array is named ‘`sys_call_table[]`’ in Linux  
[https://elixir.bootlin.com/linux/v5.0/source/arch/x86/entry/syscall\\_64.c](https://elixir.bootlin.com/linux/v5.0/source/arch/x86/entry/syscall_64.c)



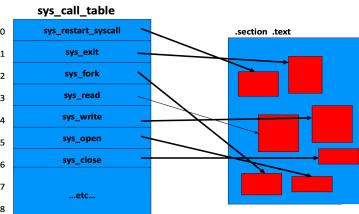
The ‘jump-table’ idea

# System call table

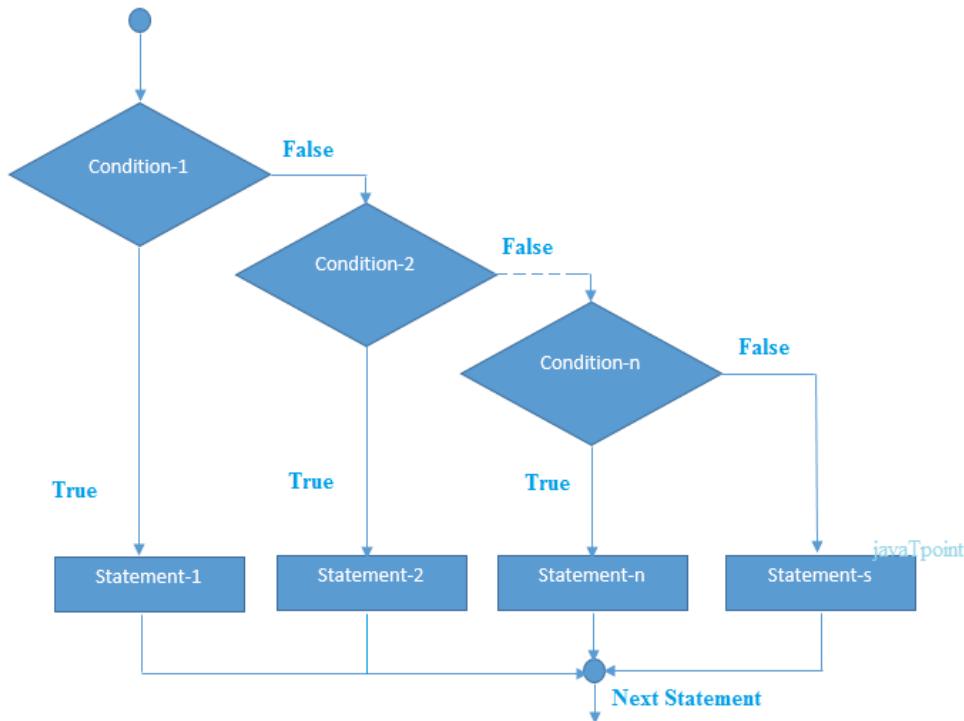
- Any specific system-call is selected by its ID-number (i.e., the system call number, which is placed into register %eax)



# Discussion



- Instead of using the approach of system table, can we use if-else tests or switch statement to transfer to the service routine's entry point?



- Functionality wise, yes.
- But it would be extremely inefficient.  $O(n)$
- System call invocations are synchronous, long system call execution is not desired.



# Syscall Naming Convention

---

- Usually a library function “foo()” will do some work and then call a system call (“sys\_foo()”)
- In Linux, all system calls begin with “sys\_”
- Often “sys\_abc()” just does some simple error checking and then calls a worker function named “do\_abc()”

[https://elixir.bootlin.com/linux/v4.14/source/arch/x86/entry/syscalls/syscall\\_64.tbl](https://elixir.bootlin.com/linux/v4.14/source/arch/x86/entry/syscalls/syscall_64.tbl)

open:

<https://elixir.bootlin.com/linux/v4.14/source/fs/open.c#L1072>

do\_sys\_open:

<https://elixir.bootlin.com/linux/v4.14/source/fs/open.c#L1044>



# Syscall return values

- Recall that library calls return **-1** on error, and place a specific error code in the global variable *errno*
- System calls return *specific negative values* to indicate an error
  - on x86, the return value is put into %eax, so that the library wrapper function can access.

```
#include <unistd.h>
#include <sys/syscall.h>
#include <errno.h>
...
int rc;

rc = syscall(SYS_chmod, "/etc/passwd", 0444);
if (rc == -1)
    fprintf(stderr, "chmod failed, errno = %d\n", e
```

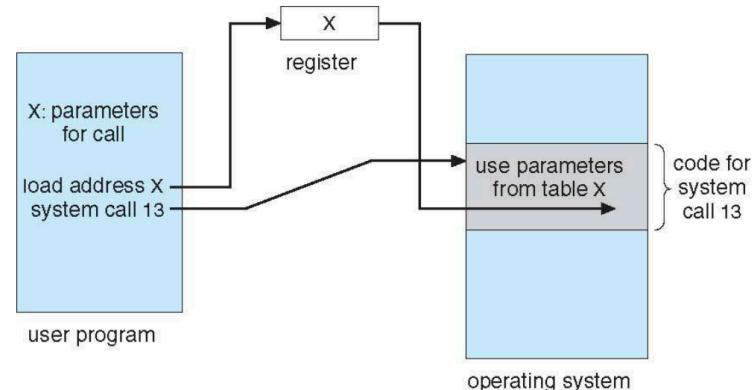
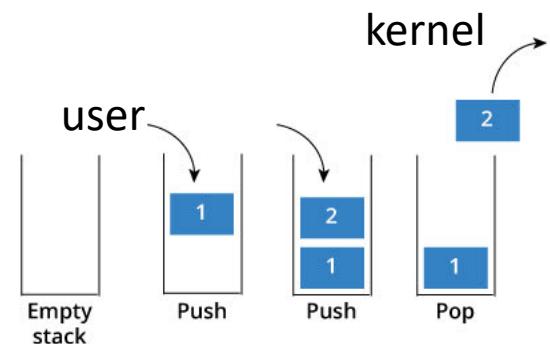
library calls

```
        if (flags & __O_SYNC)
            flags |= O_DSYNC;
        ...
        if (flags & __O_TMPFILE) {
            if ((flags & O_TMPFILE_MASK) != O_TMPFILE)
                return -EINVAL;
            if (!(acc_mode & MAY_WRITE))
                return -EINVAL;
        } else if (flags & O_PATH) {
            /*
             * If we have O_PATH in the open flag. Then we
             * cannot have anything other than the below set of
             */
            flags &= O_DIRECTORY | O_NOFOLLOW | O_PATH;
            acc_mode = 0;
```

System calls

# System call argument passing

- Three general methods used to pass arguments to the OS:
  - Method 1: pass the arguments in **registers** (simplest)
  - Method 2: arguments are placed, or pushed, onto the **stack** by the program and popped off the stack by the OS kernel code
  - Method 3: arguments are stored in a **block**, or table, **in memory**, and address of block passed as a parameter in a register, %eax
    - ▶ This approach taken by Linux and Solaris



# System call argument passing discussion

---

- Consider a system call, zeroFill, which fills a user buffer with zeroes:  
`zeroFill(char* buffer, int bufferSize);`
- The following kernel implementation of zeroFill contains a security vulnerability. What is the vulnerability, and how would you fix it?

```
void sys_zeroFill(char* buffer, int bufferSize) {  
    for (int i=0; i < bufferSize; i++) {  
        buffer[i] = 0;  
    }  
}
```



# System call argument passing discussion

```
void sys_zeroFill(char* buffer, int bufferSize) {  
    for (int i=0; i < bufferSize; i++) {  
        buffer[i] = 0;  
    }  
}
```

- The user buffer pointer is **untrusted**, and could point anywhere. In particular, it could point inside the kernel address space. This could lead to a system crash or security breakdown.
- Fix: verify the pointer is a **valid** user address

# System call argument passing discussion

---

```
void sys_zeroFill(char* buffer, int bufferSize) {  
    for (int i=0; i < bufferSize; i++) {  
        buffer[i] = 0;  
    }  
}
```

- Is it a security risk to execute the zeroFill function in user-mode?



# System call argument passing discussion

---

```
void sys_zeroFill(char* buffer, int bufferSize) {  
    for (int i=0; i < bufferSize; i++) {  
        buffer[i] = 0;  
    }  
}
```

- Is it a security risk to execute the zeroFill function in user-mode?
  - No. User-mode code **does not have permission** to access the kernel's address space. If it tries, the hardware raises an exception, which is safely handled by the OS.

# Outline

---

- What is system call?
  - Kernel space vs user space
  - System call vs library call
  - What service can system call provide?
  - System call naming, input, output
- How to design a system call
  - Example
  - Project 1



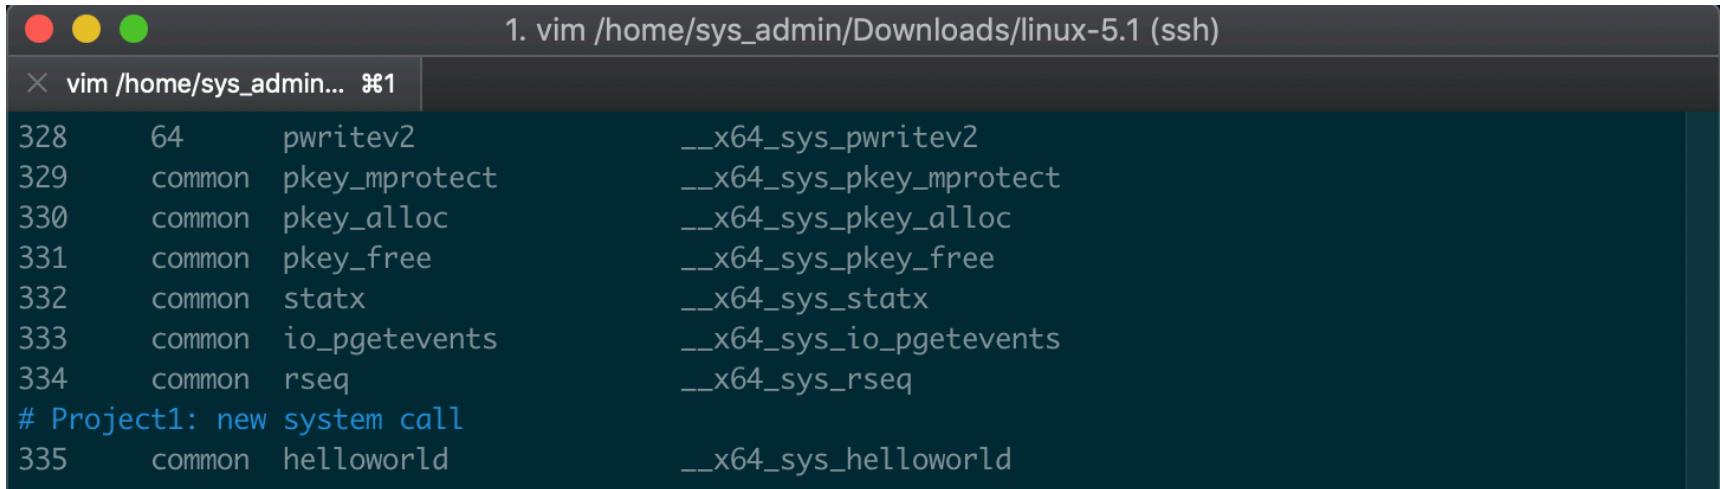
# Define your system call

---

- Step 1: register your system call
- Step 2: declare your system call in the header file
- Step 3: implement your system call
- Step 4: write user level app to call it

# Step 1: register your system call

arch/x86/entry/syscalls/syscall\_64.tbl



```
1. vim /home/sys_admin/Downloads/linux-5.1 (ssh)
vim /home/sys_admin... ⌘1

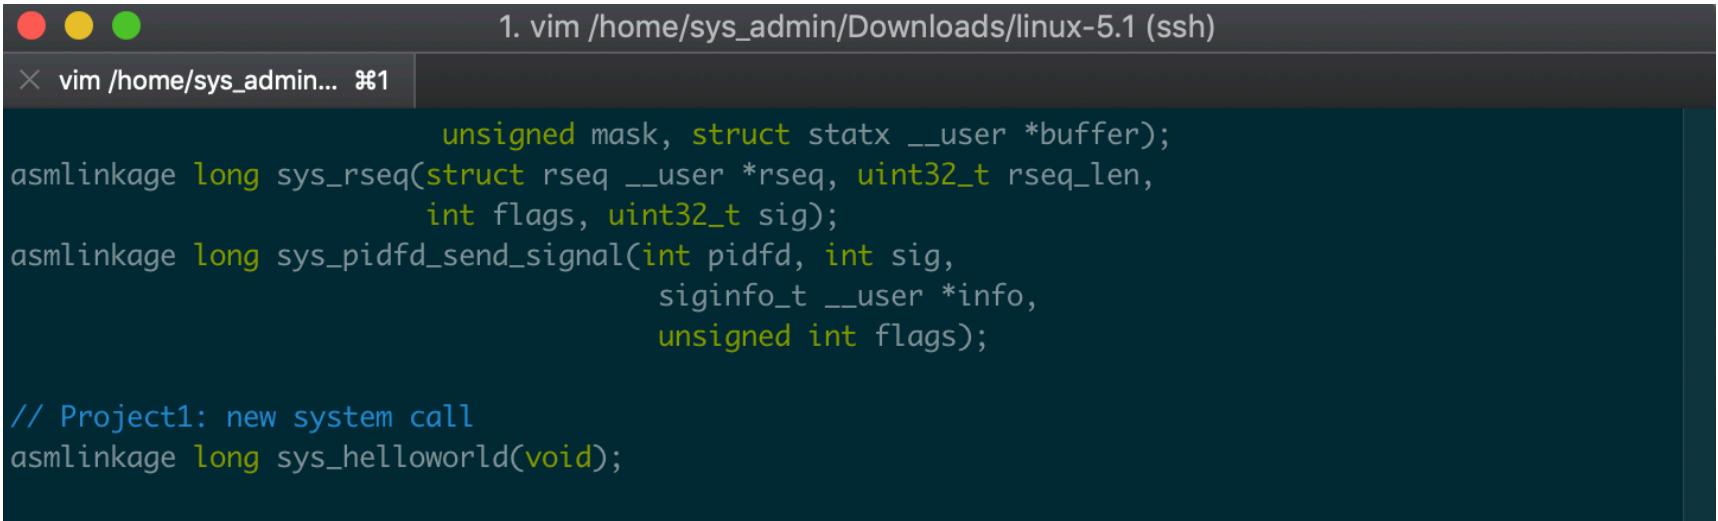
328      64      pwritev2          __x64_sys_pwritev2
329      common   pkey_mprotect    __x64_sys_pkey_mprotect
330      common   pkey_alloc       __x64_sys_pkey_alloc
331      common   pkey_free        __x64_sys_pkey_free
332      common   statx           __x64_sys_statx
333      common   io_pgetevents   __x64_sys_io_pgetevents
334      common   rseq            __x64_sys_rseq
# Project1: new system call
335      common   helloworld      __x64_sys_helloworld
```

[https://elixir.bootlin.com/linux/v5.0/source/arch/x86/entry/syscalls/syscall\\_64.tbl#L346](https://elixir.bootlin.com/linux/v5.0/source/arch/x86/entry/syscalls/syscall_64.tbl#L346)



## Step 2: declare your system call in the header file

include/linux/syscalls.h



The screenshot shows a terminal window titled "1. vim /home/sys\_admin/Downloads/linux-5.1 (ssh)". The vim editor is displaying the contents of the syscalls.h header file. The code includes several standard Linux system calls like sys\_rseq and sys\_pidfd\_send\_signal, followed by a new user-defined system call sys\_helloworld.

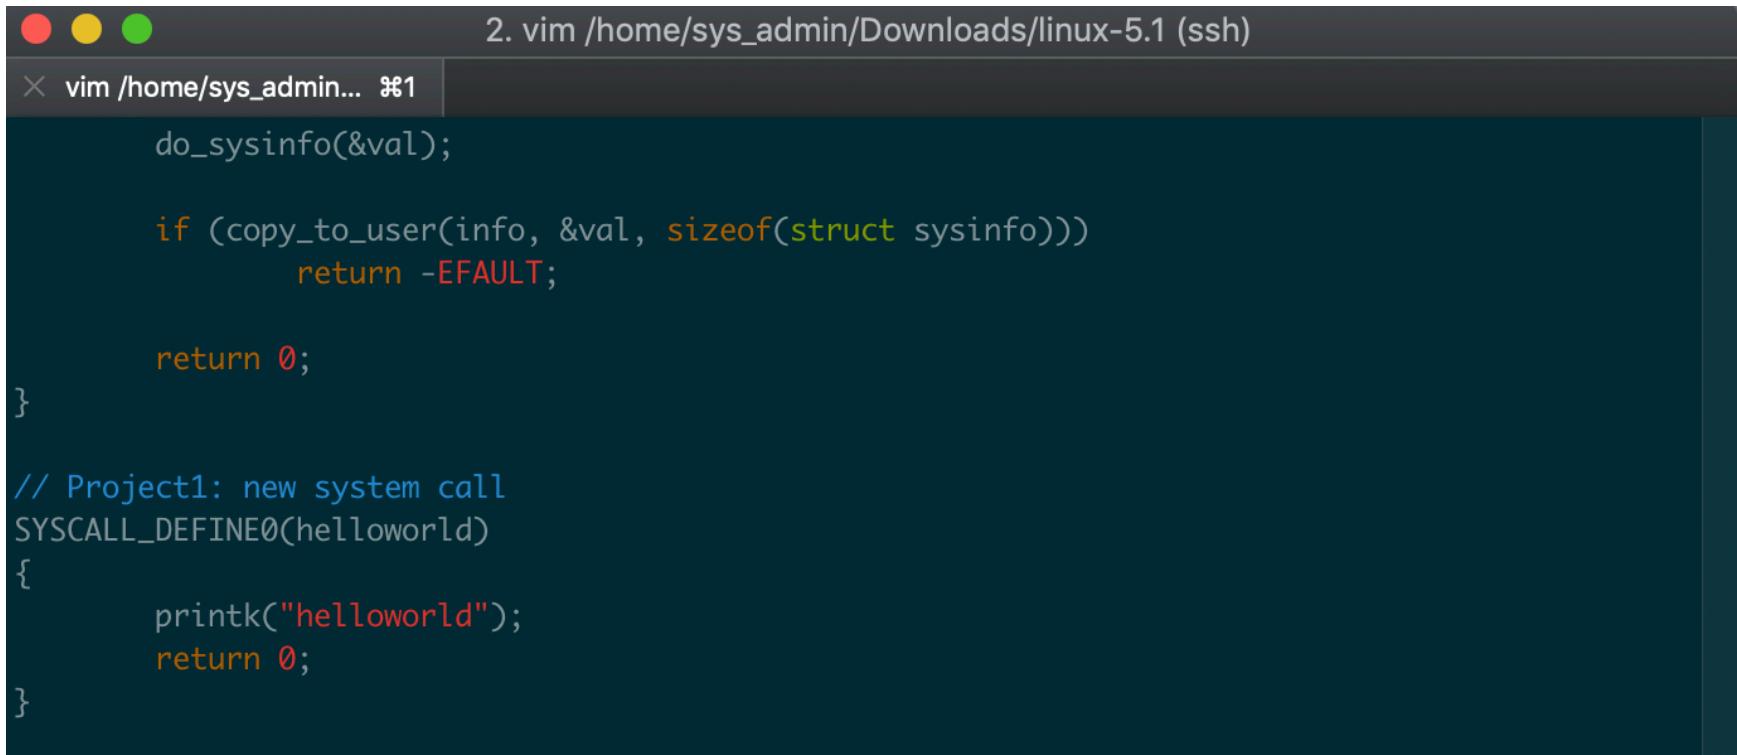
```
unsigned mask, struct statx __user *buffer);  
asmlinkage long sys_rseq(struct rseq __user *rseq, uint32_t rseq_len,  
                         int flags, uint32_t sig);  
asmlinkage long sys_pidfd_send_signal(int pidfd, int sig,  
                                       siginfo_t __user *info,  
                                       unsigned int flags);  
  
// Project1: new system call  
asmlinkage long sys_helloworld(void);
```

<https://elixir.bootlin.com/linux/v5.0/source/include/linux/syscalls.h>



# Step 3: implement your system call

kernel/sys.c



```
2. vim /home/sys_admin/Downloads/linux-5.1 (ssh)
vim /home/sys_admin... #1

do_sysinfo(&val);

if (copy_to_user(info, &val, sizeof(struct sysinfo)))
    return -EFAULT;

return 0;
}

// Project1: new system call
SYSCALL_DEFINE0(helloworld)
{
    printk("helloworld");
    return 0;
}
```

<https://elixir.bootlin.com/linux/v5.0/source/kernel/sys.c#L402>



# Step 4: write user level app to call it

test\_syscall.c:

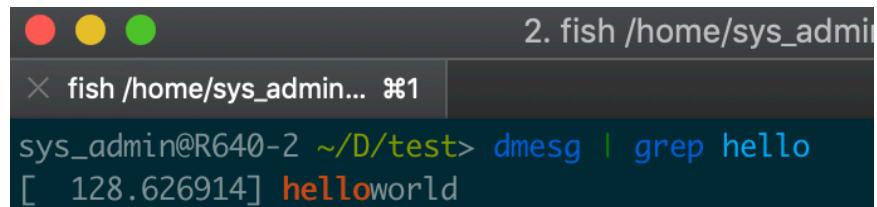
```
*****  
#include <linux/unistd .h>  
#include <sys/syscall .h>  
#include <sys/types .h>  
#include <stdio .h>  
#define __NR_helloworld 335  
  
int main(int argc, char *argv[])  
{  
    syscall (__NR_helloworld) ;  
    return 0 ;  
}  
*****
```

//If syscall needs parameter, then:  
//syscall (\_\_NR\_helloworld, a, b, c) ;

Compile and execute:

```
$ gcc test_syscall.c -o test_syscall  
$ ./test_syscall
```

The test program will call the new system call and output a helloworld message at the tail of the output of dmesg (system log).



A screenshot of a terminal window titled "fish /home/sys\_admin... #1". The command entered is "dmesg | grep hello". The output shows the message "[ 128.626914] helloworld" at the end of the log.

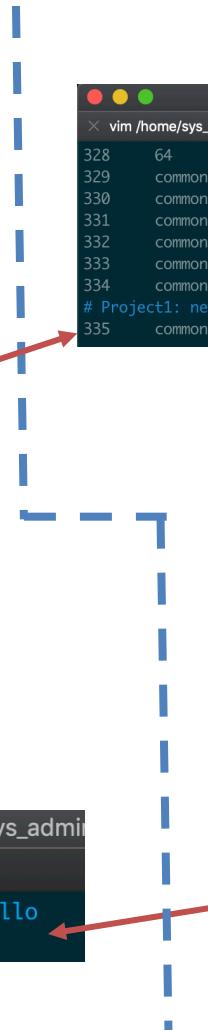
```
2. fish /home/sys_admin... #1  
x fish /home/sys_admin... #1  
sys_admin@R640-2 ~/D/test> dmesg | grep hello  
[ 128.626914] helloworld
```

# Put it all together

## User space

```
*****  
#include <linux/unistd.h>  
#include <sys/syscall.h>  
#include <sys/types.h>  
#include <stdio.h>  
#define __NR_helloworld 335  
  
int main(int argc, char *argv[]){  
    syscall (__NR_helloworld);  
    return 0;  
}
```





## Kernel space

```
1. vim /home/sys_admin/Downloads/linux-5.1 (ssh)
vim /home/sys_admin... #1
    unsigned mask, struct statx __user *buffer);
asm linkage long sys_rseq(struct rseq __user *rseq, uint32_t rseq_len,
                           int flags, uint32_t sig);
asm linkage long sys_pidfd_send_signal(int pidfd, int sig,
                                       siginfo_t __user *info,
                                       unsigned int flags);

// Project1: new system call
asm linkage long sys_helloworld(void);
```

```
2. vim /home/sys_admin/Downloads/linux

vim /home/sys_admin... 

do_sysinfo(&val);

if (copy_to_user(info, &val, sizeof(struct sysinfo)))
    return -EFAULT;

return 0;
}

// Project1: new system call
SYSCALL_DEFINE0(helloworld)
{
    printk("helloworld");
    return 0;
}
```

```
2. fish /home/sys_admin  
x fish /home/sys_admin... #1  
sys_admin@R640-2 ~ /D/test> dmesg | grep hello  
[ 128.626914] helloworld
```

# How to build one ubuntu VM?

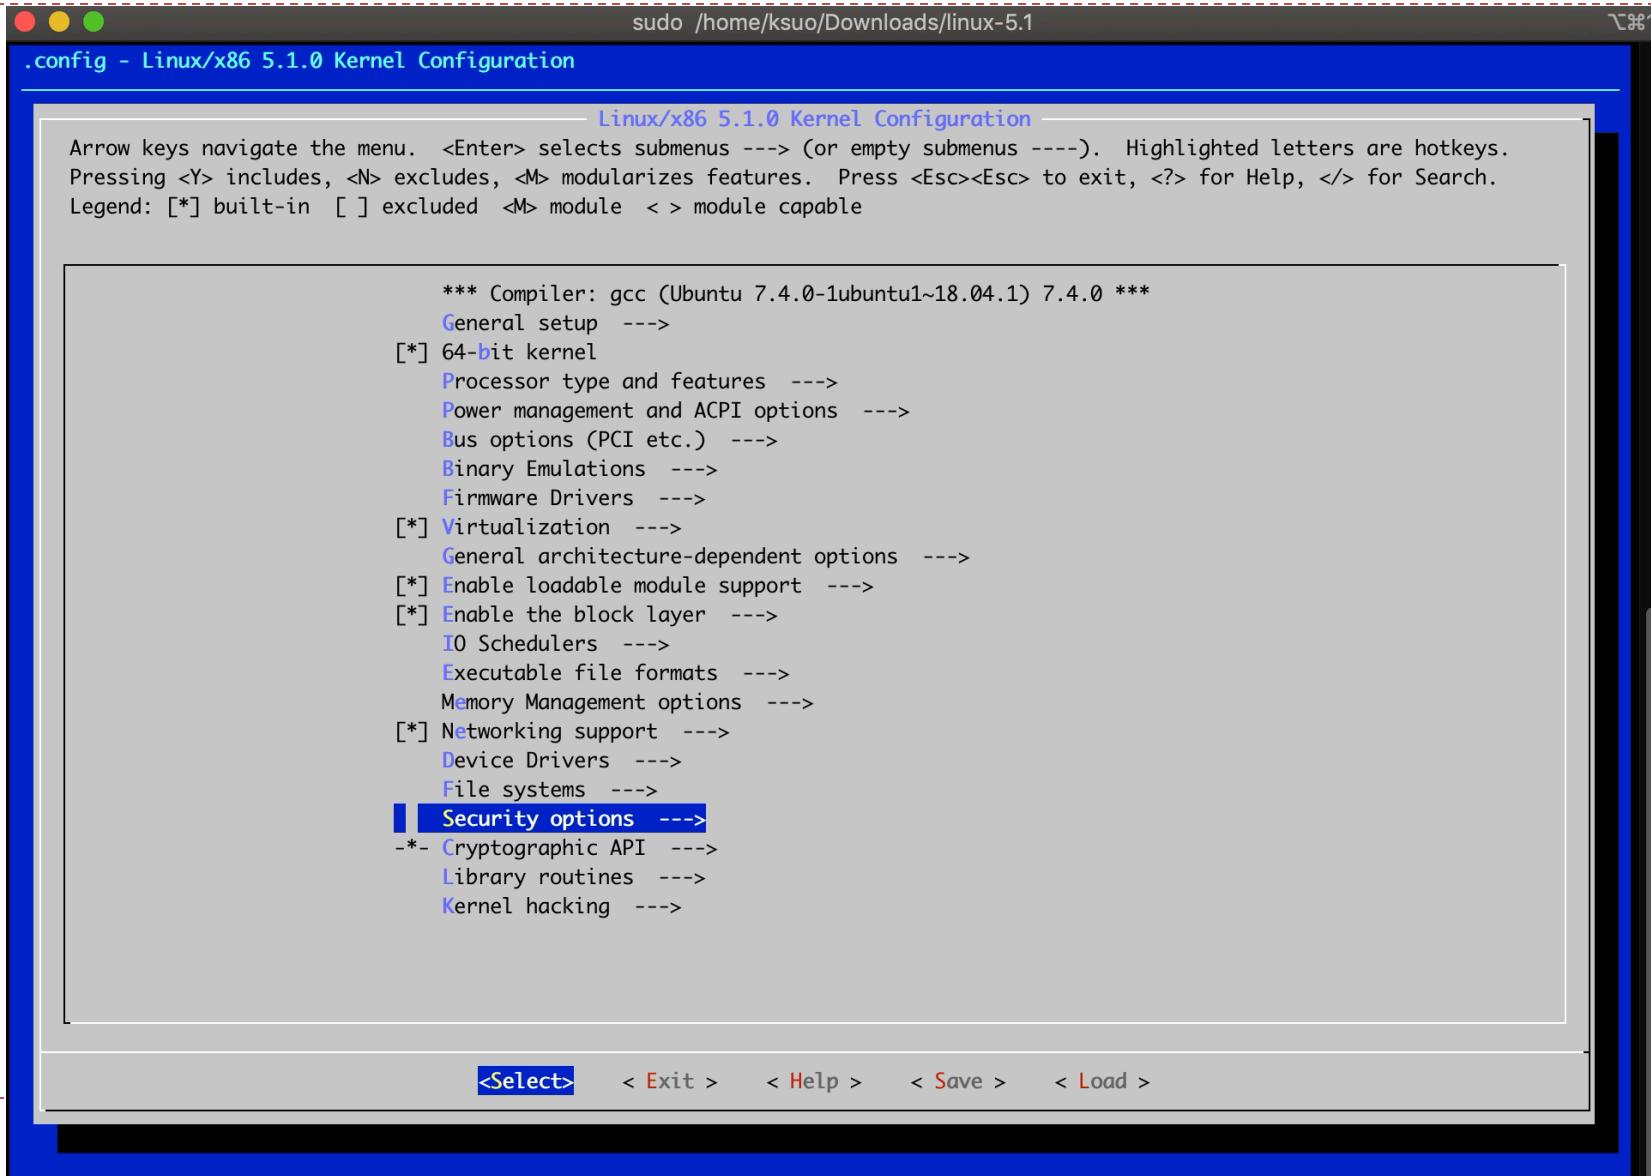
---

- HostOS
  - Windows 10:  
<https://www.youtube.com/watch?v=QbmRXJJKsvs>
  - MacOS:  
<https://www.youtube.com/watch?v=GDoCrfPma2k&t=321s>



# Project 1: menuconfig

<https://youtu.be/UyOGF4UOoR0>



# Tips

---

- Compile Linux kernel takes half to many hours, depending on machine speed
- To save the time, you can compile it before you sleep
- Run many commands in one:

\$ sudo make; sudo make modules; sudo make modules\_install; sudo make install

```
ksuo@ksuo-VirtualBox ~/D/linux-5.1>
sudo make; sudo make modules; sudo make modules_install; sudo make install
```



# Where is my kernel?

- \$ ls /boot/

Initial ramdisk: loading a temporary root file system into memory. Used for startup.

```
ksuo@ksuo-VirtualBox: ~
```

```
config-5.0.0-23-generic  
config-5.0.0-25-generic  
config-5.1.0  
grub/  
initrd.img-5.0.0-23-generic  
initrd.img-5.0.0-25-generic  
initrd.img-5.1.0  
memtest86+.bin  
ksuo@ksuo-VirtualBox ~
```

```
~/.D/linux-5.1> ls /boot/  
memtest86+.elf  
memtest86+_multiboot.bin  
System.map-5.0.0-23-generic  
System.map-5.0.0-25-generic  
System.map-5.1.0  
vmlinuz-5.0.0-23-generic  
vmlinuz-5.0.0-25-generic  
vmlinuz-5.1.0
```

vm [Running]  
Thu 11:38  
fish /home/ksuo/.D/linux-5.1>

Linux executable kernel image



# Which kernel to boot if there are many?

If you are using Ubuntu: change the grub configuration file:

```
$ sudo vim /etc/default/grub
```

The OS boots by using the first kernel by default. You have 10 seconds to choose.

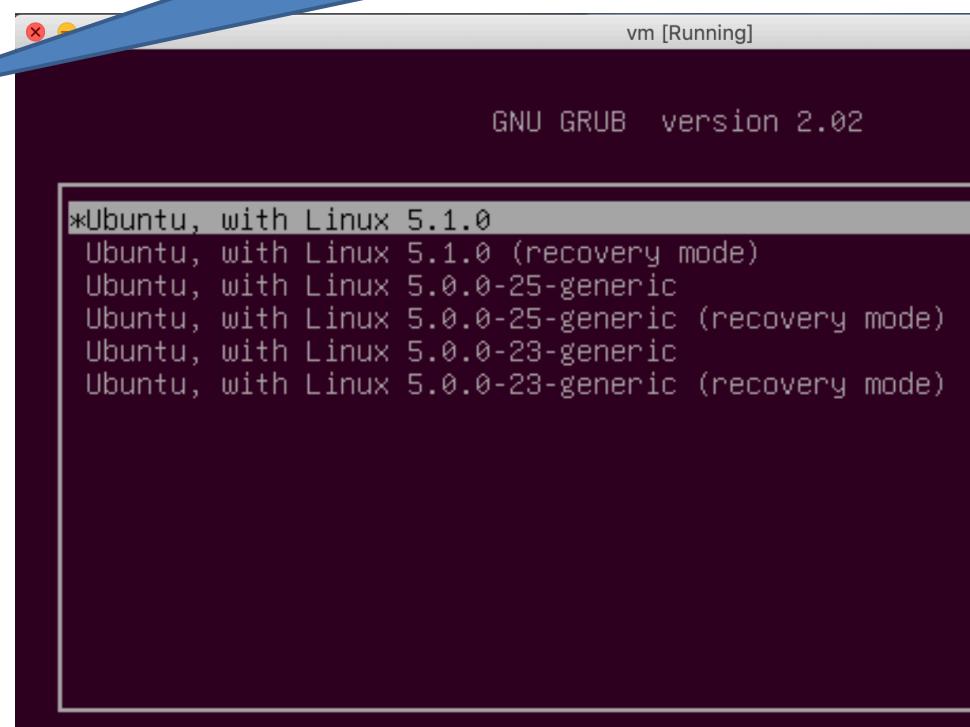
Make the following changes:

```
GRUB_DEFAULT=0
```

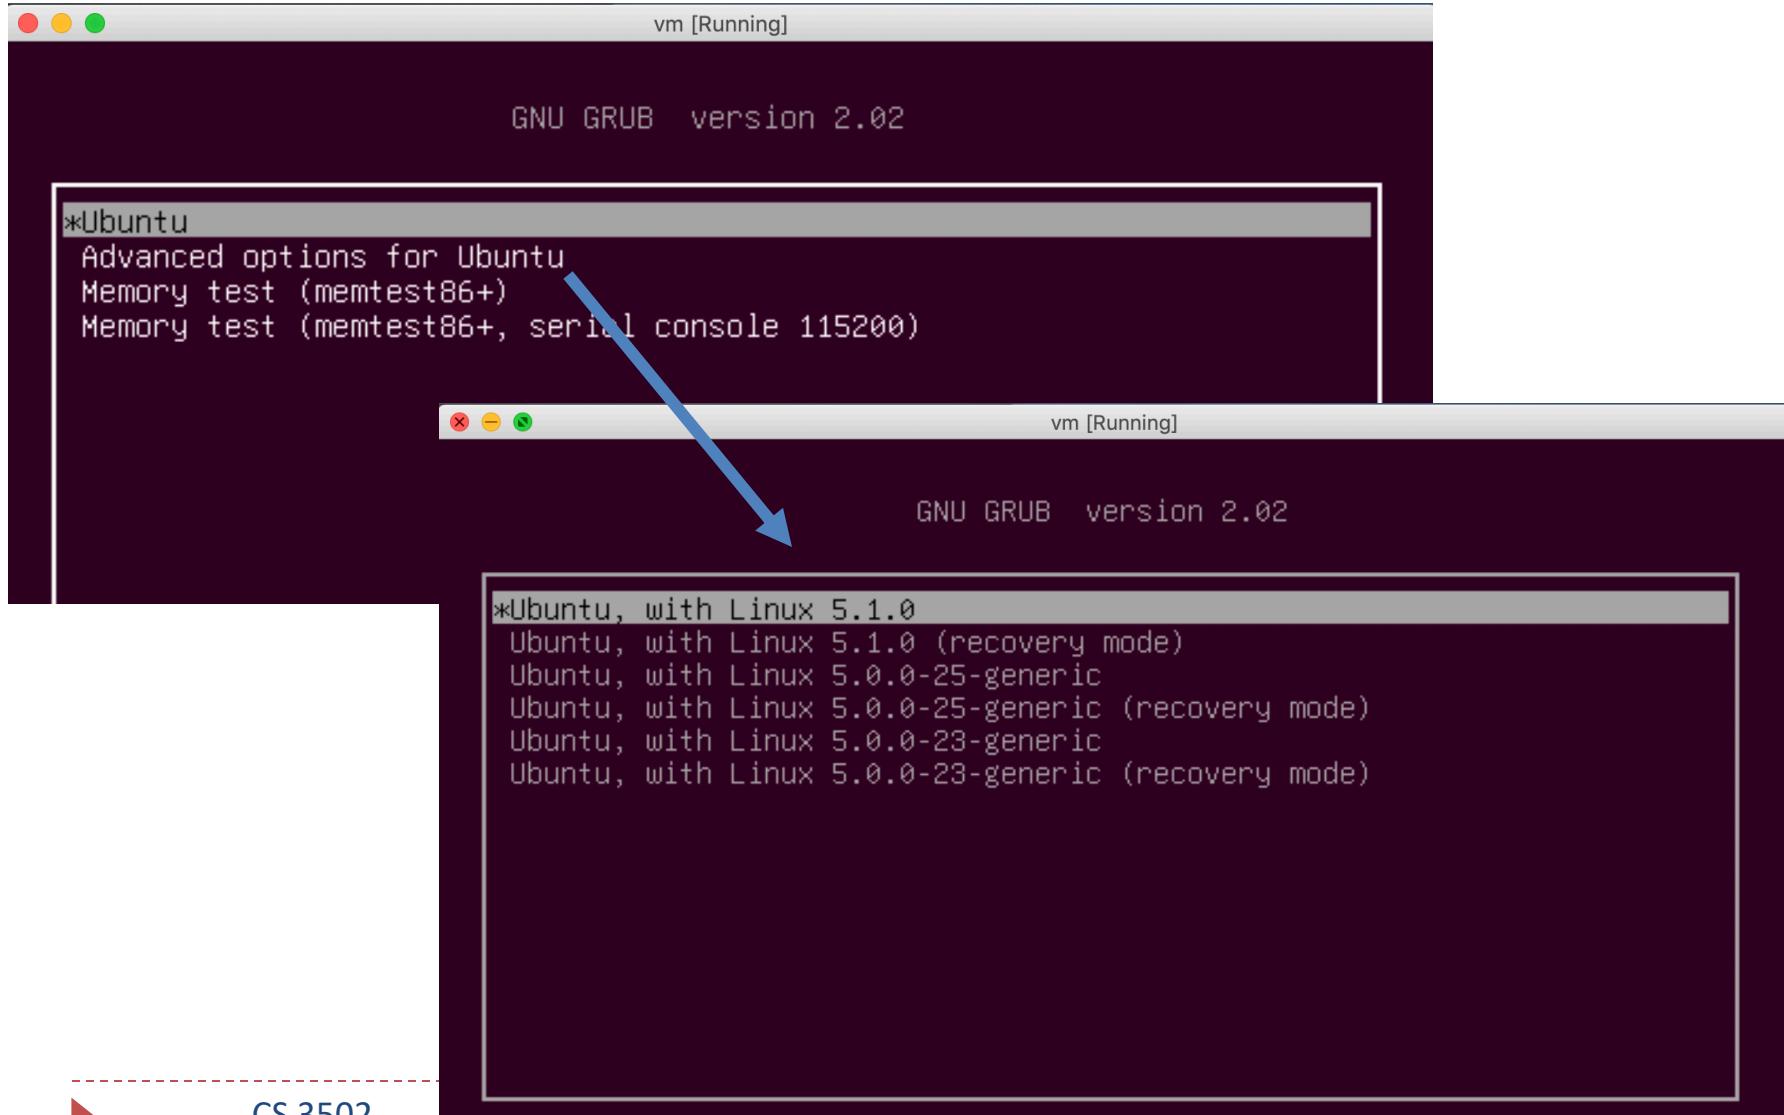
```
GRUB_TIMEOUT=10
```

Then, update the grub entry:

```
$ sudo update-grub2
```

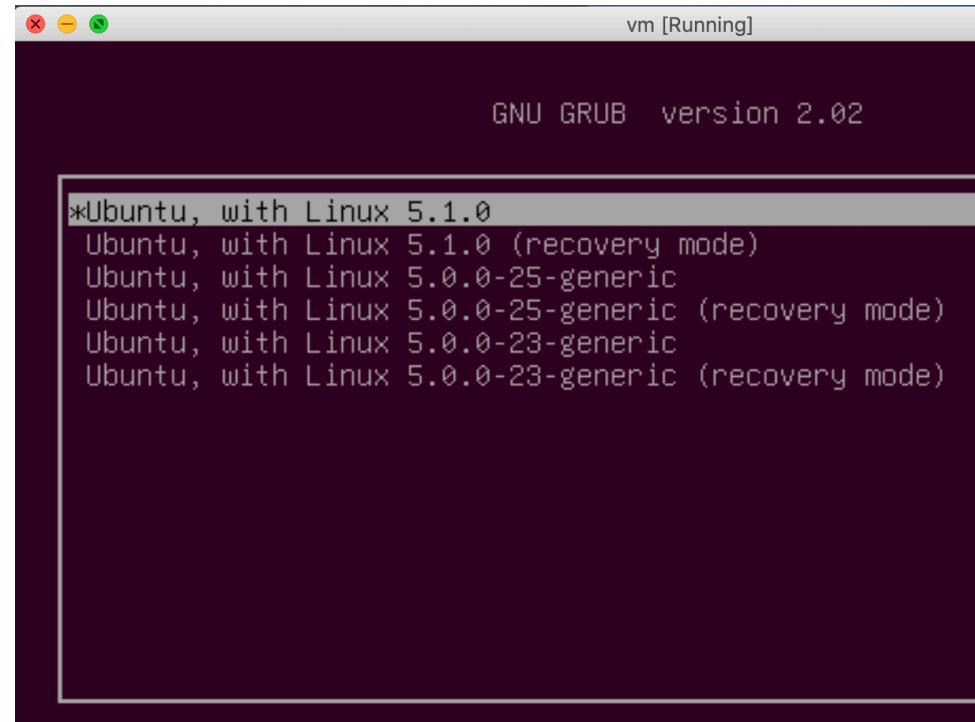


# Which kernel to boot if there are many?



# What if my kernel crashed?

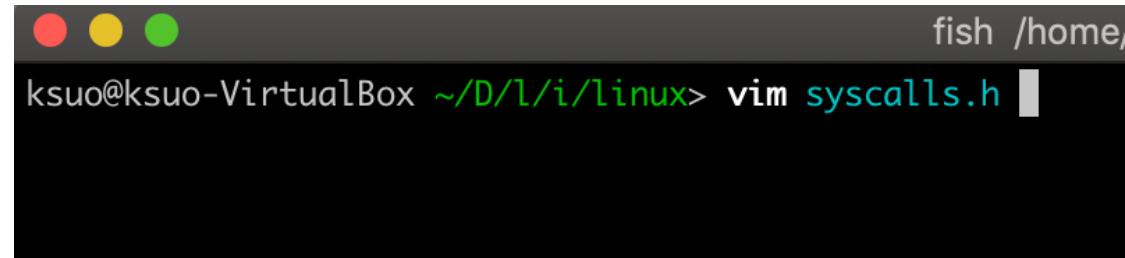
- Your kernel could crash because you might bring in some kernel bugs
- In the menu, choose the old kernel to boot the system
- Fix your bug in the source code
- Compile and reboot



# Project 1: No IDE, please use vim

- Open a file

- \$ vim test.c



```
fish /home/ksuo@ksuo-VirtualBox ~/D/l/i/linux> vim syscalls.h
```

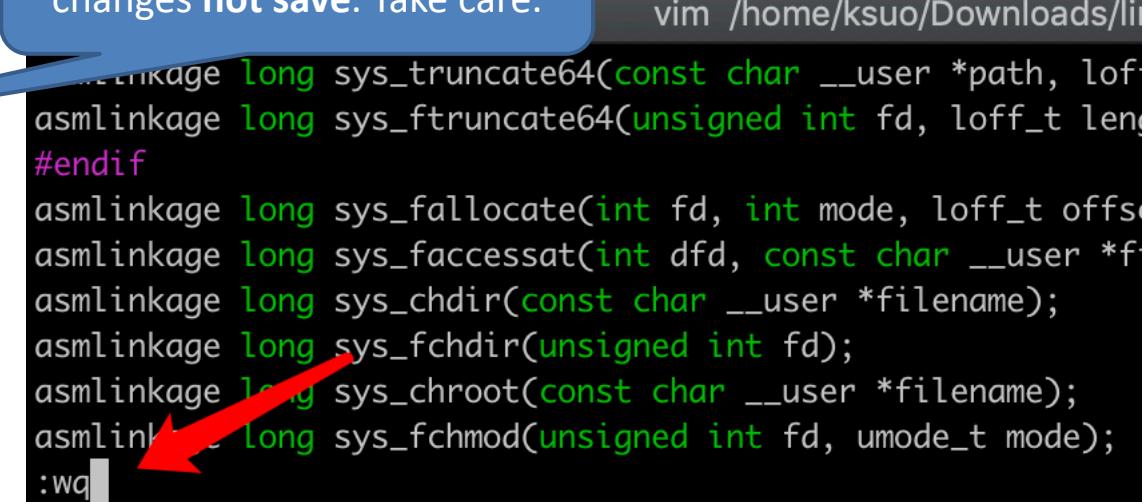
- Close a file

- Inside vim (after opening), press :q!

colon mark means menu.

Exclamation mark means changes **not save**. Take care.

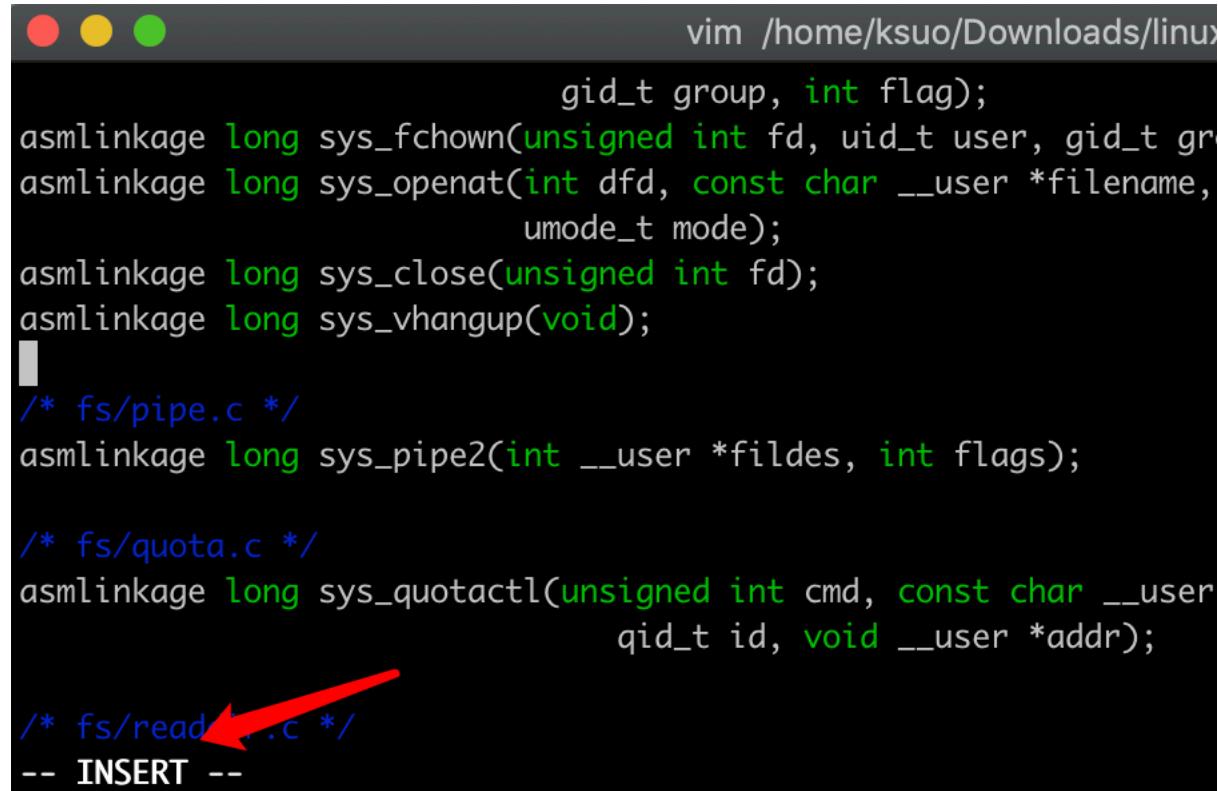
- Inside vim (after opening), press :wq, changes saved



```
asmlinkage long sys_truncate64(const char __user *path, loff_t len);  
asmlinkage long sys_ftruncate64(unsigned int fd, loff_t len);  
#endif  
asmlinkage long sys_fallocate(int fd, int mode, loff_t offset, loff_t len);  
asmlinkage long sys_faccessat(int dfd, const char __user *filename, int mode);  
asmlinkage long sys_chdir(const char __user *filename);  
asmlinkage long sys_fchdir(unsigned int fd);  
asmlinkage long sys_chroot(const char __user *filename);  
asmlinkage long sys_fchmod(unsigned int fd, umode_t mode);  
:wq
```

# Project 1: No IDE, please use vim

- Edit a file
  - Open a file
  - Press i (means enter the insert mode). then input your code
  - Press ESC to exit the insert mode
  - Close a file (:wq)



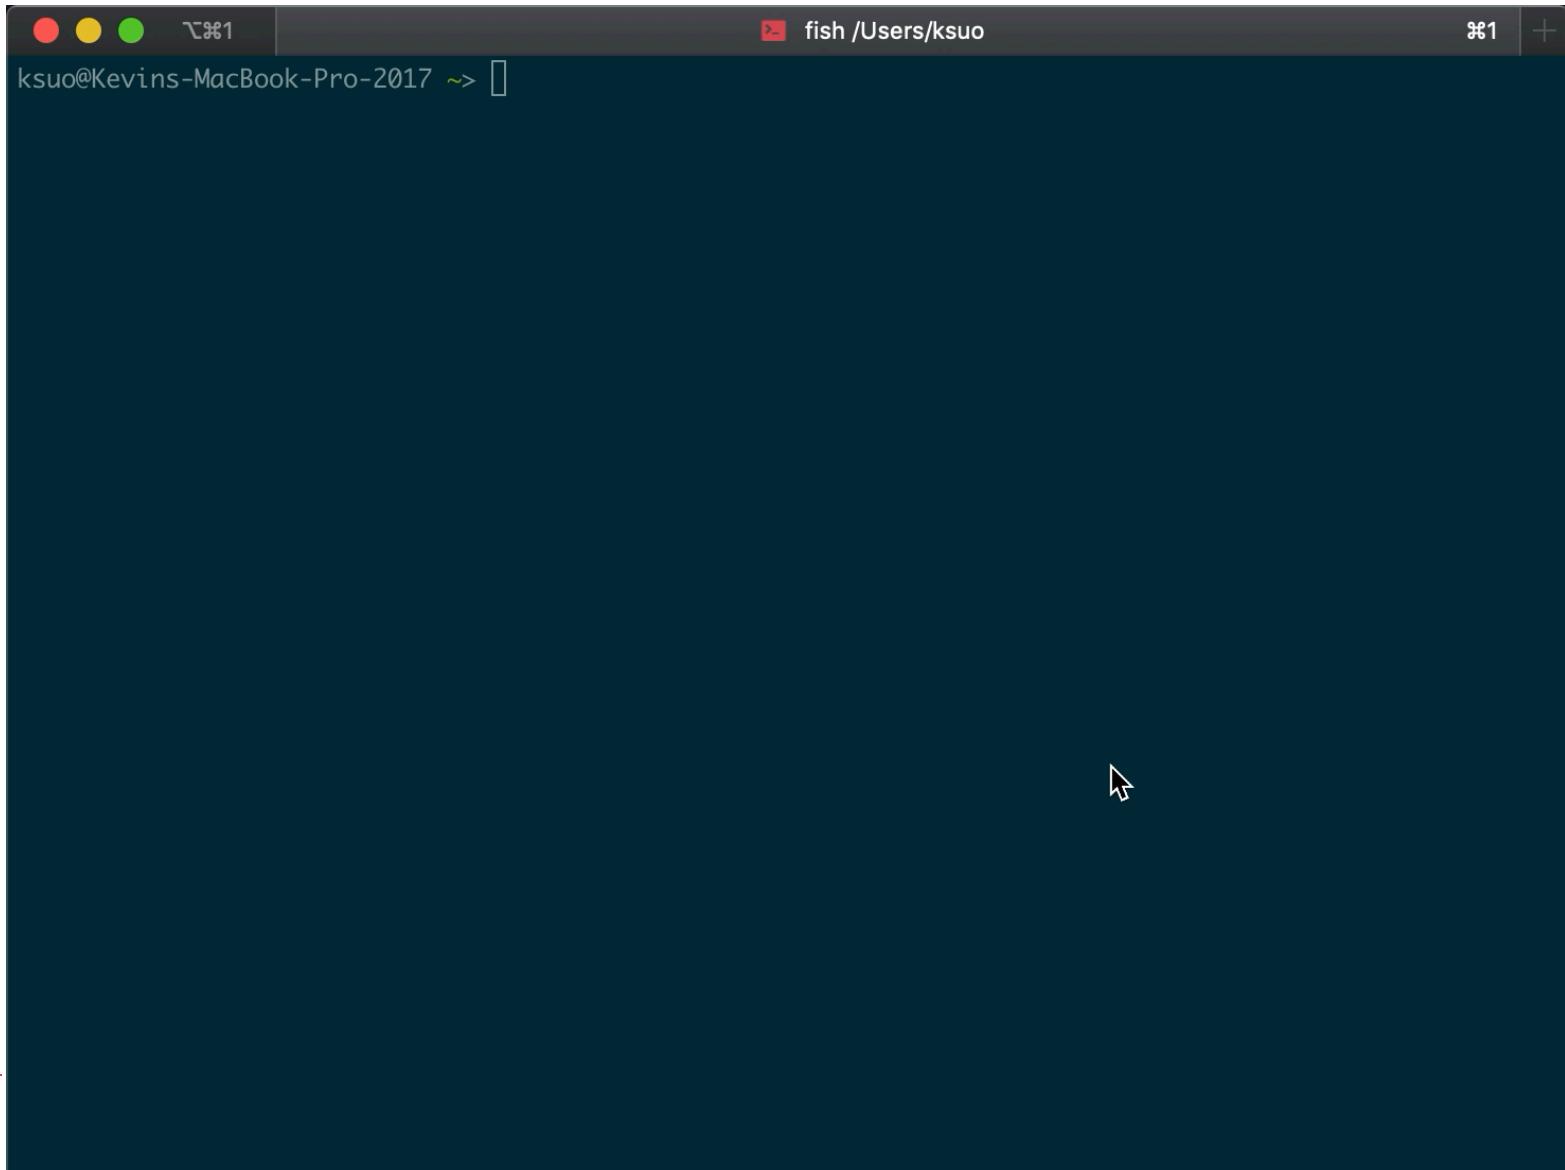
```
vim /home/ksuo/Downloads/linux-headers-4.15.0-38/include/linux/fs.h
gid_t group, int flag);
asm linkage long sys_fchown(unsigned int fd, uid_t user, gid_t gr
asm linkage long sys_openat(int dfd, const char __user *filename,
                           umode_t mode);
asm linkage long sys_close(unsigned int fd);
asm linkage long sys_vhangup(void);
/* fs/pipe.c */
asm linkage long sys_pipe2(int __user *fildes, int flags);

/* fs/quota.c */
asm linkage long sys_quotactl(unsigned int cmd, const char __user
                               qid_t id, void __user *addr);

/* fs/read.c */
-- INSERT --
```



# Project 1: No IDE, please use vim



# Edit a file with vim

---

- step 1: \$ vim file
- step 2: press **i**, enter insert mode; move the cursor to position and edit the context
- step 3: after editing, press **ESC** to exit the insert mode to normal mode
- step 4: press **:wq** to save what you edit and quit. If you do not want to save, press **:q!**



# More about vim

---

- A quick start guide for beginners to the Vim text editor
  - <https://eastmanreference.com/a-quick-start-guide-for-beginners-to-the-vim-text-editor>
- Vim basics:
  - <https://www.howtoforge.com/vim-basics>
- Learn the Basic Vim Commands [Beginners Guide]
  - [https://www.youtube.com/watch?time\\_continue=265&v=ZEGqkam-3lc](https://www.youtube.com/watch?time_continue=265&v=ZEGqkam-3lc)

