

CS 3502

Operating Systems

Memory Management

Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>


Outline

- Memory management overview
 - Memory abstraction and address spaces
 - Physical address and virtual address
 - Physical memory and virtual memory
- Memory management
 - Translation look-aside buffer
 - Page table
 - Multi-level page table

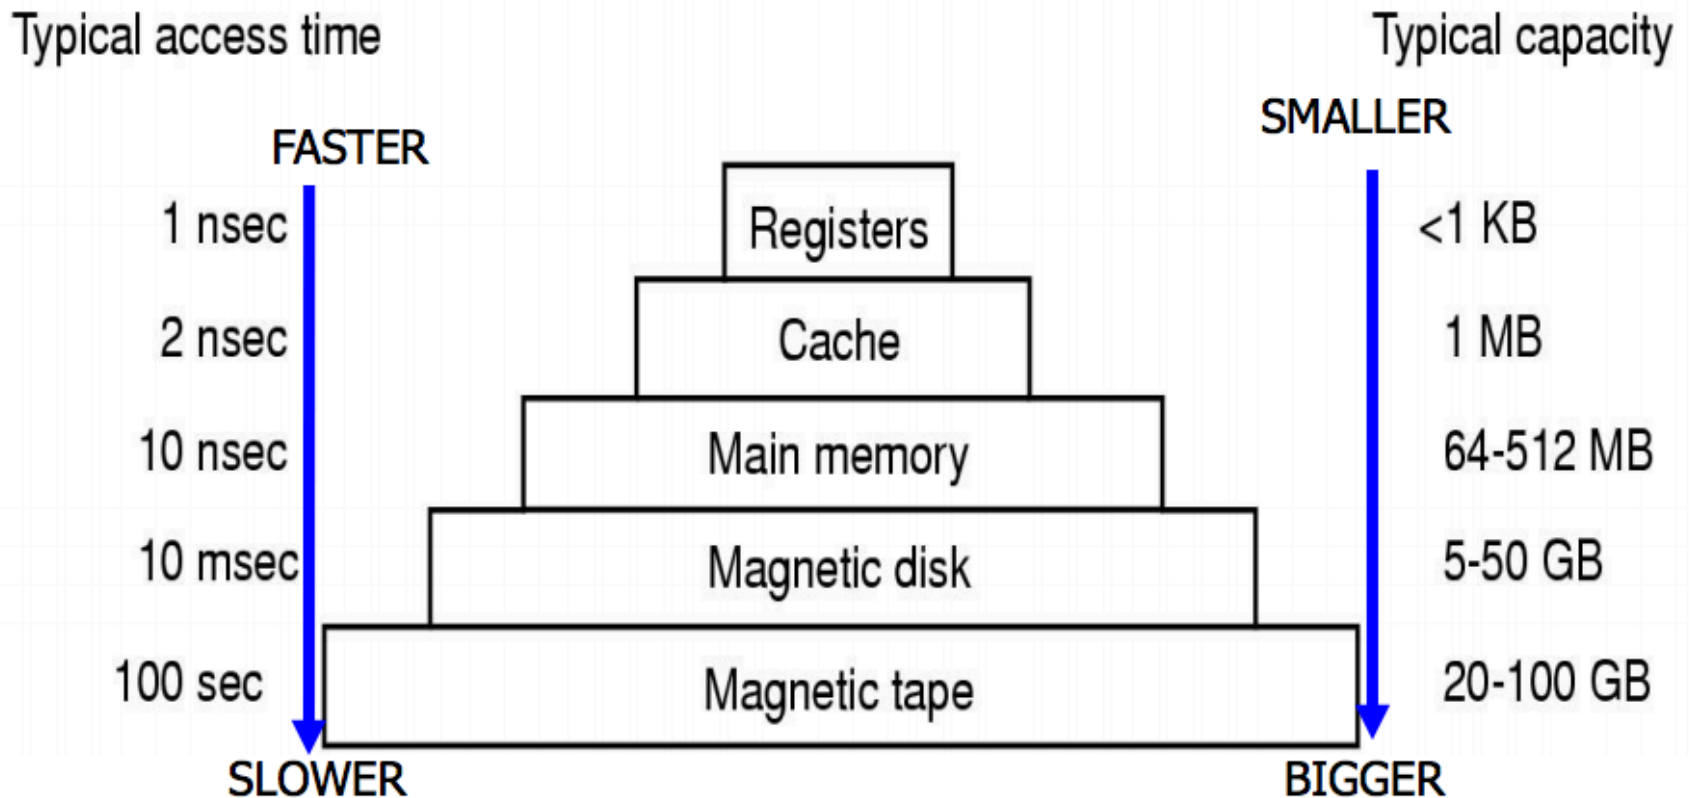


Why we need memory management

- Ideally user/programmers want memory that is
 - large --> 1GB, 2GB, 4GB, 16GB, ...
 - fast --> DDR 2, DDR 3, DDR 4, ...
 - non volatile --> not lose data when losing power



Typical Memory Hierarchy



Typical Memory Hierarchy

System Event	Actual Latency	Scaled Latency
One CPU cycle	0.4 ns	1 s
Level 1 cache access	0.9 ns	2 s
Level 2 cache access	2.8 ns	7 s
Level 3 cache access	28 ns	1 min
Main memory access (DDR DIMM)	~100 ns	4 min
Intel® Optane™ DC persistent memory access	~350 ns	15 min
Intel® Optane™ DC SSD I/O	<10 µs	7 hrs
NVMe SSD I/O	~25 µs	17 hrs
SSD I/O	50–150 µs	1.5–4 days
Rotational disk I/O	1–10 ms	1–9 months
Internet call: San Francisco to New York City	65 ms	5 years
Internet call: San Francisco to Hong Kong	141 ms	11 years

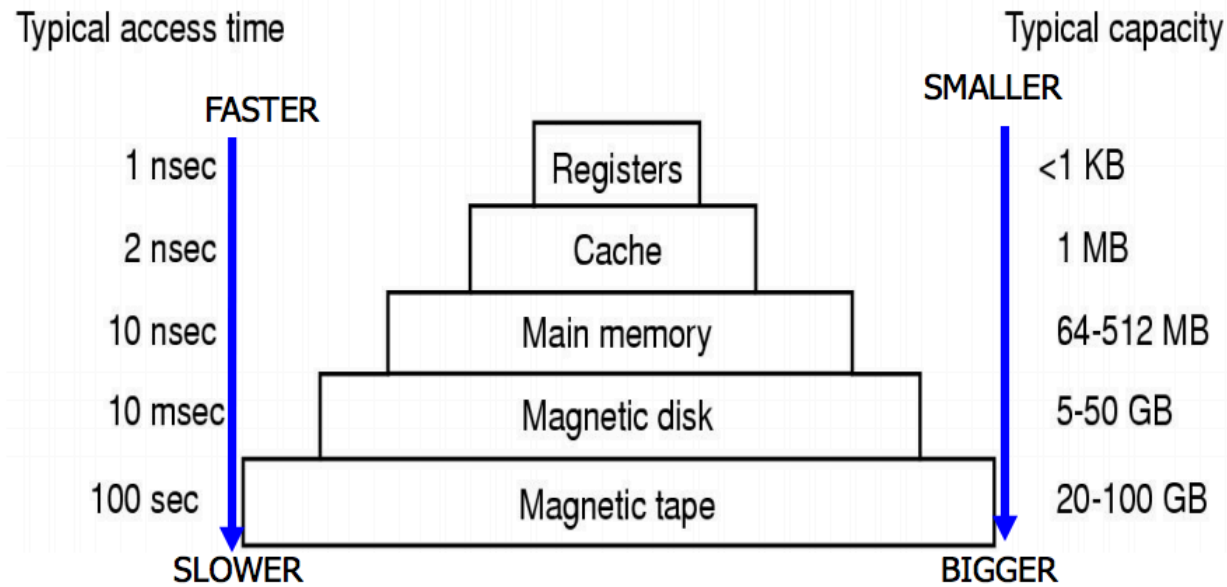
1000x than memory

20x than SSD

www.brendangregg.com/sysperfbook.html



Why we need memory hierarchy like this?



VS

Single layer memory hierarchy?

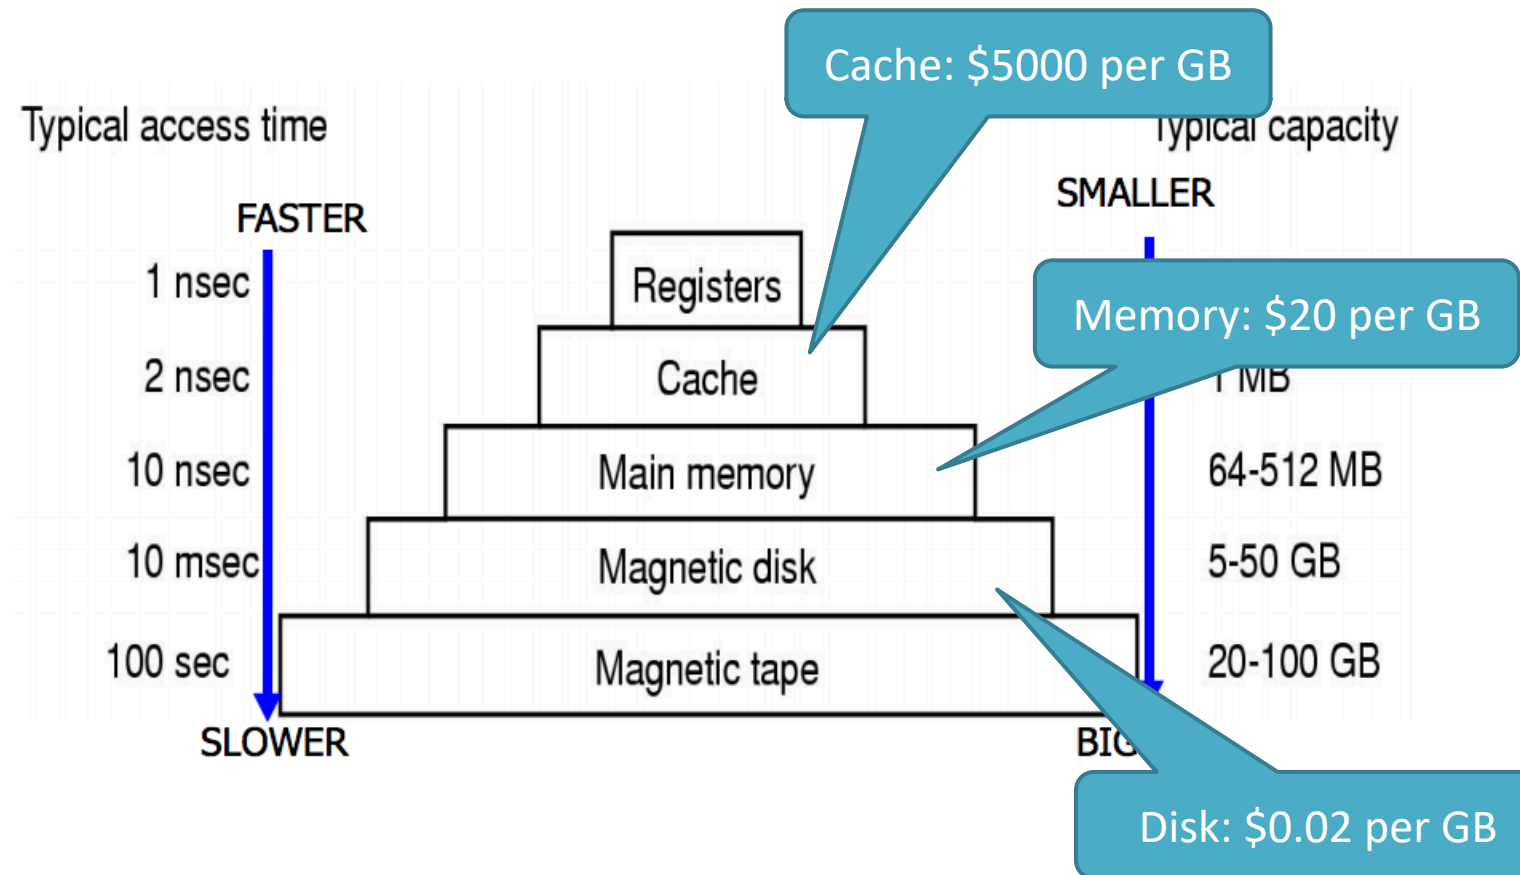
Memory



Why we need memory hierarchy like this?

- Cost


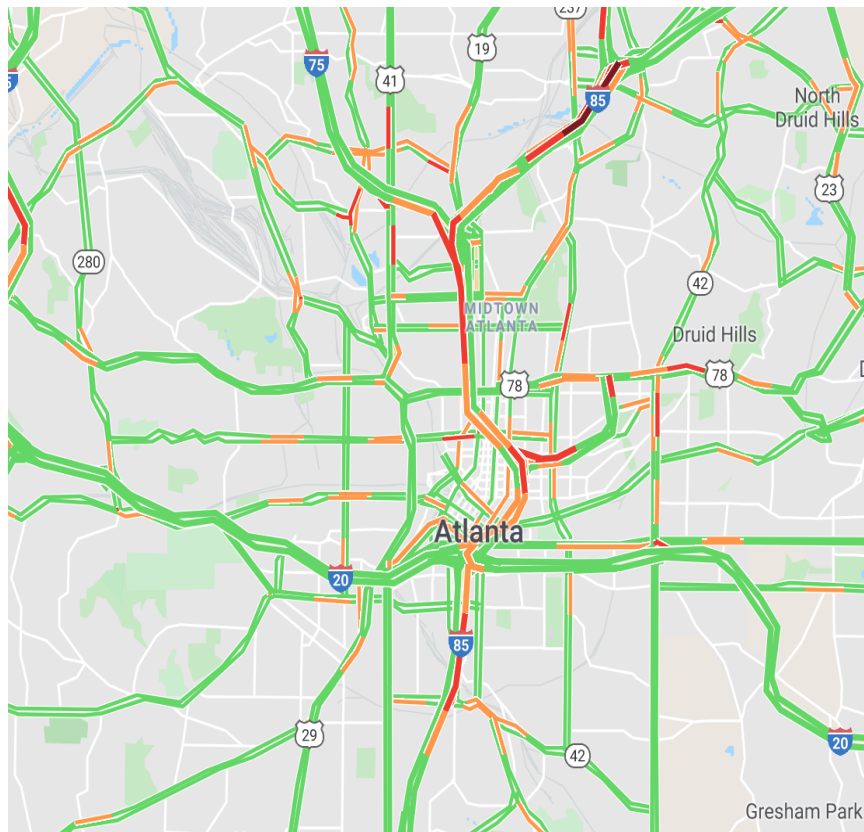
Year: 2010 data



Why we need memory hierarchy like this?

- 20/80 rule: only a small proportion of code/path is hot in programs

We only need to put the hot code inside the top hierarchy of memory



Why we need memory management

Ideal: programmers/users want memory that is	Reality: Memory hierarchy
large	small amount of fast, expensive memory e.g., register, cache
fast	some medium-speed, medium price main memory, lots of slow, cheap disk storage
Non-volatile	Volatile (e.g., top hierarchy)

Memory management handles the semantic gap

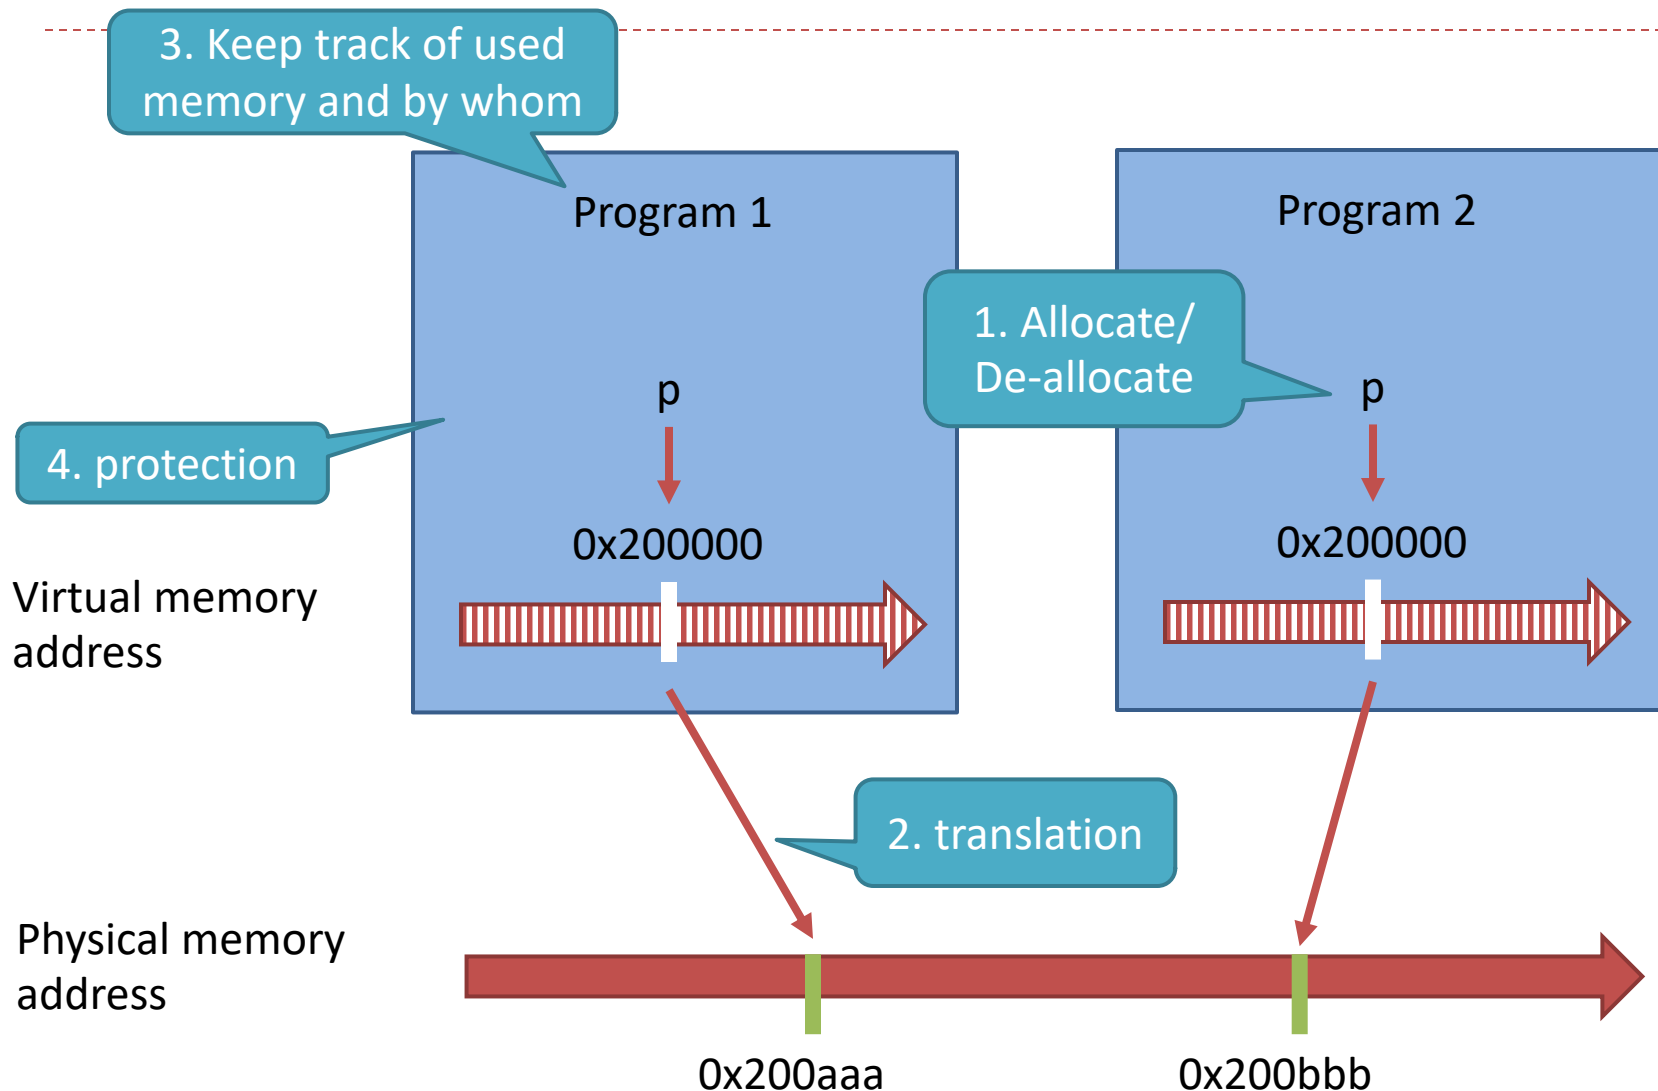


What will memory management do?

- Memory management key tasks:
 1. Allocate and de-allocate memory for processes (performed by OS and improve the programming efficiency)
 2. Address translation (e.g., between physical address and virtual address)
 3. Keep track of used memory size and used by whom
 4. Memory protection

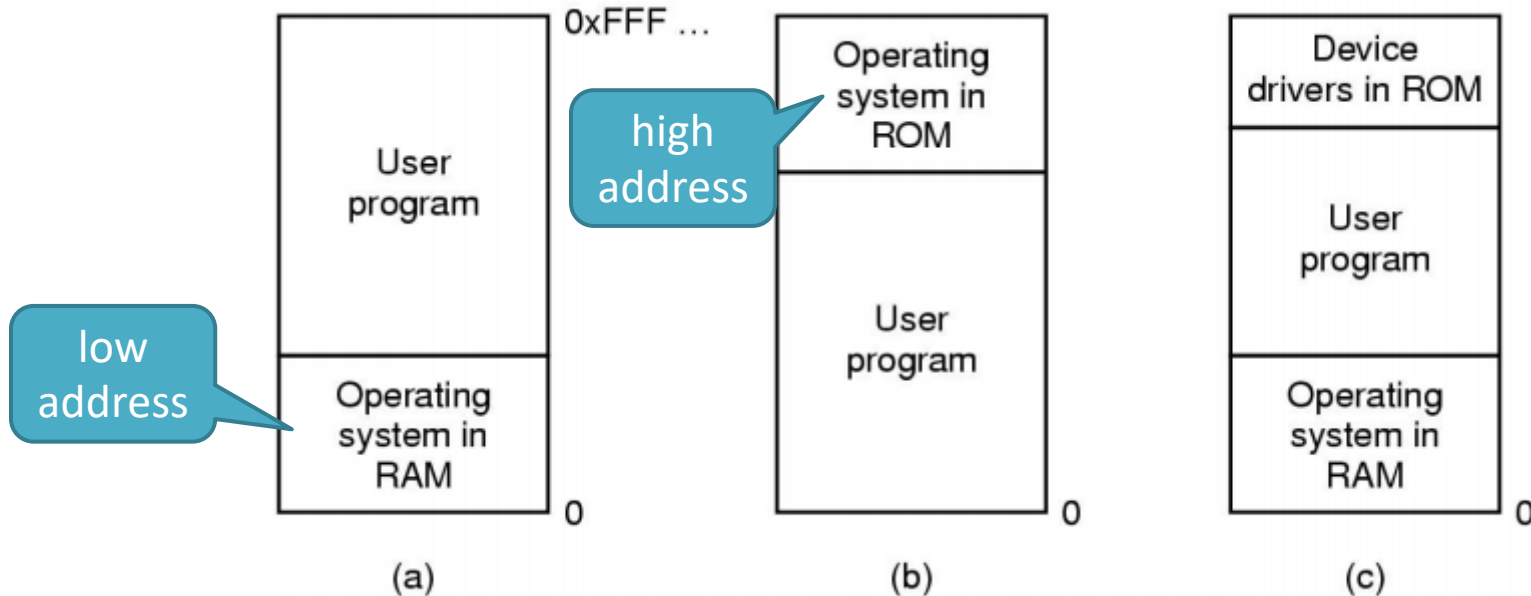


These key tasks are realized by memory abstraction



Without memory abstraction (early era)

Memory management is simple. Put small piece for OS, the rest for apps.



Three simple ways of organizing memory:

(a) early mainframes.


(b) Handheld and embedded systems.

(c) early PC.



Without memory abstraction (early era)

Memory management is simple. Put small piece for OS, the rest for apps.



Problems:

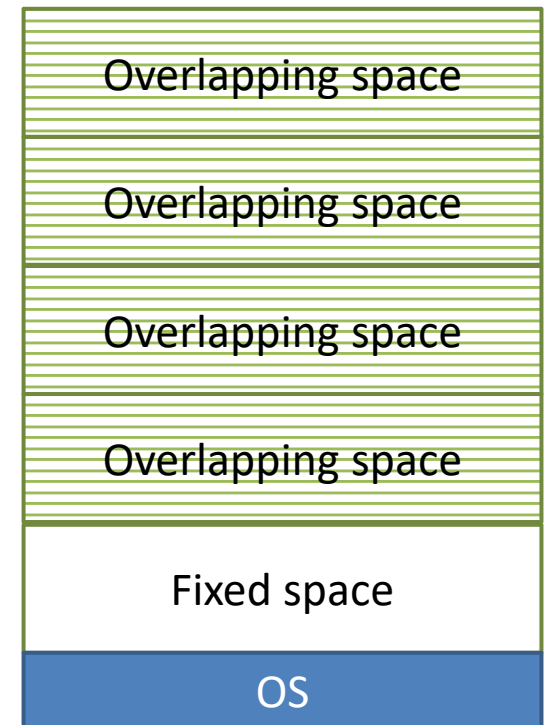
Simple memory organization: Cannot support **multi-programming** and only works for devices, like washing machines, microwaves, etc.



Without memory abstraction (What will happen when memory is not enough?)

- Overlapping



- Memory is small and cannot hold all program data
- User space is divided to one fixed space and several overlapping space
- Fixed space: hold most frequent data
- Overlapping space: the data will overlap those which never be used (after overlapping, data is gone)



Without memory abstraction (What will happen when memory is not enough?)

- Swapping



- Leave memory and are **swapped out** to disk (data still there)
- Re-enter memory by getting **swapped in** from disk



Without memory abstraction (What will happen when memory is not enough?)

- Difference

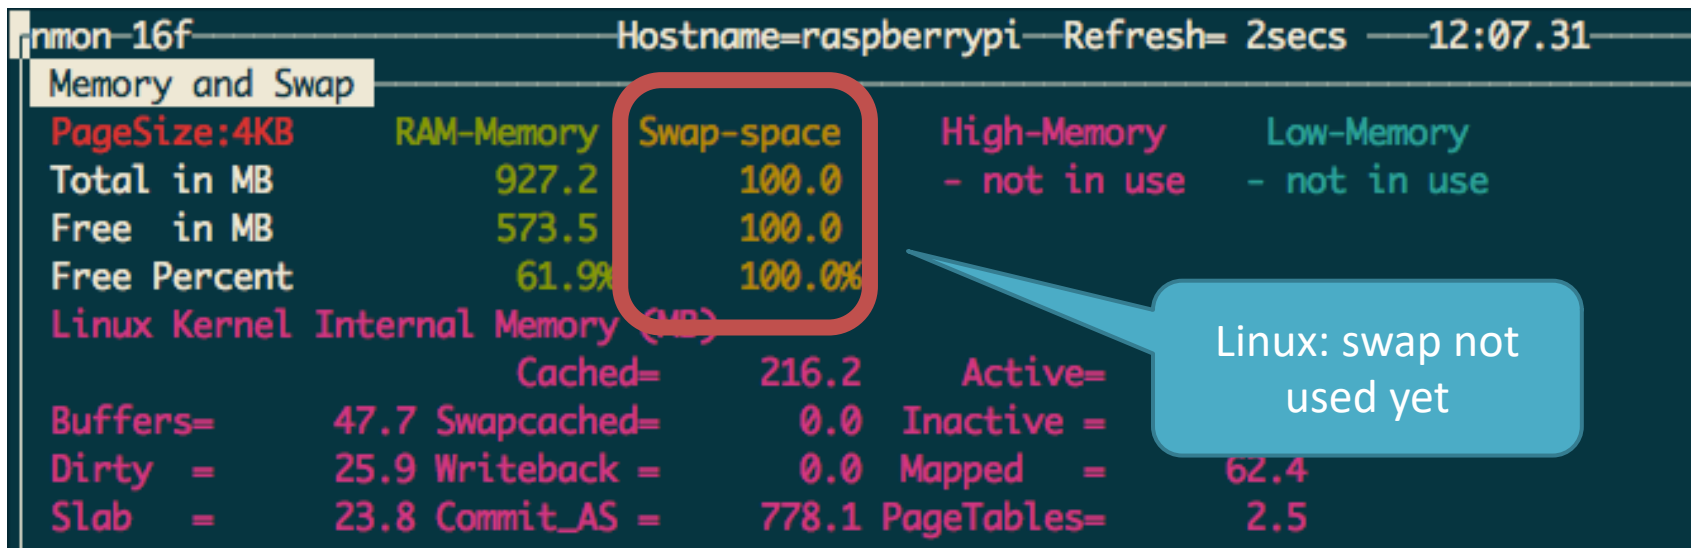
- Swapping is used for **different processes** while overlapping is used for **one program/process** (for security, otherwise causing app crash)
- Swapping is still **widely used** today while overlapping is **not**



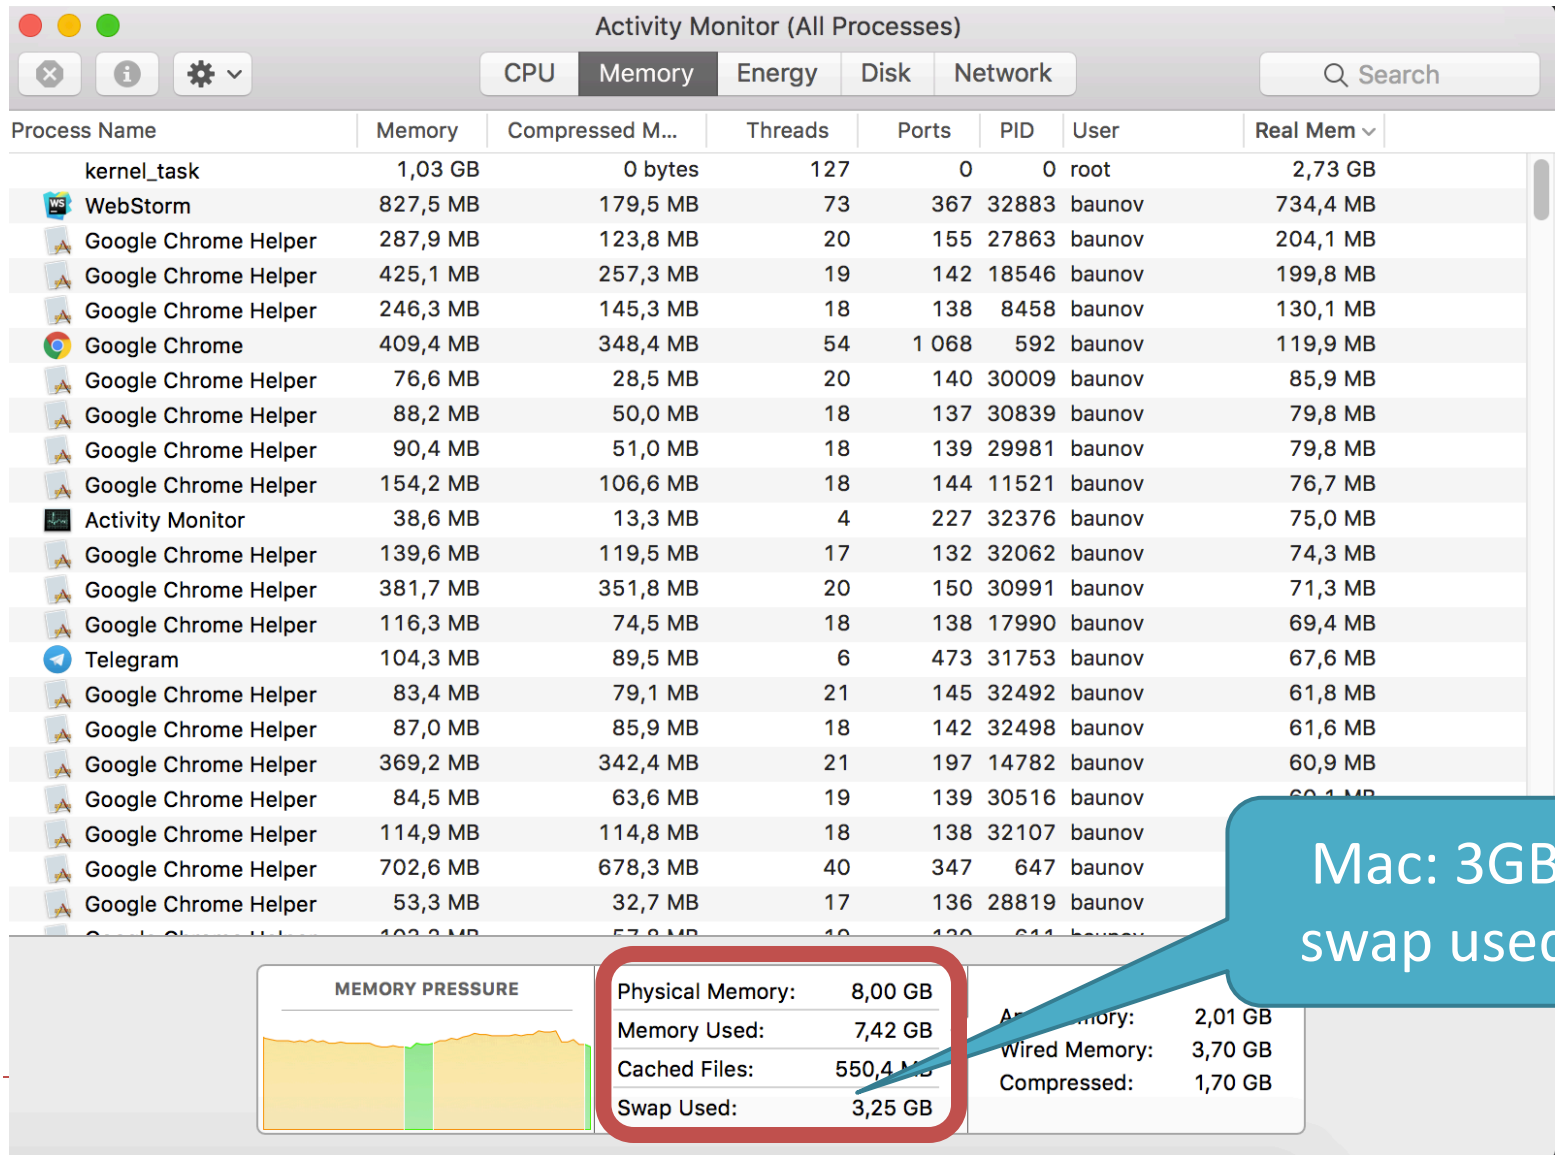
Swap example in Linux

- \$ nmon, then press 'm' to check statistics from memory

```
nmon-16f-----Hostname=raspberrypi---Refresh= 2secs ---12:07.31
Memory and Swap
PageSize:4KB    RAM-Memory    Swap-space    High-Memory    Low-Memory
Total in MB      927.2        100.0         - not in use   - not in use
Free in MB       573.5        100.0
Free Percent     61.9%       100.0%
Linux Kernel Internal Memory (MB)
                  Cached=      216.2      Active=
Buffers=         47.7 Swapcached=    0.0 Inactive =
Dirty  =         25.9 Writeback =    0.0 Mapped  =    62.4
Slab   =         23.8 Commit_AS =   778.1 PageTables=    2.5
```



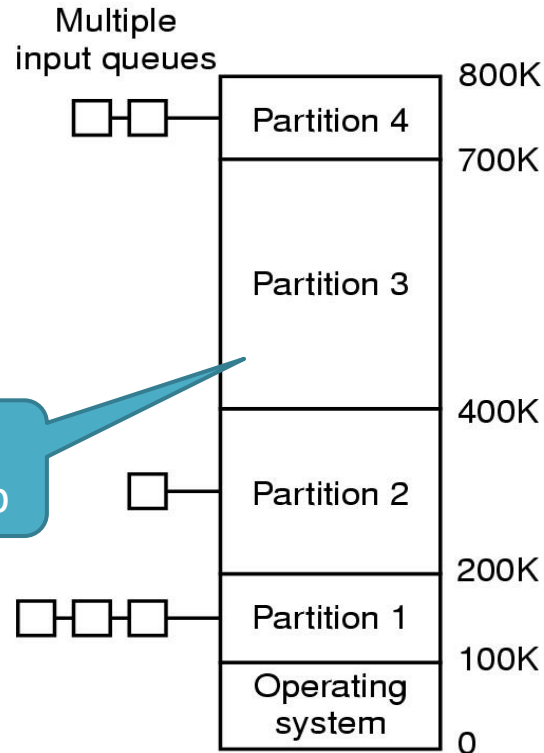
Swap example in Mac



Without memory abstraction (How to support multiprogramming)

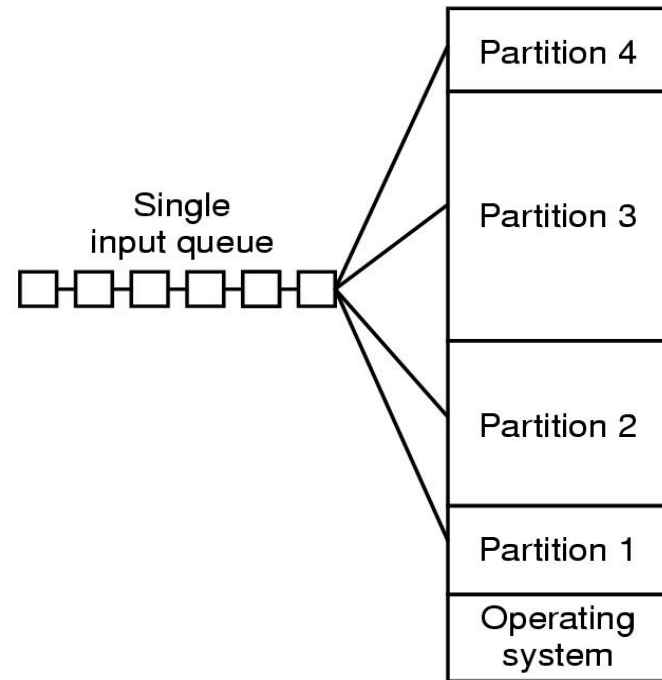
with Fixed Partitions

Each partition is used for one app



(a)

(a) separate input queues of processes for each partition



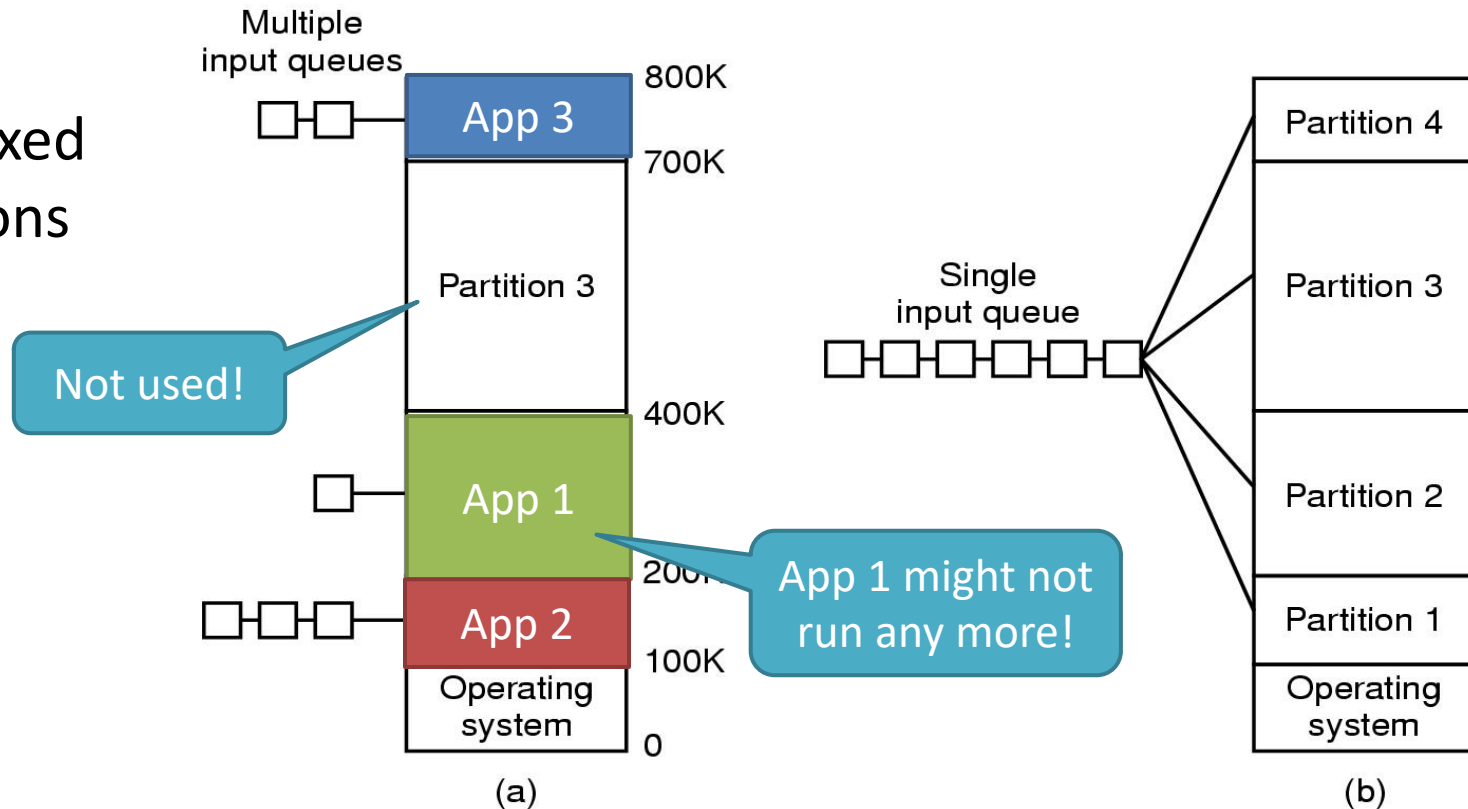
(b)

(b) single input queue



Without memory abstraction (How to support multiprogramming)

with Fixed Partitions




Problems:

- (a) fragmentation
- (b) low efficiency




Memory Abstraction: Address Spaces

- Program should have their own views of memory
 - The address space – logical address
 - Non-overlapping address spaces – protection
 - Move a program by mapping its addresses to a different place - relocation

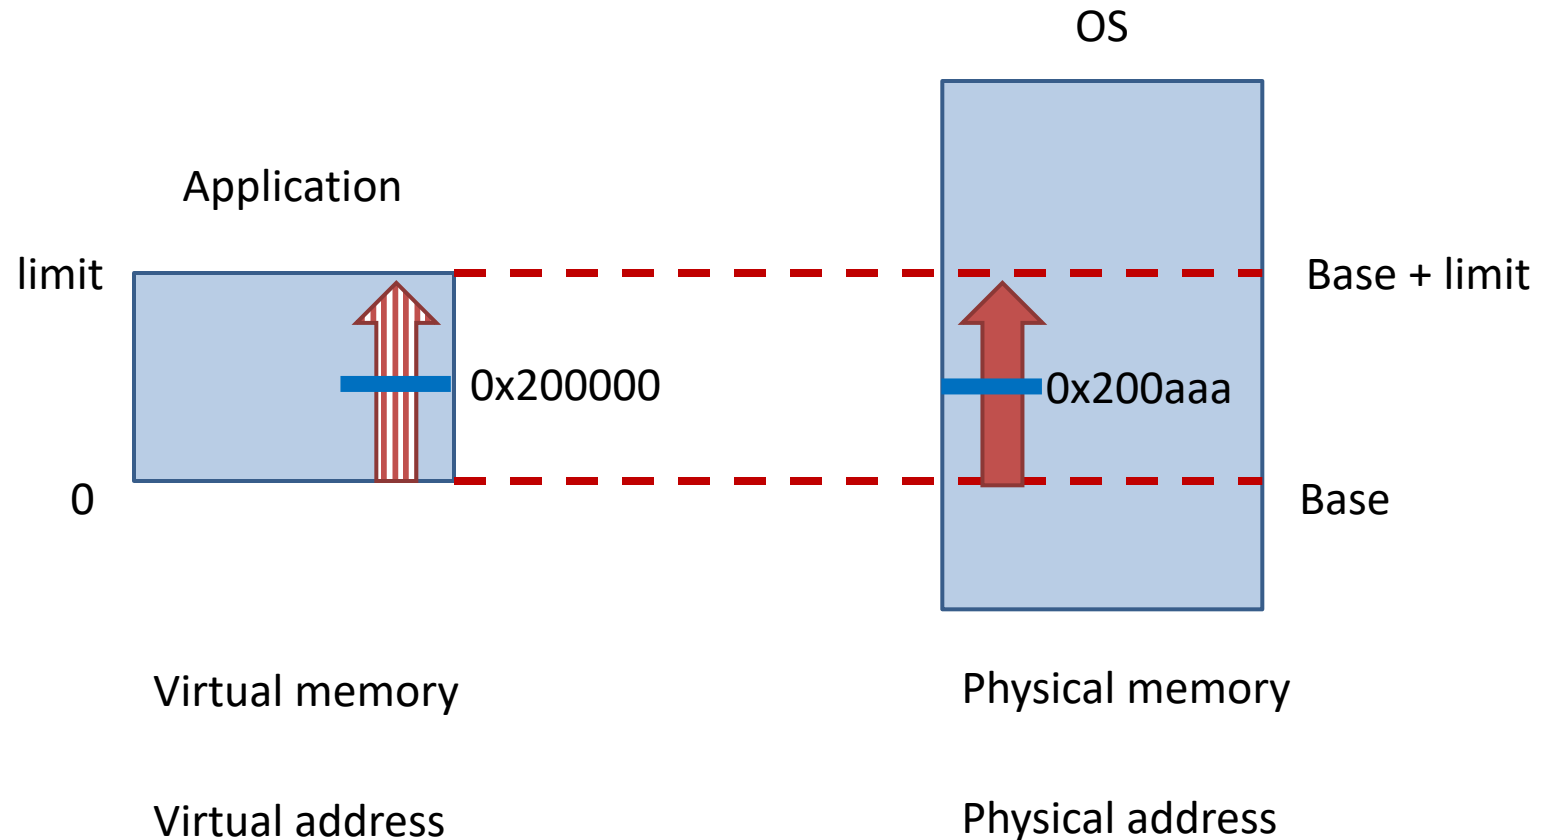


Memory Abstraction: Address Spaces

- OS dynamically relocates programs in memory
 - Two hardware registers: base and limit
 - **Base**: start address of a process
 - **Limit**: the upper bound address of a process
 - Hardware adds relocation register (base) to virtual address to get a physical address



Memory Abstraction: Address Spaces



Physical address vs. Virtual address

- Virtual address: A memory address that an operating system allows a process to use (**Program view**)
- Physical address: A unit address for memory chip level that corresponds to the address bus to which the processor and CPU are connected (**OS view**)

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    int x = 3;

    printf("location of code : %p\n", (void *) main);
    printf("location of heap : %p\n", (void *) malloc(1));
    printf("location of stack : %p\n", (void *) &x);

    return 0;
}
```


All virtual addresses
(program view)

```
pi@raspberrypi ~> ./test.o
location of code : 0x10470
location of heap : 0x156a410
location of stack : 0x7ef381c4
```




Physical memory vs. Virtual memory

- Virtual memory: the combined size of the program, data, and stack **may exceed** the amount of physical memory available (due to swapping).
- Memory is divided into pages and OS decides which pages stay in **memory** and which get moved to **disk**.



Page number and offset

- Suppose the page size is 4KB, the address is 20000
 - The virtual page number is $20000/4096 = 4$
 - The offset is $20000\%4096 = 3616$




Question

- Consider a machine that has a 32-bit virtual address space and 4K Byte page size.

1. What is the memory size?

$$2^{32} = 4\text{GB}$$

$$//2^{10} = 1\text{KB}, 2^{20} = 1\text{MB}, 2^{30} = 1\text{GB}$$



Question


- Consider a machine that has a 32-bit virtual address space and 4K Byte page size.

2. How many virtual pages could a process have?

The total memory is 4GB

The single page size is 4KB

So the page number = $4\text{GB}/4\text{KB} = 2^{20}$




Question

- Consider a machine that has a 32-bit virtual address space and 4K Byte page size.


3. Given a 4KB page, how many address bits do we need for the offset ?

$$\text{Log}_2(4\text{KB}) = \text{Log}_2(2^{12}) = 12$$



Question

- Consider a machine that has a 32-bit virtual address space and 4K Byte page size.



Question

- Consider a machine that has a 32-bit virtual address space and 8KByte page size.

1. What is the total size (in bytes) of the virtual address space for each process?

Total size (in bytes) of the virtual address space for each process
= $2^{32} = 4 * 1024 * 1024 * 1024$ bytes
= 4 GB



Question

- Consider a machine that has a 32-bit virtual address space and 8KByte page size.

2. How many bits in a 32-bit address are needed to determine the page number of the address?

Number of pages in virtual address space = $4\text{GB}/8\text{KB} = 512 * 1024 = 2^9 * 2^{10} = 2^{19}$

So the number of bits in a 32-bit address are needed to determine the page number of the address is 19 bits

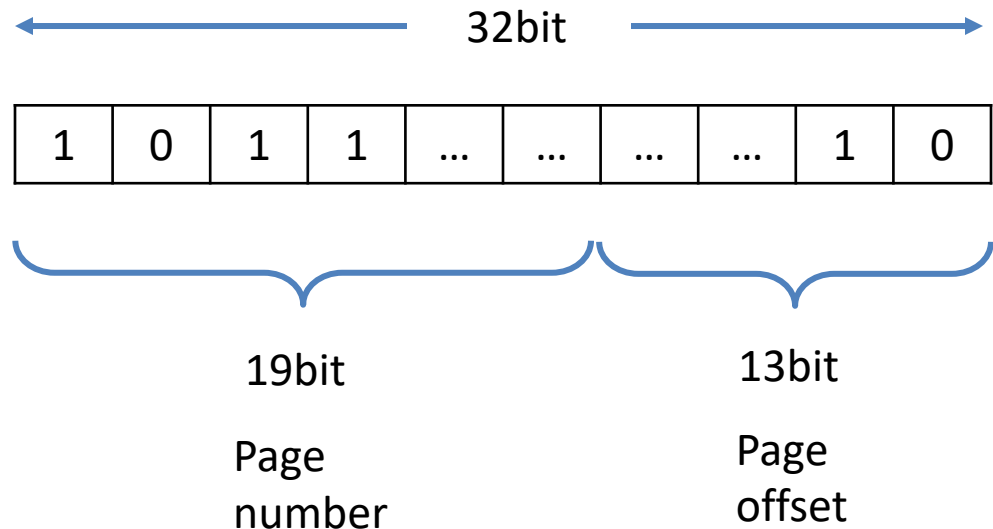


Question

- Consider a machine that has a 32-bit virtual address space and 8KByte page size.

3. How many bits in a 32-bit address represent the byte offset into a page?

$$32 - 19 = 13 \text{ bits}$$



Question

- Consider a machine that has a 32-bit virtual address space and 8KByte page size.

4. How many page-table entries are present in the page table?

Number of PTEs
= Number of pages in virtual address
= 4GB/8KB
= 2^{19} pages




Outline

- Memory management overview
 - Memory abstraction and address spaces
 - Physical address and virtual address
 - Physical memory and virtual memory
- Memory management
 - Translation look-aside buffer
 - Page table
 - Multi-level page table



Physical and virtual address translation

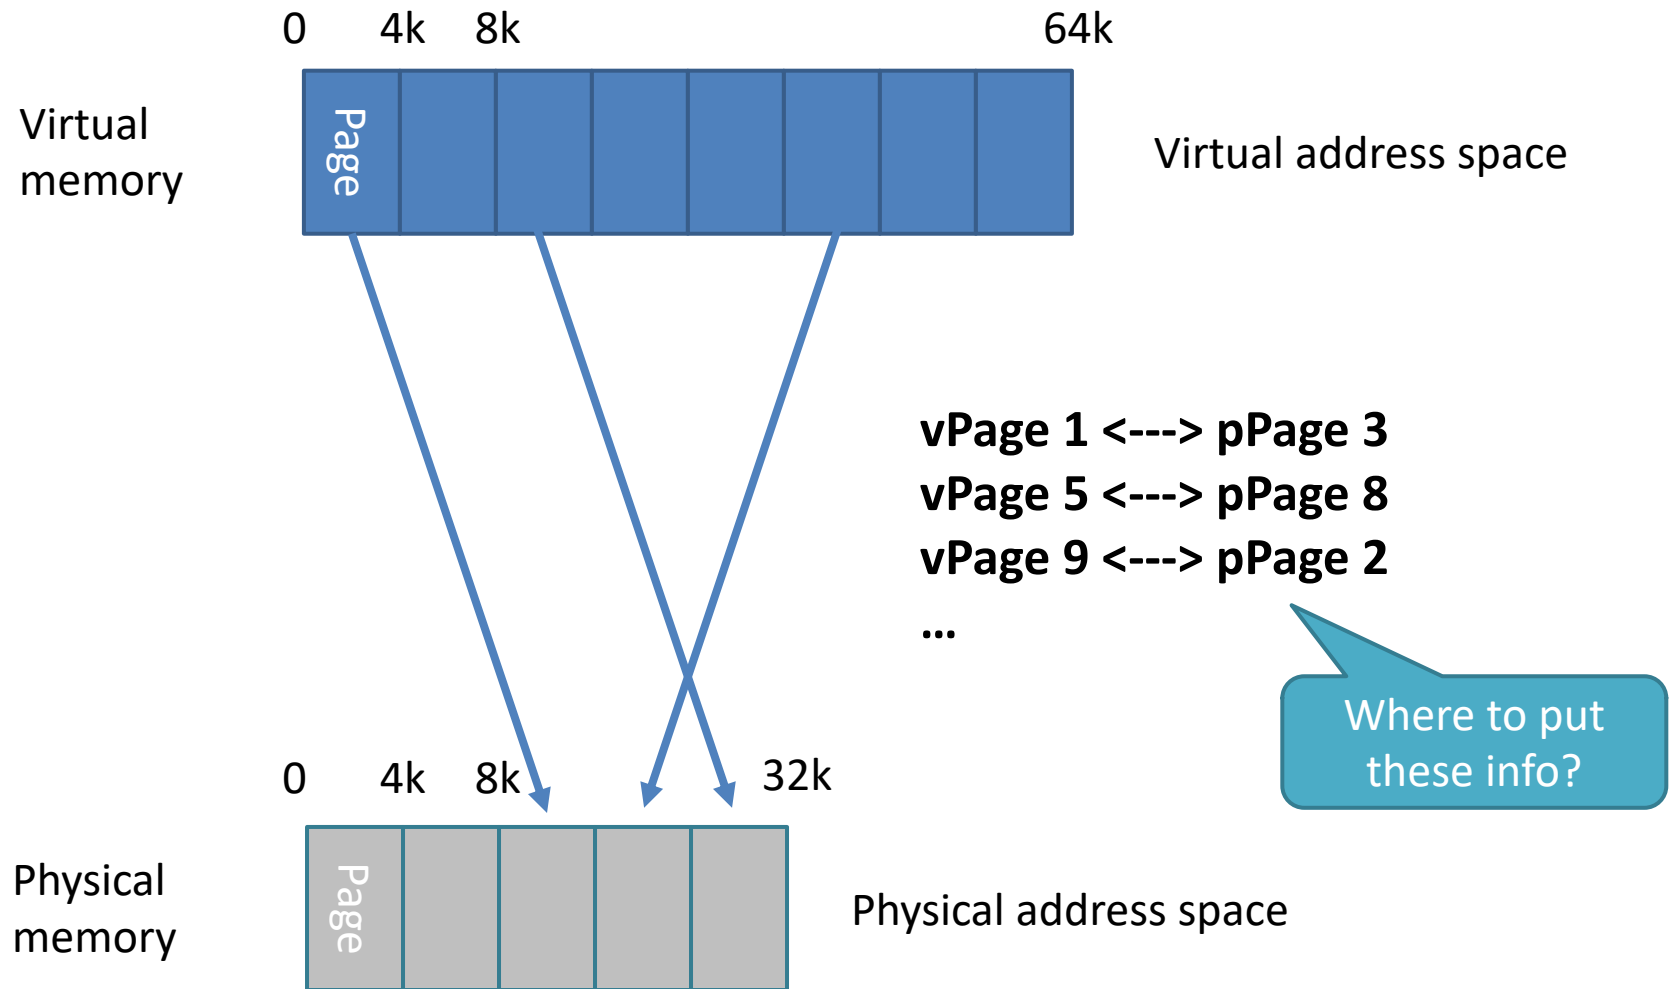


Logical program works in its
contiguous virtual address space

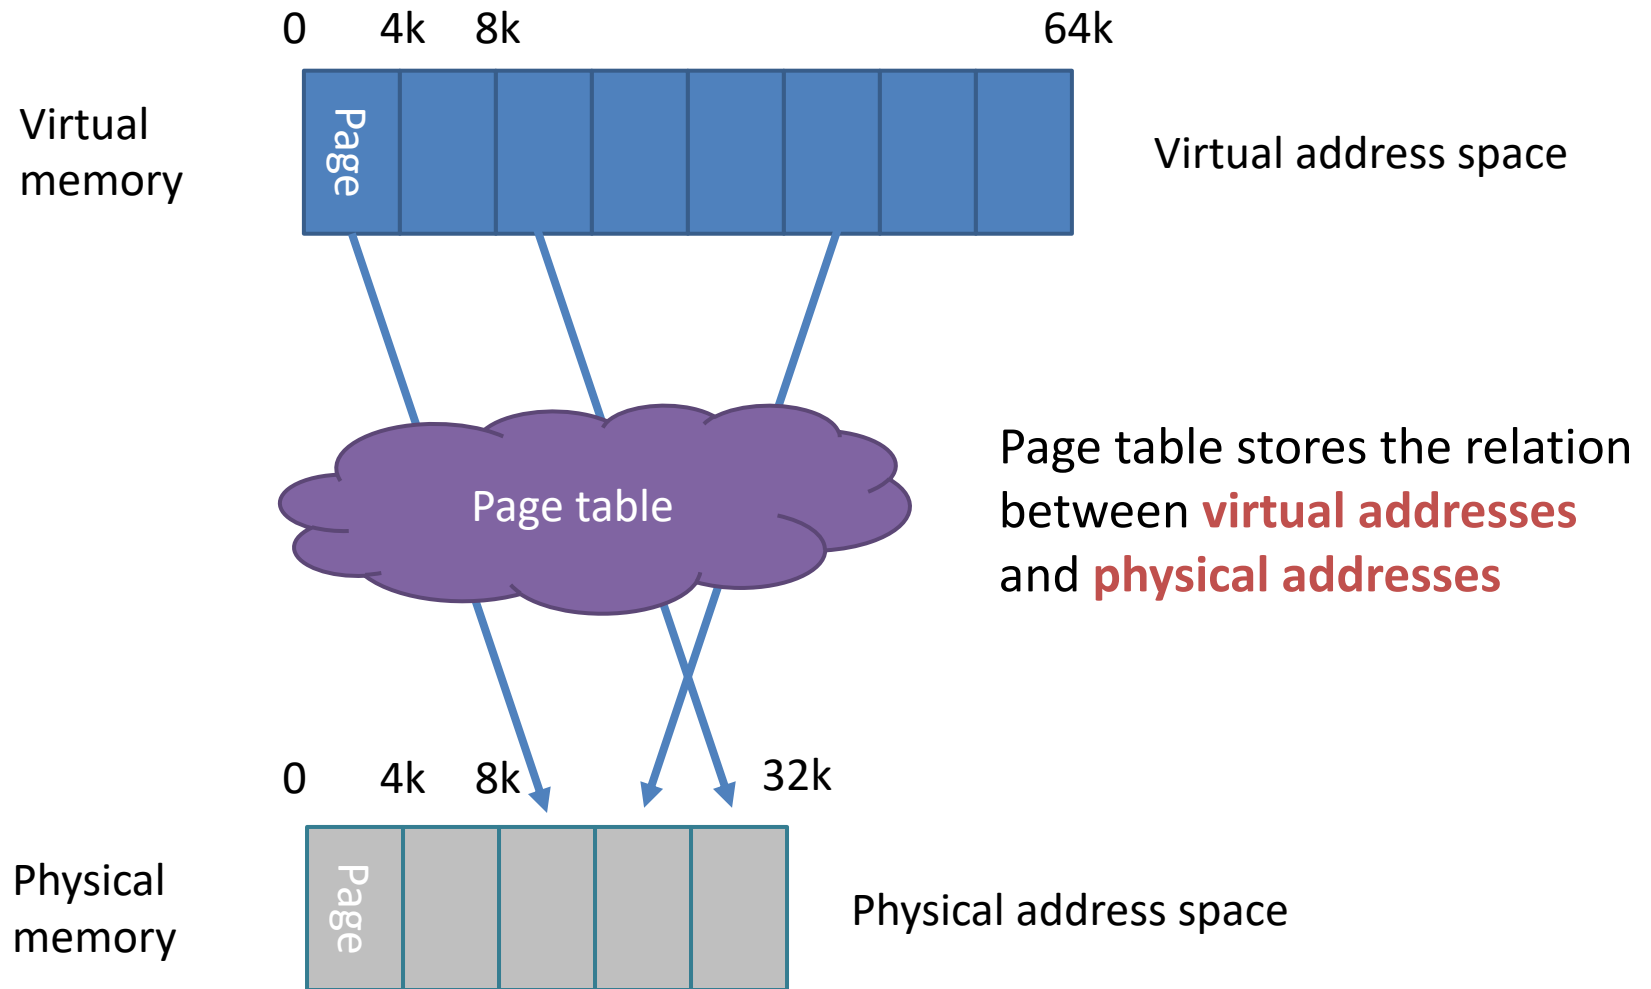
Address translation
done by MMU

Actual locations of the
data in physical memory

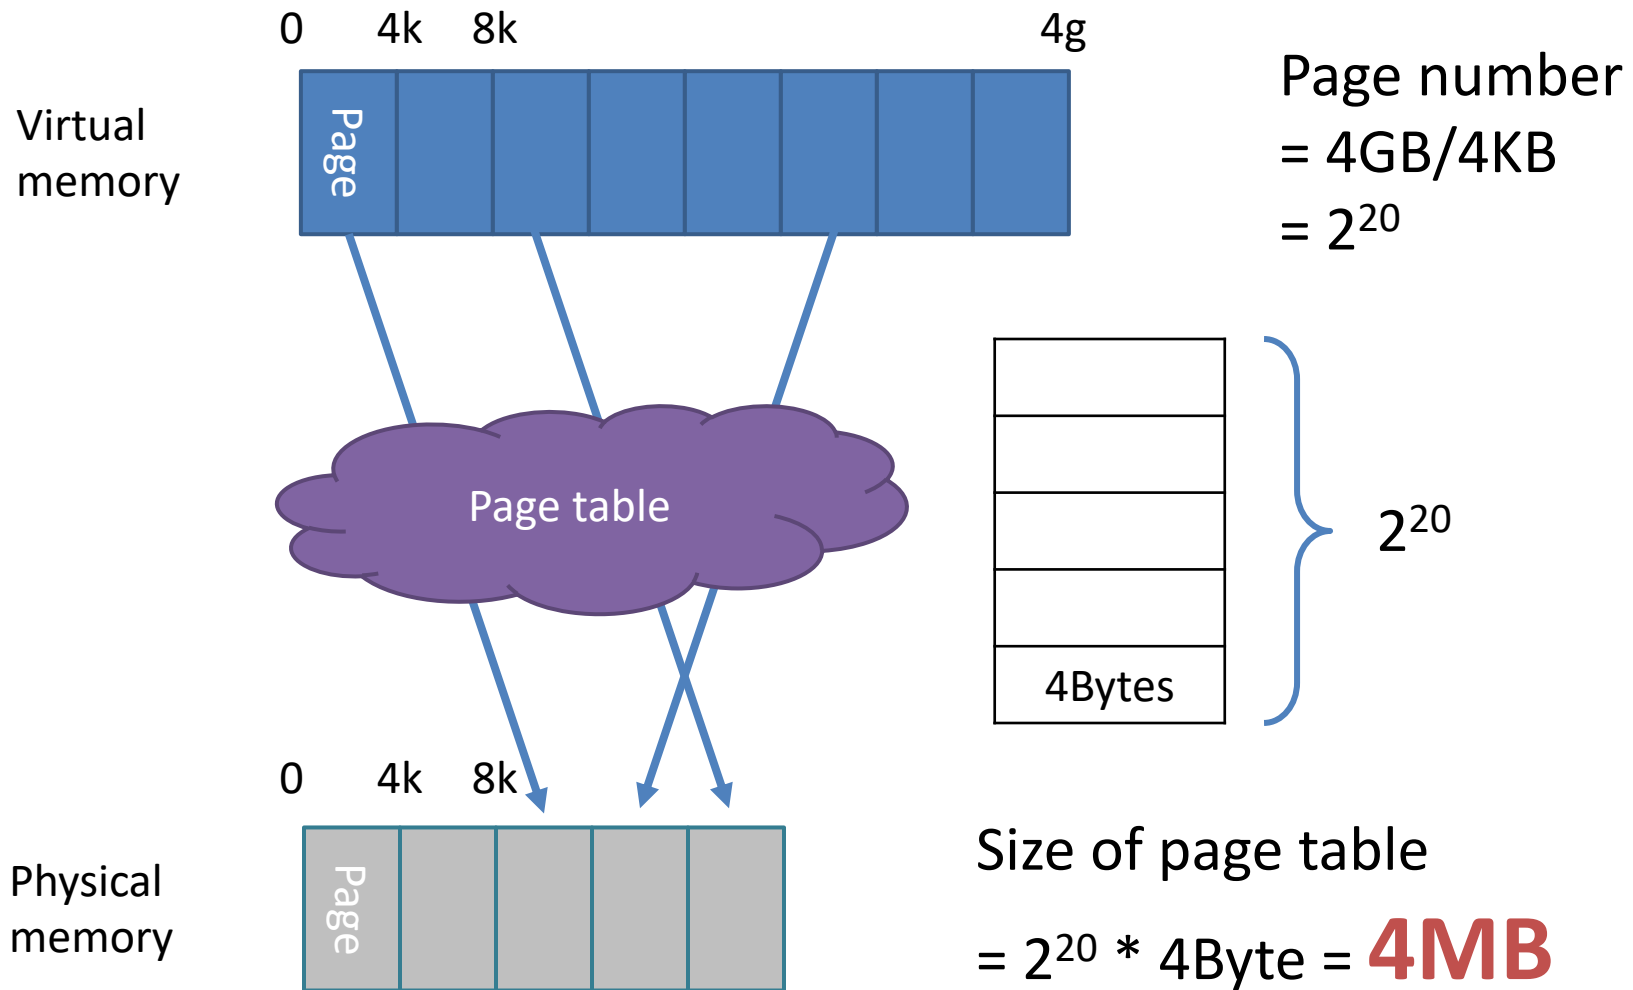
Page and page table



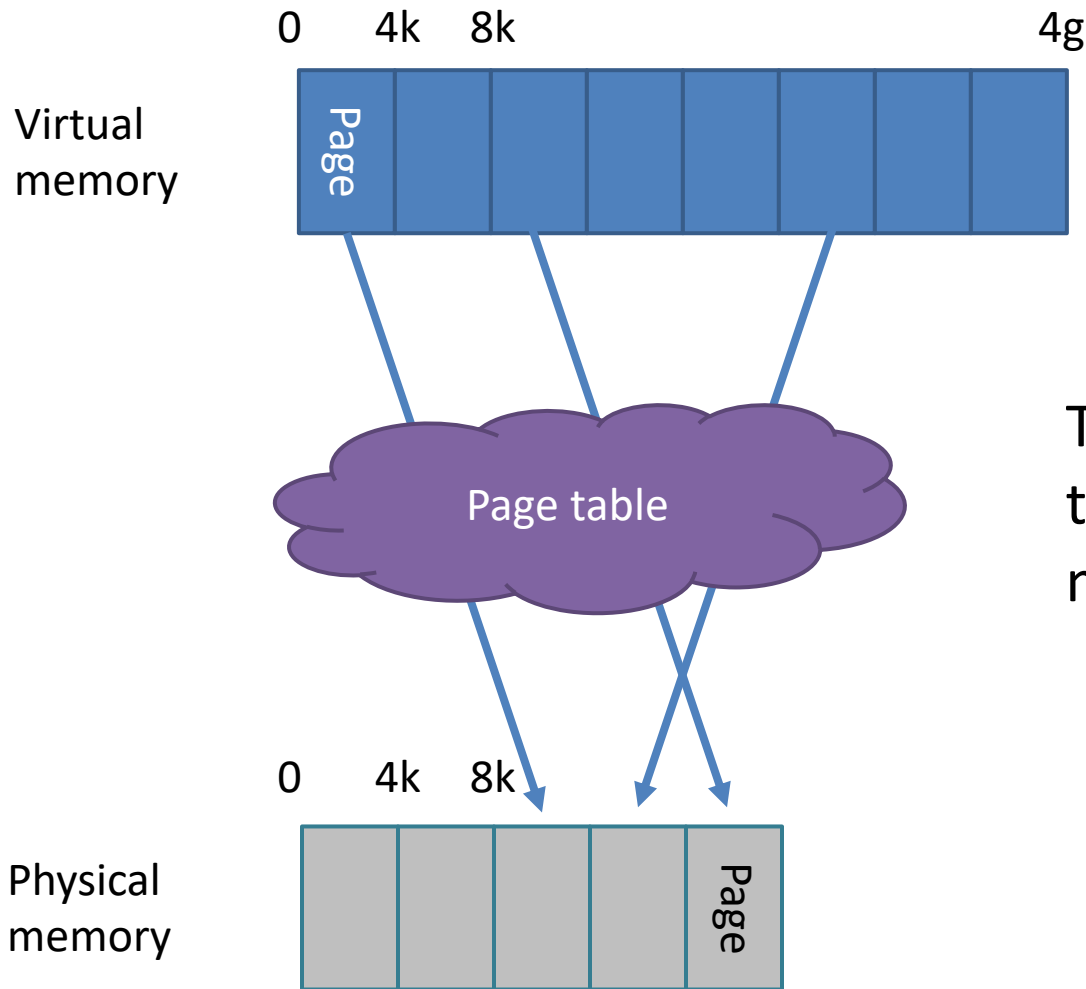
Page and page table



Suppose virtual address is 4GB and page size is 4KB. The page table item is 4 Byte. What would be the size of page table?



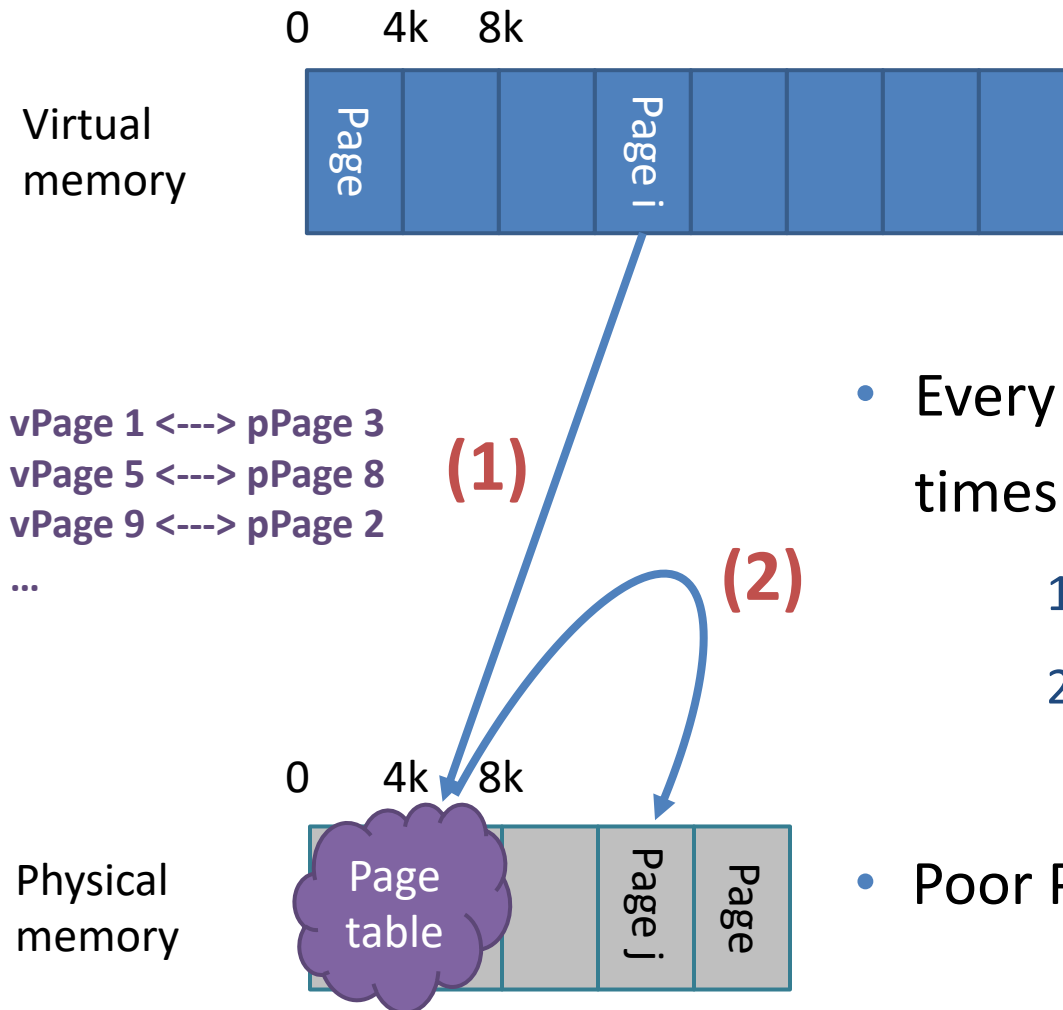
Where does the page table exist?



The table is located inside the **physical memory** and maintained by the **OS**



Page table lookup

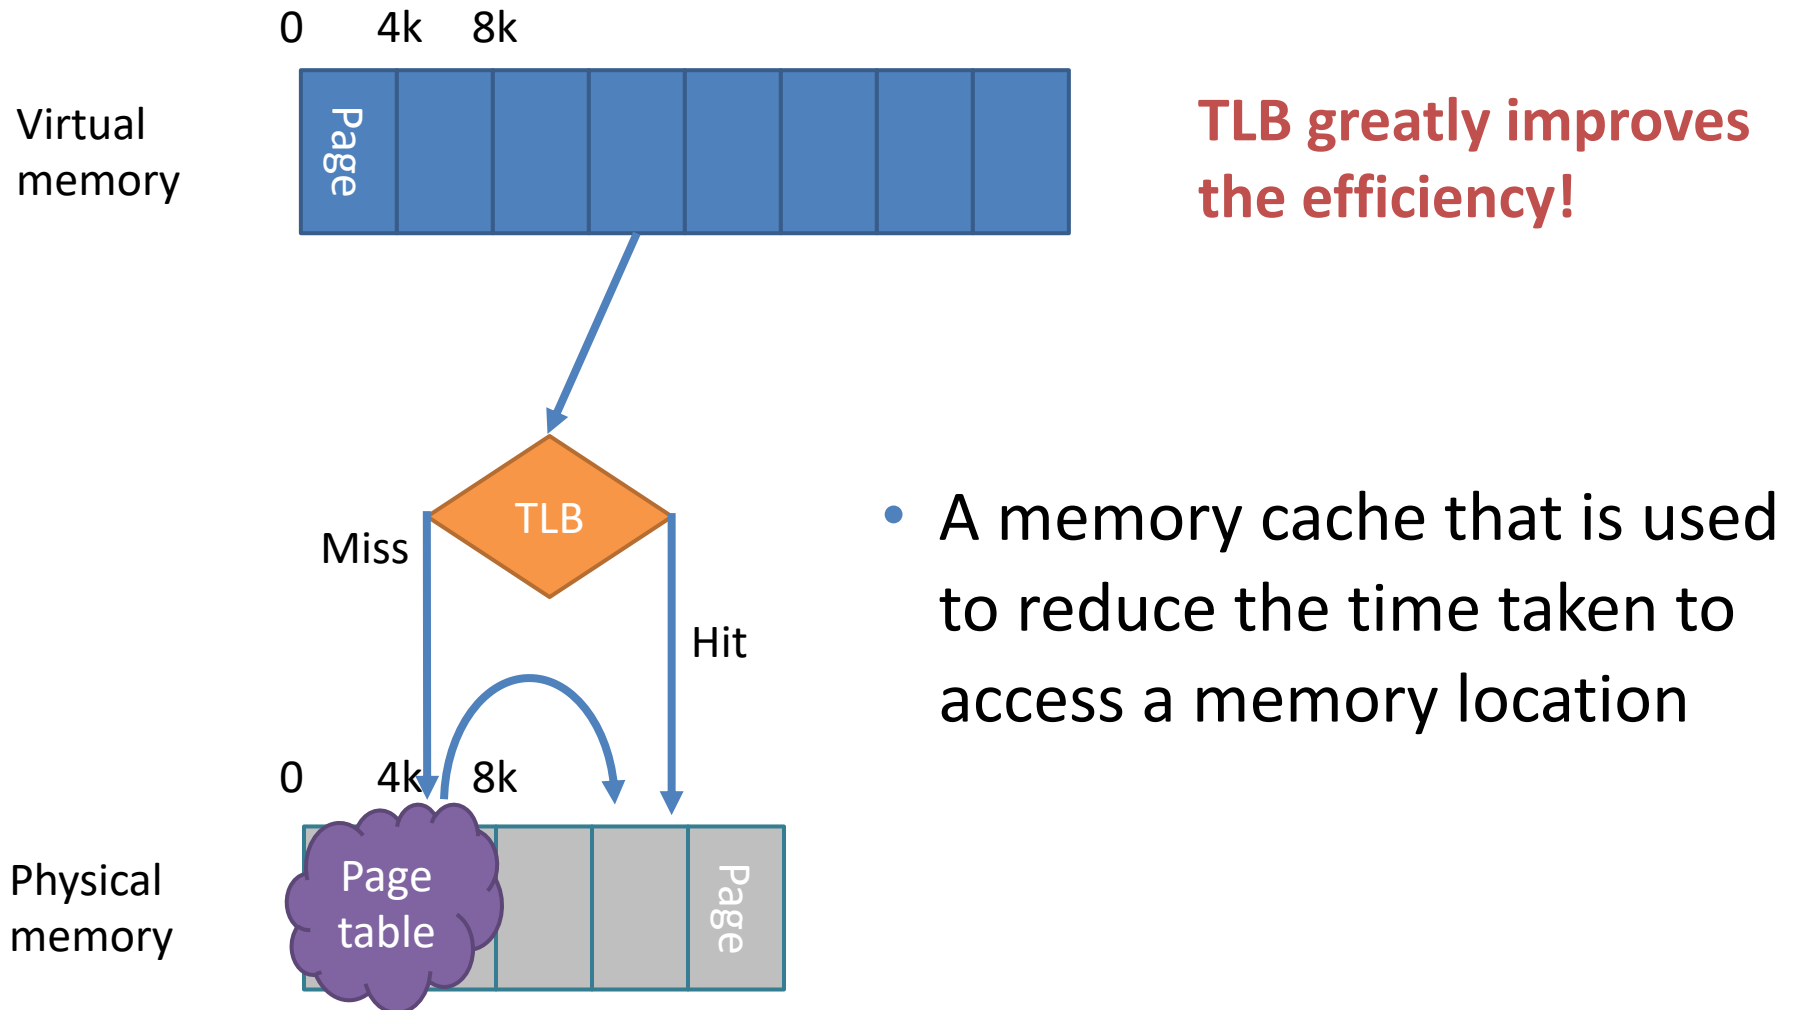


- Every visit to memory requires two times of physical memory access:
 - 1, access to page table
 - 2, visit the destination address

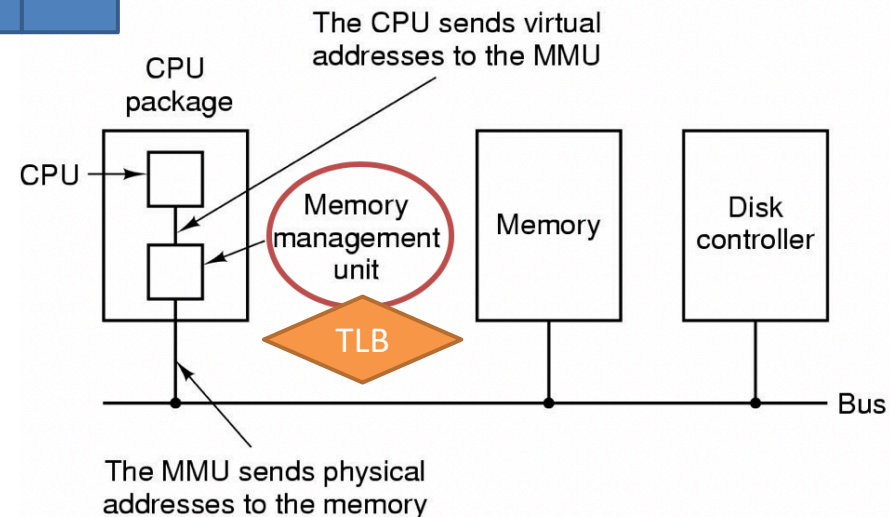
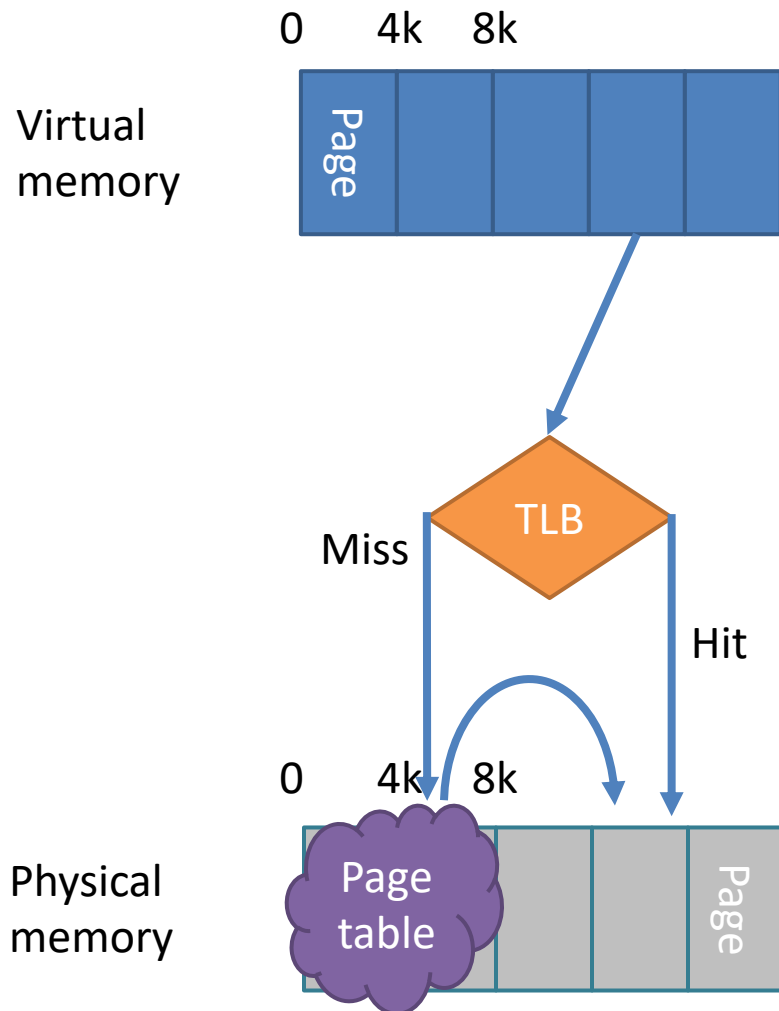
- Poor Performance



Translation Look-aside Buffers (TLB)



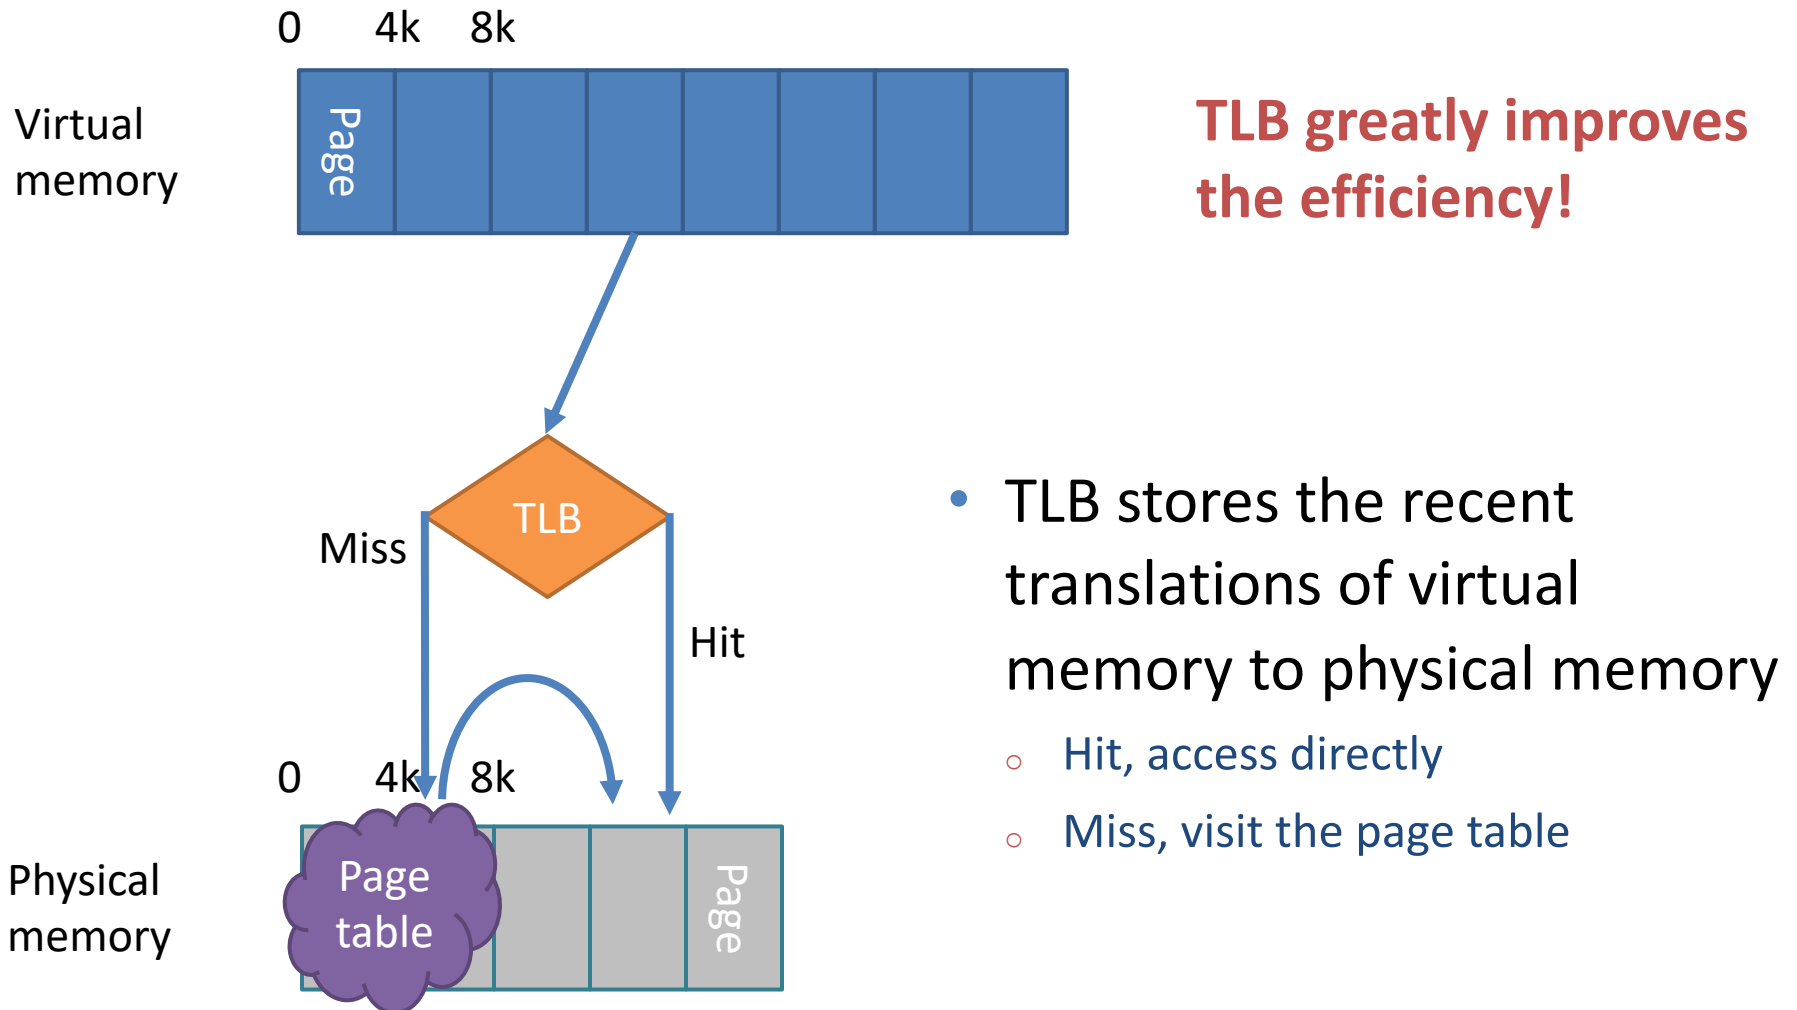
Translation Look-aside Buffers (TLB)



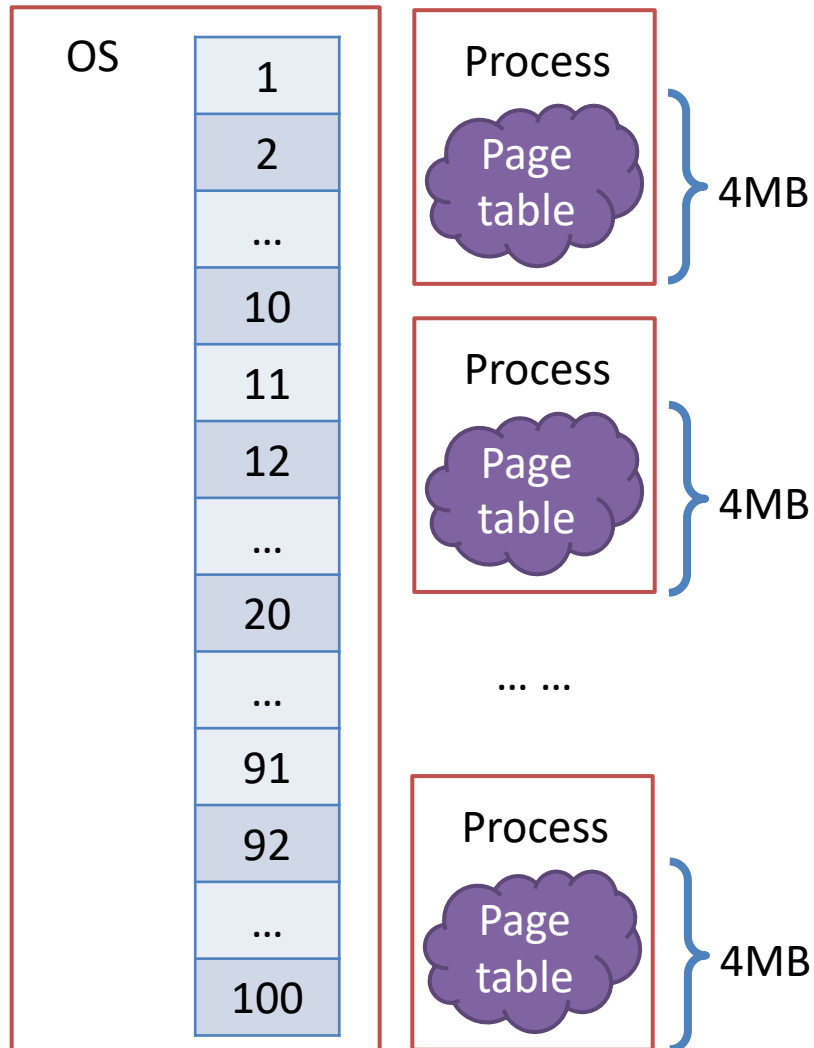
- A part of the chip's memory-management unit (MMU)



Translation Look-aside Buffers (TLB)

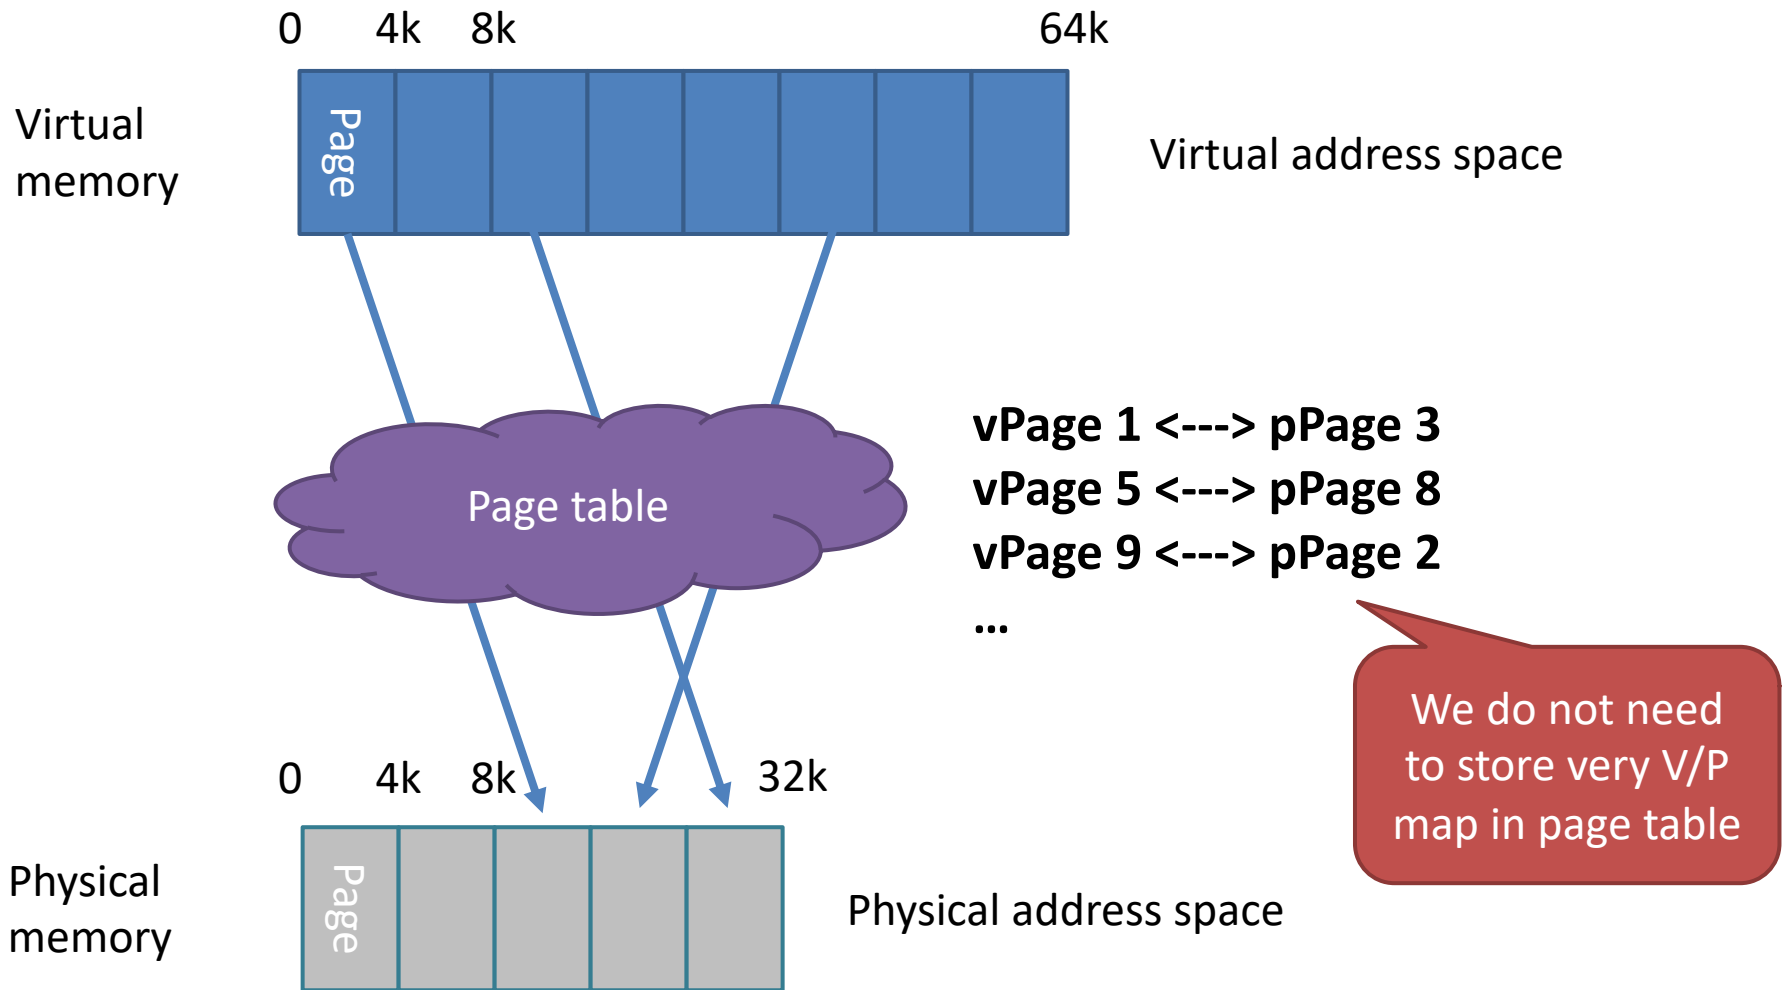


Page table problems



- Page table is a per-process data structure
- Large amount of memory could be used to store page table

Page table size optimization




Multi-level Page Tables

Single-level
page tables

Multi-level
page tables

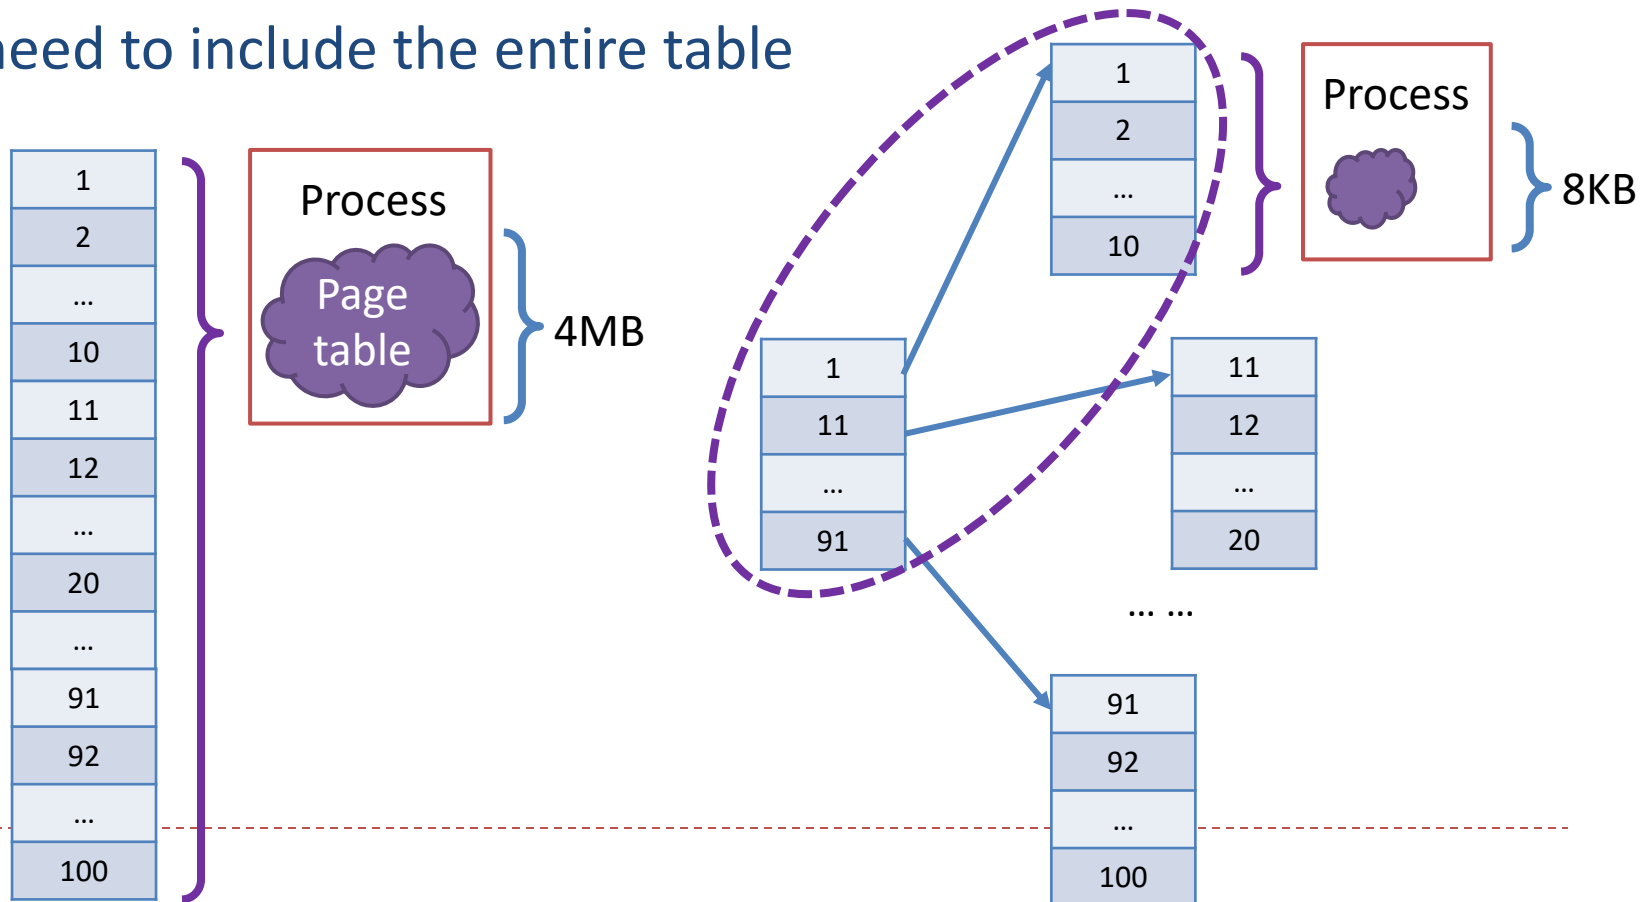
- Multi-level page tables can reduce the memory occupied by page tables and improve the memory efficiency

1
2
...
10
11
12
...
20
...
91
92
...
100




Multi-level Page Tables benefits

- Save memory space
 - 1) Far less page table to individual process (locality), not need to include the entire table



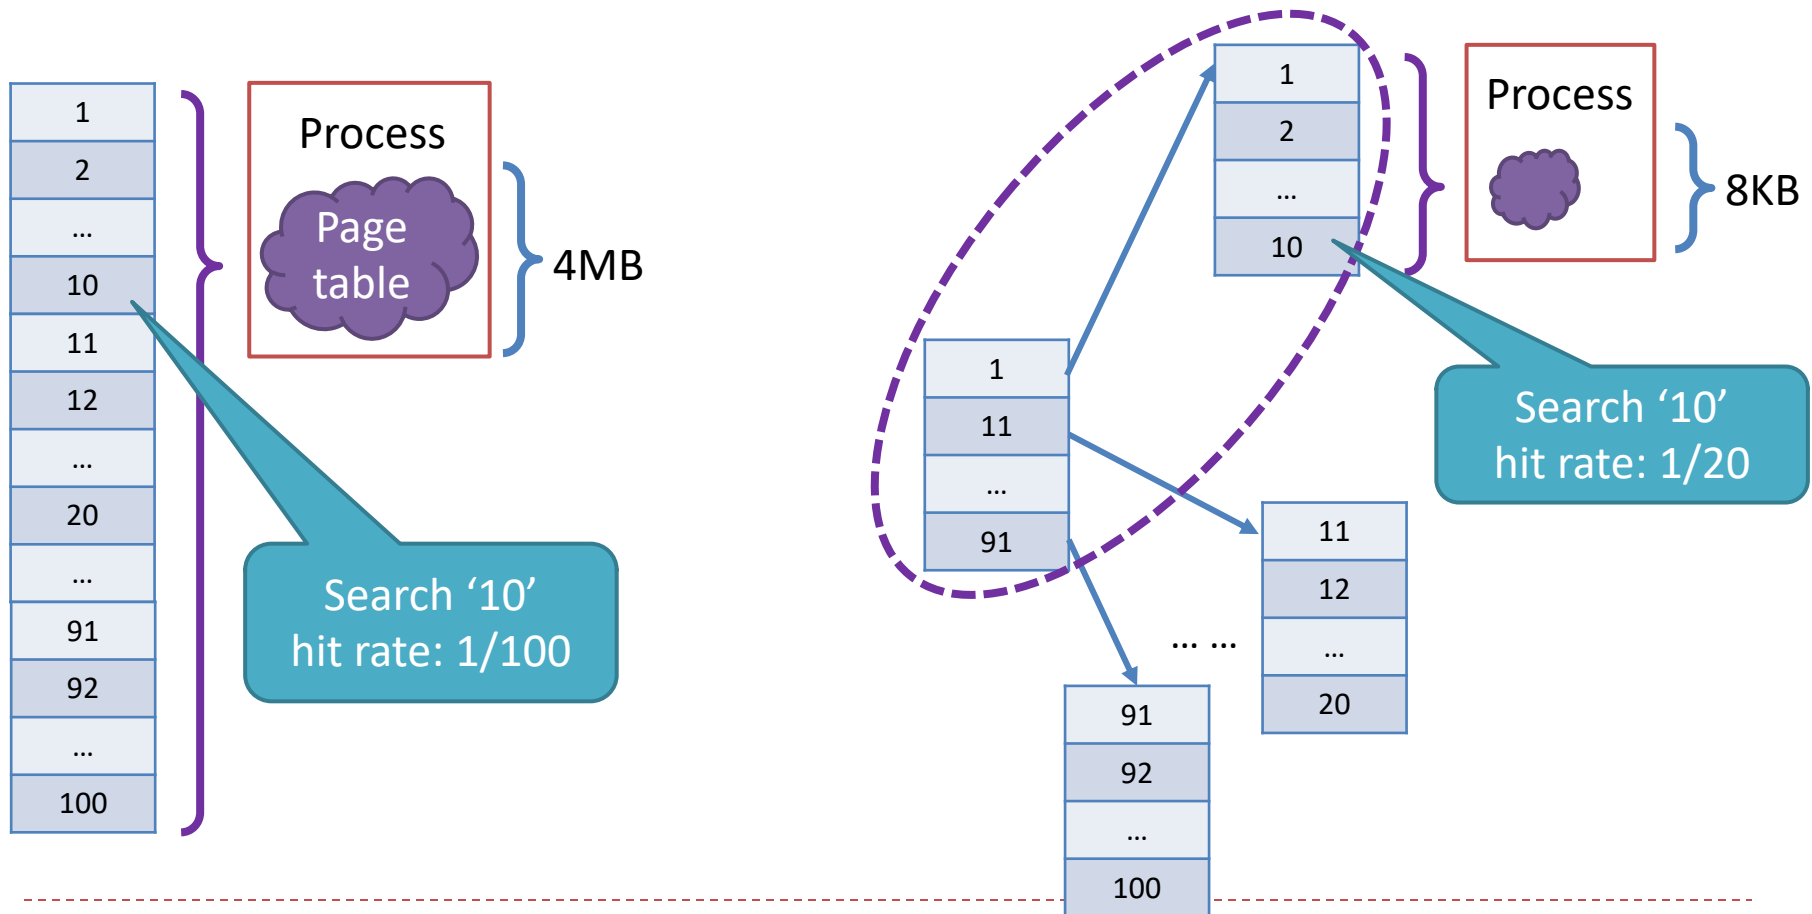
Multi-level Page Tables benefits

- Save memory space
 - 2) Second level page might not exist in memory, could be on the disk



Multi-level Page Tables benefits

- High access efficiency due to table index




Multi-level Page Tables drawbacks

- Drawbacks
 - Index space
 - Higher complexity with more levels

1
2
...
10
11
12
...
20
...
91
92
...
100

Total amount space
for single level page
table is **100**



Total amount space
for single level page
table is **110**

Multi-level Page Tables

- Benefits:
 - Save memory space
 - ▶ 1) Far less page table to individual process (locality), not need to include the entire table
 - ▶ 2) Second level page might not exist in memory, could be on the disk
 - High access efficiency due to table index
- Drawbacks:
 - Index space
 - Higher complexity with more levels



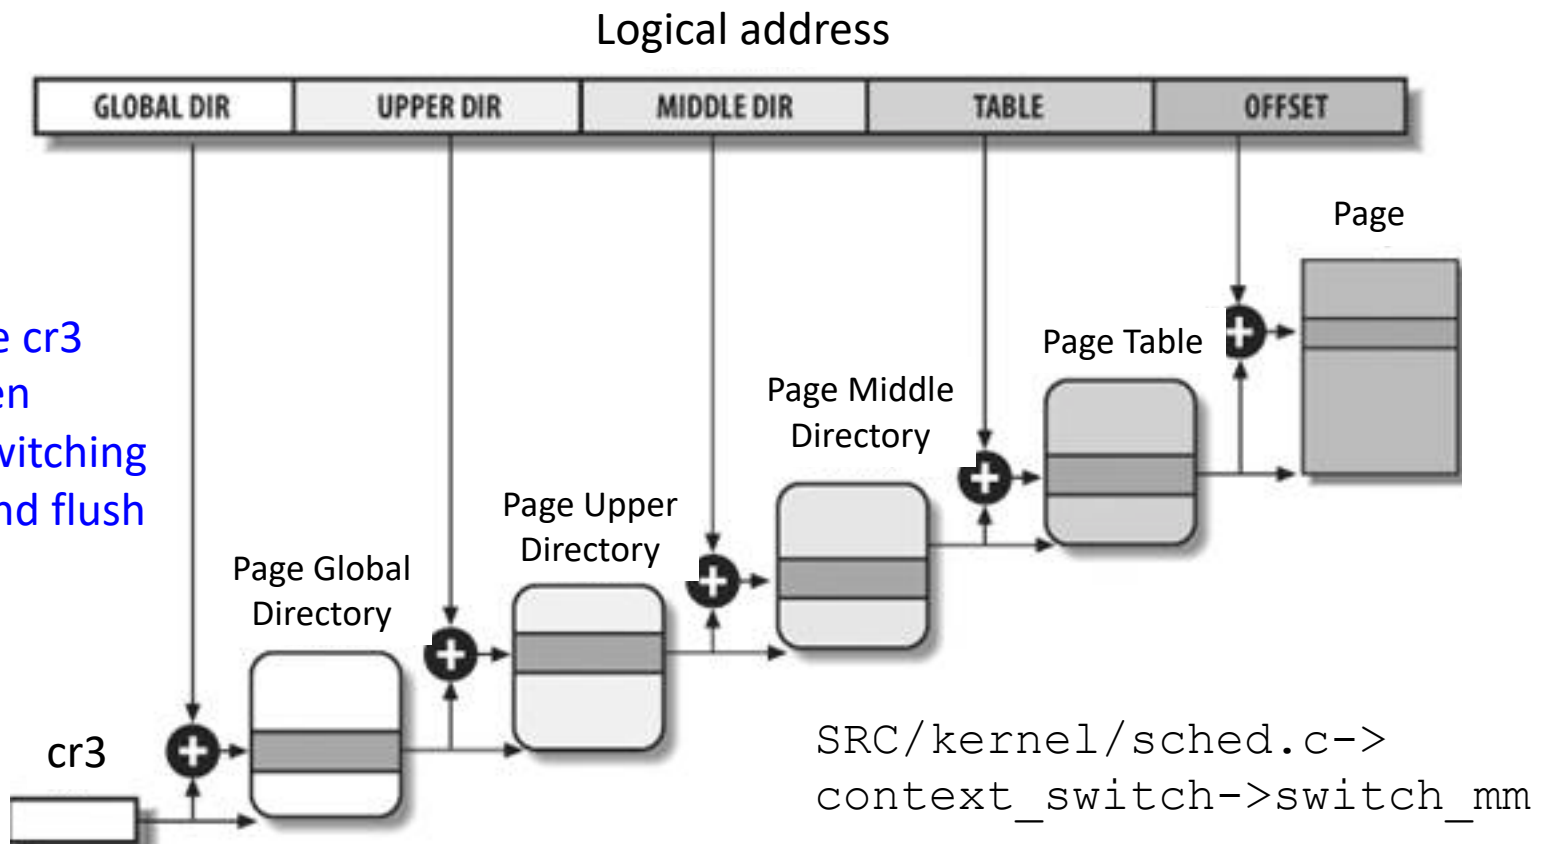
Multi-level Page Tables in Linux

https://elixir.bootlin.com/linux/latest/ident/pgdval_t

A common model for 32-bit (two-level, 4B pte) and 64-bit (four-level, 8B pte)

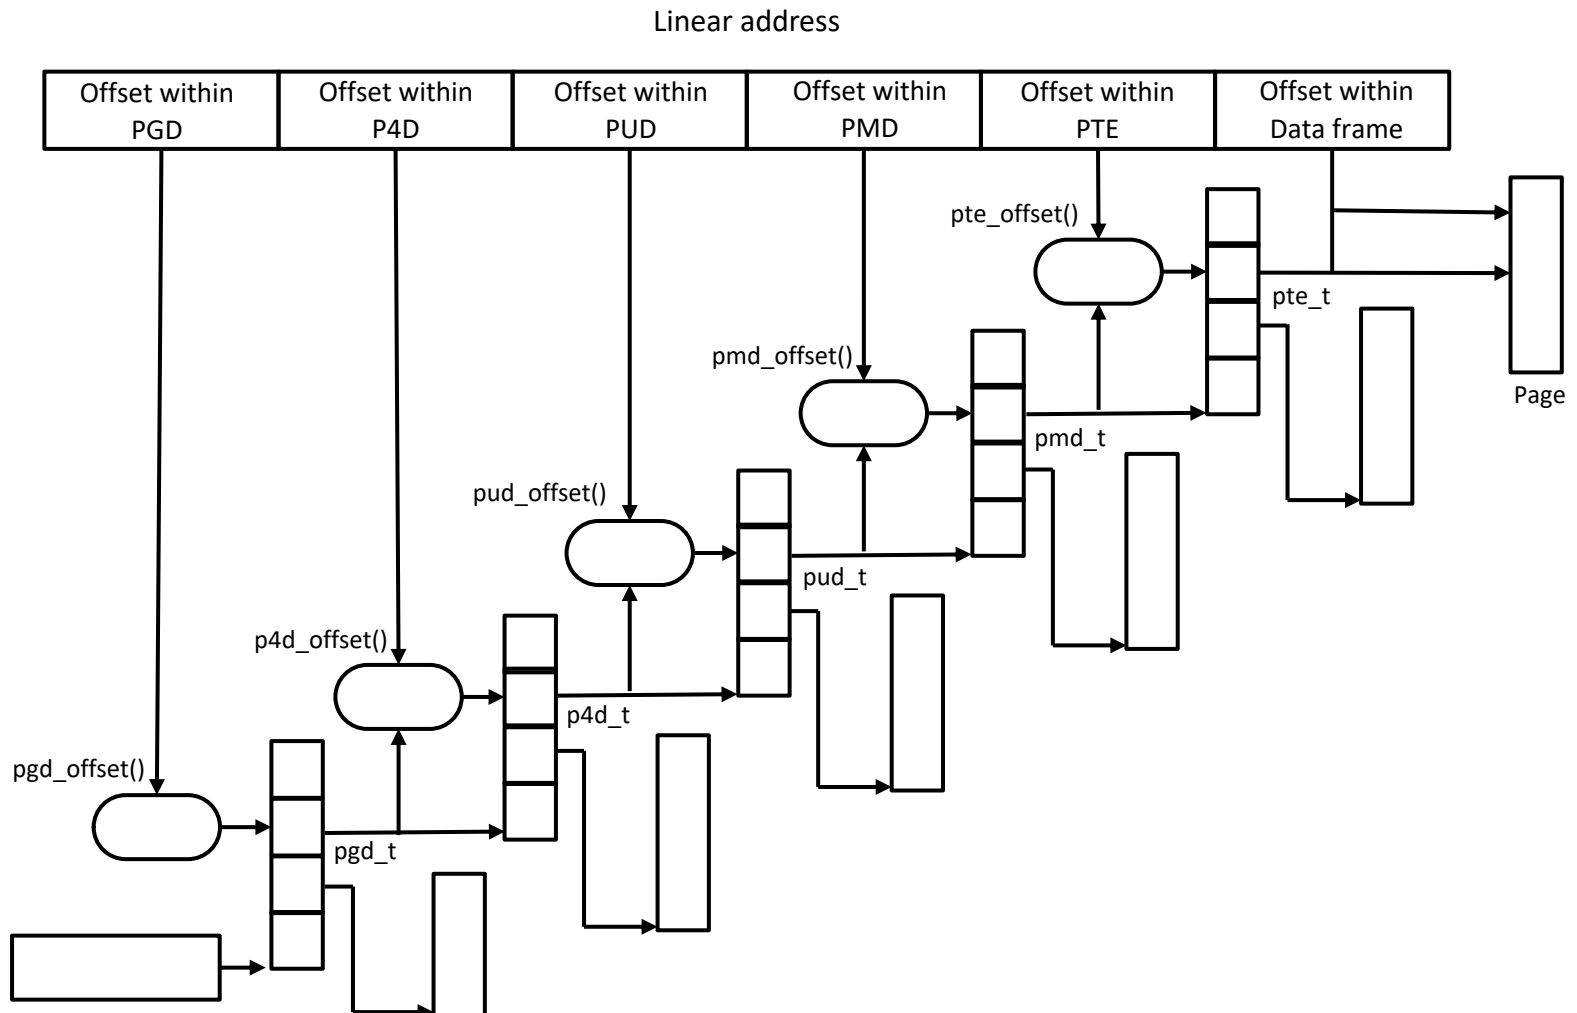
`SRC/include/linux/sched.h->task_struct->mm->pgd`

Switch the cr3 value when context switching threads and flush the TLB



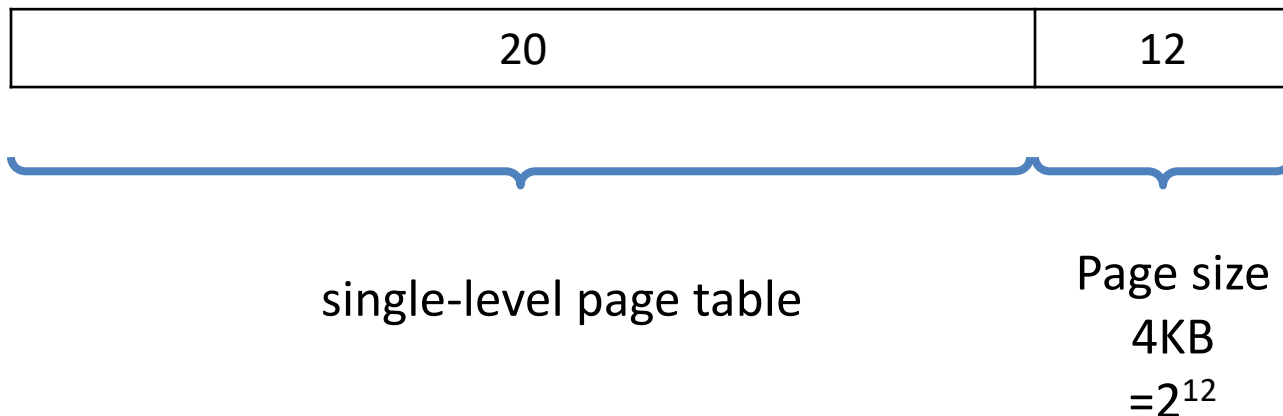
Multi-level Page Tables in Linux

https://elixir.bootlin.com/linux/latest/ident/pgdval_t




Multi-level Page Tables Example

- In a 32-bit system, a single-level page table with 4KB pages.



Multi-level Page Tables Example

- In a 32-bit system, a single-level page table with 4KB pages. Suppose PT1=1, offset=10, what is its virtual address?



single-level page table

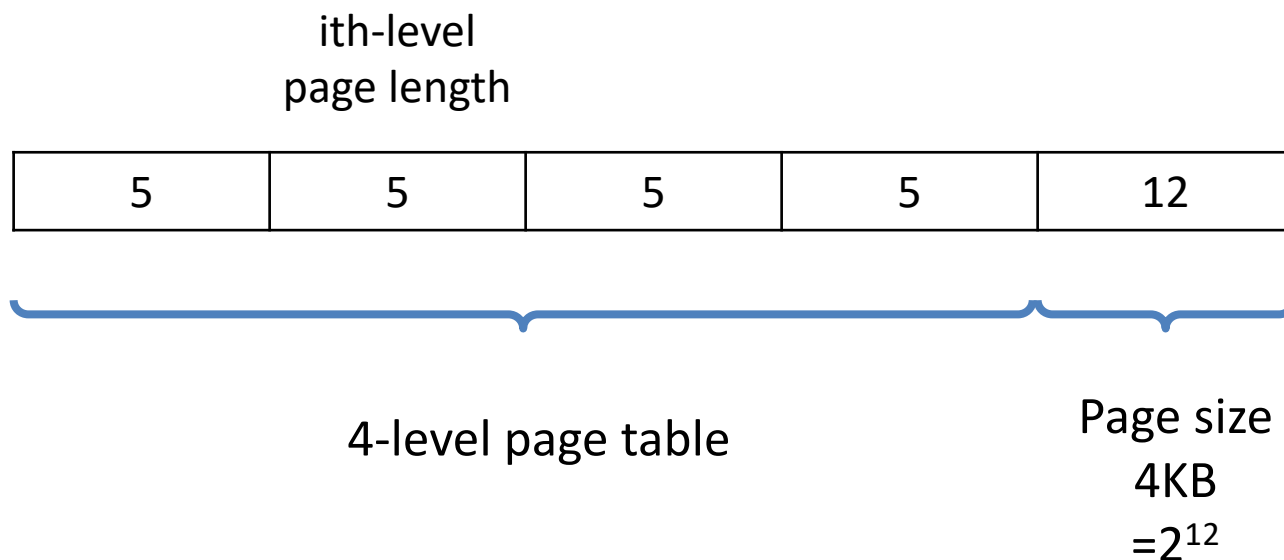
Page size
4KB
 $=2^{12}$

The virtual address = $1 * 2^{20} + 10 = 1048586$



Multi-level Page Tables Example

- In a 32-bit system, a four-level page table with 4KB pages. Each level is equal



Multi-level Page Tables Example

- In a 32-bit system, a four-level page table with 4KB pages. Each level is equal. Suppose PT1=1, PT2=2, PT3=3, PT4=4, offset=10, what is the virtual address?

ith-level
page length

5	5	5	5	12
---	---	---	---	----

The virtual address = $1 * 2^{27} + 2 * 2^{22} + 3 * 2^{17} + 4 * 2^{12} + 10 = 143015946$



Conclusion

- Memory management overview
 - Memory abstraction and address spaces
 - Physical address and virtual address
 - Physical memory and virtual memory
- Memory management
 - Translation look-aside buffer
 - Page table
 - Multi-level page table

