

# **CS 3502**

# **Operating Systems**

## **Input/Output**

**Kun Suo**

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

# Outline

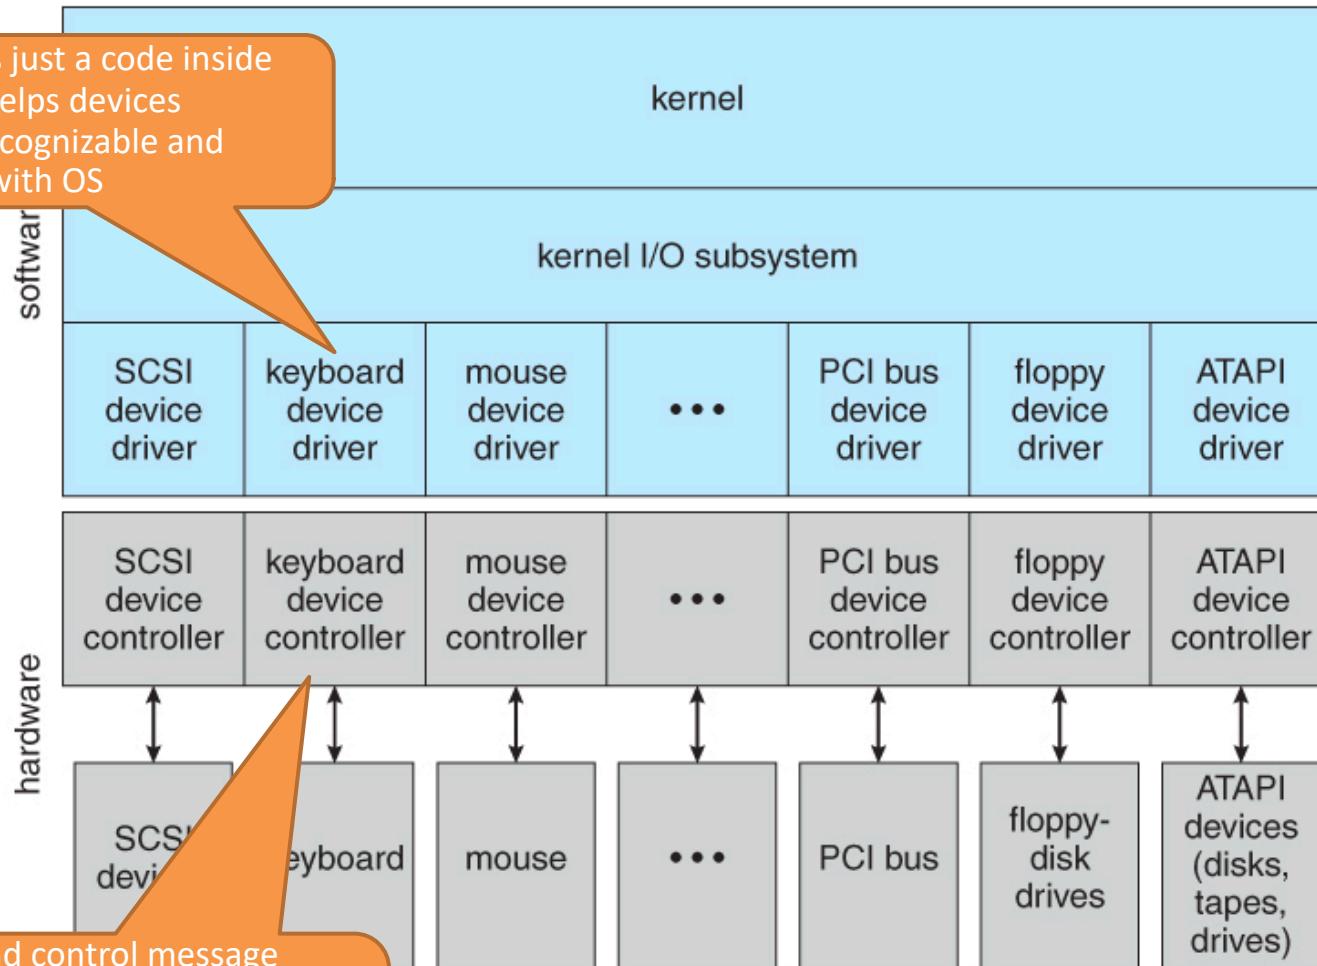
---

- I/O System
- Characteristics and Category of I/O Devices
- I/O Device Component
- How CPU works with I/O devices
- I/O Performance



# I/O System Structure

Device driver is just a code inside the OS which helps devices compatible , recognizable and communicate with OS



Receive data and control message from the OS, and send back the data and status message to the OS  
Convert digital signal into analog signal for the mechanical parts or vice versa

# I/O System Structure

---

- Device Drivers : low level **software** written by hardware manufacturers which are loaded as kernel modules and provide the kernel with the knowledge on how to control the devices. It acts bridge between **kernel** and **device**.
- Device Controllers: A device controller (**hardware**) understands "software" input. It translates software input into signal a hardware device understands. It acts bridge between **OS** and **hardware mechanical parts**.

# lsmod: list all drivers in the OS

```
ksuo@ksuo-VirtualBox ~> lsmod
Module           Size  Used by
Module
snd_intel8x0    45056  2
snd_ac97_codec  135168  1 snd_intel8x0
ac97_bus        16384  1 snd_ac97_codec
snd_pcm         102400  2 snd_intel8x0,snd_ac97_codec
crct10dif_pclmul 16384  1
crc32_pclmul   16384  0
snd_seq_midi    20480  0
ghash_clmulni_intel 16384  0
snd_seq_midi_event 16384  1 snd_seq_midi
snd_rawmidi     36864  1 snd_seq_midi
snd_seq          69632  2 snd_seq_midi,snd_seq_midi_event
aesni_intel     372736  0
aes_x86_64      20480  1 aesni_intel
crypto_simd     16384  1 aesni_intel
vmwgfx          290816  2
cryptd          24576  3 crypto_simd,ghash_clmulni_intel,aesni_intel
glue_helper     16384  1 aesni_intel
snd_seq_device   16384  3 snd_seq,snd_seq_midi,snd_rawmidi
intel_rapl_perf 16384  0
snd_timer       36864  2 snd_seq,snd_pcm
joydev          28672  0
ttm              102400  1 vmwgfx
drm_kms_helper  180224  1 vmwgfx
input_leds       16384  0
serio_raw        20480  0
drm              479232  5 vmwgfx,drm_kms_helper,ttm
snd              86016  11 snd_seq,snd_seq_device,snd_intel8x0,snd_timer,snd_ac97_codec,s
fb_sys_fops     16384  1 drm_kms_helper
syscopyarea     16384  1 drm_kms_helper
sysfillrect     16384  1 drm_kms_helper
mac_hid          16384  0
sysimgblt        16384  1 drm_kms_helper
soundcore        16384  1 snd
vboxguest        335872  0
sch fq_codel    20480  2
parport_pc       36864  0
ppdev            24576  0
hid_generic      16384  0
usbhid           53248  0
hid              126976  2 usbhid,hid_generic
ahci             40960  1
psmouse          151552  0
libahci          32768  1 ahci
i2c_piix4       28672  0
e1000            139264  0
pata_acpi        16384  0
video             49152  0

```

Keyboard driver

Mouse driver

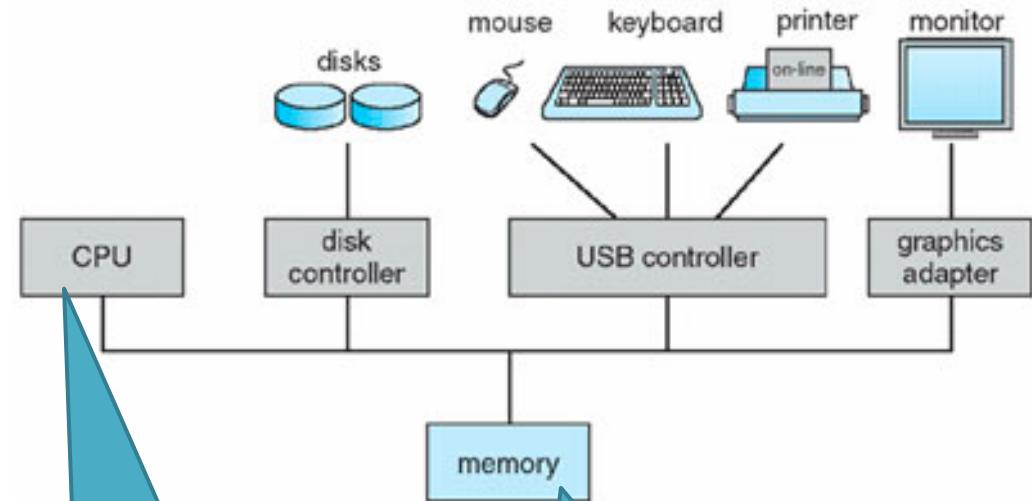
OS kernel:  
Scheduling, interrupt,  
memory, refresh display, ...

Keyboard driver:  
key code → program logic

Keyboard controller:  
Signal → key code

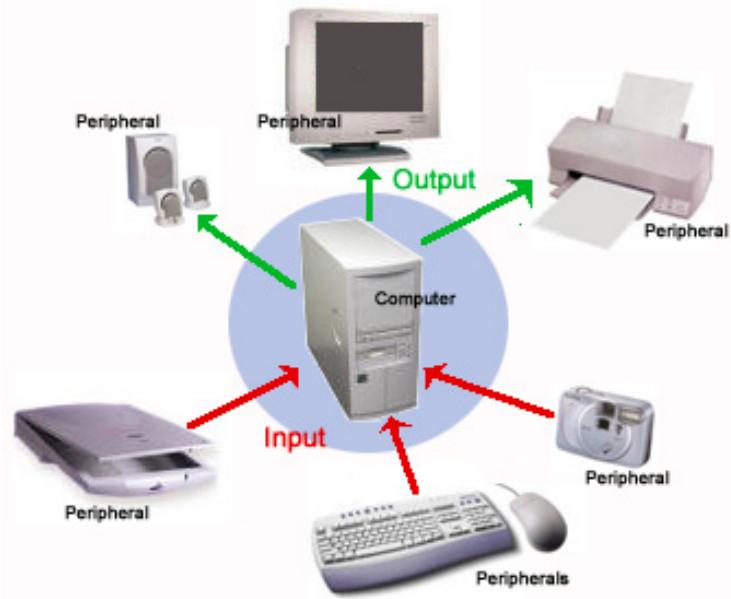


# I/O System Structure

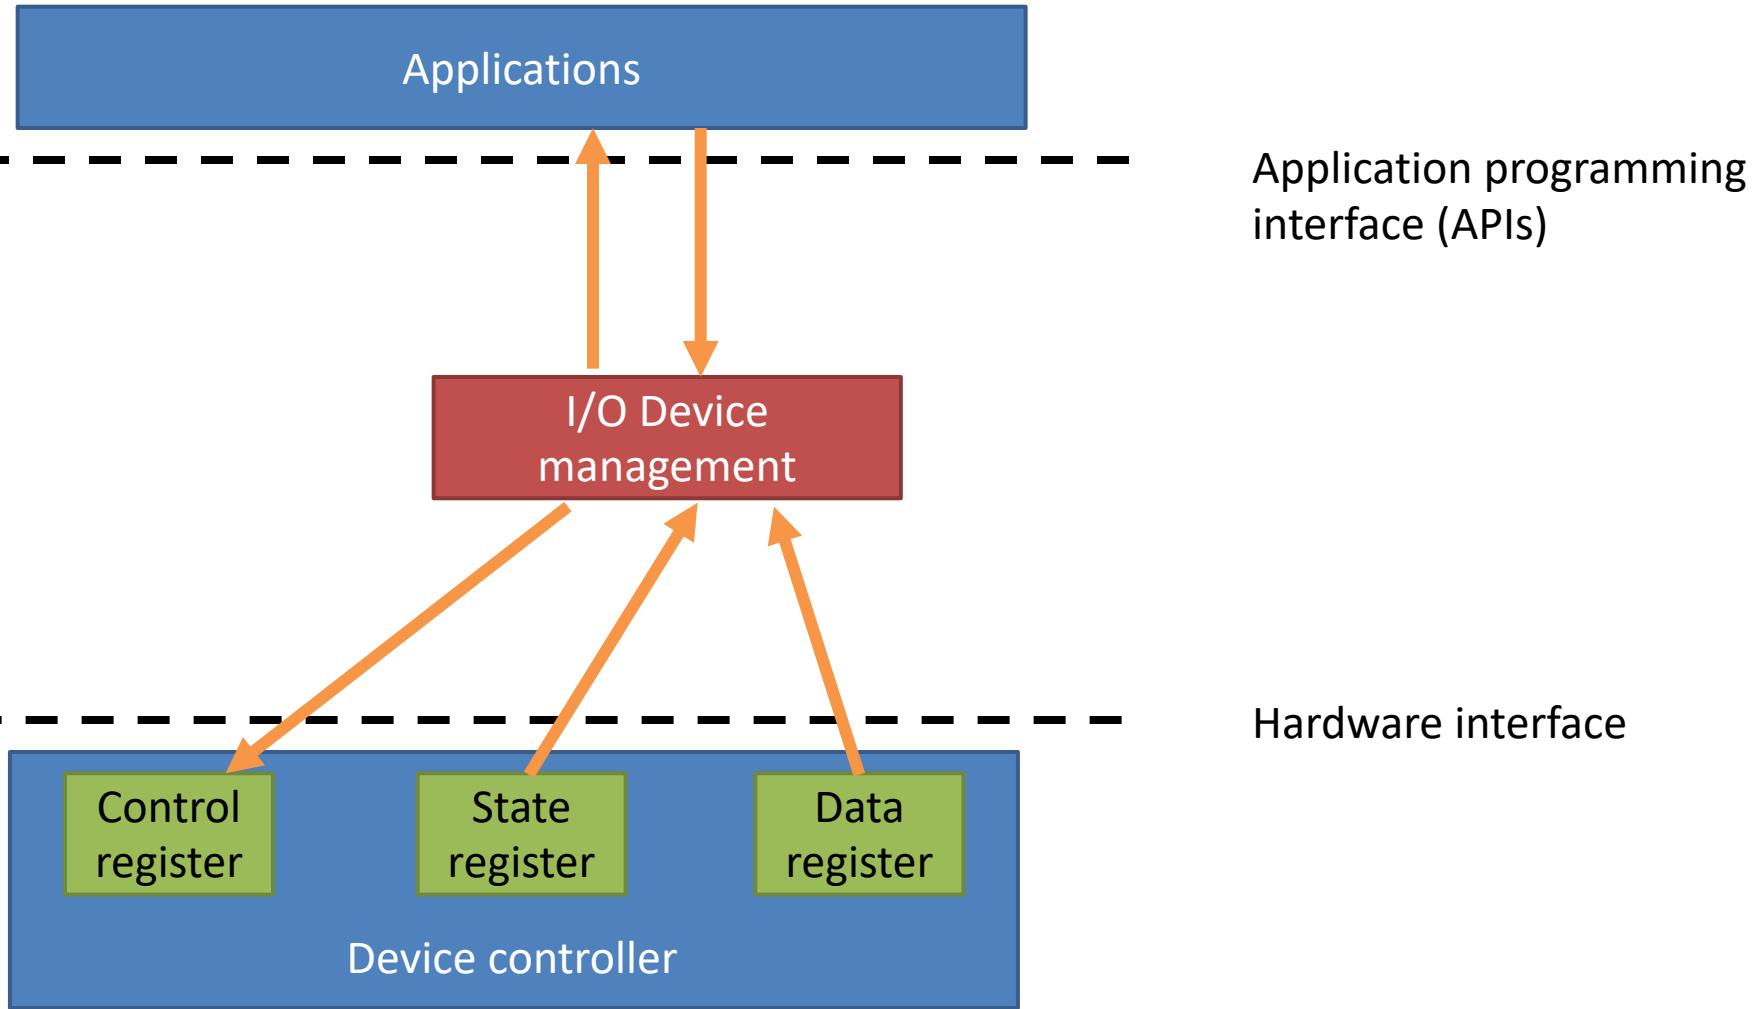


CPU is used for control and management

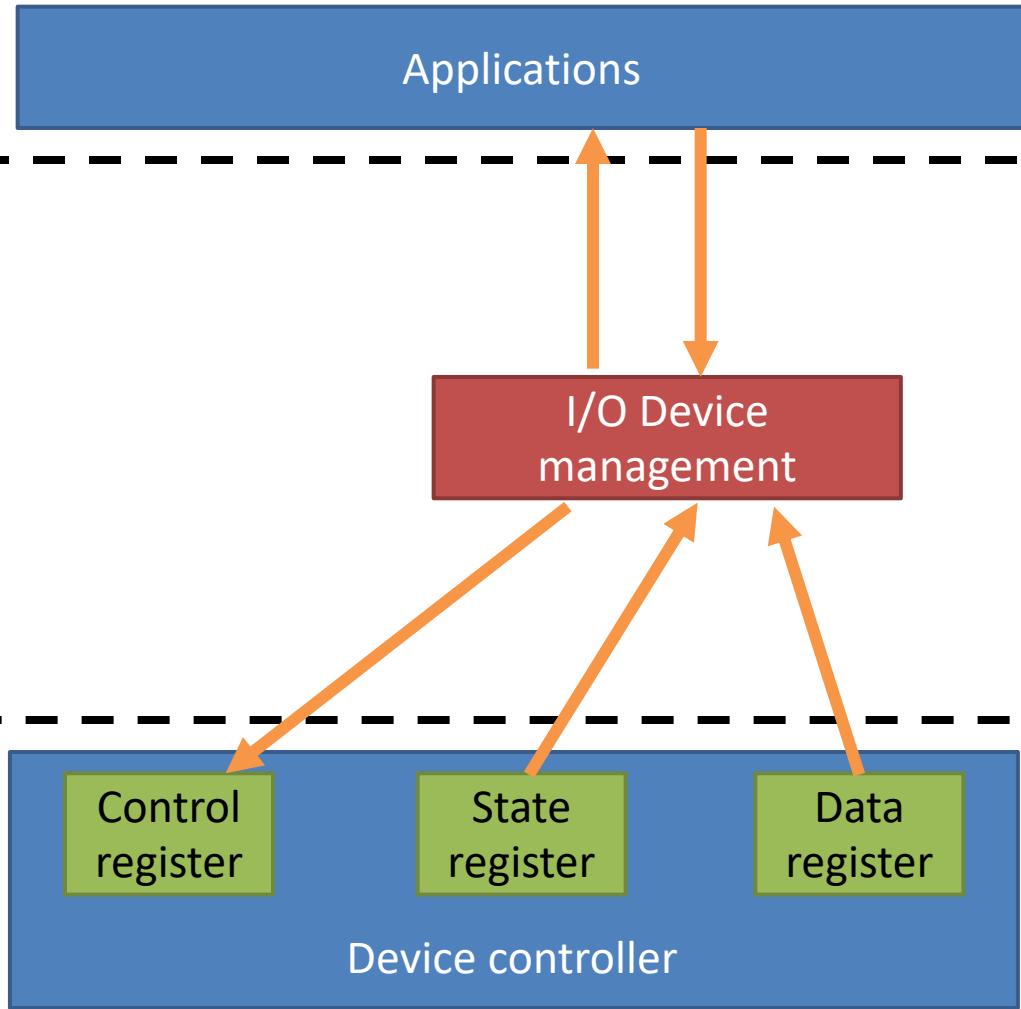
Memory is used for temporary data storage



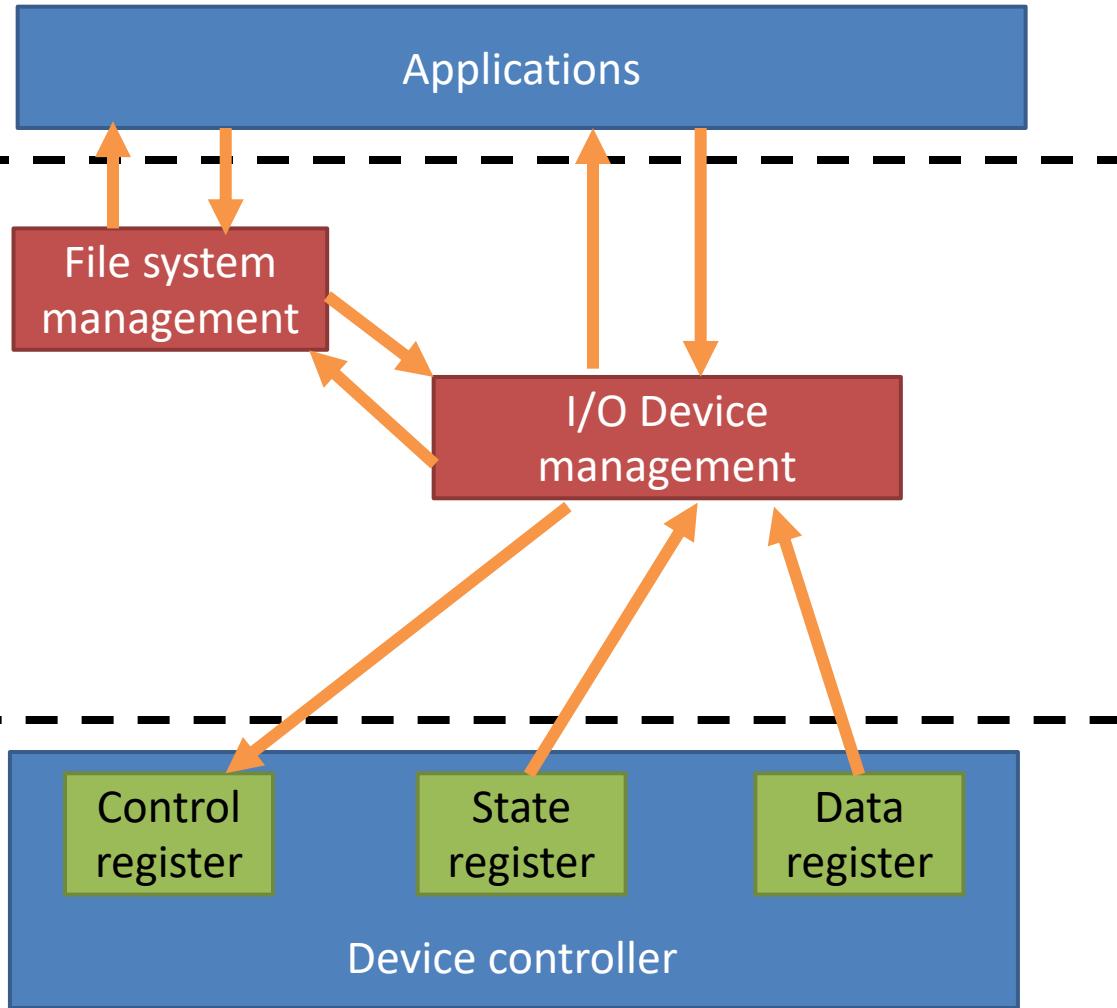
# I/O Management



# I/O Management Example



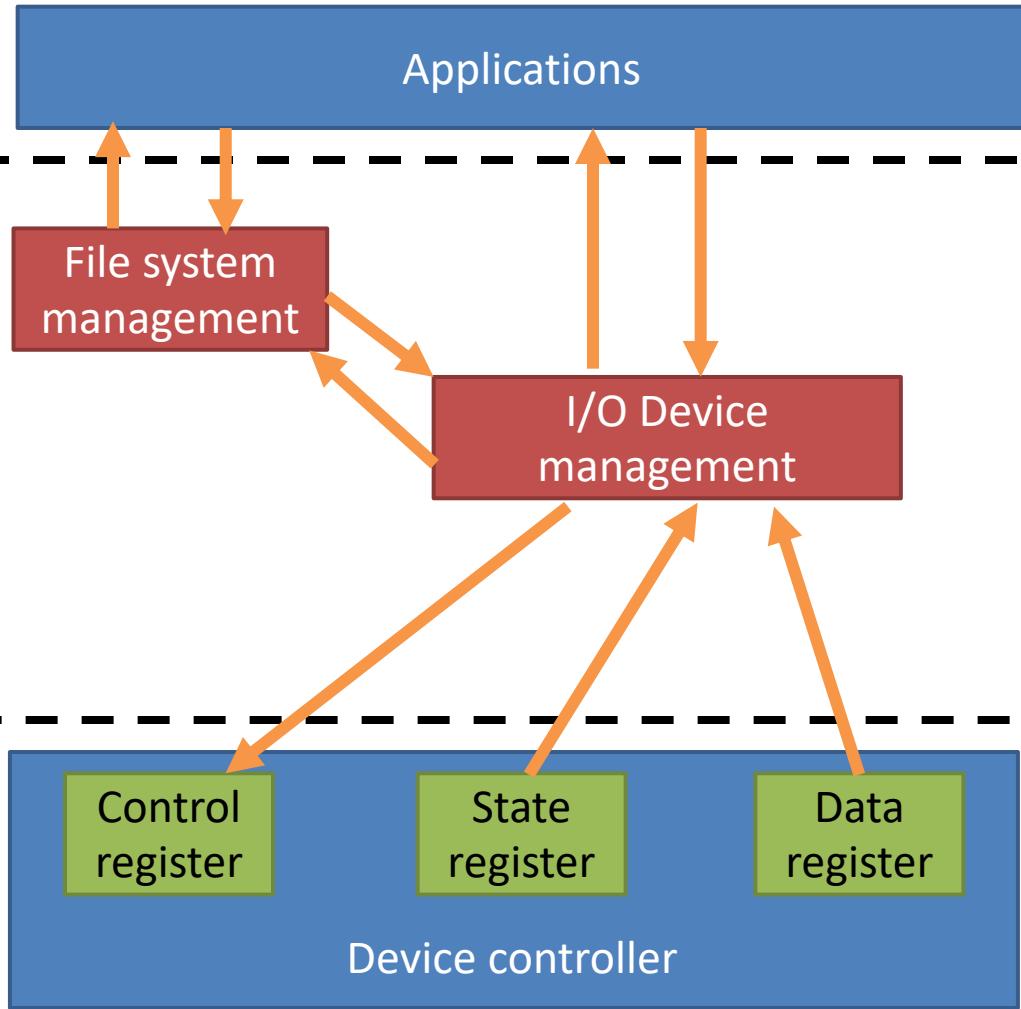
# I/O Management



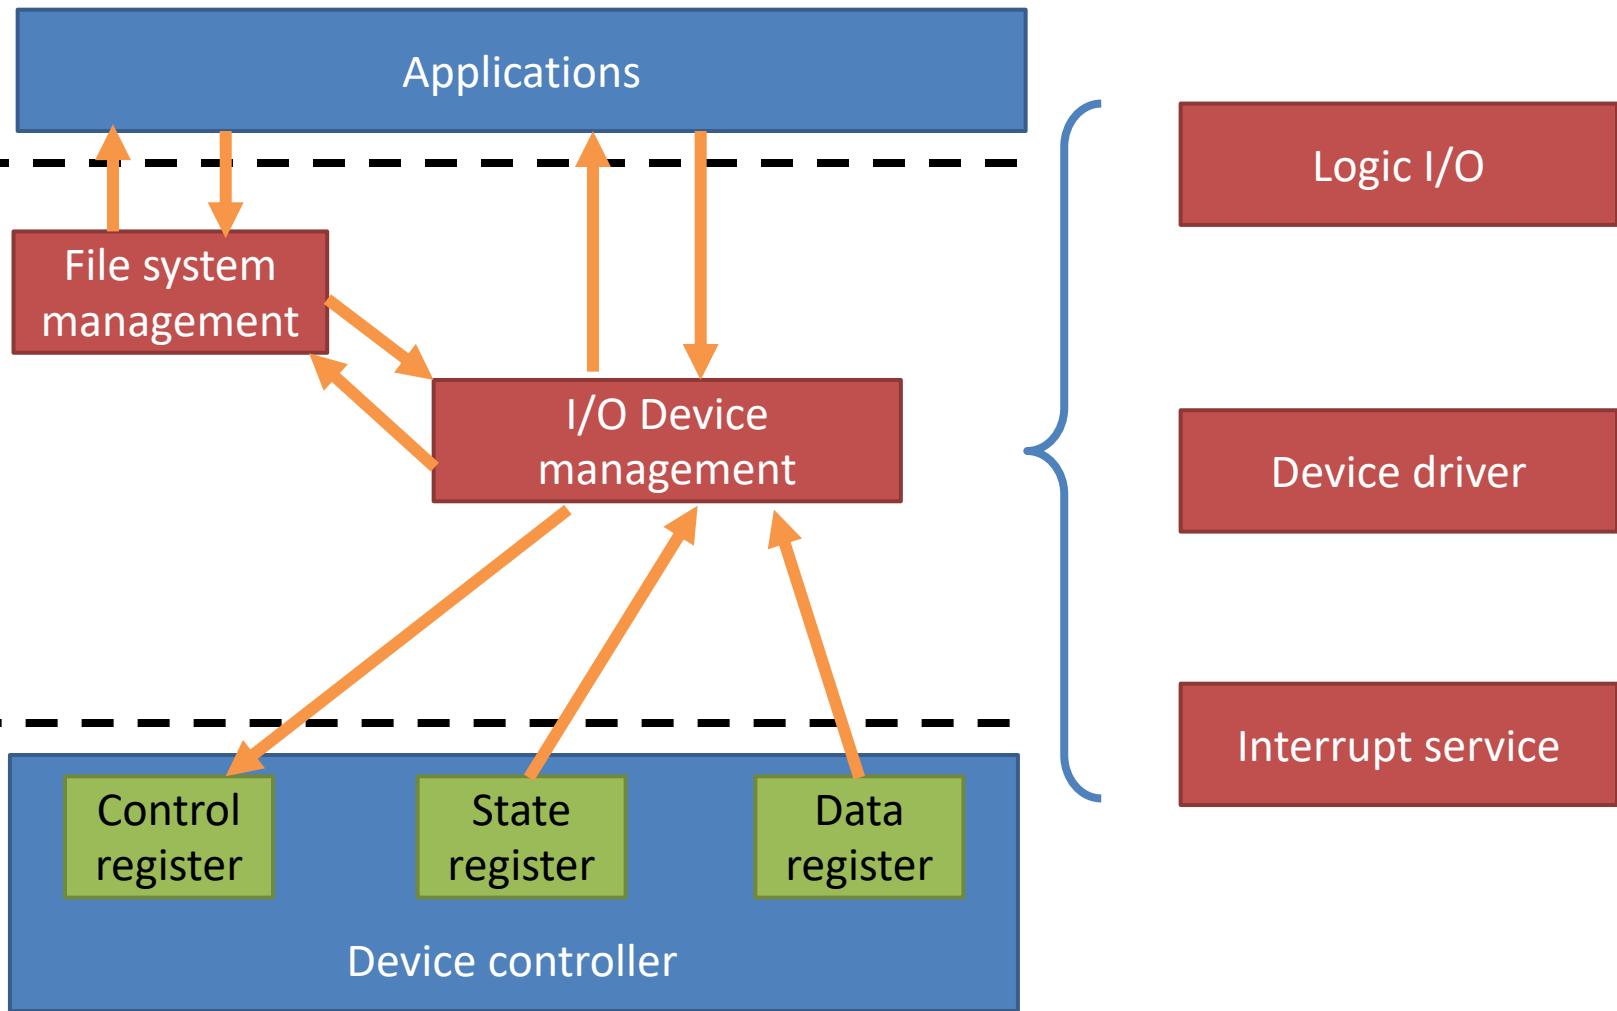
Application programming  
interface (APIs)

Hardware interface

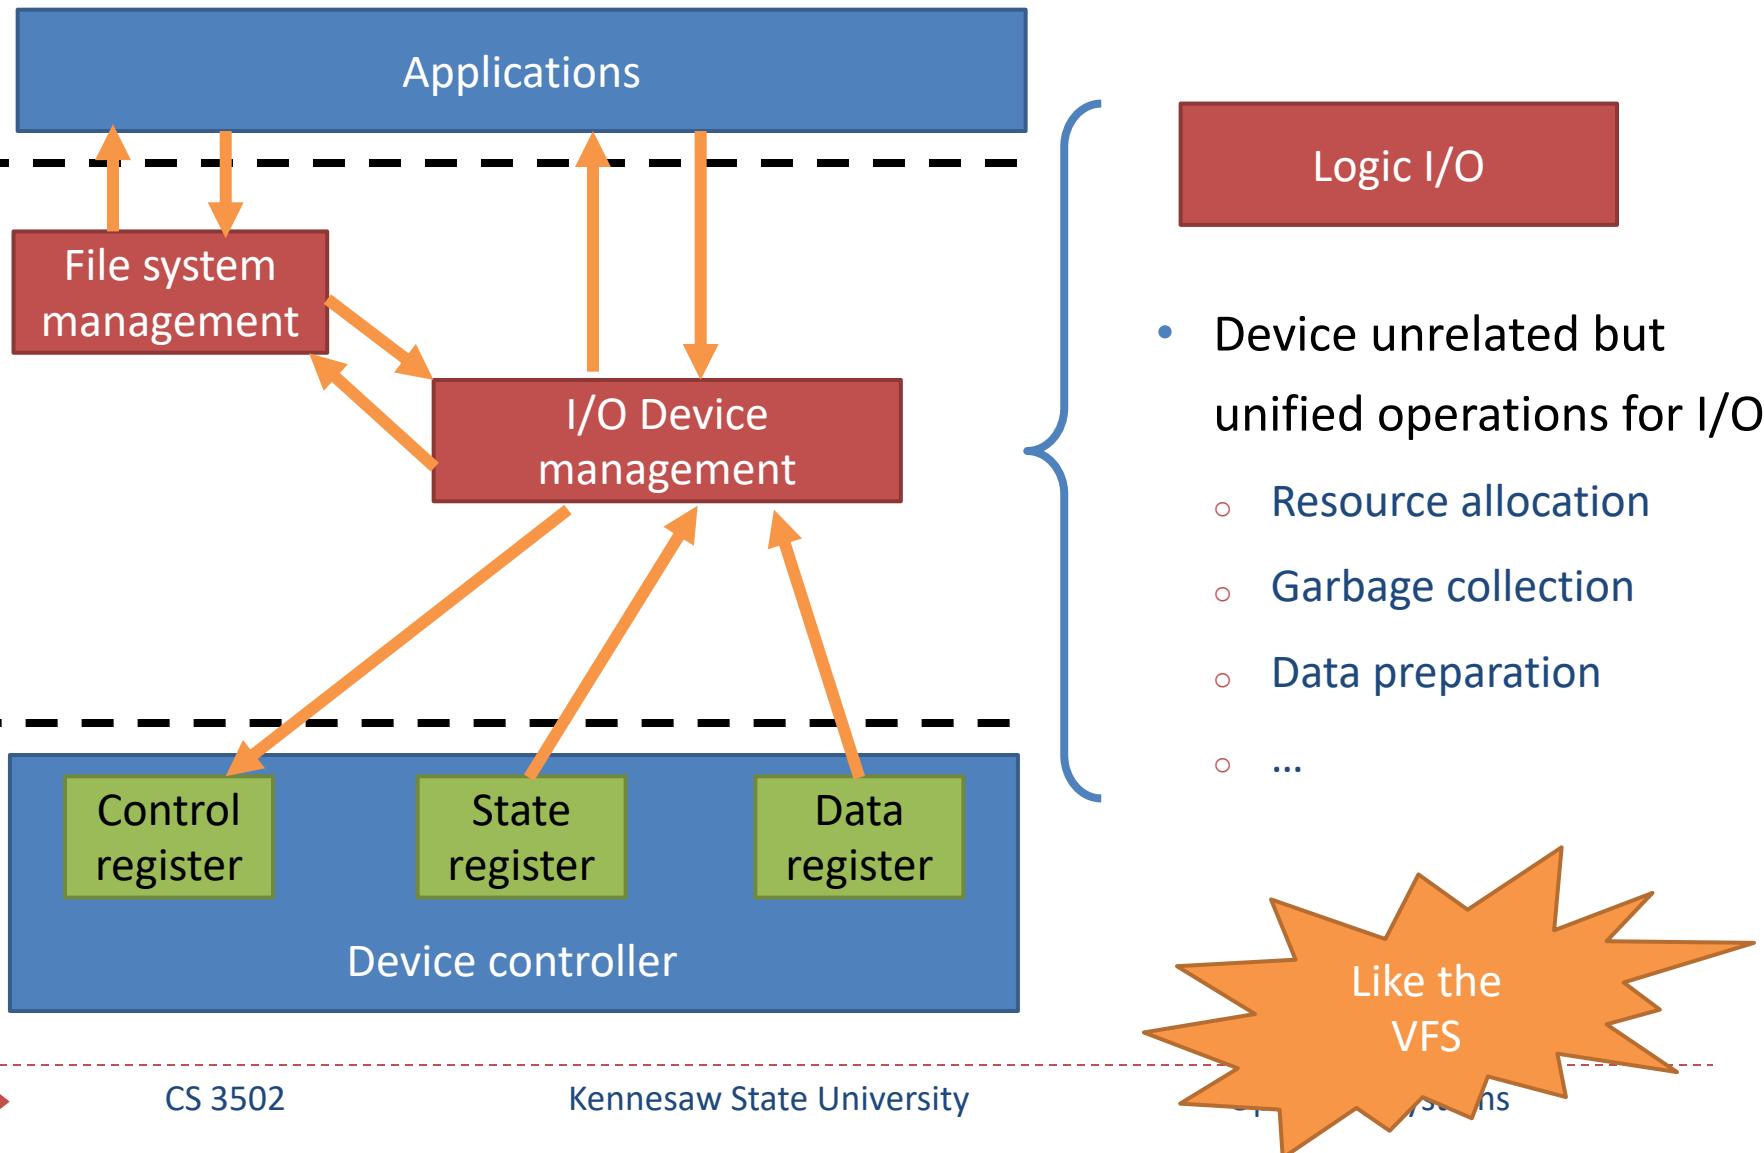
# I/O Management Example



# I/O Management



# I/O Management



# Logic I/O

---

- The logic I/O is to perform the I/O functions that are common to all devices and to provide a uniform interface to the user-level software

**Uniform interfacing for device drivers**

**Buffering**

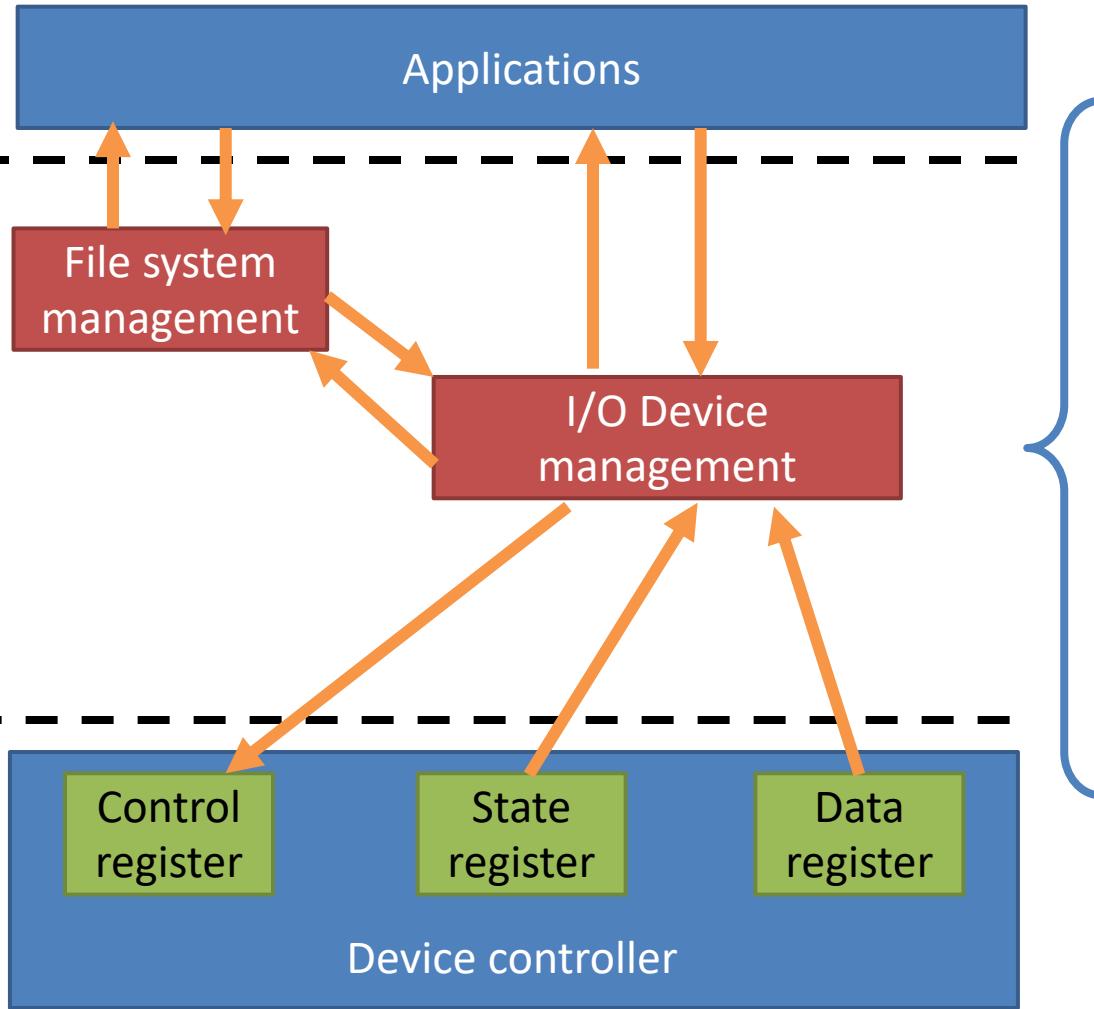
**Error reporting**

**Allocating and releasing dedicate devices**

**Providing a device-independent block size**



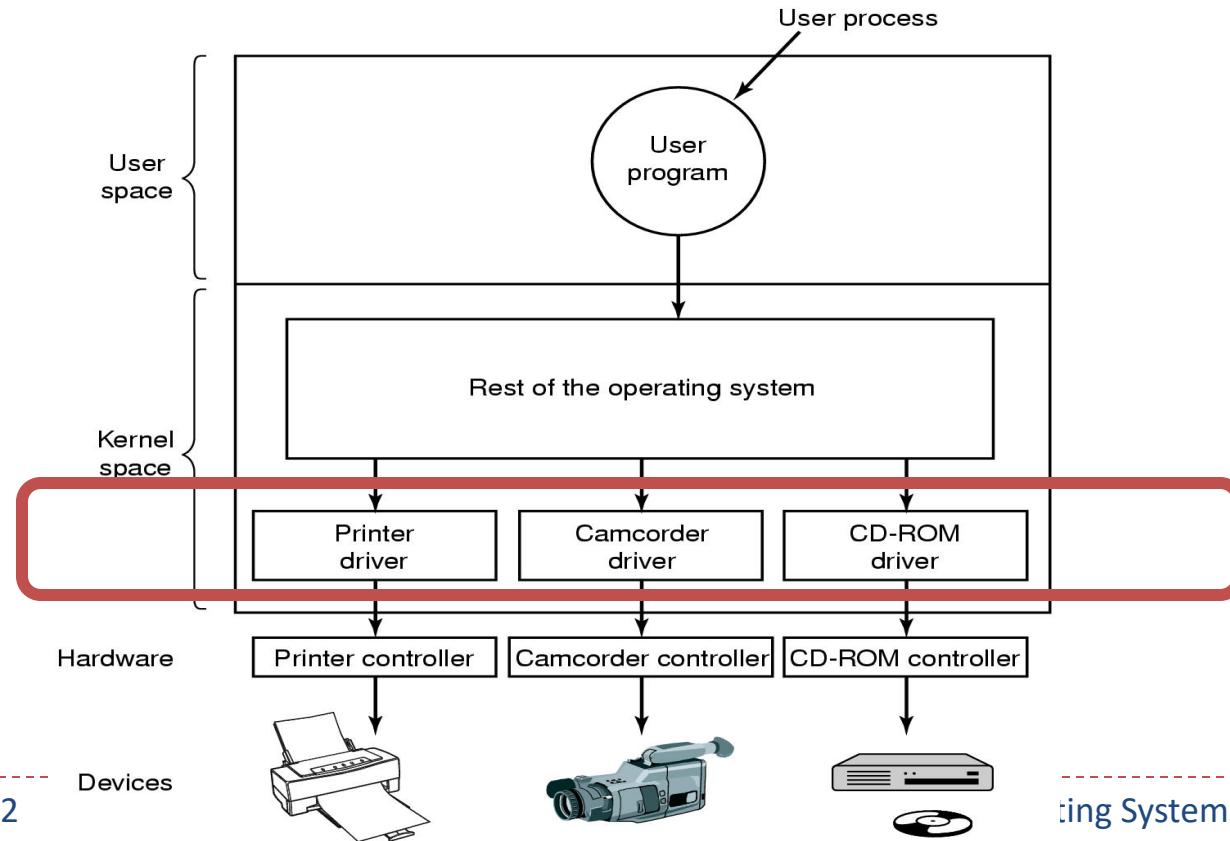
# I/O Management



- Device driver is used to interact with and control specific I/O device
- Device driver implements the operation functions on different I/O devices
- Device driver code is large

# Device driver

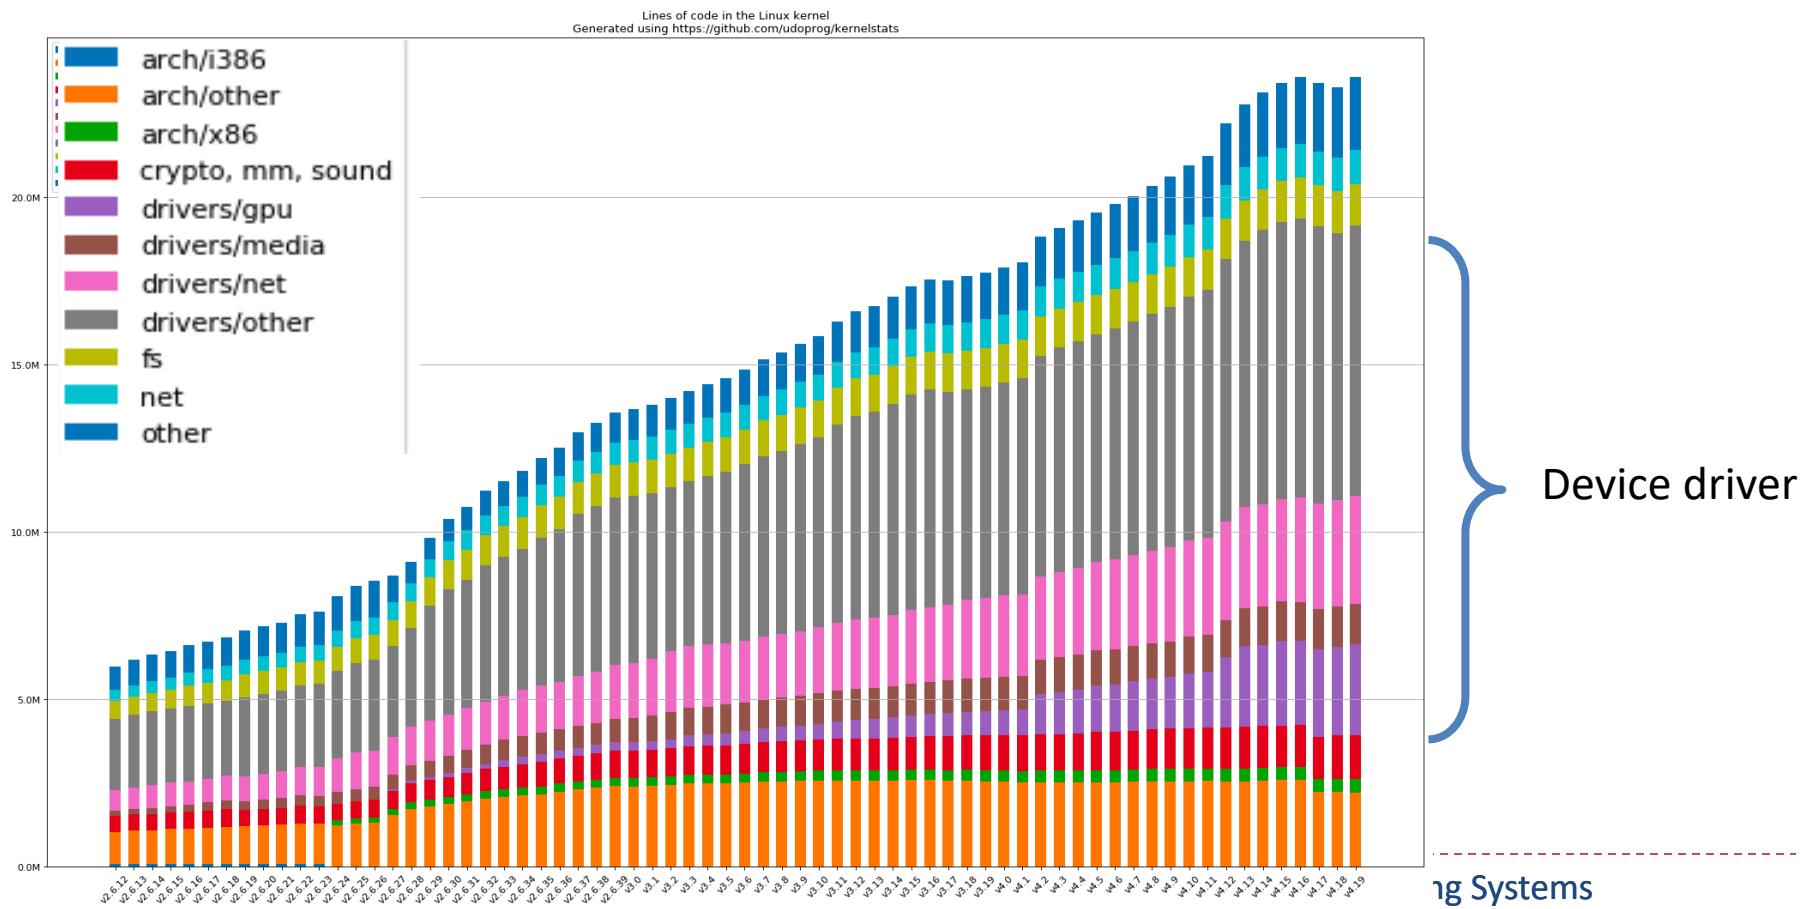
- Device drivers: provided by device's manufacturer, device-specific code for controlling the device (via device controller registers)



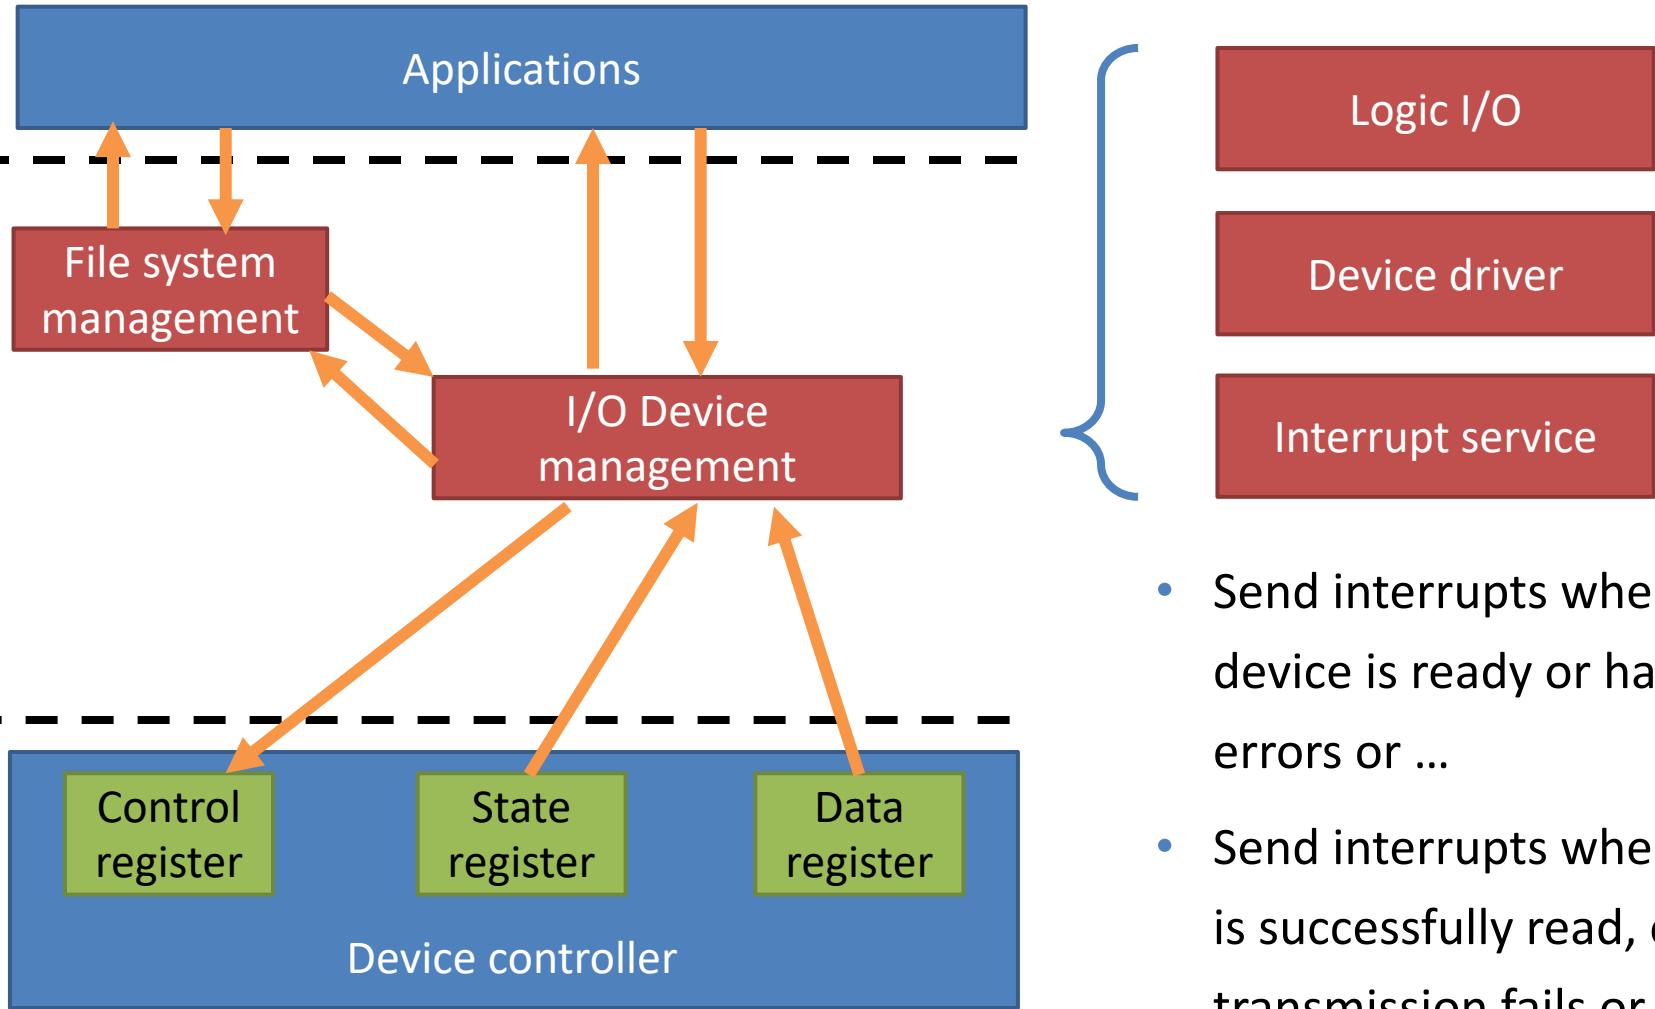
# Device driver related code in OS

- ~ 20+ million lines of code in Linux

<https://elixir.bootlin.com/linux/latest/source>



# I/O Management



- Send interrupts when device is ready or has errors or ...
- Send interrupts when data is successfully read, or transmission fails or ...

# Outline

---

- I/O System
- Characteristics and Category of I/O Devices
- I/O Device Component
- How CPU works with I/O devices
- I/O Performance



# Characteristics of I/O: complexity

---

- I/O performance is usually bottleneck of the system performance
- I/O → amounts of devices, different performance, parallel read/write → making OS more complex



# Characteristics of I/O: complexity → Speed

Device	Data rate
Keyboard	10 bytes/sec
Mouse	100 bytes/sec
56K modem	7 KB/sec
Telephone channel	8 KB/sec
Dual ISDN lines	16 KB/sec
Laser printer	100 KB/sec
Scanner	400 KB/sec
Classic Ethernet	1.25 MB/sec
USB (Universal Serial Bus)	1.5 MB/sec
Digital camcorder	4 MB/sec
IDE disk	5 MB/sec
40x CD-ROM	6 MB/sec
Fast Ethernet	12.5 MB/sec
ISA bus	16.7 MB/sec
EIDE (ATA-2) disk	16.7 MB/sec
FireWire (IEEE 1394)	50 MB/sec
XGA Monitor	60 MB/sec
SONET OC-12 network	78 MB/sec
SCSI Ultra 2 disk	80 MB/sec
Gigabit Ethernet	125 MB/sec
Ultrium tape	320 MB/sec
PCI bus	528 MB/sec
Sun Gigaplane XB backplane	20 GB/sec



Bytes/s



KBytes/s



MBytes/s



# Characteristics of I/O: complexity → Application

---



One dimension

Dynamic move



Two dimension

Static content

- Different characteristics of I/O devices make the OS interact with the drivers differently

# Characteristics of I/O: complexity → Control interface and transmission unit



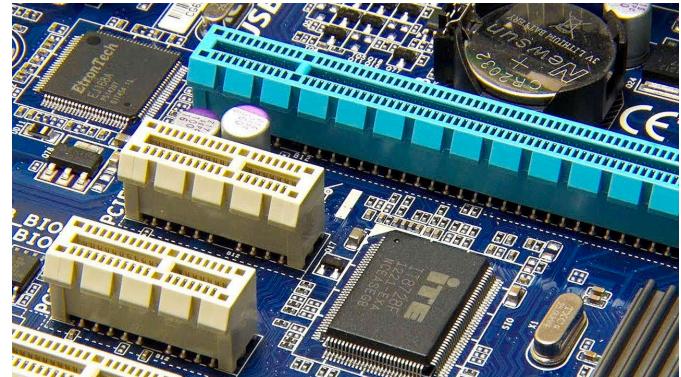
CD-Rom



CD-Rom Interface



SSD



PCIe

# Characteristics of I/O: complexity →

## Control interface and transmission unit



- For the same category of I/O device, such as HDD, the transmission unit could be different due to different manufacturer and drivers
  - Byte
  - Sector
  - Block

# Category of I/O Devices

---

- Block devices:
  - ✓ Store or transfer data in blocks
  - ✓ High transmission speed, addressable, support seeking and the random access of data
  - ✓ E.g., CD, SSD, HDD
- Character devices:
  - ✓ Store or transfer data in characters
  - ✓ Low transmission speed, not addressable
  - ✓ E.g., Keyboard, Mouse, Printer



# Category of I/O Devices

---

- Storage devices:
  - ✓ Used for temporary or persistent storage
  - ✓ E.g., Disk, Tape
- Transmission devices:
  - ✓ Used to transfer data from different locations
  - ✓ E.g., Ethernet card, Bluetooth, wireless card
- User interface devices:
  - ✓ Used to interact with the system
  - ✓ E.g., Monitor, keyboard, mouse, electronic pen, speaker



# Outline

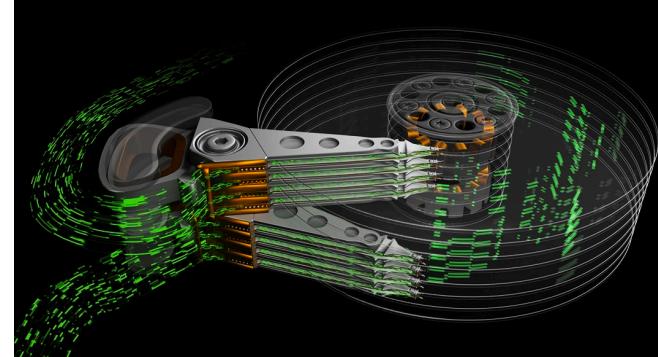
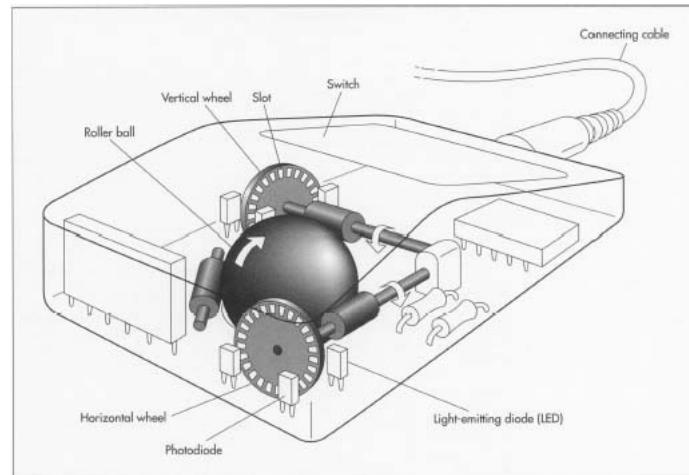
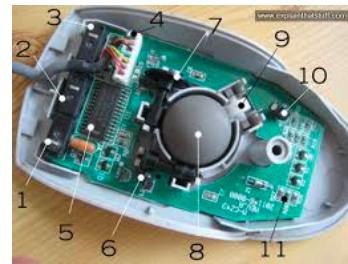
---

- I/O System
- Characteristics and Category of I/O Devices
- I/O Device Component
- How CPU works with I/O devices
- I/O Performance



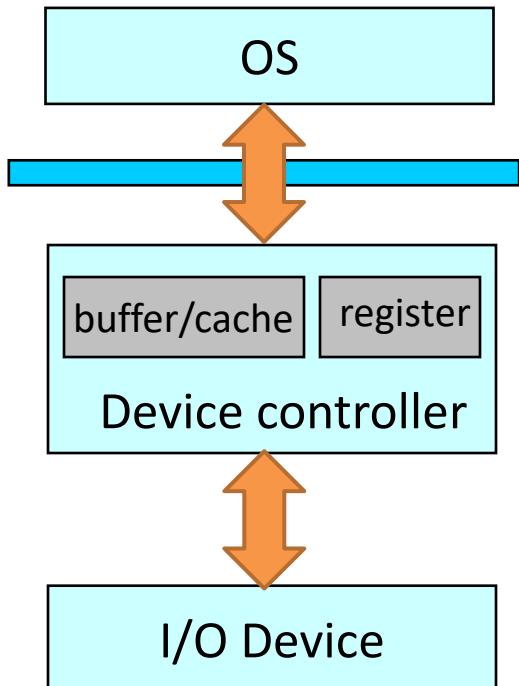
# I/O Device Component

- I/O devices have two components:
  - **mechanical component (physical part)**
  - **electronic component (device controller)**

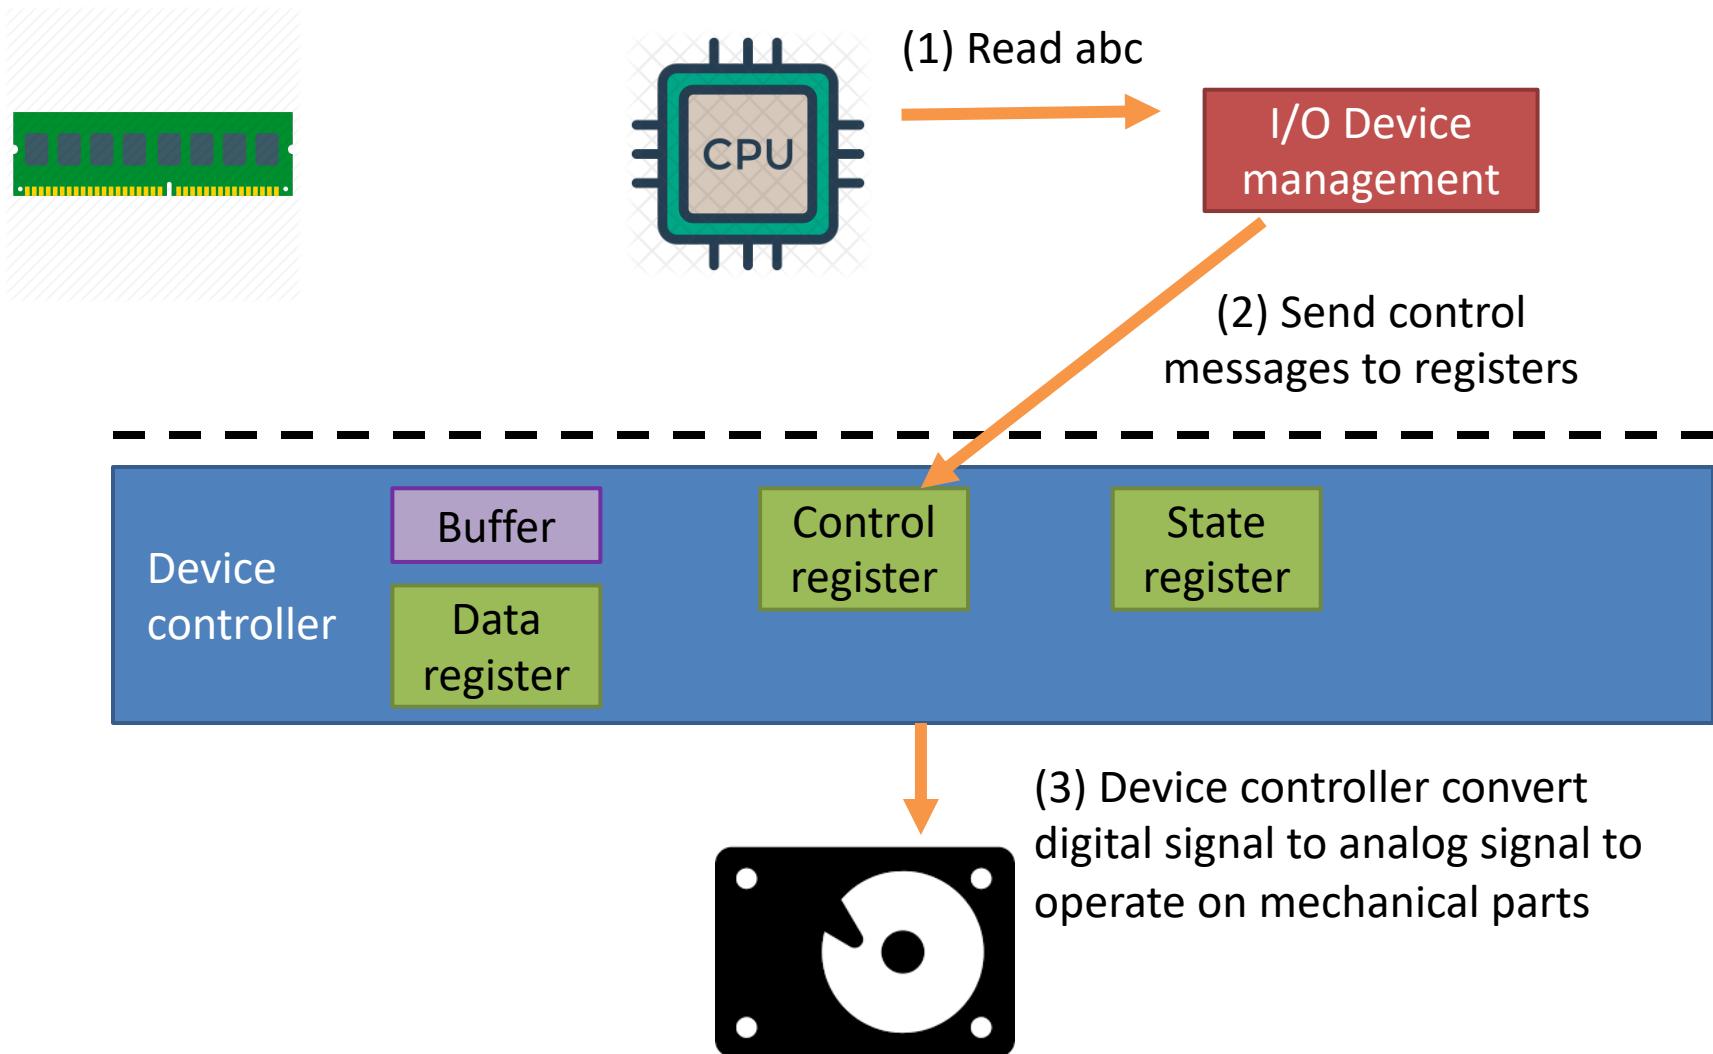


# I/O Device Component

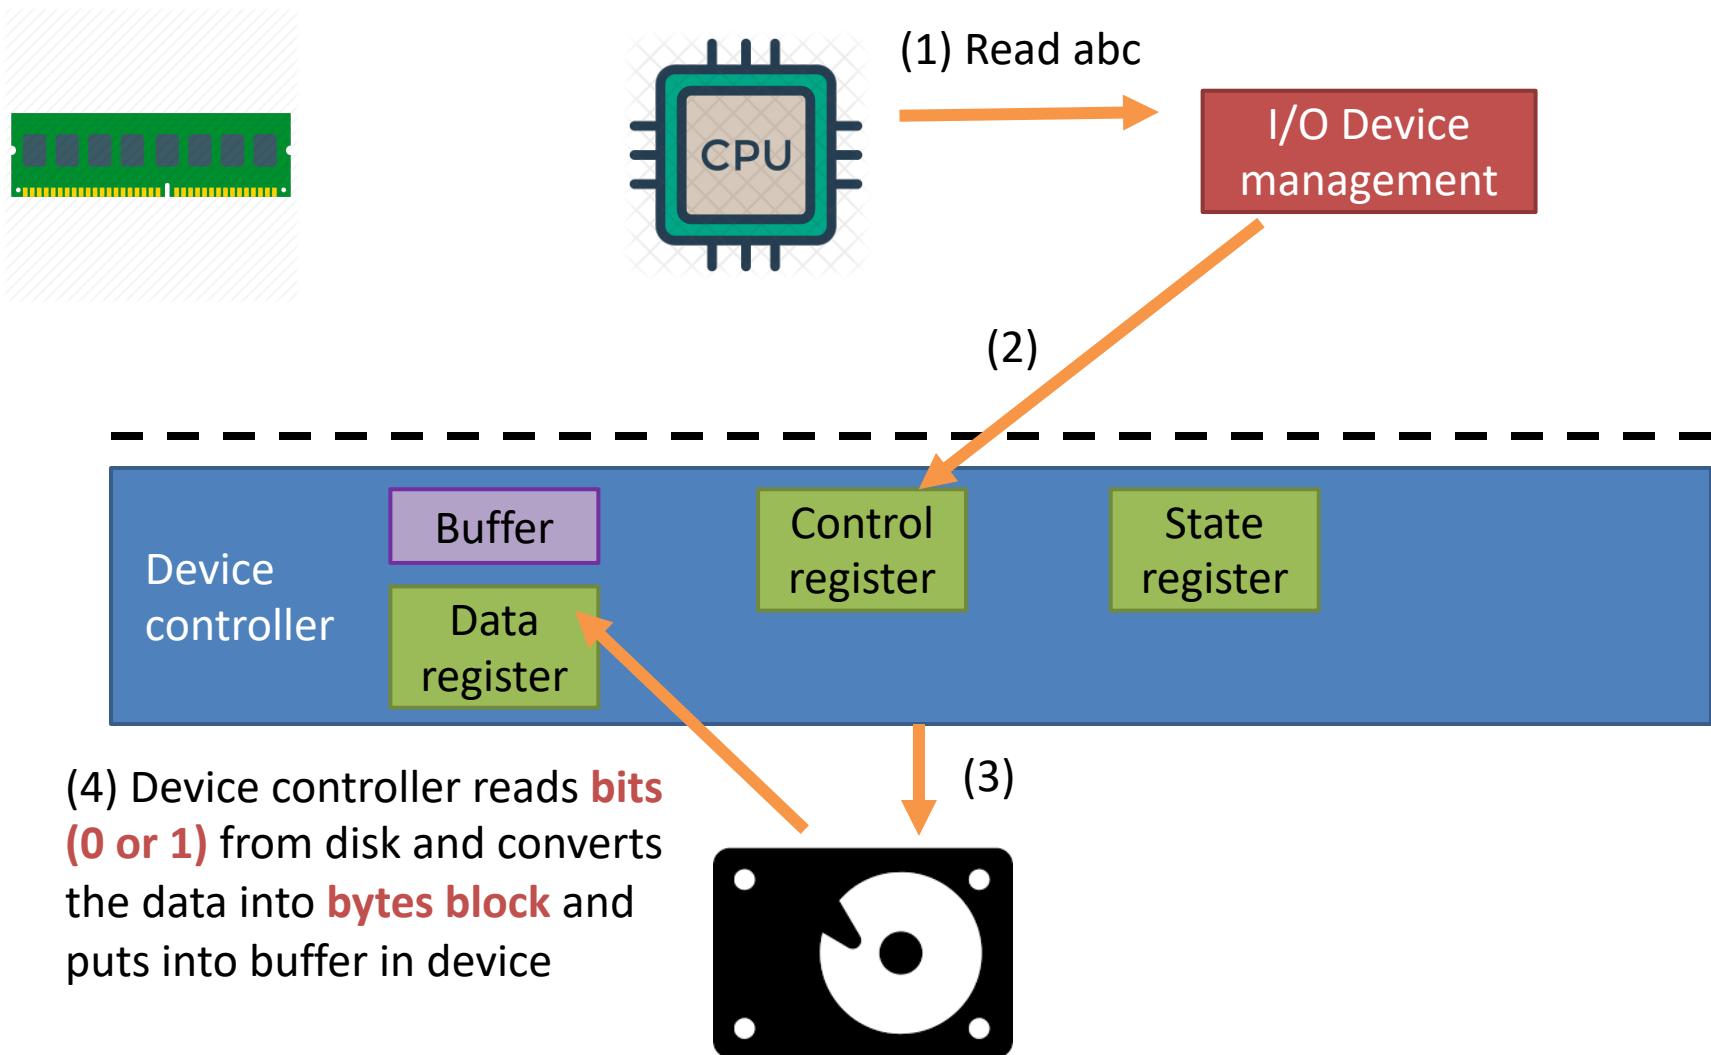
- I/O devices have two components:
  - mechanical component (physical part)
  - **electronic component (device controller)**
- Controller's tasks
  - Address transformation
  - Receive data and control message from the OS, and send back the data and status message to the OS
  - Convert digital signal into analog signal for the mechanical parts or vice versa
  - Memory buffer, data preprocess, perform error correction as necessary



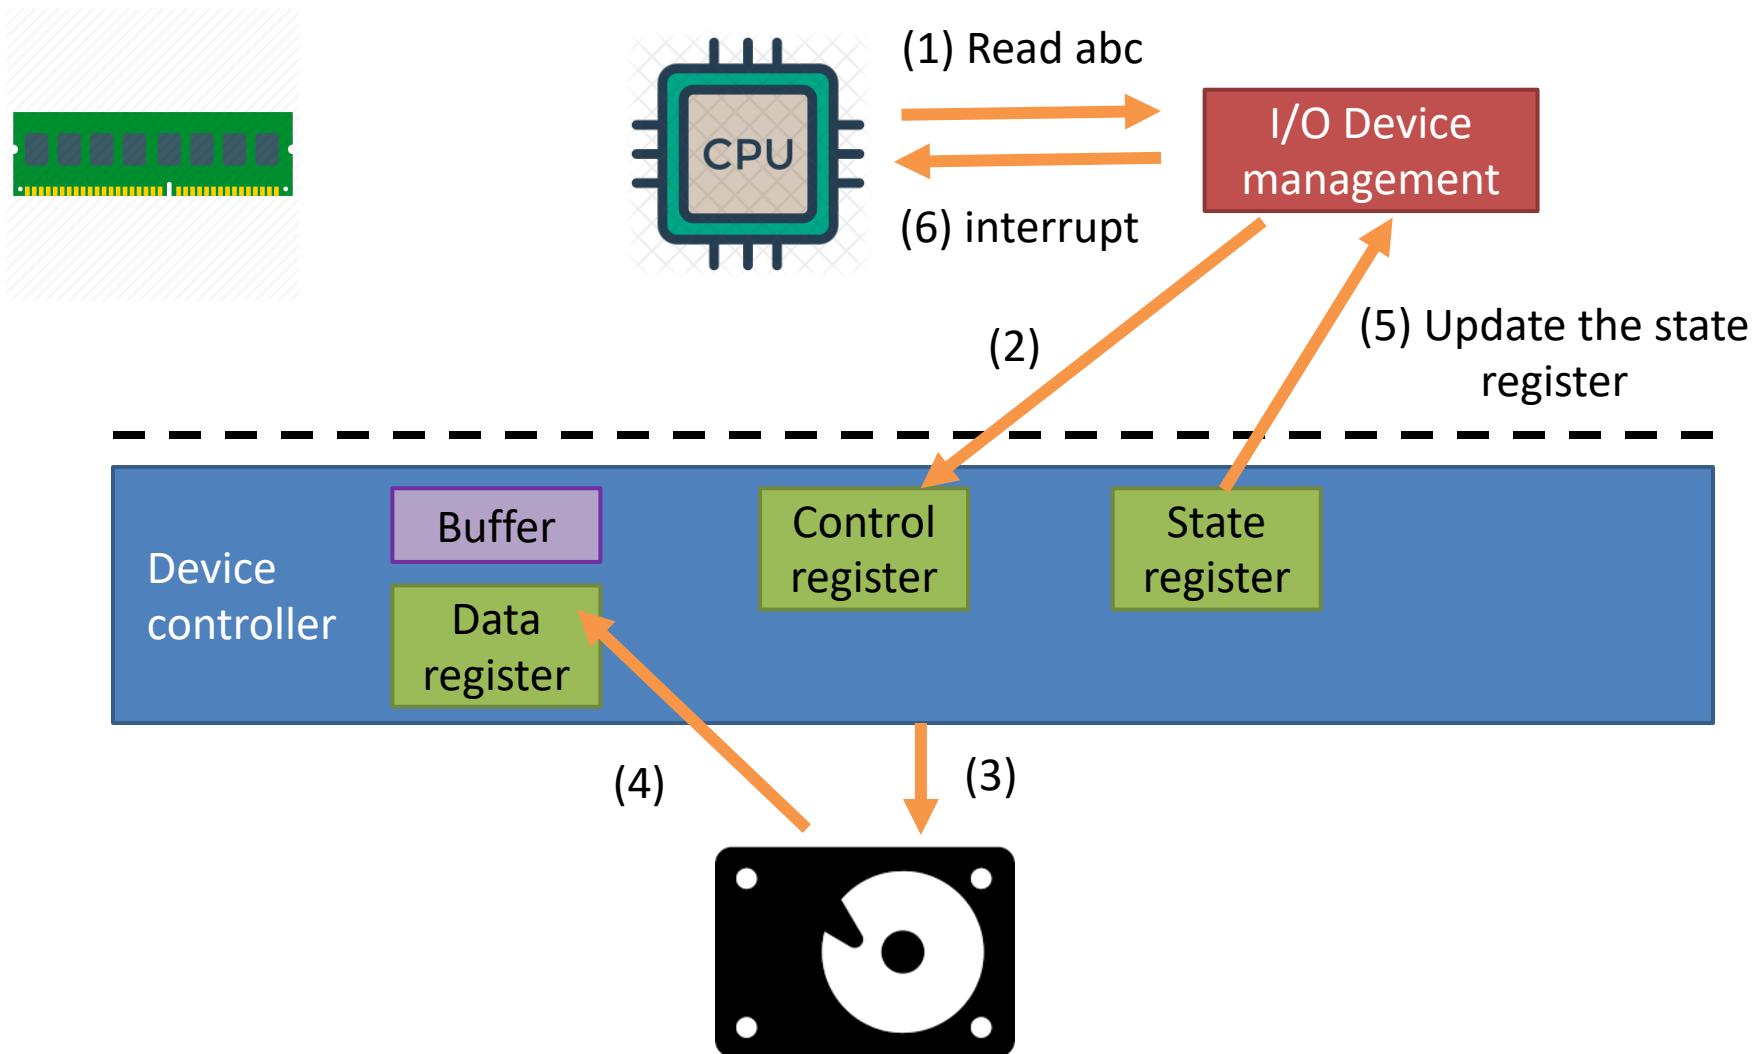
# I/O Device Component: Device controller



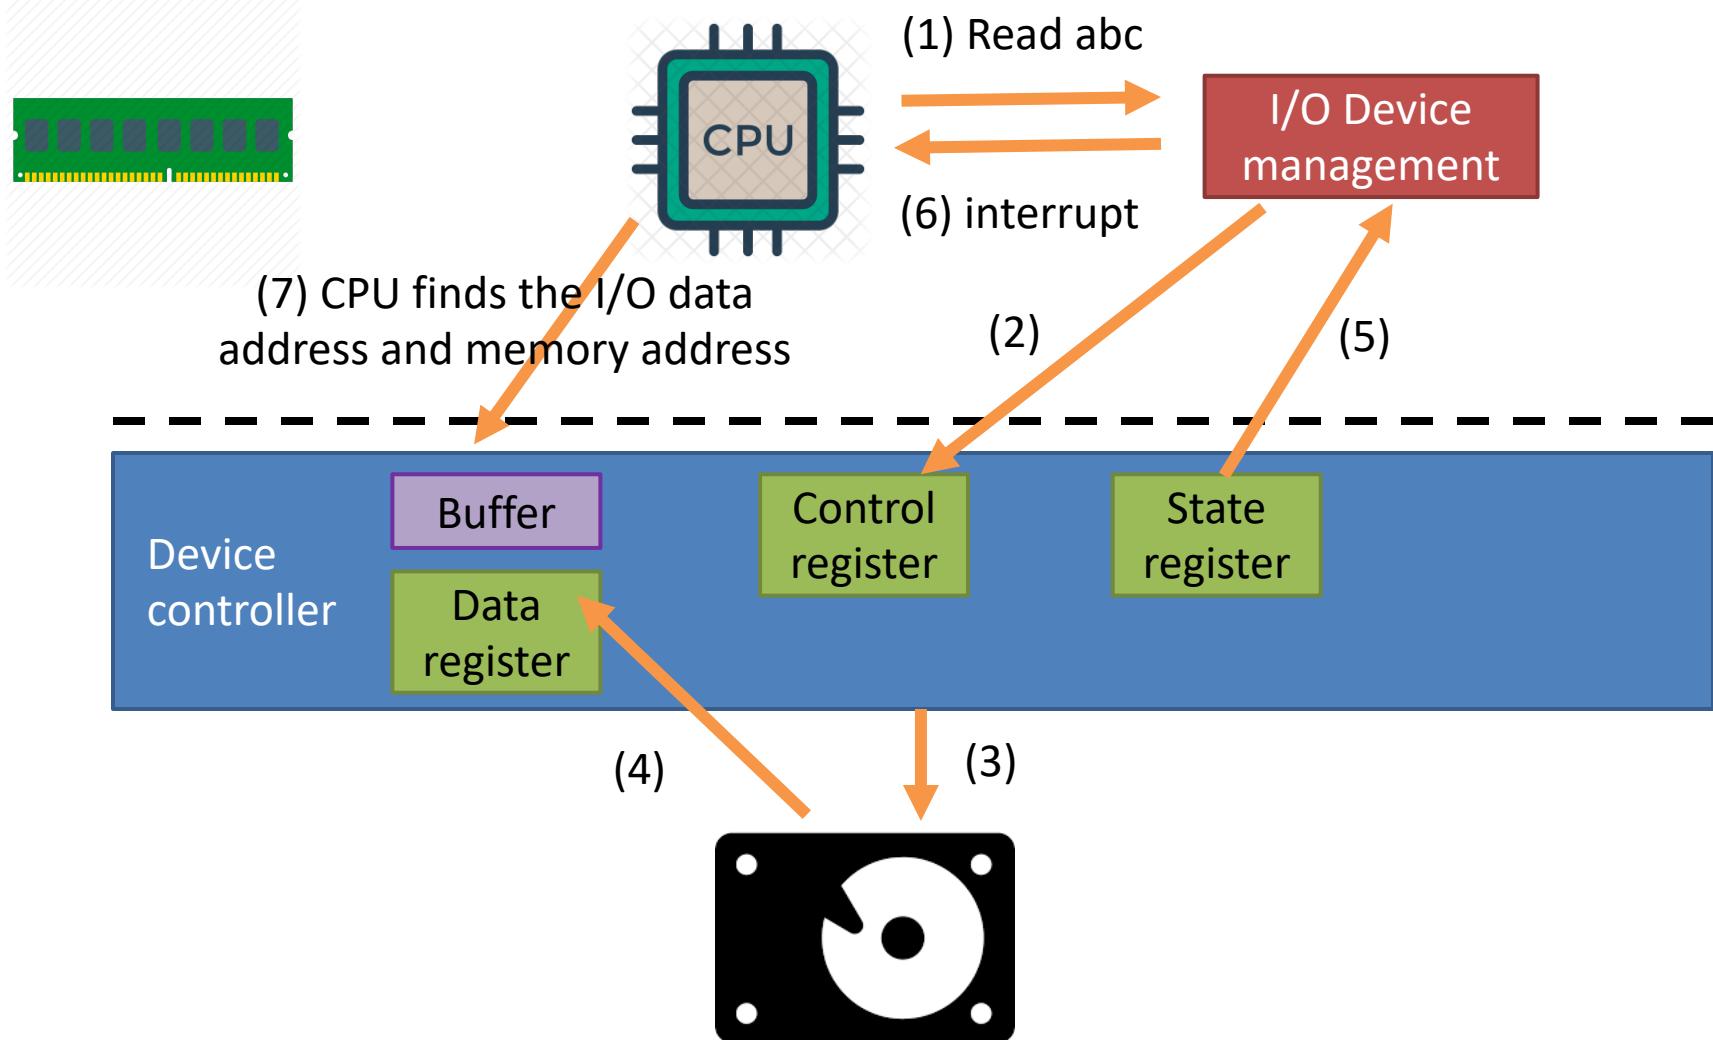
# I/O Device Component: Device controller



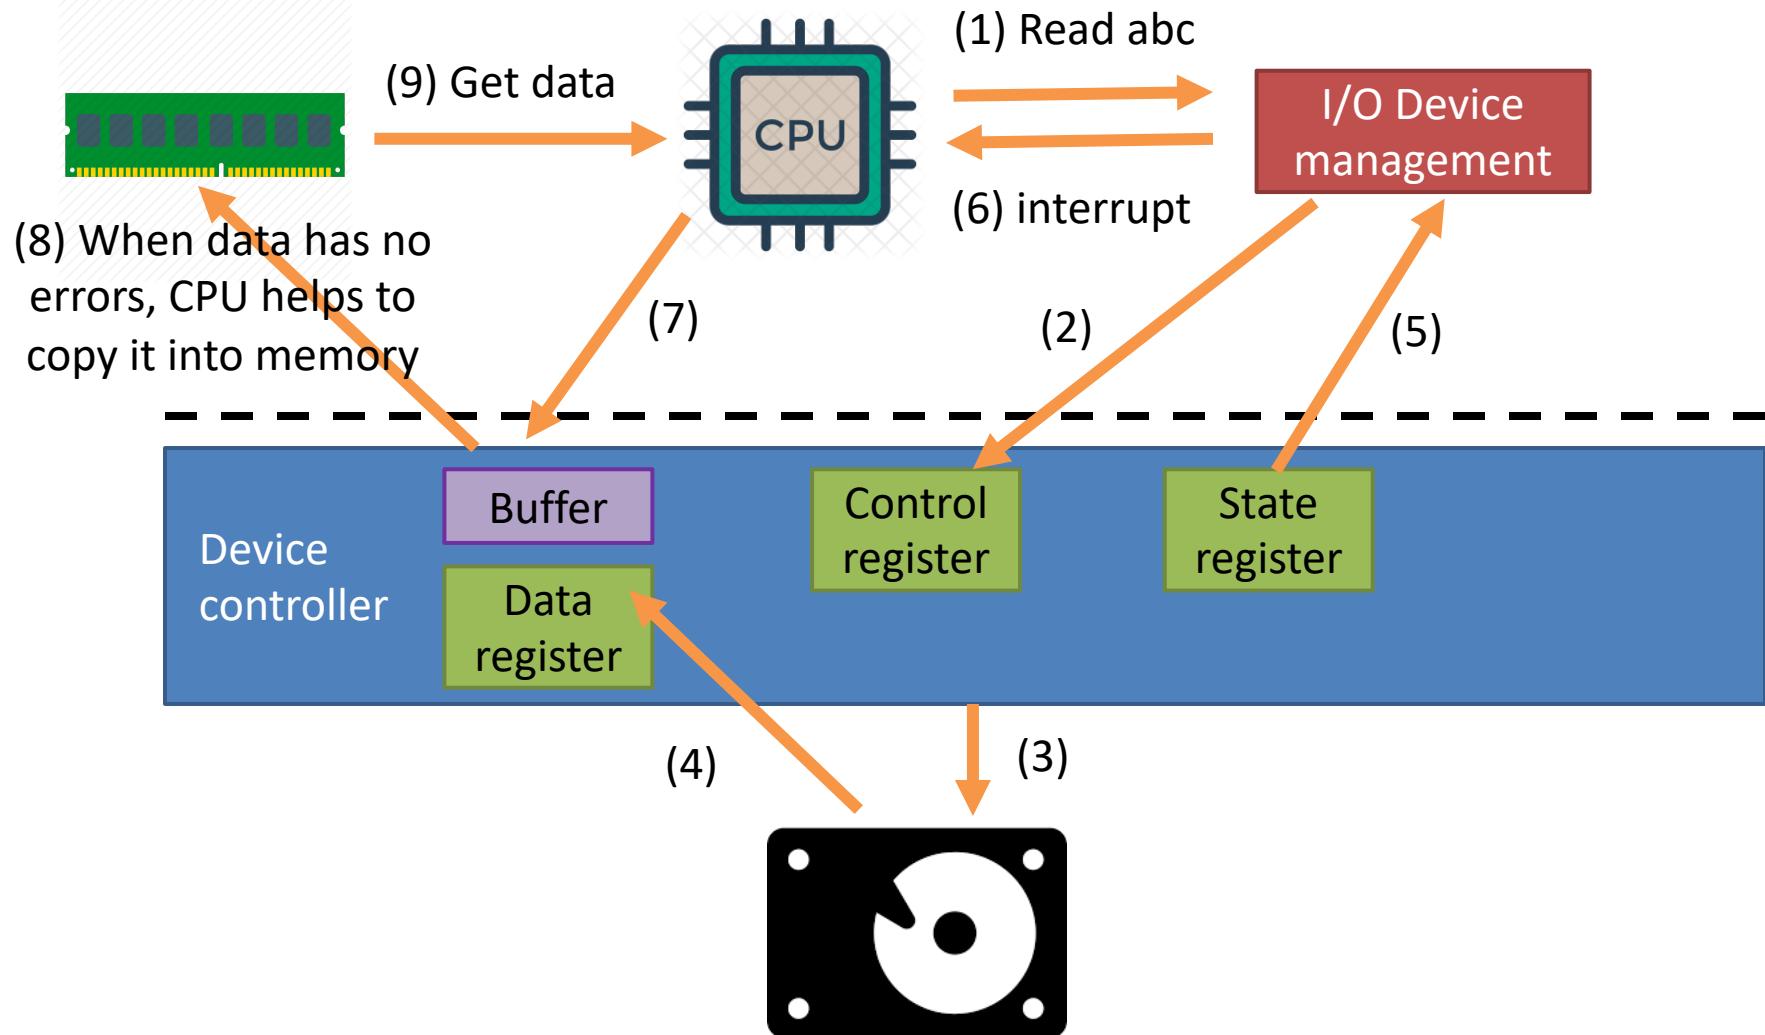
# I/O Device Component: Device controller



# I/O Device Component: Device controller



# I/O Device Component: Device controller



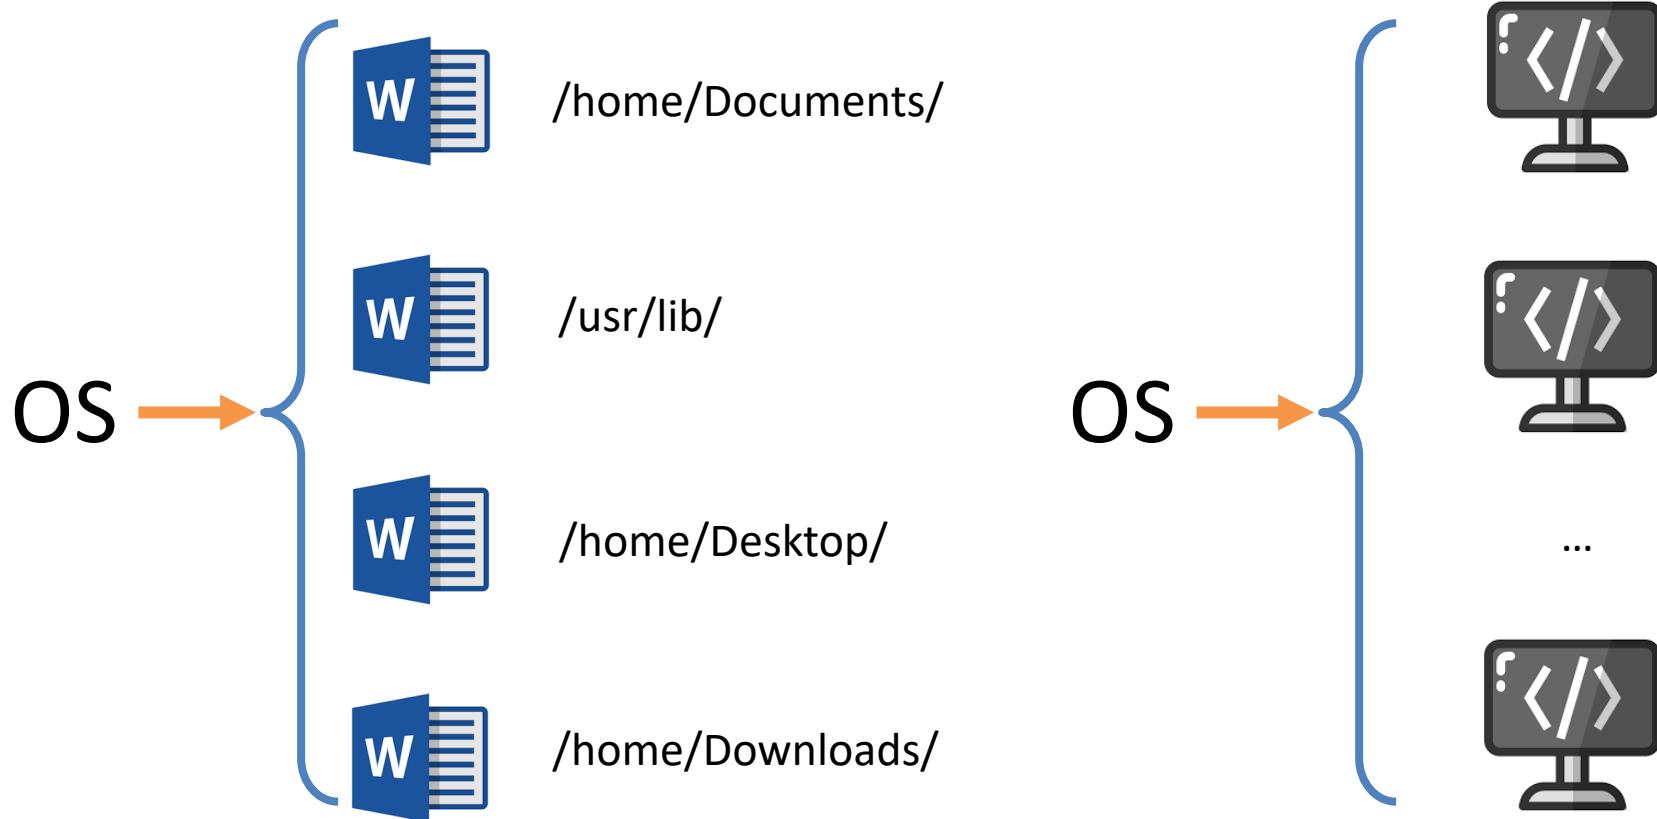
# Outline

---

- I/O System
- Characteristics and Category of I/O Devices
- I/O Device Component
- How CPU works with I/O devices
- I/O Performance



# How does the CPU(OS) address the I/O devices?



Different files have its unique address on disk

Addressing?

# How does the CPU(OS) address the I/O devices?

- Method 1: set fixed address for I/O devices

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)

# How does the CPU(OS) address the I/O devices?

- Method 1: set fixed address for I/O devices

I/O  
address

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)



memory



Data and instruction



# How does the CPU(OS) address the I/O devices?

- Method 1: set fixed address for I/O devices

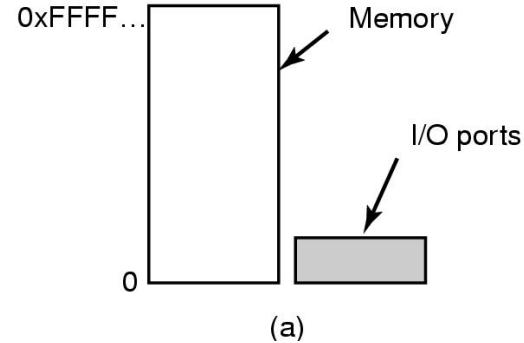
- Advantages:

- Independent from main memory address
- Easy to program on I/O

- Disadvantages:

- Fixed range and size, not flexible, extra space for I/O

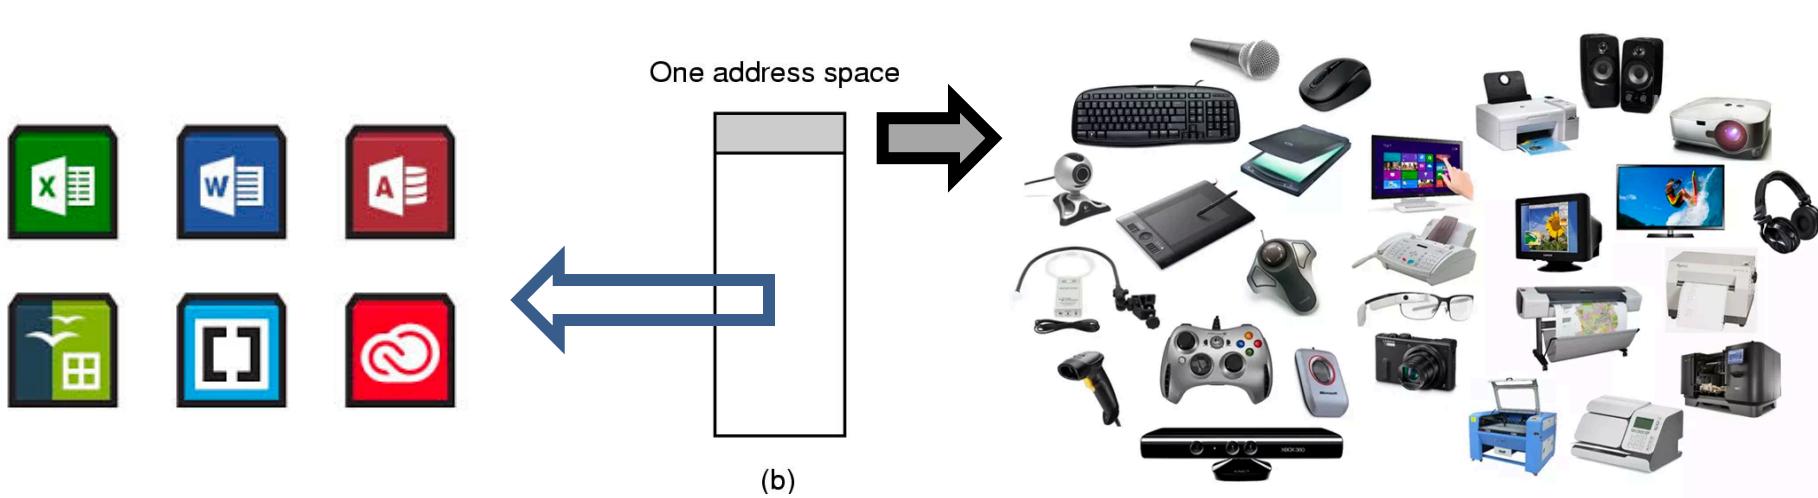
Two address



(a) Separate I/O and memory space

# How does the CPU(OS) address the I/O devices?

- Method 2: memory-mapped I/O
- A portion of memory address space is mapped to the device, and communications occur by reading and writing directly to/from those memory areas



(b) Memory-mapped I/O;

# How does the CPU(OS) address the I/O devices?

- Method 2: memory-mapped I/O

- Advantages:

- No need extra space for I/O, unified under memory operation
- I/O can take large address space

One address space



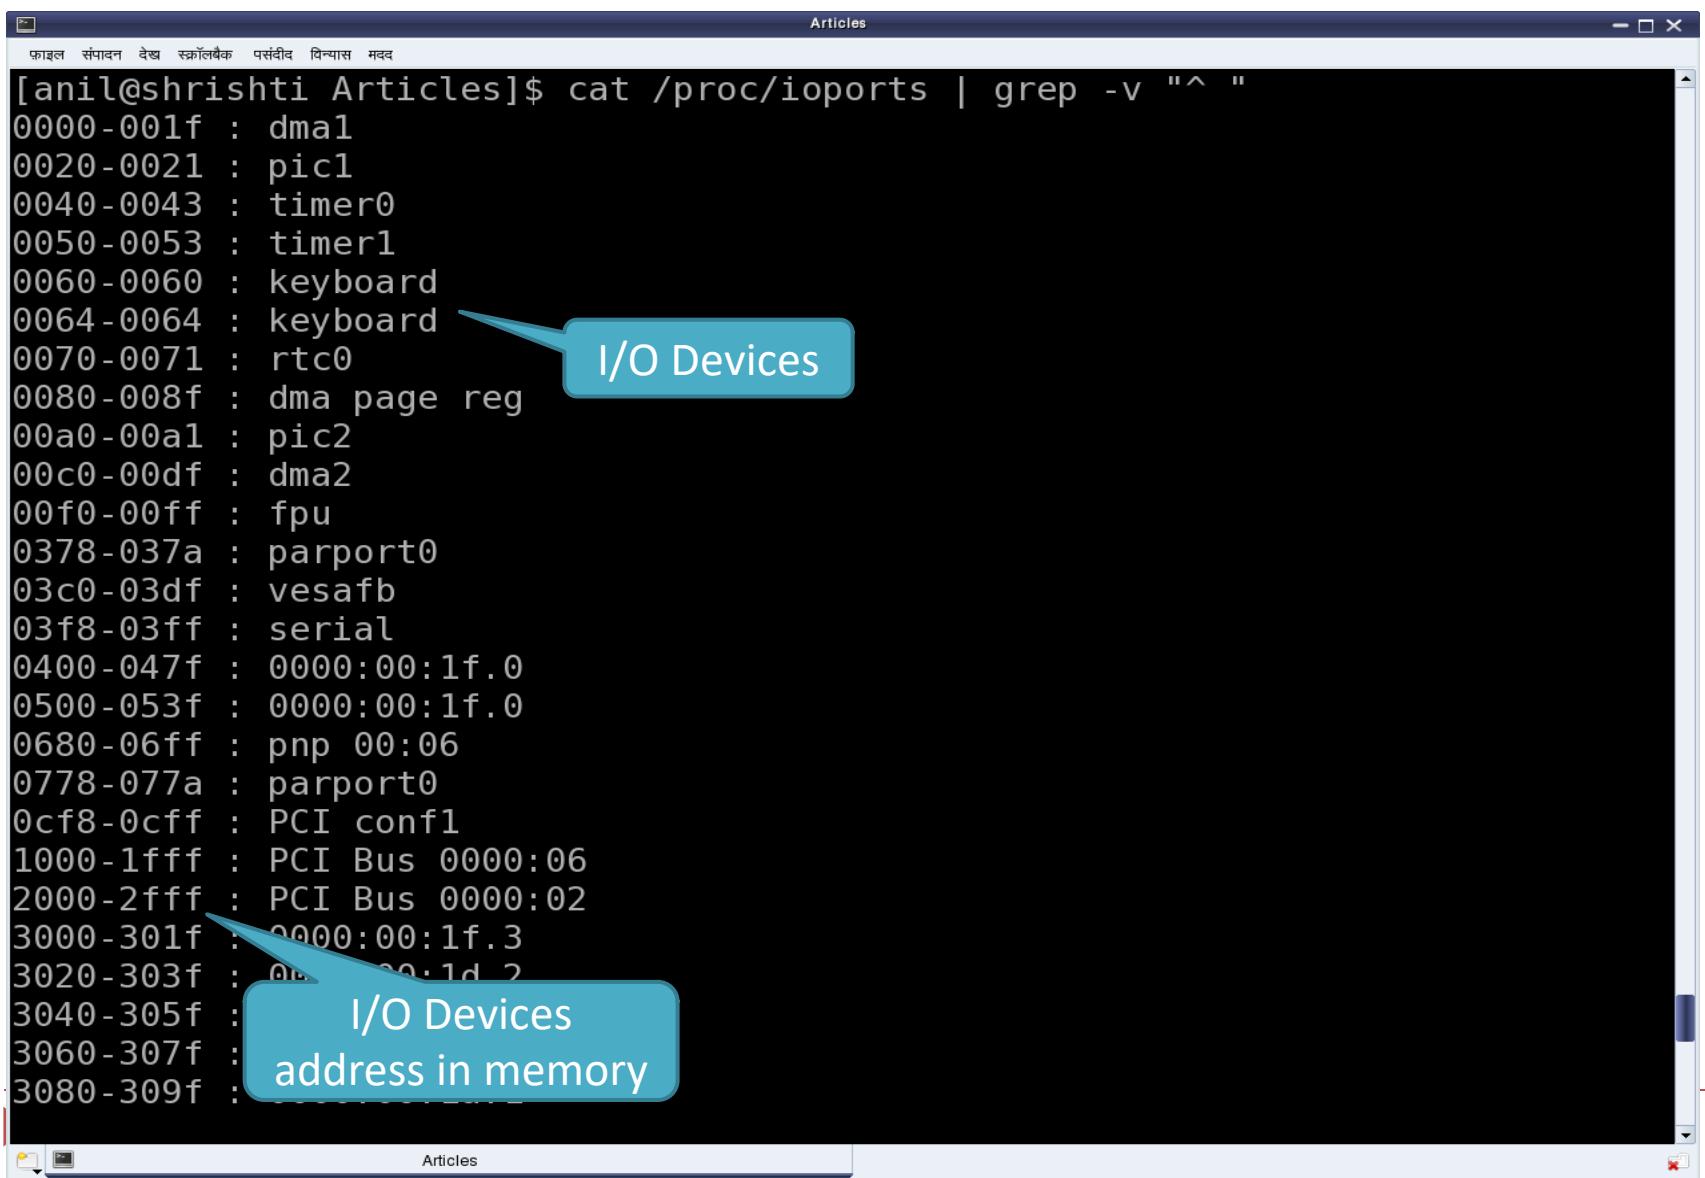
(b) Memory-mapped I/O;

- Disadvantages:

- Take part of the main memory address space

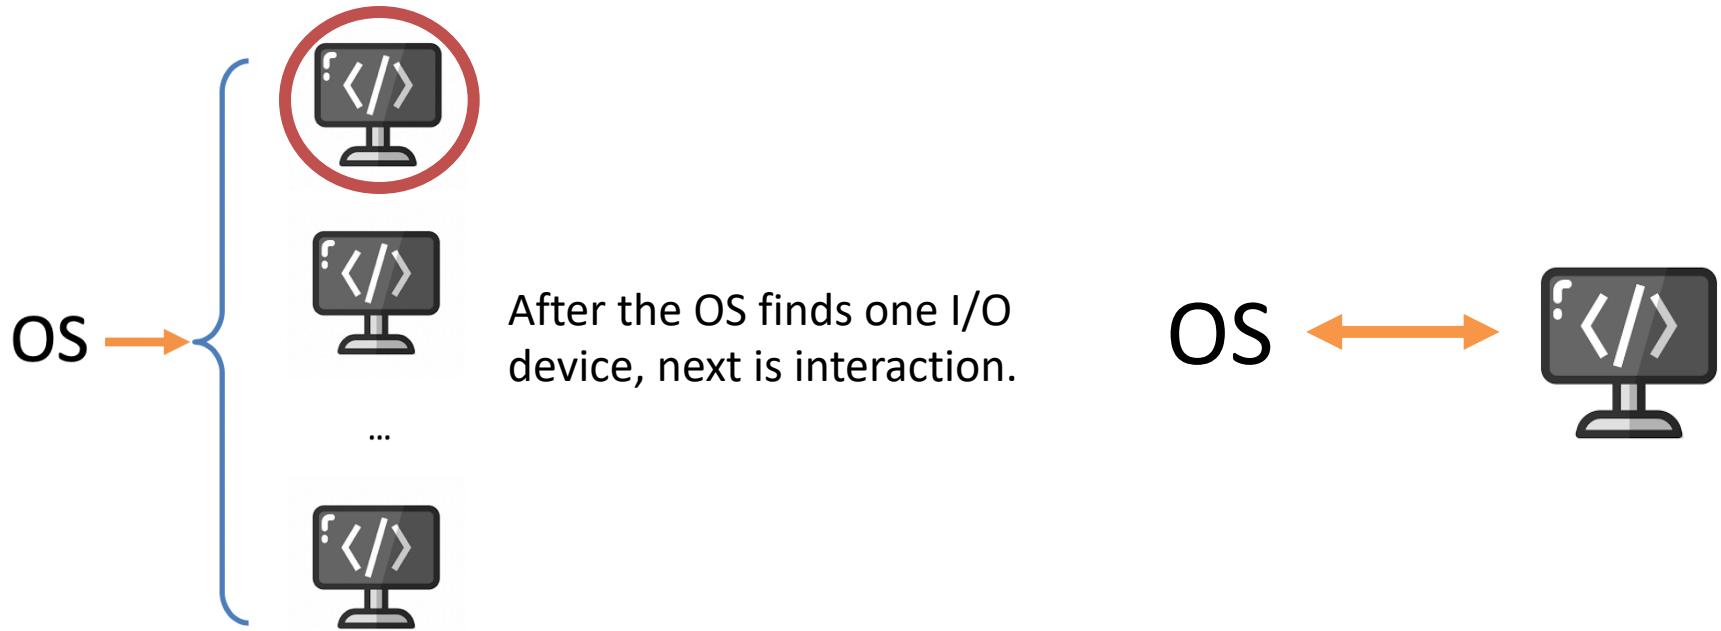


# Memory-mapped I/O Devices in Linux



```
[anil@shrishti Articles]$ cat /proc/ioports | grep -v "^\s"  
0000-001f : dma1  
0020-0021 : pic1  
0040-0043 : timer0  
0050-0053 : timer1  
0060-0060 : keyboard  
0064-0064 : keyboard  
0070-0071 : rtc0  
0080-008f : dma page reg  
00a0-00a1 : pic2  
00c0-00df : dma2  
00f0-00ff : fpu  
0378-037a : parport0  
03c0-03df : vesafb  
03f8-03ff : serial  
0400-047f : 0000:00:1f.0  
0500-053f : 0000:00:1f.0  
0680-06ff : pnp 00:06  
0778-077a : parport0  
0cf8-0cff : PCI conf1  
1000-1fff : PCI Bus 0000:06  
2000-2fff : PCI Bus 0000:02  
3000-301f : 0000:00:1f.3  
3020-303f : 0000:00:1d.2  
3040-305f :  
3060-307f :  
3080-309f :  
  
I/O Devices  
address in memory
```

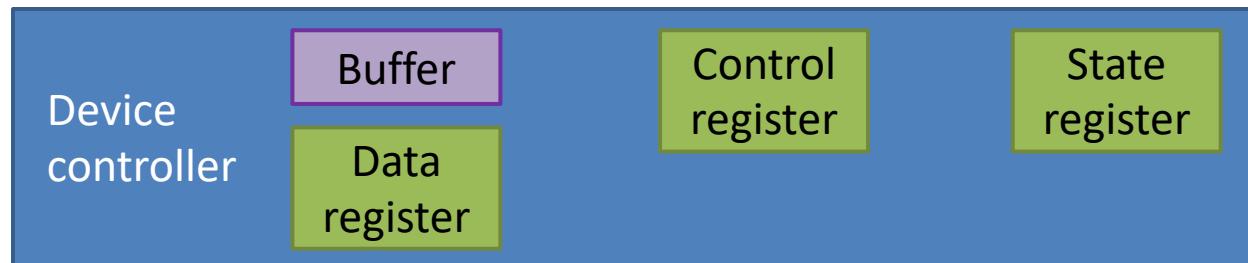
# How does the CPU(OS) communicate with the I/O devices?



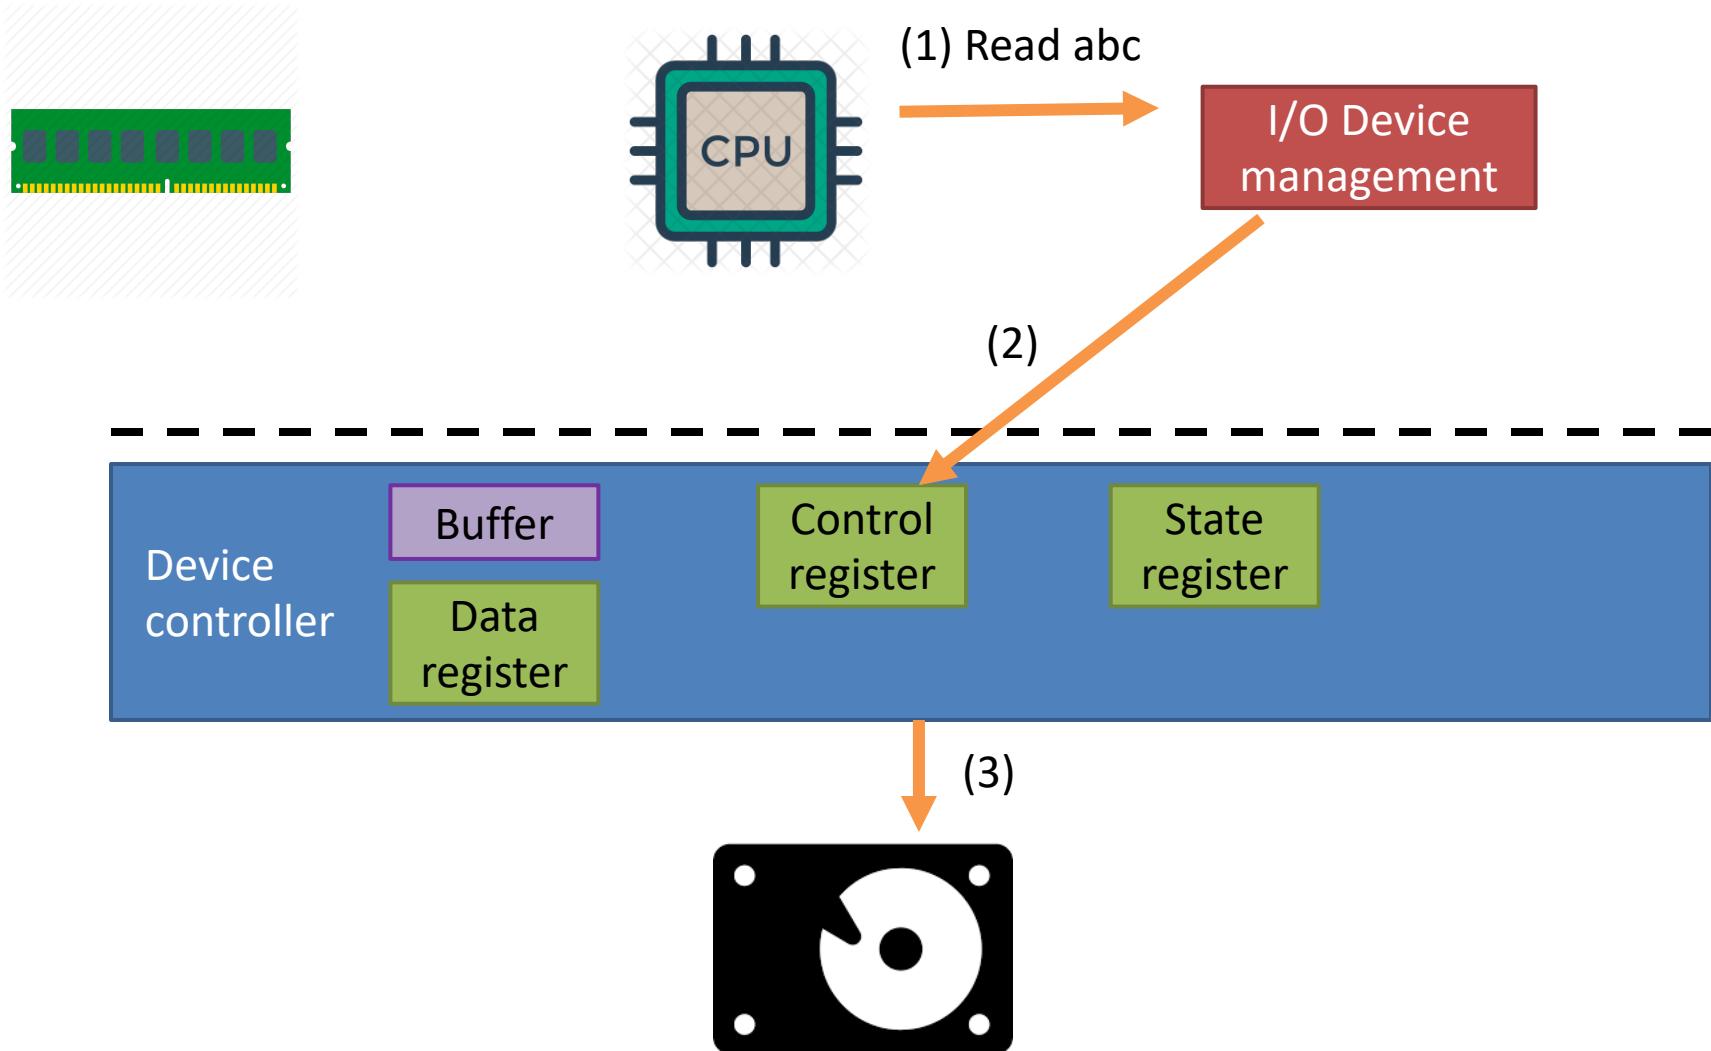
- The OS needs to know when:
  - The I/O device has completed an operation
  - The I/O operation has encountered an error

# How does the CPU(OS) communicate with the I/O devices?

- Option 1: Polling (Busy Waiting):
  - The I/O device puts information in a status register
  - The CPU periodically or continuously checks the status register

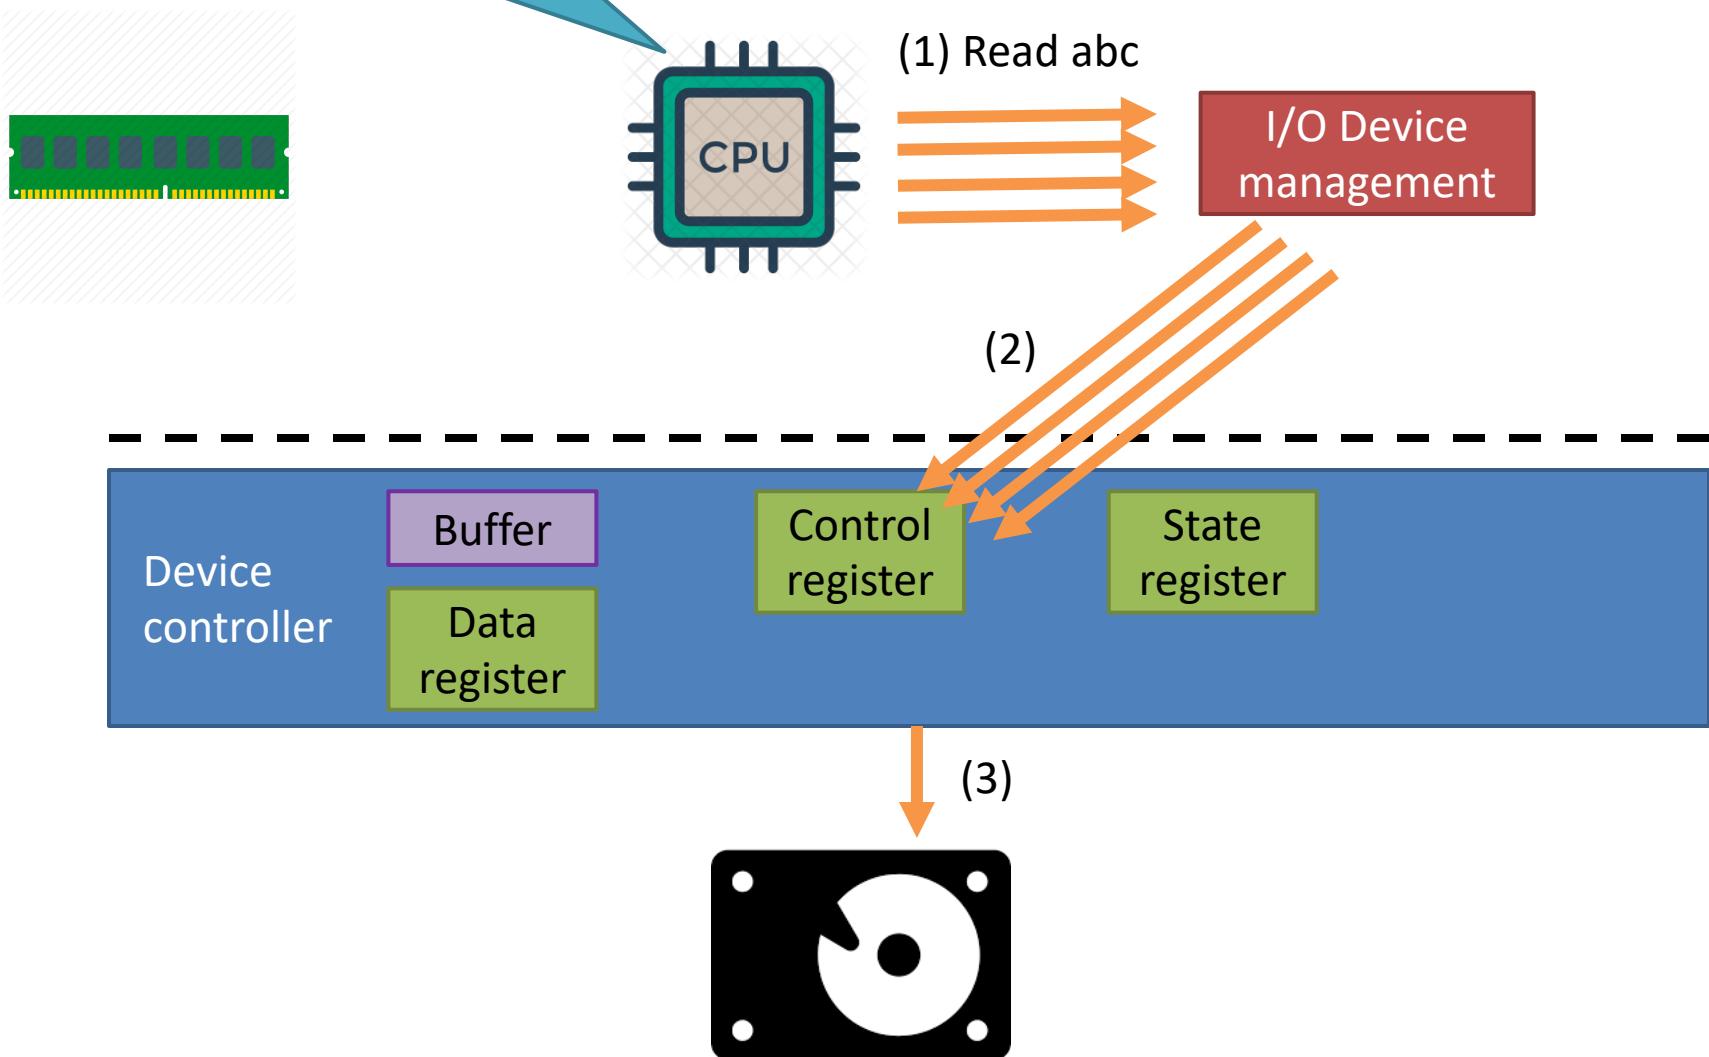


# How does the CPU(OS) communicate with the I/O devices? Polling if data not ready

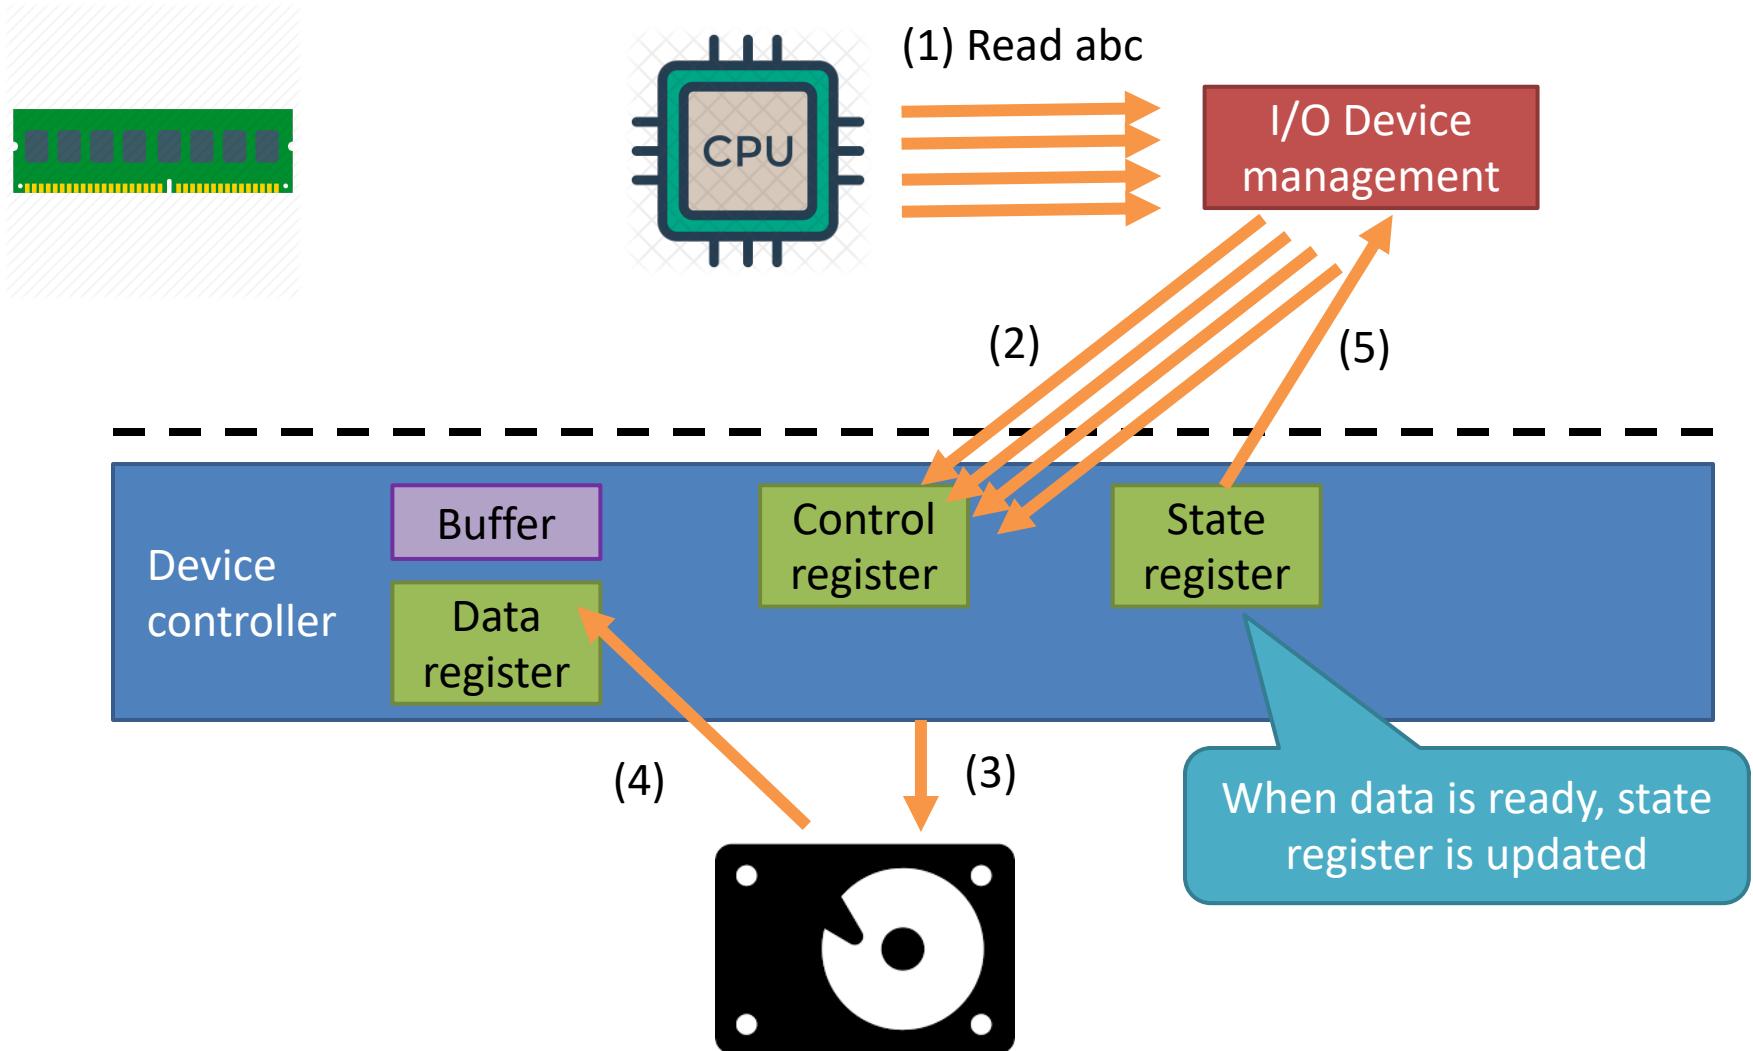


Keeps talking with I/O management and consumes lots of CPU resources

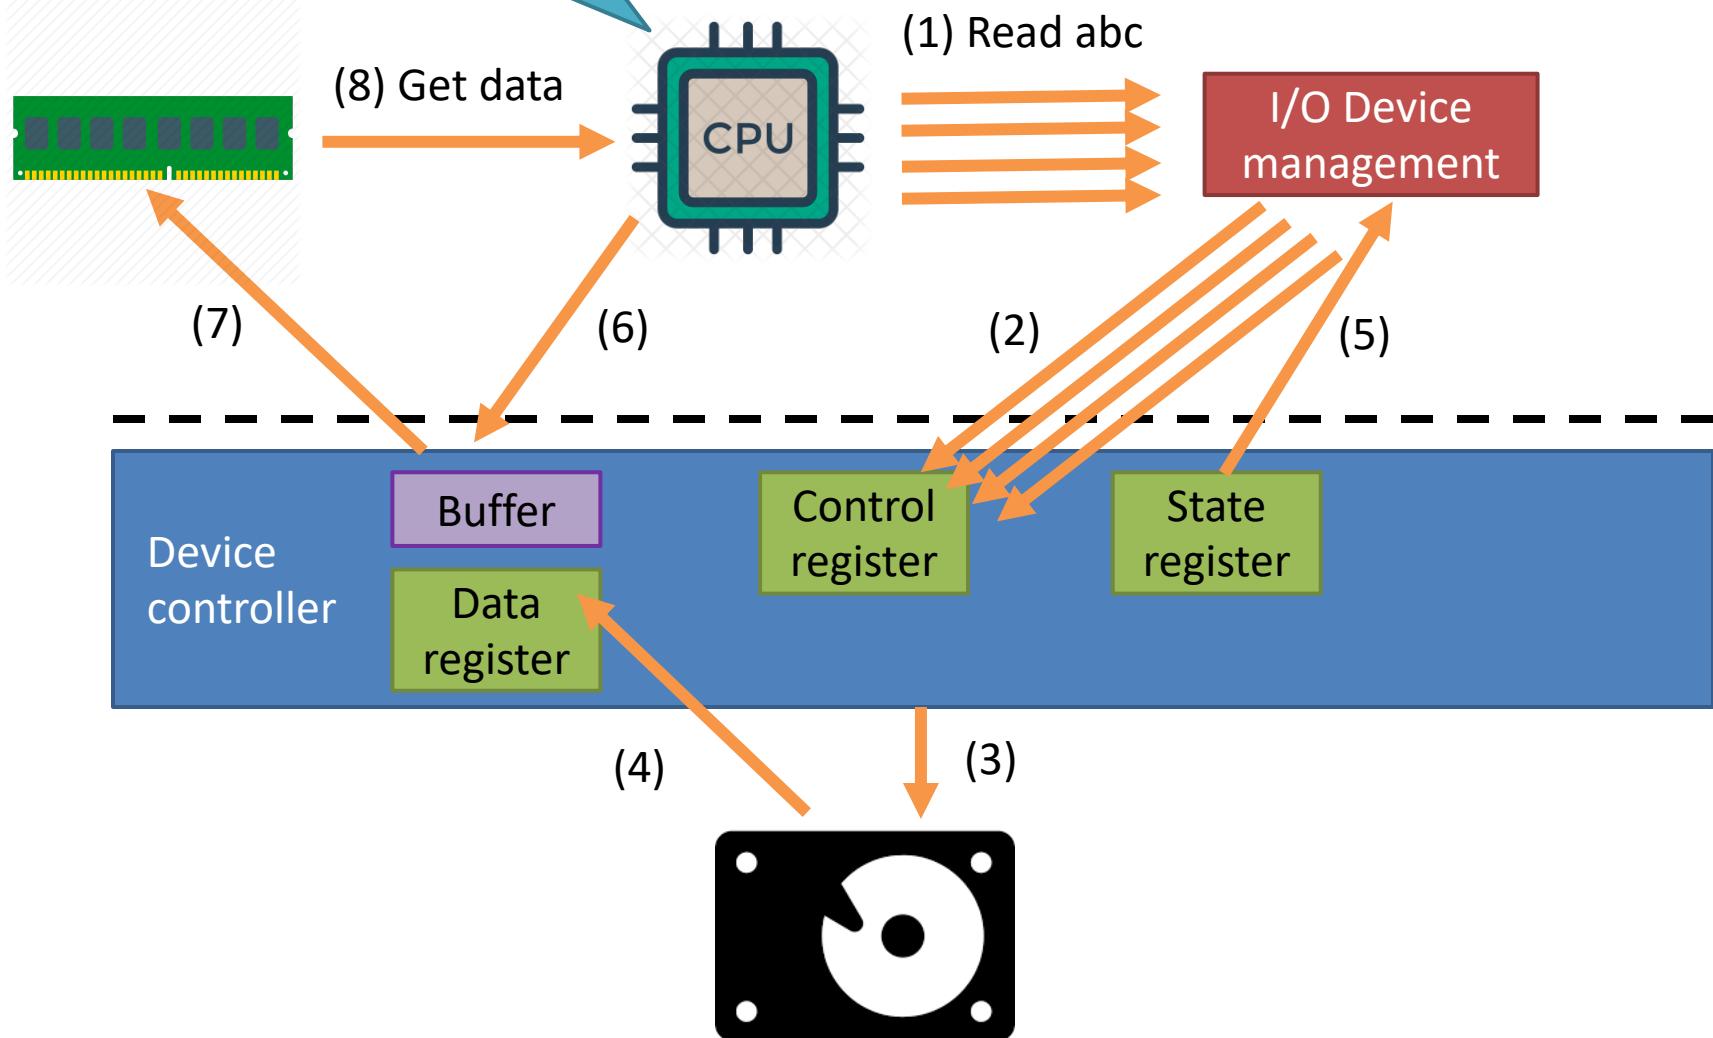
# CPU(OS) communicate with I/O devices? Polling if data not ready



# How does the CPU(OS) communicate with the I/O devices? Polling if data is ready



# How CPU(OS) communicate with the I/O device? Polling if data is ready



# How does the CPU(OS) communicate with the I/O devices?

---

- Option 1: Polling (Busy Waiting):
  - The I/O device puts information in a status register
  - The CPU periodically or continuously checks the status register
- Advantage:
  - Simple
- Disadvantage:
  - Polling overhead can consume a lot of CPU time



# How does the CPU(OS) communicate with the I/O devices?

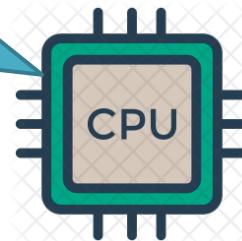
---

- Option 2: I/O Interrupt:
  - Whenever an I/O device needs attention from the processor, it interrupts the processor from what it is currently doing.



# How does the CPU(OS) communicate with the I/O devices? Interrupt if data not ready

CPU is used for other tasks when the request is sent but the data is not ready



(1) Read abc

I/O Device management

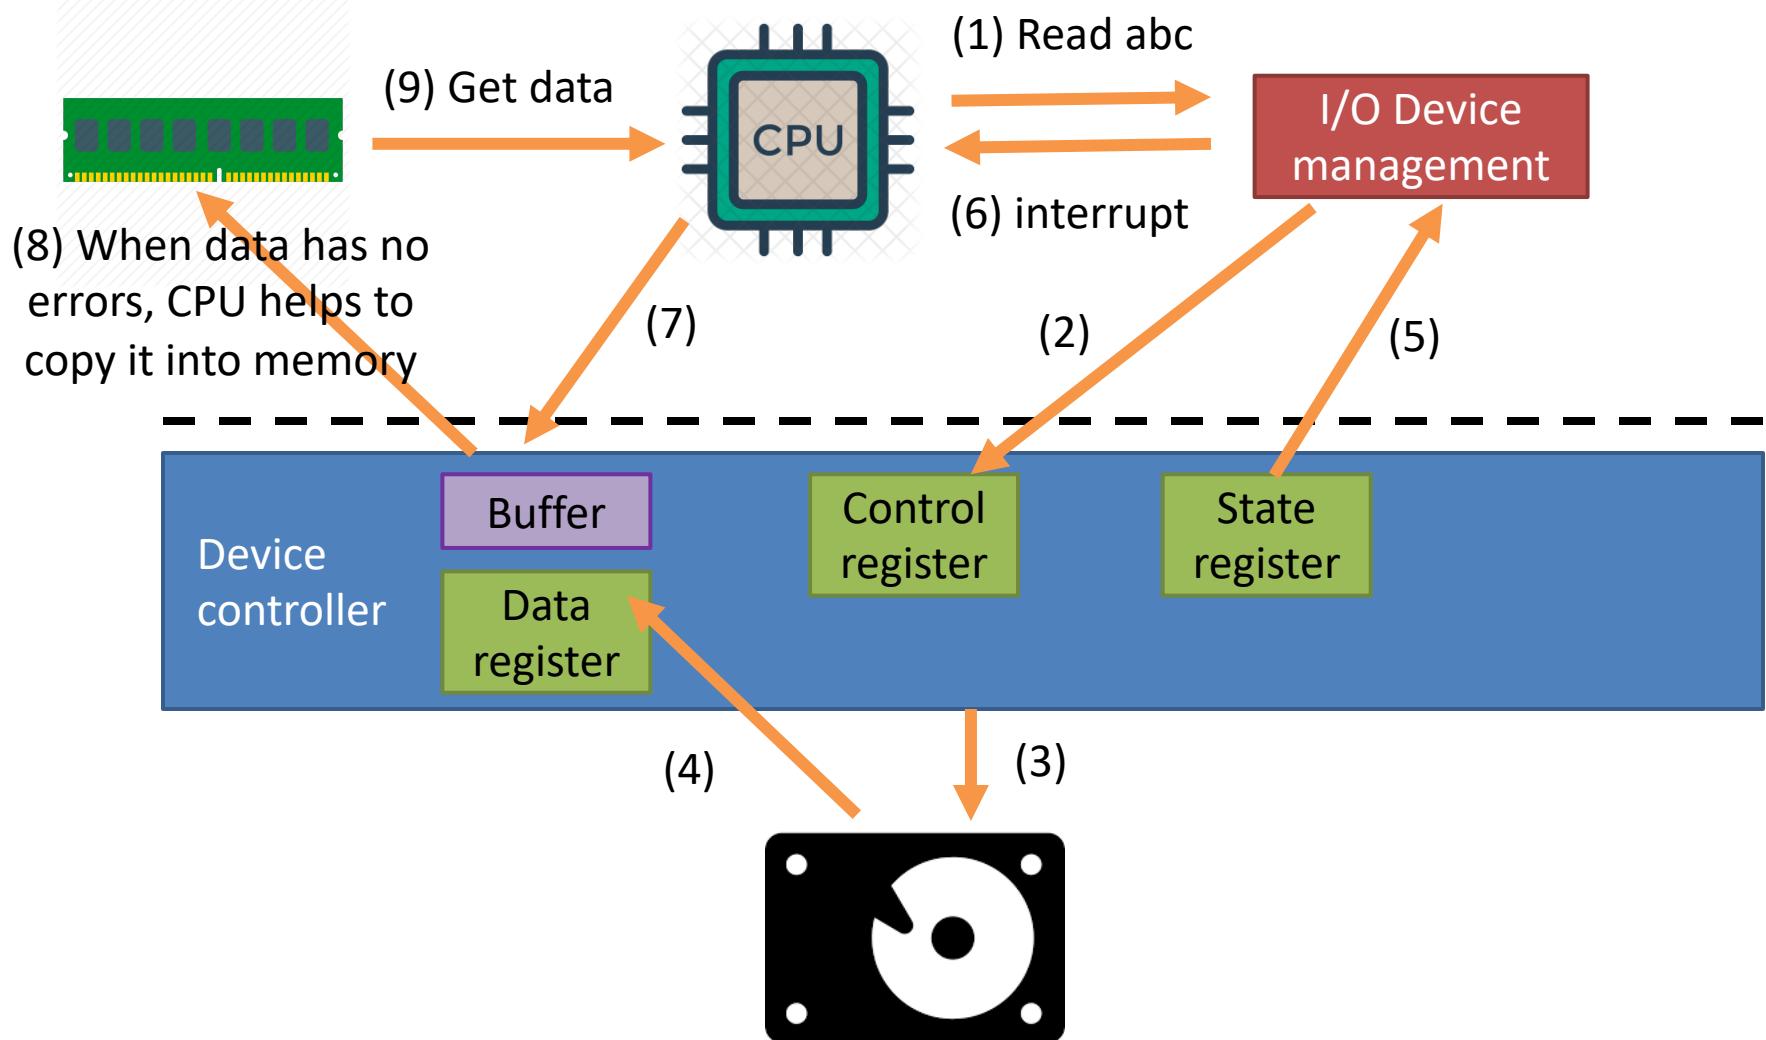
(2)



(3)



# How does the CPU(OS) communicate with the I/O devices? Interrupt if data is ready



# How does the CPU(OS) communicate with the I/O devices?

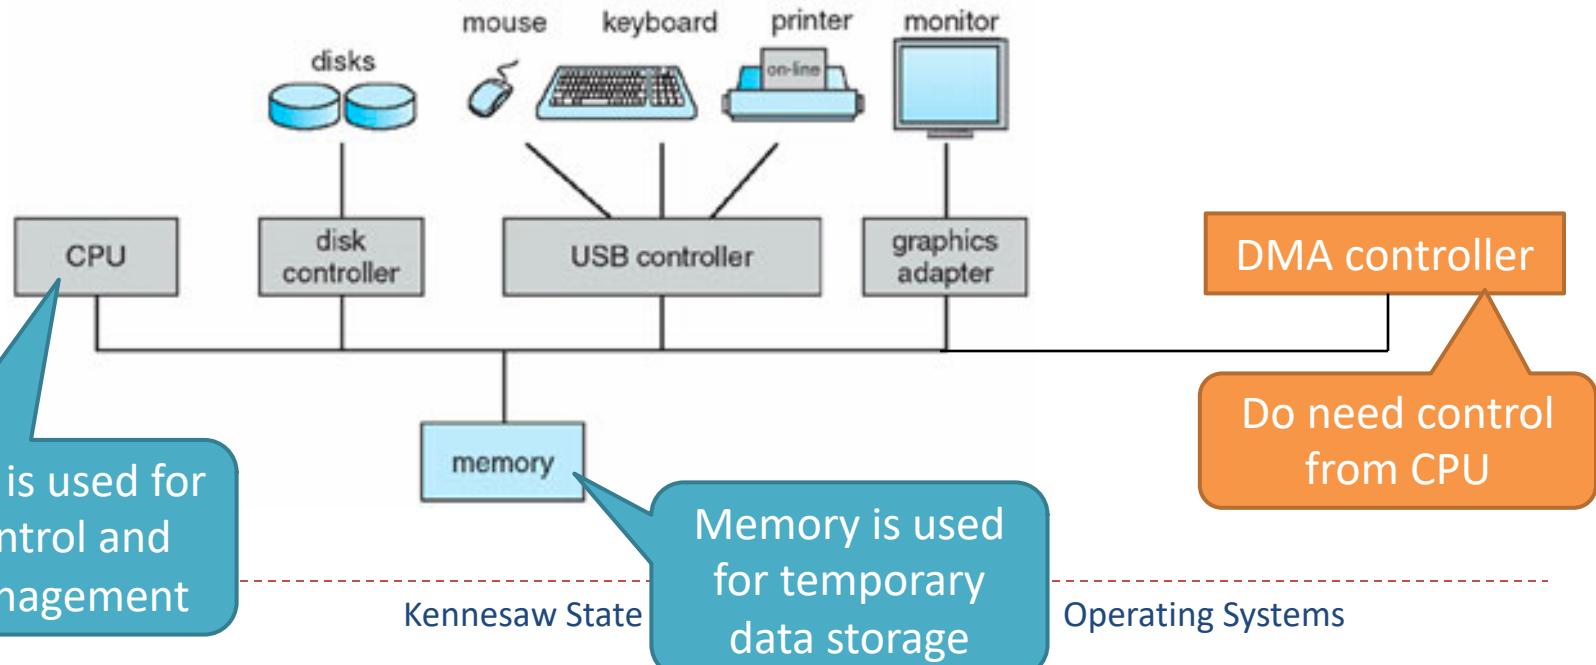
---

- Option 2: I/O Interrupt:
  - Whenever an I/O device needs attention from the processor, it interrupts the processor from what it is currently doing.
- Advantage:
  - User program progress is only halted during actual transfer
- Disadvantage
  - Save the proper states to resume after the interrupt (processor)
  - Too many interrupts will cause lots of context switch overhead

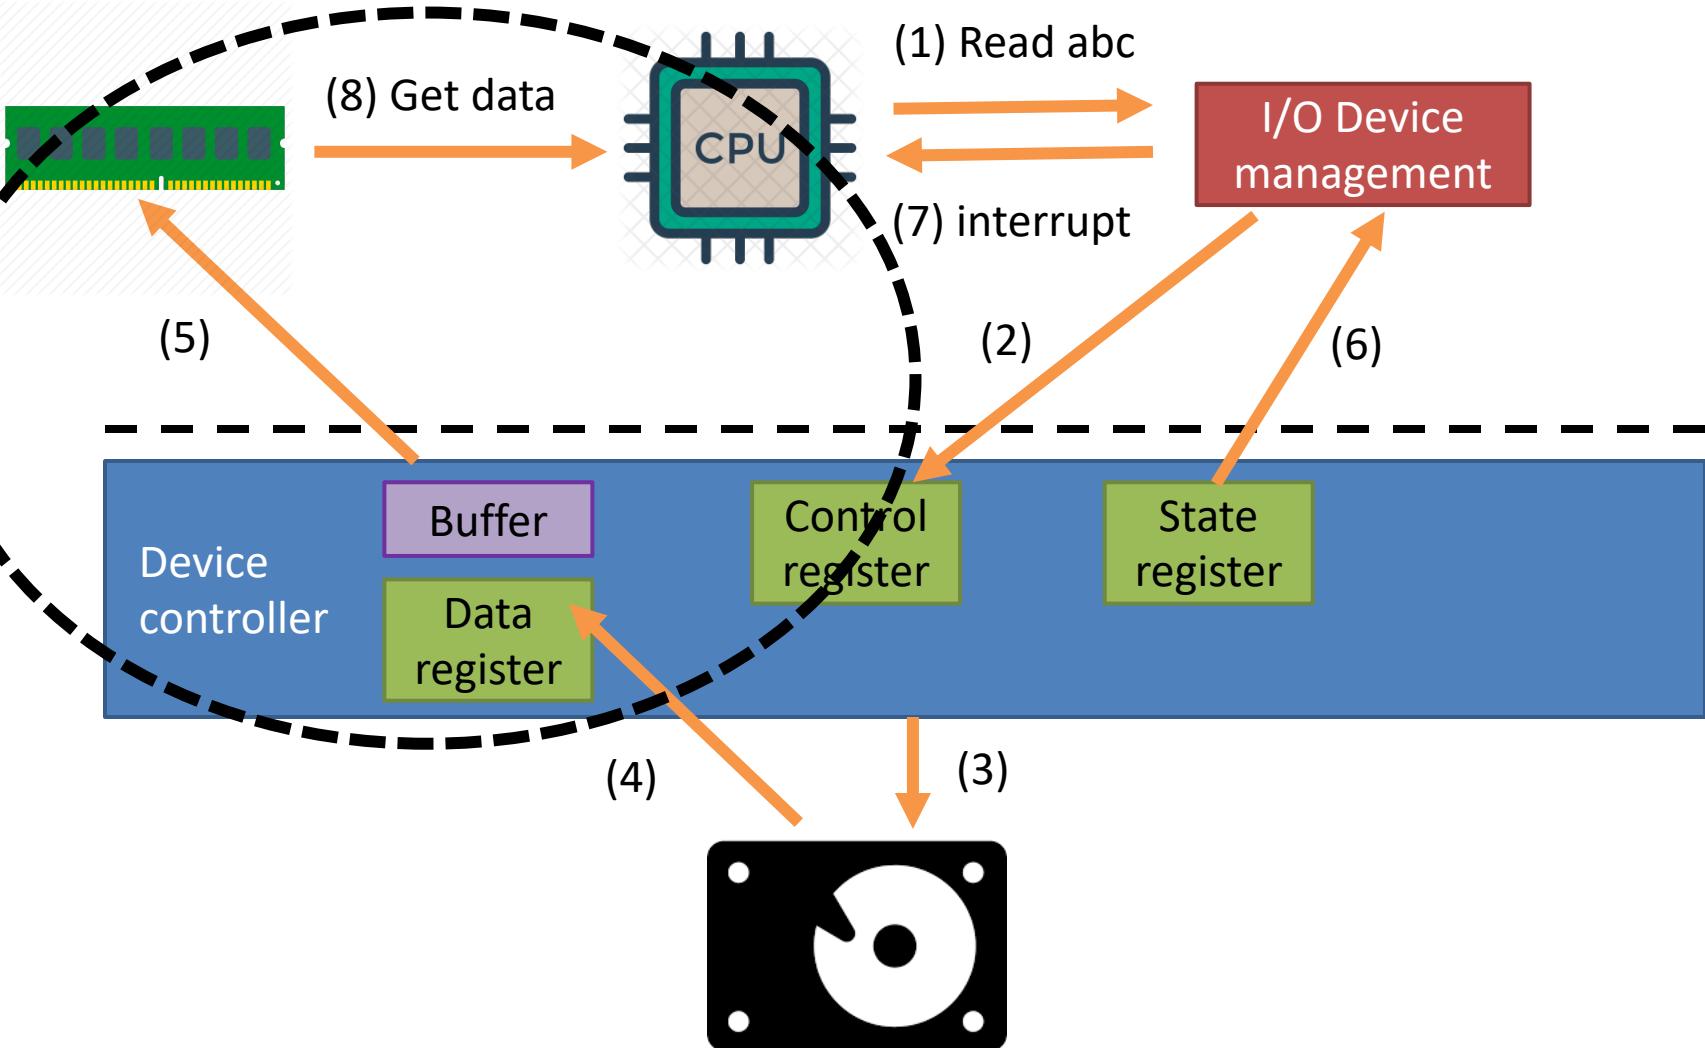


# How does the CPU(OS) communicate with the I/O devices?

- Option 3: Direct Memory Access (DMA):
  - DMA controller has access to system bus independent of CPU
  - The I/O can transfer data to memory directly

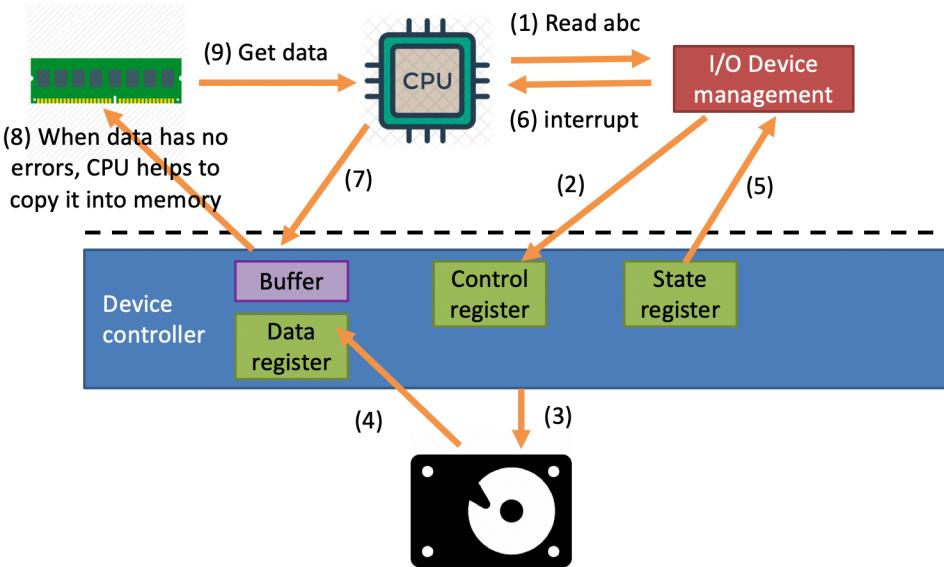


# How does the CPU(OS) communicate with the I/O devices? DMA

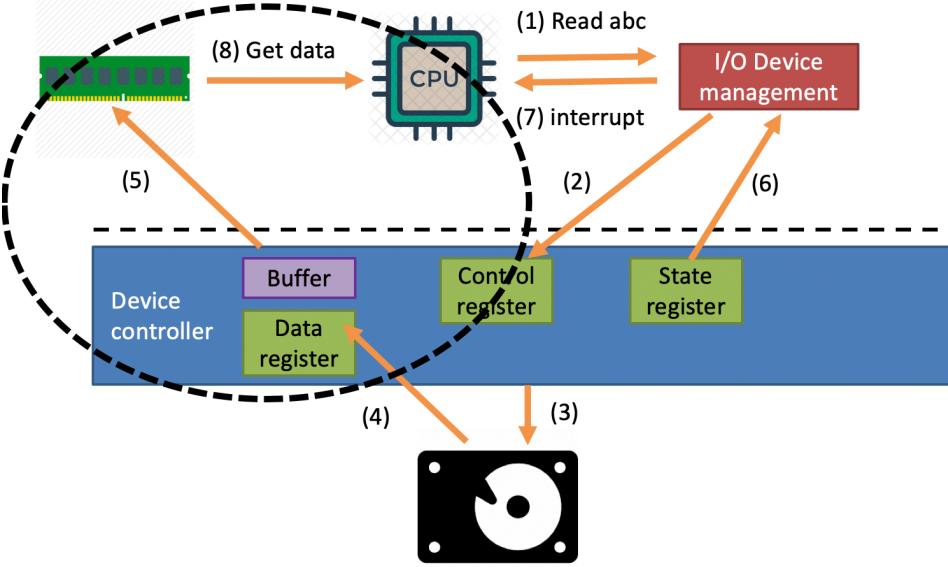


# How does the CPU(OS) communicate with the I/O devices? DMA

Without DMA



With DMA



# How does the CPU(OS) communicate with the I/O devices? DMA

	Without interrupt	With interrupt
Using CPU to transfer data between memory and I/O	Polling I/O	Interrupt I/O
Transfer data between memory and I/O directly	/	DMA



# Outline

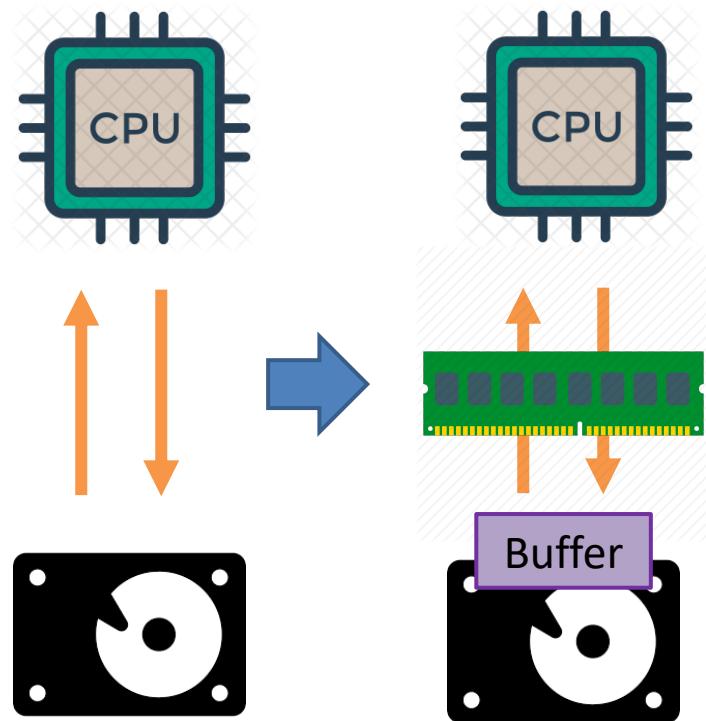
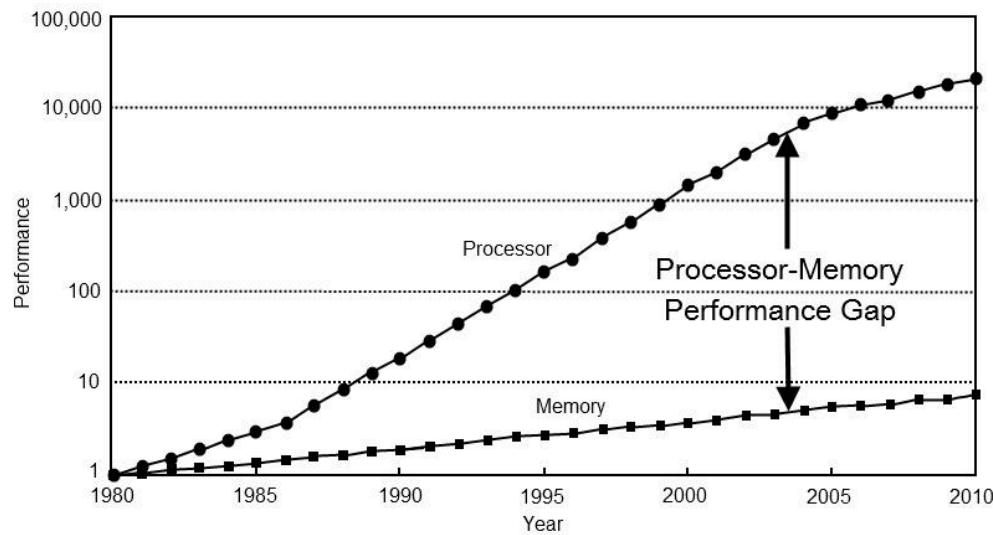
---

- I/O System
- Characteristics and Category of I/O Devices
- I/O Device Component
- How CPU works with I/O devices
- I/O Performance



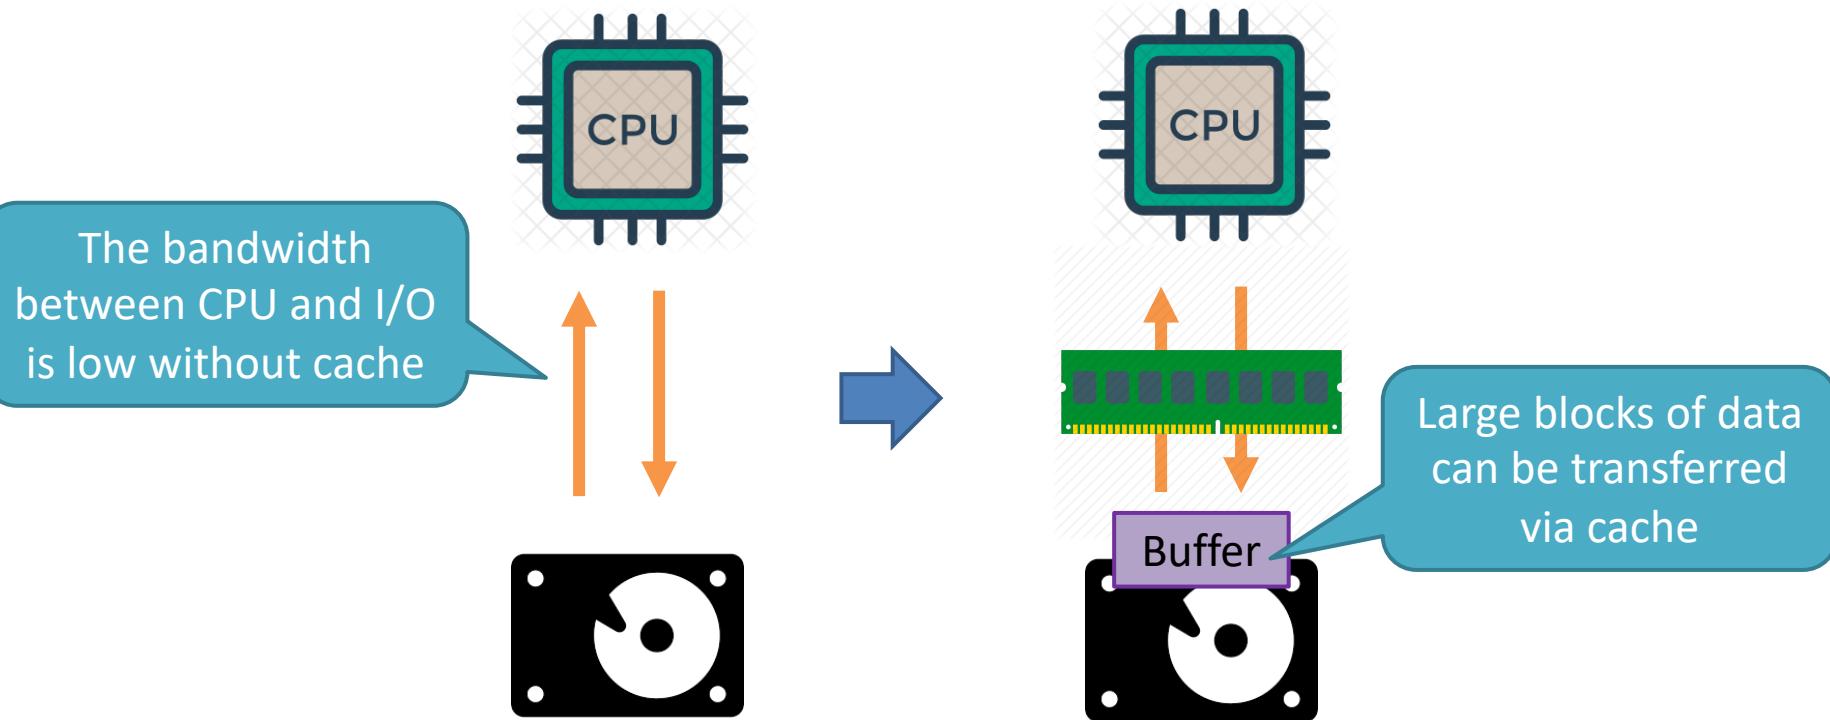
# I/O performance - Buffering

- Bridge the gap between CPU and I/O speed



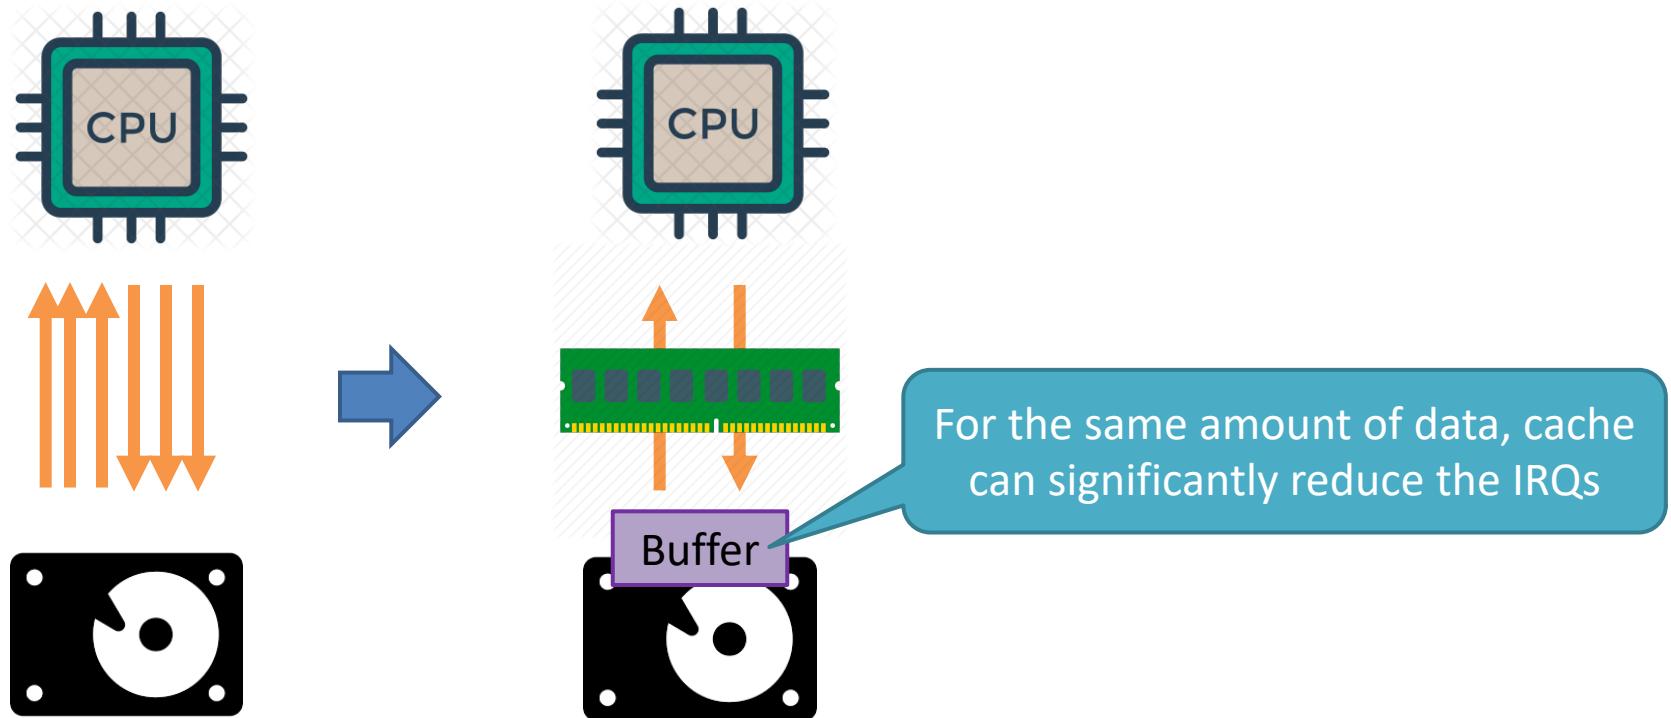
# I/O performance - Buffering

- Increase the parallelism between CPU and I/O



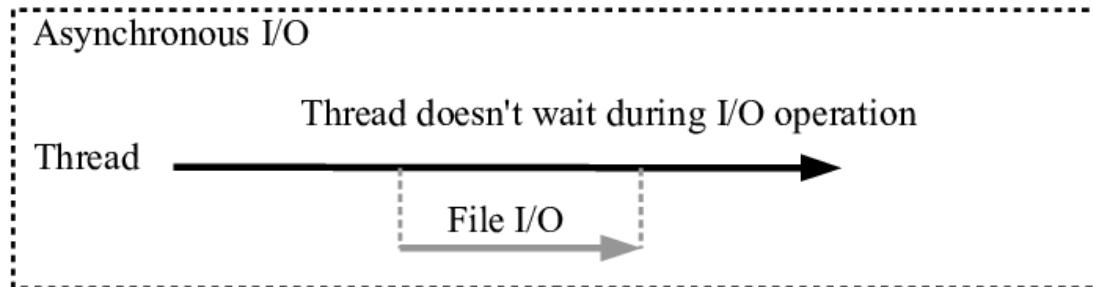
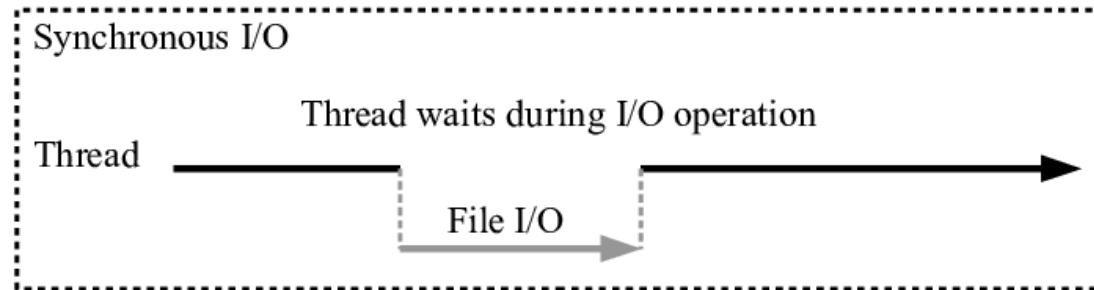
# I/O performance - Buffering

- Reduce the interrupts from I/O to CPU



# I/O performance – synchronous/asynchronous I/O

- **Synchronous I/O:** application is blocked until the I/O operation finishes
- **Asynchronous I/O:** application is not blocked and can do other tasks during the I/O operation is processing



# I/O Performance Summary

---

- I/O execution is complex and slow compared to the CPU/memory
- Ways to improve the I/O performance
  - Bridge the speed gap --> Buffering
  - Do not let CPU wait for I/O --> asynchronous I/O
  - Make I/O access to memory directly without CPU --> DMA



# Conclusion

---

- I/O System
  - Structure and management
- Characteristics and Category of I/O Devices
  - Block vs character, storage vs transmission vs UI
- I/O Device Component
  - Mechanical parts and device controller
- How CPU works with I/O devices
  - Addressing and communication (poll, interrupt, DMA)
- I/O Performance
  - Buffering, asynchronous I/O, DMA

