CS 7172
Parallel and Distributed Computation

Scheduling

Kun Suo

Computer Science, Kennesaw State University

https://kevinsuo.github.io/

Outline

- What is distributed scheduling?

- Monolithic scheduler

  o Borg

- Two-level scheduler

  o Mesos

- Shared state scheduler

  o Omega

Workloads in Datacenters (e.g., Uber)

- Stateless jobs: no state, long running services

- Stateful jobs: long running services but need local disk for storing their state. Such as Cassandra, MySQL, and Redis

**Stateless Jobs**   **Batch Jobs**   **Stateful Jobs**   **Daemon Jobs**
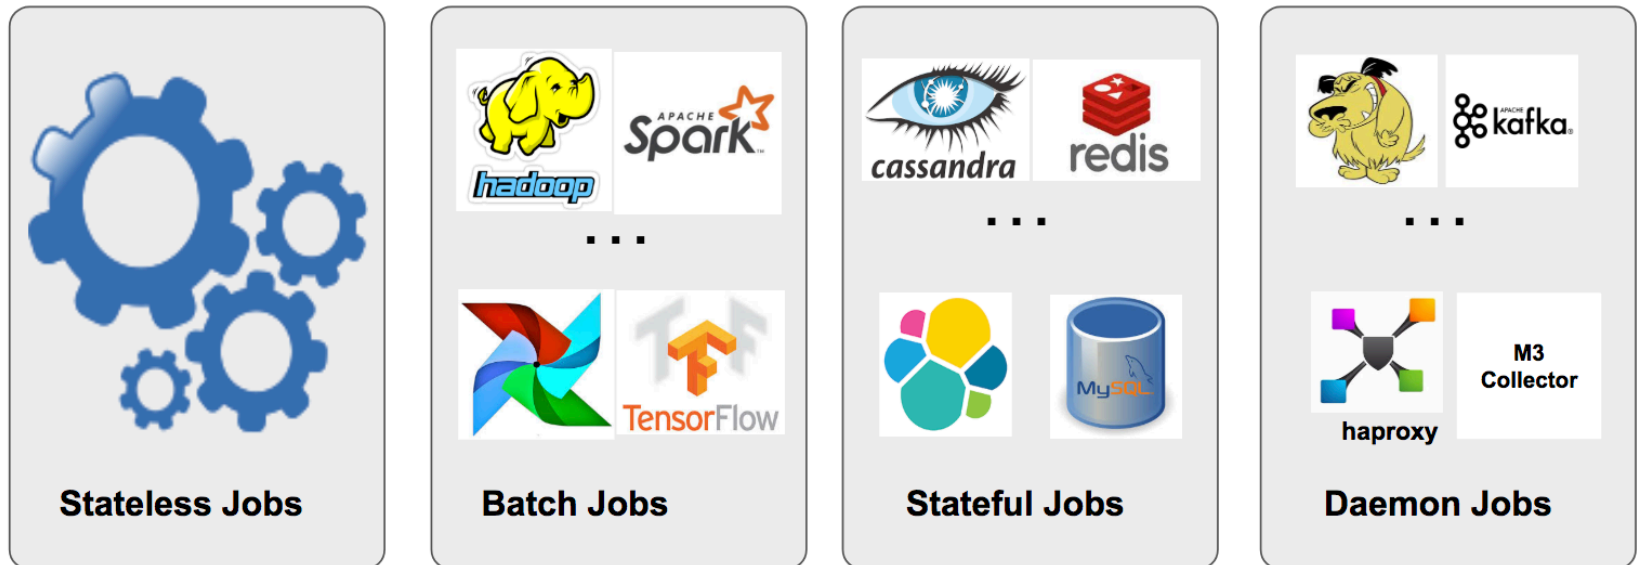
Workloads in Datacenters (e.g., Uber)

- Batch jobs: running minutes to days, not latency-sensitive, can be preempted by other jobs. Such as data analytics, machine learning, etc.

- Daemon jobs: running on each node of the cluster. Such as Kafka, HAProxy and M3 collector.

Stateless Jobs

Batch Jobs

Stateful Jobs

Daemon Jobs

Distributed Scheduling

- Distributed system architecture:
    - Centralized structure
    - Decentralized structure

- The purpose of the distributed system architecture is to manage multiple server resources and find the appropriate server to perform user tasks.

allel and Distributed Computation

Distributed Scheduling

- What is the appropriate server? Many constraints:

  o Task priority

  o Resource availability

  o Load balance

  o ...

- The process of finding the appropriate server in distributed systems for user tasks is called *distributed scheduling*

Distributed Scheduling

- The scheduler is one of the most important components in distributed system. The scheduler provides multiple scheduling strategies (FIFO, SJF, etc.) and is responsible for completing specific scheduling tasks.
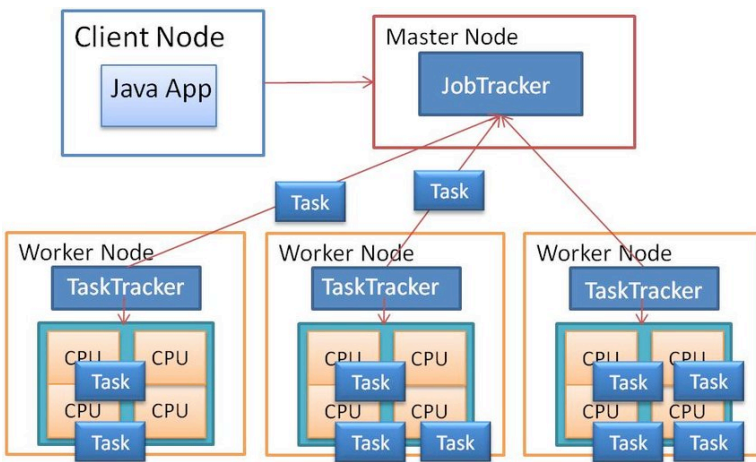
Hadoop

TABLE I
JOB 1 AND 2 SUBMISSION INFORMATION.

| Jobs	Input	CPU demand	Arrival time	Deadline
J1	18 GB	120 GHz	0_{th} min	40_{th} min
J2	9 GB	60 GHz	10_{th} min	30_{th} min

(a) Resource availability.

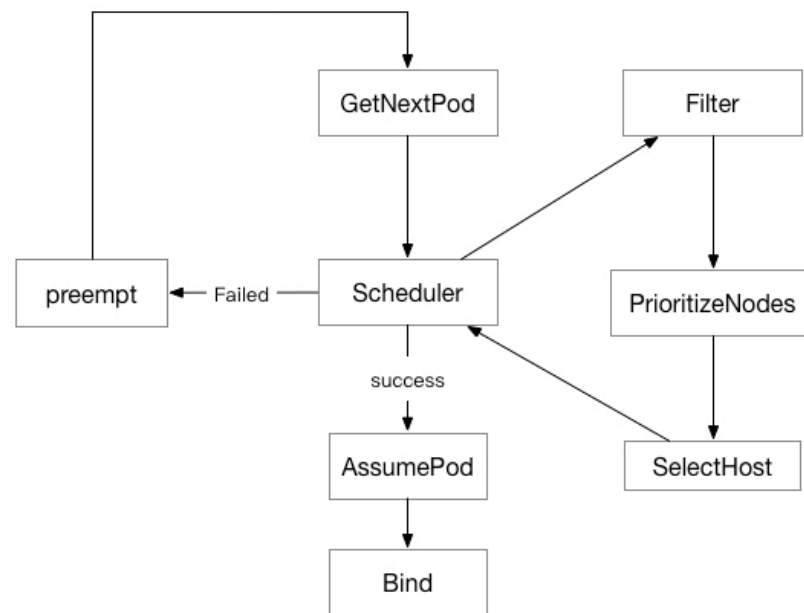(b) FIFO Scheduler.

(c) Fair Scheduler.

Distributed Scheduling

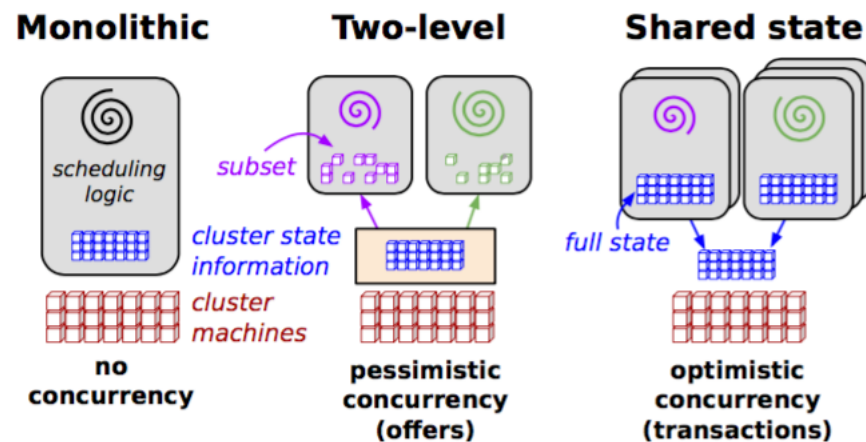- The scheduler principle of different distributed architectures is different

Kubernetes


Hadoop

Distributed Scheduling

- The most common or intuitive scheduler is the monolithic scheduler, which matches the *user task* with the *idle resources* in the distributed systems.

- Monolithic scheduler manages both tasks and resources

What is Monolithic Scheduler?

- Only one node in a cluster runs the scheduling process.

- This node can collect other node's resource information and status for unified management. According to the resource requirements of tasks, the scheduler matches tasks with available resources.

Monolithic

scheduling logic

cluster state information

cluster machines

no concurrency

- The monolithic scheduler has a global view on resource and tasks, which can easily implement task constraints and global scheduling strategies

Monolithic Scheduler Example

- The scheduling process runs on the master node (responsible for resource management, tasks, and resource matching)

- Slave nodes report "Node state" to the Cluster State on master node, which manages the resources and states of nodes in the cluster

Monolithic Scheduler Example

- Cluster State reports resource states to Scheduling Logic module, which is responsible for matching between tasks and resources

- After scheduling decision is made, the task will be sent to corresponding node in the cluster

Monolithic Scheduler

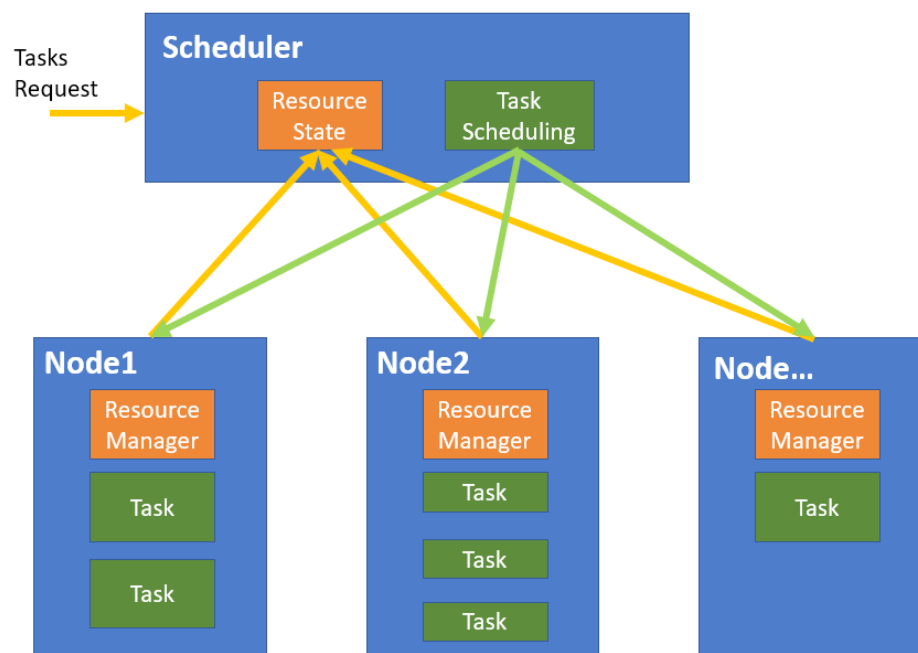- Monolithic scheduler is also called centralized scheduler, which manages the *resources* and *tasks* on a single node
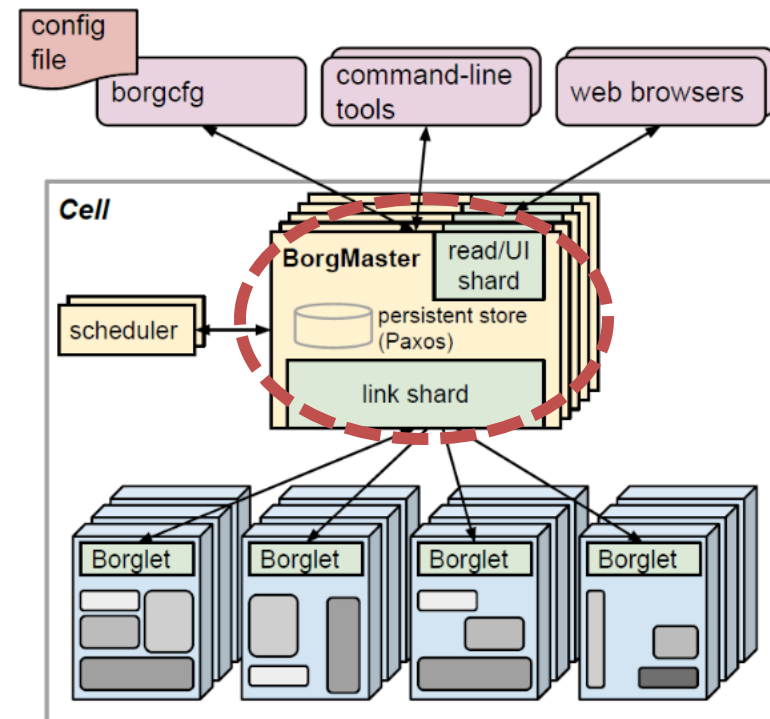
How monolithic scheduler works? Using Borg as an example

- Monolithic scheduler: find a match between tasks and resources

- Definition of jobs:
  - A job is composed of batches of tasks
  - A job can only execute on a cluster

- Definition of tasks
  - A task is composed of batches of processes
  - A task executes on a node in the cluster

How monolithic scheduler works? Using Borg as an example

1. When a user submits a job to BorgMaster, it saves the job to the Paxos warehouse and add all tasks of this job to the waiting queue.

2. The scheduler scans the waiting queue asynchronously and assigns tasks to computing nodes that meet constraints and have enough resources. Here the scheduling unit is task not job.

How monolithic scheduler works? Using Borg as an example

- Feasibility check:

  o the scheduler will find a set of machines that meet the task constraints and have enough available resources

  o Suppose Job A requires to be deployed to Node 1, 3 and 5 and consumes 0.5 CPU and 2MB memory. Scheduler will first filter Node 1, 3, 5 and then select one node has 0.5 CPU and 2MB memory.

How monolithic scheduler works? Using Borg as an example

- Rating:

  - Borg scores the selected machines in the feasibility check stage according to a scoring mechanism and selects the one that is most suitable for scheduling

  - Common scoring algorithms, including "worst fit" and "best fit"

How monolithic scheduler works? Using Borg as an example

- Worst fit:

  - distribute tasks across as many as possible different machines

  - Possible problem: it results in a small amount of unusable remaining resources for each machine. Resource fragmentation.

T1:
0.4 CPU
0.3G memory

T2:
0.3 CPU
0.5G memory

A:
1 CPU
1G memory

B:
0.8 CPU
1.2G memory

How monolithic scheduler works? Using Borg as an example

- Best fit:
  - distribute tasks across as less as possible different machines

  - T1 and T2 will be put onto the same server

  - Possible problem: is not good for applications with sudden loads and requiring immediate handling. Everything on one server causes single node failure and low reliability.

T1:
0.4 CPU
0.3G memory

T2:
0.3 CPU
0.5G memory

A:
1 CPU
1G memory

B:
0.8 CPU
1.2G memory

How monolithic scheduler works? Using Borg as an example

- For scenarios where resources are limited, and business traffic is relatively regular, and there are basically no sudden traffic, the best fit works better;

- If resources are abundant and business traffic often occurs unexpectedly, the worst fit algorithm should be used.

T1:
0.4 CPU
0.3G memory

T2:
0.3 CPU
0.5G memory

A:
1 CPU
1G memory

B:
0.8 CPU
1.2G memory

How monolithic scheduler works? Using Borg as an example

- Besides best fit and worst fit, other possible scheduling policies include:
    - Priority-based
    - Locality-based
    - …

T1:
0.4 CPU
0.3G memory

T2:
0.3 CPU
0.5G memory

A:
1 CPU
1G memory

B:
0.8 CPU
1.2G memory

Monolithic Scheduling

- ## Possible issues:

  o The central server can easily become a single point of bottleneck, which will directly lead to the limitation of the scale

  o Low reliability, single node failure

  o Limited to support jobs with different characteristics (e.g., supporting two policies at the same time? Batching jobs and latency-sensitive jobs at the same time?)

Monolithic Scheduling

- Possible issues:
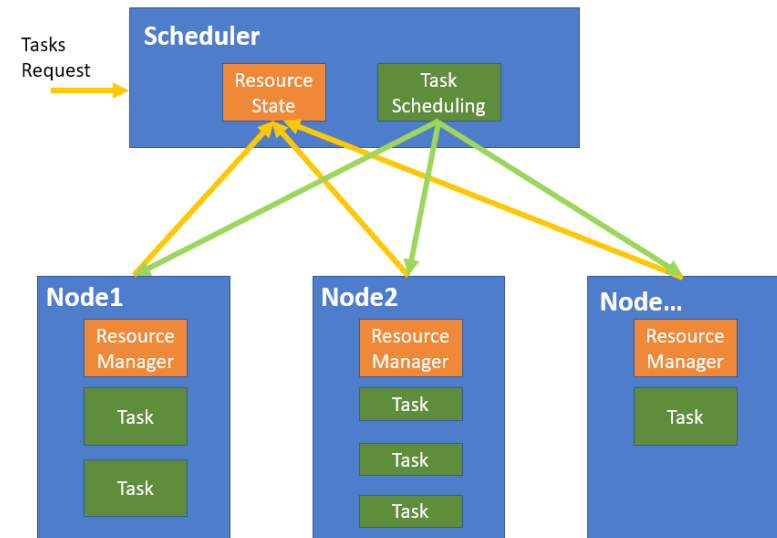 - Limited to support jobs with different characteristics
    - ▸ E.g., supporting two policies at the same time?
    - ▸ E.g., batching jobs (static content) and streaming jobs (dynamic content) at the same time?
    - ▸ Monolithic scheduling becomes more complex as the category of jobs increases.

Tasks Request →

Scheduler

| Resource State | Task Scheduling |

Node1

| Resource Manager |

| Task |

| Task |

Node2

| Resource Manager |

| Task |

| Task |

| Task |

Node...

| Resource Manager |

| Task |

Parallel and Distributed Computation

Two-level Scheduling

- Two-level Scheduling: separate the resource scheduling and task scheduling

Two-level Scheduling

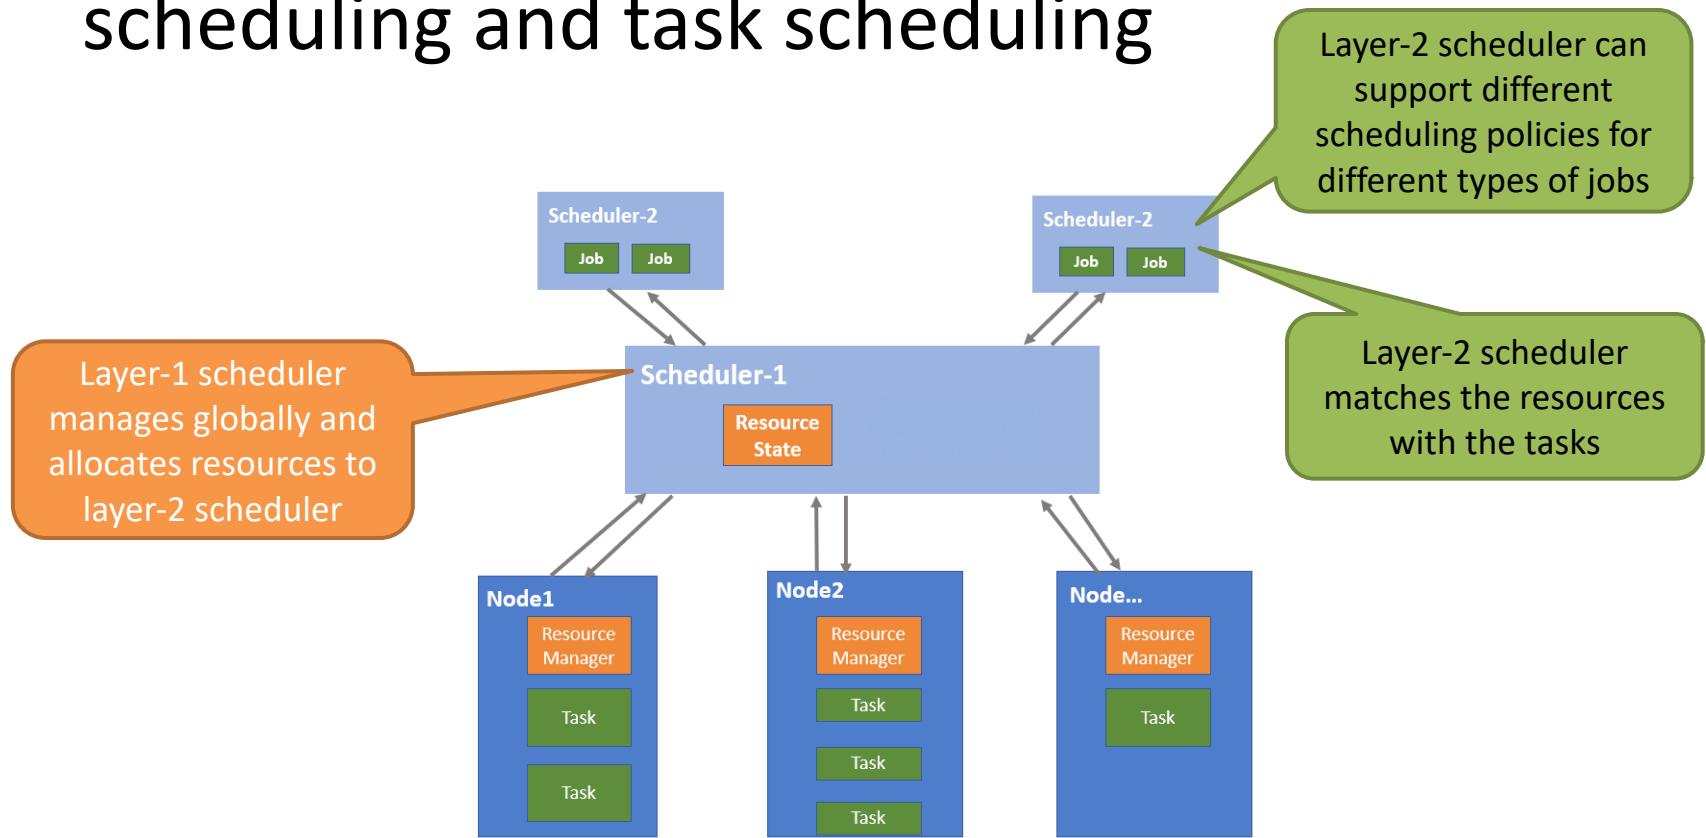- Two-level Scheduling: separate the resource scheduling and task scheduling

Layer-2 scheduler can support different scheduling policies for different types of jobs

Layer-2 scheduler matches the resources with the tasks

Layer-1 scheduler manages globally and allocates resources to layer-2 scheduler

Scheduler-2

Job   Job

Scheduler-2

Job   Job

Scheduler-1

Resource State

Node1

Resource Manager

Task

Task

Node2

Resource Manager

Task

Task

Task

Node...

Resource Manager

Task

Example of Two-level Scheduling: Mesos

- ## Two-level Scheduling

 - Mesos, first-level scheduling, is only responsible for the management and allocation of the underlying resources, and does not involve functions such as storage and task scheduling

  - The second-level task scheduling is through the framework, such as Hadoop, Spark, etc., to complete
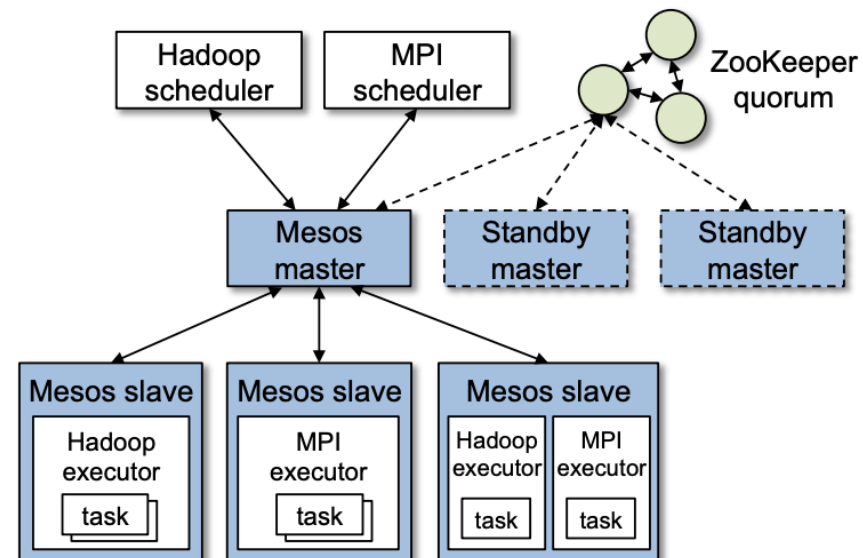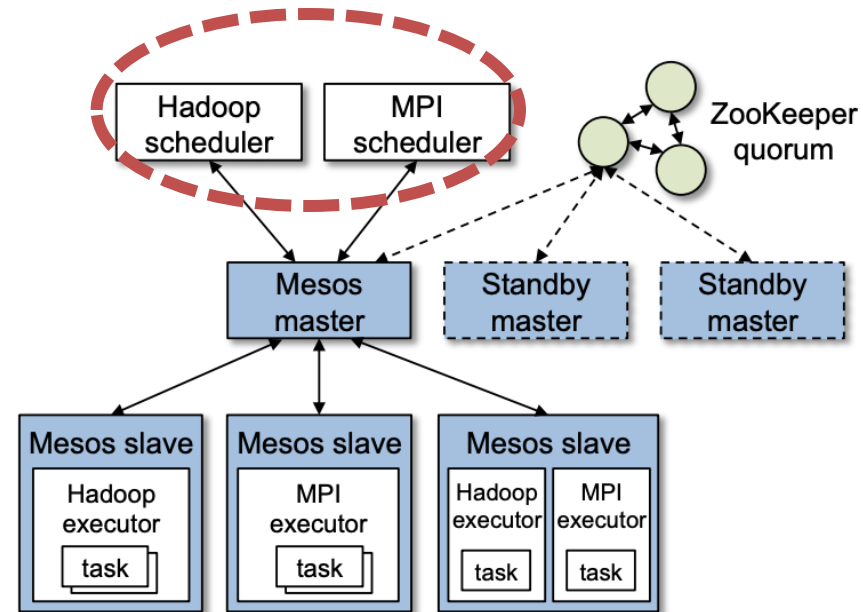
Example of Two-level Scheduling: Mesos

- First-level Scheduling

  o The resource management cluster is a centralized system composed of a master node and multiple slave nodes.

  o The cluster has only one master node, which is responsible for managing slave nodes and connecting to the upper-level framework

  o The slave nodes periodically report resource status information to the master node and execute tasks submitted by the framework.
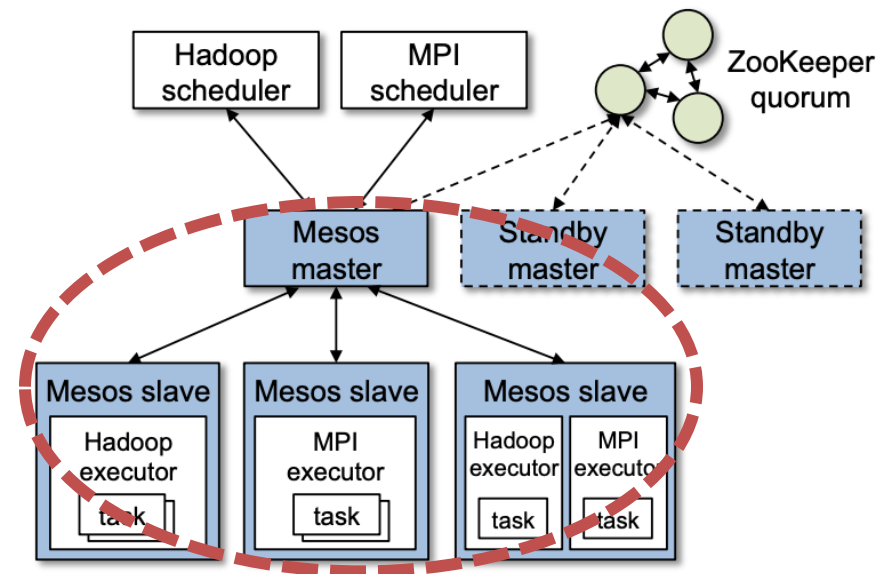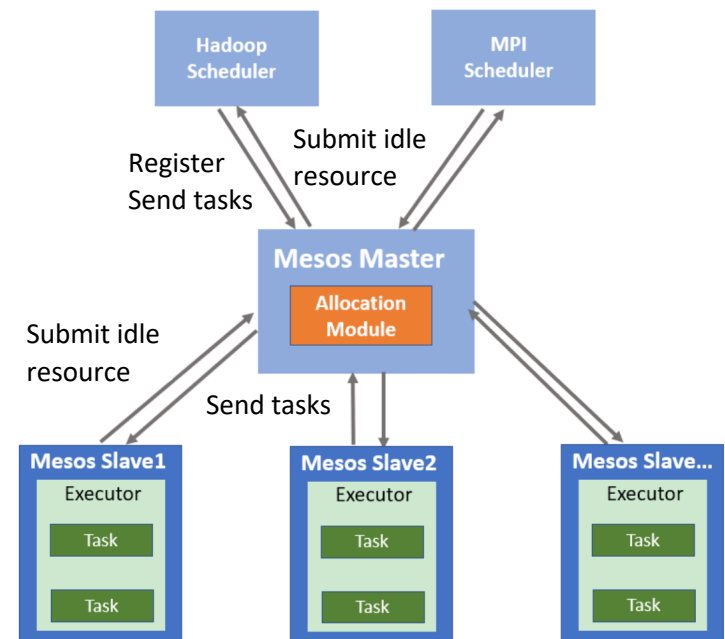
Example of Two-level Scheduling: Mesos

- ## Second-level Scheduling

  - o Frameworks, such as Hadoop, Spark, MPI, and Marathon, run on Mesos and are "components" responsible for application management and scheduling.

  - o Different frameworks are used to complete various tasks, such as batch tasks, real-time analysis tasks, etc.

Example of Two-level Scheduling: Mesos

- Workflows in Mesos

  1. The framework registers with the Mesos Master;

  2. The Mesos Slave node reports the free resources of the node to the Mesos Master periodically;

  3. The Scheduler process on the Mesos Master collects the free resource information and sends it to the registered framework;

  4. After the framework's Scheduler receives the resources, it performs task scheduling and matching. If matches, the result is sent to the Mesos Master, and then forwarded to the corresponding node's executor to execute the task.

Example of Two-level Scheduling: Mesos

- ## Scheduling policy in frameworks
  - o Hadoop: http://www.corejavaguru.com/bigdata/hadoop-tutorial/yarn-scheduler
    - ▸ FIFO
    - ▸ Capacity
    - ▸ Fair
    - ▸ …
  - o Spark: https://databricks.com/session/apache-spark-scheduler

nputation

Example of Two-level Scheduling: Mesos

- Scheduling algorithm at the first-level:
  - Max-min Fairness, MMF
  - Dominant Resource Fairness, DRF
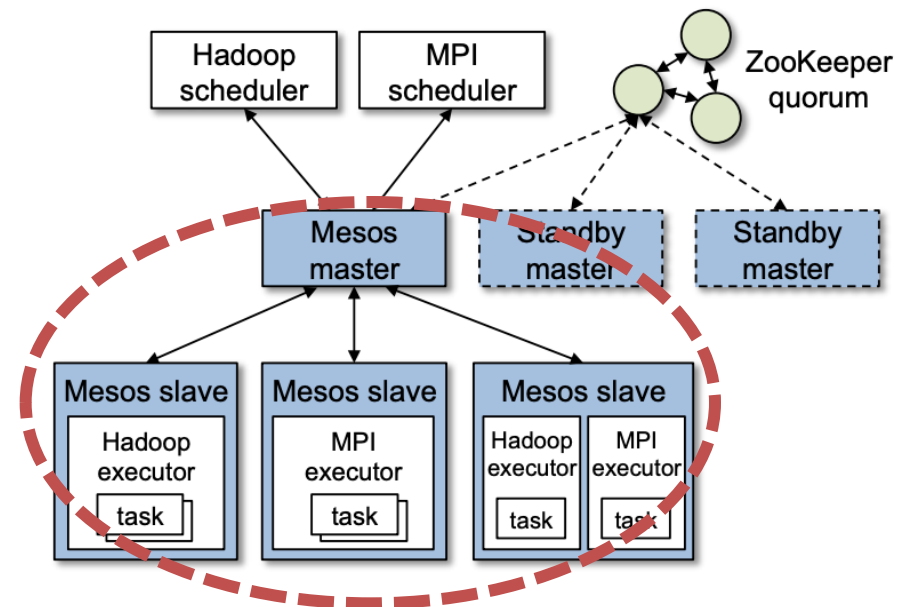
Example of Two-level Scheduling: Mesos

- ## Max-min Fairness algorithm:
  - o Fairly share the minimum requirements that each user needs, and then allocate the unused resources evenly to users who need extra resources.

- ## Example
  - o Total resource: 10
  - o User number: 4
  - o User requirement: 2, 2.6, 4, 5

Final allocation

2
2.6
4
5

2.5
2.5
2.5
2.5

2
2.67
2.67
2.67

0.5/3

2
2.6
2.7
2.7

0.07/2

2
2.6
2.7
2.7
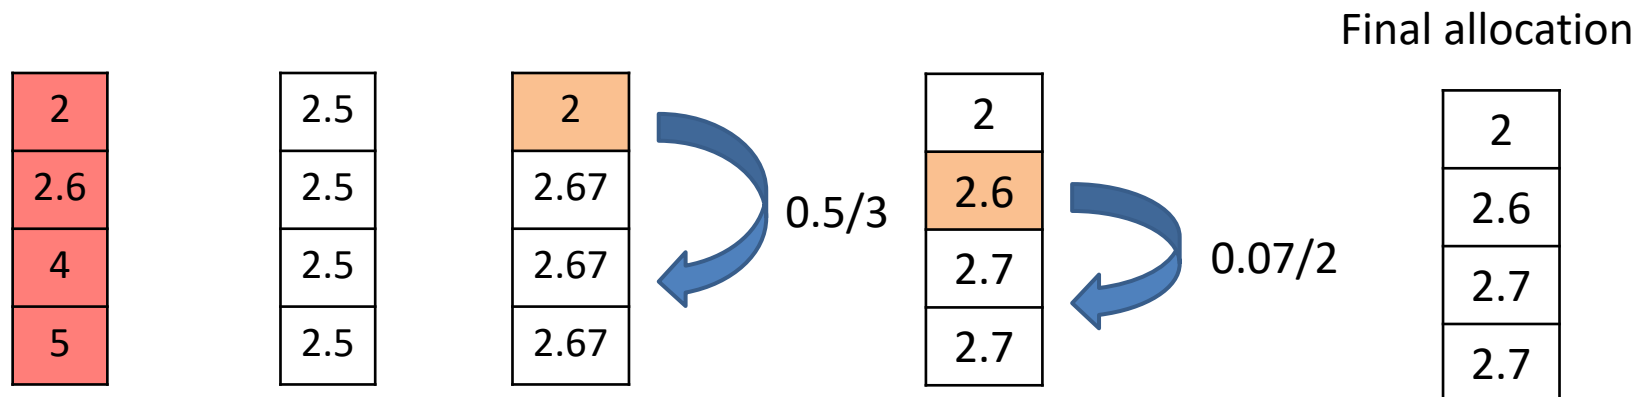
Example of Two-level Scheduling: Mesos

- ## Example of MMF

  - Total idle resource: 100

  - User number: 4

  - User requirement: 35, 10, 25, 45

Final allocation

| 35		25		30			32.5			32.5
10		25		10			10			10
25		25		30		15/3	25		5/2	25
45		25		30			32.5			32.5
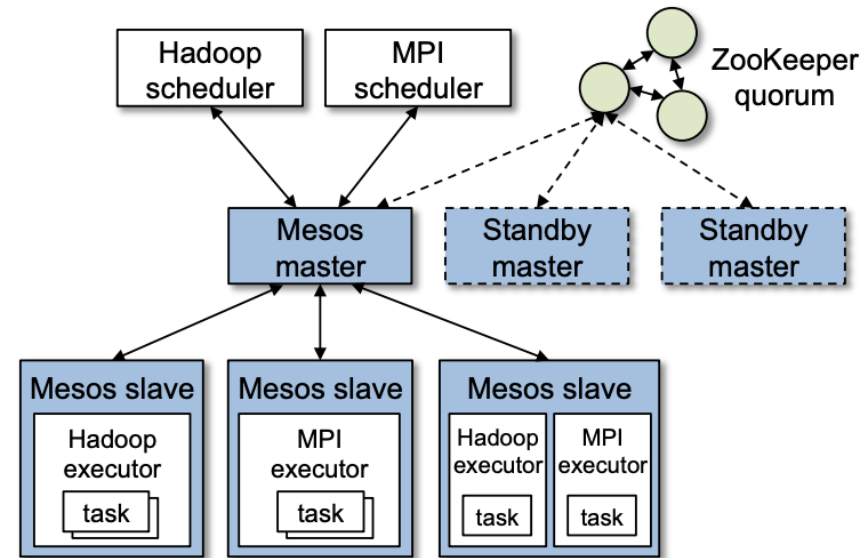
Example of Two-level Scheduling: Mesos

- ## Problems of MMF

  o Allocating resources in an absolutely fair way will cause a lot of waste of resources

    ▸ For example, user A and user D with user demand of 35 and 45 are all allocated 32.5 free resources, but because the resources do not meet the demand, both users cannot use and run tasks.

Final allocation

35
10
25
45

25
25
25
25

30
10
30
30

15/3

32.5
10
25
32.5

5/2

32.5
10
25
32.5

Example of Two-level Scheduling: Mesos

- Dominant Resource Fairness, DRF

  o Besides the fairness among users, it also considers the needs of users for different types of resources

  o Can satisfy as many users as possible

Example of Two-level Scheduling: Mesos

- ## Dominant Resource Fairness Case

  - Suppose we have 18 CPUs and 36GB memory

  - A: memory-bound task, <2 CPU, 8 GB>

  - B: CPU-bound task, <6 CPU, 2 GB>

- ## Step 1: Calculate resource allocation

  - Suppose the number of task A and B are x and y, respectively

  - A needed resources: <2x, 8x>

  - B needed resources: <6y, 2y>

  - Total resources needed: <2x+6y, 8x+2y>

  - Resource constrains: $\begin{array}{l} 2x + 6y \leq 18 \\ 8x + 2y \leq 36 \end{array}$

Example of Two-level Scheduling: Mesos

- Dominant Resource Fairness Case

  o Suppose we have 18 CPUs and 36GB memory

  o A: memory-bound task, <2 CPU, 8 GB>

  o B: CPU-bound task, <6 CPU, 2 GB>

	CPU	Memory
A	2x/18	8y/36
B	6x/18	2y/36

- Step 2: Define the main resource limitation for each task

  o For each A, it needs 2/18 system CPU, 8/36 system memory, so memory limits A's number

  o For each B, it needs 6/18 system CPU, 2/36 system memory, so CPU limits B's number

Example of Two-level Scheduling: Mesos

- Dominant Resource Fairness Case

  o Suppose we have 18 CPUs and 36GB memory

  o A: memory-bound task, <2 CPU, 8 GB>

  o B: CPU-bound task, <6 CPU, 2 GB>

	CPU	Memory
A	2x/18	8x/36
B	6y/18	2y/36

- Step 3: balance the resource utilization of all tasks, maximum the resource usage which does not limit the task number

  o
  $$2x + 6y \leq 18$$
  $$8x + 2y \leq 36$$
  $$\frac{8x}{36} = \frac{6y}{18}$$

  X=3
  Y=2

  A consumes 2/3 of memory resource
 B consumes 2/3 of CPU resource

Example of Two-level Scheduling: Mesos

- ## Max-min Fairness algorithm:

  o Suitable for a single type of resource allocation scenario

  o Assign no more resources than required to each user

- ## Dominant Resource Fairness algorithm

  o Suitable for scenarios where multiple types of resources are mixed

  o Make the most use of resources so that as many tasks as possible can be performed

Two-level Scheduling

- Possible issues:

  o Decouple the task scheduling with the resource availability might make frameworks only know part of the resource information and cannot guarantee global optimal
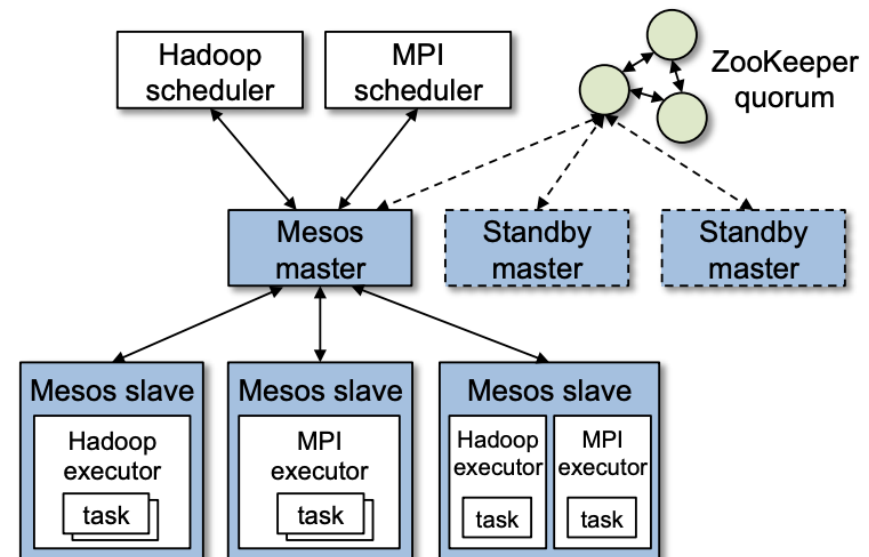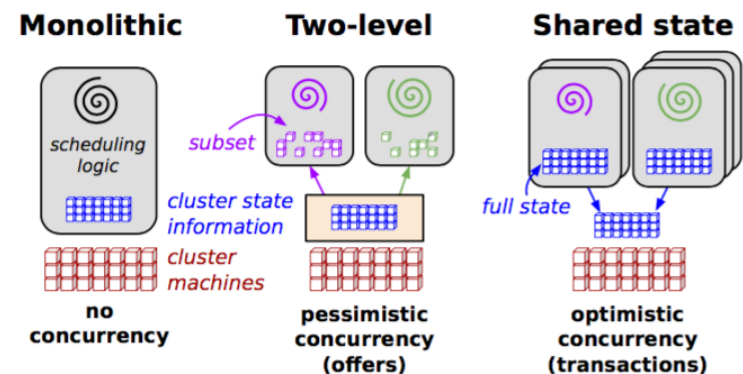
Shared state scheduling -- Why

- Two types of objects to be managed in the cluster:

  1. Resource allocation and usage status;

  2. Task scheduling and execution status;

Monolithic scheduling manages
the two at the same time:

+ Global resource management
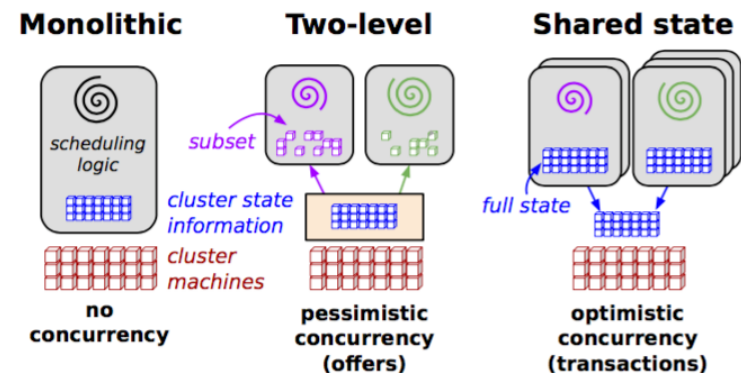
- Poor scalability and single node
bottleneck

Shared state scheduling -- Why

- Two types of objects to be managed in the cluster:

  1. Resource allocation and usage status;

 2. Task scheduling and execution status;

Two-level scheduling manages the
two at different levels:

- Cannot provide optimal resource in
global scope

+ Good scalability

| Monolithic	Two-level	Shared state
scheduling logic	subset — cluster state information	full state
cluster machines		
no concurrency	pessimistic concurrency (offers)	optimistic concurrency (transactions)

Shared state scheduling

- Break down a single scheduler into multiple schedulers

  o --> Better scalability and no single node failure

- Each scheduler has global resource status information

  o --> Optimal task scheduling
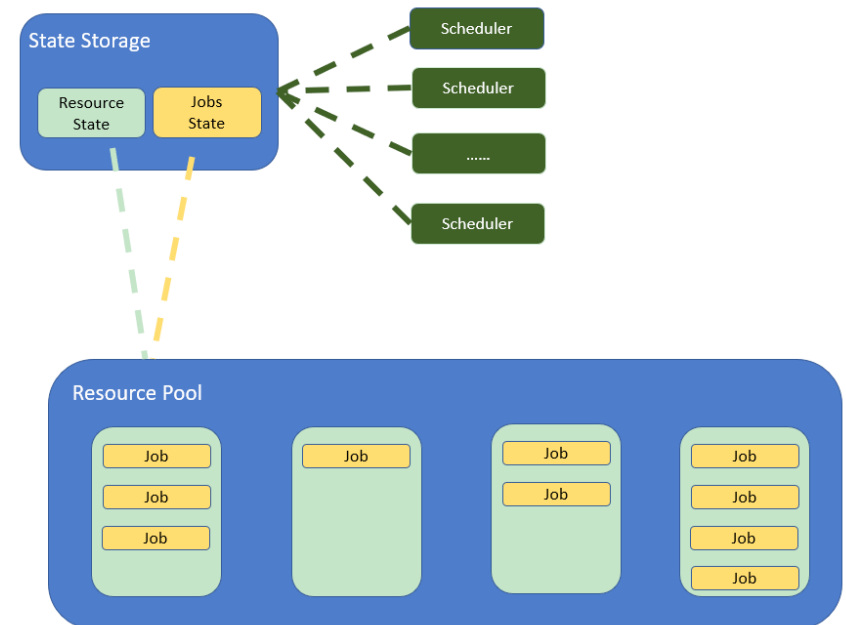
Architecture of Shared State Scheduling

- ## State storage:

  o Responsible for storing and maintaining resource and task state, so Schedulers can query available resource when scheduling tasks
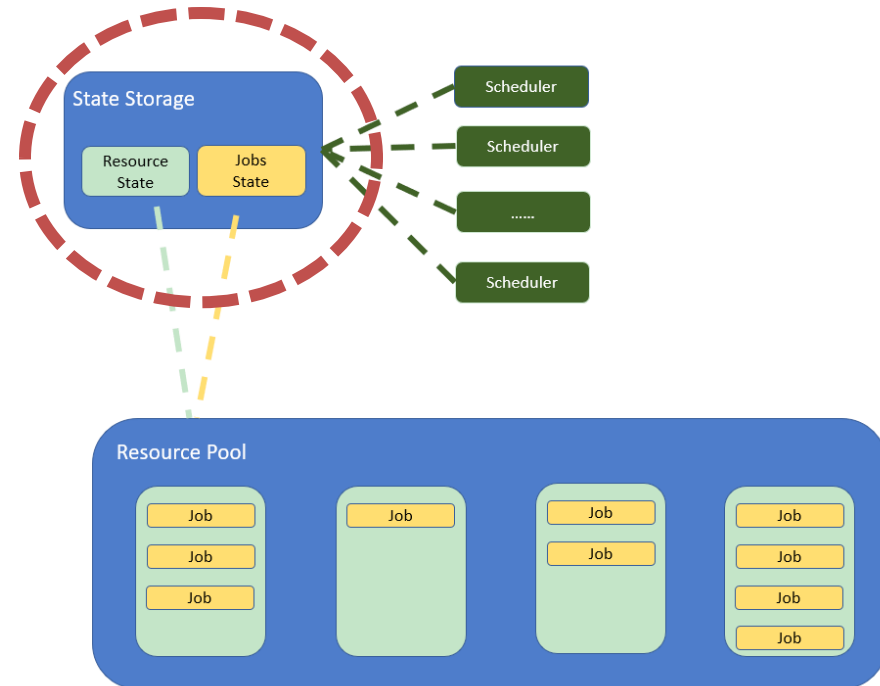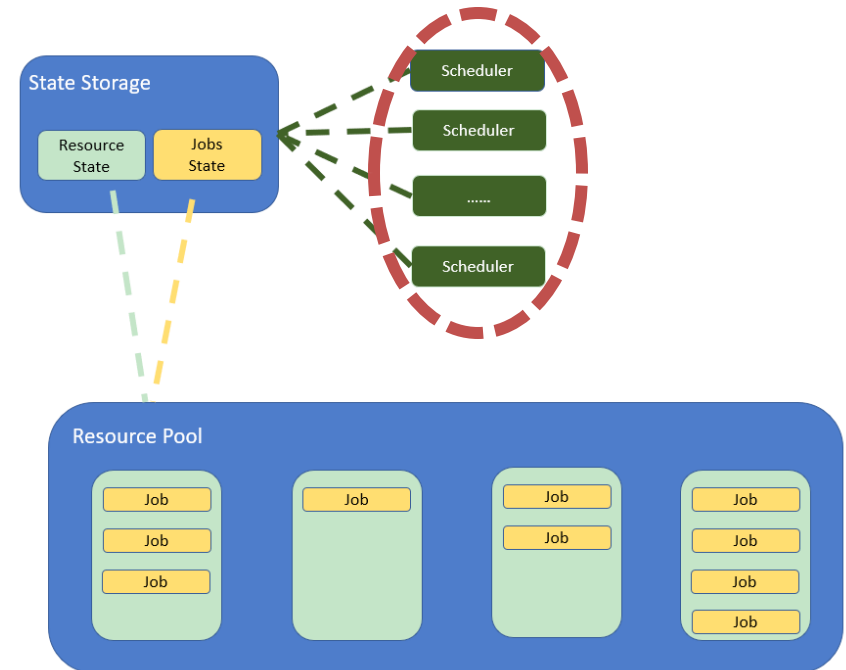
Architecture of Shared State Scheduling

- ## Resource Pool:

  o Multiple node clusters that receive and execute tasks scheduled by the Scheduler; send state of resources and job execution to the Storage
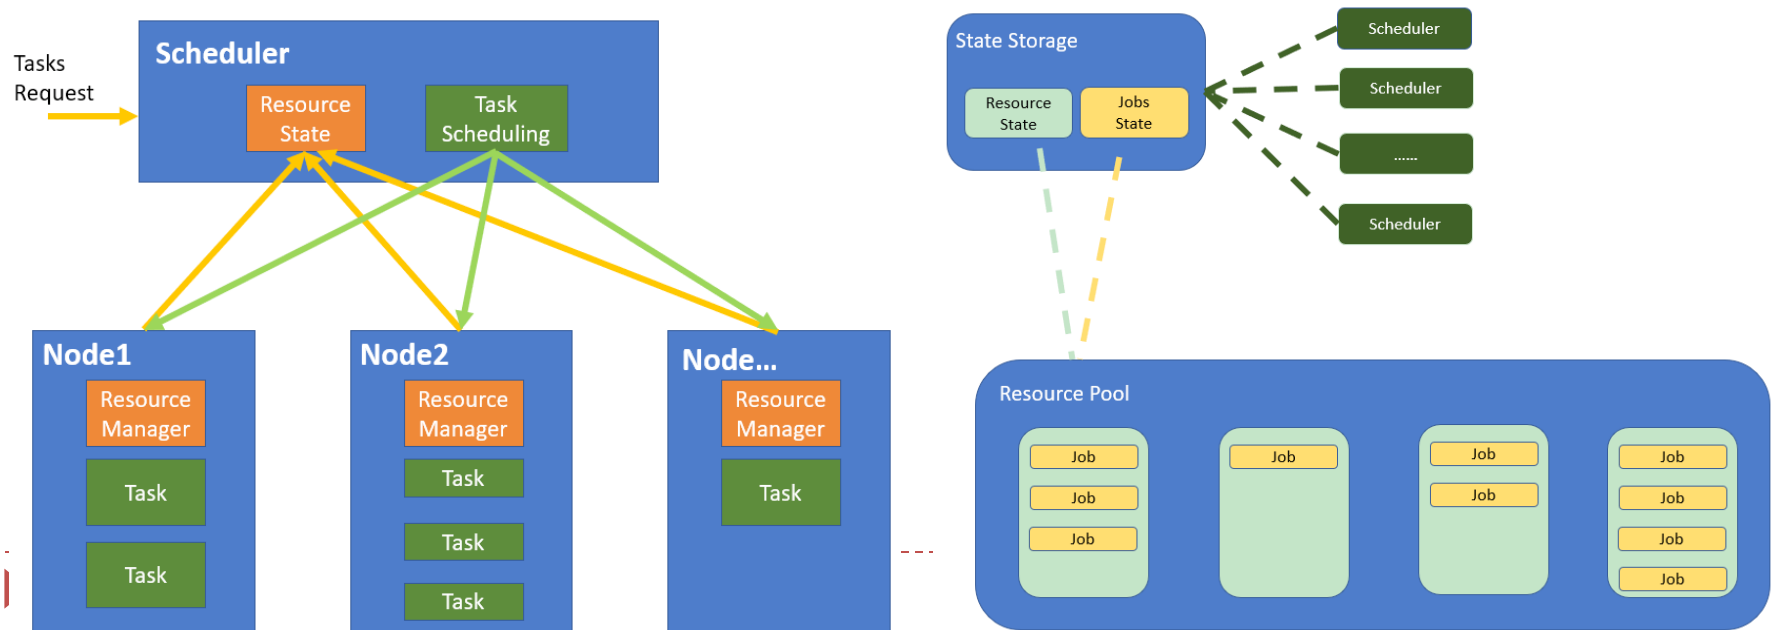
Architecture of Shared State Scheduling

- Scheduler:

 o Just for concrete task scheduling operations

Monolithic vs. Shared State Scheduling

- Different from the monolithic scheduling, the Scheduler in shared state scheduling does not need to manage the cluster resources

Two-level vs. Shared State Scheduling

- All schedulers in shared state scheduling can access the resource information directly and make optimal scheduling on a global view
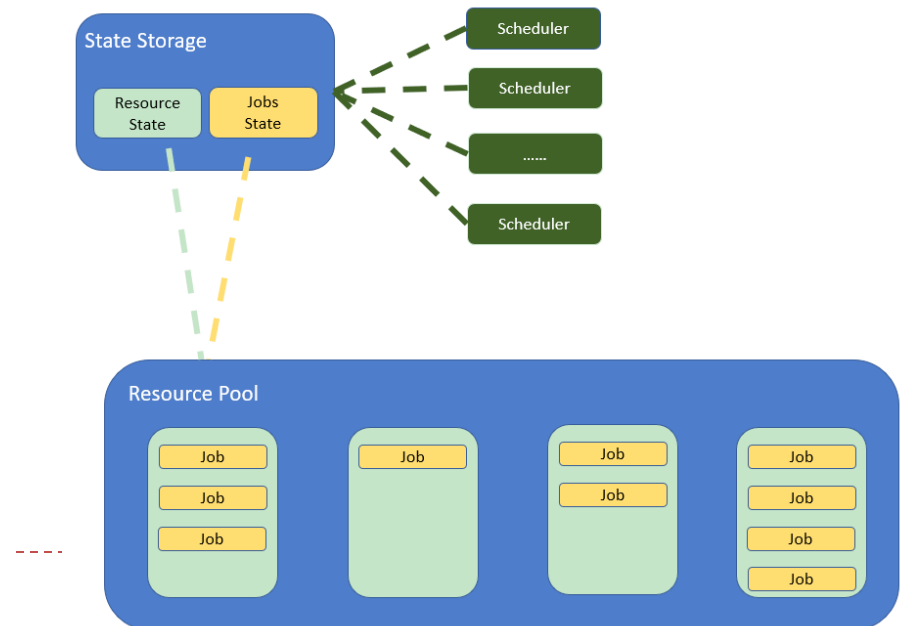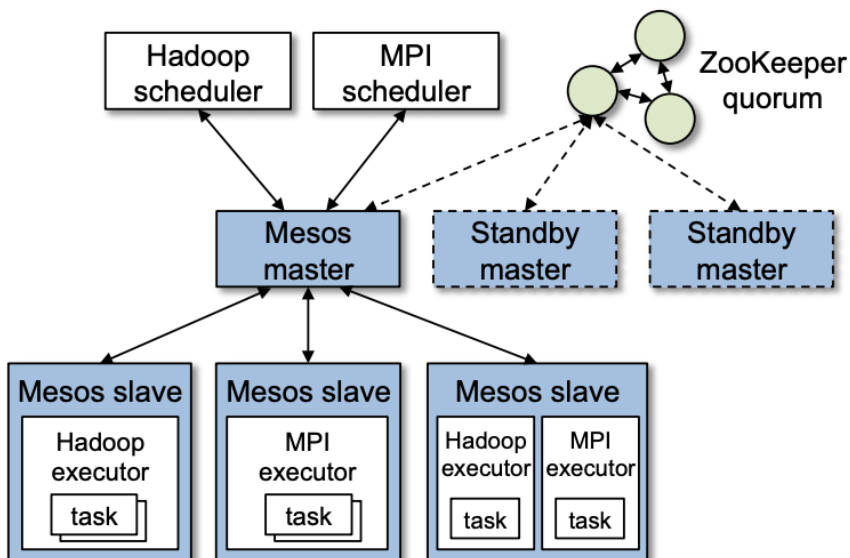
Example of Shared State Scheduling

- Google Omega:
https://storage.googleapis.com/pub-tools-public-publication-data/pdf/41684.pdf

- Microsoft Apollo:
https://www.usenix.org/system/files/conference/osdi14/osdi14-paper-boutin_0.pdf

- Hashicorp Nomad: https://www.nomadproject.io/

Example of Shared State Scheduling: Omega

- Cell: a group of servers
  - A cluster can have multiple cells

- State storage: store and maintain the states of resources and tasks

- Cell State: records the global cluster states

- Every scheduler has one copy of the Cell State → sharing the state of the cluster
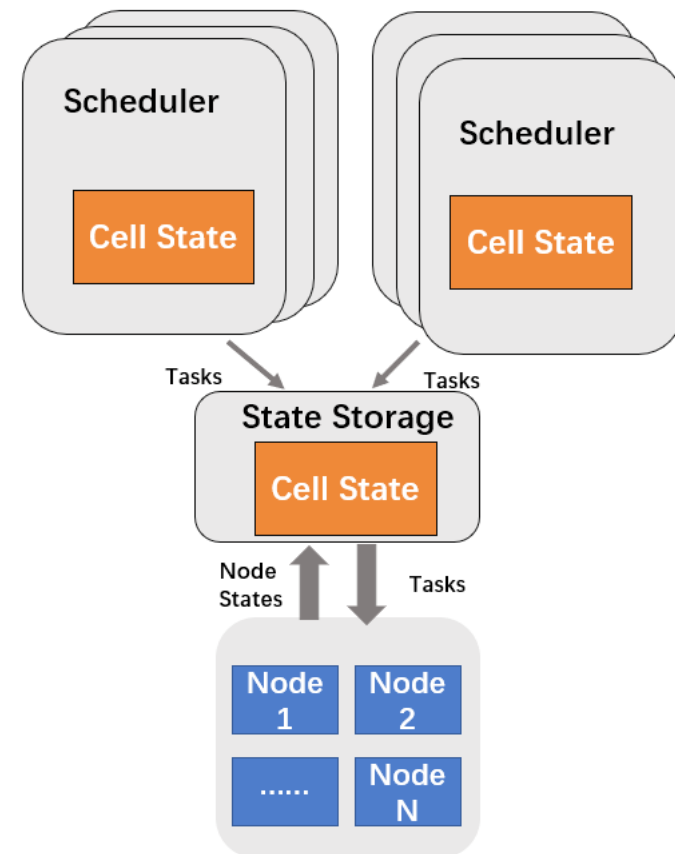
Omega architecture
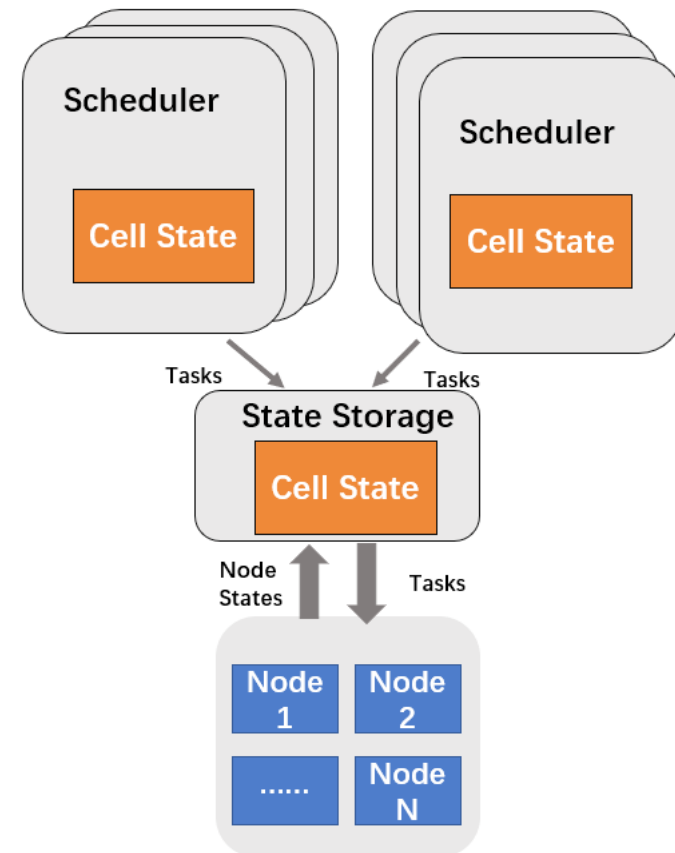
Example of Shared State Scheduling: Omega

- ## What will happen when two scheduler try to use the same resource?

  - o Omega adopts transaction model in database to maintain the consistency and atomicity for the same resource

  - o In database, only one query can update one record at one time

Omega architecture

Example of Shared State Scheduling: Omega

- The scheduler tries to match the tasks and resources
- Scheduler will check whether the resources have been used during the checkpoints
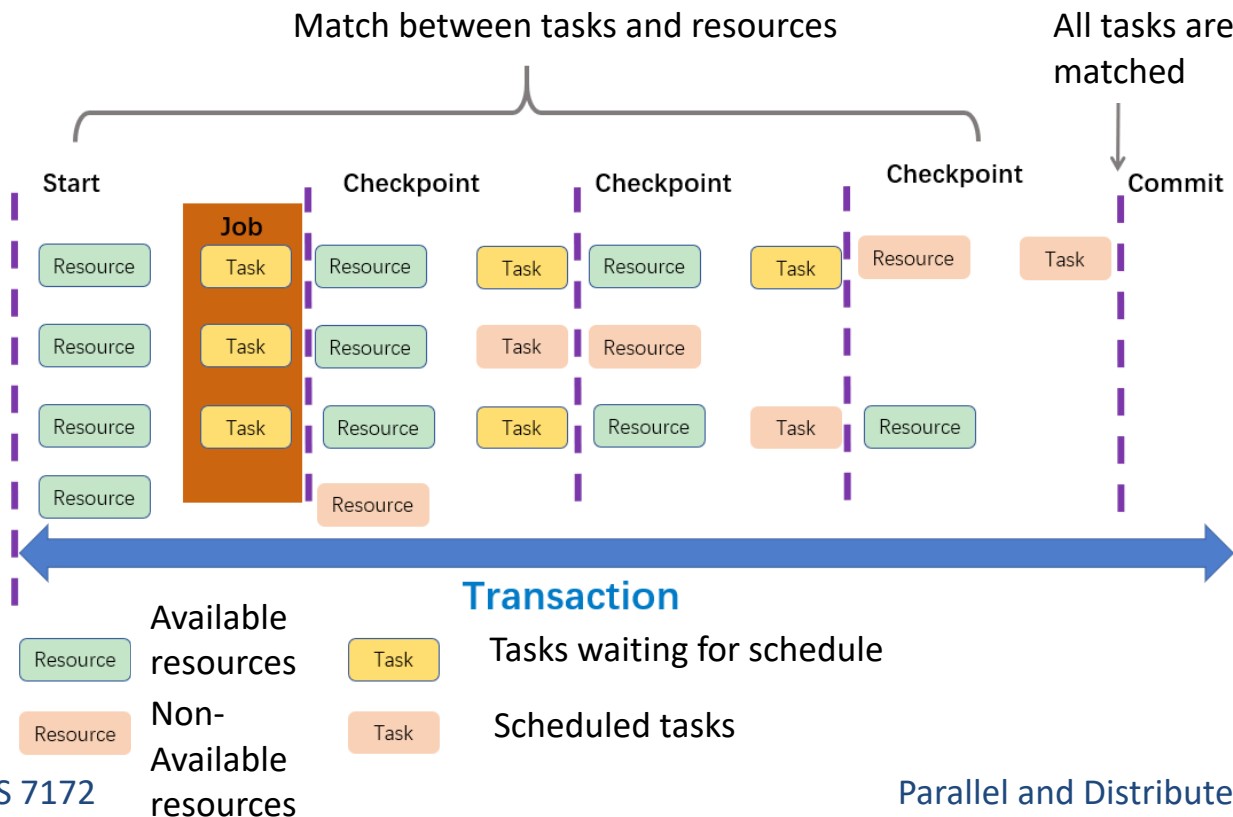
Compar...

> A central scheduler manages the resource information and task scheduling of the entire cluster. All tasks can only be scheduled by the central scheduler
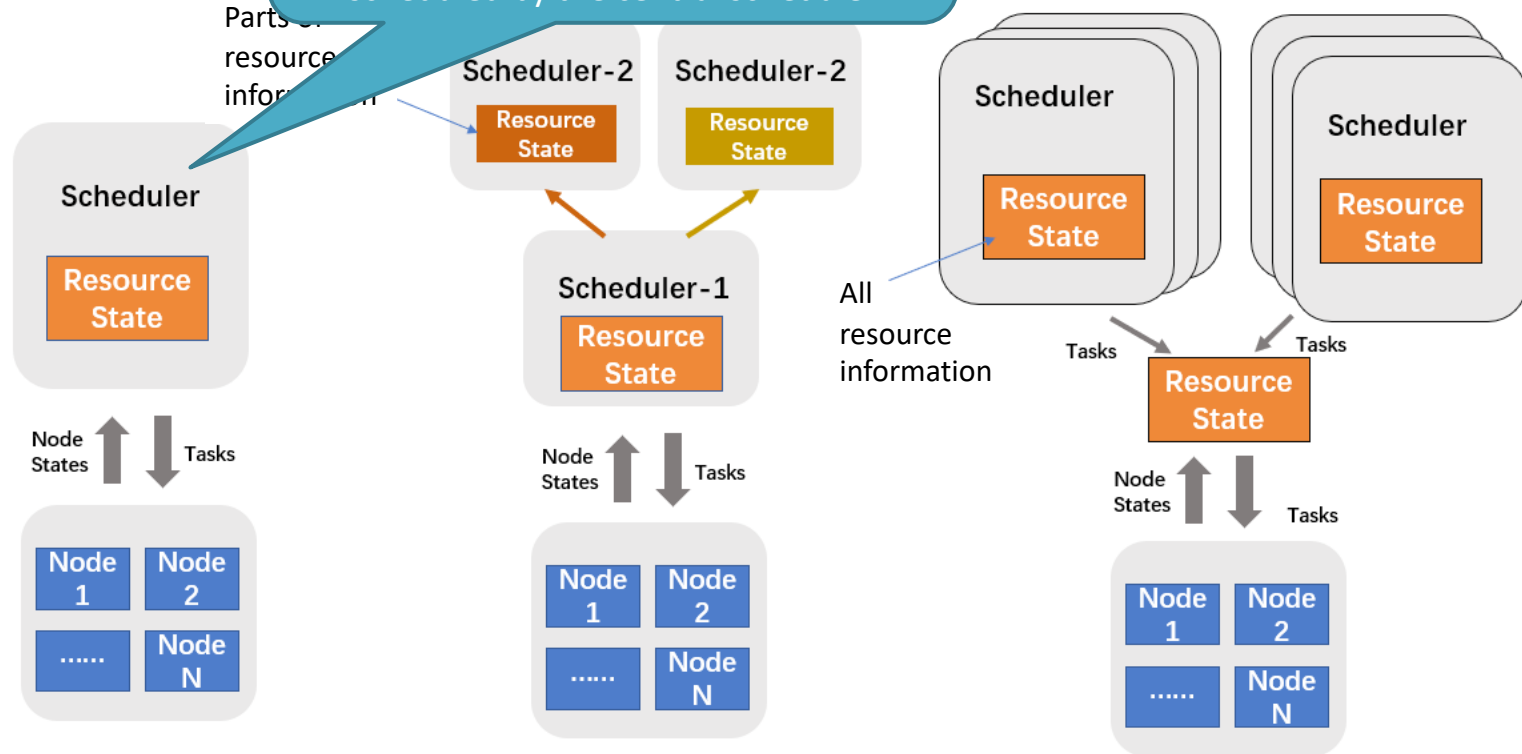
- **Adv:** global optimal

- **Disadv:** No scheduling concurrency, single point bottleneck problem, only suitable for small-scale clusters

Comparison

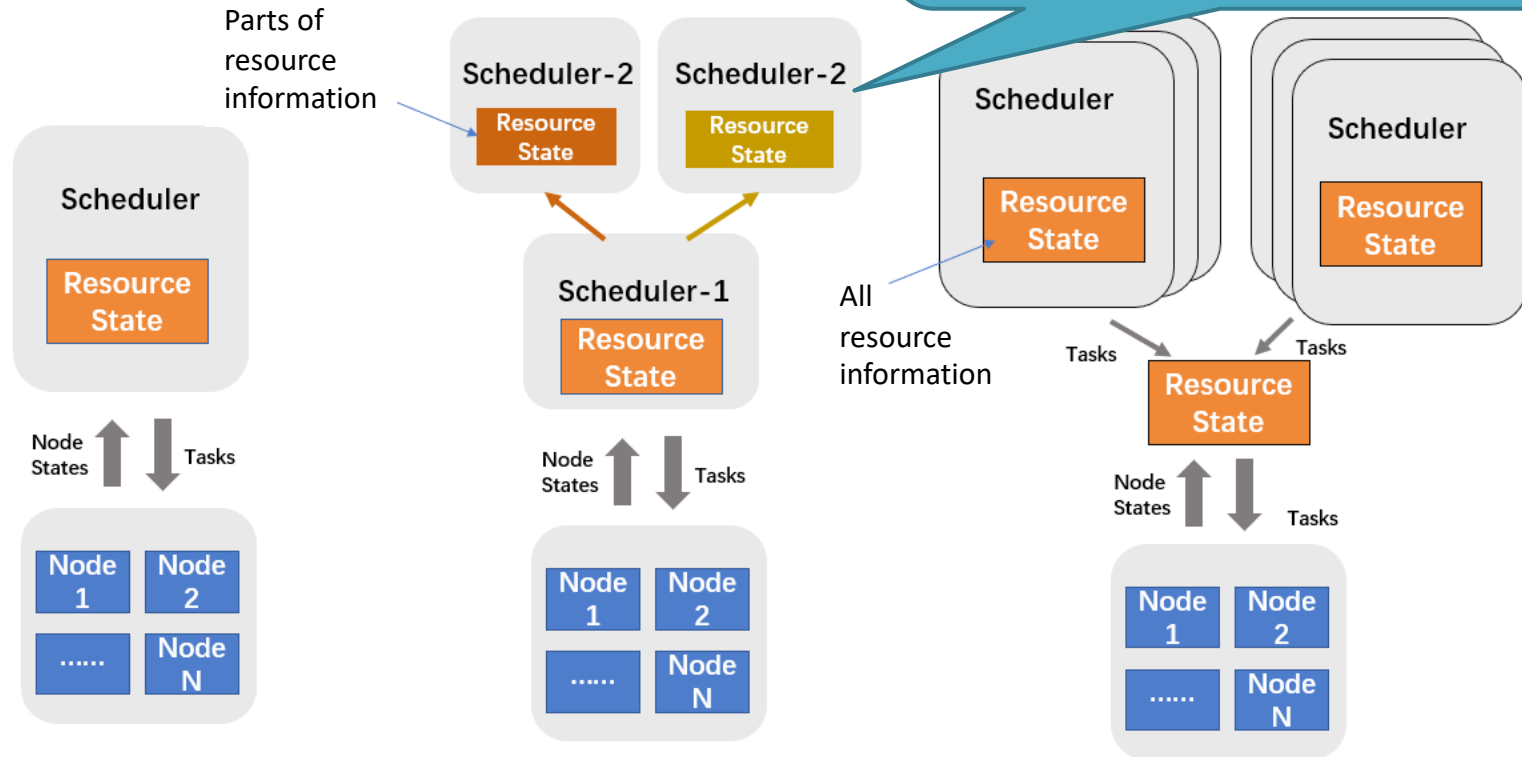The resource management and task scheduling are divided into two layers for scheduling. The first layer scheduler is responsible for cluster resource management and sends the available resources to the second layer scheduling.

- **Adv:** no single point bottleneck and higher scalability

- **Disadv:** 2rd level scheduler only has parts of the resource information and cannot provide

Comparison

Multiple schedulers, each scheduler can see the global resource information of the cluster through state sharing

Parts of resource information

Scheduler-2   Scheduler-2

Scheduler

Scheduler

All resource information
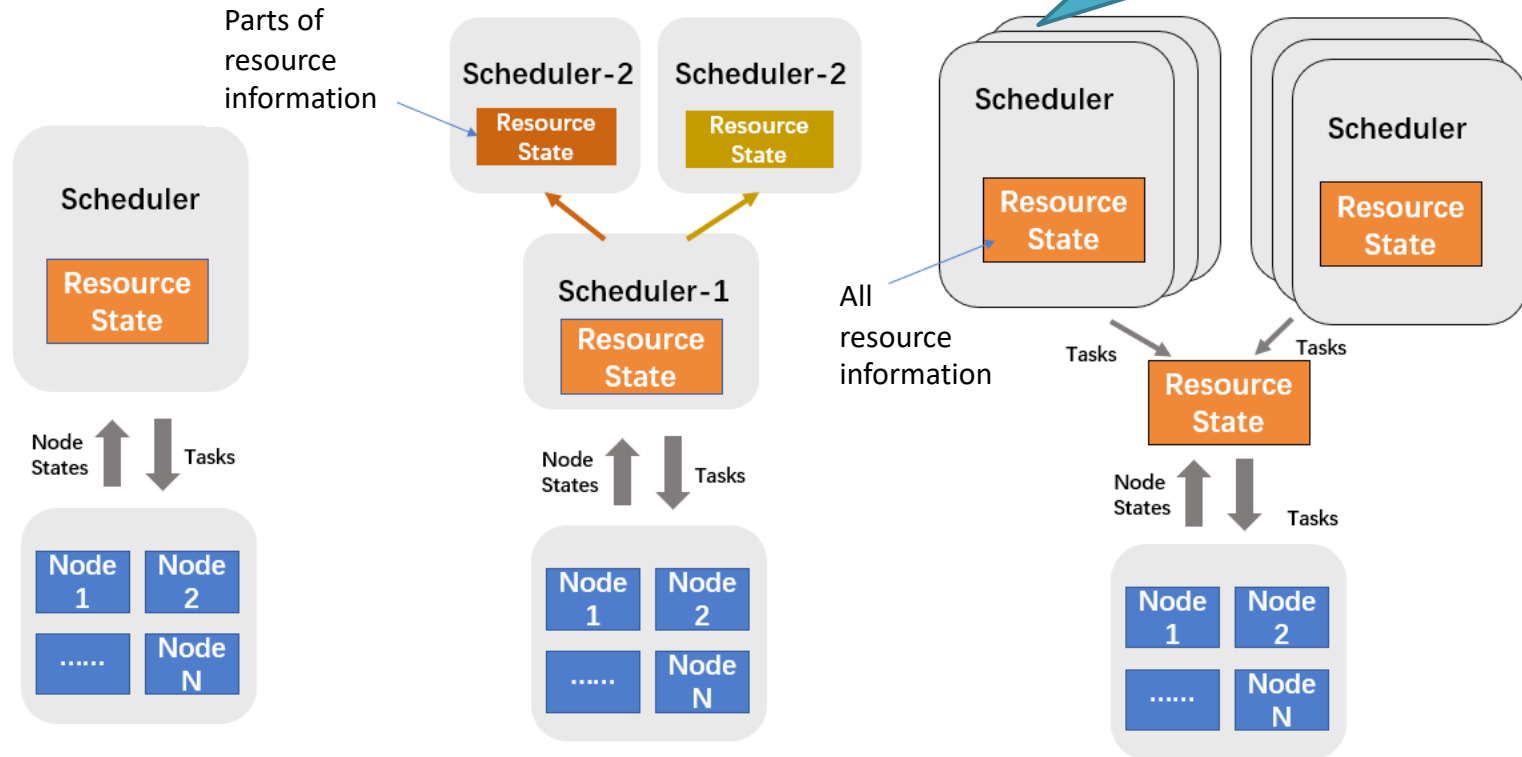
- **Adv:** Each scheduler can obtain global resource information in the cluster, so the task matching algorithm can achieve global optimality

- **Disadv:** multiple scheduler might cause resource competition and conflict

Conclusion

	Monolithic scheduling	**Two-level scheduling**	**Sharing state scheduling**
Architecture	Centralized, one node	Two schedulers	Distributed, multiple schedulers
Unit	Task	Task	Task
Parallel scheduling	No	Yes	Yes
Global optimal	Yes	No	Yes
Performance	Sharing state scheduling > Two-level scheduling > Monolithic scheduling		
Scalability	Sharing state scheduling > Two-level scheduling > Monolithic scheduling		
Suitable scenario	Small cluster	Medium cluster	Large scale cluster
Example	Borg	Mesos, YARN	Omega