

Kennesaw State University

CSE 3502 Operating Systems – Spring 2020

Project 1 - System call

Instructor: Kun Suo
Points Possible: 100

Assignments

Assignment 0: Build the Linux kernel (25 points)

Create a virtual machine using VirtualBox on your machine. As the kernel compiling is pretty large, please make sure your VM has at least 4GB memory and 40GB storage.

Step 1: Get the Linux kernel code

Before you download and compile the Linux kernel source, make sure you have development tools installed on your system. We recommend you work this project on your virtual machine.

In Ubuntu, install this software using apt:

```
$ sudo apt-get install -y gcc libncurses5-dev make wget flex bison vim libssl-dev libelf-dev
```

To obtain the version of your current kernel, type:

```
$ uname -r  
5.0
```

Then, download kernel 5.1 and extract the source:

```
$ wget https://cdn.kernel.org/pub/linux/kernel/v5.x/linux-5.1.tar.gz  
$ tar xvfz linux-5.1.tar.gz
```

We will refer LINUX_SOURCE to the top directory of the kernel source.

Step 2: Configure your new kernel

Before compiling the new kernel, a .config file needs to be generated in the top directory of the kernel source. To generate the config file and make possible changes to the default kernel configurations, type:

```
$ make menuconfig
```

No changes to the default configuration are needed at this time. Press SAVE and OK, and then exit the configuration menu and a default config file will be generated. You can check .config using the following command under kernel folder. (<https://youtu.be/UyOGF4UOoR0>)

```
$ ls -al
```

Step 3: Compile the kernel

In LINUX_SOURCE, compile to create a compressed kernel image:

```
$ make
```

You can use "make -j N" to accelerate the compiling. Here N denotes the number of CPUs on your VM.

To compile kernel modules:

```
$ make modules
```

You can use "make modules -j N" to accelerate the compiling. Here N denotes the number of CPUs on your VM.

Step 4: Install the kernel

Install kernel modules (become a root user, use the su command):

```
$ sudo make modules_install
```

Install the kernel:

```
$ sudo make install
```

If you are using Ubuntu, you need to create an init ramdisk manually:

```
$ sudo mkinitramfs -o /boot/initrd.img-5.1.0
```

```
$ sudo update-initramfs -c -k 5.1.0
```

The kernel image and other related files have been installed into the /boot directory. You can check it from /boot/grub/grub.cfg. Linux will boot by default using the 1st menu item.

Step 5: Modify grub configuration file

If you are using Ubuntu: change the grub configuration file:

```
$ sudo vim /etc/default/grub
```

Make the following changes:

```
GRUB_DEFAULT=0
```

```
GRUB_TIMEOUT=10
```

Then, update the grub entry:

```
$ sudo update-grub2
```

Step 6: Reboot your VM

Reboot to the new kernel:

```
$ sudo reboot
```

After boot, check if you have the new kernel:

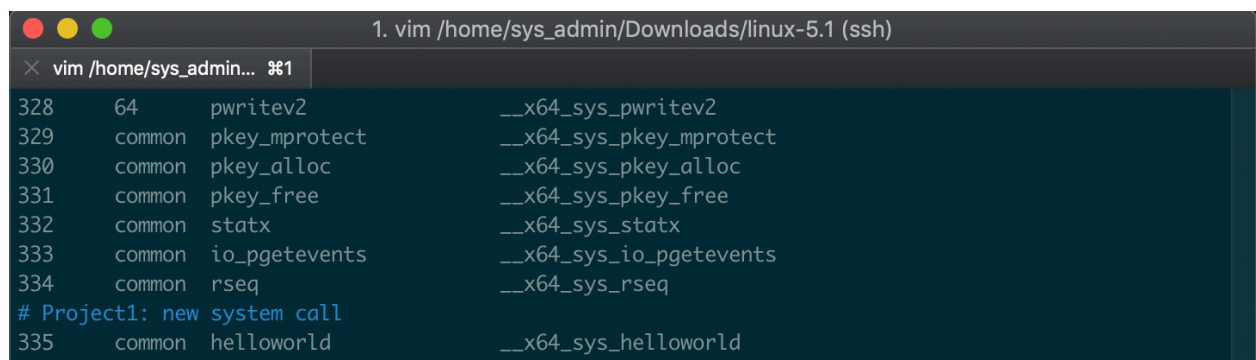
```
$ uname -r  
5.1.0
```

Assignment 1: Add a new system call into the Linux kernel (25 points)

In this assignment, we add a simple system call `helloworld` to the Linux kernel. The system call prints out a hello world message to the syslog. You need to implement the system call in the kernel and write a user-level program to test your new system call.

Step 1: register your system call

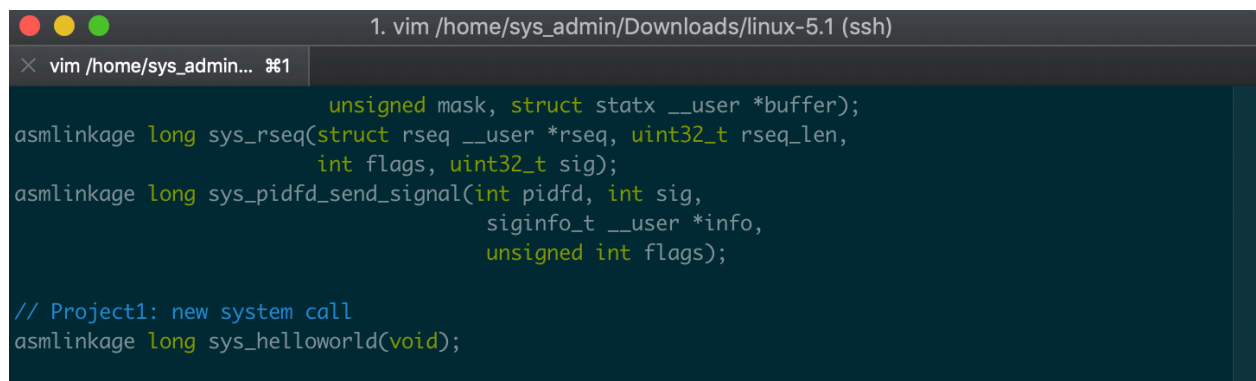
`arch/x86/entry/syscalls/syscall_64.tbl`



```
1. vim /home/sys_admin/Downloads/linux-5.1 (ssh)  
vim /home/sys_admin... 381  
328 64 pwritev2 __x64_sys_pwritev2  
329 common pkey_mprotect __x64_sys_pkey_mprotect  
330 common pkey_alloc __x64_sys_pkey_alloc  
331 common pkey_free __x64_sys_pkey_free  
332 common statx __x64_sys_statx  
333 common io_pgetevents __x64_sys_io_pgetevents  
334 common rseq __x64_sys_rseq  
# Project1: new system call  
335 common helloworld __x64_sys_helloworld
```

Step 2: declare your system call in the header file

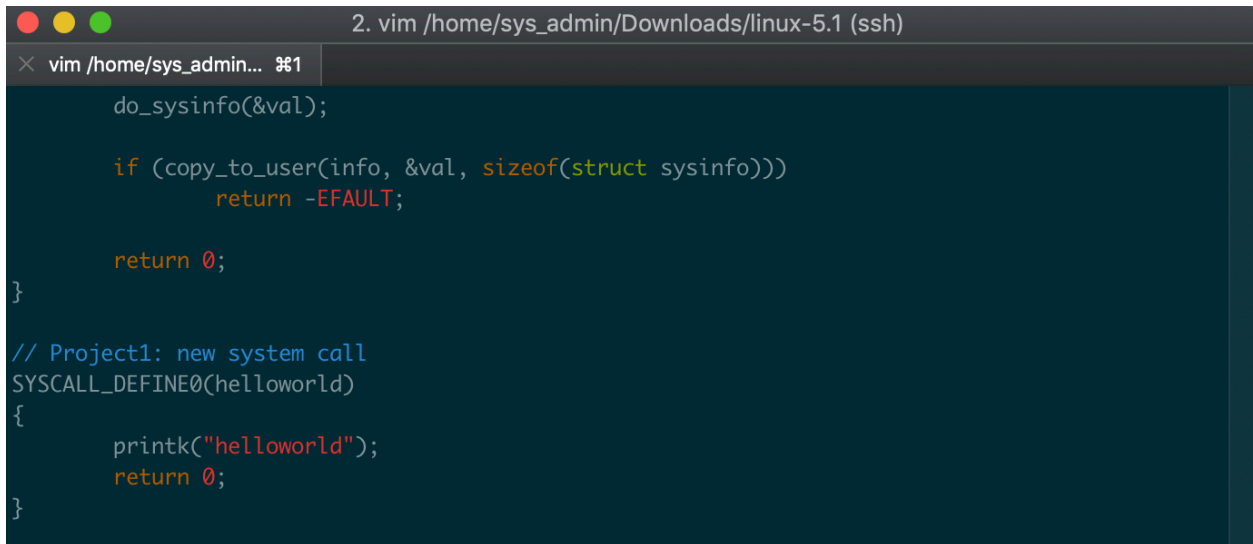
`include/linux/syscalls.h`



```
1. vim /home/sys_admin/Downloads/linux-5.1 (ssh)  
vim /home/sys_admin... 381  
    unsigned mask, struct statx __user *buffer);  
asmlinkage long sys_rseq(struct rseq __user *rseq, uint32_t rseq_len,  
    int flags, uint32_t sig);  
asmlinkage long sys_pidfd_send_signal(int pidfd, int sig,  
    siginfo_t __user *info,  
    unsigned int flags);  
  
// Project1: new system call  
asmlinkage long sys_helloworld(void);
```

Step 3: implement your system call

kernel/sys.c



```
2. vim /home/sys_admin/Downloads/linux-5.1 (ssh)
vim /home/sys_admin... 331
    do_sysinfo(&val);

    if (copy_to_user(info, &val, sizeof(struct sysinfo)))
        return -EFAULT;


    return 0;
}

// Project1: new system call
SYSCALL_DEFINE0(helloworld)
{
    printk("helloworld");
    return 0;
}
```

Repeat step 3 and 4 in assignment 0 to re-compile the kernel and reboot to the new kernel.

Step 4: write a user-level program to test your system call

Go to your home directory and create a test program test_syscall.c



```
2. vim /home/sys_admin/Downloads/test (ssh)
vim /home/sys_admin... 331
#include <linux/unistd.h>
#include <sys/syscall.h>
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

#define __NR_helloworld 335

int main(int argc, char *argv[])
{
    syscall(__NR_helloworld);
    return 0;
}
```

Compile the user level program:

```
$ gcc test_syscall.c -o test_syscall
```

Test the new system call by running:

```
$ ./test_syscall
```

The test program will call the new system call and output a helloworld message at the tail of the output of dmesg.

```
2. fish /home/sys_admin/Downloads/test (ssh)
fish /home/sys_admin... %1
sys_admin@R640-2 ~/D/test> dmesg | grep hello
[ 128.626914] helloworld
```

Submission of assignment 1:

Please have two copies of kernel source code: 1) the original kernel source code without any modification; 2) the kernel source code you modified. You can define the folder name based on your need. Here I use *linux-5.1* as the original source code without modification and *linux-5.1-modified* as the source code I worked on.

```
fish /Users/ksuo/Desktop/kernel %1
ksuo@ltkup50143mac ~/D/kernel> ls
linux-5.1          linux-5.1-modified linux-5.1.tar.gz
```

Please use diff command to highlight your modification (Here the original_file.c refers the file or file path of the original file source code; the modified_file.c refers the file or file path of the file source code you have modified):

```
$ diff -u original_file.c modified_file.c > result.txt
```

For example, to show the difference between file include/linux/syscalls.h, just use the command below:

```
fish /Users/ksuo/Desktop/kernel %1
ksuo@ltkup50143mac ~/D/kernel> diff -u ./linux-5.1/include/linux/syscalls.h ./linux-5.1-modified/include/linux/syscalls.h > result.txt
```

Assignment 2: Extend your new system call to print out the calling process's information (25 points)

Follow the instructions we discussed above and implement another system call print_self. This system call identifies the calling process at the user-level and print out various information of the process.

Implement the `print_self` system call and print out the following information of the calling process:

- Process id, running state, and program name
- Start time and virtual runtime
- Its parent processes until init (first system process)

HINT: The macro ***current*** returns a pointer to the `task_struct` of the current running process. The virtual runtime (`vruntime`) is located at the `sched_entity` of `task_struct`.

`task_struct` data structure in the kernel:

<https://elixir.bootlin.com/linux/v5.1/source/include/linux/sched.h#L585>

Please use diff command to highlight your modification:

```
$ diff -u original_file.c modified_file.c > result.txt
```

Assignment 3: Extend your new system call to print out the information of an arbitrary process identified by its PID (25 points)

Implement another system call `print_other` to print the information for an arbitrary process. The system call takes a process pid as its argument and outputs the above information of this process.

HINT: You can start from the init process and iterate over all the processes. For each process, compare its pid with the target pid. If there is a match, return the pointer to this `task_struct`.

A better approach is to use the pidhash table to look up the process in the process table. Linux provides many functions to find a task by its pid.

Please use diff command to highlight your modification:

```
$ diff -u original_file.c modified_file.c > result.txt
```

Submitting Assignment

Submit your assignment zip file through D2L using the appropriate assignment link.

Format: create three folders:

Assignment 1


Assignment 2

Assignment 3

In each folder, add the **user space source code, kernel space source code, and screenshot of output** inside.

For the kernel code, please do not add the entire kernel source code. Just add your modification code, e.g., result1.txt, result2.txt, result3.txt, ...

In order to use diff command and get the modification file result.txt. Please have two copies of kernel source code: 1) the original kernel source code without any modification; 2) the kernel source code you modified, as the below figure shows:


A terminal window titled 'fish /Users/ksuo/Desktop/kernel' showing a directory listing command and its output. The prompt is 'ksuo@ltsup50143mac ~/D/kernel>'. The command 'ls' has been executed, resulting in the output 'linux-5.1', 'linux-5.1-modified', and 'linux-5.1.tar.gz' on the same line.

```
fish /Users/ksuo/Desktop/kernel
ksuo@ltsup50143mac ~/D/kernel> ls
linux-5.1          linux-5.1-modified linux-5.1.tar.gz
```

Please use diff command to highlight your modification (Here the original_file.c refers the file or file path of the original file source code; the modified_file.c refers the file or file path of the file source code you have modified):

```
$ diff -u original_file.c modified_file.c > result.txt
```

For example, to show the difference between file include/linux/syscalls.h, just use the command below:

A terminal window titled 'fish /Users/ksuo/Desktop/kernel' showing a diff command being executed. The prompt is 'ksuo@ltsup50143mac ~/D/kernel>'. The command 'diff -u ./linux-5.1/include/linux/syscalls.h ./linux-5.1-modified/include/linux/syscalls.h > result.txt' has been entered and is partially highlighted in blue.

```
fish /Users/ksuo/Desktop/kernel
ksuo@ltsup50143mac ~/D/kernel> diff -u ./linux-5.1/include/linux/syscalls.h ./linux-5.1-modified/include/linux/syscalls.h > result.txt
```

Zip all the files and folders together into one zip file and name it as CS3502_[your D2L user name]. Such as, **CS3502_mahmed29.zip**.