

CS 3502

Operating Systems

Scheduling

Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

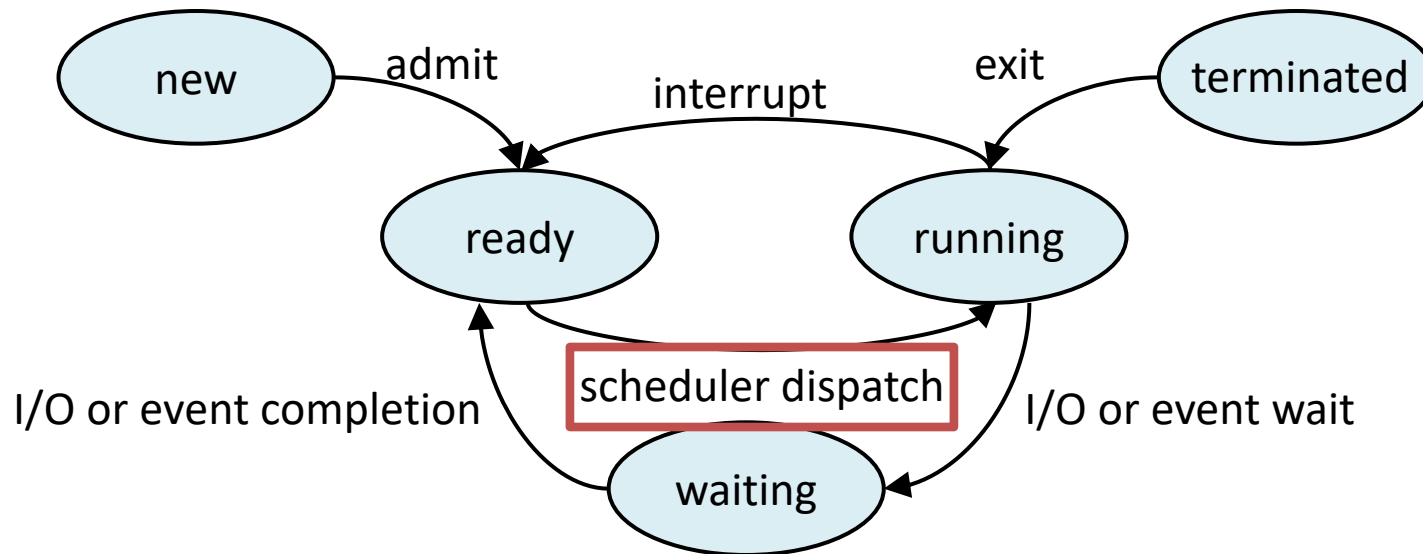
Outline

- Introduction to CPU scheduling
 - What is CPU scheduling
 - Why we need CPU scheduling
 - When scheduling happens
- Scheduling policies
 - FCFS, SJF, RR, Priority
 - Scheduling on multiple CPUs
 - Scheduling in Linux



What is CPU scheduling?

- The five-state process model






CPU scheduling

Selects from among the processes/threads that are ready to execute, and allocates the CPU to it

Why CPU scheduling?

- In support of **multiprogramming**
 - uniprocessor systems
 - ▶ Time-sharing processor
 - multiprocessor systems
 - ▶ Efficiently distributing tasks
 - Real-time systems
 - ▶ Reliably guaranteeing deadlines



Why CPU scheduling?

- It is (maybe) the most important part in a OS
 - Why some OS seems to be faster than others?
 - Why I do not see performance improvement when upgrading to a 16-core computer?



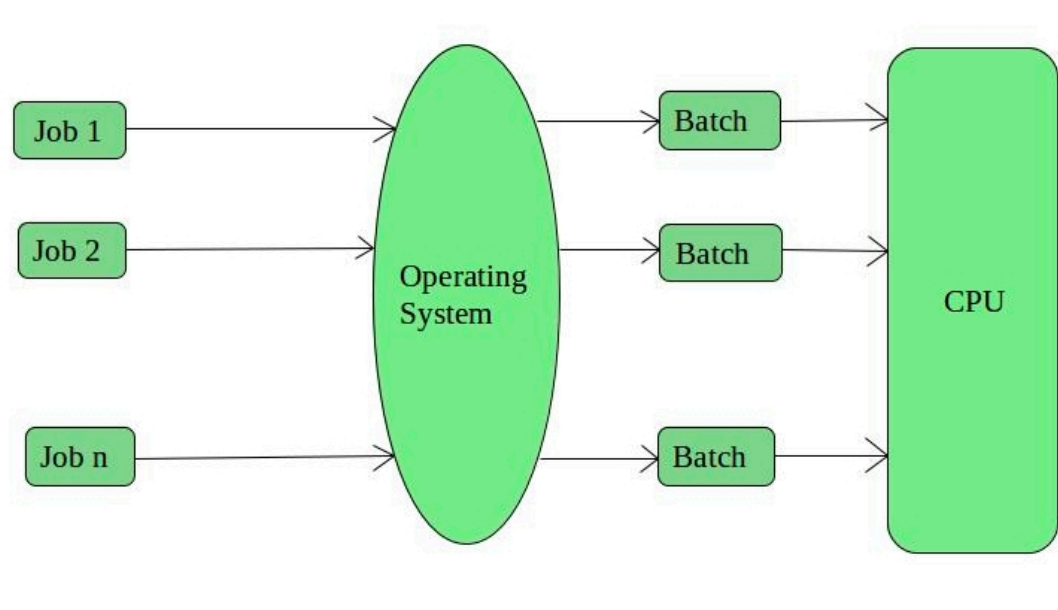
Why CPU scheduling? – Different Goals

- All systems
 - Fairness - giving each process a fair share of the CPU
 - Policy enforcement - seeing that stated policy is carried out
 - Balance - keeping all parts of the system busy



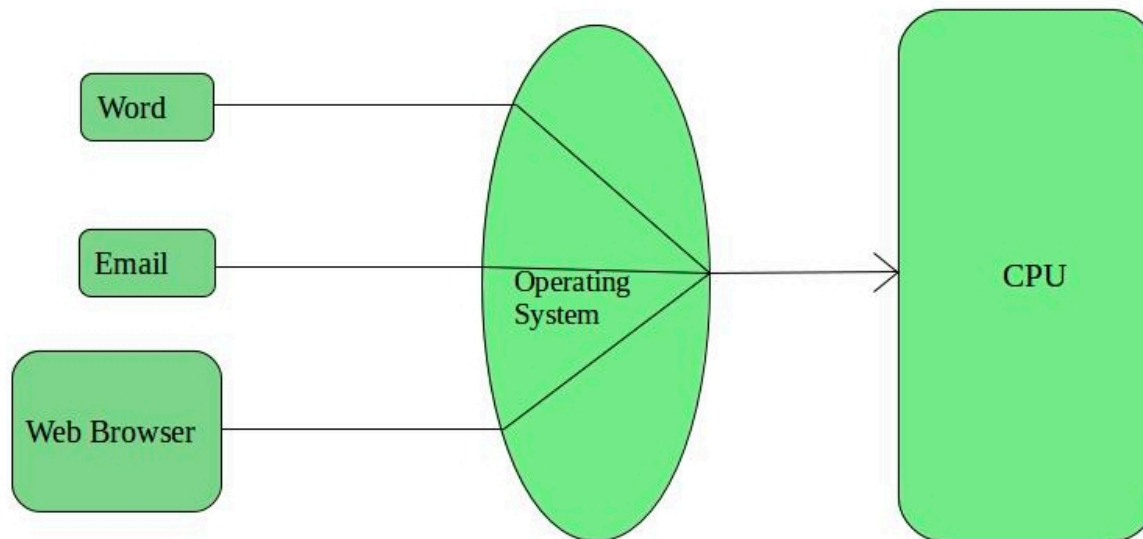
Why CPU scheduling? – Different Goals

- Batch systems
 - Throughput - maximize jobs per hour
 - Turnaround time - minimize time between submission and termination
 - CPU utilization - keep the CPU busy all the time



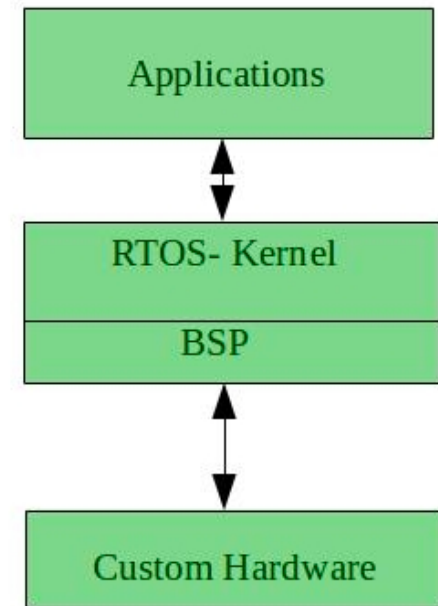
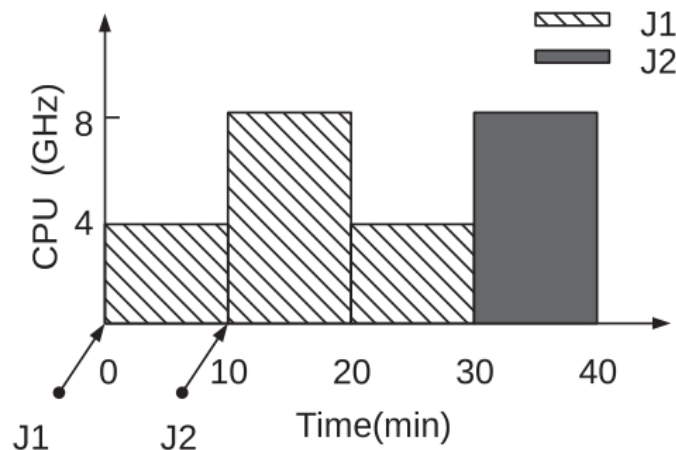
Why CPU scheduling? – Different Goals

- Interactive systems
 - Response time - respond to requests quickly
 - Proportionality - meet users' expectations



Why CPU scheduling? – Different Goals

- Real-time systems
 - Meeting deadlines - avoid losing data
 - Predictability - avoid quality degradation in multimedia systems



Why CPU scheduling? – Different Goals

- All systems
 - Fairness - giving each process a fair share of the CPU
 - Policy enforcement - seeing that stated policy is carried out
 - Balance - keeping all parts of the system busy
- Batch systems
 - Throughput - maximize jobs per hour
 - Turnaround time - minimize time between submission and termination
 - CPU utilization - keep the CPU busy all the time
- Interactive systems
 - Response time - respond to requests quickly
 - Proportionality - meet users' expectations
- Real-time systems
 - Meeting deadlines - avoid losing data
 - Predictability - avoid quality degradation in multimedia systems



Scheduling Goals: A Different Point of View

- User oriented → minimize
 - Response time (wait time): the time that the first response is received (interactivity)
 - Turnaround time: the time that the task finishes
 - Predictability: variations in different runs
- System oriented → maximize
 - Throughput: # of tasks that finish per time unit
 - Utilization: the percentage of time the CPU is busy
 - Fairness: avoid starvation




When scheduling happens?

- CPU scheduling may take place at
 - Clock interrupts
 - I/O interrupts
 - I/O completion
 - Termination

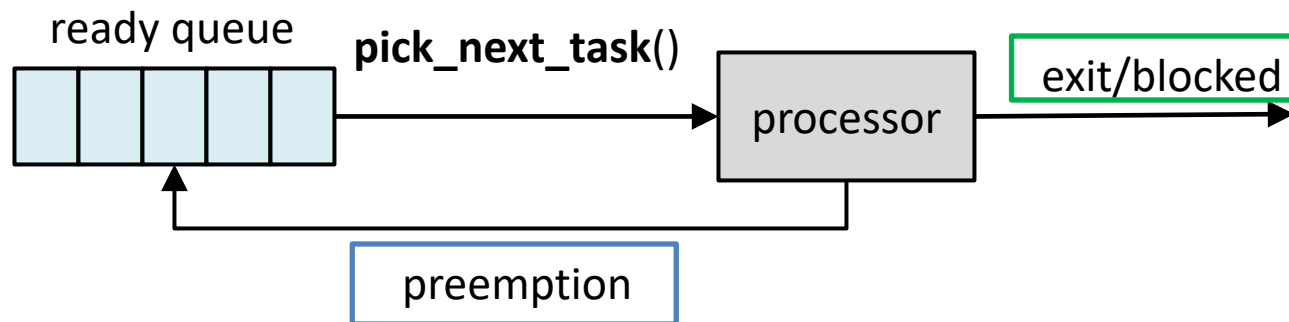
} preemptive

} non-preemptive




When scheduling happens?

- Non-preemptive
 - Scheduling only when current process **terminates** or **gives up** control
- Preemptive
 - Processes can be **forced** to give up control



CPU-bound tasks vs. I/O-bound tasks



```
#include <stdio.h>

int main() {
    int a = 1;

    while ( a < 1000000 )
    {
        a++;
    }

    return 0;
}
```

a CPU-bound/CPU-intensive process

```
#include <stdio.h>

int main()
{
    char name[20];
    printf("Enter name: ");

    scanf("%s", name);
    printf("Your name is %s.", name);

    return 0;
}
```

a I/O-bound/I/O-intensive process

I/O is when a process enters the blocked state waiting for an external device to complete its work

Outline

- Introduction to CPU scheduling
 - What is CPU scheduling
 - Why we need CPU scheduling
 - When scheduling happens
- Scheduling policies
 - FCFS, SJF, RR, Priority
 - Scheduling on multiple CPUs
 - Scheduling in Linux



Scheduling Policies


Not exist best scheduling.
It depends on your goals.

- Batch Systems
 - First-Come First-Serve (FCFS)
 - Shortest Job First
 - Shortest Remaining Time Next
- Interactive Systems
 - Round-Robin
 - Priority Scheduling
 - Multiple Queues
 - Shortest Process Next
 - Guaranteed Scheduling
 - Lottery Scheduling
- Real-time Systems
 - Rate Monotonic Scheduling
 - Earliest Deadline First Scheduling

Determine the next
ready task to run

Turnaround time = End time – Arrival time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



Response time = Start time – Arrival time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



0

1 \longleftrightarrow 8

2 \longleftrightarrow 12

3 \longleftrightarrow 21



First-Come, First-Serve (FCFS)

- CPU schedules the task that arrived earliest, non-preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



Average turnaround time = $((8-0)+(12-1)+(21-2)+(26-3)) / 4 = 15.25$



First-Come, First-Serve (FCFS)

- CPU schedules the task that arrived earliest, non-preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



Average response time = $(0 + (8 - 1) + (12 - 2) + (21 - 3)) / 4 = 8.75$




Shortest Job First (SJF)

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive




Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool




Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool




Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool




Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool




Shortest Job First (SJF)

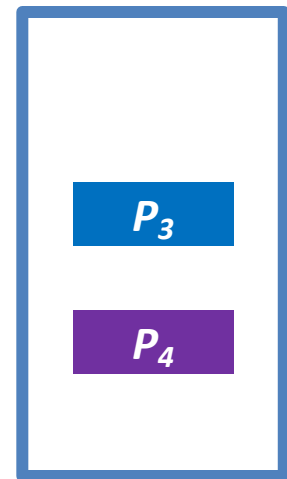
- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool



Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool



Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool



Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool



Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool



Shortest Job First (SJF)

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



$$\text{Average turnaround time} = ((8-0)+(12-1)+(26-2)+(17-3)) / 4 = 14.25$$



Shortest Job First (SJF)

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



$$\text{Average response time} = (0 + (8 - 1) + (17 - 2) + (12 - 3)) / 4 = 7.75$$




Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool



Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool




Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool

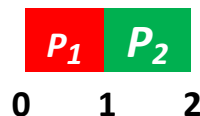
$P_1(7)$

Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



P2 remain =3

pool




Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



P2 remain =2

pool



Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool



Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool



Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool




Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool

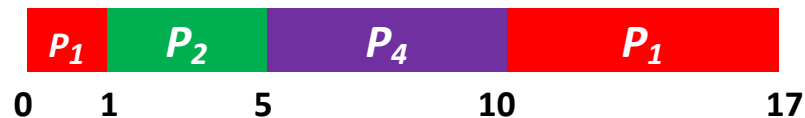


Shortest Job First (SJF) with preemption

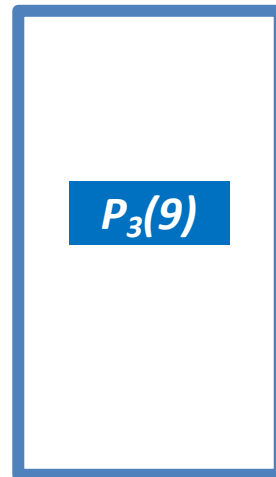
- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool



Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive




$$\text{Average turnaround time} = ((17-0)+(5-1)+(10-3)+(26-2)) / 4 = 13$$



Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



$$\text{Average response time} = (0 + (1 - 1) + (5 - 3) + (17 - 2)) / 4 = 4.25$$


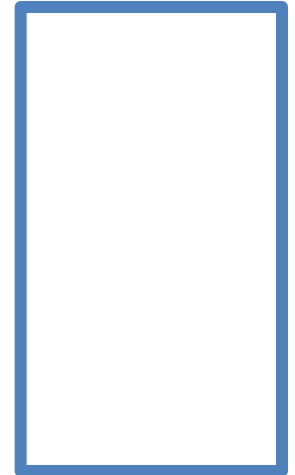


Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool




Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



P_1 is put at the end of queue after scheduled out





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool




Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool

$P_3 (1)$





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5


pool



Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5




$$\text{Average turnaround time} = ((20-0)+(8-1)+(26-2)+(25-3)) / 4 = 18.25$$

Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



$$\text{Average response time} = (0 + (4-1) + (8-2) + (12-3)) / 4 = 4.5$$

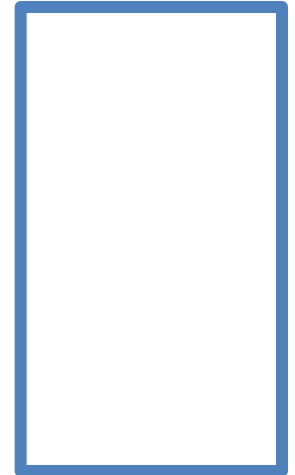


Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$



Round Robin (RR)

Select the task at the beginning

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5


pool

$P_2 (4)$

$P_3 (9)$

$P_4 (5)$

$q=5$



Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool

$P_3 (9)$

$P_4 (5)$

$P_1 (3)$

$q=5$




P_1 is put at the end of queue after scheduled out

Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool




Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$




Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$



Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool

$P_1 (3)$

$P_3 (4)$

$q=5$




Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$




Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$




Round Robin (RR)

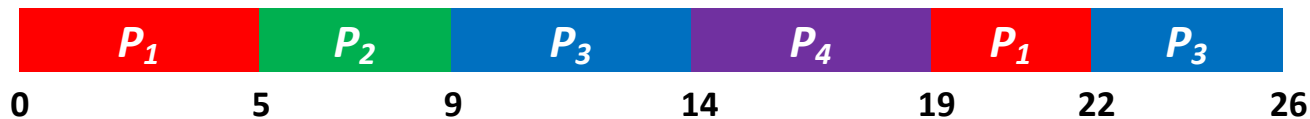
- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$

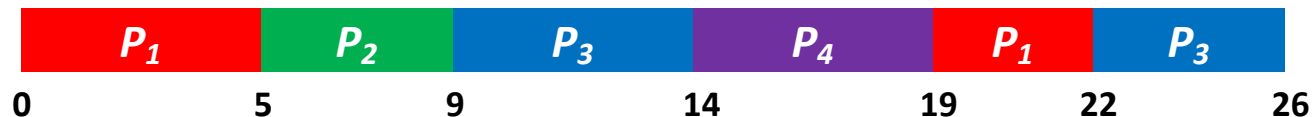


Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

$q=5$



Average turnaround time = $((22-0)+(9-1)+(26-2)+(19-2)) / 4 = 17.5$




Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

$q=5$




Average response time = $(0+(5-1)+(9-2)+(14-3)) / 4 = 5.5$




Priority Scheduling

- CPU schedules the highest priority (smaller value) first, FCFS within the same priority



Priority Scheduling

- CPU schedules the highest priority (smaller value) first, FCFS within the same priority



Comparison

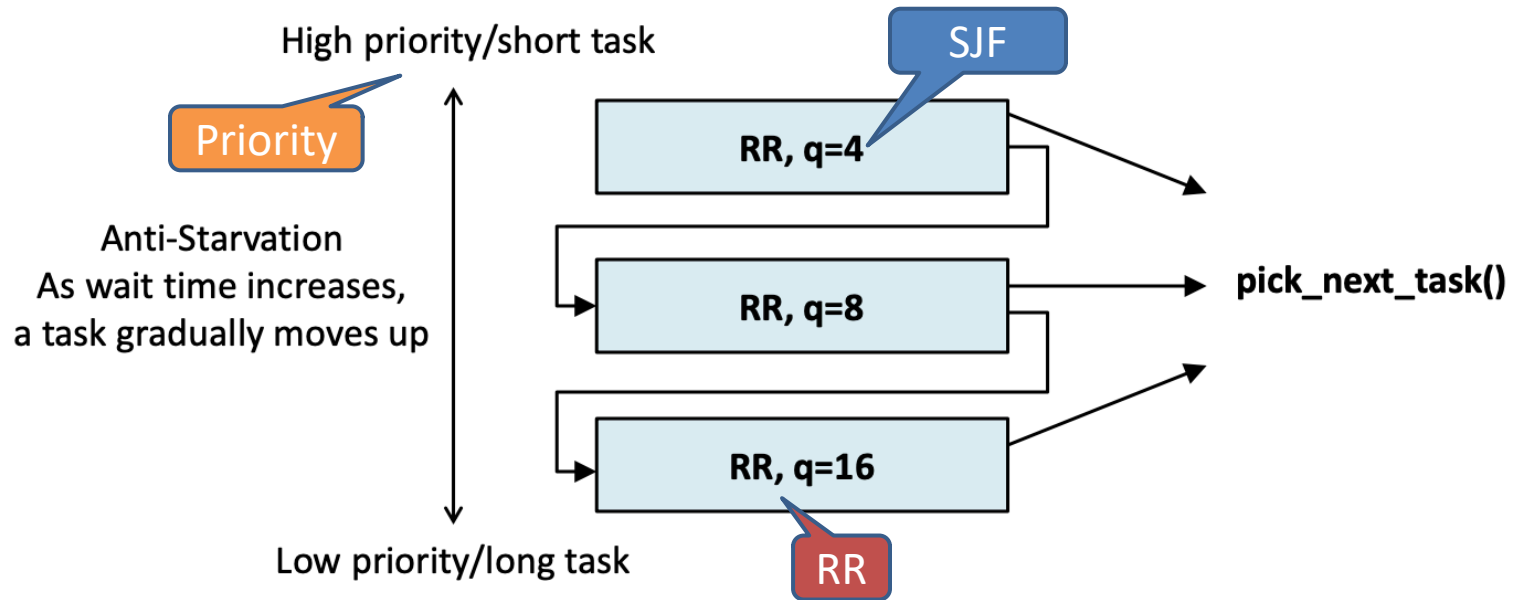
	Turnaround time	Response time
FCFS	15.25	8.75
SJF-preemptive	13	4.25
RR (q=5)	17.5	5.5
Priority scheduling	N/A	N/A

	Throughput	Response time	Starvation
FCFS	TBD	TBD	No
SJF-preemptive	High	Good	Yes
RR	Can be low	Good	No
Priority scheduling	Can be high	Can be good	Can remove



Multilevel Feedback Queue


	Throughput	Response time	Starvation
FCFS	TBD	TBD	No
SJF-preemptive	High	Good	Yes
RR	Can be low	Good	No
Priority scheduling	Can be high	Can be good	Can remove



Windows XP, Mac OS X, Linux 2.6.22 and before

Challenges on Emerging Hardware and Applications

- Multi-processor → Single queue



Multiprocessor = more powerful processor

pick_next_task() will be the bottleneck

Linux 2.4.x

Pros:

1. Easy to implement
2. Perfect load balancing


Cons:

1. Scalability issues due to centralized synchronization
2. High overhead and low efficiency
3. Hard to maintain cache hotness due to global scheduling



Challenges on Emerging Hardware and Applications


- Multi-processor → Many queues



Linux 2.4.x

Multiprocessor = more powerful processor

pick_next_task() will be the bottleneck



Linux 2.6.x

Pros:

1. Scalable to many CPUs
2. Easy to maintain cache hotness

Cons:

Self-scheduling, could have load imbalance




Overcome Load Imbalance

- Push model

Every a while, a kernel thread checks load imbalance and move threads


Made by OS




- Pull model

Whenever a queue becomes empty, steal a thread from non-empty queues


Made by local queue. Both are widely used



Load balance on SMP vs. NUMA



symmetric multiprocessing (SMP):
The distance to memory is the same



Non-uniform memory access (NUMA):
The distance to memory is different

Multi processor/core scheduling

The scheduling policy not only considers the fairness, throughput, etc., but also needs to consider the **hardware architecture** (e.g., locality)


Scheduling in Linux

- Linux Completely Fair Scheduler (CFS)
 - CFS = RR + SJF + Priority + smart data structure + hardware consideration
 - One ready queue for one processor
 - Red-black tree based ready queue

Print it
in project 1

vruntime: how much time the
task consumes the CPU cycles

pick_next_task()
Select the leftmost task
which has the least vruntime



Conclusion

- Introduction to CPU scheduling
 - What is CPU scheduling
 - Why we need CPU scheduling
 - When scheduling happens
- Scheduling policies
 - FCFS, SJF, RR, Priority
 - Scheduling on multiple CPUs
 - Scheduling in Linux

