

CS 3502

Operating Systems

Project 2 Lab

Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

Project 2

- Read and write parallel program using Pthread
- Learn to use Pthread functions
- User level projects, not kernel code



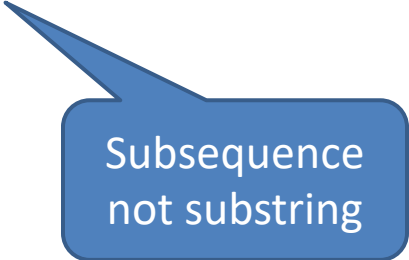
Assignment 1

- Given two character strings s1 and s2. Write a Pthread program to find out the number of substrings, in string s1, that is exactly the same as s2.
- <https://github.com/kevinsuo/CS3502/blob/master/project-2-1.c>



Assignment 1 Examples

- `number_substring("abcdab", "ab") = 2,`
- `number_substring("aaa", "a") = 3,`
- `number_substring("abac", "bc") = 0.`



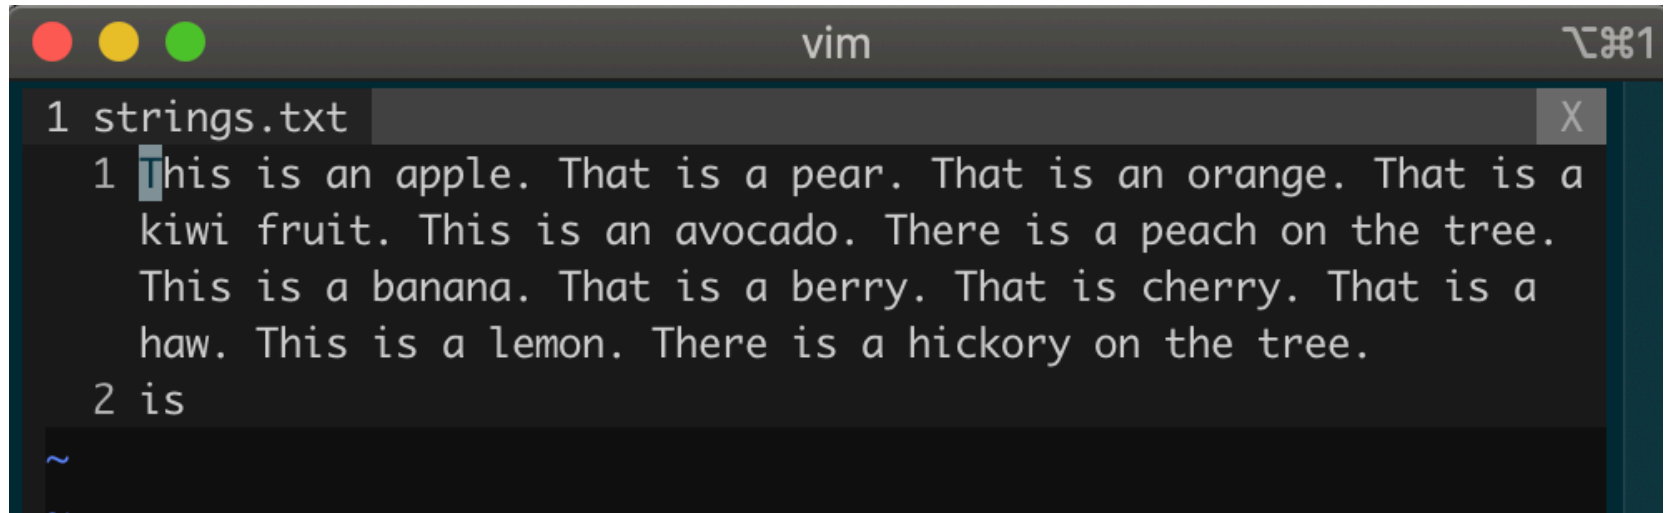
Subsequence
not substring



Assignment 1

- Input file:

<https://github.com/kevinsuo/CS3502/blob/master/strings.txt>

A screenshot of a vim editor window. The title bar shows 'vim' and a window icon. The editor is displaying a file named 'strings.txt'. The content of the file is as follows:

```
1 This is an apple. That is a pear. That is an orange. That is a  
  kiwi fruit. This is an avocado. There is a peach on the tree.  
  This is a banana. That is a berry. That is cherry. That is a  
  haw. This is a lemon. There is a hickory on the tree.  
2 is
```

```

int main(int argc, char *argv[])
{
    int count;
    readf(fp);
    count = ham_substring();
    printf("The number of substrings is: %d\n", count);
    return 1;
}

```

Use strlen(s2)-1
Because there exist a
'\n' at the end of s2

```

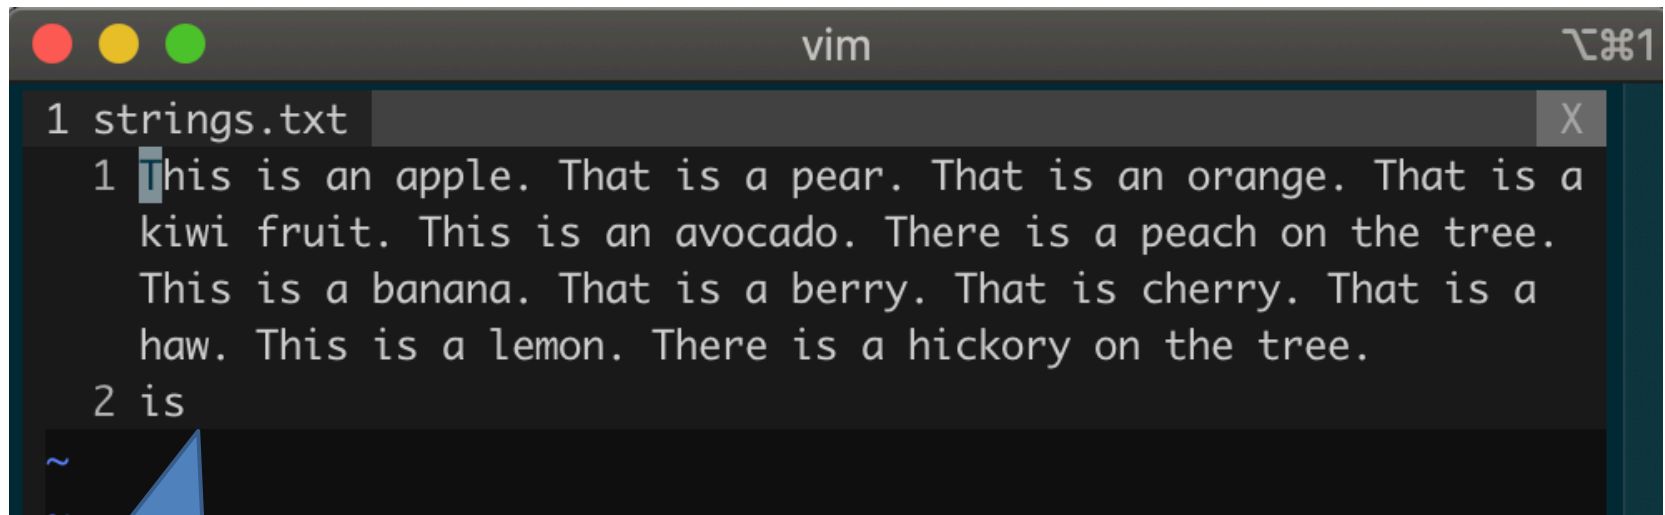
int total = 0;
int n1, n2;
char *s1, *s2;
FILE *fp;

int readf(FILE *fp)
{
    if((fp=fopen("strings.txt", "r"))==NULL){
        printf("ERROR: can't open string.txt!\n");
        return 0;
    }
    s1=(char *)malloc(sizeof(char)*MAX);
    if(s1==NULL){
        printf("ERROR: Out of memory!\n");
        return -1;
    }
    s2=(char *)malloc(sizeof(char)*MAX);
    if(s2==NULL){
        printf("ERROR: Out of memory\n");
        return -1;
    }
    /*read s1 s2 from the file*/
    s1=fgets(s1, MAX, fp);
    s2=fgets(s2, MAX, fp);
    n1=strlen(s1); /*length of s1*/
    n2=strlen(s2)-1; /*length of s2*/

    if(s1==NULL || s2==NULL || n1<n2) /*when error exit*/
        return -1;
    return 0;
}

```

'\n' at the end of s2



```
vim 1
1 strings.txt X
1 This is an apple. That is a pear. That is an orange. That is a
  kiwi fruit. This is an avocado. There is a peach on the tree.
  This is a banana. That is a berry. That is cherry. That is a
  haw. This is a lemon. There is a hickory on the tree.
2 is
```

Use `strlen(s2)-1`
Because there exist a
'\n' at the end of s2

```

int main(int argc, char *argv[])
{
    int count;

    readf(fp);
    count = num_substring();
    printf("The number of substrings is: %d\n", count);
    return 1;
}

```

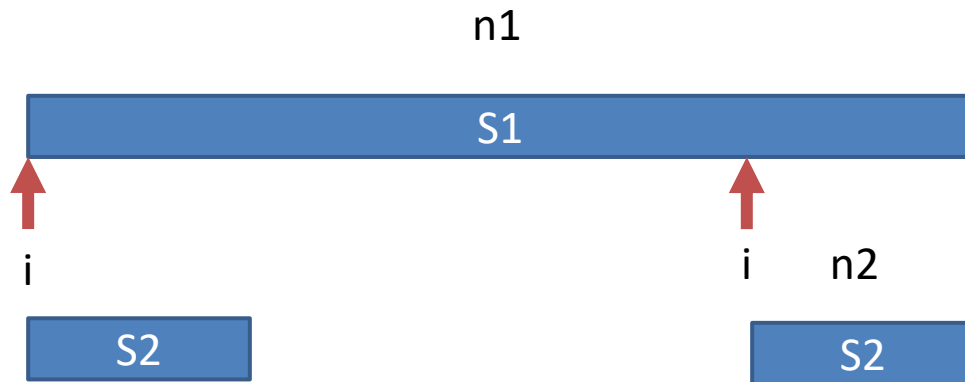
```

int num_substring(void)
{
    int i,j,k;
    int count;

    for (i = 0; i <= (n1-n2); i++){
        count=0;
        for(j = i,k = 0; k < n2; j++,k++){ /*search for the next
string of size of n2*/
            if (*(s1+j)!=*(s2+k)){
                break;
            }else{
                count++;
            }

            if(count==n2){
                total++;          /*find a substring in this
step*/
            }
        }
    }
    return total;
}

```




```

int main(int argc, char *argv[])
{
    int count;

    readf(fp);
    count = num_substring();
    printf("The number of substrings is: %d\n", count);
    return 1;
}

```

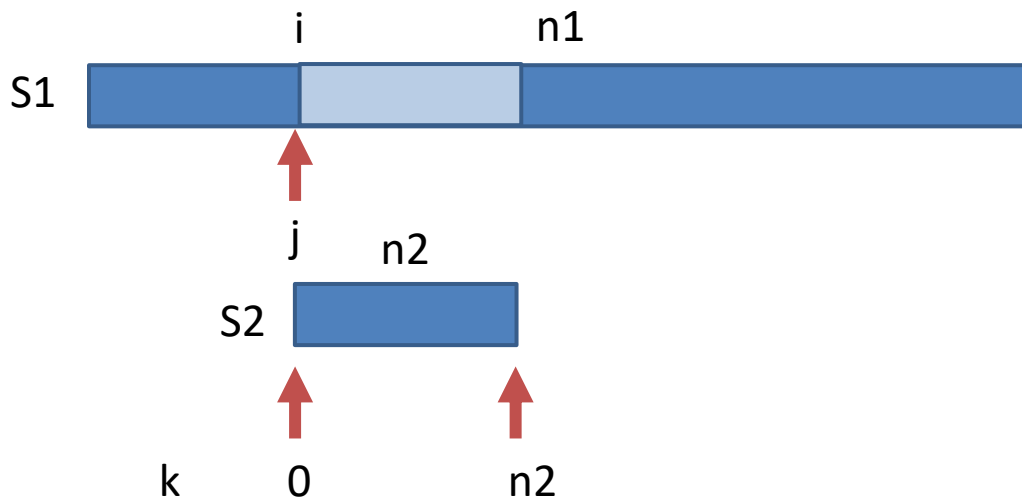
```

int num_substring(void)
{
    int i,j,k;
    int count;

    for (i = 0; i <= (n1-n2); i++){
        count=0;
        for(j = i,k = 0; k < n2; j++,k++){ /*search for the next
string of size of n2*/
            if (*(s1+j)!=*(s2+k)){
                break;
            }else{
                count++;
            }

            if(count==n2){
                total++;          /*find a substring in this
step*/
            }
        }
    }
    return total;
}

```



```

int main(int argc, char *argv[])
{
    int count;

    readf(fp);
    count = num_substring();
    printf("The number of substrings is: %d\n", count);
    return 1;
}

```

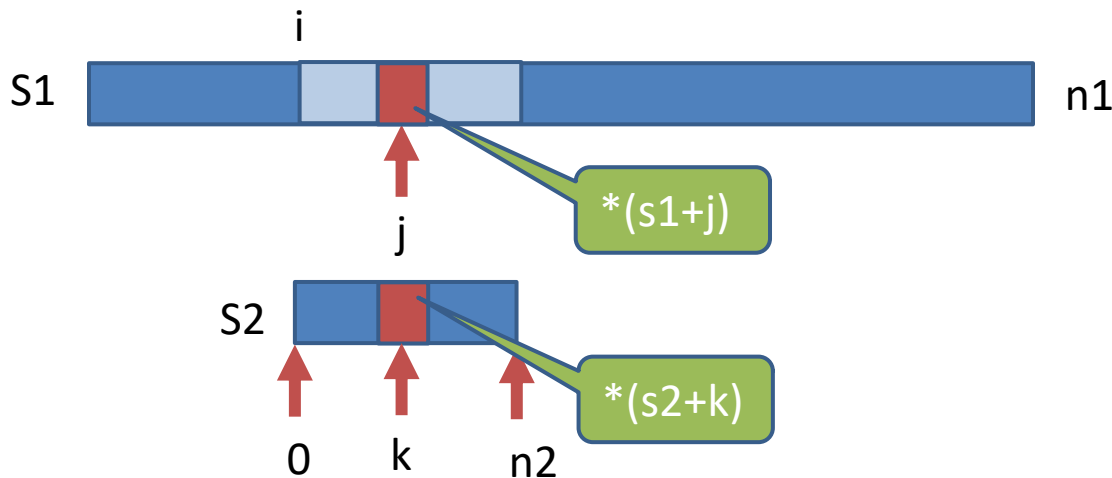
```

int num_substring(void)
{
    int i,j,k;
    int count;

    for (i = 0; i <= (n1-n2); i++){
        count=0;
        for(j = i,k = 0; k < n2; j++,k++){ /*search for the next
string of size of n2*/
            if (*(s1+j)!=*(s2+k)){
                break;
            }else{
                count++;
            }

            if(count==n2){
                total++;          /*find a substring in this
step*/
            }
        }
    }
    return total;
}

```



```

int main(int argc, char *argv[])
{
    int count;

    readf(fp);
    count = num_substring();
    printf("The number of substrings is: %d\n", count);
    return 1;
}

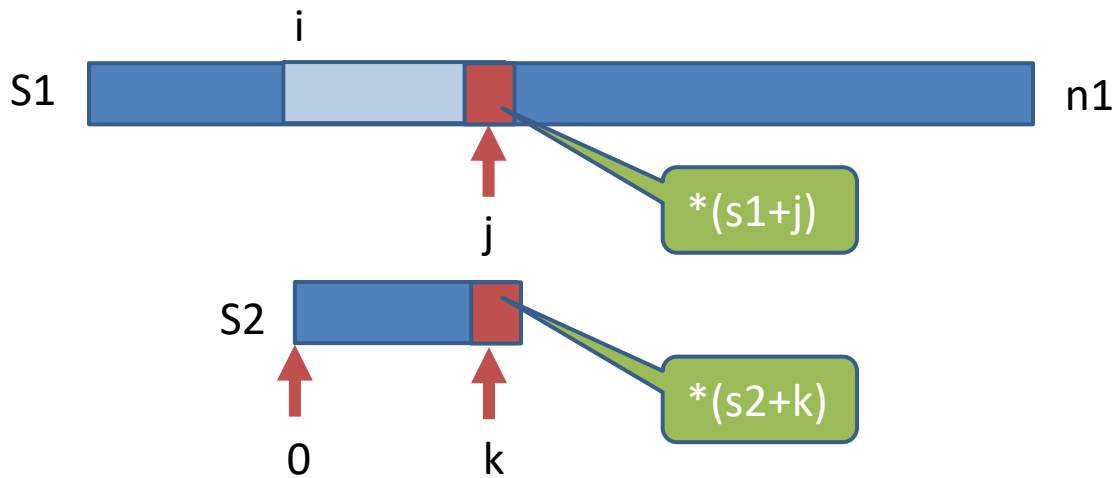
```

```

int num_substring(void)
{
    int i,j,k;
    int count;

    for (i = 0; i <= (n1-n2); i++){
        count=0;
        for(j = i,k = 0; k < n2; j++,k++){ /*search for the next
string of size of n2*/
            if (*(s1+j)!=*(s2+k)){
                break;
            }else{
                count++;
            }
        }
        if(count==n2){
            total++; /*find a substring in this
step*/
        }
    }
    return total;
}

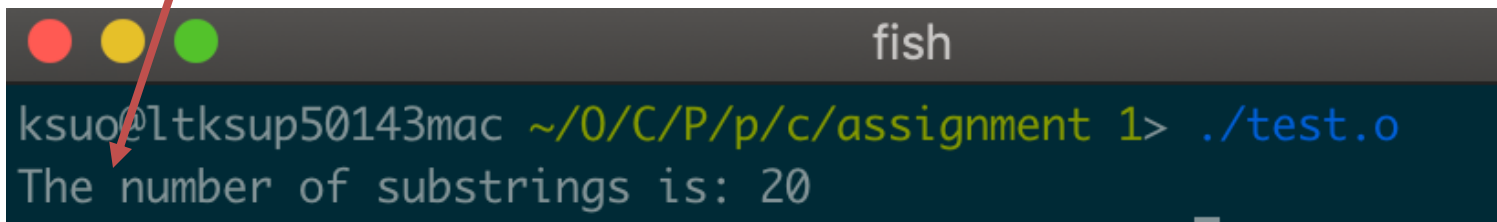
```



Assignment 1

```
int main(int argc, char *argv[])
{
    int count;

    readf(fp);
    count = num_substring();
    printf("The number of substrings is: %d\n", count);
    return 1;
}
```

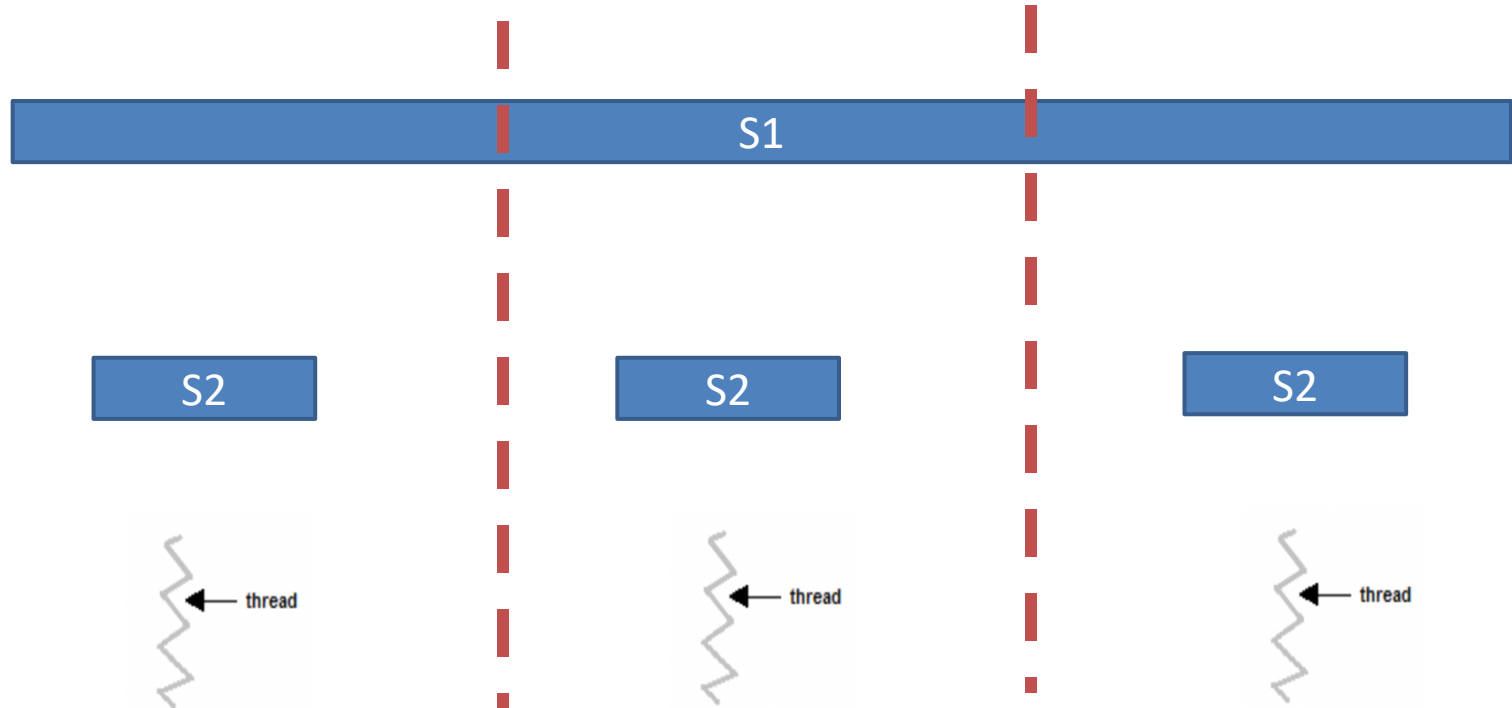


A terminal window titled 'fish' showing the execution of the program. The prompt is 'ksuo@ltsup50143mac ~/O/C/P/p/c/assignment 1>'. The command './test.o' has been entered, and the output is 'The number of substrings is: 20'. A red arrow points from the 'printf' line in the code block above to the output in the terminal.

```
fish
ksuo@ltsup50143mac ~/O/C/P/p/c/assignment 1> ./test.o
The number of substrings is: 20
```

Assignment 1

- Write a parallel program using Pthread based on this sequential solution.



Assignment 1

pthread_create(thread, attr, start_routine, arg)

- creates a thread and makes it executable;
- 1st parameter: pointer to the thread
- 2nd parameter: set attributes to threads
- 3rd parameter: the function for the thread to run
- 4th parameter: parameter for thread function

pthread_exit(status)

- If **main()** finishes before the threads it has created, and exists with the **pthread_exit()**, the other threads will continue to execute. Otherwise, they will be automatically terminated when **main()** finishes



Start from this template using pthread

```
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>
#include <stdlib.h>

int num[3];

void *myThread(void *threadid)
{
    int i = (int)threadid;
    num[i] = i;
    pthread_exit(NULL);
}

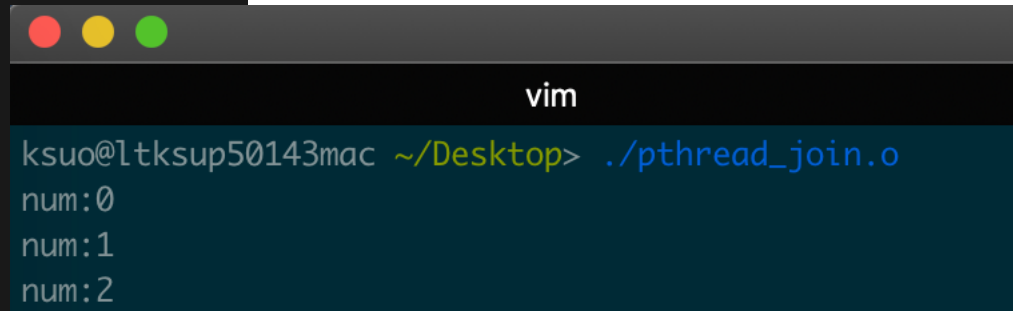
int main()
{
    int rc, t=0;
    pthread_t threads[3];

    for(t=0; t<3; t++){
        rc = pthread_create(&threads[t], NULL, myThread, (void *) (size_t)t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(1);
        }
    }

    for(t=0; t<3; t++){
        pthread_join(threads[t], NULL);
    }

    for(t=0; t<3; t++) {
        printf("num:%d\n", num[t]);
    }
    return 0;
}
```

Pass the parameter
into pthread

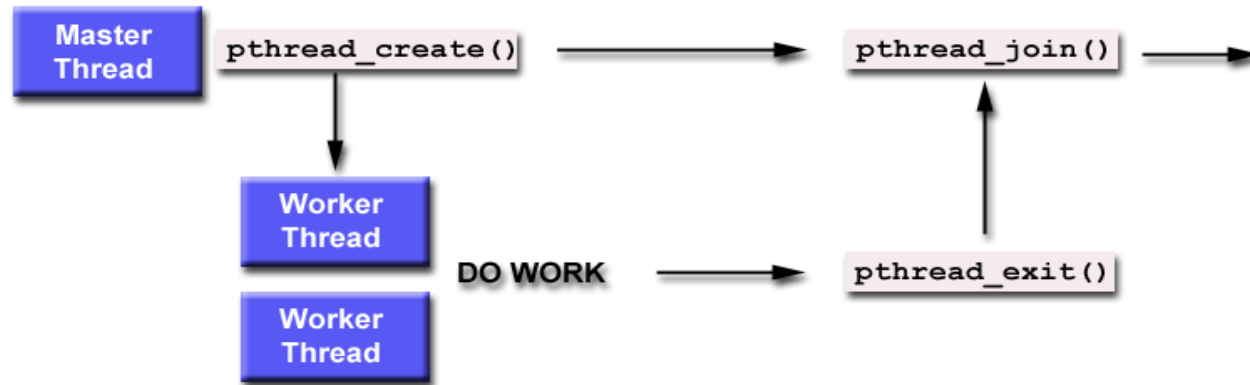


A terminal window titled 'vim' showing the execution of the program. The prompt is 'ksuo@ltsup50143mac ~/Desktop> ./pthread_join.o'. The output is 'num:0', 'num:1', and 'num:2' on separate lines.

Create 3 threads:
The n^{th} thread print number n

<https://github.com/kevinsuo/C3502/blob/master/project2-template-code.c>

Assignment 1



```
Main()
{
```

```
//create multiple threads
pthread_create
```

```
//wait thread to finish in main
pthread_join
```

```
//sum up the number from each thread
sum = num[0]+ num[1]+ ... num[k]
```

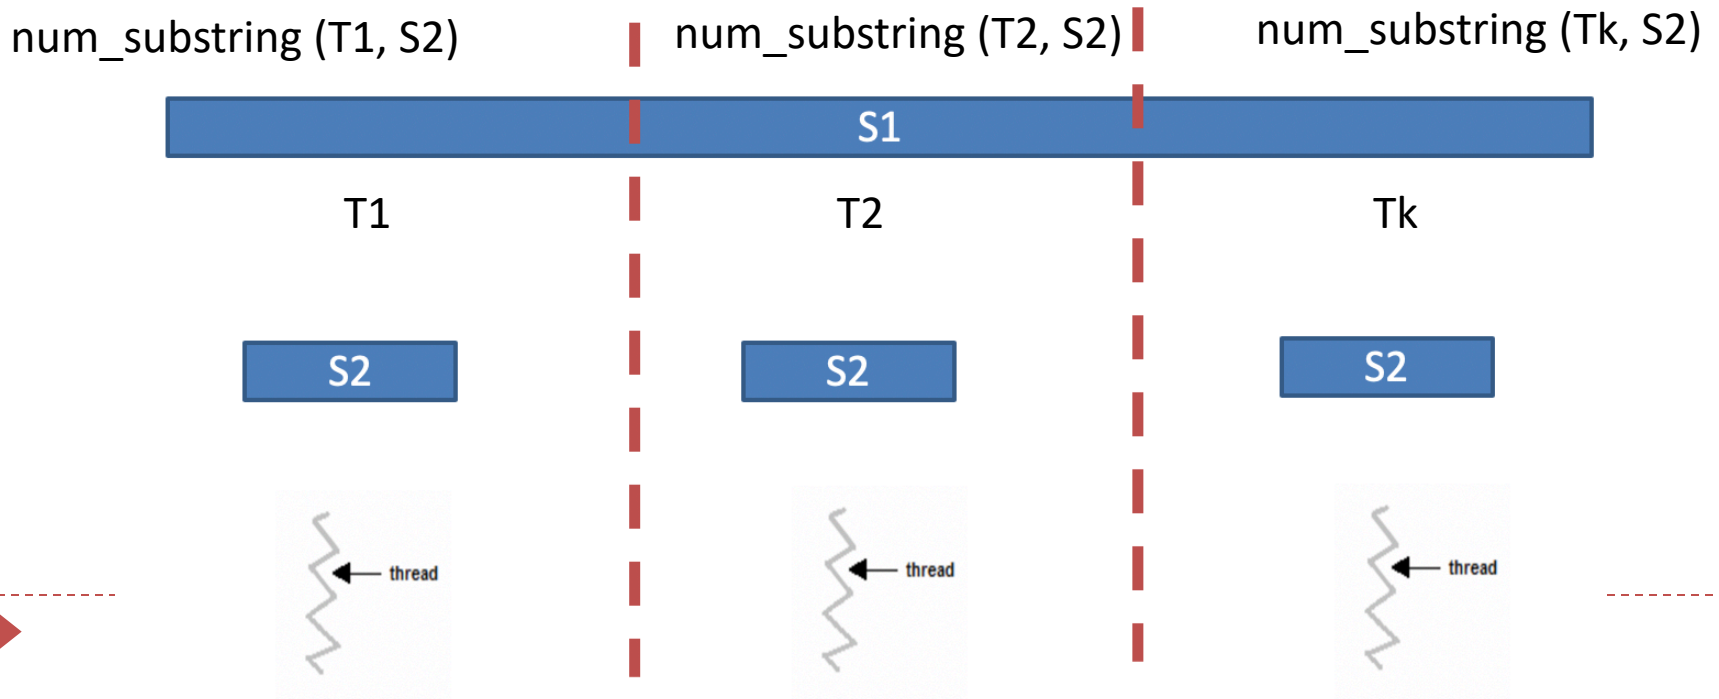
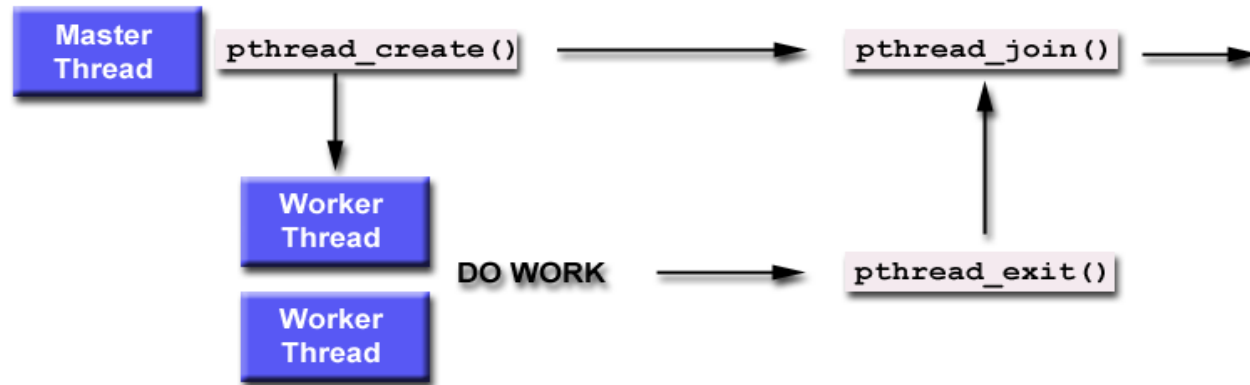
```
}
```

```
thread()
{
```

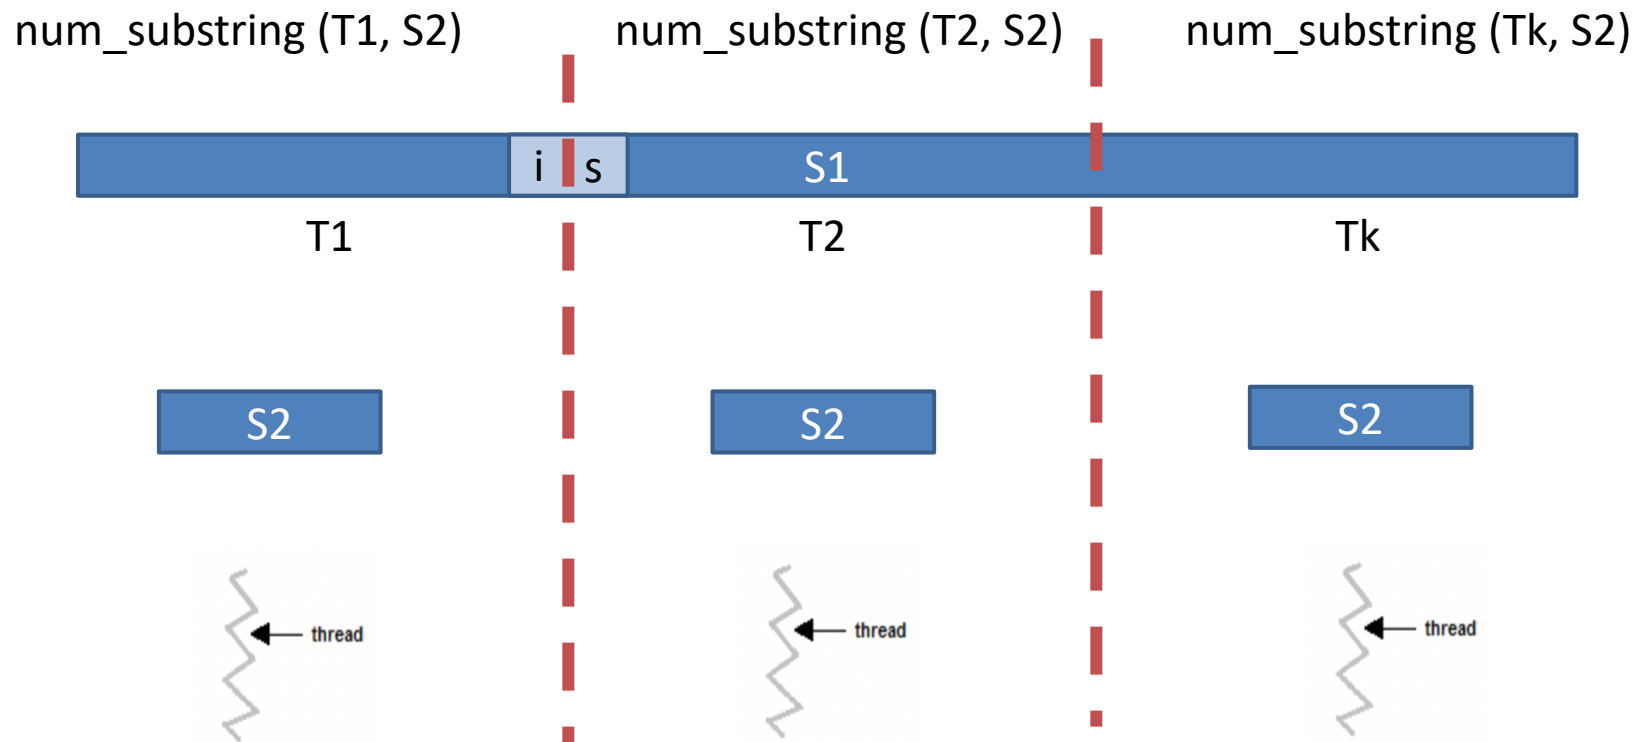
```
//do something
```

```
}
```


Assignment 1



Corner Case in Assignment 1



Verify whether your parallel thread is correct

- Modify the strings.txt by yourself
- Compare the sequential and parallel program results that whether they are the same

```
ksuo@ltkup50143mac ~/O/C/P/p/c/assignment 1> ./test-p.o
This is thread 0
This is thread 3
This is thread 4
This is thread 2
This is thread 1
The number of substrings is : 20
ksuo@ltkup50143mac ~/O/C/P/p/c/assignment 1> ./test.o
The number of substrings is: 20
```



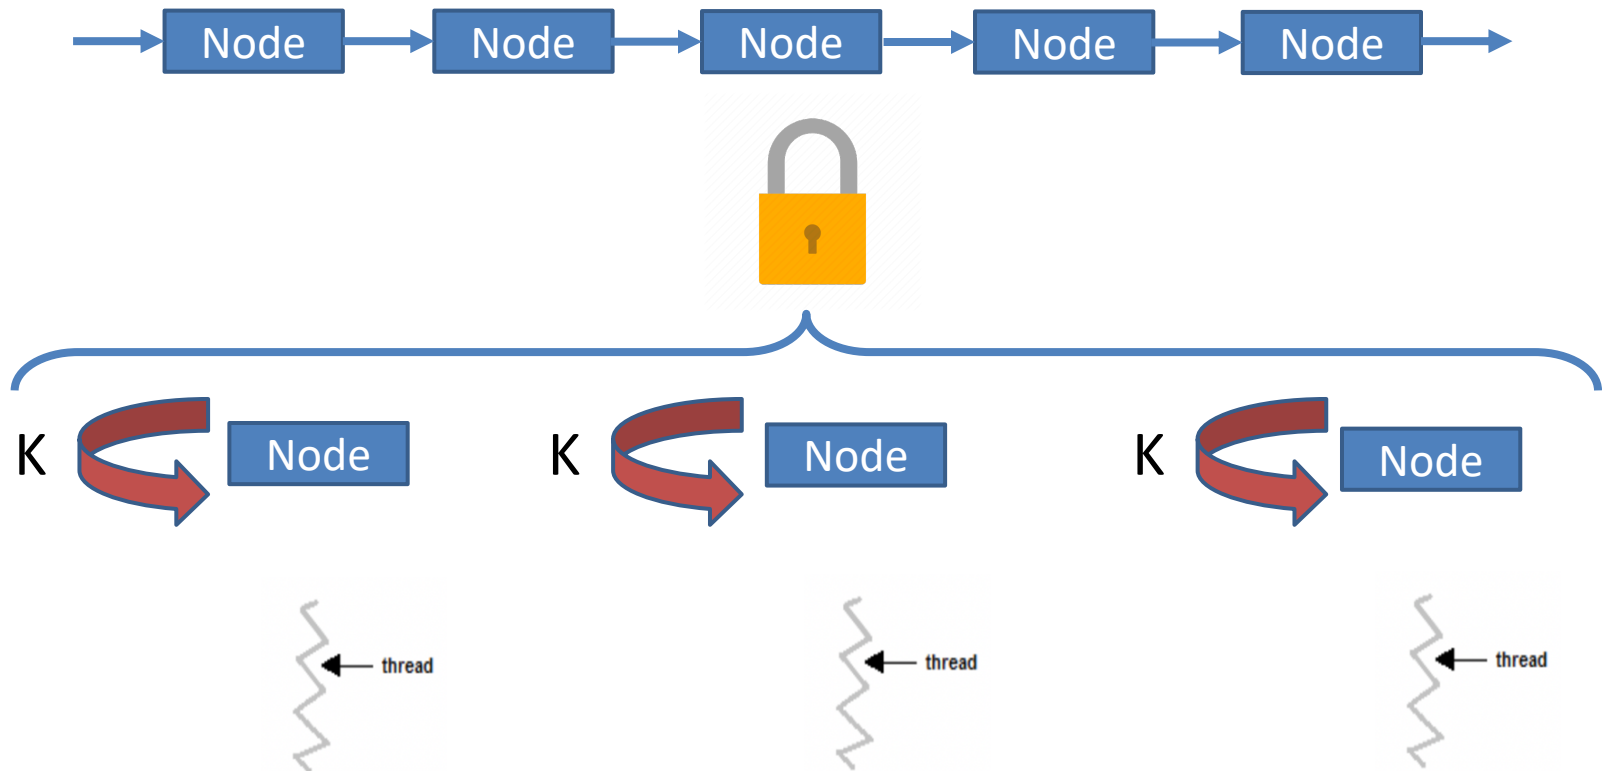
Assignment 2

- Read the following program and modify the program to improve its performance
- <https://github.com/kevinsuo/CS3502/blob/master/project-2-2.c>



Assignment 2

- Each thread creates a data node and attaches it to a global list. This operation is repeated for K times by each thread.



Assignment 2

- Make sure your VM to have more cores

(1) Shutdown your VM and change to VM setting to use 4 vCPUs.

(2) Verify that you VM has 4 vCPUs:

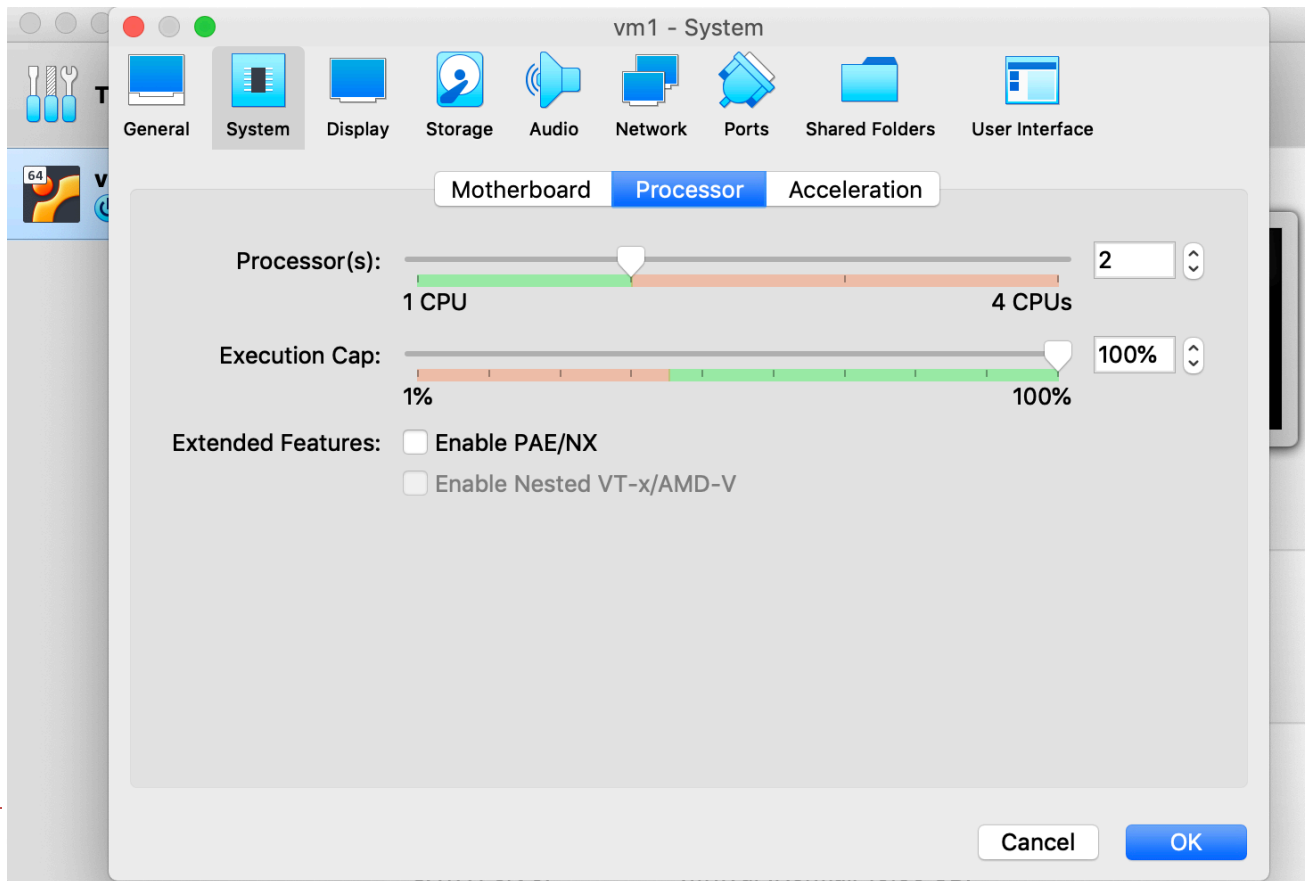
```
$ cat /proc/cpuinfo
```

You should have 4 CPUs (processor: 0-3).



Tips

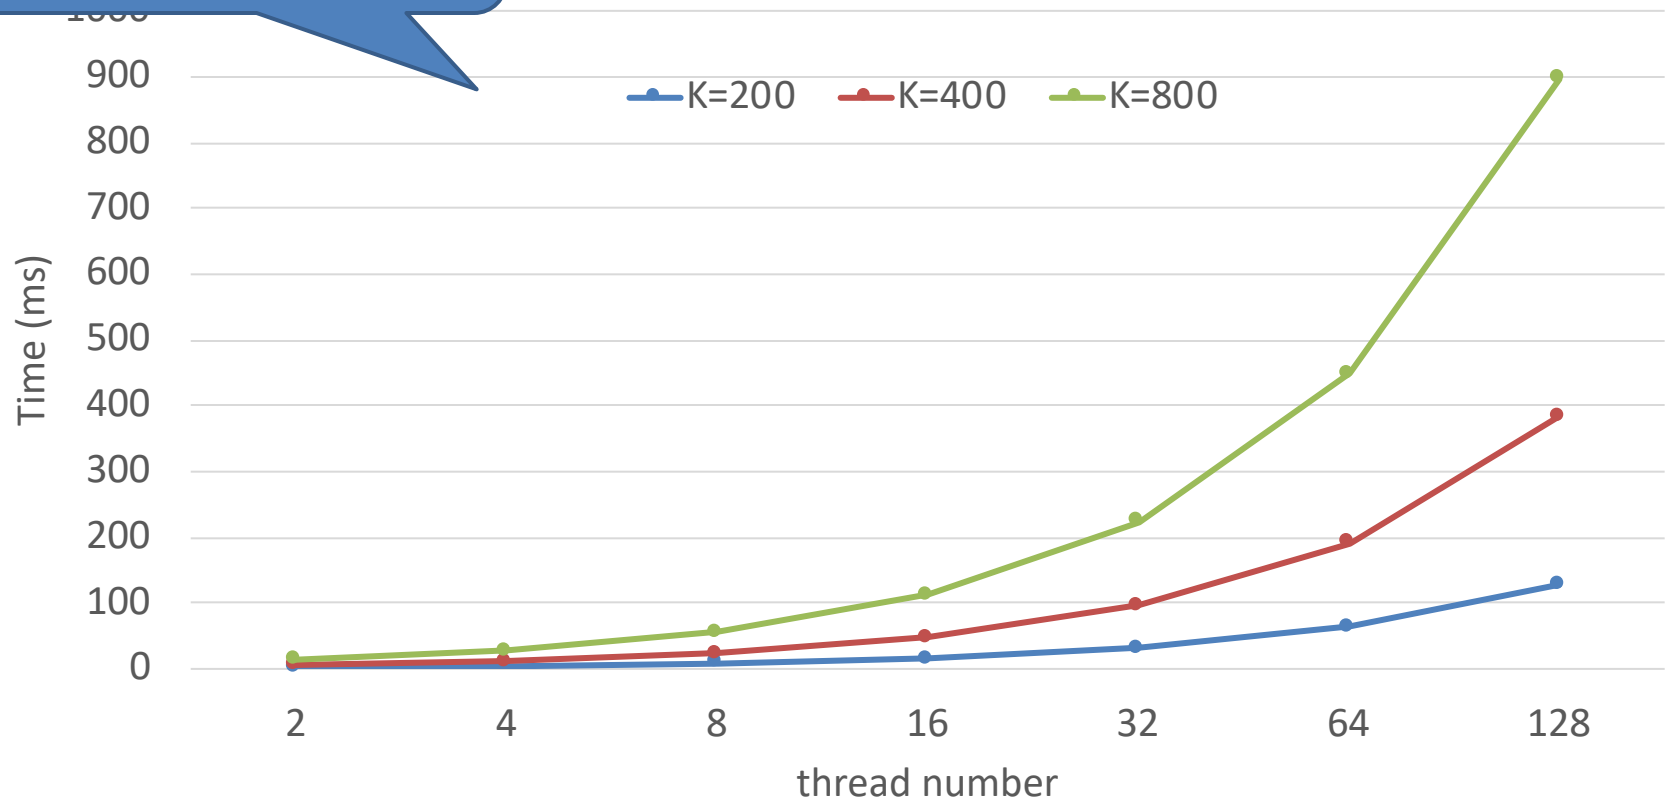
- Set up more CPUs for VM



Test different K and thread num

Could be that one is always better than another

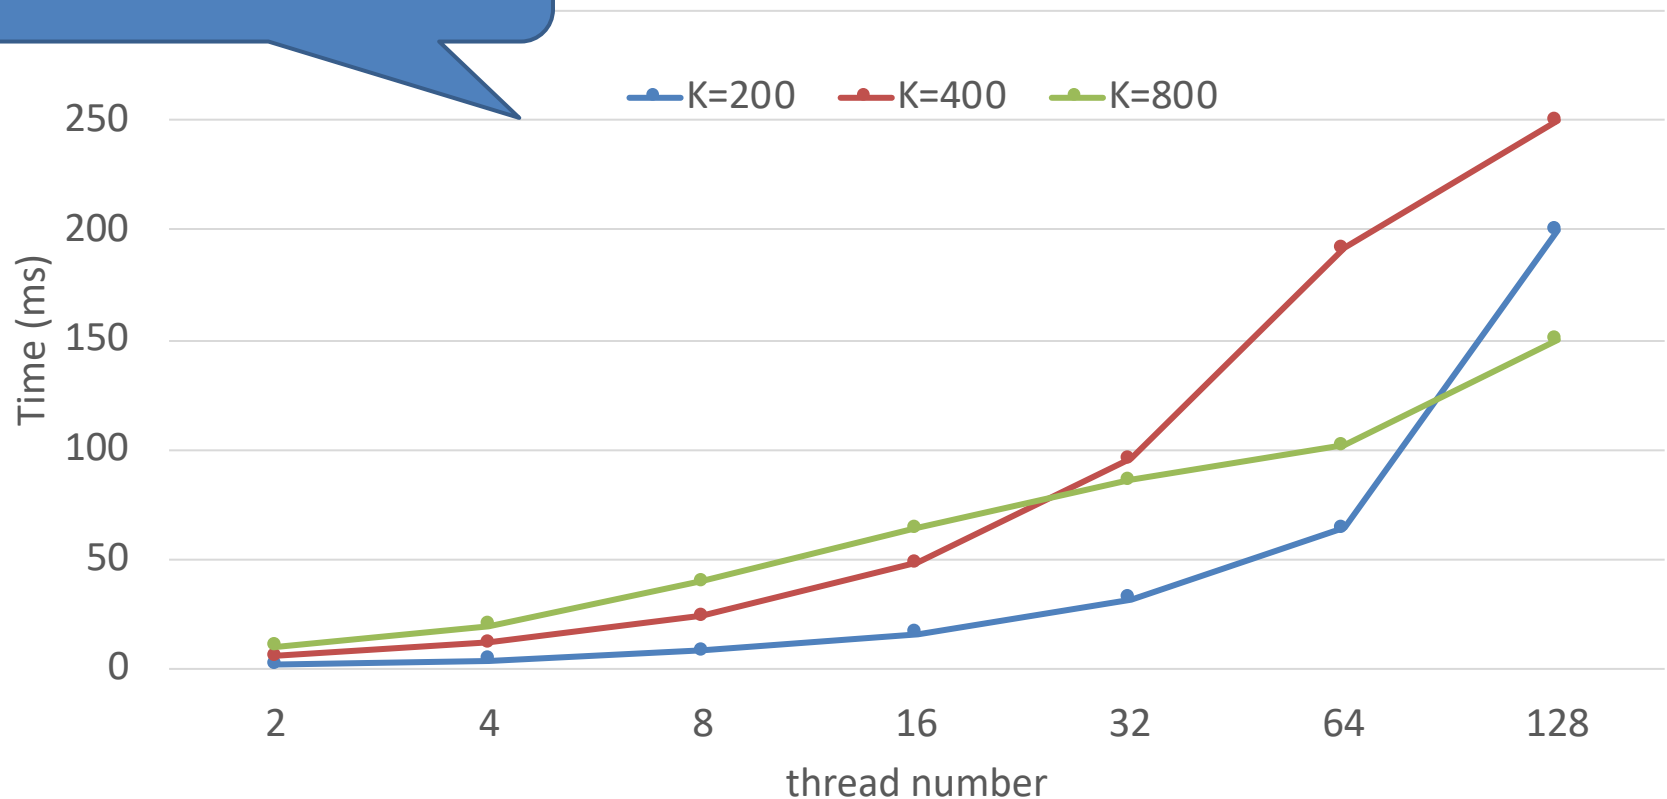
Program execution time



Test different K and thread num

Could be that one is sometimes better than another

Program execution time

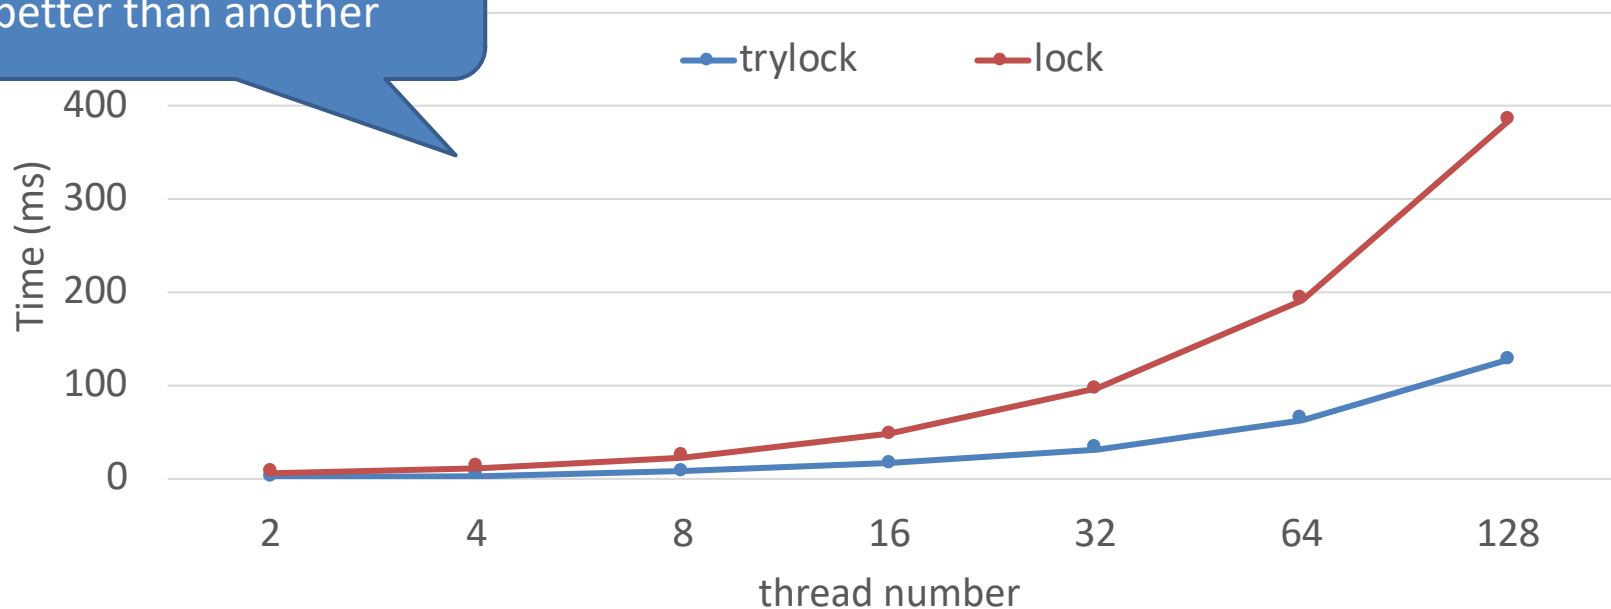


pthread_mutex_trylock vs. pthread_mutex_lock

- The original program uses pthread_mutex_trylock. Will the use of pthread_mutex_lock make a difference? Why?

Could be that one is always better than another

Program execution time

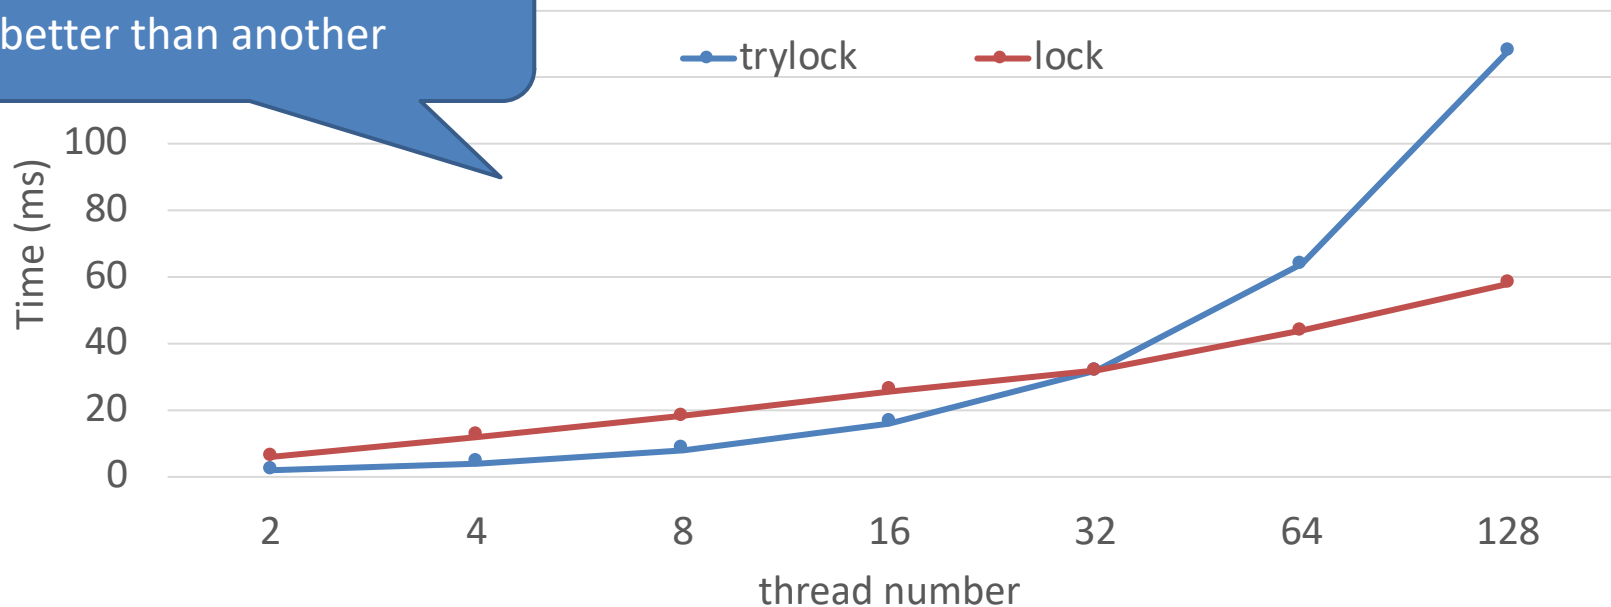


pthread_mutex_trylock vs. pthread_mutex_lock

- The original program uses pthread_mutex_trylock. Will the use of pthread_mutex_lock make a difference? Why?

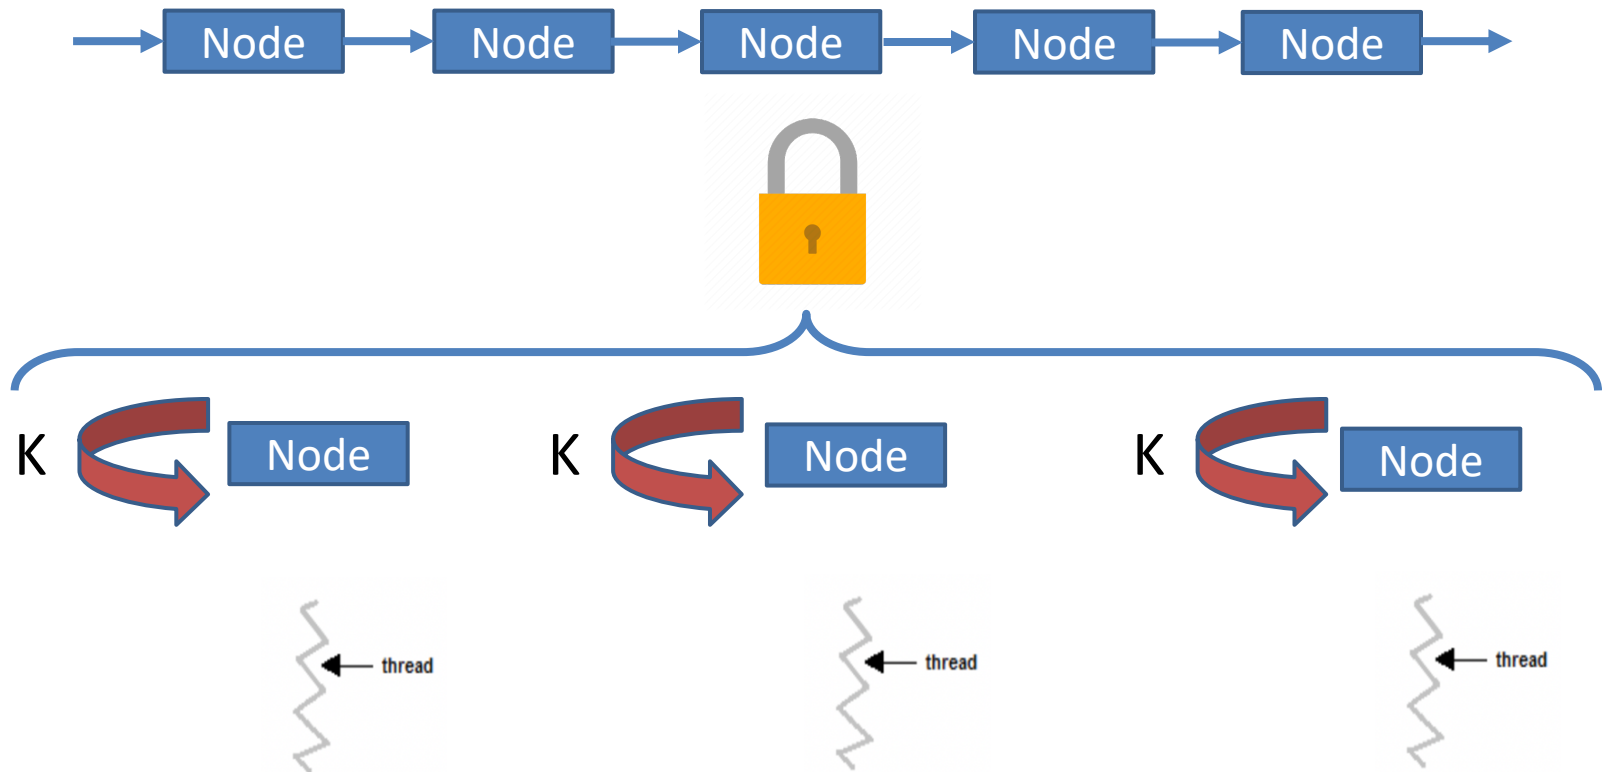
Could be that one is sometimes better than another

Program execution time



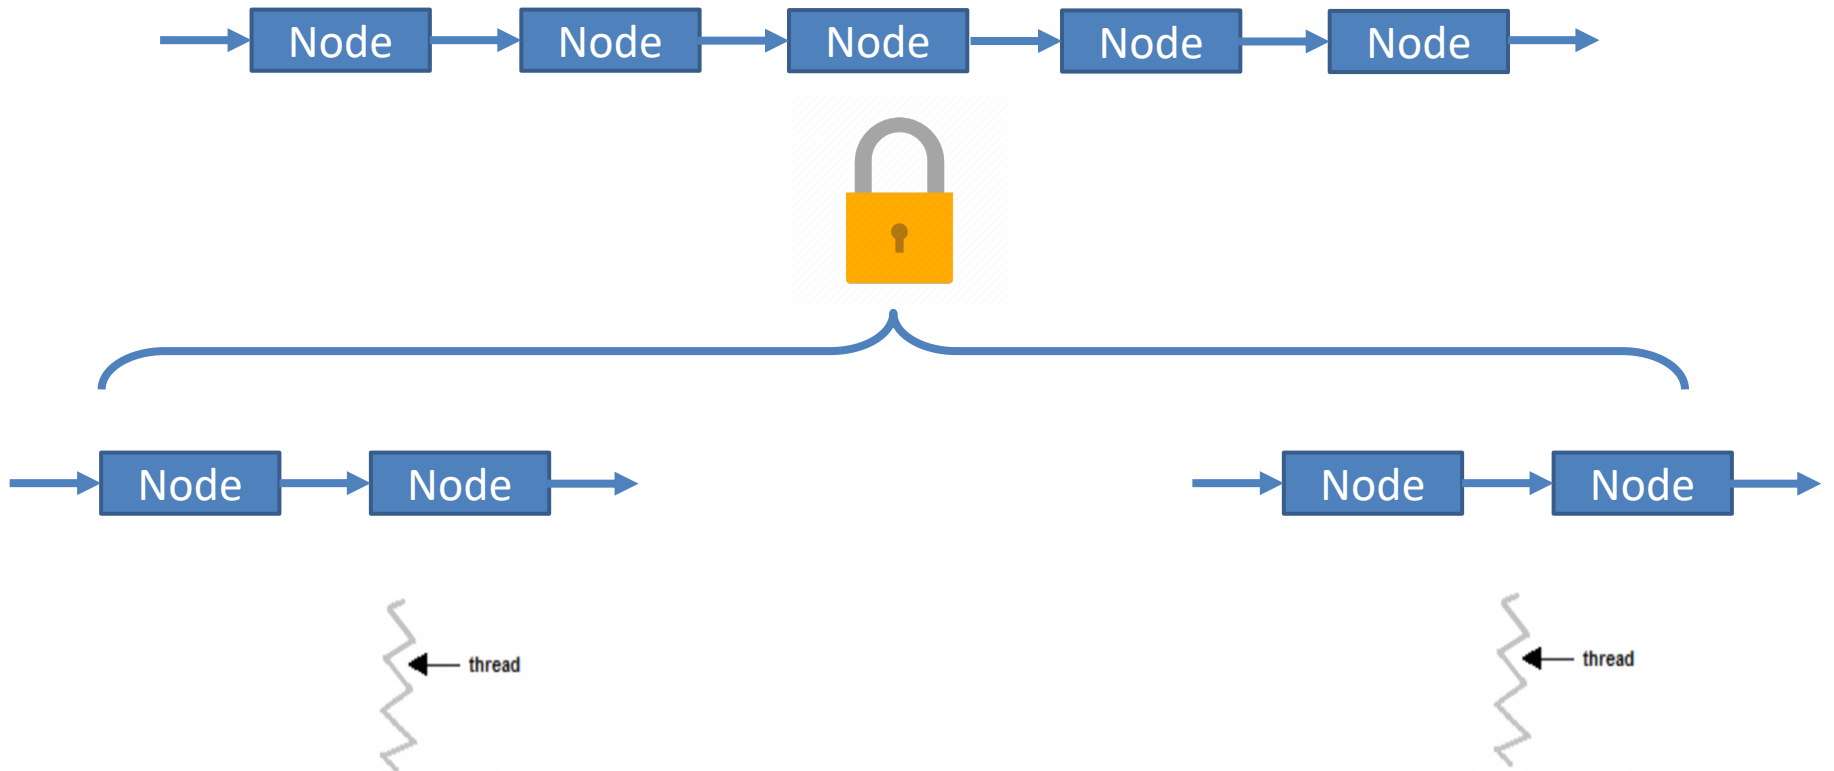
Assignment 2

Old design

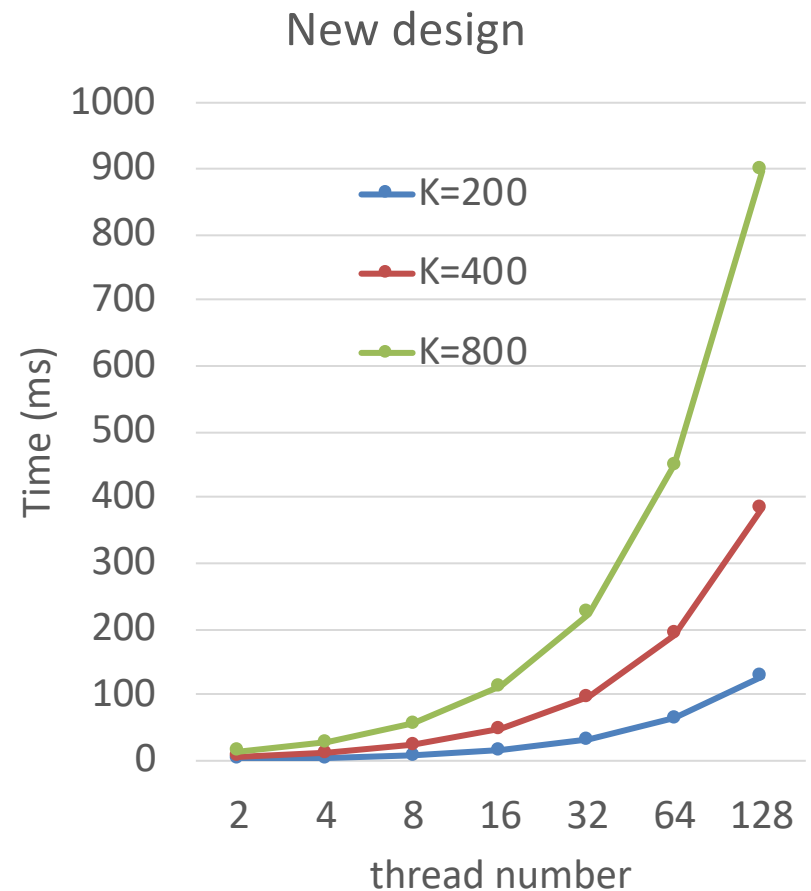
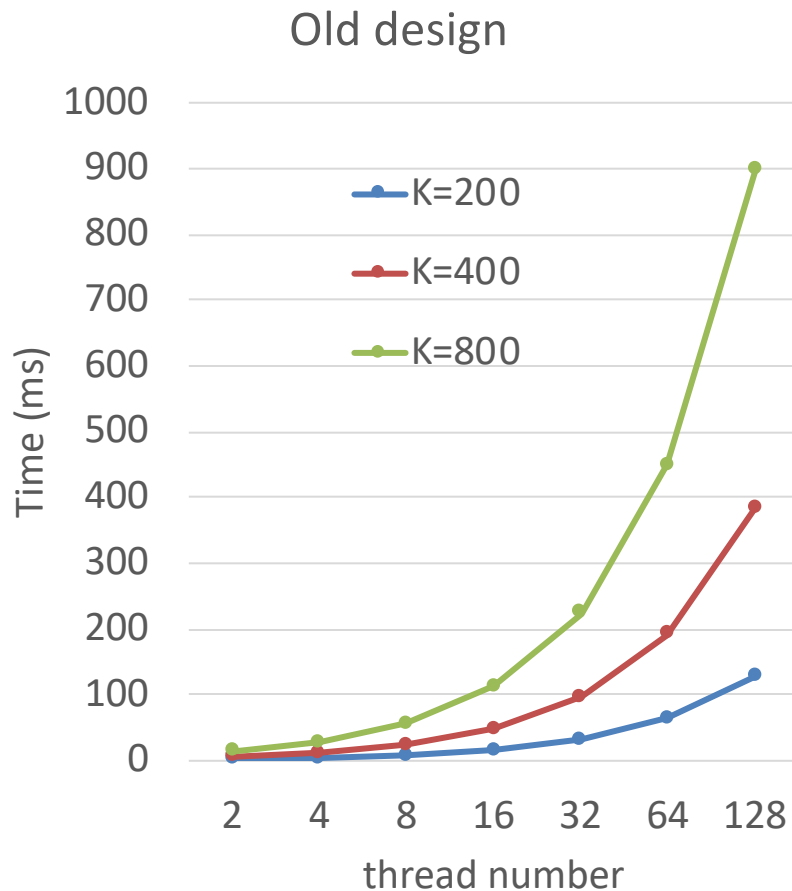


Assignment 2

New design



Performance difference and analysis



Submission

Submit your assignment zip file through D2L using the appropriate assignment link.

For assignment 1, please submit the source code;

For assignment 2, please submit a report with all the figures and analysis included.



CS3502_D2Lname.zip



Questions

- T/Th 2-3pm, J-318
- Skype ID: *suokun.nju*, share home screen, only voice, no video, no remote control and privacy issue. No time or location limitation.

