

# CS 7172

## Parallel and Distributed Computation

### Distributed Lock

**Kun Suo**

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

# Outline

---


- Computer networks, primarily from an application perspective
- Protocol layering
- Client-server architecture
- End-to-end principle
- TCP
- Socket programming



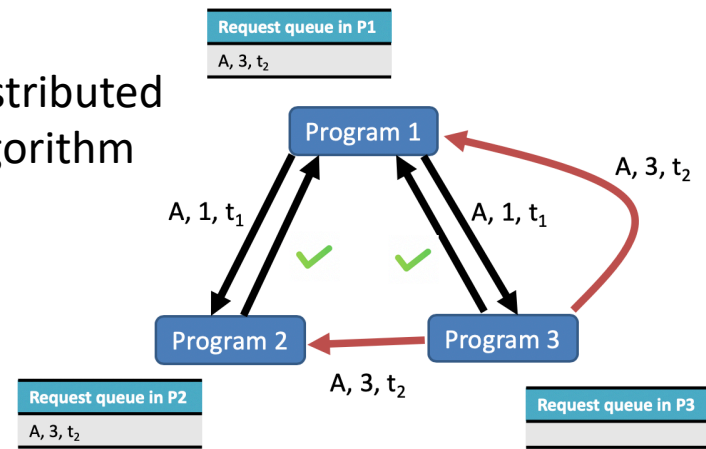
# Revisit Distributed Mutex

- In a distributed system, the method to achieve access to exclusive resource is called *Distributed Mutual Exclusion*


Centralized algorithm



Distributed algorithm



Token Ring Algorithm



mputation

# Race condition

[https://github.com/kevinsuo/CS3502/blob/master/race\\_condition.c](https://github.com/kevinsuo/CS3502/blob/master/race_condition.c)

- A race condition occurs when two or more threads **access** shared data and they try to **change** it at the **same time**.
- The **order** in which the threads attempt to access the shared data makes the results unpredictable

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

int counter = 0;

void *compute()
{
    int i = 0;
    while (i < 100) {
        counter = counter + 1;
        i++;
    }
    printf("Counter value: %d\n", counter);
}

int main()
{
    pthread_t thread1, thread2;

    pthread_create(&thread1, NULL, compute, (void *)&thread1);
    pthread_create(&thread2, NULL, compute, (void *)&thread2);

    pthread_exit(NULL);
    exit(0);
}
```

Race condition occurs for variable counter

```
pi@raspberrypi ~/Downloads> ./race_condition.o
Counter value: 100
Counter value: 200
```

Seem nothing wrong?

# Race condition example

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

int counter = 0;

void *compute()
{
    int i = 0;
    while (i < 10000) {
        counter = counter + 1;
        i++;
    }
    printf("Counter value: %d\n", counter);
}

int main()
{
    pthread_t thread1, thread2, thread3, thread4, thread5;

    pthread_create(&thread1, NULL, compute, (void *)&thread1);
    pthread_create(&thread2, NULL, compute, (void *)&thread2);
    pthread_create(&thread3, NULL, compute, (void *)&thread3);
    pthread_create(&thread4, NULL, compute, (void *)&thread4);
    pthread_create(&thread5, NULL, compute, (void *)&thread5);

    pthread_exit(NULL);
    exit(0);
}
```

Increase the loop number

Add more threads

```
pi@raspberrypi ~/Downloads> ./race_condition.o
Counter value: 14467
Counter value: 10410
Counter value: 12080
Counter value: 22745
Counter value: 32725
```


Weird results!



# Critical section

---

- A section of code in a concurrent task that **modifies or accesses** a resource shared with another task.
- Examples
  - A piece of code that reads from or writes to a shared memory region
  - Or a code that modifies or traverses a shared linked list.



# Critical section example

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

int counter = 0;

void *compute()
{
    int i = 0;
    while (i < 100) {
        counter = counter + 1;
        i++;
    }
    printf("Counter value: %d\n", counter);
}

int main()
{
    pthread_t thread1, thread2;

    pthread_create(&thread1, NULL, compute, (void *)&thread1);
    pthread_create(&thread2, NULL, compute, (void *)&thread2);



    pthread_exit(NULL);
    exit(0);
}
```

Critical section: All threads read and write the shared counter



# Lock (mutual exclusion)

- A lock (mutual exclusion) is a synchronization mechanism for enforcing limits on access to a resource in an environment where there are many threads of execution
- Types of mutual mechanism:
  - Busy-waiting, e.g., spinlock
  - Sleep and wakeup







# Deadlocks

---

- When two or more processes stop making progress indefinitely because they are **all waiting for each other** to do something.
  - If process A waits for process B to release a resource, and
  - Process B is waiting for process A to release another resource at the same time.
  - In this case, neither A nor B can proceed because both are waiting for the other to proceed.



# Deadlock example



## Thread 1

```
pthread_mutex_lock(&m1);  
/* use resource 1 */  
pthread_mutex_lock(&m2);  
/* use resources 1 and 2 */  
do_something();  
pthread_mutex_unlock(&m2);  
pthread_mutex_unlock(&m1);
```



## Thread 2


```
pthread_mutex_lock(&m2);  
/* use resource 2 */  
pthread_mutex_lock(&m1);  
/* use resources 1 and 2 */  
do_something();  
pthread_mutex_unlock(&m1);  
pthread_mutex_unlock(&m2);
```



# What is a distributed lock?

---

- Distributed lock refers to a type of lock in which the system is deployed on multiple machines to achieve mutual exclusion between multiple processes in distributed system.
- In order to ensure that multiple processes can see the lock, such the lock is stored in the public storage (such as Redis, Memcache, database and other three-party storage), so that multiple processes can concurrently access the same critical resource, and only one process can access the shared resource at the same time.



# Example

---

- Our online store sells two diamonds
- 5 users (A, B, C, D, E) order at the same time
- If you design the online system, who will get the diamond?



# Example

---

- Different e-commerce companies adopt different strategies:
  - Some e-commerce companies determine who can make a successful purchase based on the order time, First-come-first-serve
  - Some e-commerce companies determine who can make a successful purchase based on the payment time
  - ...
- No matter which policy, the online store must make sure every diamond cannot be sold to two users



# Example

---

- Different e-commerce companies adopt different strategies:
  - Some e-commerce companies determine who can make a successful purchase based on the order time, First-come-first-serve
  - Some e-commerce companies determine who can make a successful purchase based on the payment time
  - ...
- No matter which policy, the online store must make sure every diamond cannot be sold to two users. In large-scale distributed systems, we use *distributed locks* to realize that.



# Distributed Locks

---

- 3 ways to implement distributed locks:
  1. Distributed locks based on databases, here databases refer to relational databases;
  2. Distributed locks based on caches;
  3. Distributed locks based on ZooKeeper.



# Some Key Questions for Distributed Locks

---


- Mutual exclusion:
  - A distributed lock should be able to guarantee that a resource or a method can only be operated by one process of one machine at a time.
- Preventing deadlock:
  - Even if one process does not actively unlock due to crash while holding the lock, it can guarantee that other processes can obtain the lock.
- Reentrancy:
  - A critical resource can be accessed multiple times when the process has not released the lock.
- Performance:
  - Have a high availability for lock acquisition and release.





# 1. Distributed locks based on databases

- The easiest way is to create a **lock table**, and then manipulate the data in the table to achieve distributed locks
- Lock a resource → add a record to this table;  
Release the lock → delete this record.
- The database has consistent constraint on shared resources. If multiple requests are submitted to the database at the same time, the database will ensure that **only one** operation will be successful, and the thread that successfully operates will obtain a lock.



# 1. Distributed locks based on databases

---


- Problem for database:
  - The database needs to read/write the hard disk. Frequent interaction with the database will cause large I/O overhead
  - Only suitable for scenarios with low concurrency and low performance requirements



# 1. Distributed locks based on databases

- Example:

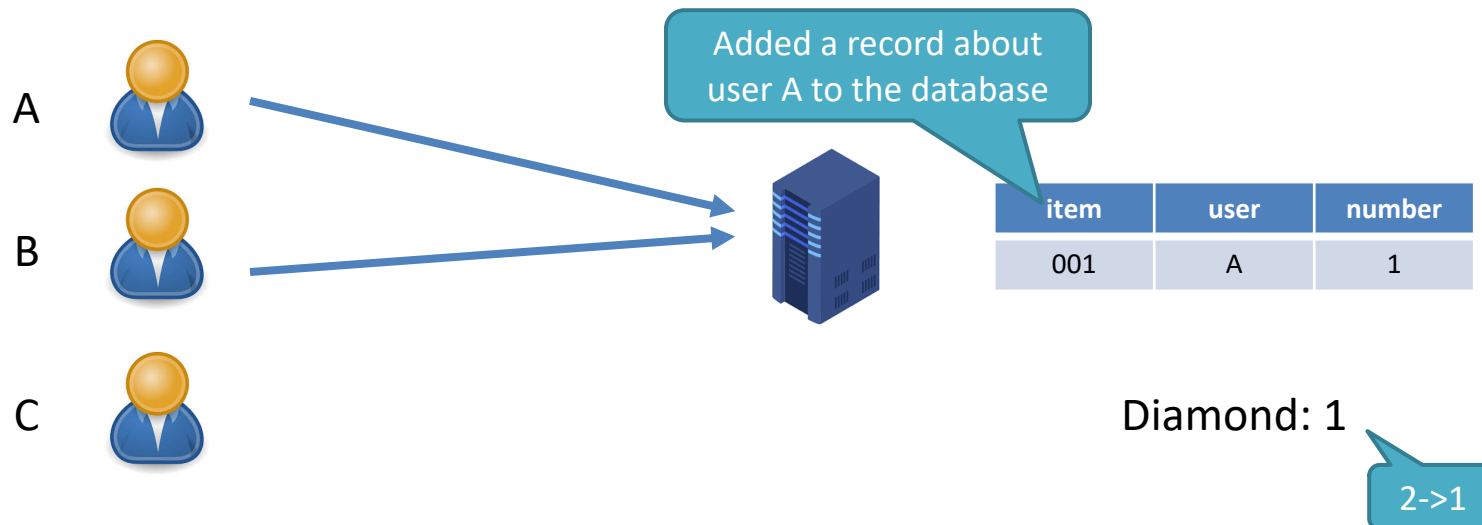
- Our online store sells 2 diamonds
- 5 users (A, B, C, D, E) order at the same time
- A orders 1 diamond, B orders 2 diamonds, C orders 1 diamond



# 1. Distributed locks based on databases

- Example:



- Our online store sells 2 diamonds
- 5 users (A, B, C, D, E) order at the same time
- A orders 1 diamond, B orders 2 diamonds, C orders 1 diamond



# 1. Distributed locks based on databases

- Example:

- Our online store sells 2 diamonds
- 5 users (A, B, C, D, E) order at the same time
- A orders 1 diamond, B orders 2 diamonds, C orders 1 diamond




Diamond: 1



# 1. Distributed locks based on databases

- Example:

- Our online store sells 2 diamonds
- 5 users (A, B, C, D, E) order at the same time
- A orders 1 diamond, B orders 2 diamonds, C orders 1 diamond

A diagram showing a database record. A speech bubble from the database table points to the 'user' column, containing the text "Write B's record into database".

item	user	number
002	B	2



Diamond: 1



# 1. Distributed locks based on databases

- Example:

- Our online store sells 2 diamonds
- 5 users (A, B, C, D, E) order at the same time
- A orders 1 diamond, B orders 2 diamonds, C orders 1 diamond




Diamond: 1



# 1. Distributed locks based on databases

- Example:

- Our online store sells 2 diamonds
- 5 users (A, B, C, D, E) order at the same time
- A orders 1 diamond, B orders 2 diamonds, C orders 1 diamond



Write C's record into database

item	user	number
003	C	1

Diamond: 0







# 1. Distributed locks based on databases

- Example:

- Our online store sells 2 diamonds
- 5 users (A, B, C, D, E) order at the same time
- A orders 1 diamond, B orders 2 diamonds, C orders 1 diamond



Diamond: 0



# 1. Distributed locks based on databases

---

- Advantages:
  - Easy to implement a distributed lock, just create a lock table and create a record in the lock table for the applicant: get the lock when the record is successfully established, release the lock when the record is deleted

# 1. Distributed locks based on databases

---

- Disadvantages:
  - Single point of failure. Once the database is unavailable, the entire system will crash.
  - Deadlock problem. The lock based on database has no expiration time and the process that has not acquired the lock has to wait.
    - ▶ If the lock holder falls into loop or the unlock operation fails, the lock record will always exist in the database, and other processes cannot acquire the lock.



## 2. Distributed locks based on caches

---

- Based on database → introduce lots of I/O operation on hard disk
- Cache-based means that the data is stored in the memory and does not need to be written to disk, which reduces I/O reading and writing overhead.

## 2. Distributed locks based on caches

---


- How the cache-based lock works?
- Key-value store:
  - Key: lock ID
  - Value: `currentTime + timeout`, after the process gets the lock, the lock will be released by the system when the time is out
  - Function `setnx(key, value)` works on the table
    - ▶ Return 1 means getting the lock and writing one record into the table
    - ▶ Return 0 means the lock is used by other processes

Key	Value
1	
2	
3	
4	
...	

## 2. Distributed locks based on caches

- Example:

- Our online store sells 2 diamonds
- 5 users (A, B, C, D, E) order at the same time
- A orders 1 diamond, B orders 2 diamonds, C orders 1 diamond



Key	Value
1	$T_1 + T_{out}$
...	


Diamond: 2



## 2. Distributed locks based on caches

- Example:

- Our online store sells 2 diamonds
- 5 users (A, B, C, D, E) order at the same time
- A orders 1 diamond, B orders 2 diamonds, C orders 1 diamond




Key	Value
1	$T_1 + T_{out}$
...	

Diamond: 2

## 2. Distributed locks based on caches

- Example:

- Our online store sells 2 diamonds
- 5 users (A, B, C, D, E) order at the same time
- A orders 1 diamond, B orders 2 diamonds, C orders 1 diamond






## 2. Distributed locks based on caches

- Example:

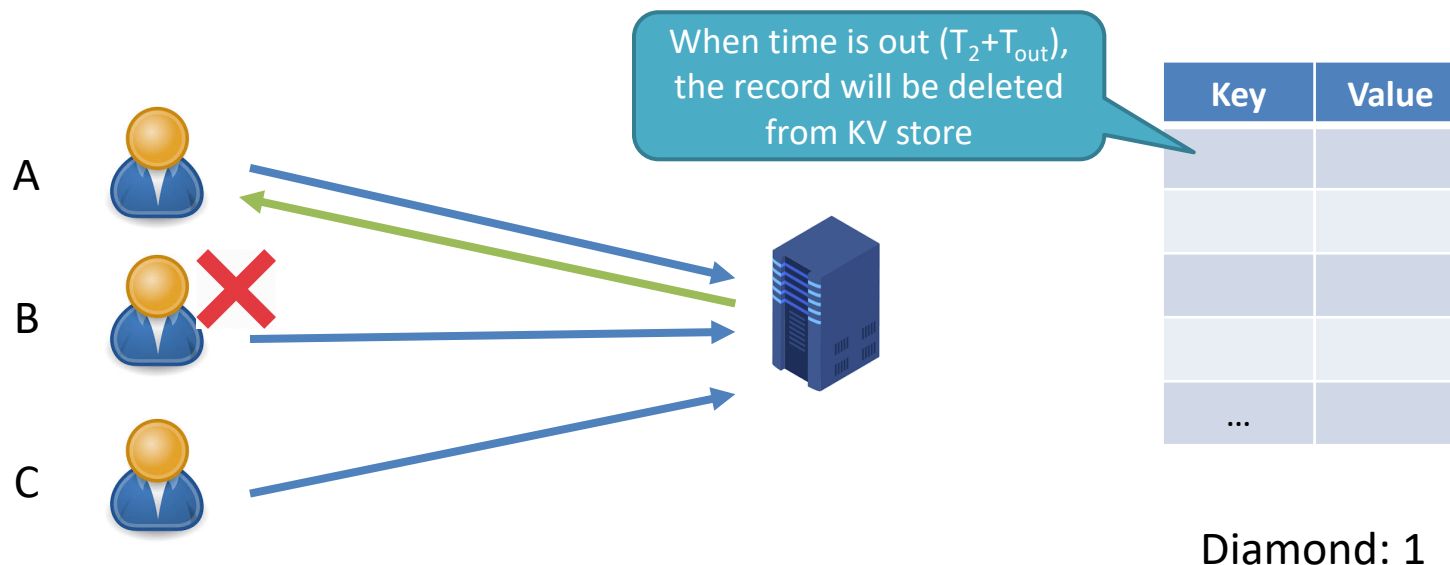
- Our online store sells 2 diamonds
- 5 users (A, B, C, D, E) order at the same time
- A orders 1 diamond, B orders 2 diamonds, C orders 1 diamond



## 2. Distributed locks based on caches

- Example:


- Our online store sells 2 diamonds
- 5 users (A, B, C, D, E) order at the same time
- A orders 1 diamond, B orders 2 diamonds, C orders 1 diamond



## 2. Distributed locks based on caches

- Example:

- Our online store sells 2 diamonds
- 5 users (A, B, C, D, E) order at the same time
- A orders 1 diamond, B orders 2 diamonds, C orders 1 diamond



Key	Value
1	$T_3 + T_{out}$
...	

Diamond: 0

## 2. Distributed locks based on caches

---

- Advantages:
  - Better performance. Data is stored in memory instead of disk, avoiding frequent IO operations.
  - Many caches can be deployed across clusters, avoiding single points of failure.
  - It can directly set the timeout to control the release of locks. More flexible.



## 2. Distributed locks based on caches

---

- Disadvantages:
  - Simply using timeout might cause failure when releasing locks
    - ▶ One process executes longer (100ms) than it expects (50ms), while the lock is timeout (in 50ms). It will cause error when releasing locks earlier.

### 3. Distributed locks based on ZooKeeper

---

- ZooKeeper implements distributed lock based on tree data storage structure
- ZooKeeper consists 4 nodes:
  - Persistent node. This will always exist in ZooKeeper.
  - Persistent order node. When a node is created, ZooKeeper assigns numbers to the nodes according to when the node is created.
  - Temporary node. This will be deleted when not used.
  - Temporary sequence node: When a node is created, ZooKeeper assigns numbers to the nodes according to when the node is created.




# 3. Distributed locks based on ZooKeeper

---

- Example:

Step 1: Under the persistent node `shared_lock`, create a temporary sequential node for each process (LockNode i).




### 3. Distributed locks based on ZooKeeper

---

- Example:

Step 2: Each process gets a list of all temporary nodes in the shared\_lock directory

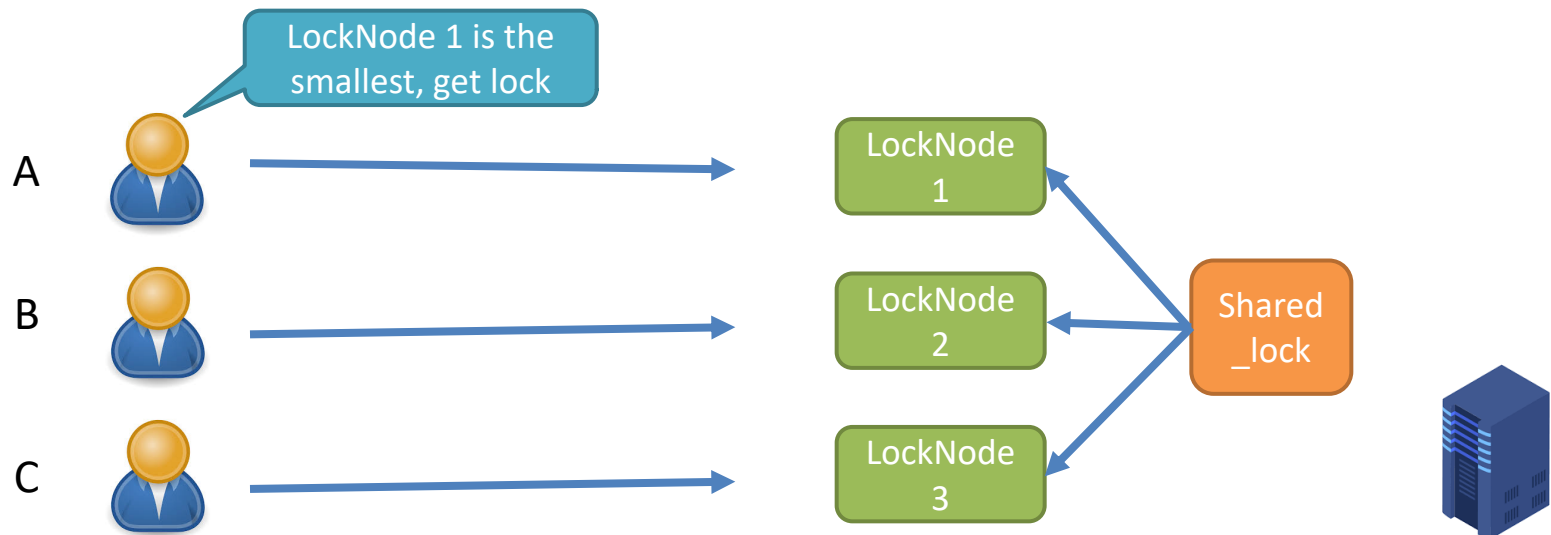




# 3. Distributed locks based on ZooKeeper

- Example:

Step 3: Each node determines whether its own number is the smallest of all nodes under `shared_lock`. If it is the smallest, just acquires the lock.

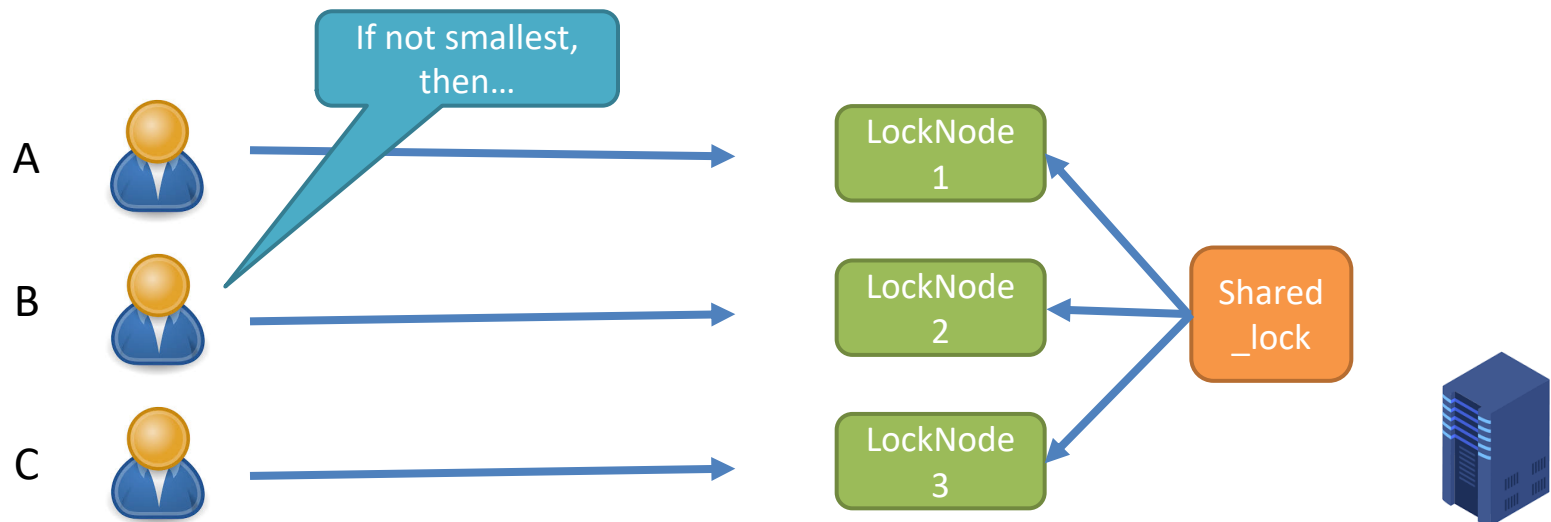


# 3. Distributed locks based on ZooKeeper

- Example:

Step 4: If the temporary node number to this process is not the smallest:

- ▶ If the request is to read: check whether there exist write requests with smaller number, if so, wait;
- ▶ If the request is to write: check whether there exist read requests with smaller number, if so, wait;



# 3. Distributed locks based on ZooKeeper


---

- Advantages:
  - ZooKeeper locks can solve problems such as single points of failure, non-reentrant, deadlock, etc.
  - ZooKeeper provides APIs for distributed locks, easy to implement

# 3. Distributed locks based on ZooKeeper

---

- Disadvantages:
  - Need to add and delete temporary nodes frequently
  - Low performance compared to Cache locks



# Comparison

---

	Distributed locks implementation
Design idea (from easy to hard)	Databased based locks > Cached based locks > ZooKeeper based locks
Implementation difficulty (from low to high)	ZooKeeper based locks < Cached based locks < Databased based locks
Performance (from low to high)	Databased based locks < ZooKeeper based locks < Cached based locks
Reliability (from low to high)	Databased based locks < Cached based locks < ZooKeeper based locks

