

CS 7172

Parallel and Distributed Computation

MapReduce

Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

Outline

- Computer networks, primarily from an application perspective
- Protocol layering
- Client-server architecture
- End-to-end principle
- TCP
- Socket programming



Processing Data on a single Machine

- Data on file system on disk
- Simple processing flow:
 1. Read into memory
 2. Process data (apply some function)
 3. Write back to disk



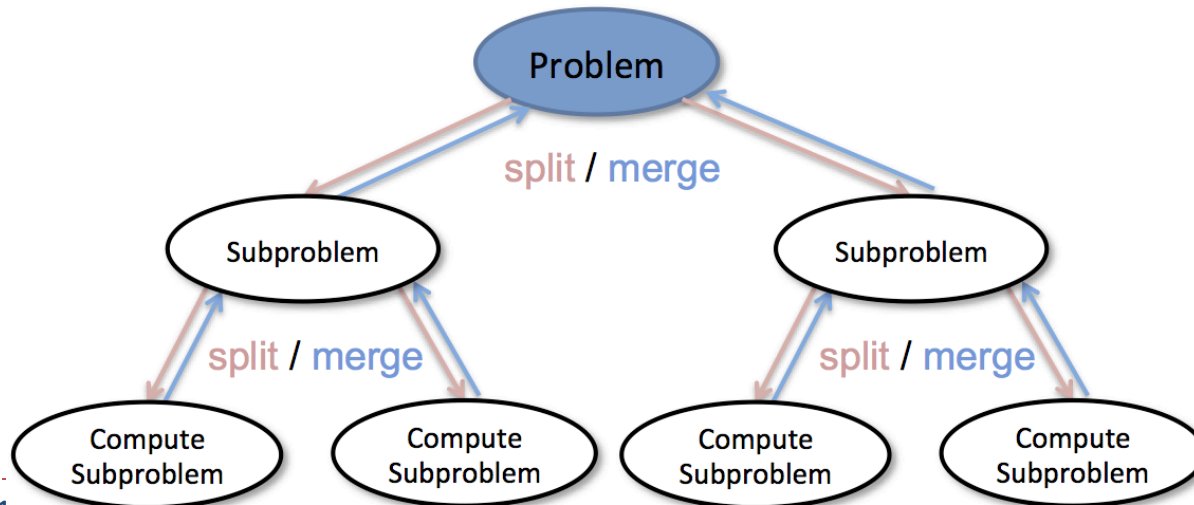
Why Multiple Servers?

- Data size larger than disk capacity
- Sequentially processing data can be slow
 - Limited by disk read/write speeds
 - Parallel processing can significantly reduce time
- Single point of failure with a single server



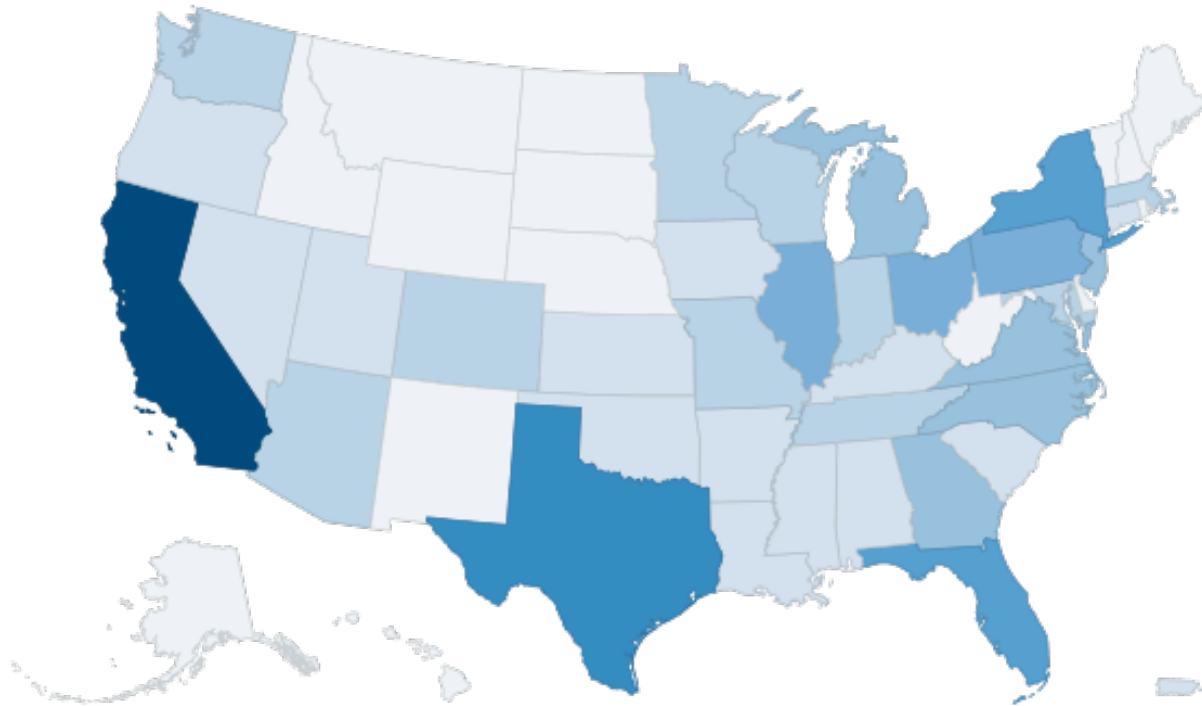
Divide-and-Conquer

- A divide-and-conquer algorithm works by recursively breaking down a problem into two or more sub-problems of the same or related type, until these become simple enough to be solved directly. The solutions to the sub-problems are then combined to give a solution to the original problem.



Divide-and-Conquer

- How to get the total population of USA?



What kind of problem can Divide-and-Conquer solve?

- Problem characteristics:
 - The scale of the problem is large or complex, and the problem can be broken down into several smaller, simple problems of the same type to solve
 - The sub-problems are independent of each other
 - The solutions of the sub-problems can be combined to obtain the solution of original problem



How to Use Divide-and-Conquer

- Decompose the original problem.
 - The original problem is decomposed into several smaller, independent sub-problems that have the same form as the original problem.
- Solve subproblems.
 - If the sub-problems are small and easy to solve, solve them directly;
 - Otherwise, solve the sub-problems recursively.
- Merging solutions
 - Merge the solutions of the sub-problems into solutions of the original problem.



MapReduce

- <https://static.googleusercontent.com/media/research.google.com/en//archive/mapreduce-osdi04.pdf>

MapReduce: Simplified Data Processing on Large Clusters

Jeffrey Dean and Sanjay Ghemawat


jeff@google.com, sanjay@google.com

Google, Inc.

Abstract

MapReduce is a programming model and an associated implementation for processing and generating large data sets. Users specify a *map* function that processes a key/value pair to generate a set of intermediate key/value

given day, etc. Most such computations are conceptually straightforward. However, the input data is usually large and the computations have to be distributed across



Published at OSDI 2004.
>25,000 citations

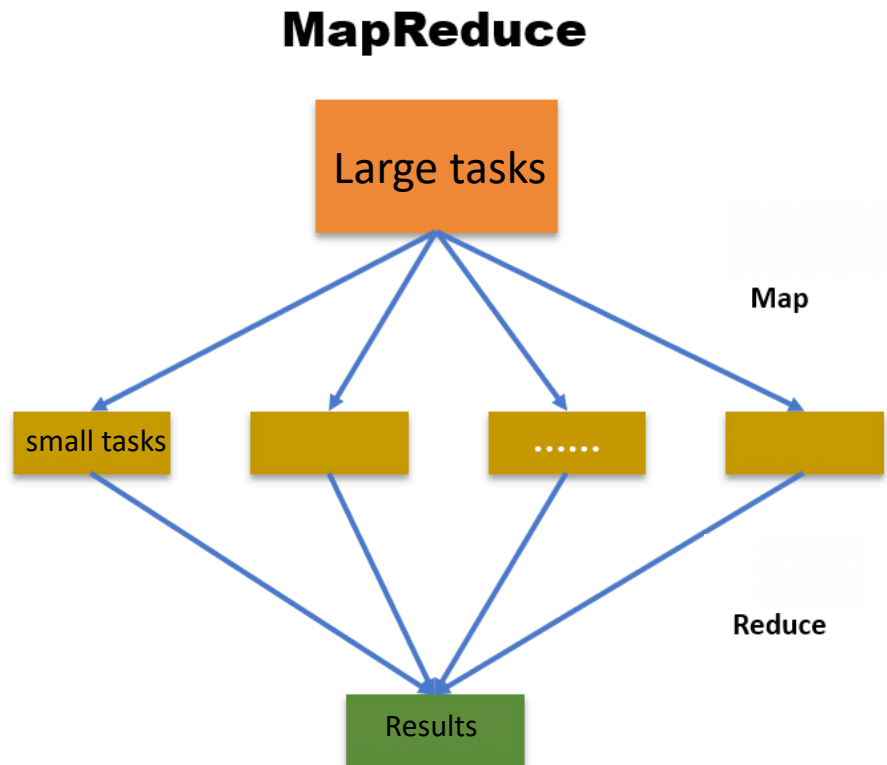
Why MapReduce?

- Automated parallelization and distribution of data and processing
- Clean, powerful, well-understood abstraction (Map and Reduce)
- Fault-tolerance
- Scalability: 1000s of servers, TBs of data, ...
- Apache Hadoop: widely used open source Java implementation




MapReduce

- MapReduce is divided into two core phases: Map and Reduce
- Map corresponds to "divide": complex tasks are broken down into several simple tasks for execution
- Reduce corresponds to "conquer": the results of the Map phase are summarized.



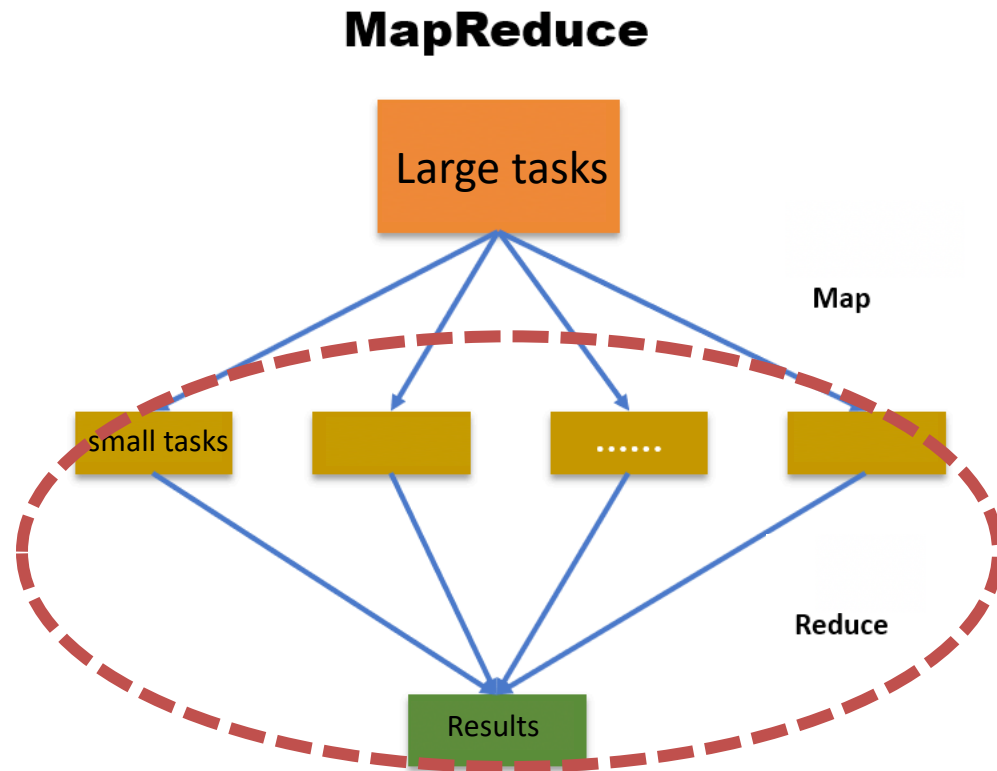
Map

- The small tasks after map have the following characteristics:
 - Compared to the original task, the divided sub-tasks are homogeneous with the original task
 - The sub-task data size and calculation size is much smaller
 - There are no dependencies between different sub-tasks, which can run independently and run in parallel



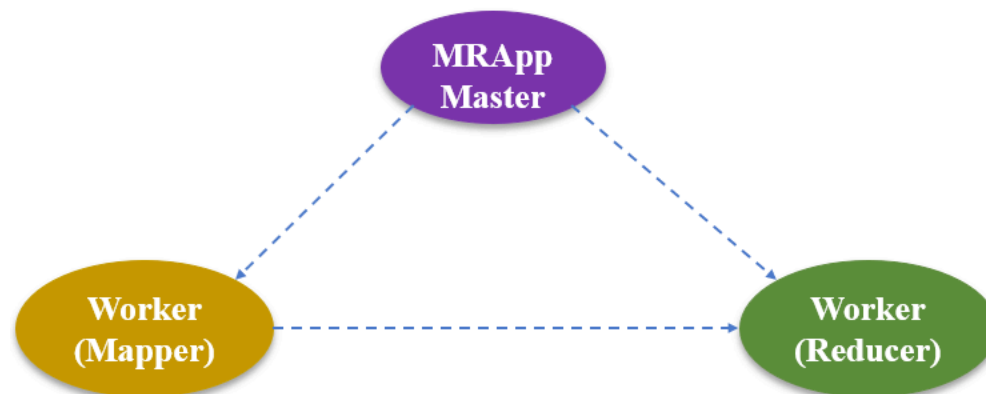
Reduce

- After the division of the sub-tasks in the first stage is completed, the calculation results of all sub-tasks are summarized to obtain the final results.



How the MapReduce works

- MapReduce includes the three components:
 - Master (MRAppMaster): is responsible for assigning tasks, coordinating the operation of tasks, and assigning map function operations to Mappers, and reduce function operations to the Reducers.
 - Mapper worker: is responsible for Map function and performing sub-tasks.
 - Reducer worker: is responsible for Reduce function and summarizing the results of each sub-task

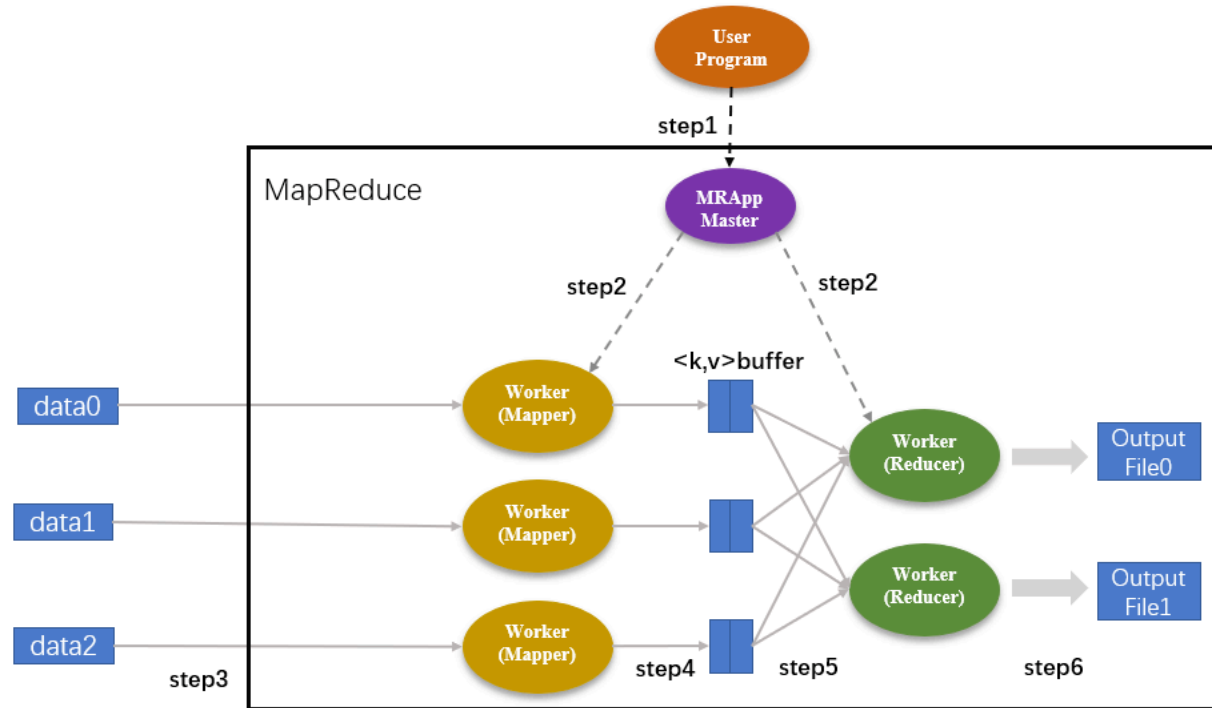


How the MapReduce works

- Workflow (step 1):

User Program sends the task to MRAppMaster. Then, MRAppMaster splits the task into M subtasks (M is a user-defined value).

Example: Assume that the MapReduce function divides the task into 5, of which there are 3 Map jobs and 2 Reduce jobs.




How the MapReduce works

- Workflow (step 2):

MRAppMaster assigns Map and Reduce jobs to Mapper and Reducer, respectively.

The number of Map jobs is the number of divided subtasks, which is 3; Reduce jobs are 2

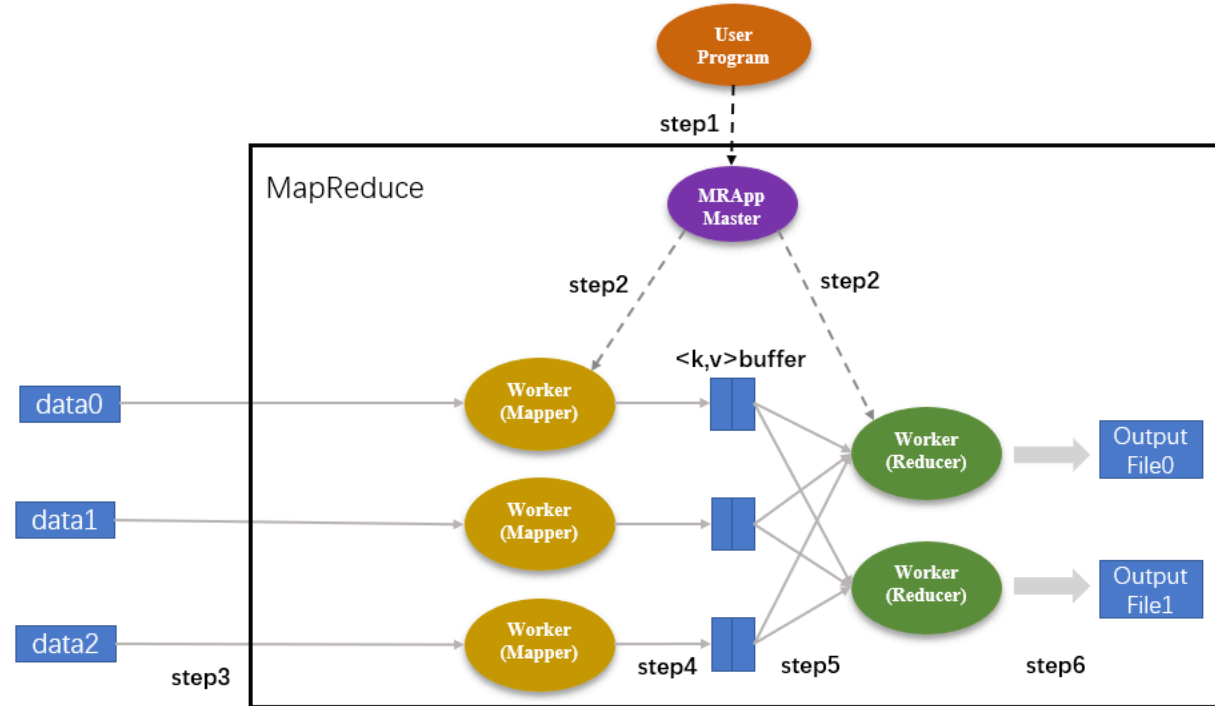


How the MapReduce works

- Workflow (step 3):

The Mapper worker starts reading the input data of the subtask and extracts the $\langle \text{key}, \text{value} \rangle$ pairs from the input data.

Each key-value pair is passed to the `map ()` function as a parameter.

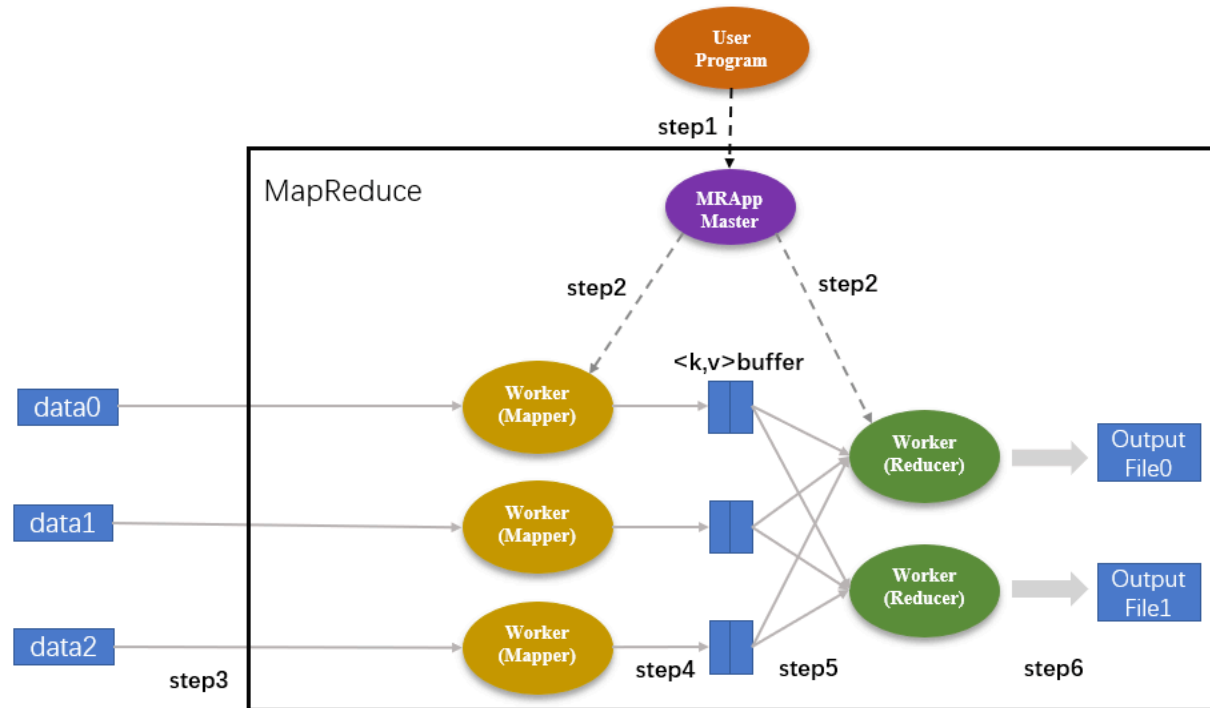


How the MapReduce works

- Workflow (step 4):

The output of the map () function is stored in the ring buffer $\langle k, v \rangle$ Buffer.

These results are periodically written to the local disk and stored in R different disk areas. Here R represents the number of Reduce jobs. In this case, $R = 2$. Also, the storage location of each Map result is reported to the MRAppMaster.

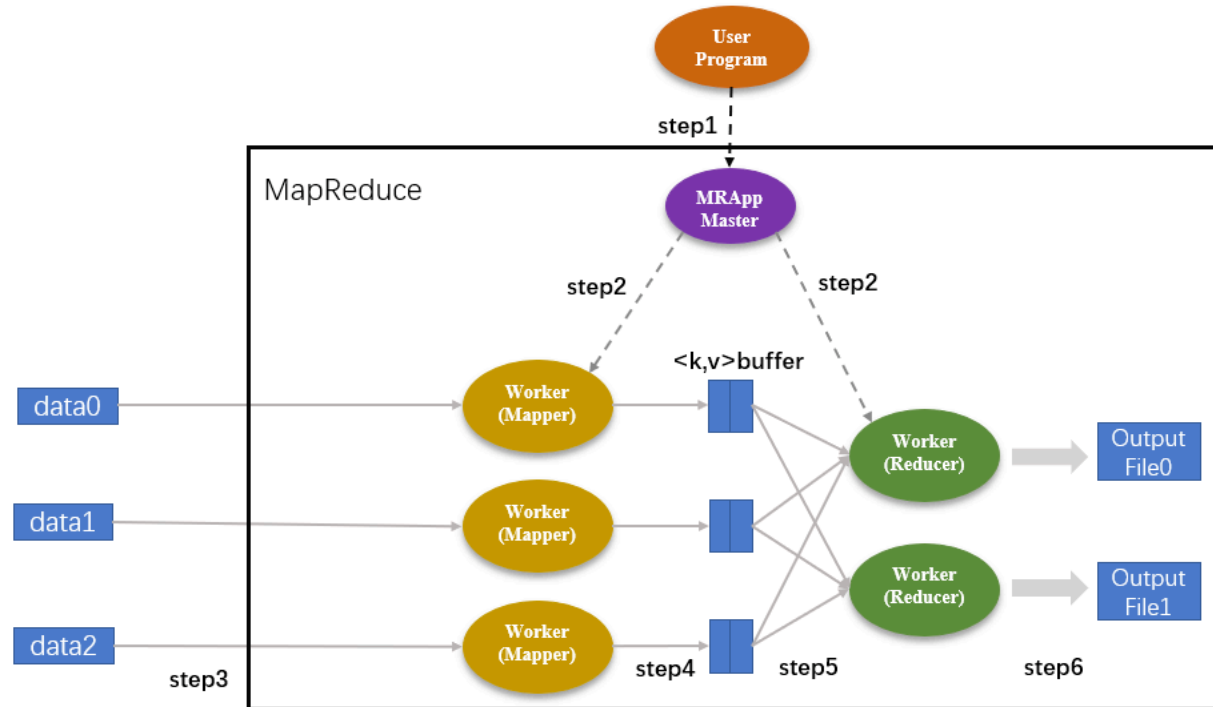


How the MapReduce works

- Workflow (step 5):

The MRAppMaster informs the Reducer which partition it is responsible for, and the Reducer remotely reads the corresponding Map result, which is the intermediate key-value pair.


After the Reducers read all the intermediate key-value pairs it is responsible for, it sorts and aggregates the key-value pairs of the same key value.



How the MapReduce works

- Workflow (step 6):

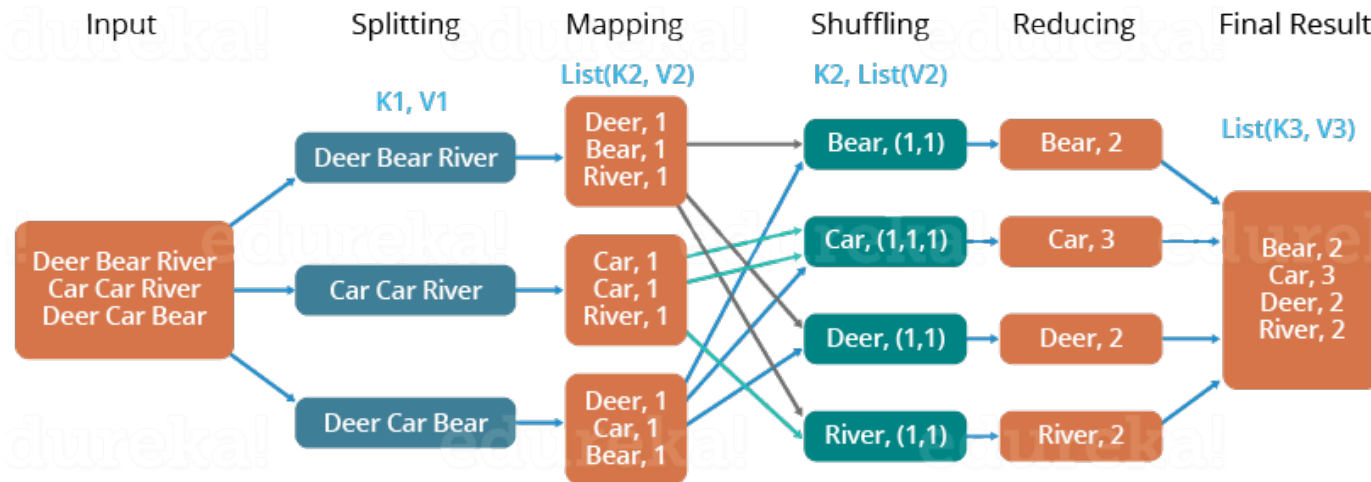
The reducer traverses the sorted key-value pairs, merges the key-value pairs with the same key value, and stores the statistical results as output files in the corresponding partition.



How the MapReduce works

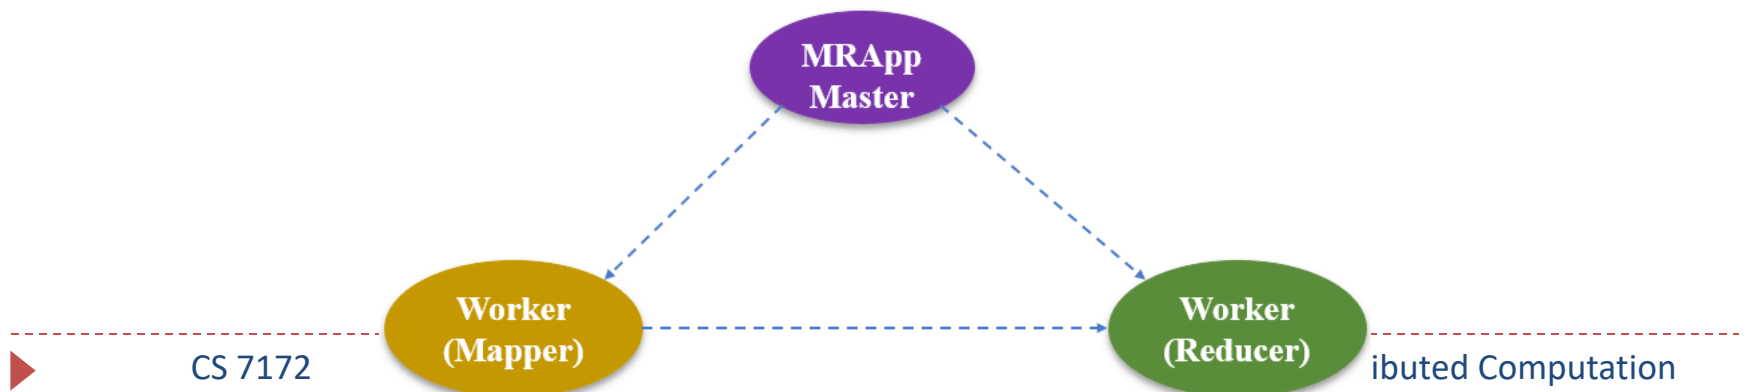
- The entire workflow of MapReduce can be divided into 5 parts:
 - Input, splitting, mapping, reducing and final result

The Overall MapReduce Word Count Process



Fault-Tolerance in MapReduce


- Master fails
 - Switch to secondary master
 - Restart entire job
- Worker fails
 - If running map: restart map tasks on available/free worker nodes
 - If running reduce: restart reducer tasks



MapReduce Example

- We want to know the top 3 selling smartphones in Alabama, Georgia and Florida area

AL	GA	FL
LG	Sumsang	LG
Sumsang	apple	apple
apple	LG	Sumsang
...



MapReduce Example

- Splitting: divide the task into 3 sub-tasks

AL
LG
sumsang
...

GA
sumsang
apple
...

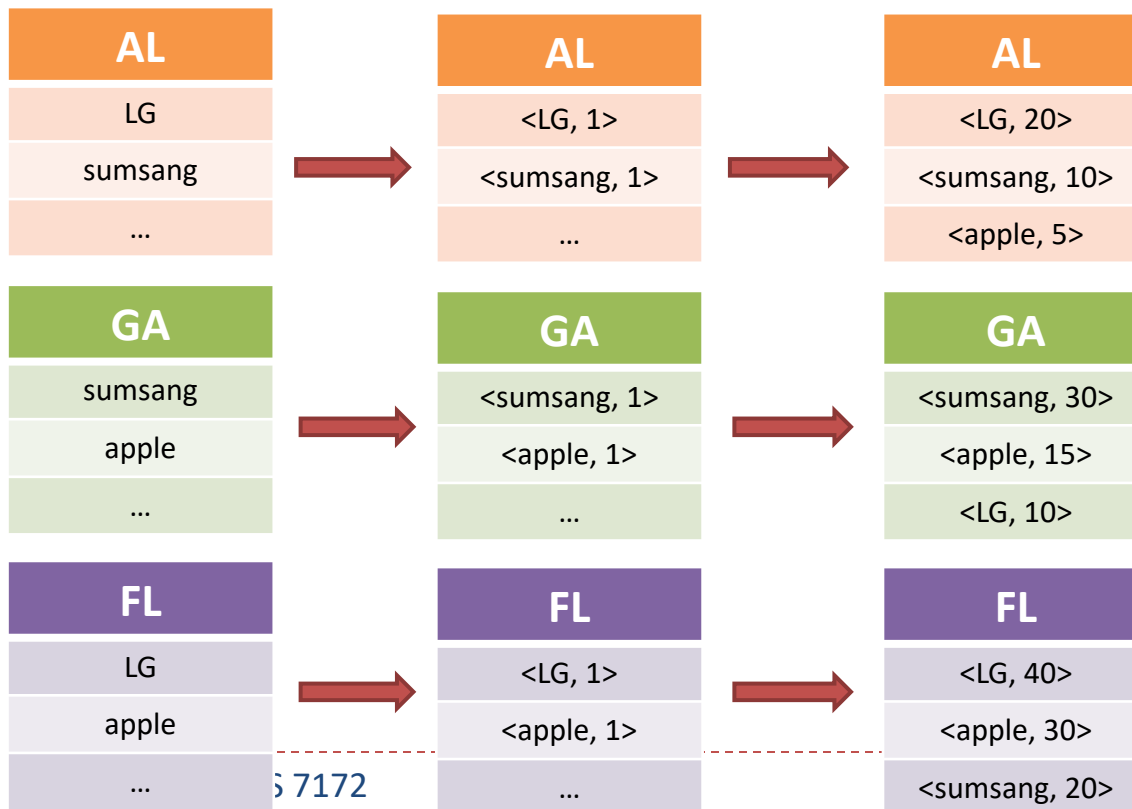
FL
LG
apple
...



AL	GA	FL
LG	Sumsang	LG
Sumsang	apple	apple
apple	LG	Sumsang
...

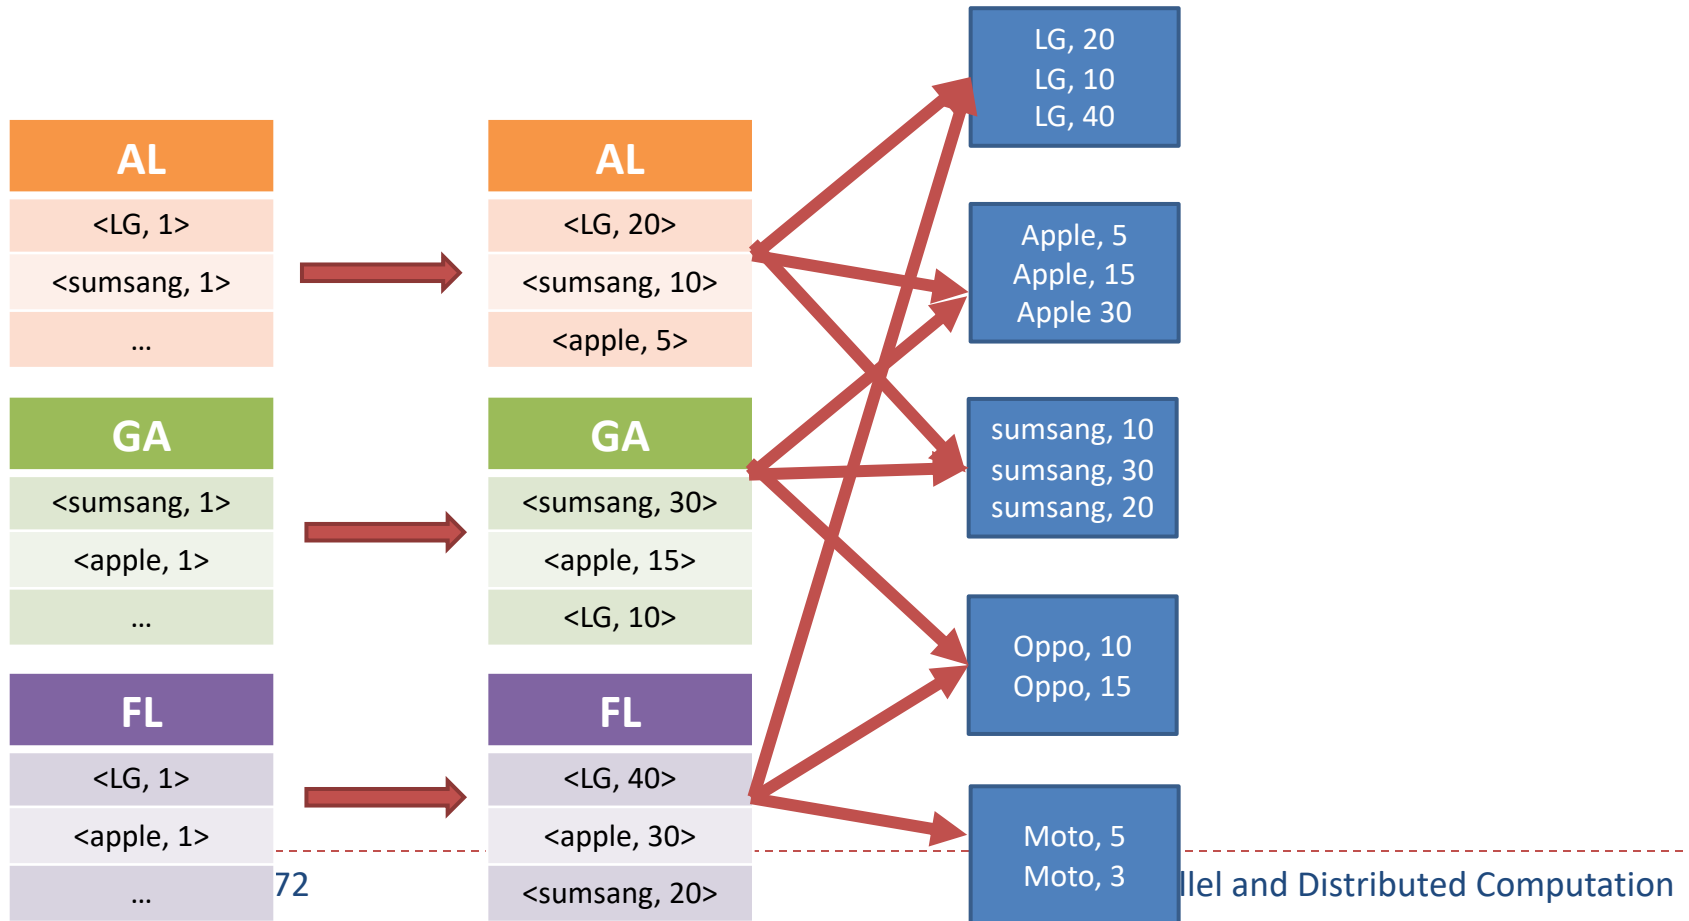
MapReduce Example

- Mapping: keep calling map() function and get the selling statistics of each smartphone brand. Here the key is the smartphone brand and value is the number



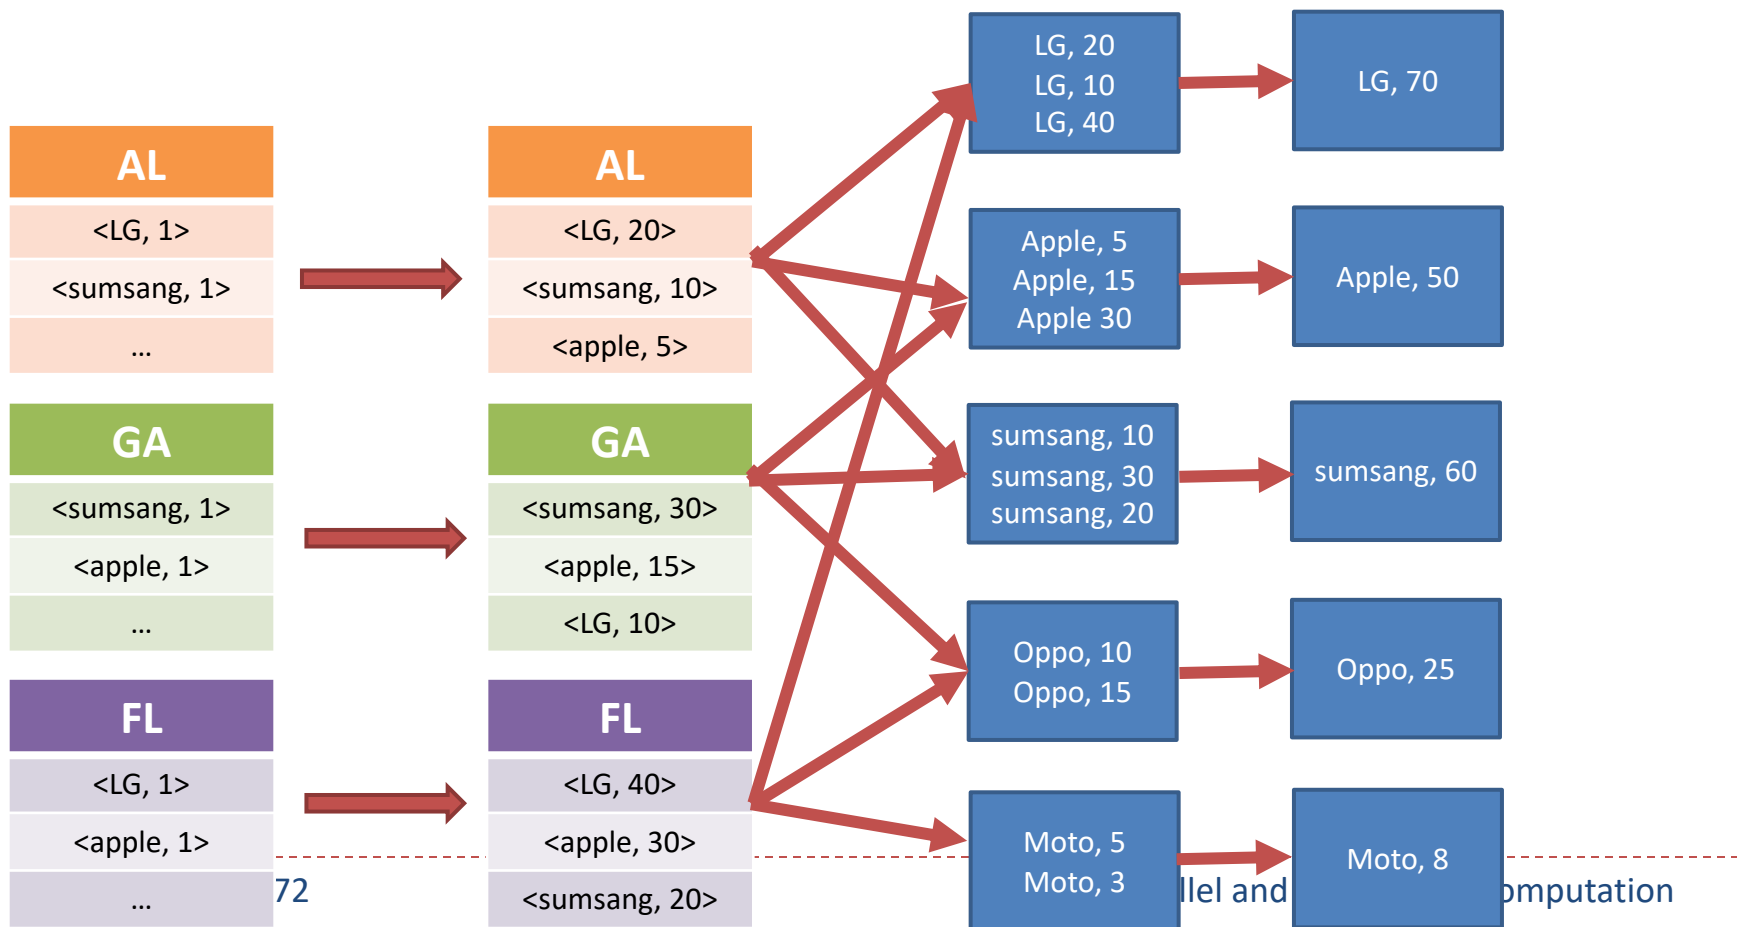
MapReduce Example

- Shuffling: read the results of the Mapping phase and divide the different results into different areas



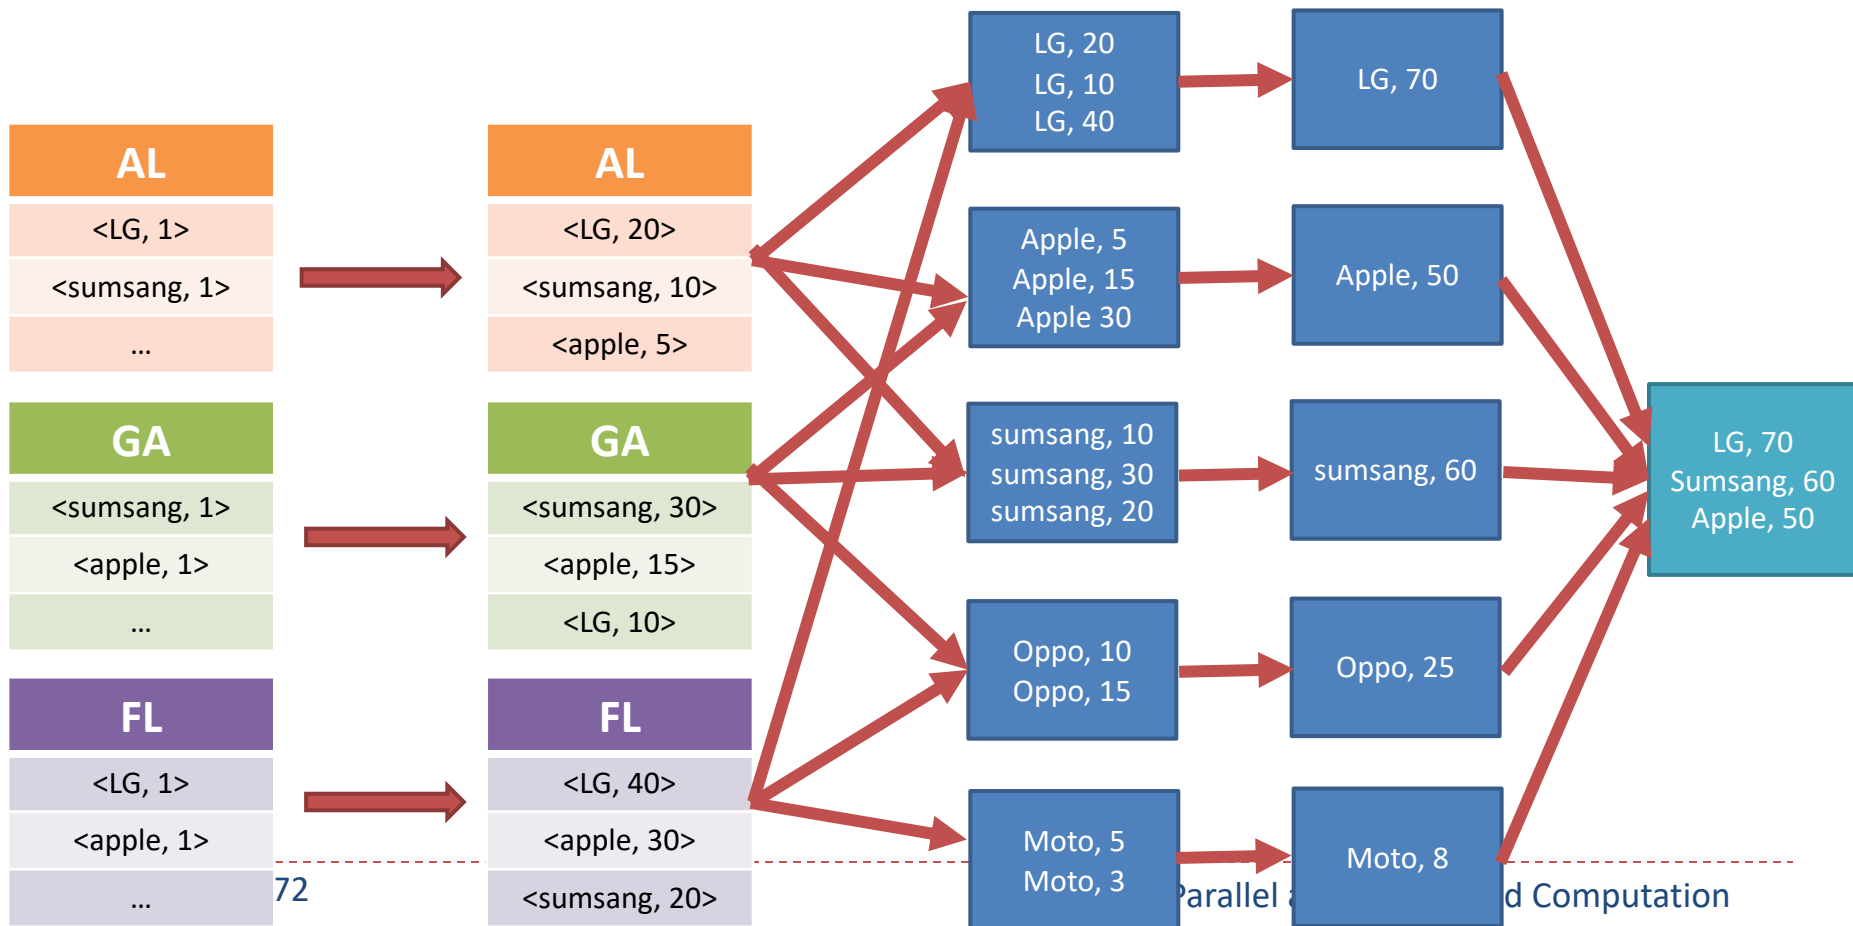
MapReduce Example

- Reducing: merge the results of the same brand




MapReduce Example

- Final results: get the top 3 selling smartphones in Alabama, Georgia and Florida area



MapReduce Limitations

- Static data
 - Restrictive programming model
 - No support for dataflows (not dynamic)
 - Purely batch processing (not flow)
 - Jobs can take hours to complete
 - No streaming, interactive analytics



- How many mappers and reducers?
- How much memory to allocate?

When NOT To Use MapReduce

- Modern computing hardware has plenty of computing power:
 - A typical laptop: 8 cores, 16 GB RAM, 512 GB SSD (usually NVMe)
 - A typical server: 64 cores, 256 GB RAM, 2 TB SSD, ...
- Is your dataset really that large?
 - Wikipedia: 30 GB
- In many cases, an optimized non-distributed implementation beats a large cluster (!)

