

CS 3502

Operating Systems

Scheduling

Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

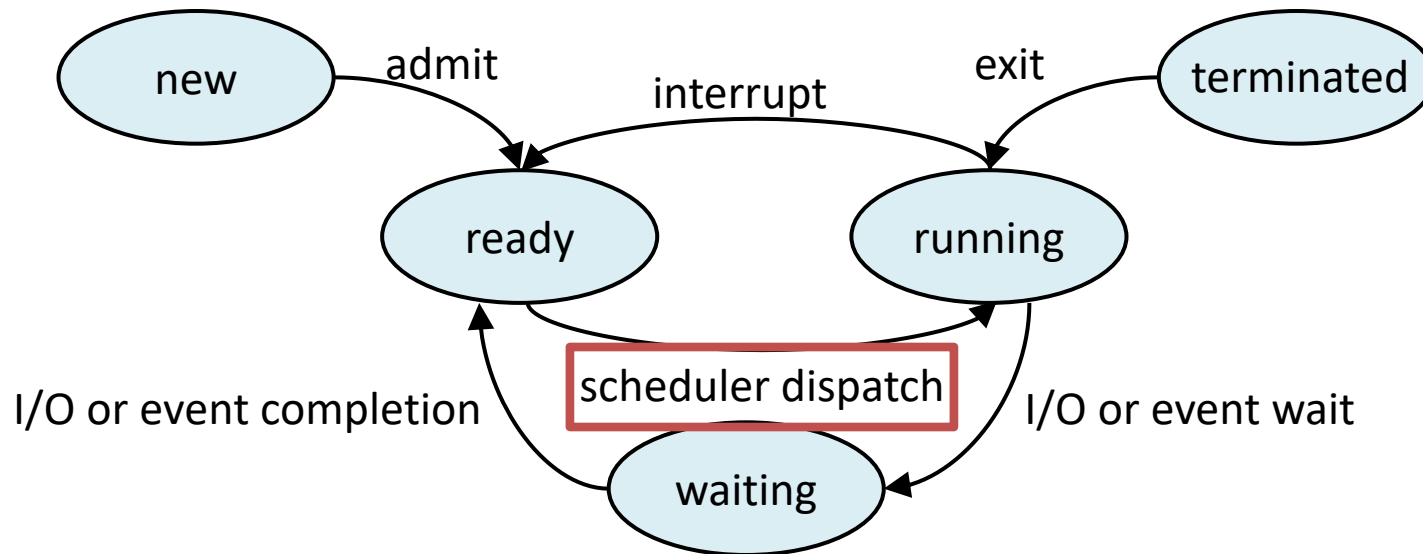
Outline

- Introduction to CPU scheduling
 - What is CPU scheduling
 - Why we need CPU scheduling
 - When scheduling happens
- Scheduling policies
 - FCFS, SJF, RR, Priority
 - Scheduling on multiple CPUs



What is CPU scheduling?

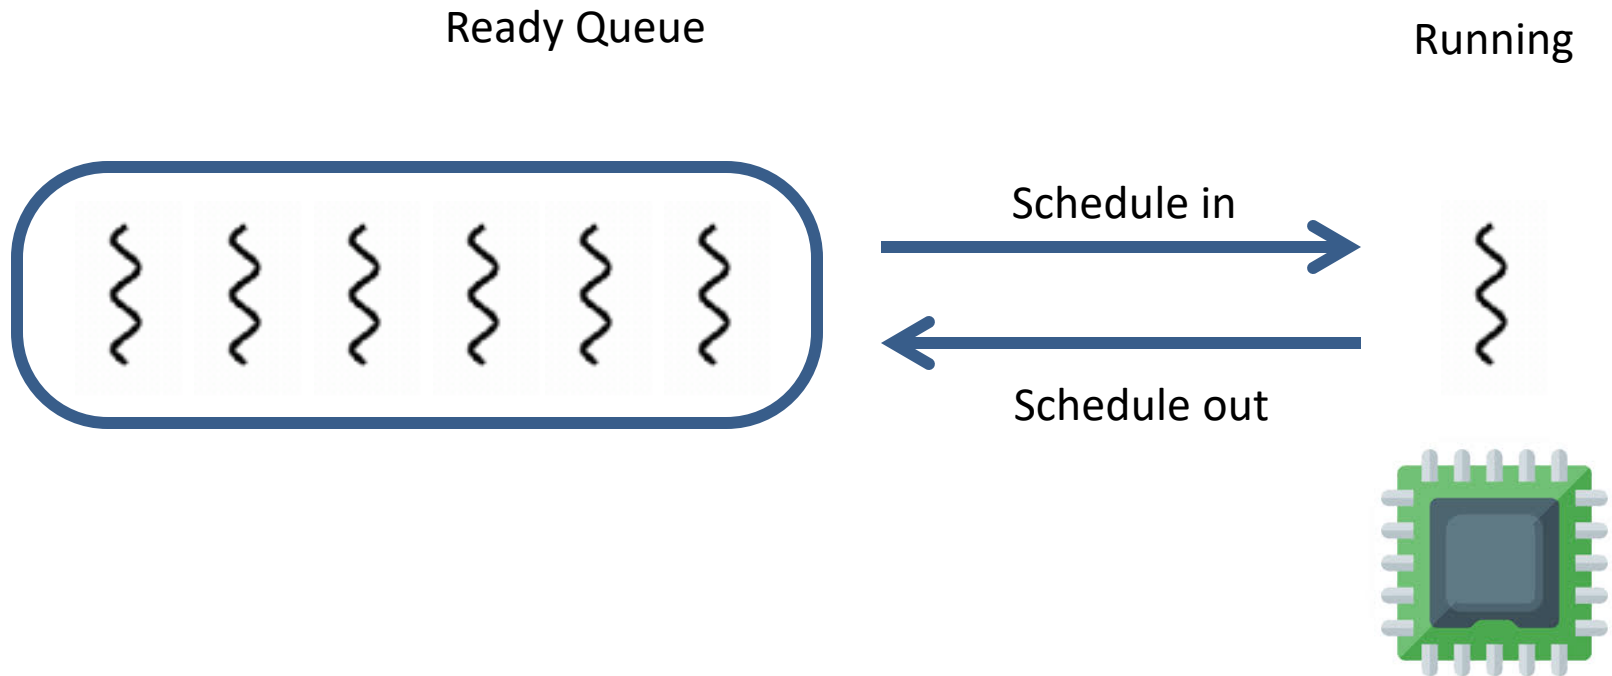
- The five-state process model



CPU scheduling


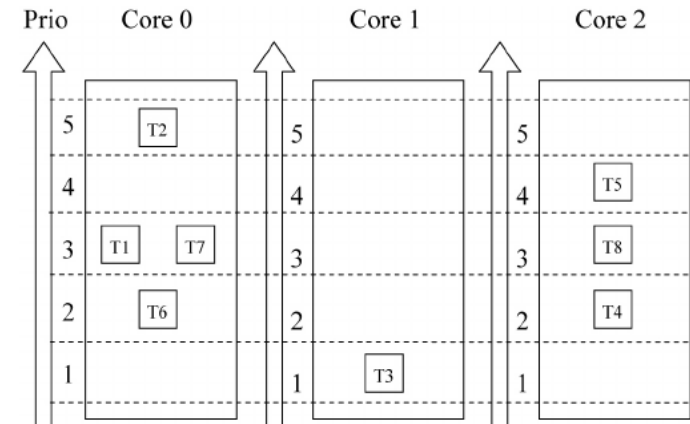
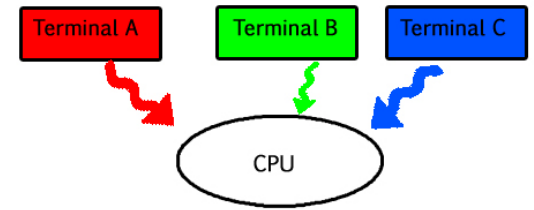
Selects from among the processes/threads that are ready to execute, and allocates the CPU to it

What is CPU scheduling?



Why CPU scheduling?

- In support of **multiprogramming**
 - uniprocessor systems
 - ▶ Time-sharing processor
 - multiprocessor systems
 - ▶ Efficiently distributing tasks
 - Real-time systems
 - ▶ Reliably guaranteeing deadlines




Why CPU scheduling? – Different Goals

- All systems
 - Fairness - giving each process a fair share of the CPU
 - Policy enforcement - seeing that stated policy is carried out
 - Balance - keeping all parts of the system busy



Why CPU scheduling? – Different Goals

- Batch systems
 - Throughput - maximize jobs per hour
 - Turnaround time - minimize time between submission and termination
 - CPU utilization - keep the CPU busy all the time




National lab:
Scientific research
Energy, nuclear, climate, ...



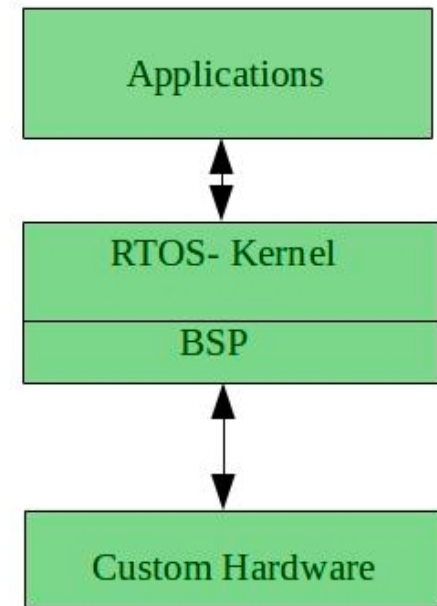
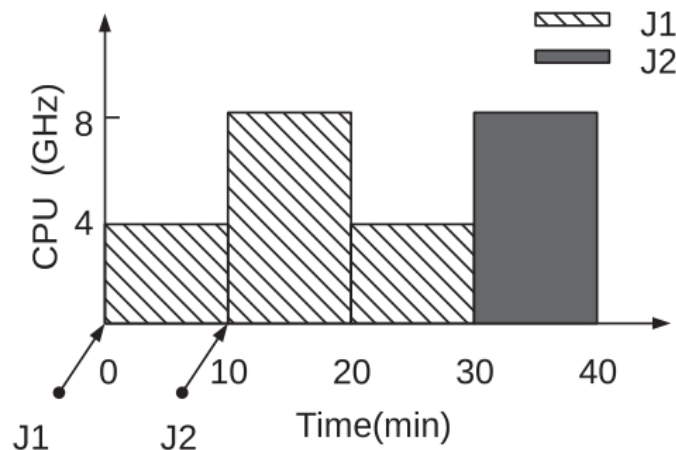
Why CPU scheduling? – Different Goals

- Interactive systems
 - Response time - respond to requests quickly
 - Proportionality - meet users' expectations



Why CPU scheduling? – Different Goals

- Real-time systems
 - Meeting deadlines - avoid losing data
 - Predictability - avoid quality degradation in multimedia systems



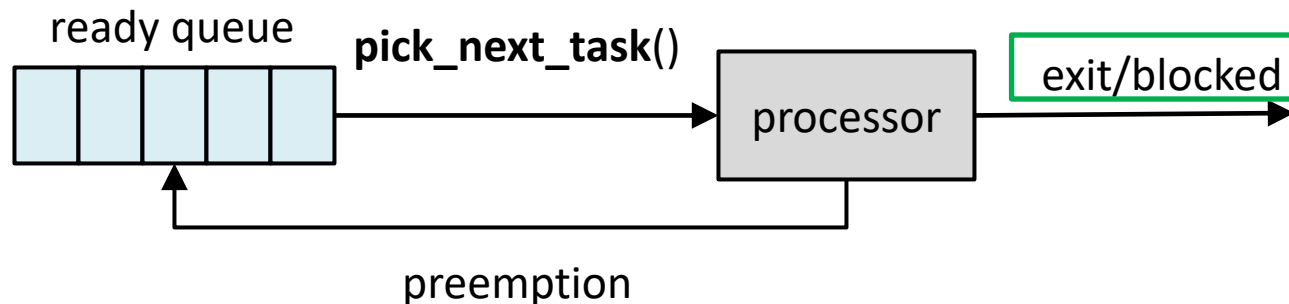
Why CPU scheduling? – Different Goals

- All systems
 - Fairness - giving each process a fair share of the CPU
 - Policy enforcement - seeing that stated policy is carried out
 - Balance - keeping all parts of the system busy
- Batch systems
 - Throughput - maximize jobs per hour
 - Turnaround time - minimize time between submission and termination
 - CPU utilization - keep the CPU busy all the time
- Interactive systems
 - Response time - respond to requests quickly
 - Proportionality - meet users' expectations
- Real-time systems
 - Meeting deadlines - avoid losing data
 - Predictability - avoid quality degradation in multimedia systems



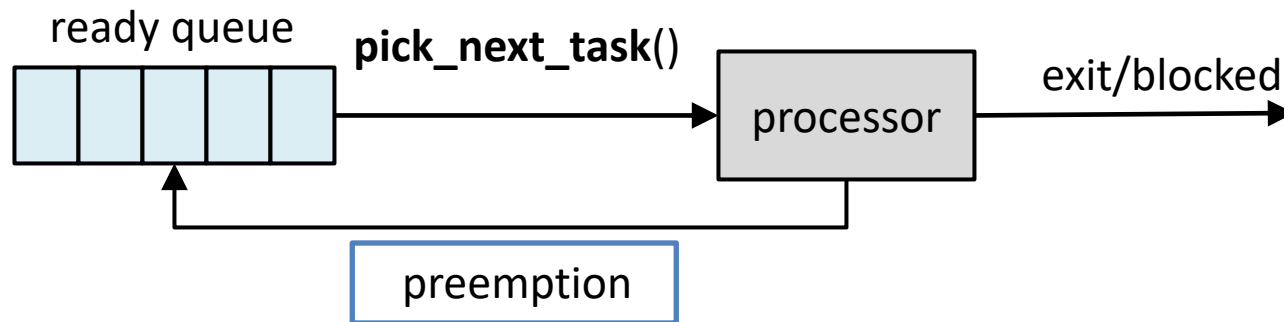
When scheduling happens?

- Non-preemptive
 - Scheduling only when current process **terminates** or **gives up** control



When scheduling happens?

- Preemptive
 - Processes can be **forced** to give up control

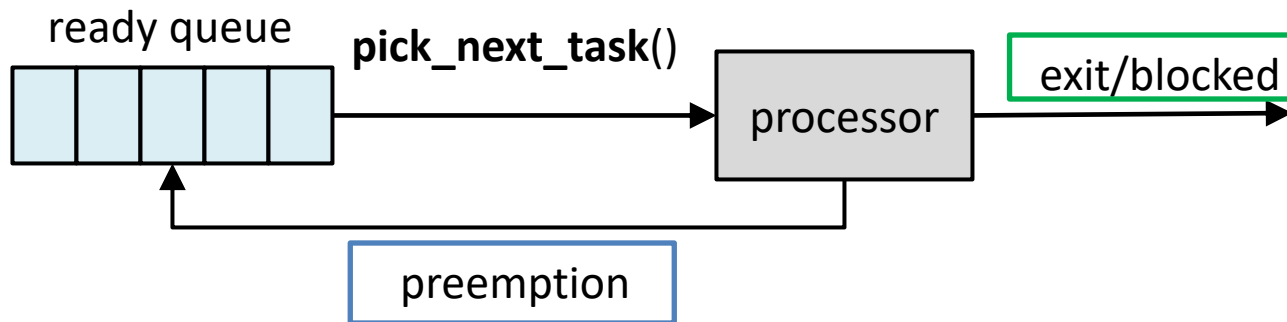


When scheduling happens?

- CPU scheduling may take place at
 - Clock interrupts
 - I/O interrupts
 - I/O completion
 - Termination

} preemptive

} non-preemptive



Outline

- Introduction to CPU scheduling
 - What is CPU scheduling
 - Why we need CPU scheduling
 - When scheduling happens
- Scheduling policies
 - FCFS, SJF, RR, Priority
 - Scheduling on multiple CPUs



Scheduling Policies


Not exist best scheduling.
It depends on your goals.

- Batch Systems
 - First-Come First-Serve (FCFS)
 - Shortest Job First
 - Shortest Remaining Time Next
- Interactive Systems
 - Round-Robin
 - Priority Scheduling
 - Multiple Queues
 - Shortest Process Next
 - Guaranteed Scheduling
 - Lottery Scheduling
- Real-time Systems
 - Rate Monotonic Scheduling
 - Earliest Deadline First Scheduling

Determine the next
ready task to run

Turnaround time = End time – Arrival time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



Response time = Start time – Arrival time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



0

1 \longleftrightarrow 8

2 \longleftrightarrow 12

3 \longleftrightarrow 21



First-Come, First-Serve (FCFS)

Turnaround time = End time – Arrival time
Response time = Start time – Arrival time

- CPU schedules the task that arrived earliest, non-preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



Average turnaround time = ?



First-Come, First-Serve (FCFS)

- CPU schedules the task that arrived earliest, non-preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



Average turnaround time = $((8-0)+(12-1)+(21-2)+(26-3)) / 4 = 15.25$



First-Come, First-Serve (FCFS)

Turnaround time = End time – Arrival time
Response time = Start time – Arrival time

- CPU schedules the task that arrived earliest, non-preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



Average response time = ?



First-Come, First-Serve (FCFS)

- CPU schedules the task that arrived earliest, non-preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



$$\text{Average response time} = (0 + (8 - 1) + (12 - 2) + (21 - 3)) / 4 = 8.75$$




Shortest Job First (SJF)

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive




Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool




Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool




Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool




Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool




Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool



Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool



Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool



Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool



Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



pool



Shortest Job First (SJF)

Turnaround time = End time – Arrival time

Response time = Start time – Arrival time

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



Average turnaround time = ?



Shortest Job First (SJF)

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



$$\text{Average turnaround time} = ((8-0)+(12-1)+(26-2)+(17-3)) / 4 = 14.25$$



Shortest Job First (SJF)

Turnaround time = End time – Arrival time

Response time = Start time – Arrival time

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



Average response time = ?



Shortest Job First (SJF)

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

non-preemptive



$$\text{Average response time} = (0 + (8 - 1) + (17 - 2) + (12 - 3)) / 4 = 7.75$$




Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool



Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool



Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool


$P_1(7)$

Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



P2 remain =3

pool




Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



P2 remain =2

pool



Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool



Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool



Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool




Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool




Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

preemptive



pool



Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

Turnaround time = End time – Arrival time
Response time = Start time – Arrival time

preemptive



Average turnaround time = ?



Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

Turnaround time = End time – Arrival time

Response time = Start time – Arrival time

preemptive



$$\text{Average turnaround time} = ((17-0)+(5-1)+(10-3)+(26-2)) / 4 = 13$$



Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

Turnaround time = End time – Arrival time
Response time = Start time – Arrival time

preemptive



Average response time = ?



Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

Turnaround time = End time – Arrival time

Response time = Start time – Arrival time

preemptive



$$\text{Average response time} = (0 + (1 - 1) + (5 - 3) + (17 - 2)) / 4 = 4.25$$


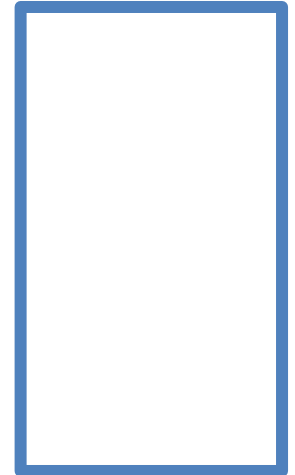


Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool




Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



P_1 is put at the end of queue after scheduled out





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool

$P_1 (4)$

$P_3 (5)$





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool




Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool

$P_3 (1)$





Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool




Round Robin (RR)

Turnaround time = End time – Arrival time

Response time = Start time – Arrival time

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



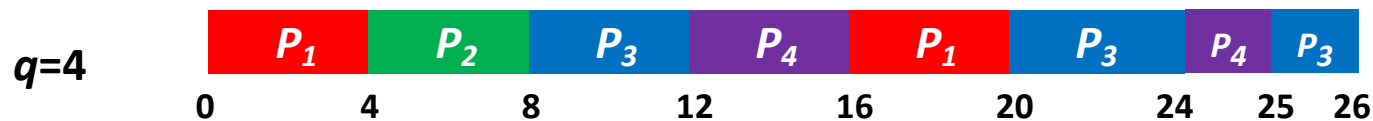
Average turnaround time = ?



Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5




$$\text{Average turnaround time} = ((20-0)+(8-1)+(26-2)+(25-3)) / 4 = 18.25$$

Round Robin (RR)

Turnaround time = End time – Arrival time
Response time = Start time – Arrival time

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5




Average response time = ?



Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



$$\text{Average response time} = (0 + (4 - 1) + (8 - 2) + (12 - 3)) / 4 = 4.5$$




Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$



Round Robin (RR)

Select the task at the beginning

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5


pool

$P_2 (4)$

$P_3 (9)$

$P_4 (5)$

$q=5$



Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool

$P_3 (9)$

$P_4 (5)$

$P_1 (3)$

$q=5$




P1 is put at the end of queue after scheduled out

Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$




Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$




Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$



Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool

$P_1 (3)$

$P_3 (4)$

$q=5$



Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool

$P_1 (3)$

$P_3 (4)$

$q=5$




Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$




Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$




Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

pool



$q=5$



Round Robin (RR)


Turnaround time = End time – Arrival time

Response time = Start time – Arrival time

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

$q=5$



Average turnaround time = ?

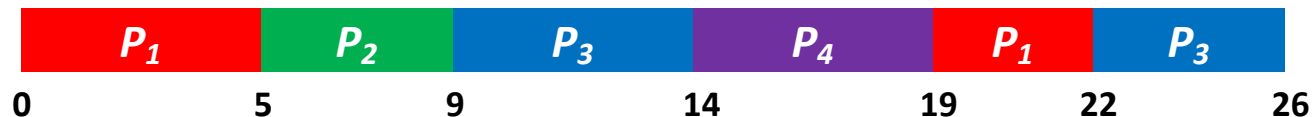


Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

$q=5$



$$\text{Average turnaround time} = ((22-0)+(9-1)+(26-2)+(19-2)) / 4 = 17.5$$



Round Robin (RR)


Turnaround time = End time – Arrival time

Response time = Start time – Arrival time

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

$q=5$



Average response time = ?




Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

$q=5$




$$\text{Average response time} = (0 + (5 - 1) + (9 - 2) + (14 - 3)) / 4 = 5.5$$




Priority Scheduling

- CPU schedules the highest priority (smaller value) first, FCFS within the same priority



Priority Scheduling

- CPU schedules the highest priority (smaller value) first, FCFS within the same priority



Comparison


	Turnaround time	Response time
FCFS	15.25	8.75
SJF-preemptive	13	4.25
RR (q=5)	17.5	5.5
Priority scheduling	N/A	N/A

	Throughput	Response time	Starvation
FCFS	TBD	TBD	No
SJF-preemptive	High	Good	Yes
RR	Can be low	Good	No
Priority scheduling	Can be high	Can be good	Can remove



Challenges on Emerging Hardware and Applications

- Multi-processor → Single queue



Multiprocessor = more powerful processor

pick_next_task() will be the bottleneck

Linux 2.4.x

Pros:

1. Easy to implement
2. Perfect load balancing

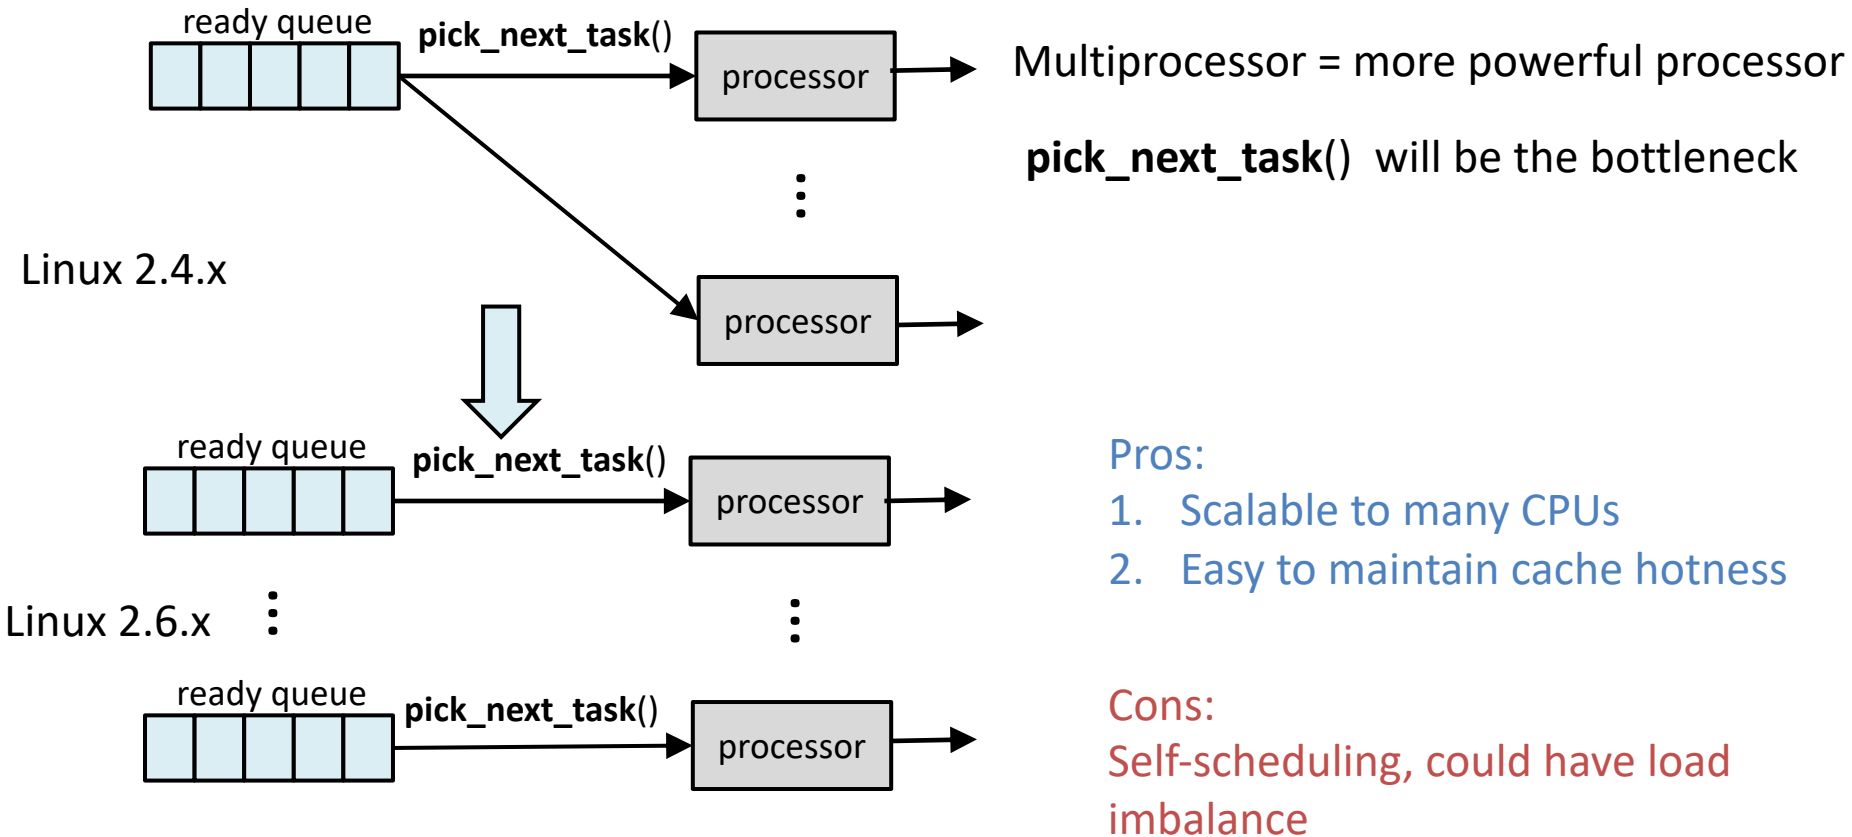
Cons:

1. Scalability issues due to centralized synchronization
2. High overhead and low efficiency
3. Hard to maintain cache hotness due to global scheduling



Challenges on Emerging Hardware and Applications

- Multi-processor → Many queues




Overcome Load Imbalance

- Push model

Every a while, a kernel thread checks load imbalance and move threads


Made by OS




- Pull model

Whenever a queue becomes empty, steal a thread from non-empty queues


Made by local queue. Both are widely used



Load balance on SMP vs. NUMA



symmetric multiprocessing (SMP):
The distance to memory is the same



Non-uniform memory access (NUMA):
The distance to memory is different

Multi processor/core scheduling

The scheduling policy not only considers the fairness, throughput, etc., but also needs to consider the **hardware architecture** (e.g., locality)

Conclusion

- Introduction to CPU scheduling
 - What is CPU scheduling
 - Why we need CPU scheduling
 - When scheduling happens
- Scheduling policies
 - FCFS, SJF, RR, Priority
 - Scheduling on multiple CPUs

