

CS 3502

Operating Systems

Page Design and Segmentation

Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

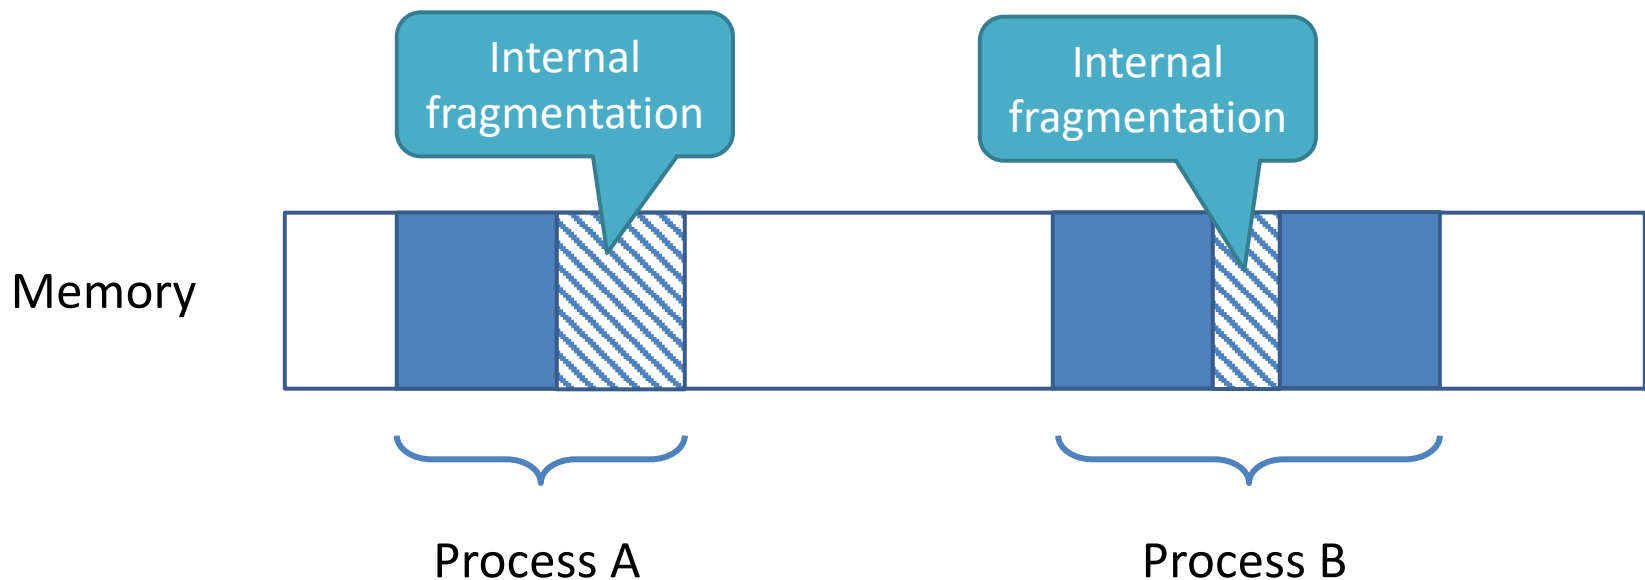
Outline

- Page design
 - Internal fragmentation vs. External fragmentation
 - Local page replacement vs. Global page replacement
 - Page size small vs. Large
 - Shared page
 - Paging with Process life cycle
- Segmentation
 - Page vs. Segmentation




Internal fragmentation

- **Internal fragmentation**: when memory allocated to a process is larger than requested memory, the difference between these two numbers is internal fragmentation.



How internal fragmentation is generated?

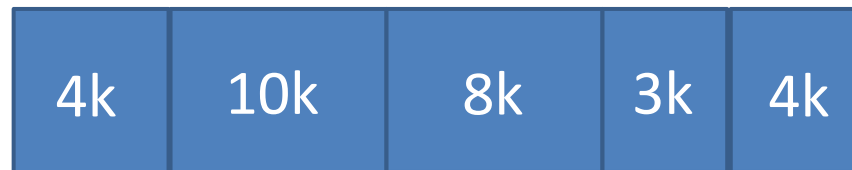
- Many internal fragmentation is caused by **fixed-sized blocks** of memory
- Whenever a process requests for the memory, the fixed sized block is allocated to the process



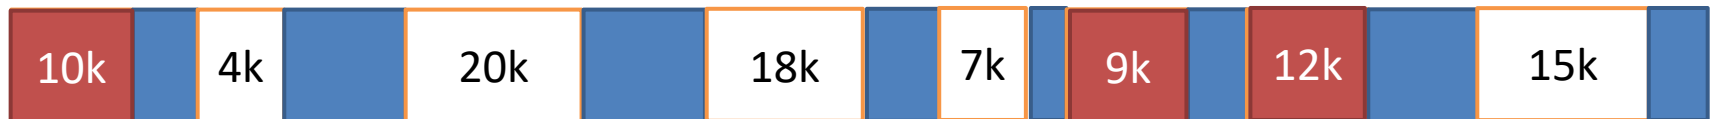
How to deal with the internal fragmentations?

- Allocate dynamic size blocks of memory based on process requirement

Virtual
memory

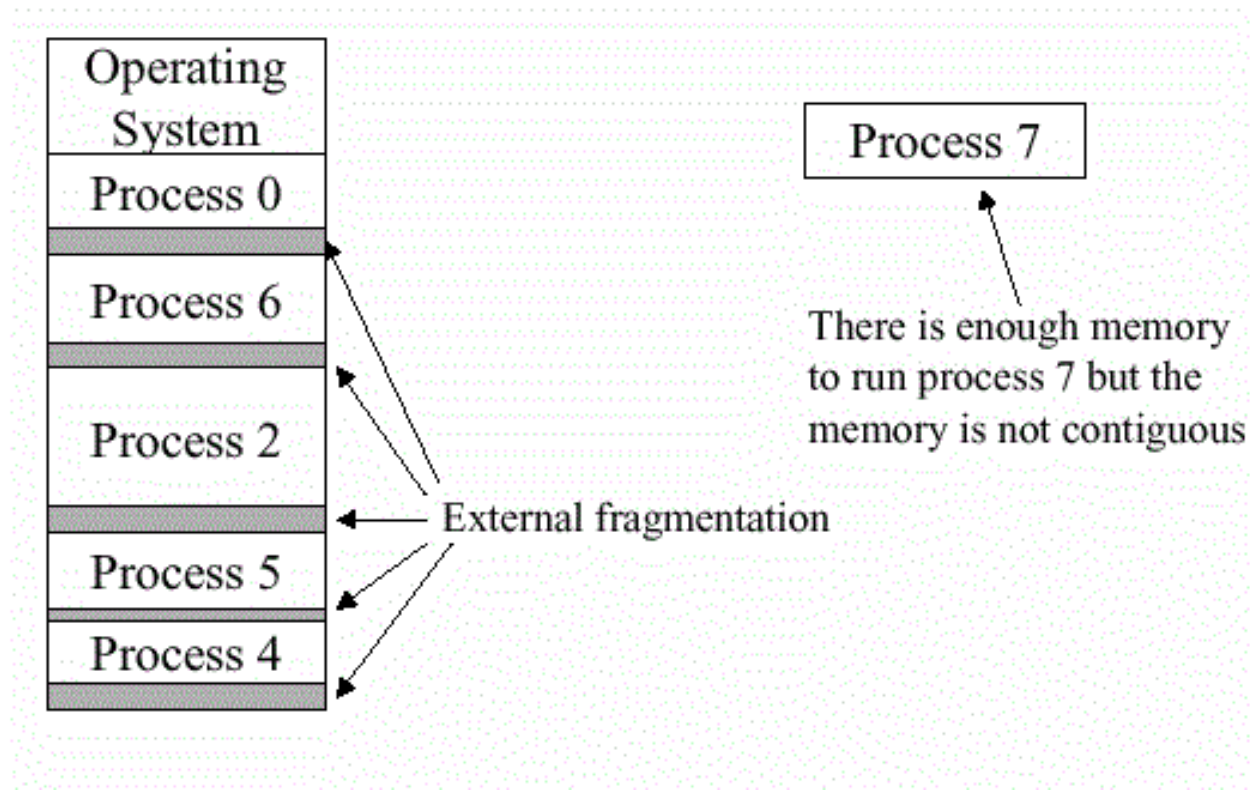


- Best-fit allocation algorithm



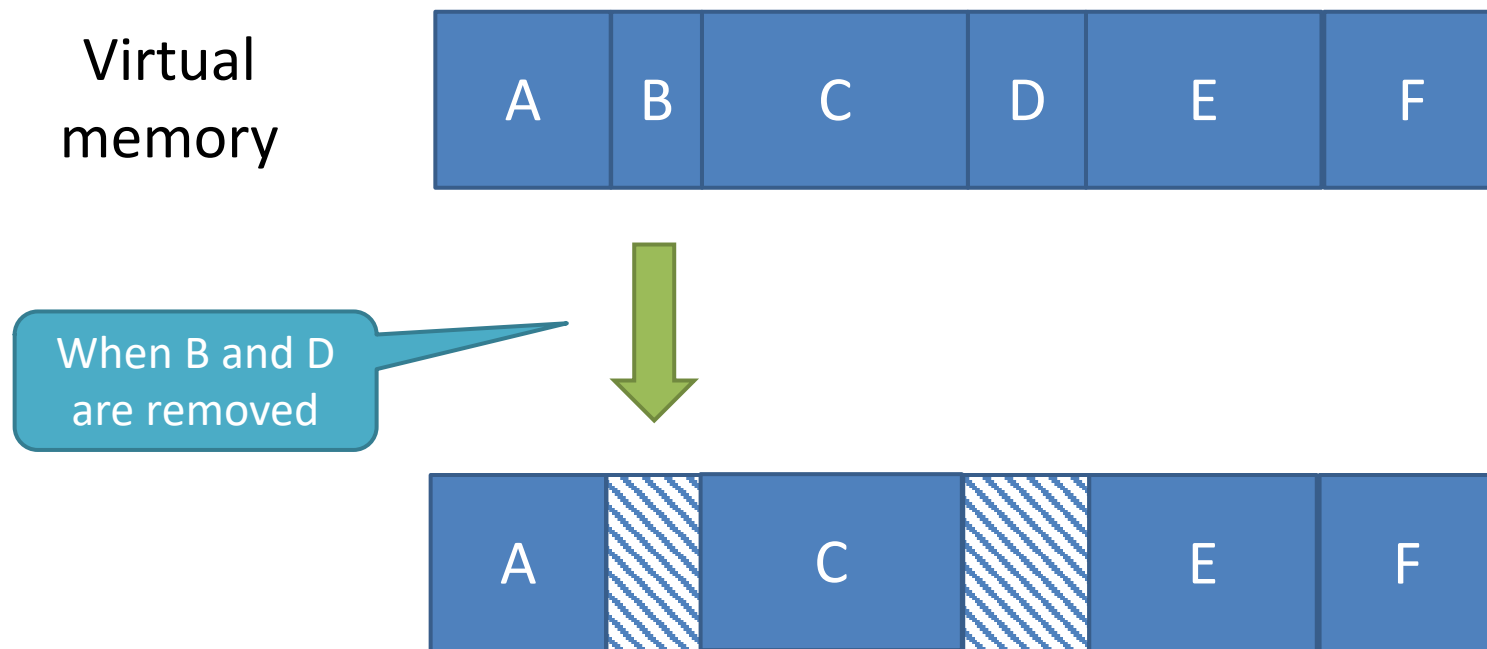
External fragmentation

- **External fragmentation**: Total memory space is enough to satisfy a request or to reside a process in it. However, it is not contiguous and can not be used.



How external fragmentation is generated?

- When a process is removed from the memory, the free space creates the hole in the memory



How to deal with the external fragmentations?


- Best-fit allocation algorithm



Memory compaction in MacOS

- Memory compaction

https://youtu.be/hligp_bxUcQ?t=1763

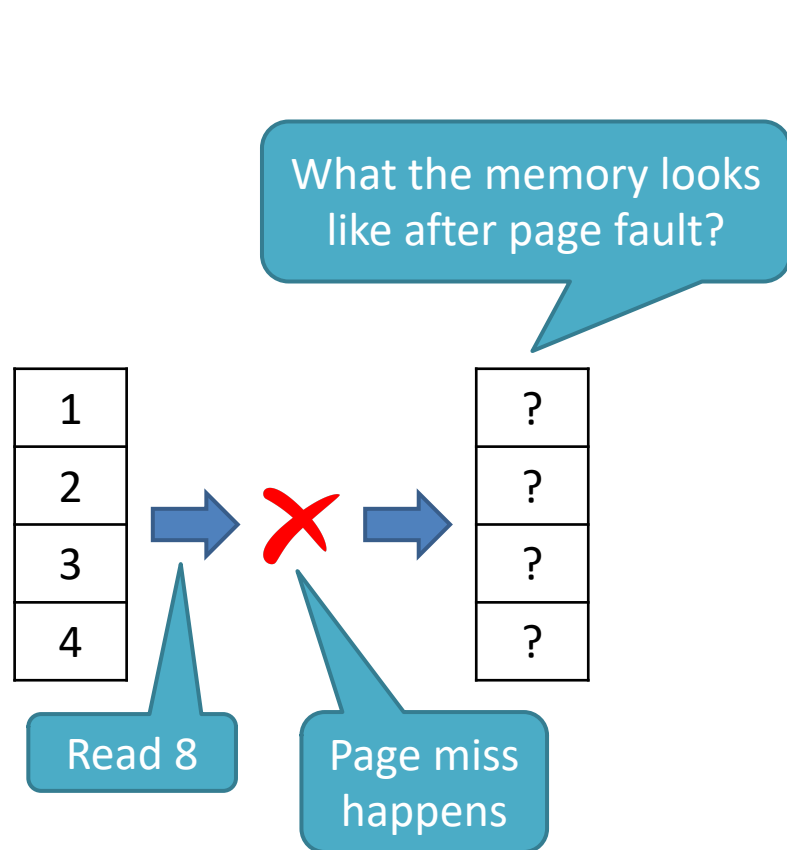


Internal fragmentation vs. External fragmentation

	Internal fragmentation	External fragmentation
Definition	A form of fragmentation that arises when there are sections of memory remaining because of allocating large blocks of memory for a process than required	A form of fragmentation that arises when there is enough memory available to allocate for the process, but that available memory is not contiguous
Reason	Memory block assigned to a process is large – the remaining portion is left unused as it cannot be assigned to another process	Memory space is enough to reside a process, but it is not contiguous. Therefore, that space cannot be used for allocation
Solution	Best fit Dynamic block size	Best fit Memory Compaction



Page replacement



- Page replacement algorithm
 - OPR
 - FIFO
 - LRU
 - NFU
 - NRU
 - Second chance
 - Clock
 - Aging

Local Page replacement

		Age
A	A0	10
	A1	7
	A2	5
	A3	4
	A4	6
	A5	3
B	B0	9
	B1	4
	B2	6
	B3	2
	B4	5
	B5	6
	B6	12
C	C1	3
	C2	5
	C3	6

(a)

(a) Original configuration.

A0
A1
A2
A3
A4
A6
B0
B1
B2
B3
B4
B5
B6
C1
C2
C3

(b)

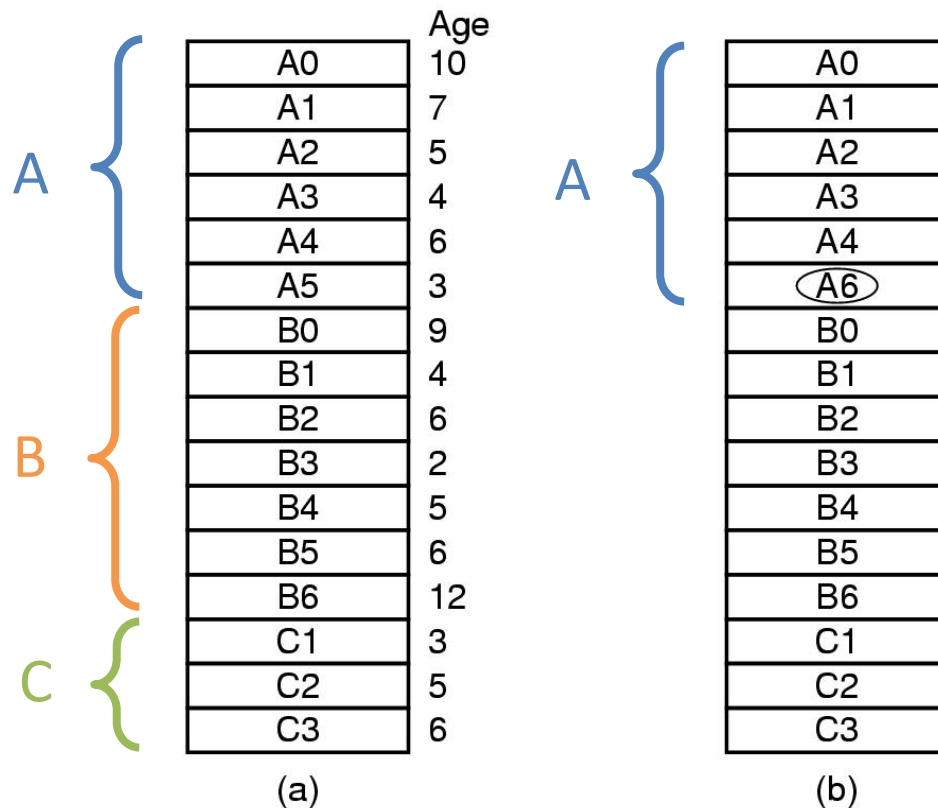
(b) Local page replacement.

A5 will be replaced
by local algorithm

Suppose the algorithm
replaces the page which
has *the least age*.

e.g., A6 comes

Local Page replacement



- Local page replacement requires **static allocation**
- The **page number** of one process does **not change** during replacement

Local Page replacement problem

- Thrashing
 - Physical memory is too small to hold the process work set
 - Large page faults happen and swap frequently
 - Slowdown the process speed

		Age
A	A0	10
	A1	7
	A2	5
	A3	4
	A4	6
	A5	3
B	B0	9
	B1	4
	B2	6
	B3	2
	B4	5
	B5	6
	B6	12
C	C1	3
	C2	5
	C3	6


Suppose A needs **9** pages during execution:

A0 A1 A2 A3 A4 A5 **A6 A7 A8 ...**



Thrashing
happens

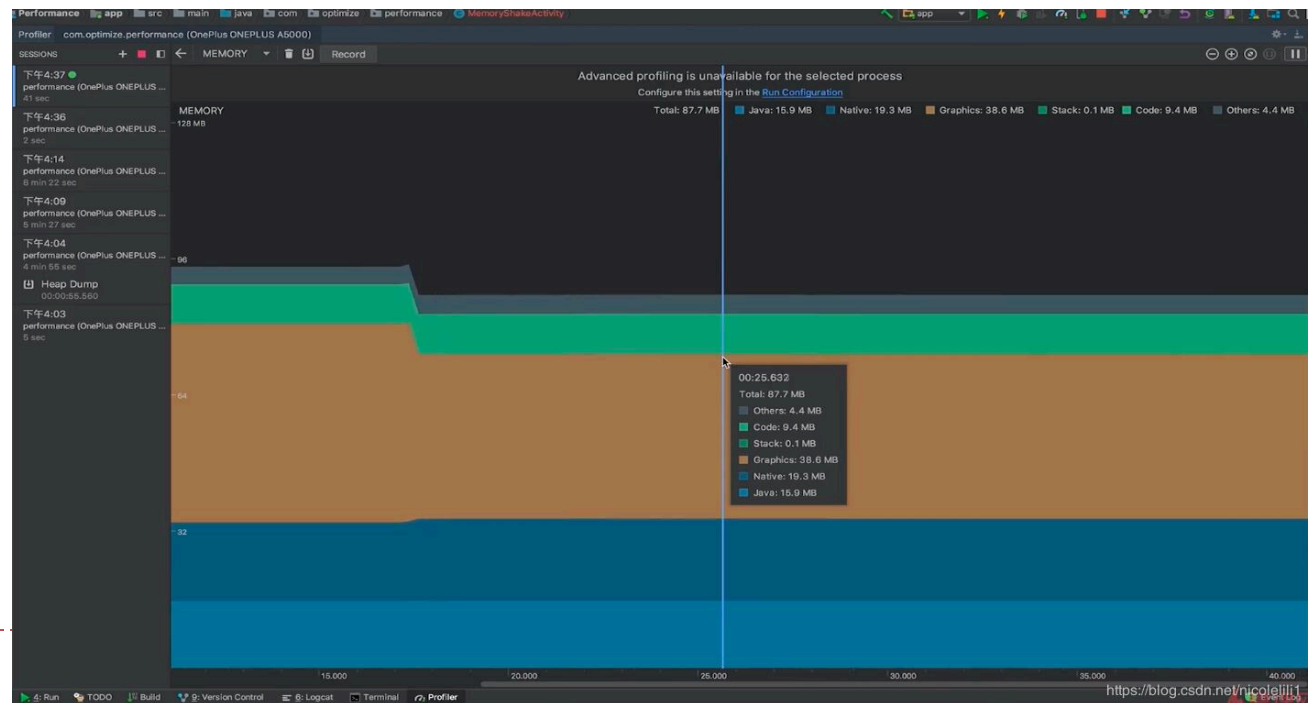
Local Page replacement problem: Thrashing




A0
A1
A2
A3
A4

Suppose A needs ≤ 6 pages during execution
and the memory space for A is 6

Memory performance is stable

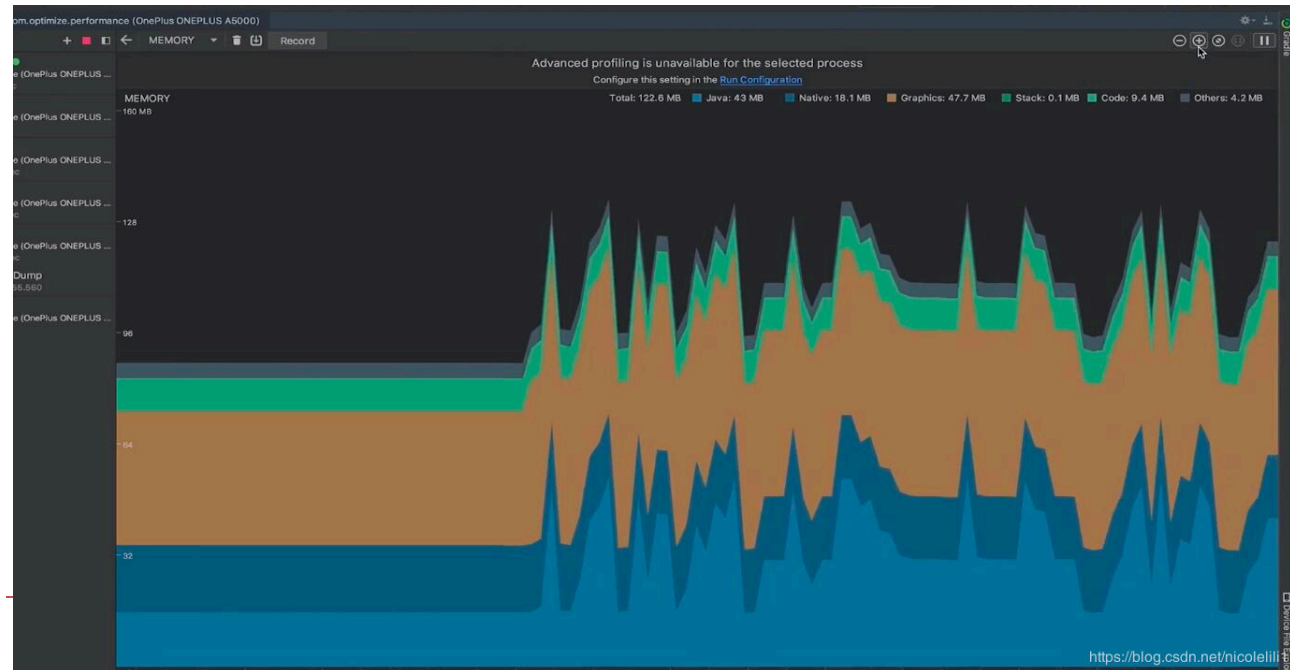


Local Page replacement problem: Thrashing



Suppose A needs **9** pages during execution and the memory space for A is only **6**


Memory performance is unstable and lots of CPU resources will be wasted on swapping



Local Page replacement problem

- Thrashing

- As the number of process in memory increases, the memory for each process decreases and page faults could also increase
- OS needs a *tradeoff* between parallelism and page fault rate



Global Page Replacement

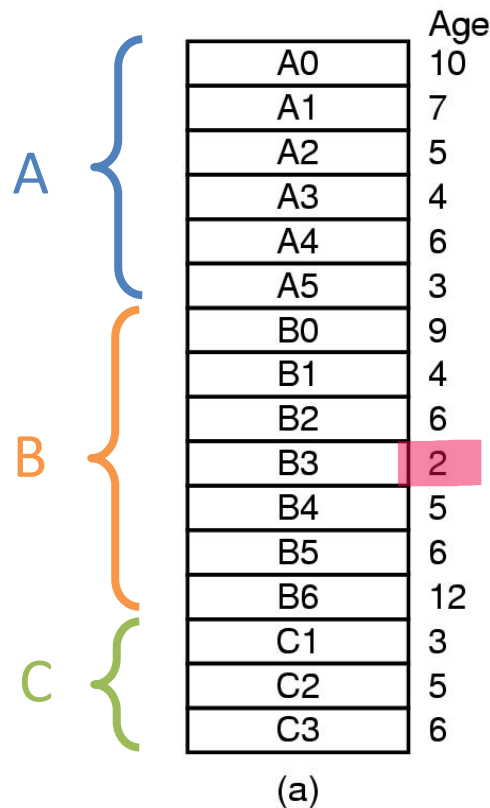


Diagram (a) shows the original configuration of memory frames. The frames are grouped into three sets: A (blue bracket), B (orange bracket), and C (green bracket). Each frame contains a page identifier and an age value. The age values are listed to the right of the frames.

	Page	Age
A	A0	10
	A1	7
	A2	5
	A3	4
	A4	6
	A5	3
B	B0	9
	B1	4
	B2	6
	B3	2
	B4	5
	B5	6
	B6	12
C	C1	3
	C2	5
	C3	6

(a)

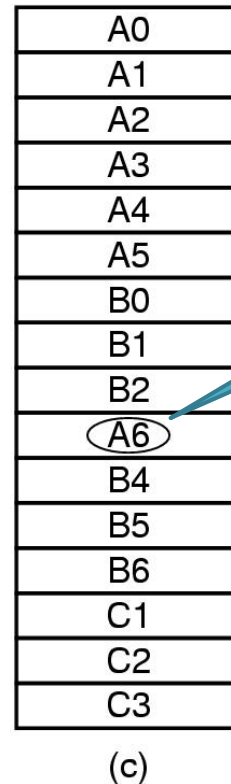


Diagram (c) shows the result of global page replacement. The frame that previously held B3 now holds A6. The age of A6 is not yet updated.

A0
A1
A2
A3
A4
A5
B0
B1
B2
A6
B4
B5
B6
C1
C2
C3

(c)

B3 will be replaced
by global algorithm

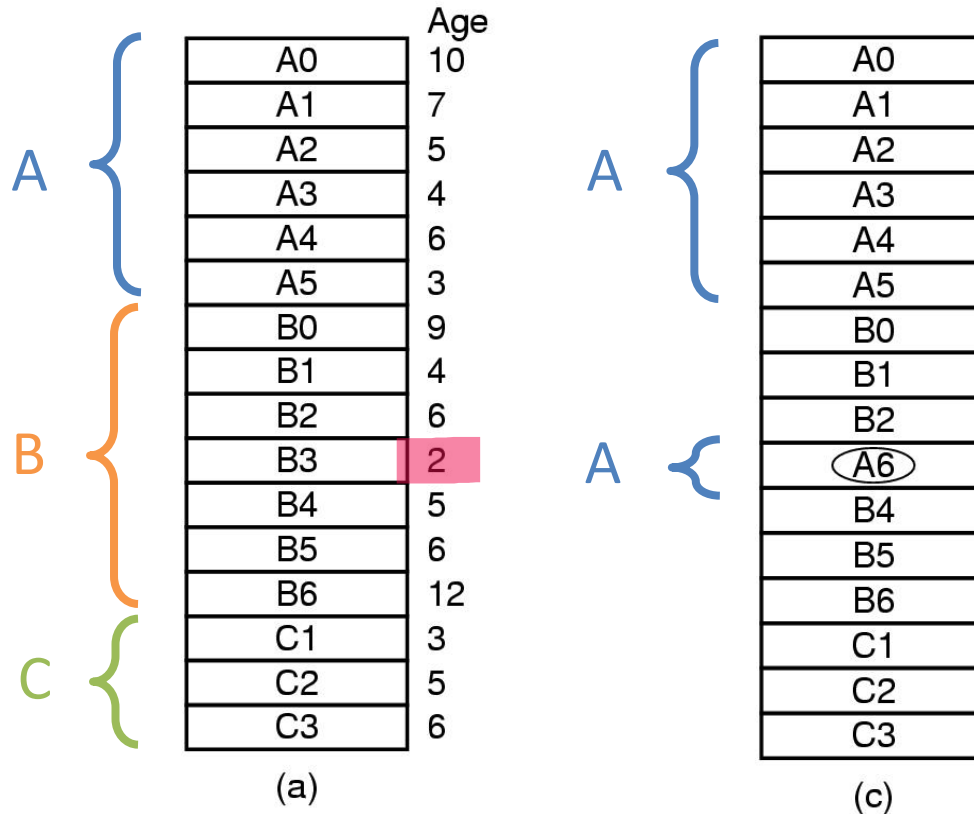
Suppose the algorithm
replaces the page which
has the least age.

e.g., A6 comes

(a) Original configuration.

(c) Global page replacement.

Global Page Replacement

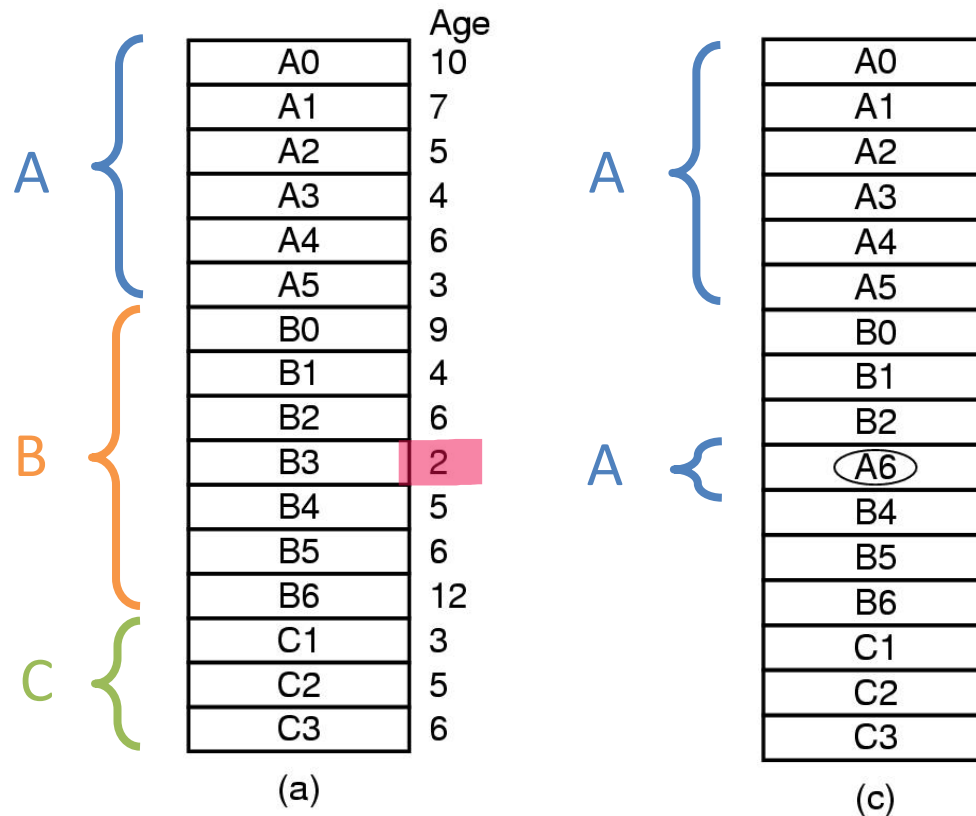


(a) Original configuration.

(c) Global page replacement.

- Global page replacement requires **dynamic allocation**
- The **page number** of one process will **change** during replacement (**A++** , **B--**)

Global Page Replacement Problem



(a) Original configuration.

(c) Global page replacement.

- How to control the page frames assigned to each process?
- Otherwise, some processes will take much more memories than others

Page Fault Frequency (PFF) algorithm

- PFF: control the size of allocation set of a process
 - when and how much to increase or decrease a process' page frame allocation

		Age
A	A0	10
	A1	7
	A2	5
	A3	4
	A4	6
	A5	3
B	B0	9
	B1	4
	B2	6
	B3	2
	B4	5
	B5	6
	B6	12
C	C1	3
	C2	5
	C3	6

Keep track of page fault
rate for each process


A: 20% → 25% → 30%

B: 20% → 15% → 10%

C: 50% → 50% → 50%



Page Fault Frequency (PFF) algorithm



Keep page faults not too high or too low


- Keep monitoring the page fault for each process and set the threshold
 - If one process page fault rate is too high, allocate more memory pages for it
 - If one process page fault rate is too low, allocate less memory pages for it

Local Page Replacement vs. Global Page Replacement

	Definition	Allocation	Potential problem	Possible solution
Local Page Replacement	Replace the page within one process	static	Thrashing	OS needs a tradeoff between parallelism and page fault rate
Global Page Replacement	Replace the page entire the system	dynamic	Page allocation imbalance among processes	Page Fault Frequency (PFF) and load control




Page Size



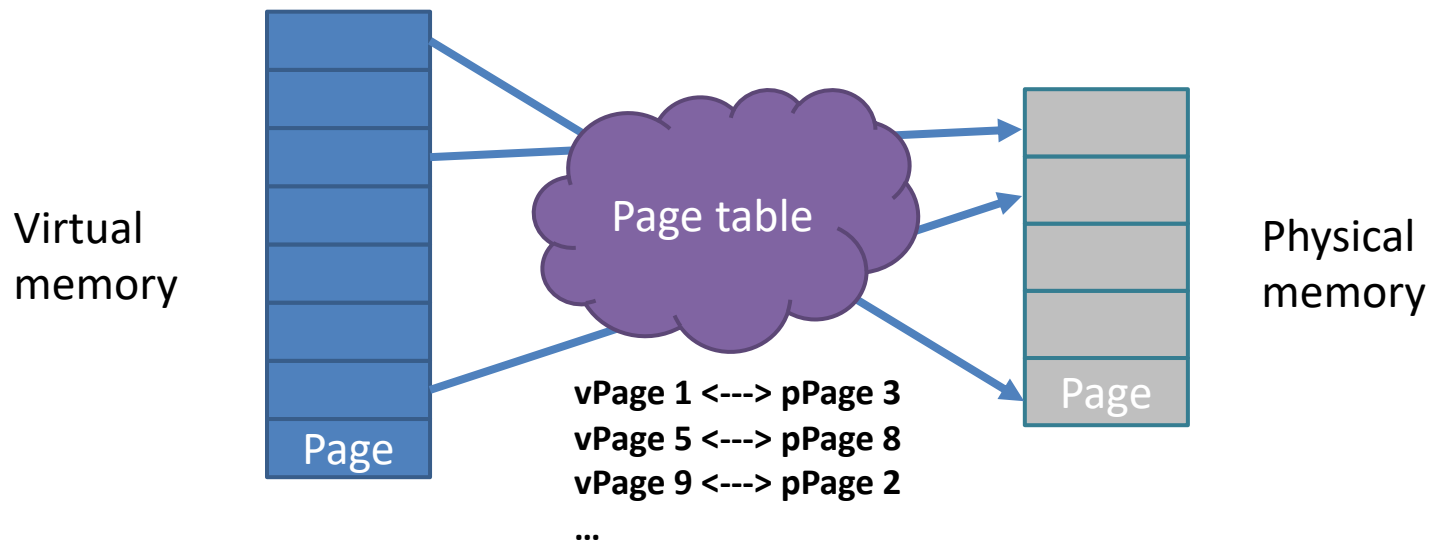
Page Size is small

- Advantages
 - less unused program in memory (due to *internal fragmentation*)
 - better fit for various data structures, code sections (e.g., 80% of the data structures or codes are small)




Page Size is small

- Disadvantages
 - Programs need many pages, larger page tables
 - Longer access time of page due to more pages
 - More page faults could happen due to more pages



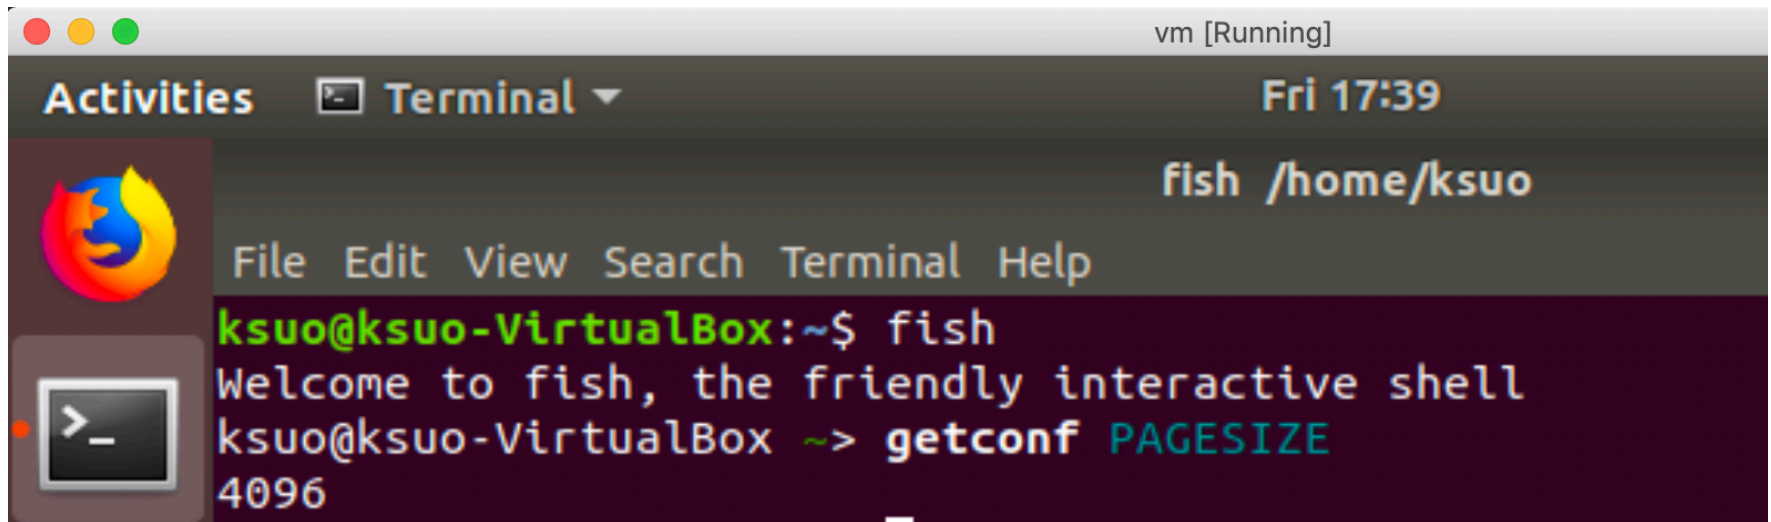
Page Size is large

- Disadvantages
 - More internal fragmentation and less efficiency
- Tradeoff between page size and memory efficiency
 - Normal we choose 4k for page size



\$ getconf PAGESIZE

- How to get page size in Linux

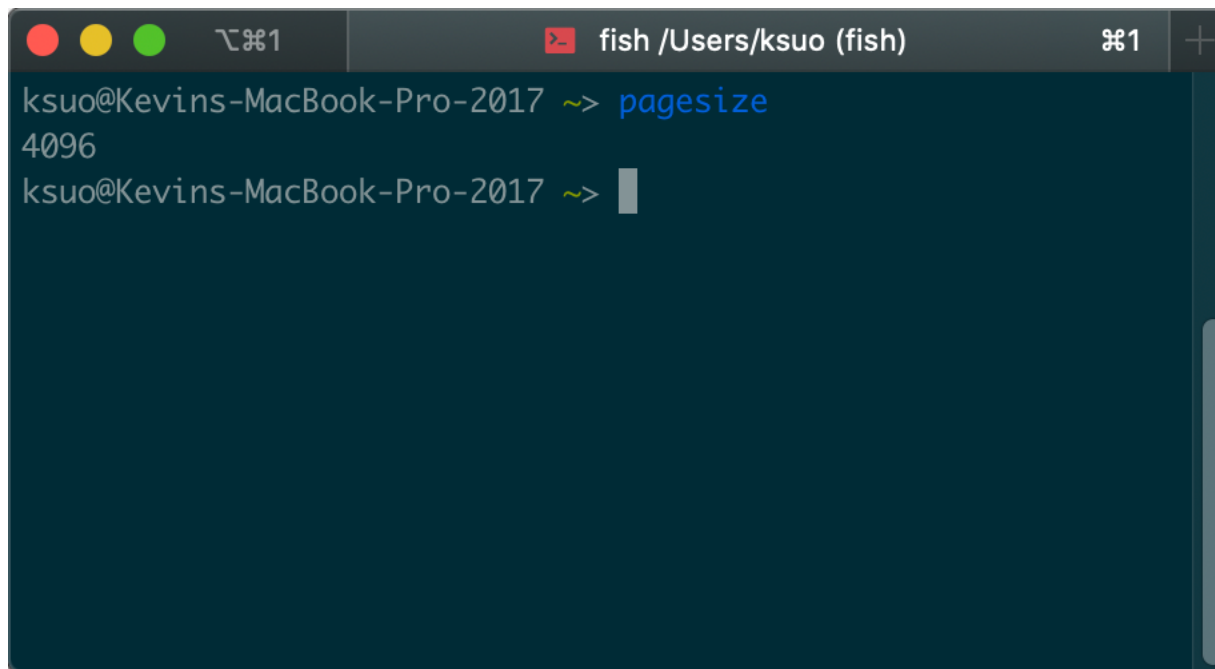


The screenshot shows a terminal window titled 'vm [Running]' with a timestamp of 'Fri 17:39'. The window has a menu bar with 'Activities', 'Terminal', and a dropdown arrow. Below the menu bar is a toolbar with icons for Firefox and a terminal. The terminal content shows the prompt 'ksuo@ksuo-VirtualBox:~\$' followed by the command 'fish'. Below that is a welcome message: 'Welcome to fish, the friendly interactive shell'. The prompt changes to 'ksuo@ksuo-VirtualBox ~>' and the command 'getconf PAGESIZE' is entered. The output '4096' is displayed below the command.

```
vm [Running]
Fri 17:39
fish /home/ksuo
File Edit View Search Terminal Help
ksuo@ksuo-VirtualBox:~$ fish
Welcome to fish, the friendly interactive shell
ksuo@ksuo-VirtualBox ~> getconf PAGESIZE
4096
```

\$ pagesize

- How to get page size on Mac

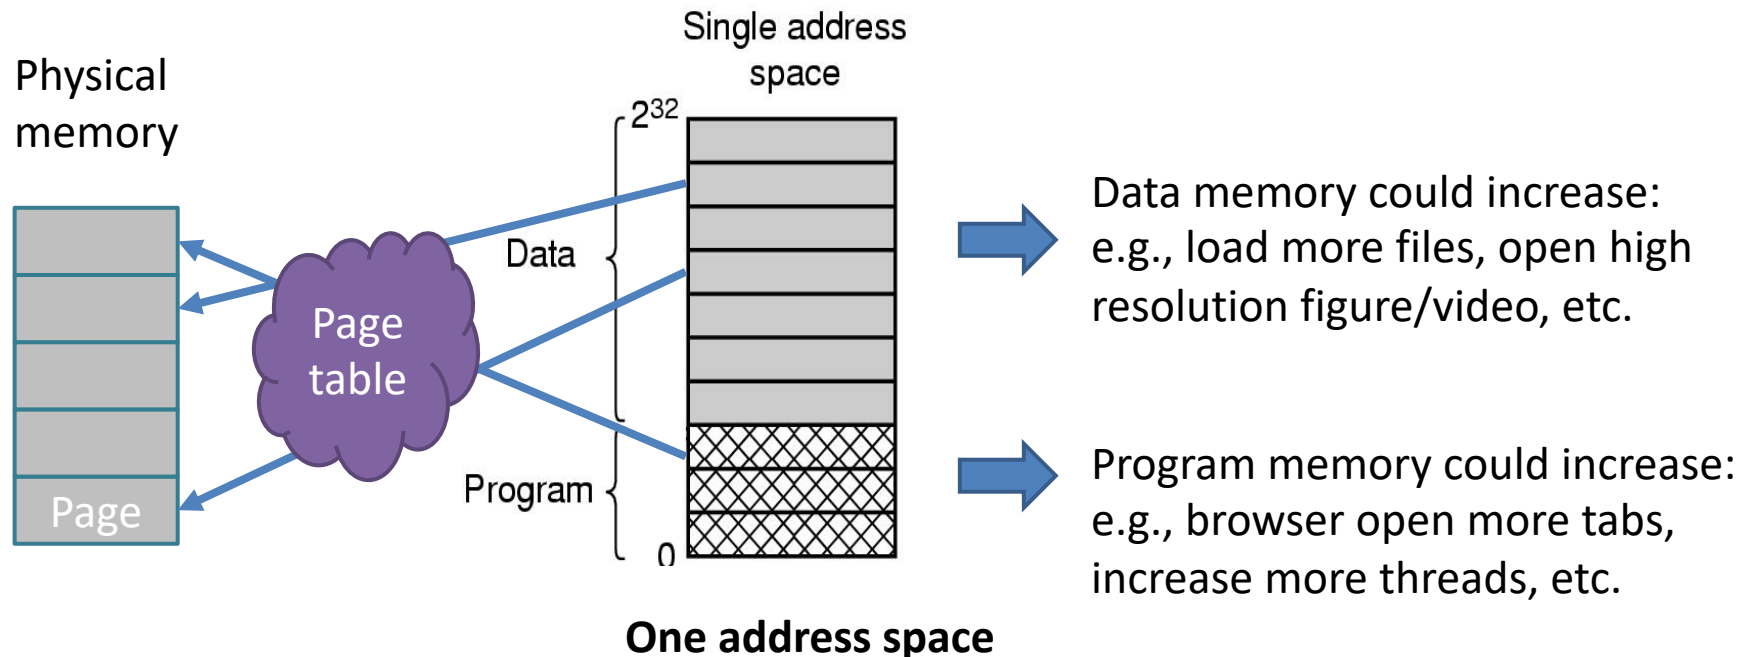


```
fish /Users/ksuo (fish)
ksuo@Kevins-MacBook-Pro-2017 ~> pagesize
4096
ksuo@Kevins-MacBook-Pro-2017 ~> 
```

A screenshot of a macOS terminal window. The window title bar shows three colored window control buttons (red, yellow, green) on the left, followed by a menu icon and the text 'fish /Users/ksuo (fish)'. The terminal content shows the prompt 'ksuo@Kevins-MacBook-Pro-2017 ~>' followed by the command 'pagesize' in blue. The output '4096' is displayed on the next line. The prompt is then shown again with a cursor, indicating the command has finished execution.


Separate Instruction and Data Spaces

- Normally, the memory address stores instruction and data of program together
 - Address space is limited
 - Interference between program and data memory space (e.g., security issue)





Separate Instruction and Data Spaces

- Separate memory address: I-space, D-space
 - Address space of instruction and data is independent
 - Both addresses have its own pages and page tables mapping for physical address to virtual address (sacrifice space for performance)




Shared Pages

- It is common that multiple users execute a same application
- In the cloud, multiple Oses usually run on the same hypervisor



Shared Pages

- To avoid multiple duplicates in memory, shared pages have more efficiency in memory design




Two processes sharing same program sharing its **I-page table**

Data pages are independent



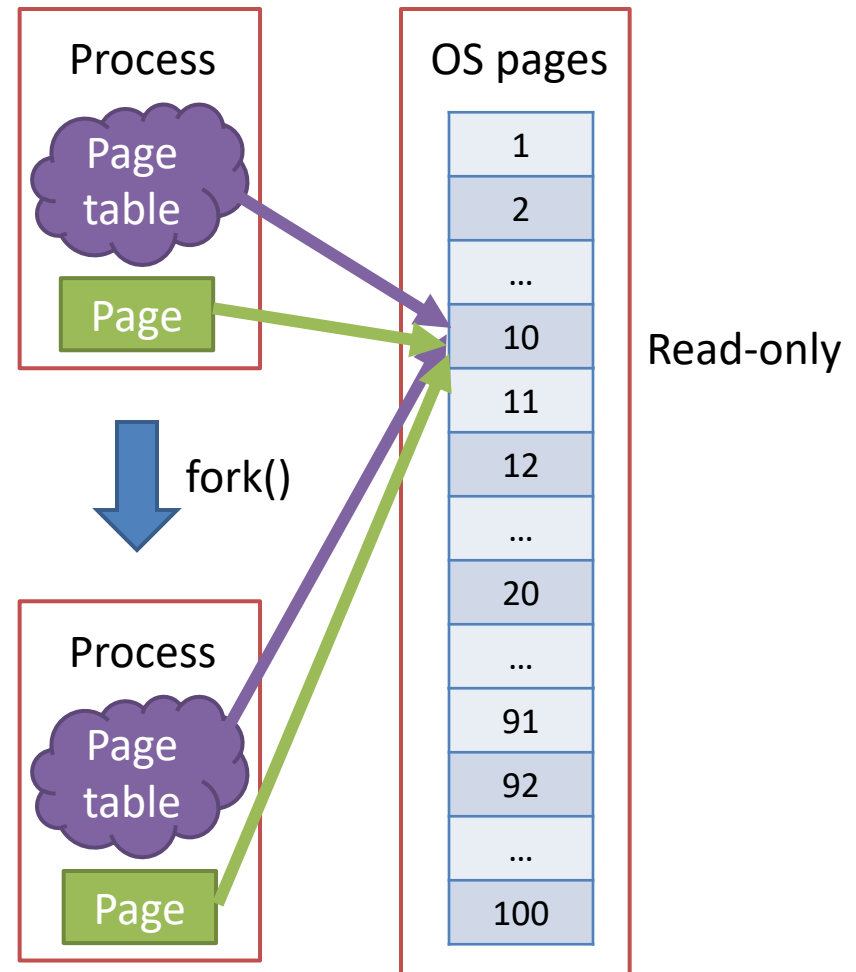
Shared Pages Problems

- Suppose Process A and B share the same I-Page. Then OS scheduler schedules out A and releases all pages of A. What will happen to B?

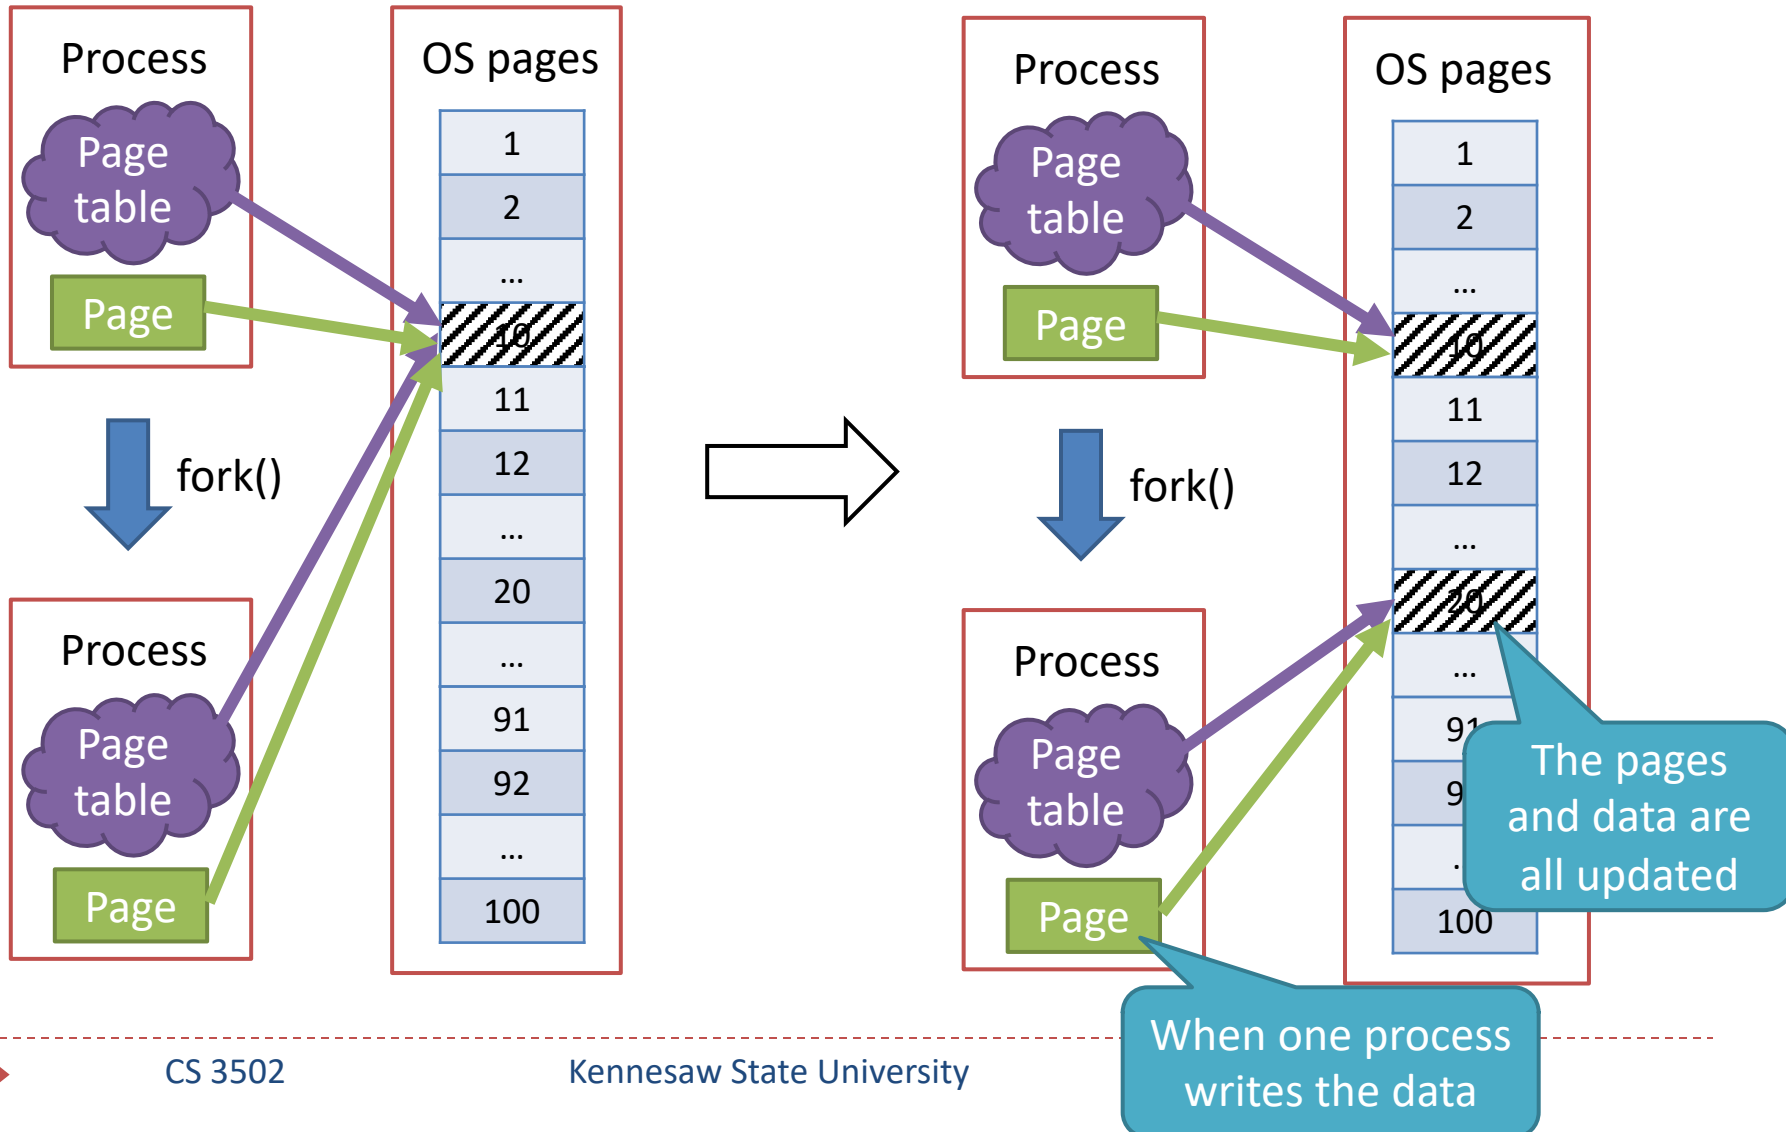


Shared Pages Example: fork()

- In `fork()`, two processes share program instruction and data in memory
- Two processes have their own page table but mapping to the same page
- All data pages are read-only




Copy-on-write(COW) for Shared Pages



Paging with Process Life Cycle

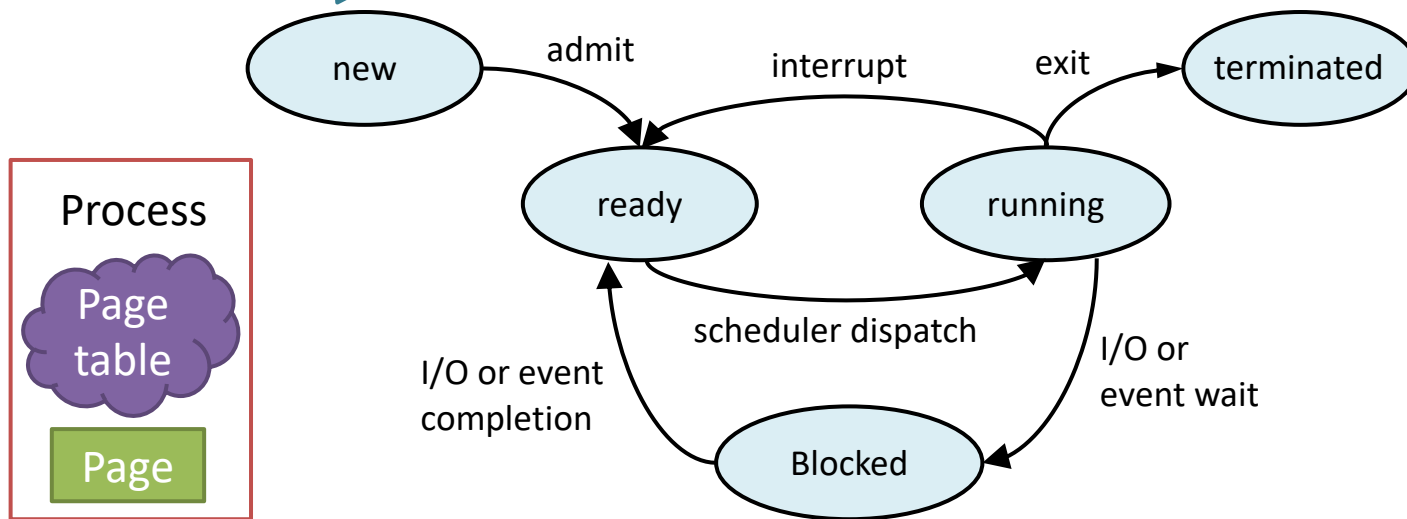
- During the process execution, how does the process interact with the page memory system?



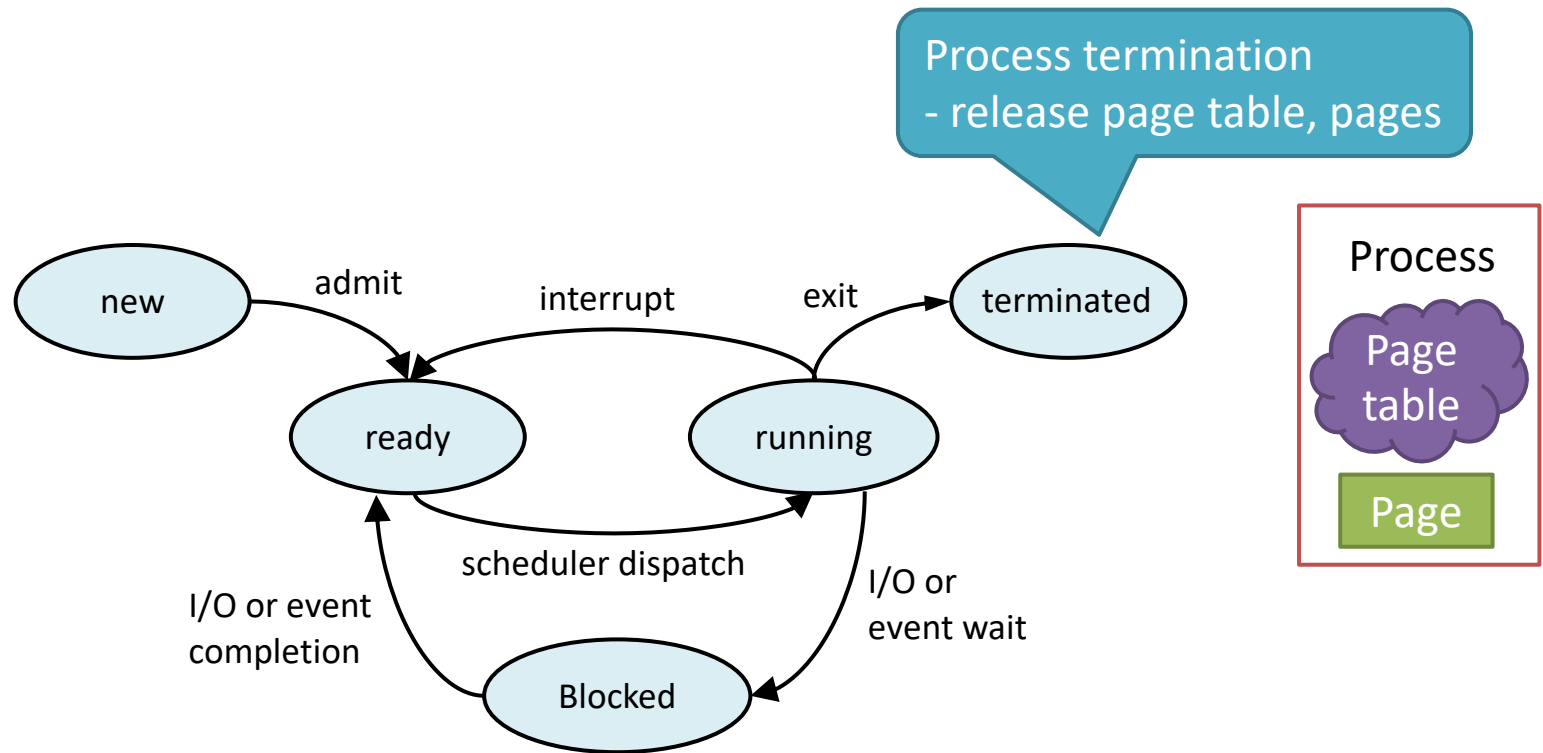
Paging with Process Life Cycle


Process creation:

- determine program size
- create page table



Paging with Process Life Cycle






virtual
address
space

Address translation
done by MMU

TLB

physical
memory

Keep updating as
address changes




Process execution

- MMU reset for new process
- TLB flushed

Reference string:

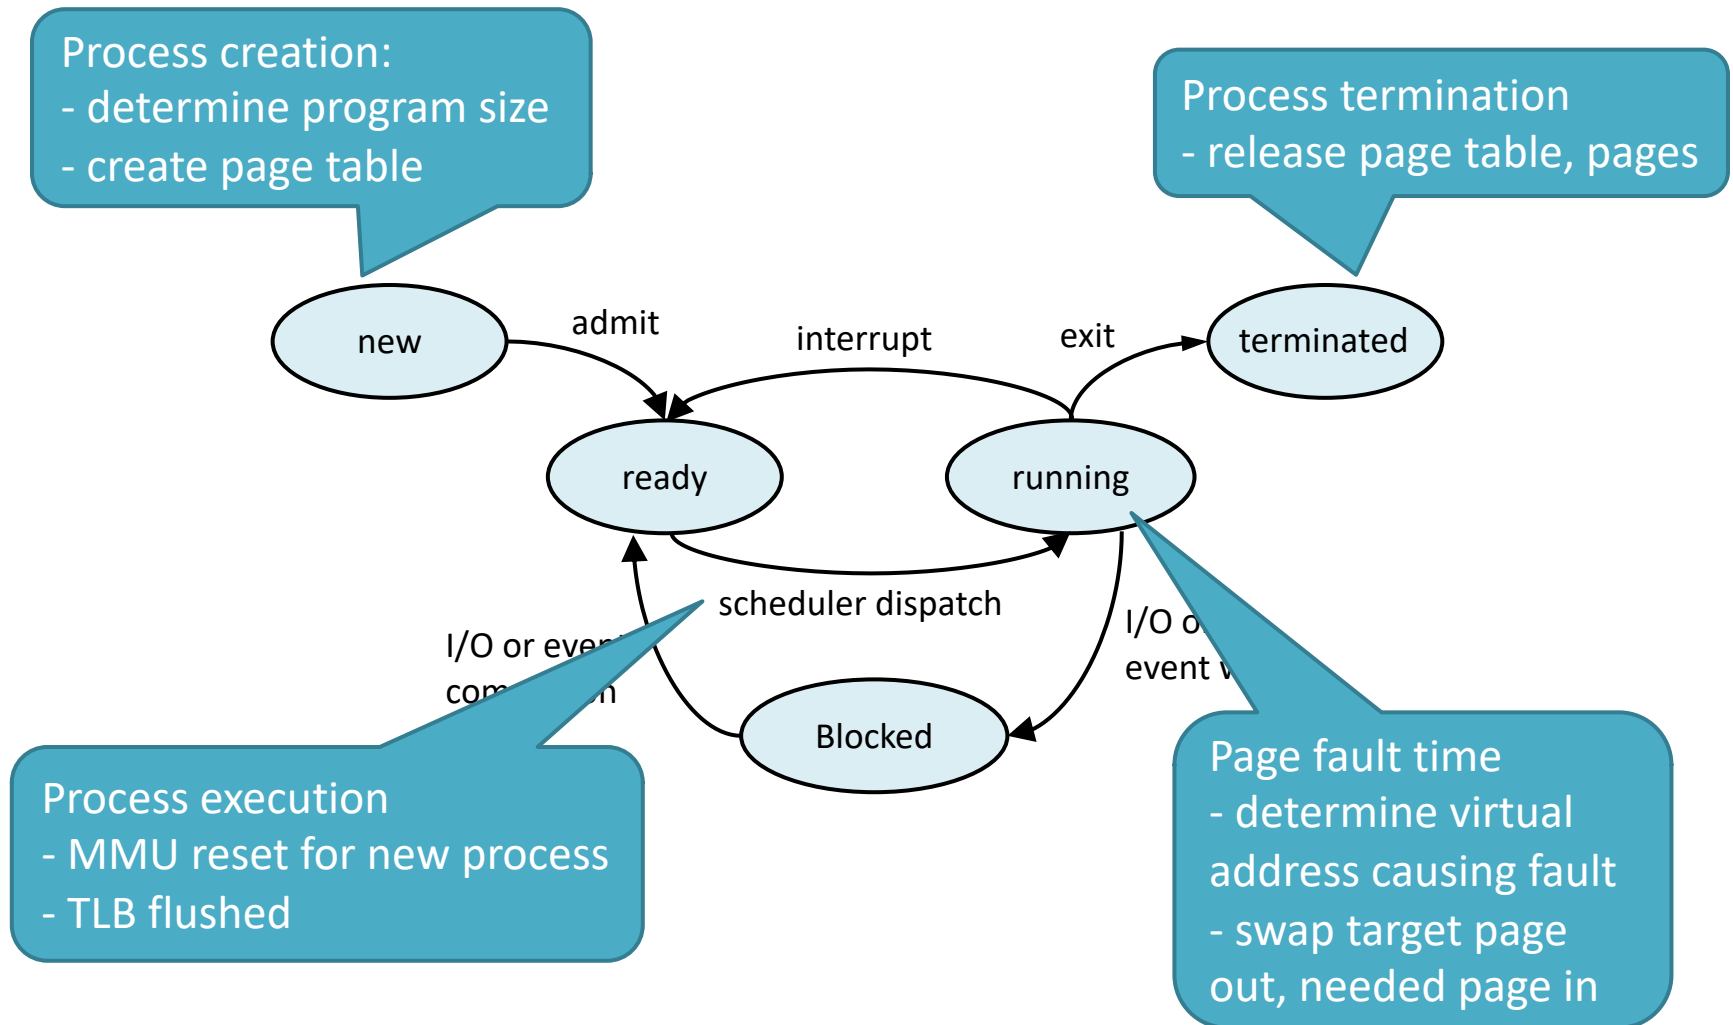
1	2	3	4	1	2	5	1	2	3	4	5
1	1	1	1	1	1	1	1	1	1	4	4
	2	2	2	2	2	2	2	2	2	2	2
		3	3	3	3	3	3	3	3	3	3
			4	4	4	5	5	5	5	5	5
X	X	X	X			X				X	

6 page faults



Page fault time
 - determine virtual address causing fault
 - swap target page out, needed page in

Paging with Process Life Cycle




Outline

- Page design
 - Internal fragmentation vs. External fragmentation
 - Local page replacement vs. Global page replacement
 - Page size small vs. Large
 - Shared page
 - Paging with Process life cycle
- Segmentation
 - Page vs. Segmentation



Segmentation: Rethink Pages

- Divide memory into pages
 - Fix size --> easy to build and management, but also its problem
 - Page has no relation with program logic



Segmentation

The user program address is divided into several segments of different sizes

Each segment can define a relatively complete set of logical information

Segment 0

Segment 2

Segment 4

```
vim /Users/ksuo (vim)
1 2,3,4
#include<stdio.h>
void quicksort(int number[25],int first,int last){
    int i, j, pivot, temp;


    if(first<last){
        pivot=first;
        i=first;
        j=last;

        while(i<j){
            while(number[i]<=number[pivot]&& i<last)
                i++;
            while(number[j]>number[pivot])
                j--;
            if(i<j){
                temp=number[i];
                number[i]=number[j];
                number[j]=temp;
            }

            temp=number[pivot];
            number[pivot]=number[j];
            number[j]=temp;
            quicksort(number,first,j-1);
            quicksort(number,j+1,last);
        }
    }

    int main(){
        int i, count, number[25];

        printf("How many elements are u going to enter?: ");
        scanf("%d",&count);
```




Logical Address Space

Segmentation

Allows each segment to grow or shrink, independently

- When storing allocation in segments, they can be non-contiguous in memory, and under discrete allocation



Logical Address Space


Segment Number

	base address	Limit
0	500	600
1	2500	800
2	1500	400
3	4600	200
4	3800	400

Segment Table

start

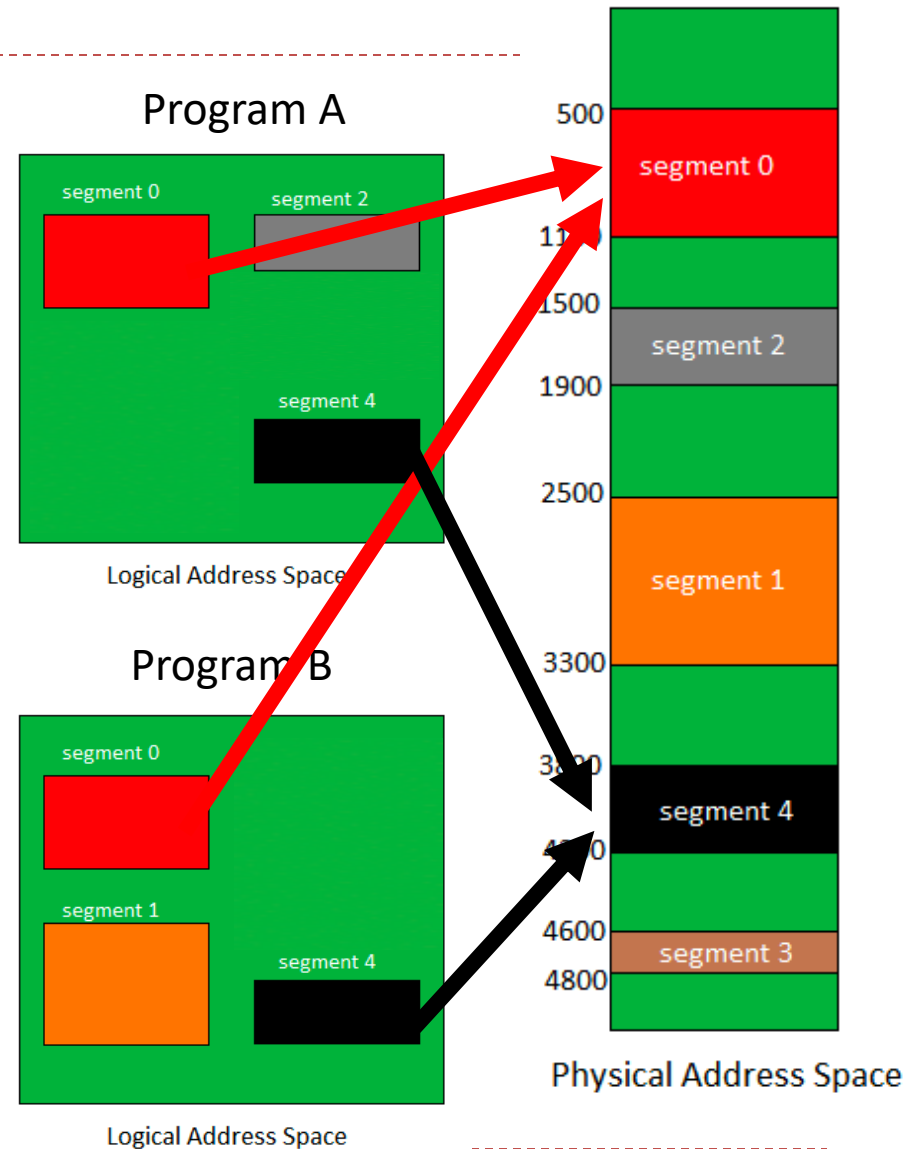
length



Physical Address Space


Segmentation

- Segmentation is convenient for multi-program sharing
 - E.g., Program A and B share the same piece of code or data in segment 0 and 4
 - Convenient for programming



Segmentation

- Advantages:
 - The **logical independence** of the segments makes it easy to compile, manage, modify, and protect, and is also convenient for multi-program sharing
 - The segment length can be **dynamically** changed as needed, allowing free scheduling to make efficient use of the main memory space
 - **Convenient** for programming, including segment sharing, segmentation protection, dynamic linking, dynamic growth




Logical Address Space

Segment Number

	base address	Limit
0	500	600
1	2500	800
2	1500	400
3	4600	200
4	3800	400

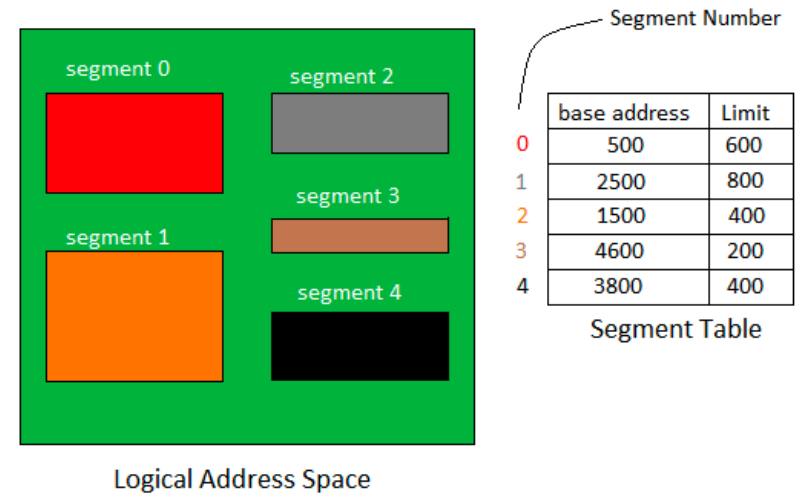
Segment Table



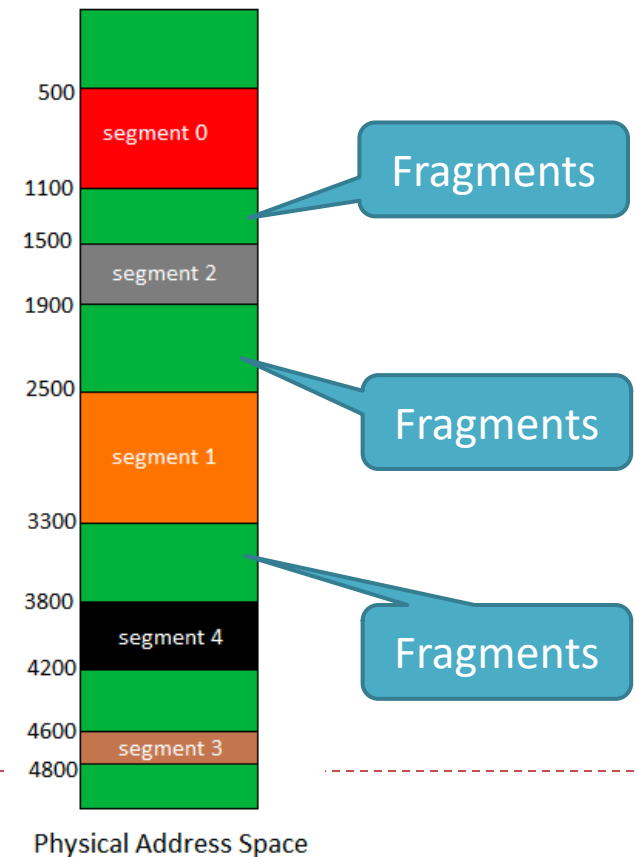
Physical Address Space

Segmentation

- Disadvantages:
 - The **allocation** of memory space is difficult (what size, where, etc.)
 - It is easy to leave a lot of **fragments** between the segments, resulting in a decrease in space utilization




Logical Address Space



Paged vs. Segmented Virtual Memory

- Paged virtual memory
 - Memory divided into fixed sized pages
- Segmented virtual memory
 - Memory divided into variable length segments




Comparison of Segmentation and Paging

	Page	Segment
Definition	Main memory is partitioned into same sized pages	Main memory is partitioned into various segments
Address	One word (start address)	Two words (start and end address)
Programmer visible	No	Yes
Block replacement	Easy	Hard
Fragmentation	Internal	External



Combination of Page and Segment



Conclusion

- Page design
 - Internal fragmentation vs. External fragmentation
 - Local page replacement vs. Global page replacement
 - Page size small vs. Large
 - Shared page
 - Paging with Process life cycle
- Segmentation
 - Page vs. Segmentation

