

CS 4504

Parallel and Distributed Computing

Process

Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

Outline

- What is process?
 - Process vs Program
 - Linux Process Control Block
- Process related System calls
 - Fork
 - Exec
 - Wait
- Process on Distributed OSes



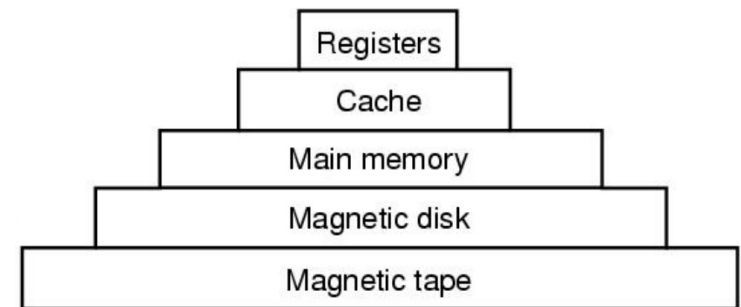
Process

- Definition

- An instance of a *program* running on a computer
- An *abstraction* that supports running programs - -> cpu virtualization
- An *execution stream* in the context of a particular *process state* - -> dynamic unit
- A *sequential stream* of execution in its *own address space* - -> execution code line by line

Process

- Two parts of a process
 - Sequential execution of instructions
 - Process state
 - ▶ registers: PC (program counter), SP (stack pointer),...
 - ▶ Memory: address space, code, data, stack, heap ...
 - ▶ I/O status: open files ...

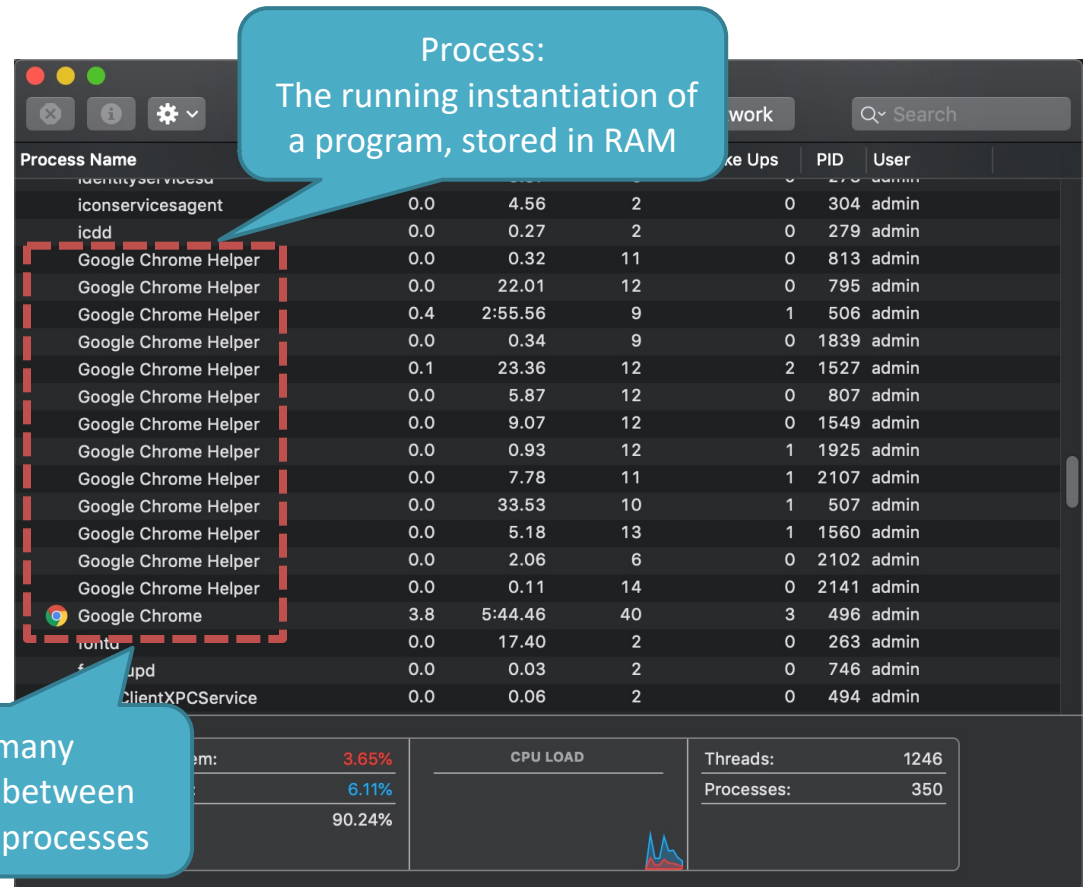
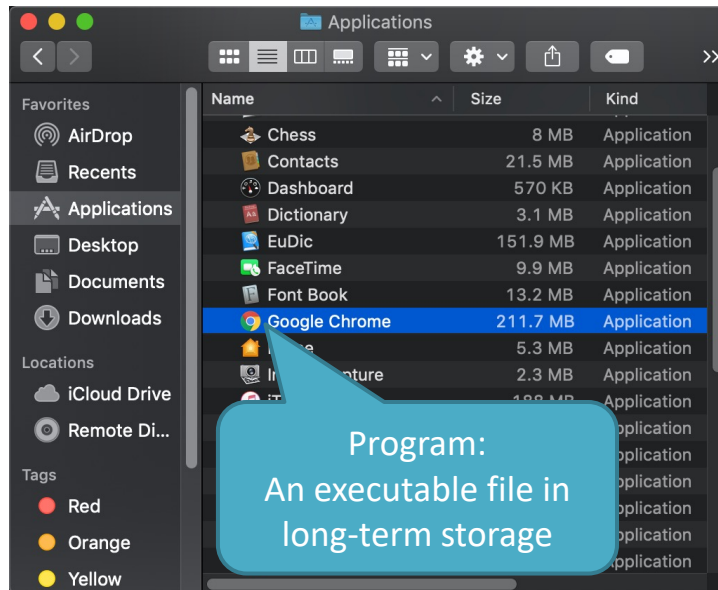


Program vs. Process

- Program != Process
 - Program = static code + data
 - Process = dynamic instantiation of code + data + files ...
- No 1:1 mapping
 - Program : process = 1:N
 - ▶ A program can invoke many processes



Program vs. Process



One-to-many
relationship between
program and processes

Program vs. Process

BASIS FOR COMPARISON	PROGRAM	PROCESS
Basic	Program is a set of instruction.	When a program is executed, it is known as process.
Nature	Passive	Active
Lifespan	Longer	Limited
Required resources	Program is stored on disk in some file and does not require any other resources.	Process holds resources such as CPU, memory address, disk, I/O etc.

<https://techdifferences.com/difference-between-program-and-process.html>



Process Descriptor



- Driving license

- ID
- Name
- Address
- Birth
- Time
- ...

- Process control block (PCB)

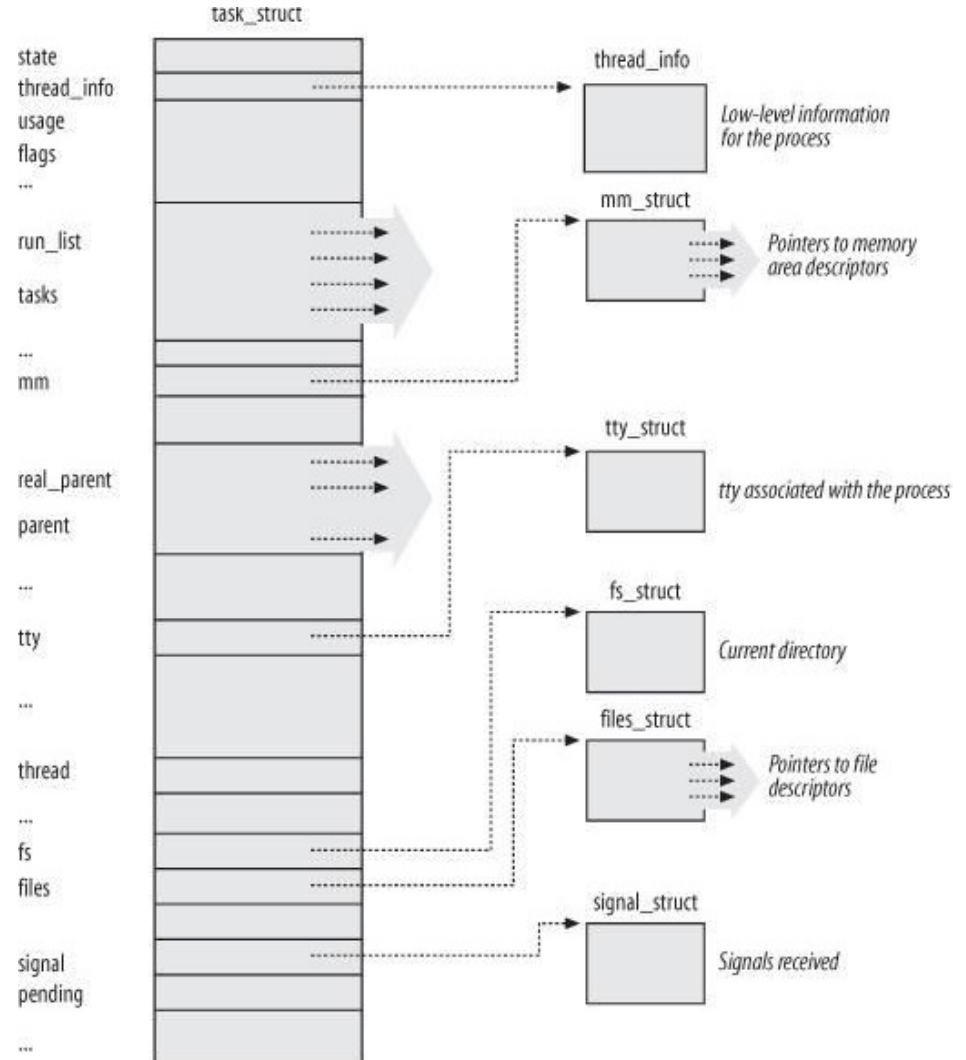
- State
- Identifiers
- Scheduling info
- File system
- Virtual memory
- Process specific context
- ...




Process in Linux

<https://elixir.bootlin.com/linux/v5.4/source/include/linux/sched.h#L624>

- Process control block (PCB)
 - State
 - Identifiers
 - Scheduling info
 - File system
 - Virtual memory
 - Process specific context
 - ...



Process in Linux

- Process control block (PCB)
 - State 
 - Identifiers
 - Scheduling info
 - File system
 - Virtual memory
 - Process specific context
 - ...

Linux PCB (5-state model to describe lifecycle of one process)

- State

- TASK_RUNNING

- ▶ Running, ready

- TASK_INTERRUPTIBLE

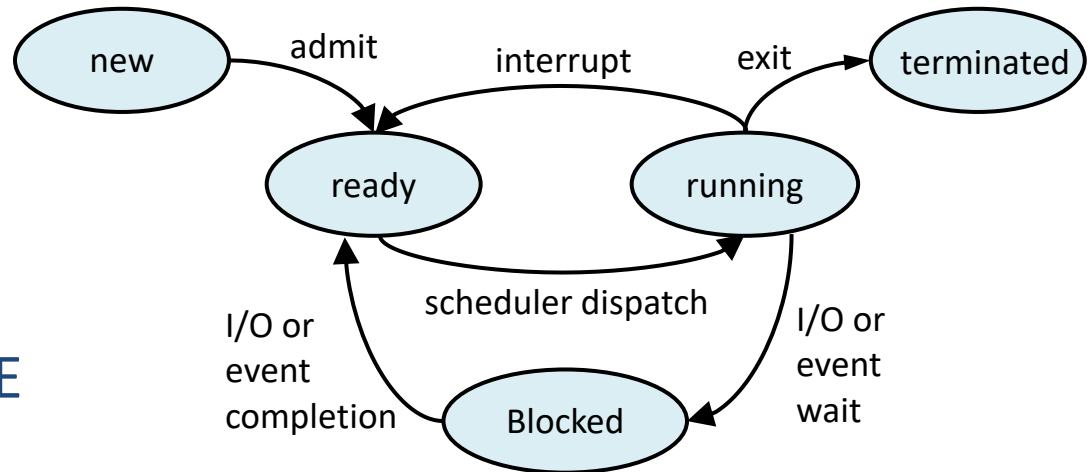
- ▶ Blocked

- EXIT_ZOMBIE

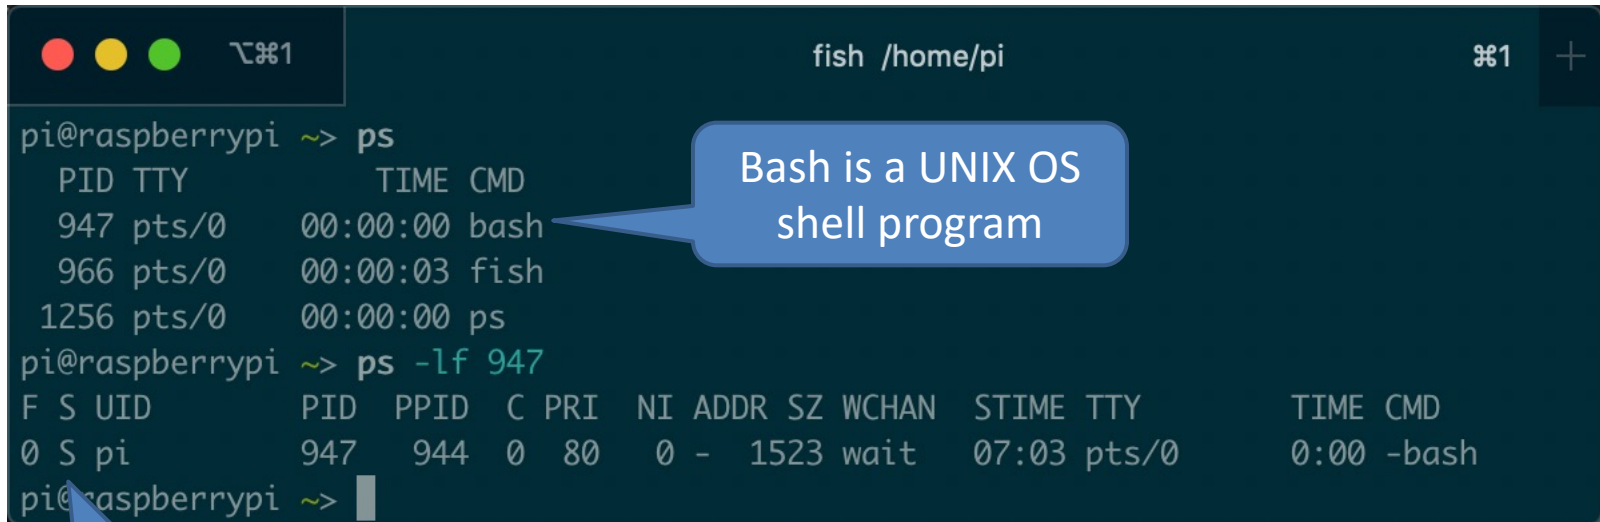
- ▶ Terminated by not deallocated

- EXIT_DEAD

- ▶ Completely terminated



\$ ps -lf : process information



The image shows a terminal window on a Raspberry Pi. The top bar indicates the window title is 'fish /home/pi' and the user is 'pi'. The terminal shows the command 'ps' being executed, which lists running processes. A blue callout bubble points to the 'bash' process, stating 'Bash is a UNIX OS shell program'. Below this, the command 'ps -lf 947' is executed, showing detailed information for PID 947. Another blue callout bubble points to the 'S' state in the output, stating 'The state of the process'.

```
pi@raspberrypi ~> ps
  PID TTY          TIME CMD
  947 pts/0    00:00:00 bash
  966 pts/0    00:00:03 fish
 1256 pts/0    00:00:00 ps
pi@raspberrypi ~> ps -lf 947
F S UID          PID  PPID  C PRI  NI ADDR SZ WCHAN  STIME TTY          TIME CMD
0 S pi            947   944  0  80   0 -  1523 wait  07:03 pts/0    0:00 -bash
pi@raspberrypi ~>
```

The state of the process

- R : The process is running
- S : The process is sleeping/idle
- T : The process is terminated
- Z : The process is in zombie state

\$ ps -lf : process information

<https://github.com/kevinsuo/CS7172/blob/master/sleep.c>

<https://github.com/kevinsuo/CS7172/blob/master/loop.c>



What is the state of the process?

Download: wget Raw-file-URL


Compile: gcc [file-name].c -o [file-name].o

Run: ./[file-name].o

```
pi@raspberrypi ~> ps -lf 8021
F S UID          PID  PPID  C PRI  NI ADDR SZ WCHAN  STIME TTY          TIME CMD
0 S pi            8021  6190  0  80   0 -   450 hrtime 18:02 pts/0    0:00 ./test.o
```

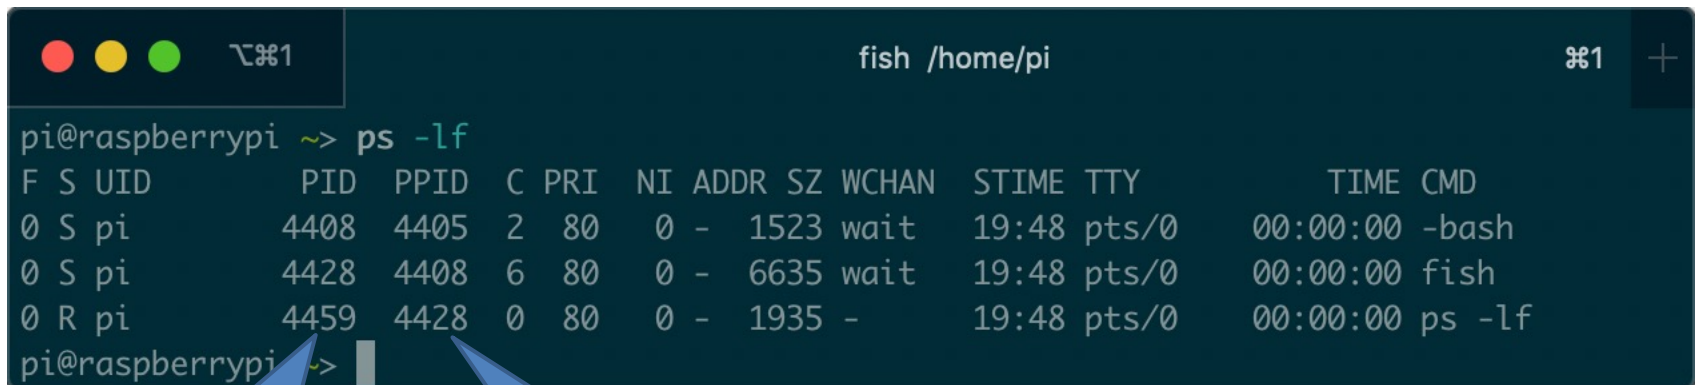
```
pi@raspberrypi ~> ps -lf 10057
F S UID          PID  PPID  C PRI  NI ADDR SZ WCHAN  STIME TTY          TIME CMD
0 R pi           10057  6190 99  80   0 -   450 -      18:06 pts/0    0:14 ./test.o
```

Process in Linux

- Process control block (PCB)
 - State
 - Identifiers 
 - Scheduling info
 - File system
 - Virtual memory
 - Process specific context
 - ...

Linux Process Control Block (cont')

- Identifiers
 - pid: ID of the process




```
pi@raspberrypi ~> ps -lf
F S UID          PID  PPID  C PRI  NI ADDR SZ WCHAN  STIME TTY          TIME CMD
0 S pi            4408  4405  2  80   0 - 1523 wait   19:48 pts/0    00:00:00 -bash
0 S pi            4428  4408  6  80   0 - 6635 wait   19:48 pts/0    00:00:00 fish
0 R pi            4459  4428  0  80   0 - 1935 -      19:48 pts/0    00:00:00 ps -lf
pi@raspberrypi ~>
```

ID for this process

Parent process ID



Process in Linux

- Process control block (PCB)
 - State
 - Identifiers
 - Scheduling info 
 - File system
 - Virtual memory
 - Process specific context
 - ...

Linux Process Control Block (cont')

- Scheduling information
 - prio, static_prio, normal_prio
 - rt_priority
 - sched_class



Linux Process Control Block (cont')

- Scheduling information
 - `prio`, `static_prio`, `normal_prio`



- (1) Static priority: $P_1 > P_2 = P_3 = P_4$, P_1 can execute whenever it needs;
- (2) Normal priority: $P_1 = P_2 = P_3 = P_4$, P_1 execute depending on the scheduling algorithm;
- (3) Prio: dynamic priority, will change over the time

Linux Process Control Block (cont')

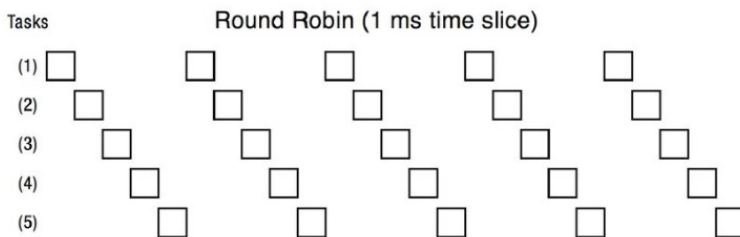
- Scheduling information
 - `rt_priority`



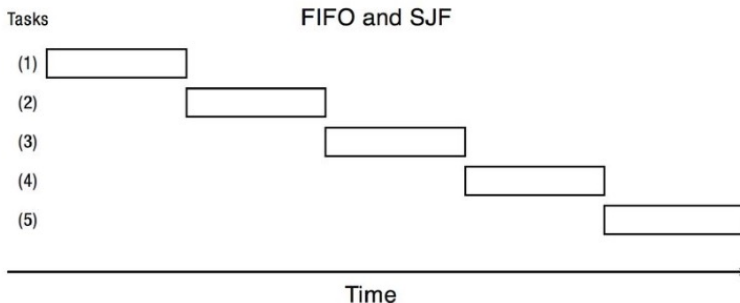
Rt_priority process is always higher than other priority of processes and will be scheduled immediately when it needs

Linux Process Control Block (cont')

- Scheduling information
 - `sched_class`: different scheduling policy implementations, e.g., FIFO, SJF, RR...
 - ▶ Task->sched_class->pick_next_task(runqueue)



`pick_next_task` of RR:
pick based on time cycle



`pick_next_task` of FIFO:
pick based on order



\$ ps -lf : process information

```
pi@raspberrypi ~> ps
  PID TTY          TIME CMD
  947 pts/0        00:00:00 bash
  966 pts/0        00:00:03 fish
 1256 pts/0        00:00:00 ps
pi@raspberrypi ~> ps -lf 947
F S UID        PID  PPID  C PRI  NI ADDR SZ WCHAN  STIME TTY          TIME CMD
0 S pi          947   944  0  80   0 -  1523 wait   07:03 pts/0    0:00 -bash
pi@raspberrypi ~>
```

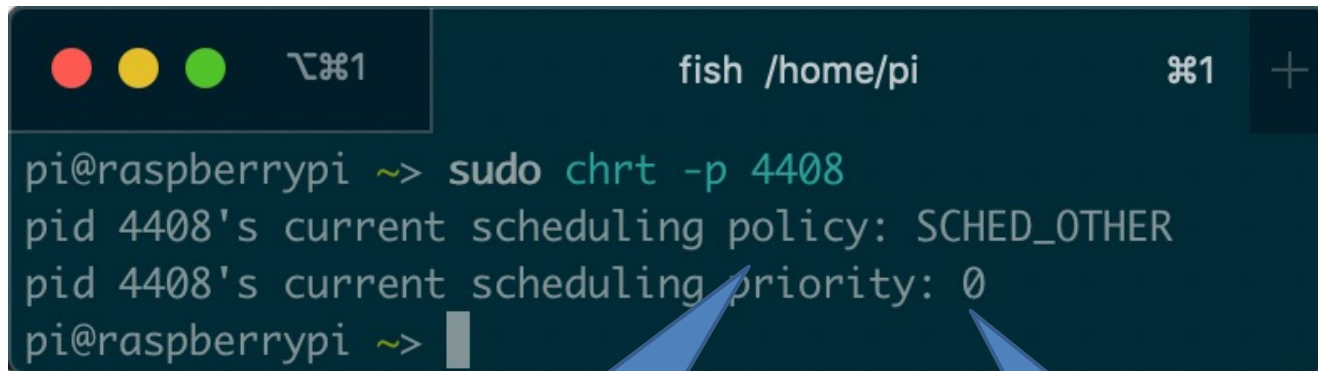
Bash is a UNIX OS shell program

How many cpus it consumes

Process priority value

Nice value: default is 0, could be modified to adjust the priority

\$ chrt: process scheduling info



```
pi@raspberrypi ~> sudo chrt -p 4408
pid 4408's current scheduling policy: SCHED_OTHER
pid 4408's current scheduling priority: 0
pi@raspberrypi ~>
```

The image shows a terminal window with a dark background. The window title bar includes window control buttons (red, yellow, green) and the text 'fish /home/pi'. The terminal content shows the command 'sudo chrt -p 4408' being executed, followed by two lines of output: 'pid 4408's current scheduling policy: SCHED_OTHER' and 'pid 4408's current scheduling priority: 0'. The prompt 'pi@raspberrypi ~>' is visible at the end of the output.


scheduling policy

scheduling priority

SCHED_OTHER
SCHED_FIFO
SCHED_RR
SCHED_BATCH

min/max priority : 0/0
min/max priority : 1/99
min/max priority : 1/99
min/max priority : 0/0

Process in Linux

- Process control block (PCB)
 - State
 - Identifiers
 - Scheduling info
 - File system 
 - Virtual memory
 - Process specific context
 - ...

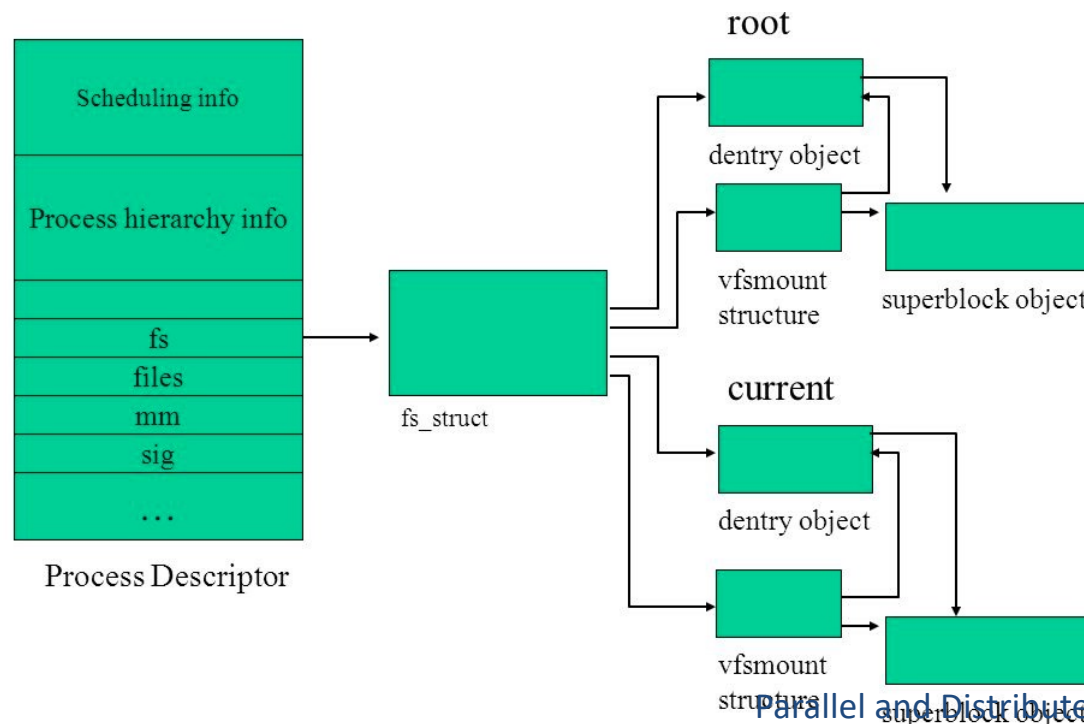
Linux Process Control Block (cont')

- Files

- fs_struct

<https://elixir.bootlin.com/linux/v4.2/source/include/linux/sched.h#L1525>

- file system information: root directory, current directory



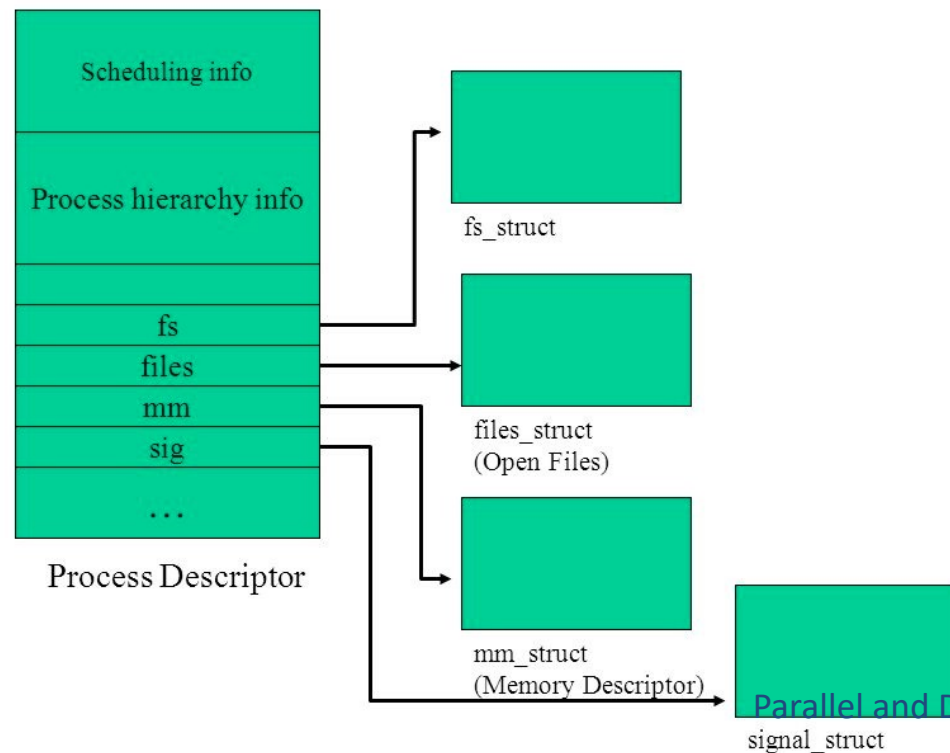
Linux Process Control Block (cont')

- Files

- files_struct

<https://elixir.bootlin.com/linux/v4.2/source/include/linux/sched.h#L1528>

- Information on opened files




\$lsof: list all open files

```
pi@raspberrypi ~$ ps
  PID TTY          TIME CMD
  947 pts/0    00:00:00 bash
  966 pts/0    00:00:02 fish
 1209 pts/0    00:00:00 ps
pi@raspberrypi ~$ lsof -p 947
COMMAND PID  USER   FD   TYPE    DEVICE  SIZE/OFF      NODE NAME
bash    947   pi     cwd   DIR     179,7    4096  1572867 /home/pi
bash    947   pi     rtd   DIR     179,7    4096        2 /
bash    947   pi     txt   REG     179,7   912712  524329  /bin/bash
bash    947   pi     mem   REG     179,7   38560  1445488 /lib/arm-linux-gnueabi/libnss_files-2.24.so
bash    947   pi     mem   REG     179,7   38588  1445507 /lib/arm-linux-gnueabi/libnss_nis-2.24.so
bash    947   pi     mem   REG     179,7   71604  1445594 /lib/arm-linux-gnueabi/libnsl-2.24.so
bash    947   pi     mem   REG     179,7   26456  1445612 /lib/arm-linux-gnueabi/libnss_compat-2.24.so
bash    947   pi     mem   REG     179,7 1679776  670175 /usr/lib/locale/locale-archive
bash    947   pi     mem   REG     179,7 1234700  1445500 /lib/arm-linux-gnueabi/libc-2.24.so
bash    947   pi     mem   REG     179,7   9800  1445460 /lib/arm-linux-gnueabi/libdl-2.24.so
bash    947   pi     mem   REG     179,7  124808  1445519 /lib/arm-linux-gnueabi/libtinfo.so.5.9
bash    947   pi     mem   REG     179,7   21868  144001  /usr/lib/arm-linux-gnueabi/libarmmem.so
bash    947   pi     mem   REG     179,7  138576  1445547 /lib/arm-linux-gnueabi/ld-2.24.so
bash    947   pi     mem   REG     179,7   26262  145746  /usr/lib/arm-linux-gnueabi/gconv/gconv-modules.cache
bash    947   pi      0u   CHR    136,0    0t0        3 /dev/pts/0
bash    947   pi      1u   CHR    136,0    0t0        3 /dev/pts/0
bash    947   pi      2u   CHR    136,0    0t0        3 /dev/pts/0
bash    947   pi    255u   CHR    136,0    0t0        3 /dev/pts/0
pi@raspberrypi ~$
```

All files opened by bash

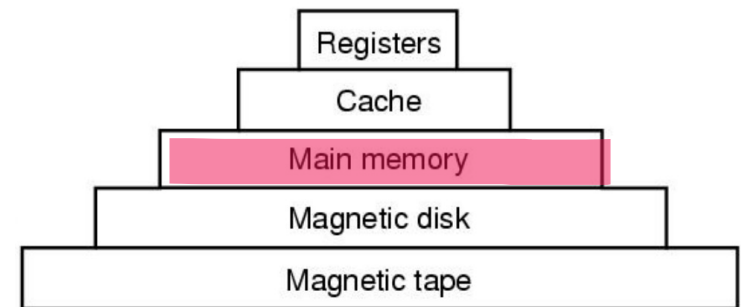
File descriptor, size, name, location, ...

Process in Linux

- Process control block (PCB)
 - State
 - Identifiers
 - Scheduling info
 - File system
 - Virtual memory 
 - Process specific context
 - ...

Linux Process Control Block (cont')

- Virtual memory
 - `mm_struct`: describes the content of a process's virtual memory
 - The pointer to the page table and the virtual memory areas



\$pmap: memory mapping

```
pi@raspberrypi ~> ps
  PID TTY          TIME CMD
  947 pts/0        00:00:00 bash
  966 pts/0        00:00:01 fish
 1026 pts/0        00:00:00 ps
pi@raspberrypi ~> pmap 947
947:  -bash
00010000   872K r-x-- bash
000f9000    4K r---- bash
000fa000   20K rw--- bash
000ff000   36K rw--- [ anon ]
00533000  1088K rw--- [ anon ]
76bac000   36K r-x-- libnss_files-2.24.so
76bb5000   60K ----- libnss_files-2.24.so
76bc4000    4K r---- libnss_files-2.24.so
76bc5000    4K rw--- libnss_files-2.24.so
76bc6000   24K rw--- [ anon ]
76bcc000   36K r-x-- libnss_nis-2.24.so
76bd5000   60K ----- libnss_nis-2.24.so
76be4000    4K r---- libnss_nis-2.24.so
76be5000    4K rw--- libnss_nis-2.24.so
76be6000   68K r-x-- libnsl-2.24.so
76bf7000   60K ----- libnsl-2.24.so
76c06000    4K r---- libnsl-2.24.so
76c07000    4K rw--- libnsl-2.24.so
76c08000    8K rw--- [ anon ]
76c0a000   24K r-x-- libnss_compat-2.24.so
76c0b000   60K ----- libnss_compat-2.24.so
76c1f000    4K r---- libnss_compat-2.24.so
```

All memory accessed by bash

Memory address

size

Read/write/
execution
permission

Lib/execut
ion file

memory used by
process bash

\$pmap: memory mapping

```
pi@raspberrypi ~> ps
  PID TTY          TIME CMD
  947 pts/0        00:00:00 bash
  966 pts/0        00:00:01 fish
 1026 pts/0        00:00:00 ps
pi@raspberrypi ~> pmap 947
947:  -bash
00010000   872K r-x-- bash
000f9000     4K r---- bash
000fa000    20K rw--- bash
000ff000    36K rw--- [ anon ]
00533000  1088K rw--- [ anon ]
76bac000    36K r-x-- libnss_files-2.24.so
76bb5000    60K ----- libnss_files-2.24.so
76bc4000     4K r---- libnss_files-2.24.so
76bc5000     4K rw--- libnss_files-2.24.so
76bc6000    24K rw--- [ anon ]
76bcc000    36K r-x-- libnss_nis-2.24.so
76bd5000    60K ----- libnss_nis-2.24.so
76be4000     4K r---- libnss_nis-2.24.so
76be5000     4K rw--- libnss_nis-2.24.so
76bf0000    68K r-x-- libnsl-2.24.so
76bf7000    60K ----- libnsl-2.24.so
76c06000     4K r---- libnsl-2.24.so
76c07000     4K rw--- libnsl-2.24.so
76c08000     8K rw--- [ anon ]
76c0a000    24K r-x-- libnss_compat-2.24.so
76c0b000    60K ----- libnss_compat-2.24.so
76c1f000     4K r---- libnss_compat-2.24.so
```

Memory
address

size

Read/write/
execution
permission

Lib/execut
ion file

Code loaded by bash

Constant/static variable
loaded by bash

Data loaded by bash

\$pmap: memory mapping

Bash in
memory

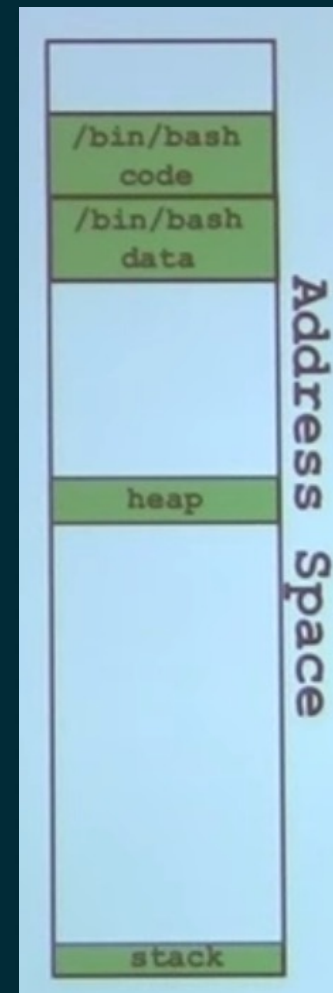
Memory
address

size

Read/write/
execution
permission

Lib/execut
ion file

```
pi@raspberrypi ~> ps
  PID TTY          TIME CMD
  947 pts/0        00:00:00 bash
  966 pts/0        00:00:01 fish
 1026 pts/0        00:00:00 ps
pi@raspberrypi ~> pmap 947
947:  -bash
00010000   872K r-x-- bash
000f9000     4K r---- bash
000fa000    20K rw--- bash
000ff000    36K rw--- [ anon ]
00533000   1088K rw--- [ anon ]
76bac000    36K r-x-- libnss_files-2.24.so
76bb5000    60K ----- libnss_files-2.24.so
76bc4000     4K r---- libnss_files-2.24.so
76bc5000     4K rw--- libnss_files-2.24.so
76bc6000    24K rw--- [ anon ]
76bcc000    36K r-x-- libnss_nis-2.24.so
76bd5000    60K ----- libnss_nis-2.24.so
76be4000     4K r---- libnss_nis-2.24.so
76be5000     4K rw--- libnss_nis-2.24.so
76bf0000    68K r-x-- libnsl-2.24.so
76bf7000    60K ----- libnsl-2.24.so
76c06000     4K r---- libnsl-2.24.so
76c07000     4K rw--- libnsl-2.24.so
76c08000     8K rw--- [ anon ]
76c0a000    24K r-x-- libnss_compat-2.24.so
76c0b000    60K ----- libnss_compat-2.24.so
76c1f000     4K r---- libnss_compat-2.24.so
```



Outline

- What is process?
 - Process vs Program
 - Linux Process Control Block
- Process related System calls
 - Fork
 - Exec
 - Wait
- Process on Distributed OSes



Test: create a process and show the following information of it

<https://github.com/kevinsuo/CS7172/blob/master/sleep.c>

Download: wget Raw-file-URL

Compile: gcc [file-name].c -o [file-name].o

Run: ./[file-name].o

- What is the process ID?
- What is the process state?
- What is the scheduling policy for the process?

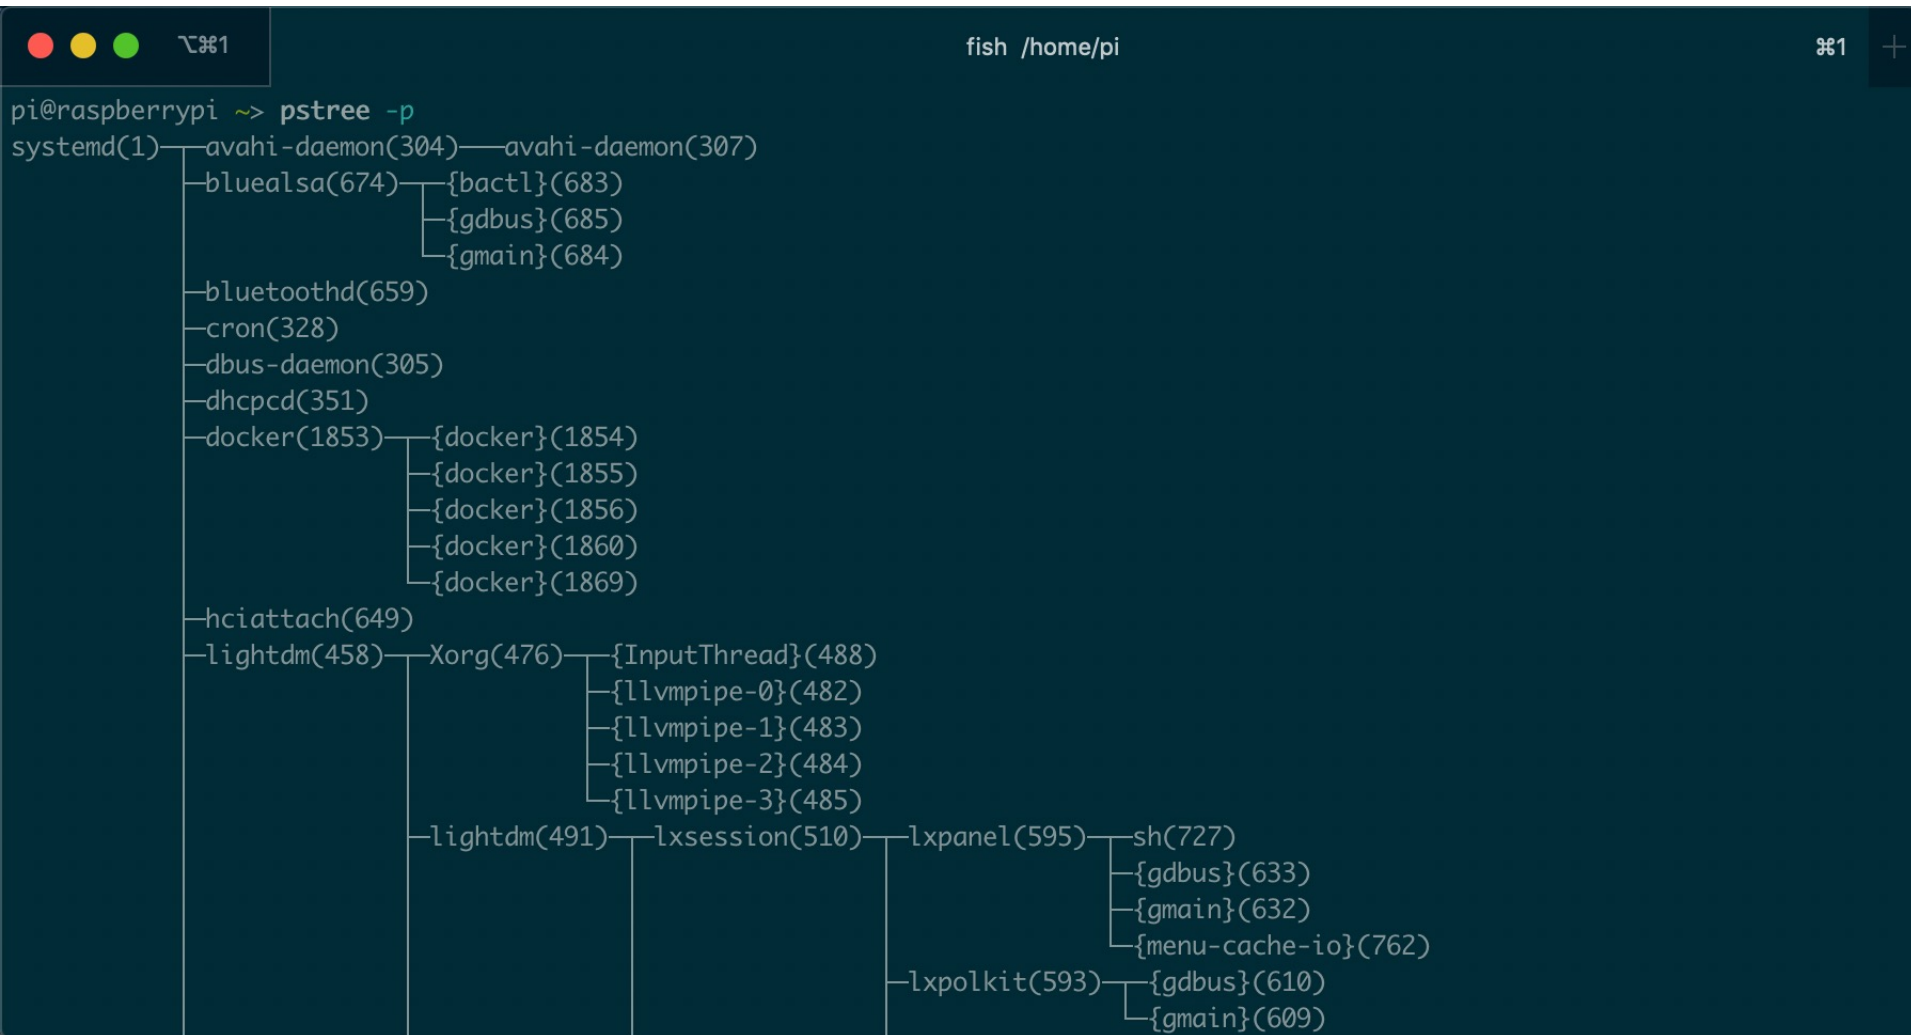
SCHED_OTHER	min/max priority : 0/0
SCHED_FIFO	min/max priority : 1/99
SCHED_RR	min/max priority : 1/99
SCHED_BATCH	min/max priority : 0/0

- Show all files the process it is using.
- Show all memory the process it is using.



Where do processes come from?

- Process creation always uses `fork()` system call



Where do processes come from? First process in the kernel

```
asmlinkage __visible void __init start_kernel(void)
{
    char *command_line;
    char *after_dashes;

    set_task_stack_end_magic(&init_task);
    smp_setup_processor_id();
    debug_objects_early_init();

    cgroup_init_early();

    local_irq_disable();
    early_boot_irqs_disabled = true;

    /*
     * Interrupts are still disabled. Do necessary setups, then
     * enable them.
     */
    boot_cpu_init();
    page_address_init();
    pr_notice("%s", linux_banner);
    setup_arch(&command_line);
```

<https://elixir.bootlin.com/linux/v5.4/source/init/main.c#L580>

- The start of linux kernel begins from `start_kernel()` function, it is equal to the main function of kernel
- `set_task_stack_end_magic()` creates the first process in the OS
- The first process is the only one which is not created by `fork` function



Fork() system call

- Process creation always uses fork() system call
- When?
 - User runs a program at command line
 - ▶ `$/test.o`
 - OS creates a process to provide a service
 - ▶ Timer, networking, load-balance, daemon, etc.
 - One process starts another process
 - ▶ Parents and child process



Fork() system call

```
1. vim /home/pi/Downloads (ssh)
vim /home/pi/Dow... %1

#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main(void) {
    int pid;
    pid = fork();

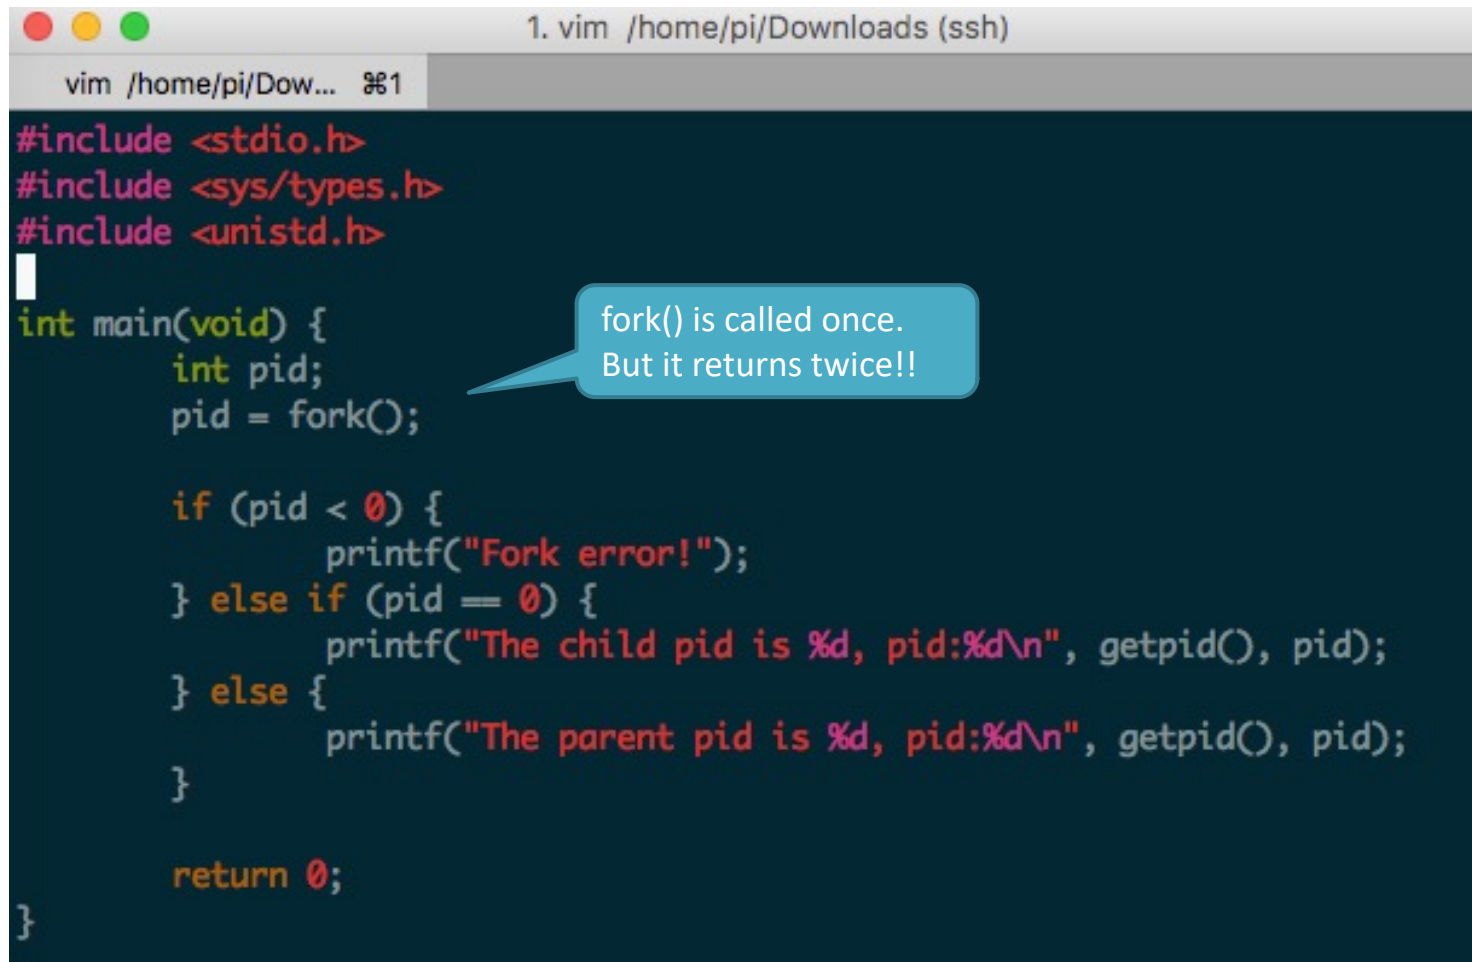
    if (pid < 0) {
        printf("Fork error!");
    } else if (pid == 0) {
        printf("The child pid is %d, pid: %d\n", getpid(), pid);
    } else {
        printf("The parent pid is %d, pid: %d\n", getpid(), pid);
    }

    return 0;
}
```

Fork() system call

- fork() is called once. But it returns twice!!
 - Once in the parent (return child id > 0)
 - Once in the child (return 0)

Fork() system call



```
1. vim /home/pi/Downloads (ssh)
vim /home/pi/Dow... %1

#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main(void) {
    int pid;
    pid = fork();

    if (pid < 0) {
        printf("Fork error!");
    } else if (pid == 0) {
        printf("The child pid is %d, pid:%d\n", getpid(), pid);
    } else {
        printf("The parent pid is %d, pid:%d\n", getpid(), pid);
    }

    return 0;
}
```

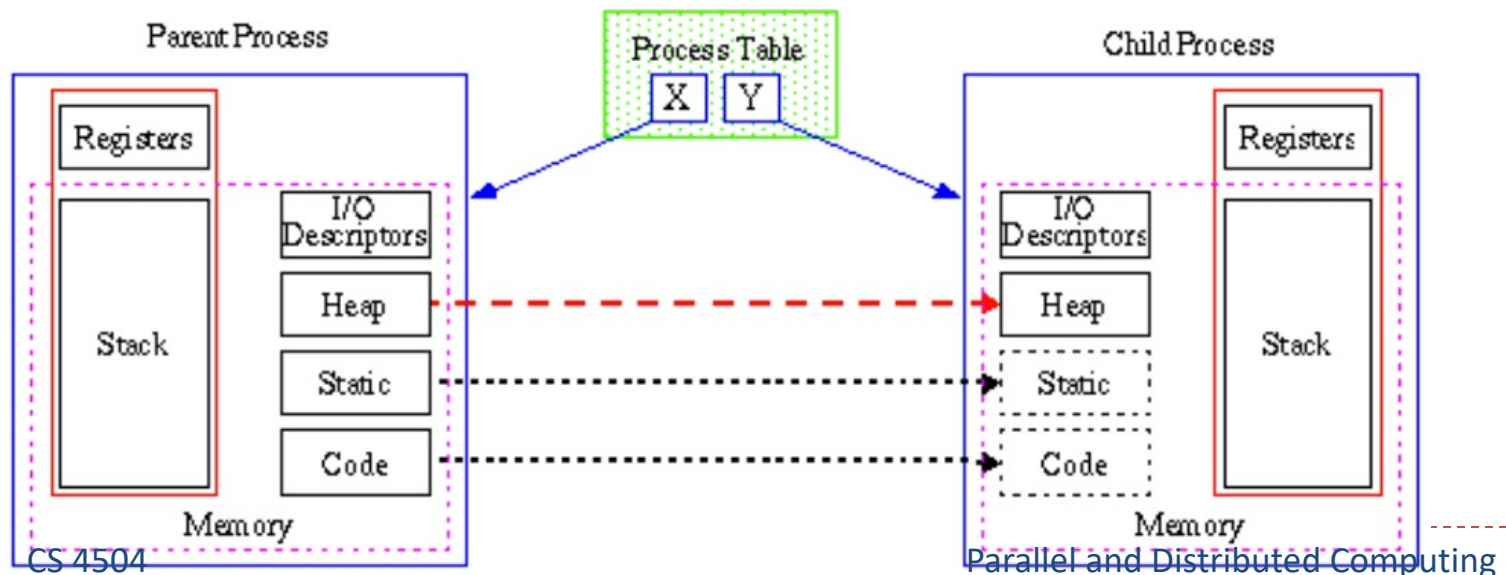
```
pi@raspberrypi ~/Downloads> ./test.o
The parent pid is 2510, pid:2511
The child pid is 2511, pid:0
```

Fork() system call

- fork() is the UNIX system call that creates a new process.
- fork() creates a new process that is a **copy** of the calling process.
- After fork () we refer to the caller as the **parent** and the newly-created process as the **child**. They have a special relationship and special responsibilities.

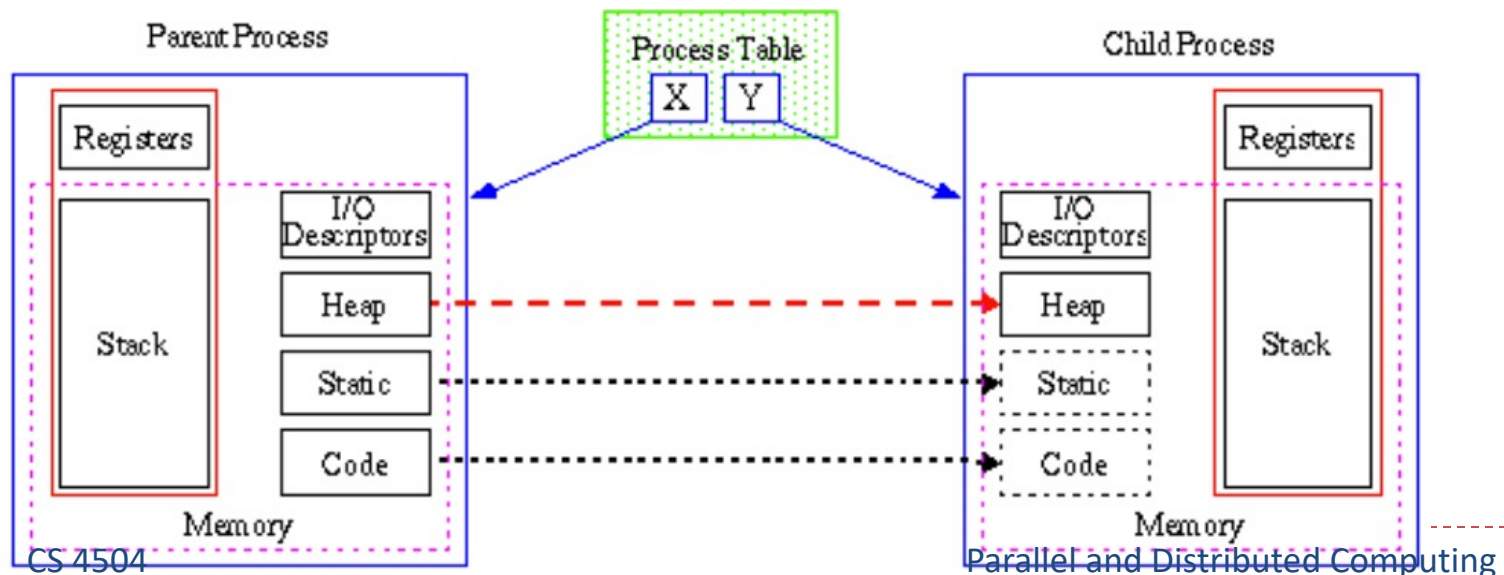
Parent process and child process

- When a parent process uses `fork()` to create a child process, the two processes have
 - the **same** program text.
 - but **separate** copies of the data, stack, and heap segments.



Parent process and child process

- The child's stack, data, and heap segments are **initially** exact **duplicates** of the corresponding parts the parent's memory.
- After the fork(), each process can modify the variables in its **own** data, stack, and heap segments without affecting the other process.



Fork() example

```
1. vim /home/pi/Downlo
vim /home/pi/Dow... %1
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main()
{
    fork();

    printf("Hello world!\n");

    return 0;
}

pi@raspberrypi ~/Downloads> ./a.o
Hello world!
Hello world!
```

```
1. vim /home/pi/Dowr
vim /home/pi/Dow... %1
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main()
{
    fork();

    fork();

    fork();

    printf("Hello world!\n");

    return 0;
}

pi@raspberrypi ~/Downloads> ./a.o
Hello world!
Hello world!
Hello world!
Hello world!
Hello world!
Hello world!
Hello world!
Hello world!
```

Fork() example

- How many a it will output?

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main(){
    int i;
    pid_t pid;
    for(i = 0; i < 2; i++){
        printf("a\n");
        pid = fork();
    }
    return 0;
}
```

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main(){
    int i;
    pid_t pid;
    for(i = 0; i < 2; i++){
        printf("a\n");
        pid = fork();
    }
    return 0;
```

```
}
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

Fork() example

- How many a it will output?

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main(){
    int i;
    pid_t pid;
    for(i = 0; i < 2; i++){
        printf("a\n");
        pid = fork();
    }
    return 0;
}
```

i=0	Main:	a Create a process named 111
i=1	Main:	a Create a process named 222
	111:	a Create a process named 333



Fork() example

- How many a it will output?

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main(){
    int i;
    pid_t pid;
    for(i = 0; i < 2; i++){
        printf("a");
        pid = fork();
    }
    return 0;
}
```



The image shows a Vim editor window with a dark blue background. The title bar at the top reads "vim /home/ubuntu". Below the title bar, there are two tabs: "fish /Users/ksuo/OneDrive - Kennesaw State Univer..." and "vim /home/ubuntu (ssh)". The main editing area contains the following C code:

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main(){
    int i;
    pid_t pid;
    for(i = 0; i < 2; i++){
        printf("a");
        pid = fork();
    }
    return 0;
}
```

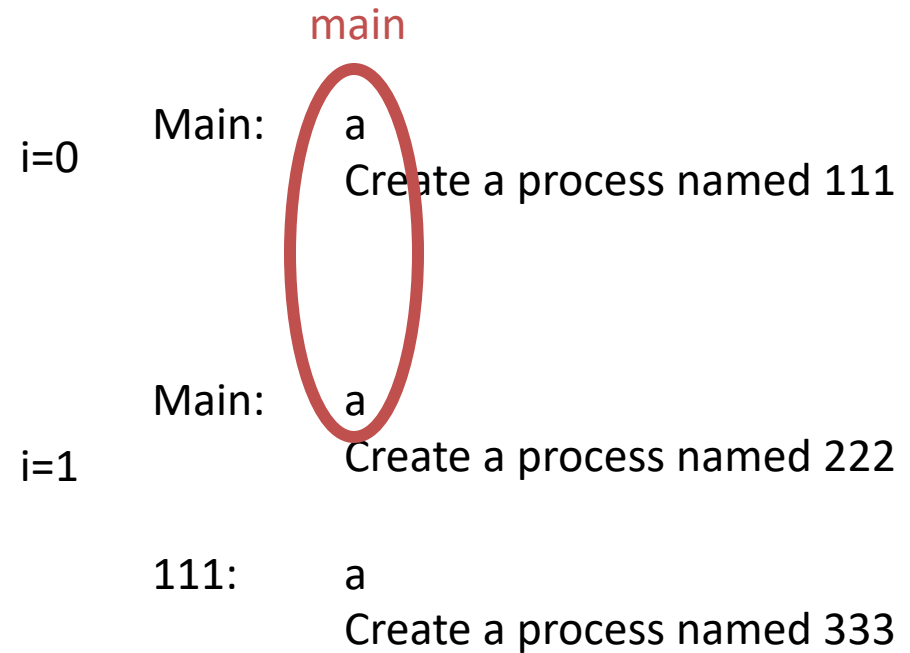
The cursor is positioned at the end of the closing brace of the `main` function. The status bar at the bottom left shows "test.c" 13L, 218C. The status bar at the bottom right shows "13,3" and "All".

Fork() example

- Printf (without /n) will not flush the buffer until the program finished

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main(){
    int i;
    pid_t pid;
    for(i = 0; i < 2; i++){
        printf("a");
        pid = fork();
    }
    return 0;
}
```

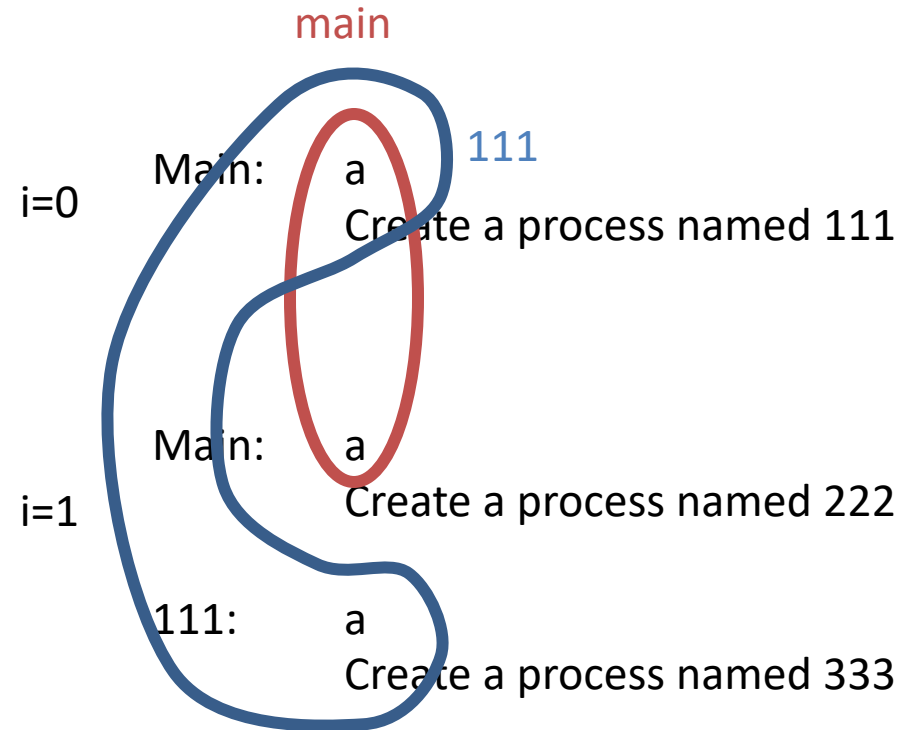


Fork() example

- Printf (without /n) will not flush the buffer until the program finished

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main(){
    int i;
    pid_t pid;
    for(i = 0; i < 2; i++){
        printf("a");
        pid = fork();
    }
    return 0;
}
```

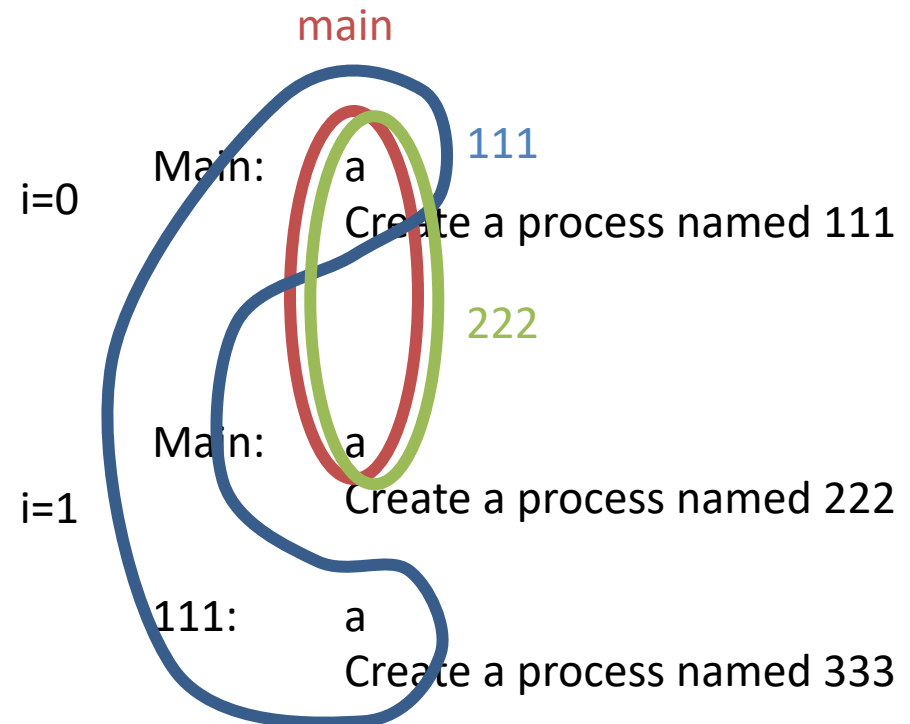


Fork() example

- Printf (without /n) will not flush the buffer until the program finished

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main(){
    int i;
    pid_t pid;
    for(i = 0; i < 2; i++){
        printf("a");
        pid = fork();
    }
    return 0;
}
```

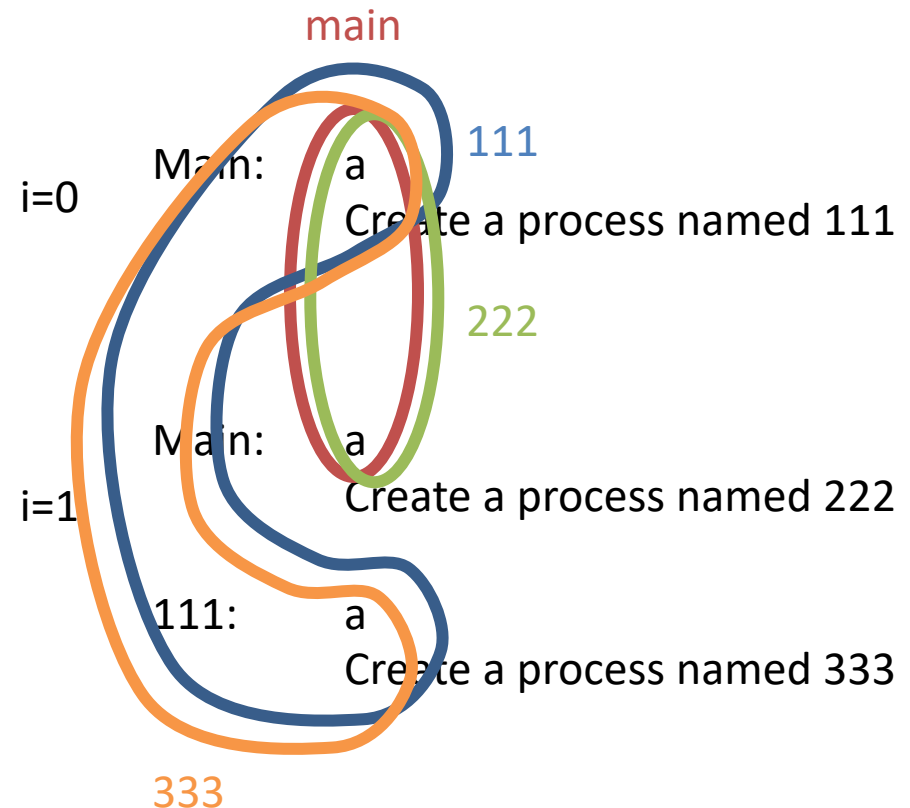


Fork() example

- Printf (without /n) will not flush the buffer until the program finished

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main(){
    int i;
    pid_t pid;
    for(i = 0; i < 2; i++){
        printf("a");
        pid = fork();
    }
    return 0;
}
```



A fork() bomb

- What does this code do?

```
• • • 1
while(1) {
    fork();
}
```



Agent smith

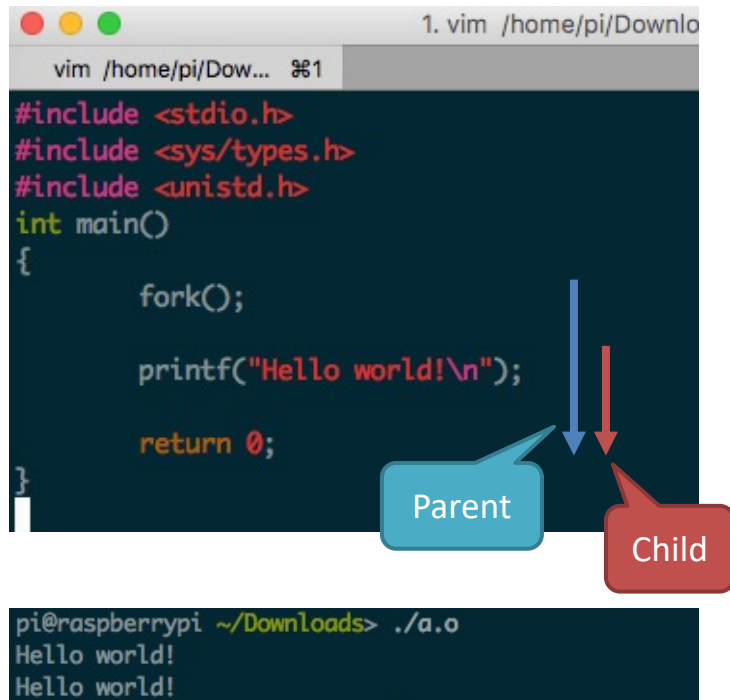
Exec() system call

- Replaces current process image with new program image.
- Exec system call is a collection of functions and in C programming language, the standard names for these functions are as follows:
 - `int execl(const char* path, const char* arg, ...)`
 - `int execlp(const char* file, const char* arg, ...)`
 - `int execlxe(const char* path, const char* arg, ..., char* const envp[])`
 - `int execv(const char* path, const char* argv[])`
 - `int execvp(const char* file, const char* argv[])`
 - `int execvpe(const char* file, const char* argv[], char *const envp[])`



Exec() system call

- Replaces current process image with new program image.



```
1. vim /home/pi/Downlo
vim /home/pi/Dow... %1
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main()
{
    fork();

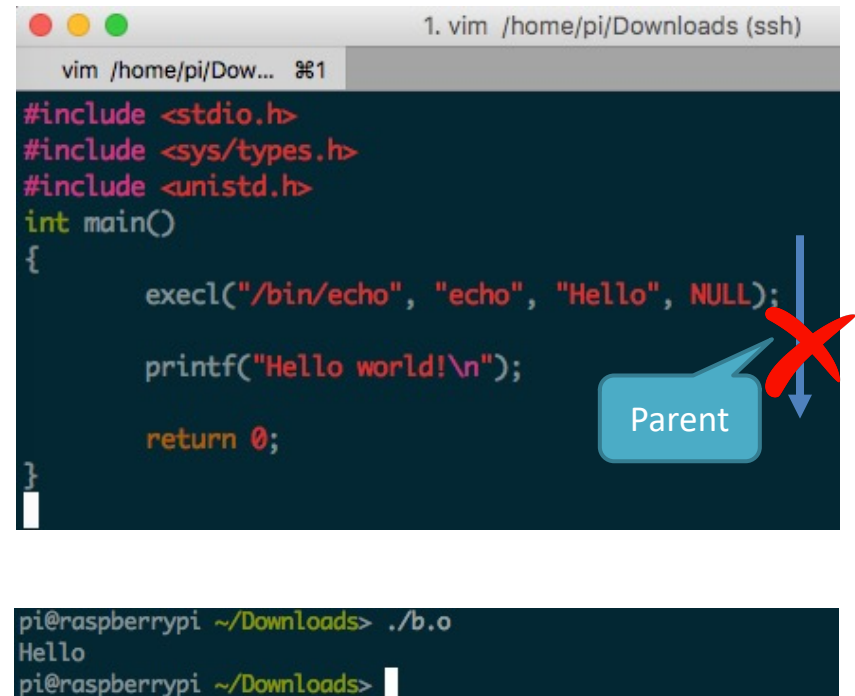
    printf("Hello world!\n");

    return 0;
}
```

Parent

Child

```
pi@raspberrypi ~/Downloads> ./a.o
Hello world!
Hello world!
```



```
1. vim /home/pi/Downloads (ssh)
vim /home/pi/Dow... %1
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main()
{
    execl("/bin/echo", "echo", "Hello", NULL);

    printf("Hello world!\n");

    return 0;
}
```

Parent

```
pi@raspberrypi ~/Downloads> ./b.o
Hello
pi@raspberrypi ~/Downloads>
```

Wait() system call

<https://github.com/kevinsuo/CS7172/blob/master/wait.c>

- Helps the parent process
 - to know when a child completes
 - to check the return status of child

```
#include<stdio.h>
#include<stdlib.h>
#include<sys/wait.h>
#include<unistd.h>

int main()
{
    pid_t cpid;
    if (fork()== 0) {
        printf("Child pid = %d\n", getpid());
    } else {
        // cpid = wait(NULL); /* returns a process ID of dead children */
        printf("Parent pid = %d\n", getpid());
    }

    return 0;
}
```

```
pi@raspberrypi ~/Downloads> ./wait.o
Parent pid = 3425
Child pid = 3426
```

```
#include<stdio.h>
#include<stdlib.h>
#include<sys/wait.h>
#include<unistd.h>

int main()
{
    pid_t cpid;
    if (fork()== 0) {
        printf("Child pid = %d\n", getpid());
    } else {
        cpid = wait(NULL); /* returns a process ID of dead children */
        printf("Parent pid = %d\n", getpid());
    }

    return 0;
}
```

```
pi@raspberrypi ~/Downloads> ./wait.o
Child pid = 3395
Parent pid = 3394
```



Few other useful syscalls

- `sleep(seconds)`
 - suspend execution for certain time
- `exit(status)`
 - Exit the program.
 - Status is retrieved by the parent using `wait()`.
 - 0 for normal status, non-zero for error
- `kill(pid_t pid, int sig)`
 - Kill certain process



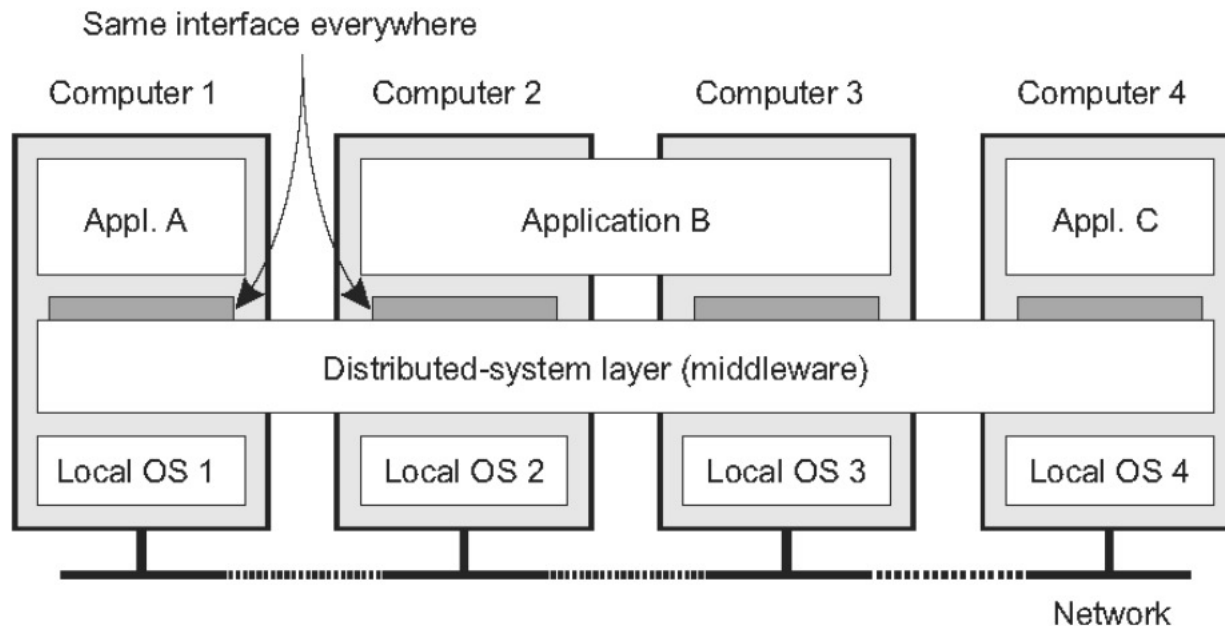
Outline

- What is process?
 - Process vs Program
 - Linux Process Control Block
- Process related System calls
 - Fork
 - Exec
 - Wait
- Process on Distributed OSes



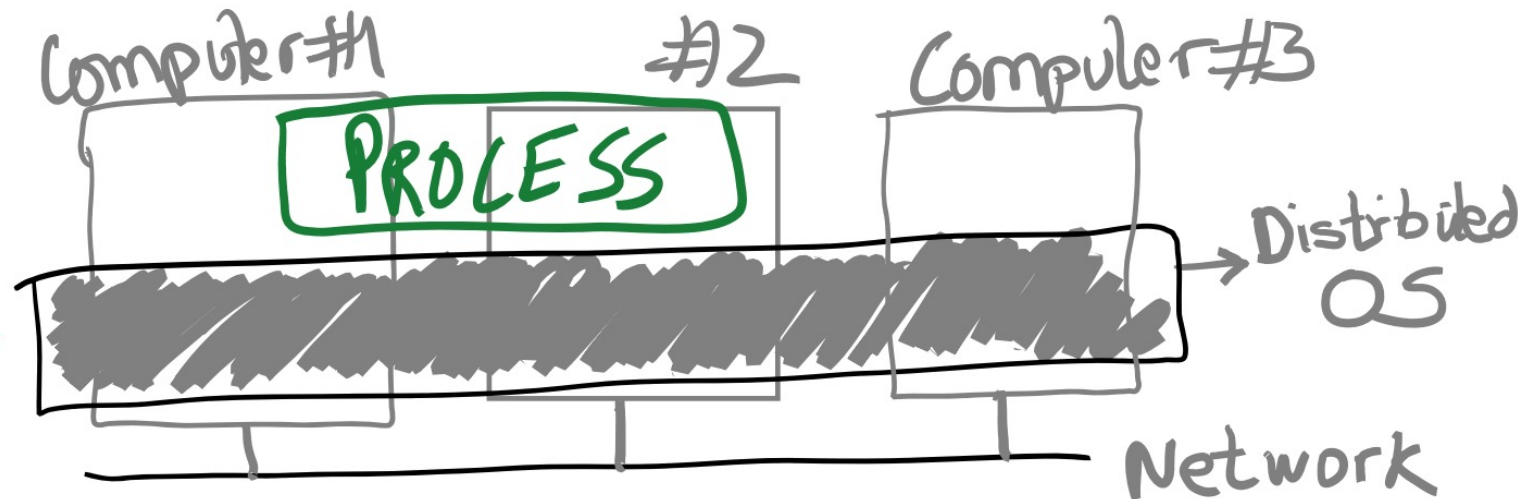
The OS of Distributed Systems

- Commonly used components and functions for distributed applications



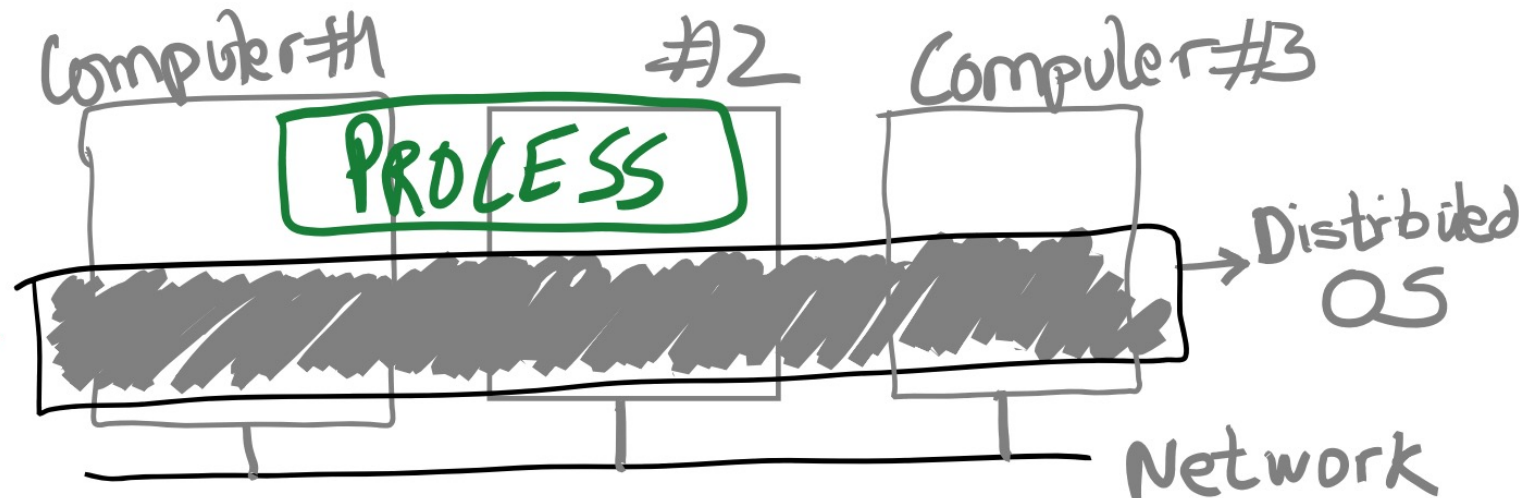
The OS of Distributed Systems

- An OS that spans multiple computers
- Same OS services, functionality, and abstractions as single-machine OS



Discussion

- What could be the challenges for distributed OS?



Distributed OS Challenges

- Providing the process abstraction and resource virtualization is hard
- Resource virtualization must be transparent
 - But in distributed settings, there's always a distinction between local and remote resources
- In a single-machine OS, processes don't care where their resources are coming from:
 - Which CPU cores, when they are scheduled, which physical memory pages they use, etc.
- In fact, providing abstract, virtual resources is one of the main OS services



Transparency Issues In Distributed OS

PROCESS

Process state:

- Code segment
- Memory pages
- Files
- Sockets
- Security permissions

- Where does code run?
- Which memory is used?
Local vs. remote
- How are files accessed?

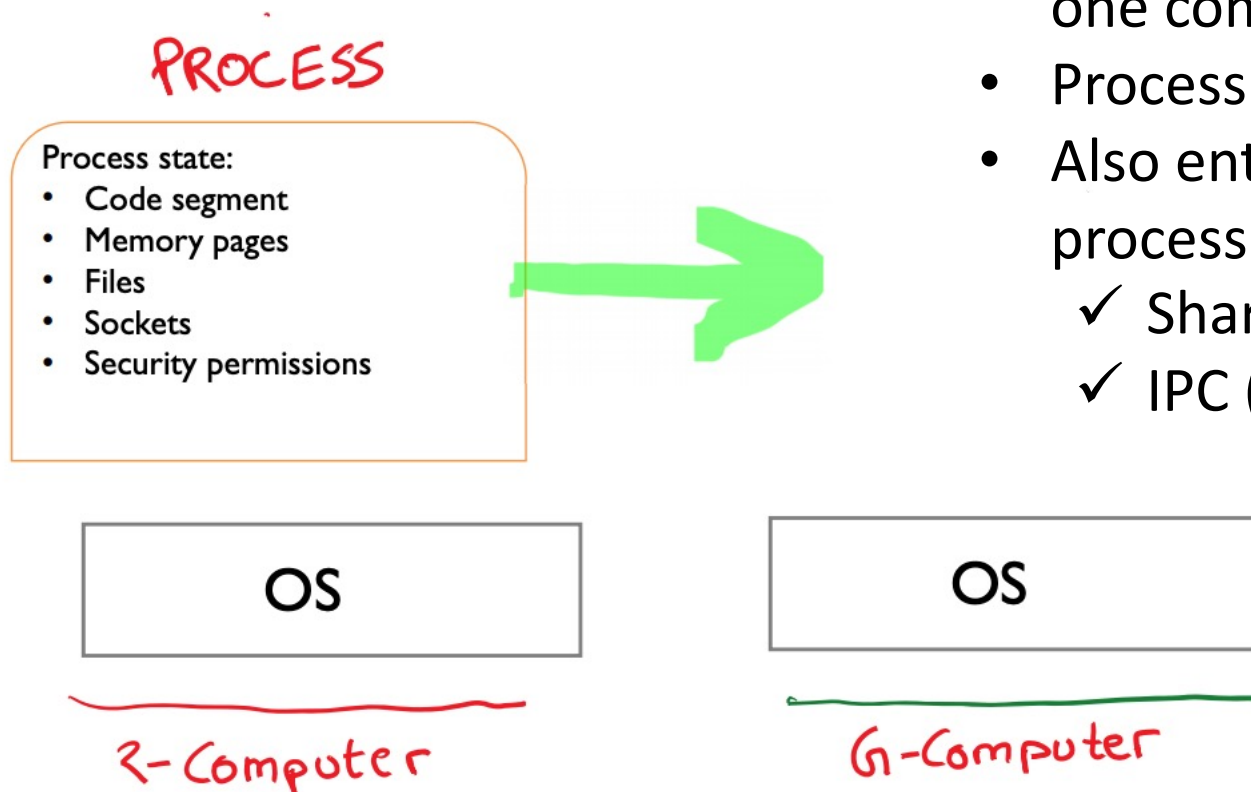
Distributed OS

2-Computer

G-Computer

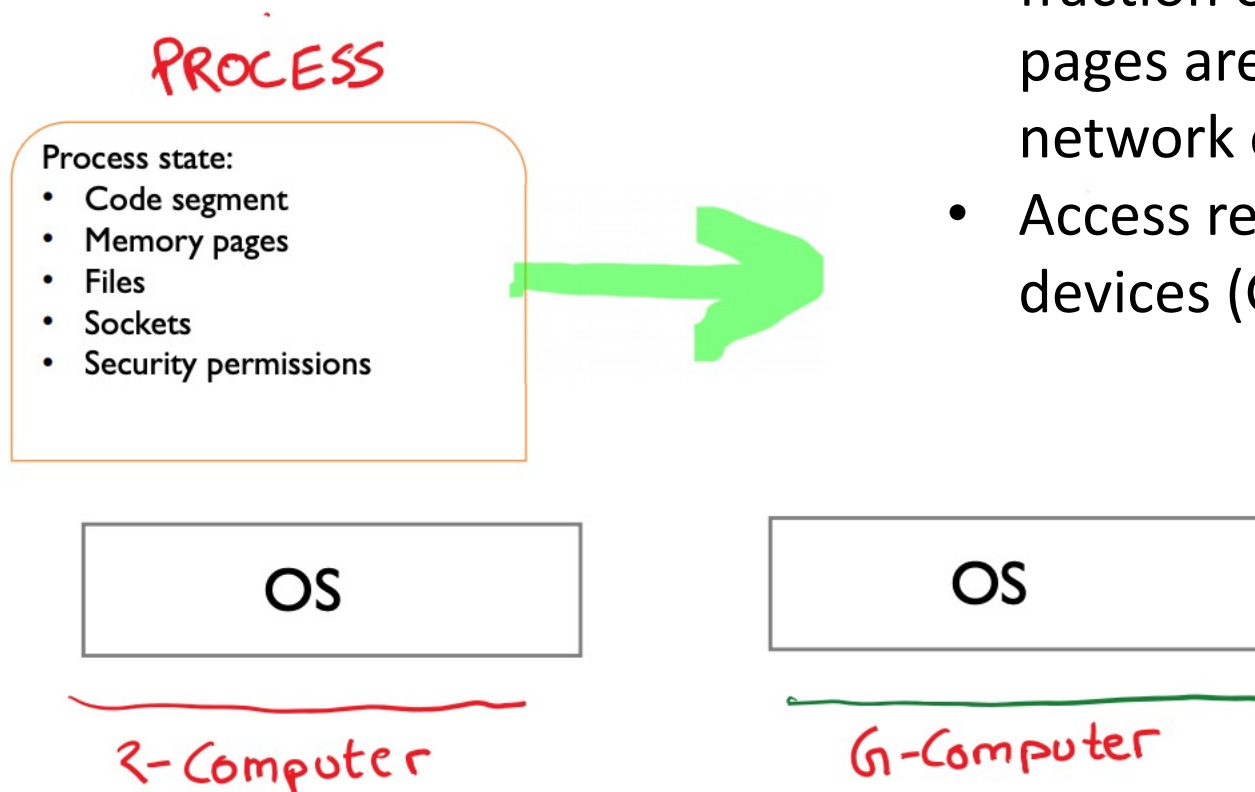
Process Migration

- Move all process state from one computer to another
- Process state can be vast
- Also entangled with other process states
 - ✓ Shared files?
 - ✓ IPC (pipes etc)



Process Migration

- Migrate some state? Other state, if required, is accessed over the network?
- Example: migrate only fraction of pages. Other pages are copied over the network on access?
- Access remote hardware devices (GPUs)?



Conclusion

- What is process?
 - Process vs Program
 - Linux Process Control Block
- Process related System calls
 - Fork, etc.
- Process on Distributed OSes

