

# CS 3502

# Operating Systems

## Project 1

**Kun Suo**

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

# Command: fish

- <https://fishshell.com/>

## Autosuggestions

fish suggests commands as you type based on history and completions, just like a web browser. Watch out, Netscape Navigator 4.0!

```
tion to do it  
nber it and s  
it.  
> sh -l den
```

## Sane Scripting

fish is fully scriptable, and its syntax is simple, clean, and consistent. You'll never write esac again.

```
function de  
switch  
cas  
cas
```

## Man Page Completions

Other shells support programmable completions, but only fish generates them automatically by parsing your installed man pages.

```
demo > fish_  
Parsing man  
demo > rwsno  
-P (print p  
-Z  
-j (
```

## Glorious VGA Color

fish supports 24 bit true color, the state of the art in terminal technology. Behold the monospaced rainbow.

```
mo > echo "Lo  
mo > █
```

## Web Based configuration

For those lucky few with a graphical computer, you can set your colors and view functions, variables, and history all from a web page.

```
Current  
/bright/vi  
echo 'Erro  
# This is  
Th is an
```

## Works Out Of The Box


fish will delight you with features like tab completions and syntax highlighting that just work, with nothing new to learn or configure.

```
git checkout  
(Tag) 1.21  
(Tag) 1.21  
(Tag) 1.21  
(Tag) 1.21  
(Tag) 1.21
```

# Project 1: No IDE, please use vim

- Open a file

- `$ vim test.c`



```
fish /home/ksuo@ksuo-VirtualBox ~/D/l/i/linux> vim syscalls.h
```

- Close a file

- Inside vim (after opening), press

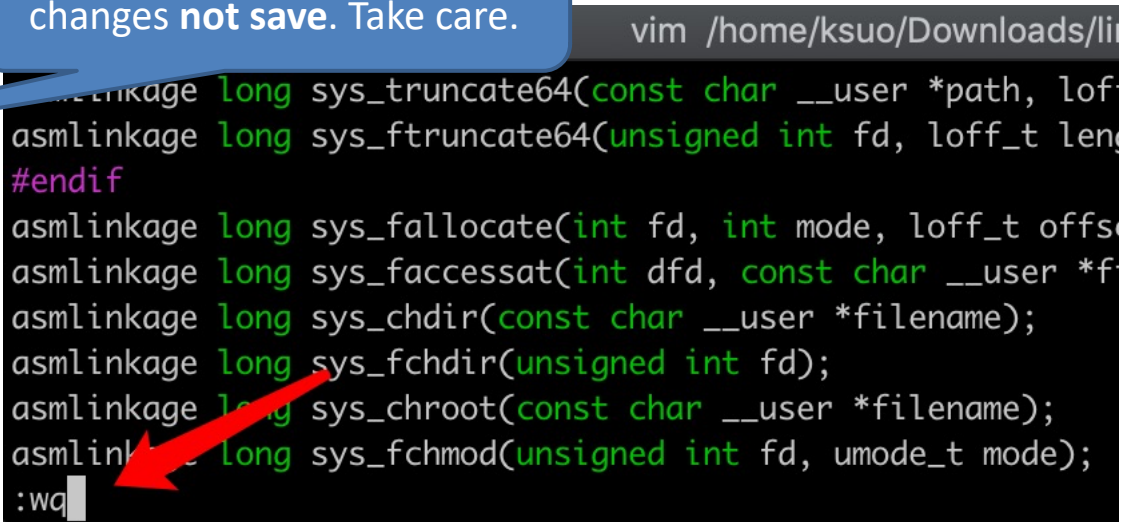
`:q!`

- Inside vim (after opening), press

`:wq`, changes saved

colon mark means menu.

Exclamation mark means changes **not** save. Take care.



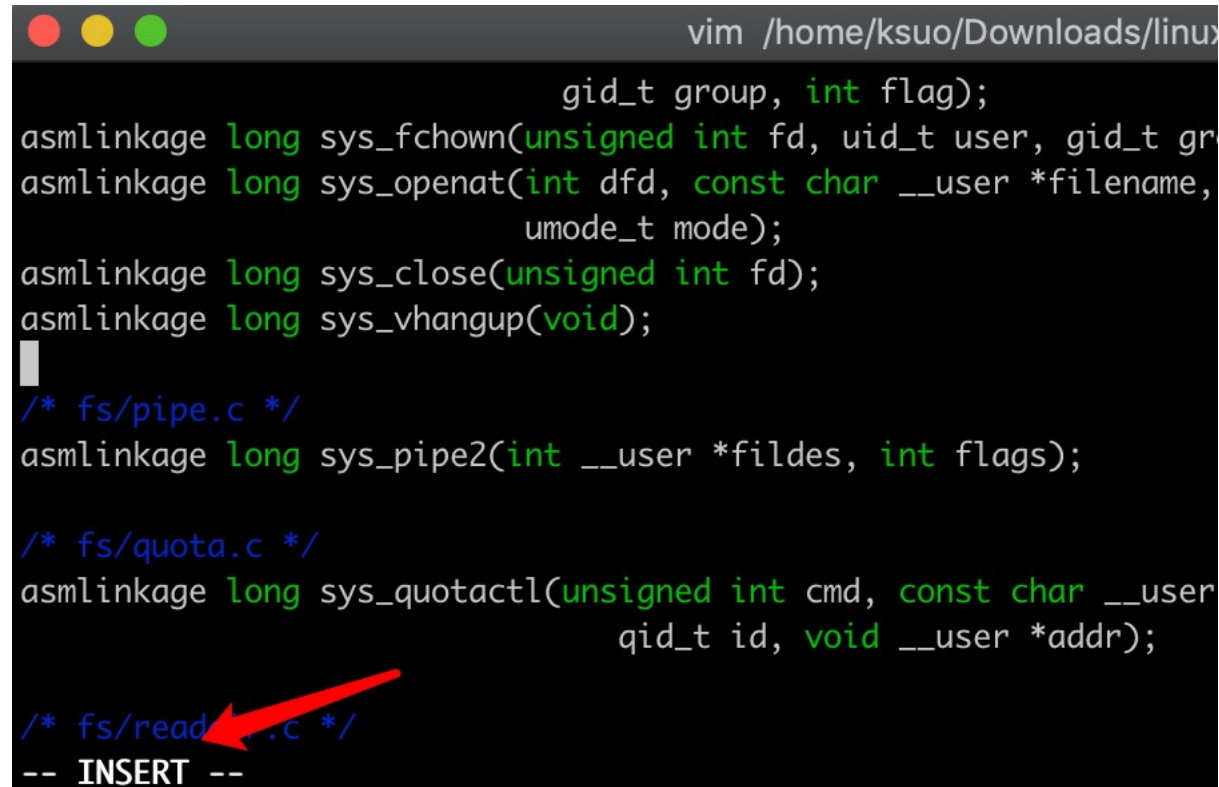
```
vim /home/ksuo/Downloads/li
asmlinkage long sys_truncate64(const char __user *path, loff_t len
asmlinkage long sys_ftruncate64(unsigned int fd, loff_t len
#endif
asmlinkage long sys_fallocate(int fd, int mode, loff_t offs
asmlinkage long sys_faccessat(int dfd, const char __user *f
asmlinkage long sys_chdir(const char __user *filename);
asmlinkage long sys_fchdir(unsigned int fd);
asmlinkage long sys_chroot(const char __user *filename);
asmlinkage long sys_fchmod(unsigned int fd, umode_t mode);
:wq
```



# Project 1: No IDE, please use vim

---

- Edit a file
  - Open a file
  - Press i (means enter the insert mode). then input your code
  - Press ESC to exit the insert mode
  - Close a file (:wq)



```
vim /home/ksuo/Downloads/linux
                                gid_t group, int flag);
asmlinkage long sys_fchown(unsigned int fd, uid_t user, gid_t gr
asmlinkage long sys_openat(int dfd, const char __user *filename,
                                umode_t mode);
asmlinkage long sys_close(unsigned int fd);
asmlinkage long sys_vhangup(void);
/* fs/pipe.c */
asmlinkage long sys_pipe2(int __user *fildes, int flags);

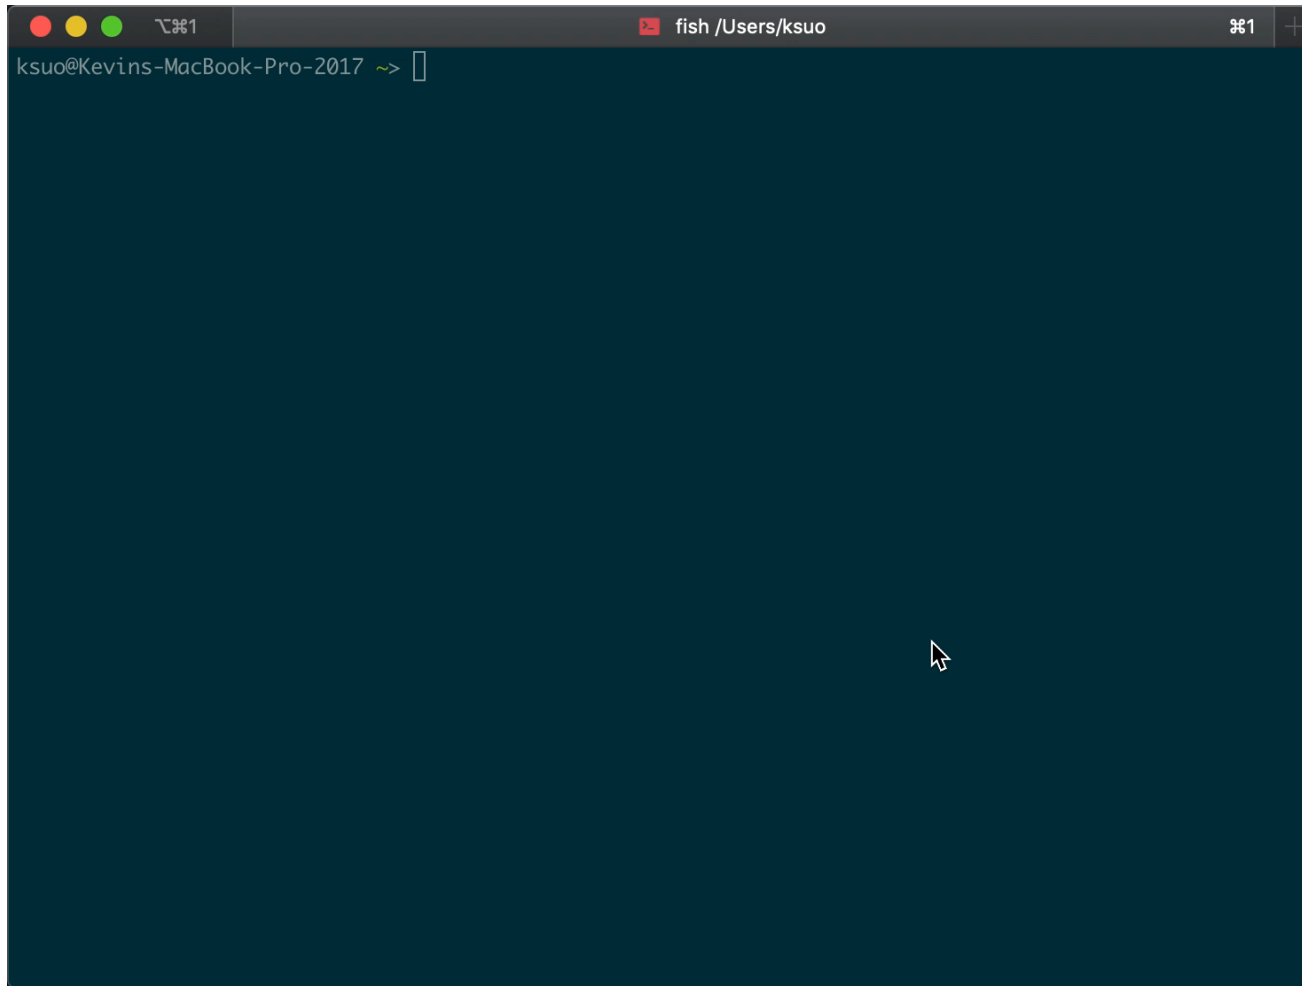
/* fs/quota.c */
asmlinkage long sys_quotactl(unsigned int cmd, const char __user
                                gid_t id, void __user *addr);

/* fs/read.c */
-- INSERT --
```

# Project 1: No IDE, please use vim

---

<https://youtu.be/rS0XZpLtVlo>



# Edit a file with vim

---

- step 1: \$ **vim** file
- step 2: press **i**, enter insert mode; move the cursor to position and edit the context
- step 3: after editing, press **ESC** to exit the insert mode to normal mode
- step 4: press **:wq** to save what you edit and quit. If you do not want to save, press **:q!**



# More about vim

---

- A quick start guide for beginners to the Vim text editor
  - <https://eastmanreference.com/a-quick-start-guide-for-beginners-to-the-vim-text-editor>
- Vim basics:
  - <https://www.howtoforge.com/vim-basics>
- Learn the Basic Vim Commands [Beginners Guide]
  - [https://www.youtube.com/watch?time\\_continue=265&v=ZEGqkam-3lc](https://www.youtube.com/watch?time_continue=265&v=ZEGqkam-3lc)



# Outline

---

- **Part A: create a VM and compile your kernel**
- Part B: add a new system call into the Linux kernel





# How to build one ubuntu VM?

---

- HostOS

Windows 10 (x86):

<https://www.youtube.com/watch?v=QbmRXJJKsvs>

MacOS (x86):

<https://www.youtube.com/watch?v=GDoCrfPma2k&t=321s>

MacOS (arm):

<https://youtu.be/O19mv1pe76M?si=4cYayFigPNoHoY1w>



# Tips

---

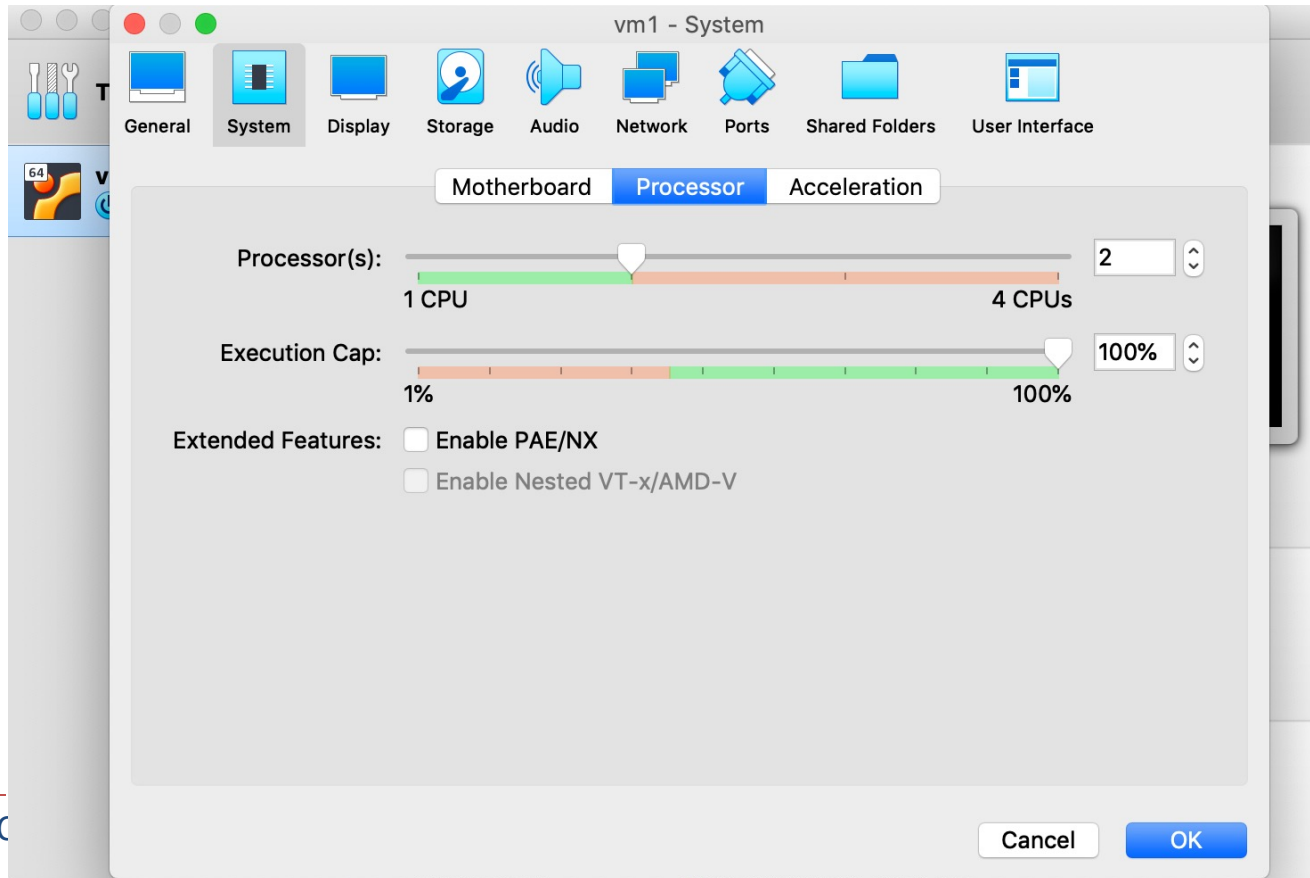
- Compile Linux kernel takes half to two hours, depending on machine speed
- To save the time, use `$ make localmodconfig`
- Run many commands in one:

`$ sudo make; sudo make modules; sudo make modules_install; sudo make install`

```
ksuo@ksuo-VirtualBox ~/D/linux-5.1>  
sudo make; sudo make modules; sudo make modules_install; sudo make install
```

# Tips

- Set up more CPUs for VM
- Make `-j N`, to accelerate compiling (N is your CPU number)



# Where is my kernel?

- `$ ls /boot/`

Initial ramdisk: loading a temporary root file system into memory. Used for startup.

Linux executable kernel image

initrd.img-5.0.0-23-generic  
initrd.img-5.0.0-25-generic  
initrd.img-5.1.0

vmlinuz-5.0.0-23-generic  
vmlinuz-5.0.0-25-generic  
vmlinuz-5.1.0

# Which kernel to boot if there are many?

If you are using Ubuntu: change the grub configuration file:

```
$ sudo vim /etc/default/grub
```

The OS boots by using the first kernel by default. You have 10 seconds to choose.

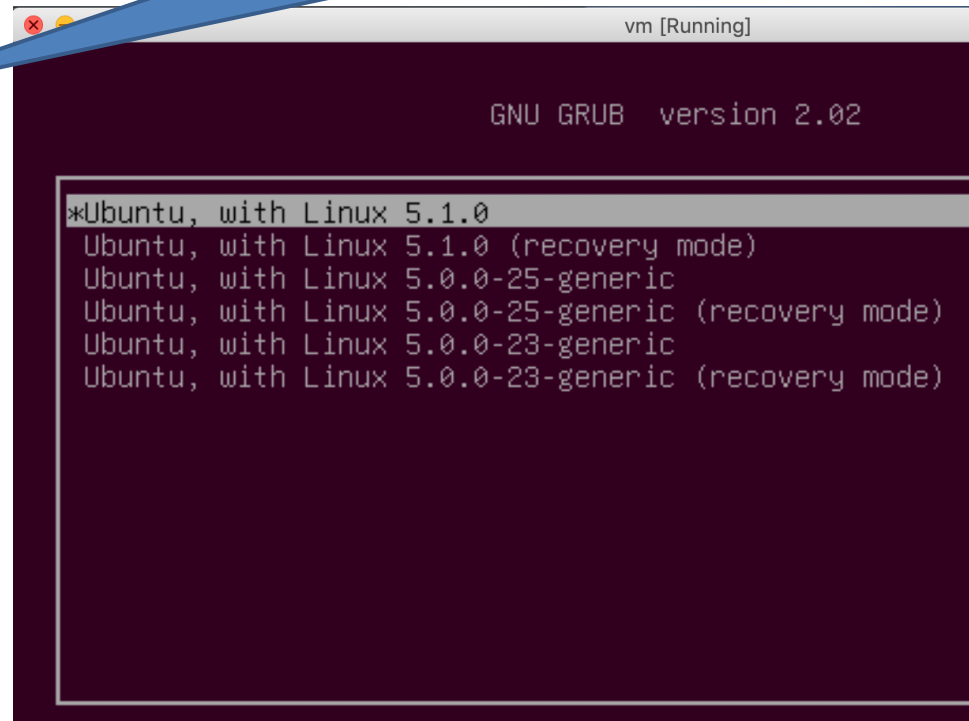
Make the following changes:

```
GRUB_DEFAULT=0
```

```
GRUB_TIMEOUT=10
```

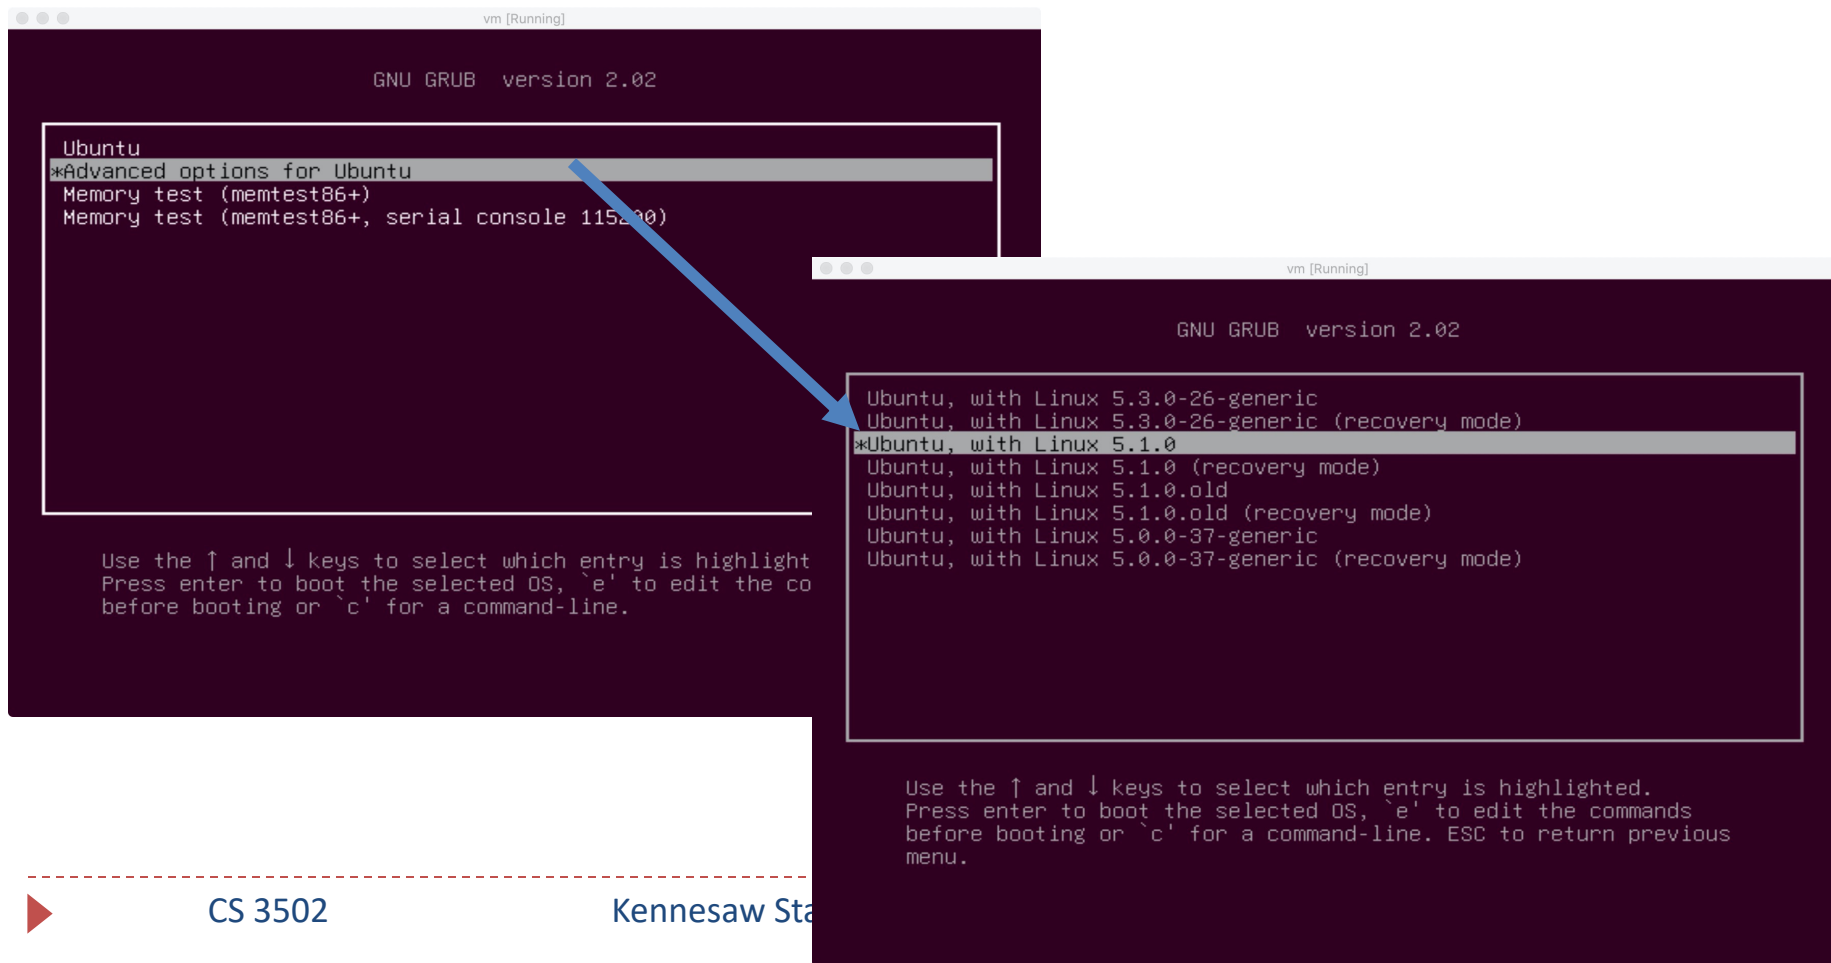
Then, update the grub entry:

```
$ sudo update-grub2
```



# Which kernel to boot if there are many?

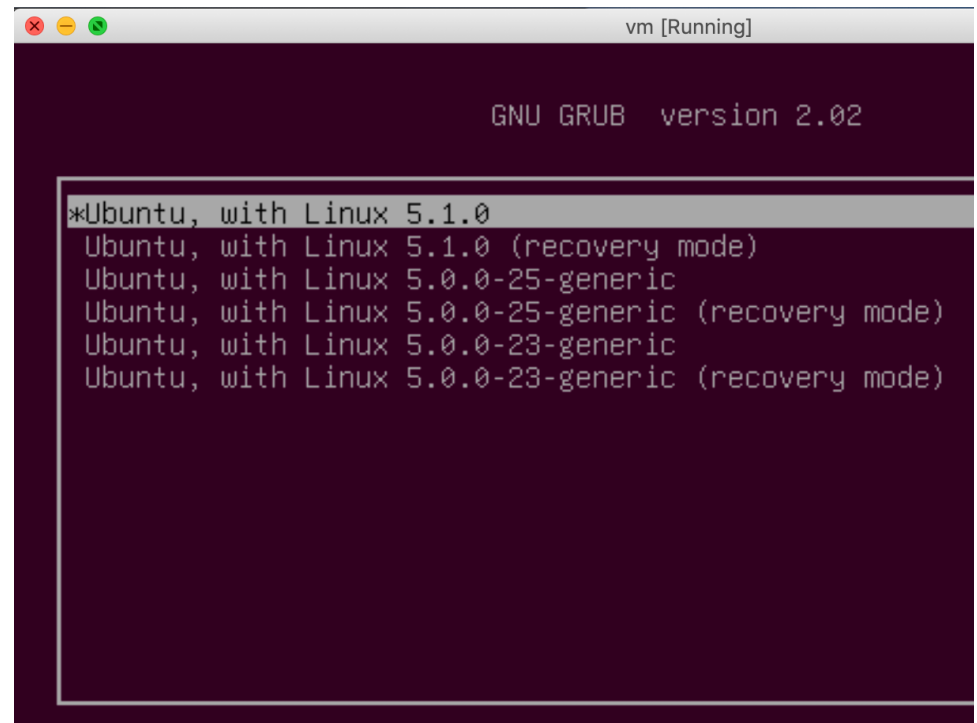
Immediately after the BIOS/UEFI splash screen during boot, with BIOS, quickly press and hold the **Shift** key, which will bring up the GNU GRUB menu. (If you see the Ubuntu logo, you've missed the point where you can enter the GRUB menu.)



# What if my kernel crashed?

---

- Your kernel could crash because you might bring in some kernel bugs
- In the menu, choose the old kernel to boot the system
- Fix your bug in the source code
- Compile and reboot



# Compile the kernel

---

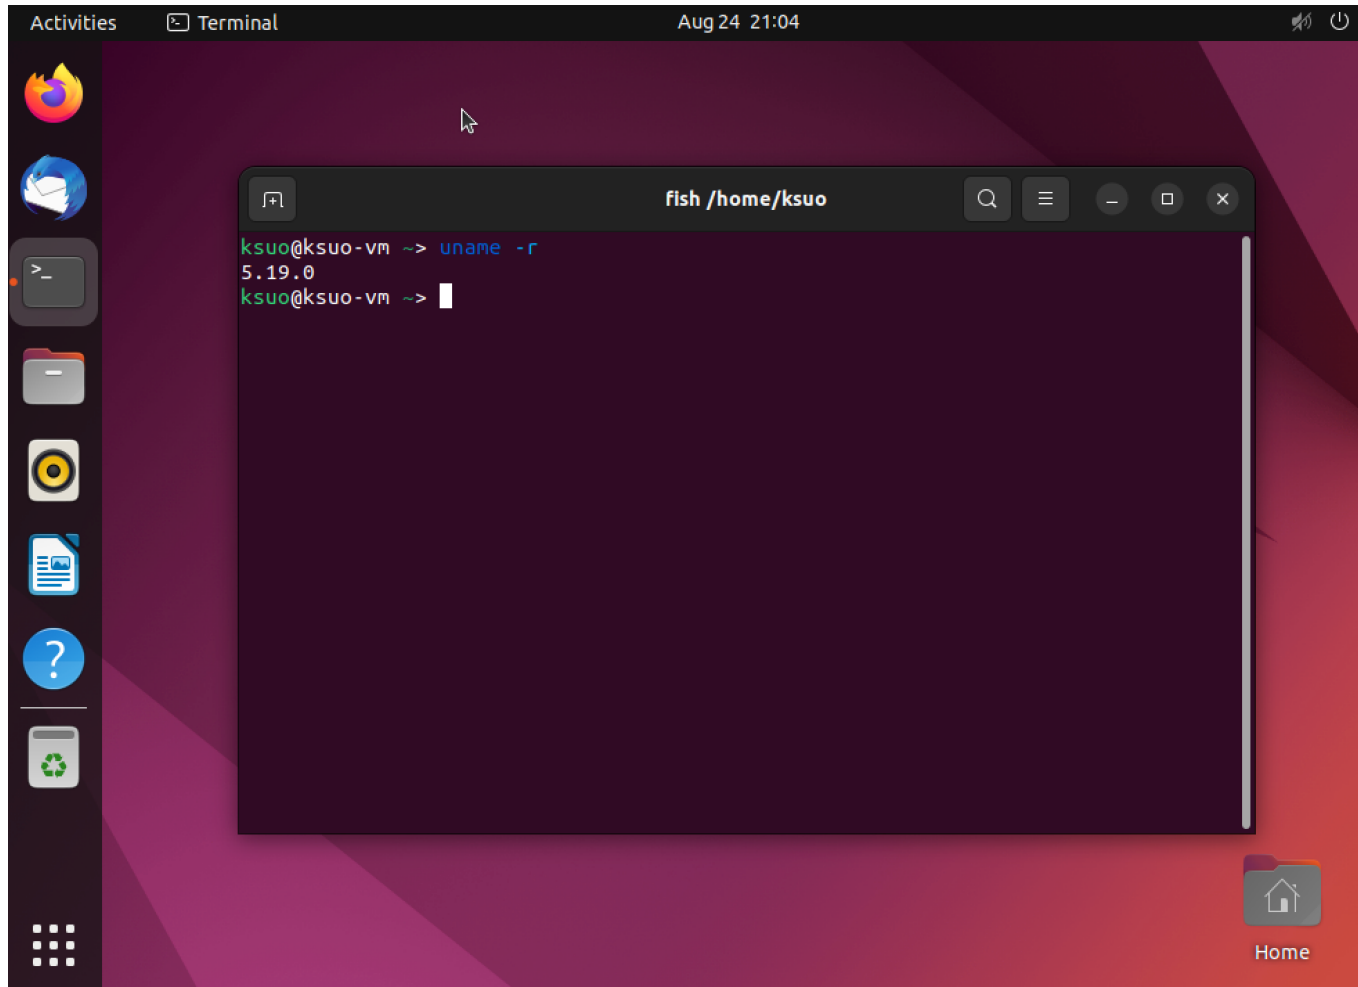
1. download the kernel source code
2. unzip the file
3. use *\$ make localmodconfig* to create a config file
4. use *\$ make/make modules* to compile the code
5. use *\$ make modules\_install/make install* to install the kernel
6. reboot the VM





# Screenshot of my new kernel

---



# Outline

---

- Part A: create a VM and compile your kernel
- **Part B: add a new system call into the Linux kernel**



# Define your system call

---

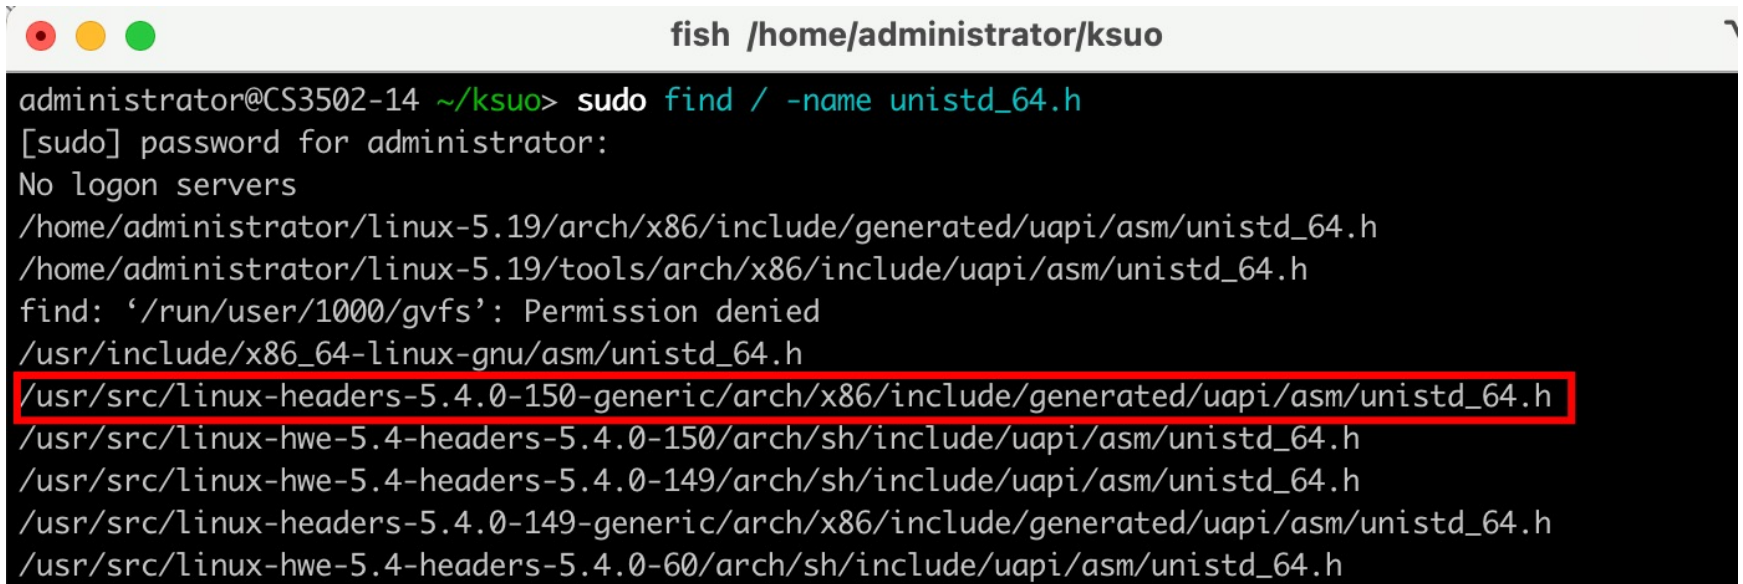
- Step 1: Check the available system call number
- Step 2: Create a kernel module syscall
- Step 3: Define the Makefile
- Step 4: Compile and enable the module syscall
- Step 5: write a test program to test your system call



# Step 1: Check the available system call number

---

```
$ sudo find / -name unistd_64.h
```

A terminal window titled 'fish /home/administrator/ksuo' showing the execution of the command 'sudo find / -name unistd\_64.h'. The output lists several file paths, with the path '/usr/src/linux-headers-5.4.0-150-generic/arch/x86/include/generated/uapi/asm/unistd\_64.h' highlighted by a red rectangular box.

```
administrator@CS3502-14 ~/ksuo> sudo find / -name unistd_64.h
[sudo] password for administrator:
No logon servers
/home/administrator/linux-5.19/arch/x86/include/generated/uapi/asm/unistd_64.h
/home/administrator/linux-5.19/tools/arch/x86/include/uapi/asm/unistd_64.h
find: '/run/user/1000/gvfs': Permission denied
/usr/include/x86_64-linux-gnu/asm/unistd_64.h
/usr/src/linux-headers-5.4.0-150-generic/arch/x86/include/generated/uapi/asm/unistd_64.h
/usr/src/linux-hwe-5.4-headers-5.4.0-150/arch/sh/include/uapi/asm/unistd_64.h
/usr/src/linux-hwe-5.4-headers-5.4.0-149/arch/sh/include/uapi/asm/unistd_64.h
/usr/src/linux-headers-5.4.0-149-generic/arch/x86/include/generated/uapi/asm/unistd_64.h
/usr/src/linux-hwe-5.4-headers-5.4.0-60/arch/sh/include/uapi/asm/unistd_64.h
```

As my current Linux version is 5.4.0-150-generic, so I check the above one.

# Step 1: Check the available system call number

---

```
$ sudo cat /usr/src/linux-headers-5.4.0-150-  
generic/arch/x86/include/generated/uapi/asm/unistd_64.h
```

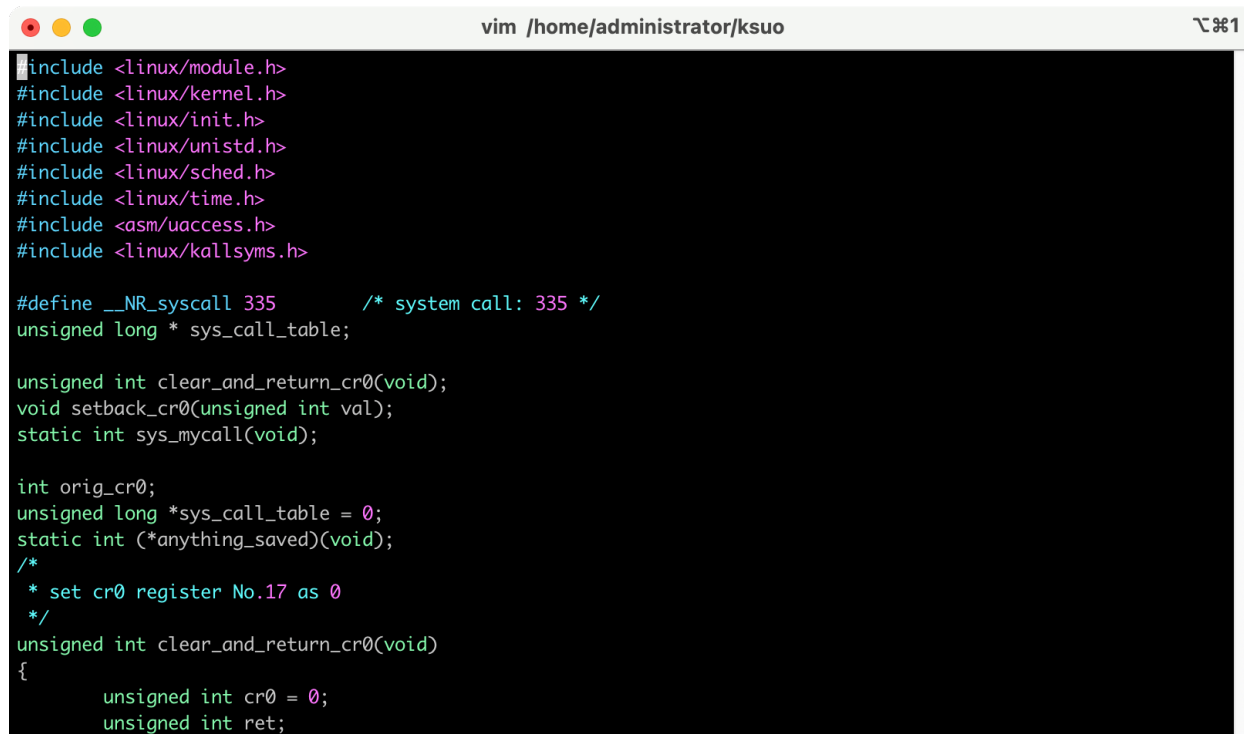
```
fish /home/administrator/ksuo  
administrator@CS3502-14 ~/ksuo> sudo cat /usr/src/linux-headers-5.4.0-150-generic/arch/x86/include/ge  
nerated/uapi/asm/unistd_64.h  
#ifndef _ASM_X86_UNISTD_64_H  
#define _ASM_X86_UNISTD_64_H 1  
  
#define __NR_read 0  
#define __NR_write 1  
#define __NR_open 2  
#define __NR_close 3  
#define __NR_stat 4  
#define __NR_fstat 5  
  
#define __NR_pkey_alloc 330  
#define __NR_pkey_free 331  
#define __NR_statx 332  
#define __NR_io_pgetevents 333  
#define __NR_rseq 334  
#define __NR_pidfd_send_signal 424  
#define __NR_io_uring_setup 425  
#define __NR_io_uring_enter 426  
#define __NR_io_uring_register 427  
#define __NR_open_tree 428
```

No.335 is not used yet

## Step 2: Create a kernel module syscall

---

\$ vim syscall.c



```
vim /home/administrator/ksuo ㄟ#1
#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/init.h>
#include <linux/unistd.h>
#include <linux/sched.h>
#include <linux/time.h>
#include <asm/uaccess.h>
#include <linux/kallsyms.h>

#define __NR_syscall 335      /* system call: 335 */
unsigned long * sys_call_table;

unsigned int clear_and_return_cr0(void);
void setback_cr0(unsigned int val);
static int sys_mycall(void);

int orig_cr0;
unsigned long *sys_call_table = 0;
static int (*anything_saved)(void);
/*
 * set cr0 register No.17 as 0
 */
unsigned int clear_and_return_cr0(void)
{
    unsigned int cr0 = 0;
    unsigned int ret;
```

<https://github.com/kevinsuo/CS3502/blob/master/syscall.c>



## Step 2: Create a kernel module syscall

```
/* The former is used in 32-bit systems. The latter is used in 64-bit systems, this system is 64-bit */
//asm volatile ("movl %%cr0, %%eax" : "=a"(cr0));
asm volatile ("movq %%cr0, %%rax" : "=a"(cr0));

ret = cr0;
cr0 &= 0xfffffff;

//asm volatile ("movl %%eax, %%cr0" :: "a"(cr0));
asm volatile ("movq %%rax, %%cr0" :: "a"(cr0));
return ret;
}

/* Read the value of val to the rax register, and then put the value of the rax register into cr0 */
void setback_cr0(unsigned int val)
{

    //asm volatile ("movl %%eax, %%cr0" :: "a"(val));
    asm volatile ("movq %%rax, %%cr0" :: "a"(val));
}

/* Our system call is here */
static int sys_mycall(void)
{
    int ret = 12345;
    printk("Here is my syscall in OS kernel!\n");
    return ret;
}
```

Here is the message to be printed  
when it is called

<https://github.com/kevinsuo/CS3502/blob/master/syscall.c>

# Step 3: Define the Makefile

---

```
vim /home/administrator/ksuo  
  
obj-m:=syscall.o  
PWD:= $(shell pwd)  
KERNELDIR:= /lib/modules/$(shell uname -r)/build  
EXTRA_CFLAGS= -O0  
  
all:  
    make -C $(KERNELDIR) M=$(PWD) modules  
  
clean:  
    make -C $(KERNELDIR) M=$(PWD) clean
```

<https://github.com/kevinsuo/CS3502/blob/master/Makefile>

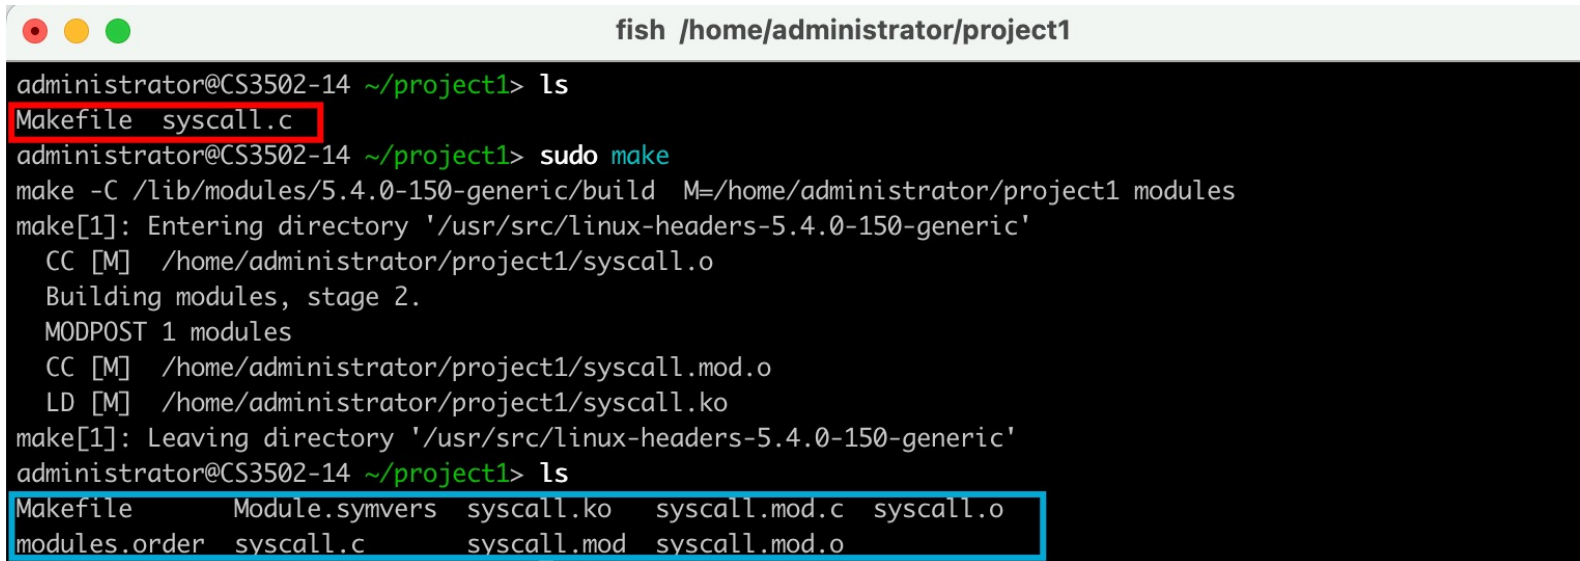




# Step 4: Compile and enable the module syscall

---

\$ sudo make



```
fish /home/administrator/project1
administrator@CS3502-14 ~/project1> ls
Makefile  syscall.c
administrator@CS3502-14 ~/project1> sudo make
make -C /lib/modules/5.4.0-150-generic/build M=/home/administrator/project1 modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-150-generic'
  CC [M]  /home/administrator/project1/syscall.o
Building modules, stage 2.
MODPOST 1 modules
  CC [M]  /home/administrator/project1/syscall.mod.o
  LD [M]  /home/administrator/project1/syscall.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-150-generic'
administrator@CS3502-14 ~/project1> ls
Makefile      Module.symvers  syscall.ko  syscall.mod.c  syscall.o
modules.order  syscall.c       syscall.mod  syscall.mod.o
```

The red and blue parts are before and after the compiling



# Step 4: Compile and enable the module syscall

---

`$ sudo insmod syscall.ko`

```
fish /home/administrator/project1
administrator@CS3502-14 ~/project1> sudo insmod syscall.ko
```

You can use the `$ lsmod` to check the enabled modules:

```
fish /home/administrator/project1
administrator@CS3502-14 ~/project1> lsmod
Module                Size  Used by
syscall               16384  0
test                  16384  0
ipt_REJECT             16384  2
nf_reject_ipv4         16384  1 ipt_REJECT
xt_tcpudp              20480  2
iptable_filter         16384  1
bpfILTER              24576  0
```



# Step 5: write user level app to call it

---

```
vim /home/administrator/ksuo
#include <syscall.h>
#include <stdio.h>
int main(void)
{
    /* call system call No.335 */
    printf("%d\n",syscall(335));
    return 0;
}
```

Compile and execute:

```
$ gcc test.c -o test.o
```

```
$ sudo ./test.o
```

The test program will call the new system call and output a message that you defined in step 2 at the tail of the output of dmesg (system log).

```
fish /home/administrator/project1
administrator@CS3502-14 ~/project1> sudo ./test.o
12345
administrator@CS3502-14 ~/project1> dmesg | grep my
[178116.531246] Here is my syscall in OS kernel!
administrator@CS3502-14 ~/project1>
```

# Put it all together

## User space

*Step 5: user call it*

```
#include <syscall.h>
#include <stdio.h>

int main(void)
{
    /* call system call No.335 */
    printf("%d\n", syscall(335));
    return 0;
}
```

```
fish /home/ad
administrator@CS3502-14 ~/project1> sudo ./test.o
12345
administrator@CS3502-14 ~/project1> dmesg | grep my
[178116.531246] Here is my syscall in OS kernel!
administrator@CS3502-14 ~/project1>
```

## Kernel space

*Step 2: system call defined in kernel module*

```
#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/init.h>
#include <linux/unistd.h>
#include <linux/sched.h>
#include <linux/time.h>
#include <asm/uaccess.h>
#include <linux/kallsyms.h>
```

```
#define __NR_syscall 335 /* system call: 335 */
unsigned long * sys_call_table;
```

```
static int __init init_addsyscall(void)
```

```
{
    printk("My syscall is starting. . . \n");
    sys_call_table = (unsigned long *)kallsyms_lookup_name("sys_call_table");
    printk("sys_call_table: 0x%p\n", sys_call_table);
    anything_saved = (int (*)(void))(sys_call_table[__NR_syscall]);
    orig_cr0 = clear_and_return_cr0();
    sys_call_table[__NR_syscall] = (unsigned long)&sys_mycall;
    setback_cr0(orig_cr0);
    return 0;
}
```

*Step 4: init and enable the kernel module*

```
/* Our system call is here */
static int sys_mycall(void)
{
    int ret = 12345;
    printk("Here is my syscall in OS kernel!\n");
    return ret;
}
```