

CS 3502

Operating Systems

Interrupt

Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

Outline

- What is interrupt?
 - Interrupt vs Polling
 - Advanced programmable interrupt controller
 - Interrupt processing
- Interrupt types and affinity
 - Hardware and software interrupt
 - Interrupt affinity

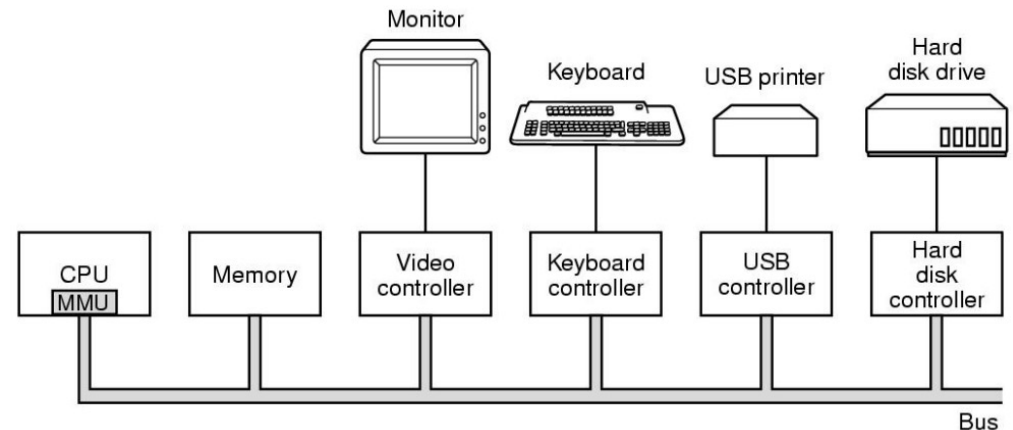


How can the processor work with hardware



When event happens
When task finishes
When job is ready
When user inputs something
When something is broken

...



What is Polling?

- How can the processor work with hardware without impacting the machine's overall performance?
 - **Polling: repeatedly check whether the hardware is ready or not**

Polling



Polling Drawbacks


- How can the processor work with hardware without impacting the machine's overall performance?
 - **Polling: repeatedly check whether the hardware is ready or not**

Polling: overhead

Randomly

Every N_s/min

1, Increase system overhead. Whether it is task polling or timer polling, it needs to consume the corresponding system resources.




What is Polling?

- How can the processor work with hardware without impacting the machine's overall performance?
 - **Polling: repeatedly check whether the hardware is ready or not**

Polling: (timely)

2, Unable to sense device status changes in a timely manner. Device state changes during the polling interval can only be discovered on the next poll, which will not be able to meet real-time sensitive applications.




What is Polling?

- How can the processor work with hardware without impacting the machine's overall performance?
 - **Polling: repeatedly check whether the hardware is ready or not**

Polling: resource efficiency

3, waste CPU resources. Polling is always in progress regardless of whether the device has changed state. In the real world, the state change of most devices is usually not so frequent, and polling idle will waste CPU time slices.




What is Polling?

- How can the processor work with hardware without impacting the machine's overall performance?
 - **Polling: repeatedly check whether the hardware is ready or not**

Polling:

- 1, increase system overhead
- 2, cannot detect in time
- 3, waste of CPU resources




What is interrupt?


- How can the processor work with hardware without impacting the machine's overall performance?
 - Polling: repeatedly check whether the hardware is ready or not
 - **Interrupt: hardware signals to the kernel when attention is needed**

Polling:

- 1, increase system overhead
- 2, cannot detect in time
- 3, waste of CPU resources



Interrupt



Interrupt vs. Polling

Example for Polling and Interrupt



CPU (Mario)



I/O Device (Princess Peach)

Unregistered PowerVideoMaker

Reference: https://www.youtube.com/watch?v=M3nXI_86uLE

Interrupt vs. Polling

- Polling:
 - increase system overhead
 - cannot detect in time
 - waste of CPU resources
- Interrupt
 - lightweight and save CPU resources
 - handle in time




What will happen if you have too many phone calls?



1. do not have time to work on valuable things
2. might miss some important calls and tasks
3. loss data

Interrupt problem



Interrupt:

Too many **interrupt**

-> Less resources for **IRQ handler**

-> miss some interrupt and data


Phone sounds



Answer the call



Combine Interrupt & Polling




Interrupt:

Too many **interrupt**

-> Less resources for **IRQ handler**


-> miss some interrupt and data



NAPI: network receiving

Combine Interrupt & Polling

- Benefits
 - Avoid frequent interruptions when receiving large amounts of data or requests
 - No need to spend a lot of CPU resources for polling



NAPI: network receiving



Interrupt: Quick look

- `$ cat /proc/interrupts`
- `$ watch -n 1 cat /proc/interrupts`

```

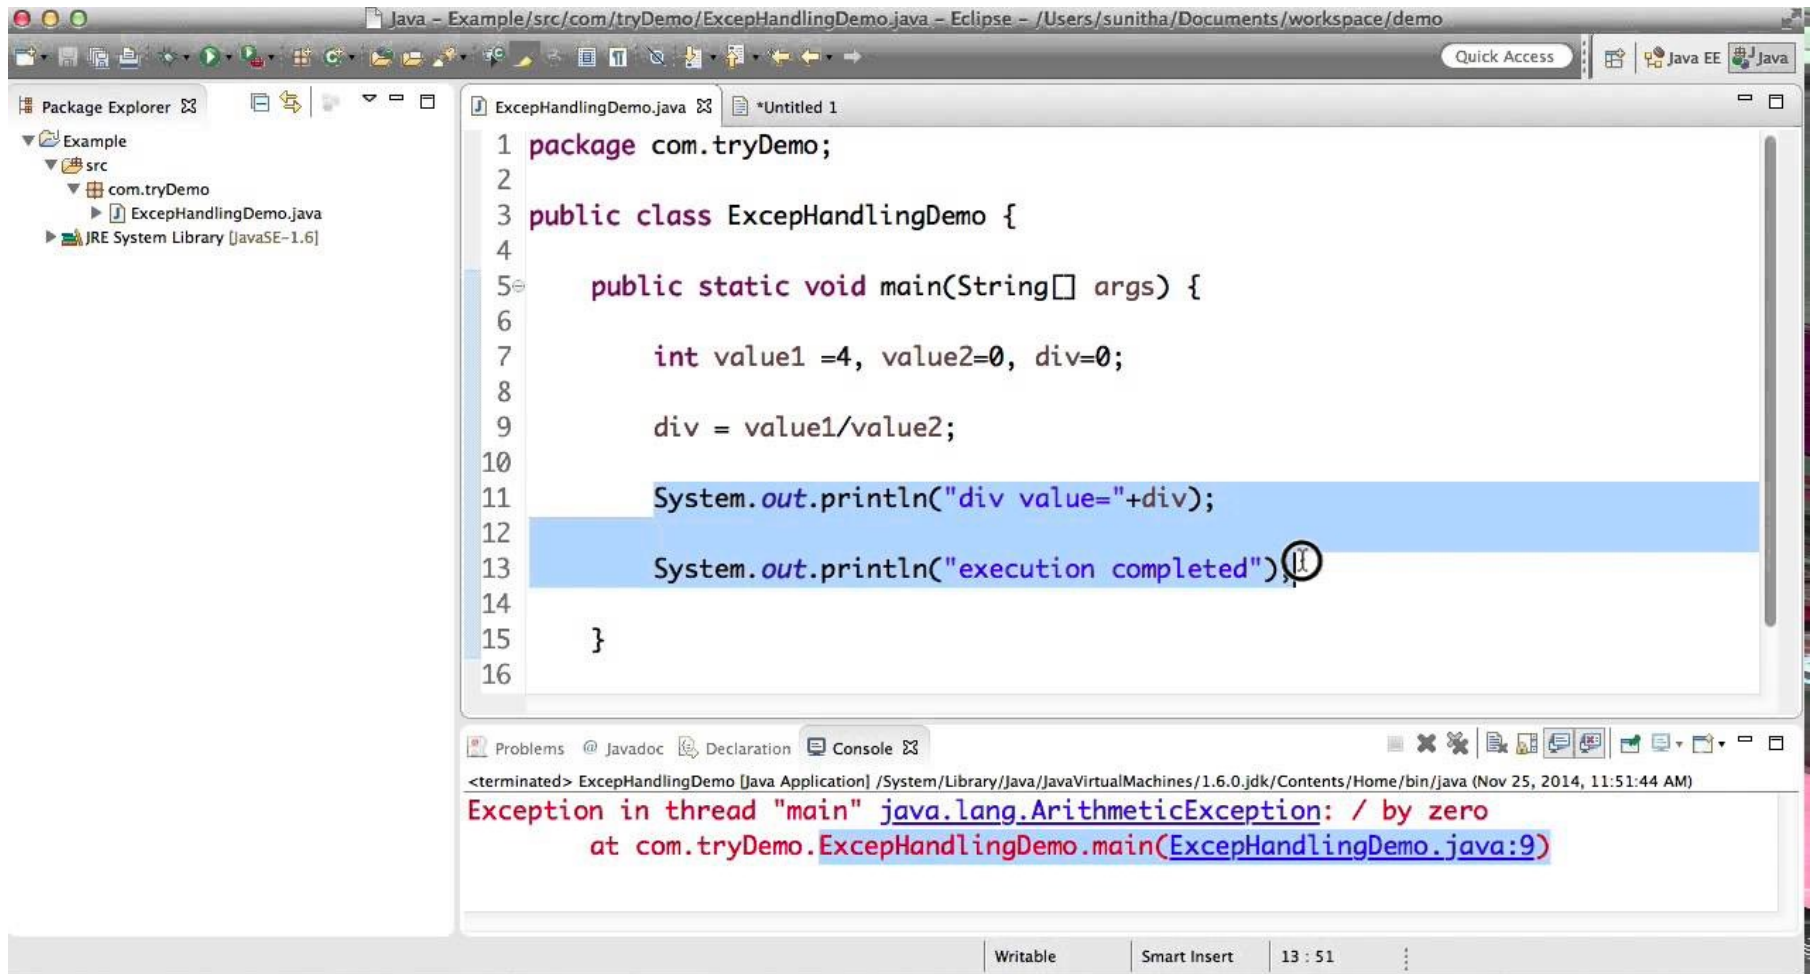
          CPU0      CPU1
0:         1         0  IO-APIC  2-edge  timer
1:         0        77  IO-APIC  1-edge  i8042
8:         1         0  IO-APIC  8-edge  rtc0
9:         0         0  IO-APIC  9-fasteoi acpi
12:        15         0  IO-APIC 12-edge  i8042
14:         0         0  IO-APIC 14-edge  ata_piix
15:         0    1028751  IO-APIC 15-edge  ata_piix
16:         1        218  IO-APIC 16-fasteoi vmwgfx
17:         0    477023  IO-APIC 17-fasteoi ioc0
24:         0         0  PCI-MSI 344064-edge  PCIe PME, pciehp
25:         0         0  PCI-MSI 346112-edge  PCIe PME, pciehp
26:         0         0  PCI-MSI 348160-edge  PCIe PME, pciehp
27:         0         0  PCI-MSI 350208-edge  PCIe PME, pciehp
```

Exceptions

- An exception indicates that code running on the CPU has created a situation that the processor needs help to address.
- Can you think of examples of software exceptions?
 - Divide by zero -- probably kills the process.
 - Attempt to use a privileged instruction - - also probably kills the process.
 - Attempt to use a virtual address that the CPU does not know how to translate -- a common exception handled transparently as part of virtual memory management.



Exceptions example




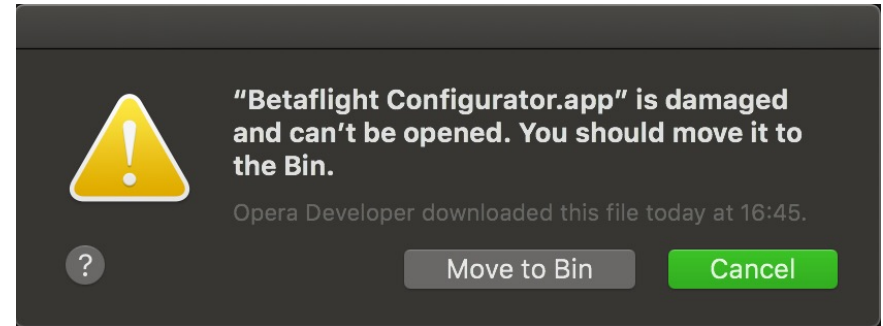

The screenshot shows the Eclipse IDE with a Java project named 'Example'. The package explorer on the left shows the source code structure: 'src' containing 'com.tryDemo' which has 'ExcepHandlingDemo.java'. The main editor displays the code for 'ExcepHandlingDemo.java'.

```
1 package com.tryDemo;
2
3 public class ExcepHandlingDemo {
4
5     public static void main(String[] args) {
6
7         int value1 =4, value2=0, div=0;
8
9         div = value1/value2;
10
11         System.out.println("div value="+div);
12
13         System.out.println("execution completed");
14
15     }
16 }
```

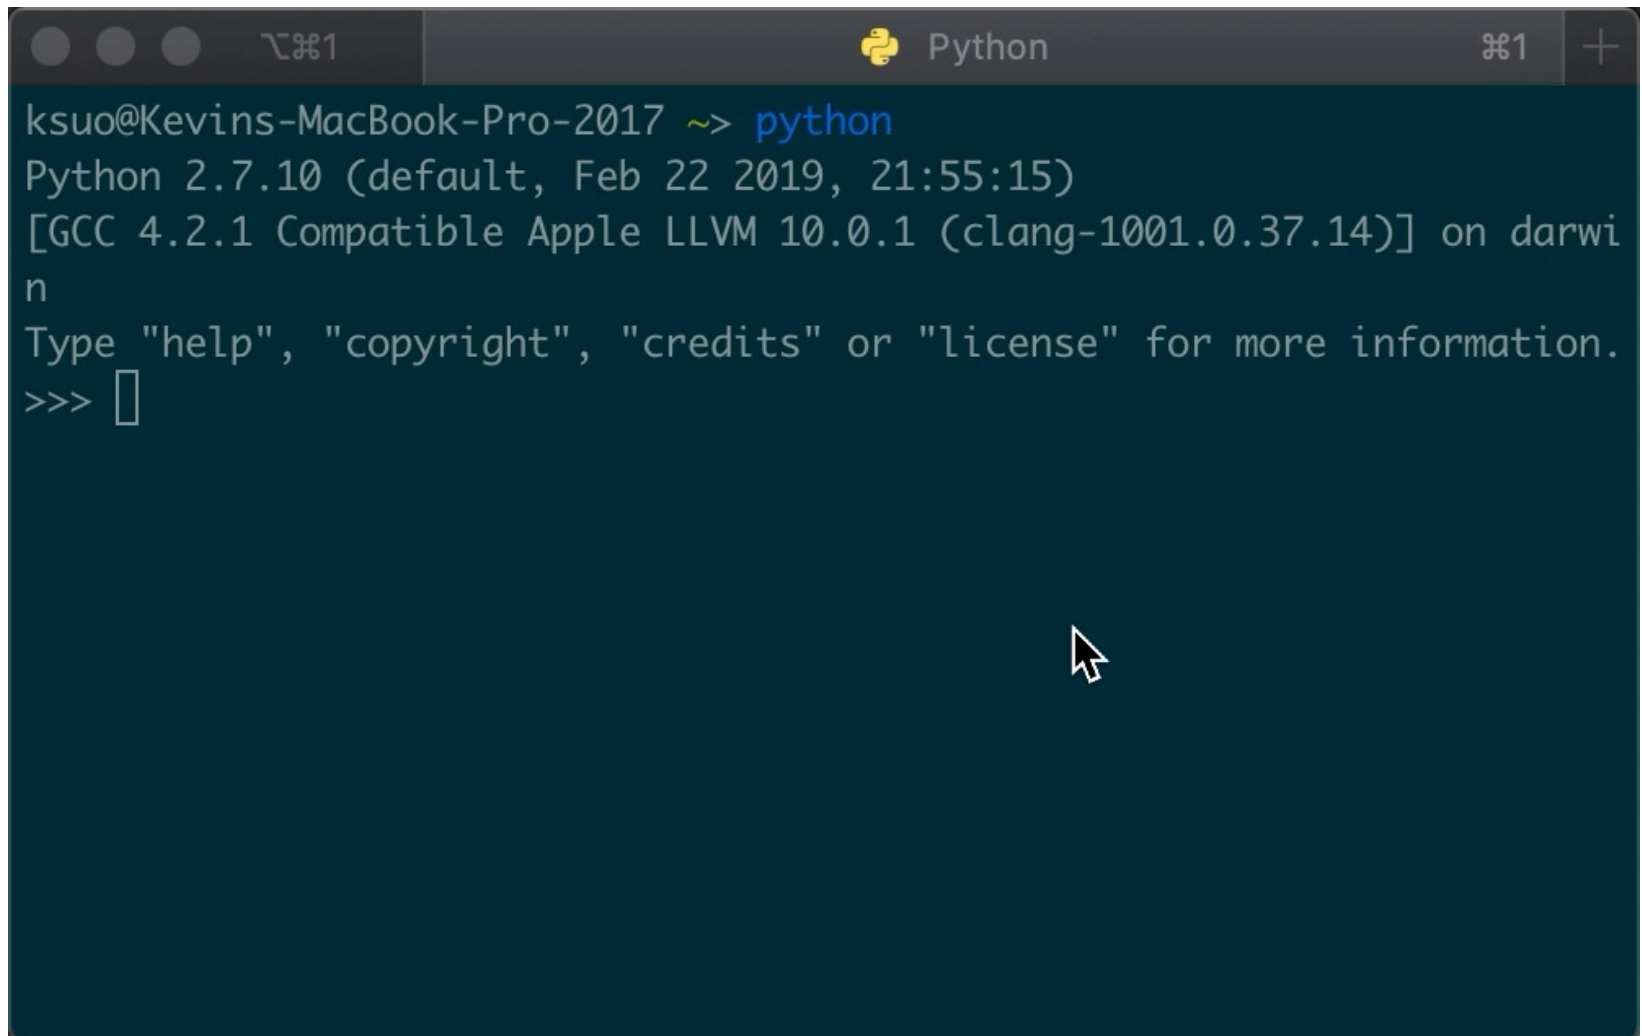
The console output at the bottom shows the program's execution and the resulting exception:

```
<terminated> ExcepHandlingDemo [Java Application] /System/Library/Java/JavaVirtualMachines/1.6.0.jdk/Contents/Home/bin/java (Nov 25, 2014, 11:51:44 AM)
Exception in thread "main" java.lang.ArithmeticException: / by zero
at com.tryDemo.ExcepHandlingDemo.main(ExcepHandlingDemo.java:9)
```

Exceptions example



Exceptions example

A screenshot of a macOS terminal window. The title bar at the top shows three window control buttons (red, yellow, green) on the left, a tab labeled 'Python' with a yellow Python logo icon in the center, and a window zoom button on the right. The terminal content shows a user prompt 'ksuo@Kevins-MacBook-Pro-2017 ~>' followed by the command 'python' in blue. The output shows 'Python 2.7.10 (default, Feb 22 2019, 21:55:15)' and '[GCC 4.2.1 Compatible Apple LLVM 10.0.1 (clang-1001.0.37.14)] on darwin'. Below this is the instruction 'Type "help", "copyright", "credits" or "license" for more information.' and the prompt '>>>' followed by a cursor. A mouse cursor is visible in the lower right area of the terminal window.

```
ksuo@Kevins-MacBook-Pro-2017 ~> python
Python 2.7.10 (default, Feb 22 2019, 21:55:15)
[GCC 4.2.1 Compatible Apple LLVM 10.0.1 (clang-1001.0.37.14)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
>>> 
```

Interrupt vs. Exception

- Interrupts are **voluntary**.
 - “The process actively asks for assistance.”
- Exceptions are **non-voluntary**.
 - "It just tried to divide by zero, and I think it needs to be terminated. I need some help with this process. "




Outline

- What is interrupt?
 - Interrupt vs Polling
 - Advanced programmable interrupt controller
 - Interrupt processing
- Interrupt types and affinity
 - Hardware and software interrupt
 - Interrupt affinity




APIC: Advanced programmable interrupt controller



APIC: Advanced programmable interrupt controller

1. Responsible for telling the CPU when a specific external device wishes to 'interrupt'
 - Needs to tell the CPU which one among several devices is the one needing service




APIC: Advanced programmable interrupt controller

2. APIC translates IRQ to interrupt number

- Raises interrupt to CPU
- Interrupt # available in register


Resource	Device	Status
IRQ 0	System timer	OK
IRQ 3		OK
IRQ 4	Communications Port (COM1)	OK
IRQ 6	Standard floppy disk controller	OK
IRQ 8	System CMOS/real time clock	OK
IRQ 10	NVIDIA nForce PCI System Management	OK
IRQ 10	Multimedia Audio Controller	OK
IRQ 13	Numeric data processor	OK
IRQ 14	ATA Channel 0	OK
IRQ 15	ATA Channel 1	OK
IRQ 18	Realtek RTL8130/810x Family Fast Ethernet NIC	OK



APIC: Advanced programmable interrupt controller


3. Interrupts can have varying priorities

- APIC also needs to prioritize multiple requests



APIC: Advanced programmable interrupt controller

4. Possible to “mask” (disable) interrupts at PIC or CPU at PIC or CPU



APIC: Advanced programmable interrupt controller

1. Responsible for telling the CPU when a specific external device wishes to ‘interrupt’
 - Needs to tell the CPU which one among several devices is the one needing service
2. APIC translates IRQ to interrupt number
 - Raises interrupt to CPU
 - Interrupt # available in register
3. Interrupts can have varying priorities
 - APIC also needs to prioritize multiple requests
4. Possible to “mask” (disable) interrupts at PIC or CPU




APIC:

cat /proc/interrupts

VM on Amazon EC2

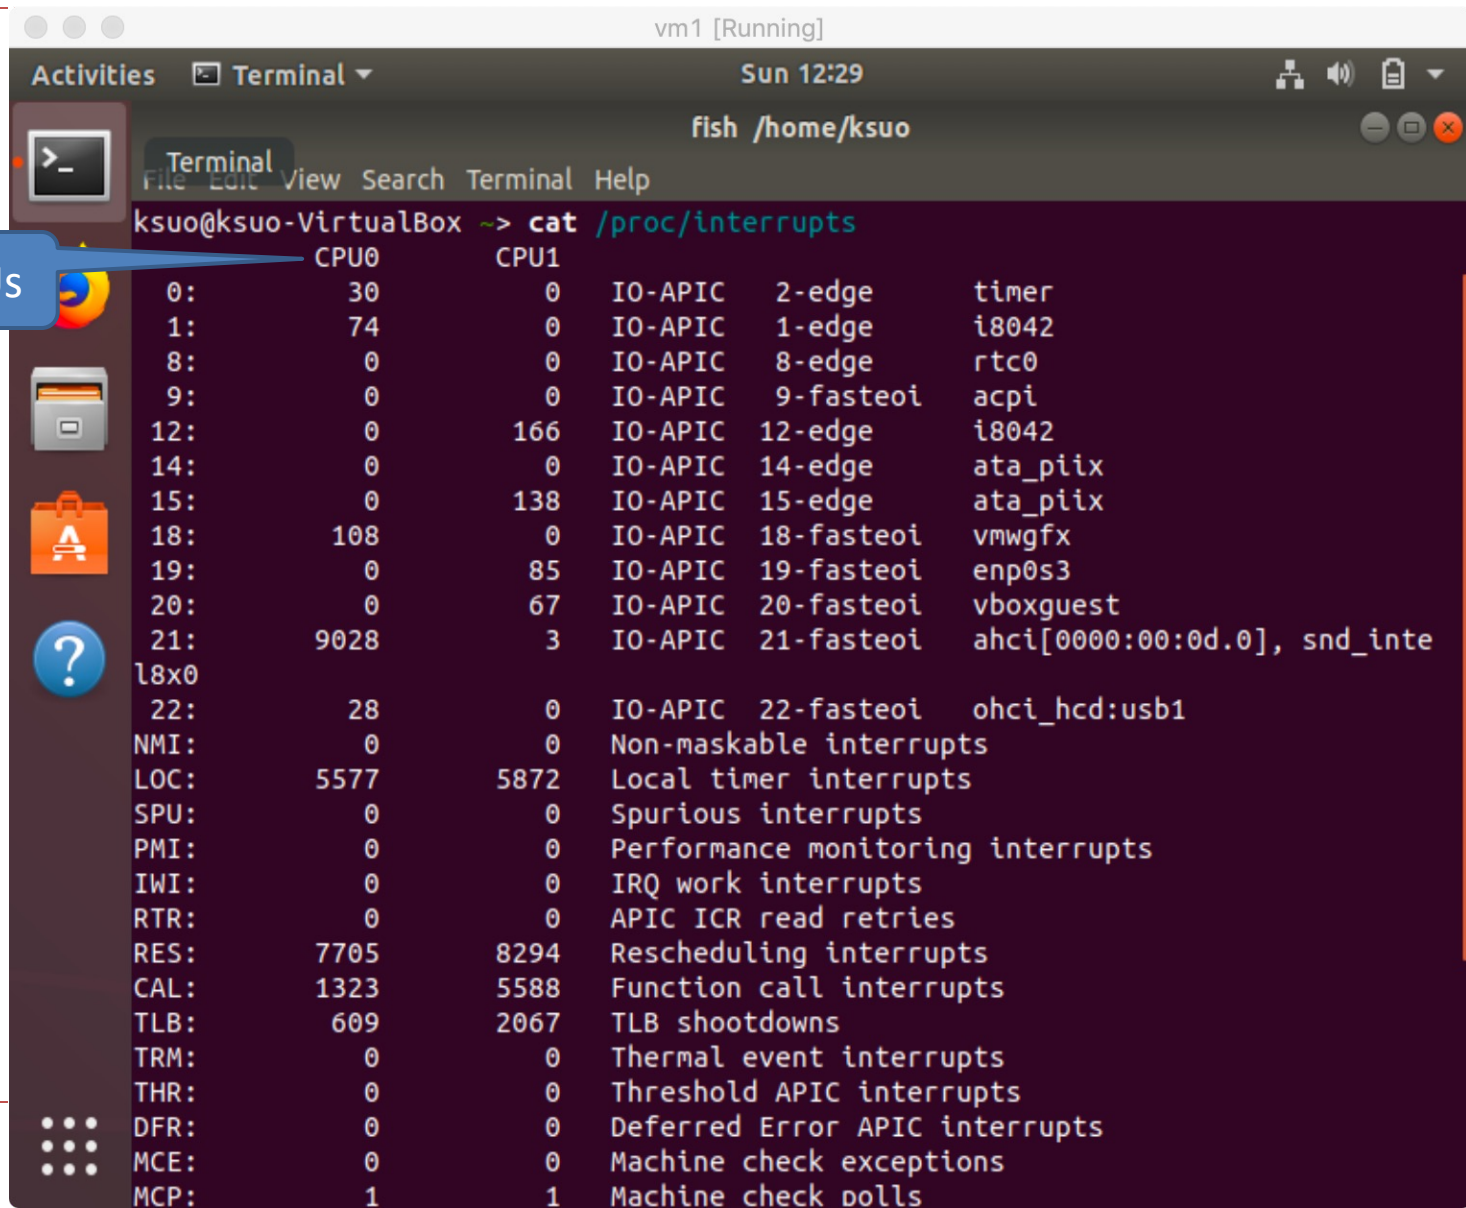
One CPU



```
1. ec2-user@ip-172-31-36-115:~ (ssh)
fish /home/pi/D...  1  ec2-user@ip-172-...  2
[ec2-user@ip-172-31-36-115 ~]$ cat /proc/interrupts
```

	CPU0			
0:	47	I/O-APIC	2-edge	timer
1:	9	xen-pirq	1-ioapic-edge	i8042
7:	2599	xen-pirq	4-ioapic-edge	ttyS0
8:	2	xen-pirq	8-ioapic-edge	rtc0
9:	0	xen-pirq	9-ioapic-level	acpi
12:	86	xen-pirq	12-ioapic-edge	i8042
14:	0	I/O-APIC	14-edge	ata_piix
15:	0	I/O-APIC	15-edge	ata_piix
48:	4827	xen-percpu	-virq	timer0
49:	0	xen-percpu	-ipi	resched0
50:	0	xen-percpu	-ipi	callfunc0
51:	0	xen-percpu	-virq	debug0
52:	0	xen-percpu	-ipi	callfuncsingle0
53:	0	xen-percpu	-ipi	spinlock0
54:	254	xen-dyn	-event	xenbus
55:	19103	xen-dyn	-event	eth0
57:	6784	xen-dyn	-event	blkif
NMI:	0	Non-maskable interrupts		
LOC:	0	Local timer interrupts		
SPU:	0	Spurious interrupts		
PMI:	0	Performance monitoring interrupts		
IWI:	0	IRQ work interrupts		
RTR:	0	APIC ICR read retries		
RES:	0	Rescheduling interrupts		
CAL:	0	Function call interrupts		
TLB:	0	TLB shutdowns		
TRM:	0	Thermal event interrupts		
THR:	0	Threshold APIC interrupts		
DFR:	0	Deferred Error APIC interrupts		
MCE:	0	Machine check exceptions		
MCP:	1	Machine check polls		
HYP:	32046	Hypervisor callback interrupts		
ERR:	0			
MIS:	0			
PIN:	0	Posted-interrupt notification event		
NPI:	0	Nested posted-interrupt event		
PIW:	0	Posted-interrupt wakeup event		

APIC: cat /proc/interrupts



```
vm1 [Running]
Sun 12:29
fish /home/ksuo
Terminal
ksuo@ksuo-VirtualBox ~-> cat /proc/interrupts
```

	CPU0	CPU1			
0:	30	0	IO-APIC	2-edge	timer
1:	74	0	IO-APIC	1-edge	i8042
8:	0	0	IO-APIC	8-edge	rtc0
9:	0	0	IO-APIC	9-fastest	acpi
12:	0	166	IO-APIC	12-edge	i8042
14:	0	0	IO-APIC	14-edge	ata_piix
15:	0	138	IO-APIC	15-edge	ata_piix
18:	108	0	IO-APIC	18-fastest	vmwgfx
19:	0	85	IO-APIC	19-fastest	enp0s3
20:	0	67	IO-APIC	20-fastest	vboxguest
21:	9028	3	IO-APIC	21-fastest	ahci[0000:00:0d.0], snd_intel8x0
22:	28	0	IO-APIC	22-fastest	ohci_hcd:usb1
NMI:	0	0			Non-maskable interrupts
LOC:	5577	5872			Local timer interrupts
SPU:	0	0			Spurious interrupts
PMI:	0	0			Performance monitoring interrupts
IWI:	0	0			IRQ work interrupts
RTR:	0	0			APIC ICR read retries
RES:	7705	8294			Rescheduling interrupts
CAL:	1323	5588			Function call interrupts
TLB:	609	2067			TLB shutdowns
TRM:	0	0			Thermal event interrupts
THR:	0	0			Threshold APIC interrupts
DFR:	0	0			Deferred Error APIC interrupts
MCE:	0	0			Machine check exceptions
MCP:	1	1			Machine check polls

two CPUs


Outline

- What is interrupt?
 - Interrupt vs Polling
 - Advanced programmable interrupt controller
 - Interrupt processing
- Interrupt types and affinity
 - Hardware and software interrupt
 - Interrupt affinity



Similar as System call table


- There are approximately 300 system-calls in Linux 2.6.
- An array of function-pointers (identified by the ID number)
- This array is named 'sys_call_table[]' in Linux https://elixir.bootlin.com/linux/v5.0/source/arch/x86/entry/syscall_64.c





The 'jump-table' idea

Interrupt Descriptor Table (IDT)

- IDT is in memory, initialized by OS at boot
- IDT associates with each interrupt vector and stores the entry address of the corresponding interrupt handler




The process of interrupt



\$0.52

Texas Instruments
NE555N Timers &...




APIC chip



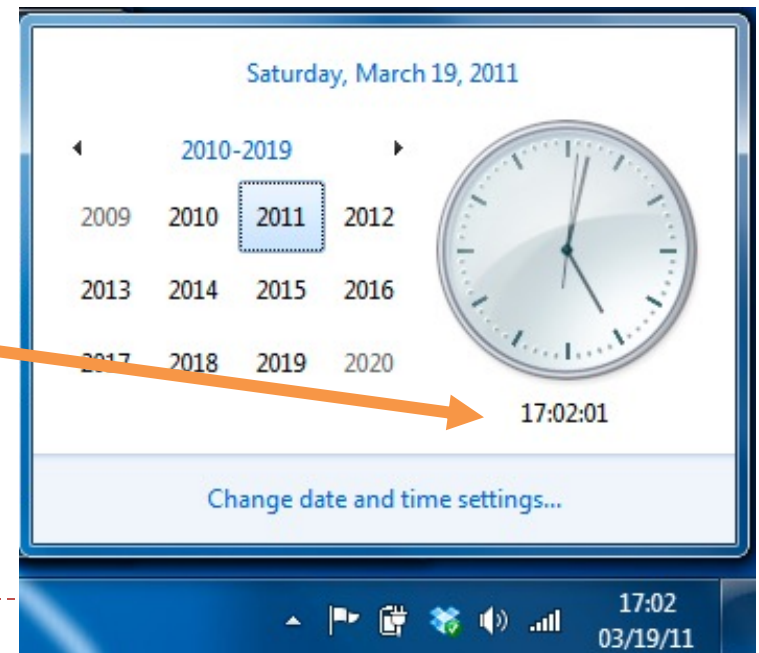
The process of interrupt

Take the *time interrupt* as an example



```
[ec2-user@ip-172-31-36-115 ~]$ cat /proc/interrupts
```

CPU0				
0:	47	I/O-APIC	2-edge	timer
1:	9	xen-pirq	1-ioapic-edge	i8042
4:	2599	xen-pirq	4-ioapic-edge	ttyS0
8:	2	xen-pirq	8-ioapic-edge	rtc0
9:	0	xen-pirq	9-ioapic-level	acpi
12:	86	xen-pirq	12-ioapic-edge	i8042
14:	0	I/O-APIC	14-edge	ata_piix
15:	0	I/O-APIC	15-edge	ata_piix
48:	43857	xen-percpu	-virq	timer0
49:	0	xen-percpu	-ipi	resched0
50:	0	xen-percpu	-ipi	callfunc0
51:	0	xen-percpu	-virq	debug0
52:	0	xen-percpu	-ipi	callfuncsingle0
53:	0	xen-percpu	-ipi	spinlock0
54:	254	xen-dyn	-event	xenbus



Interrupt vs. Exception vs. System call

	Source	Handling	Mechanism
Interrupt	Device, etc.	Asynchronous	Interrupts are handled by the processor after finishing the current instruction. If it finds a signal on its interrupt pin, it will look up the address of the interrupt handler in the interrupt table and pass that routine control. After returning from the interrupt handler routine, it will resume program execution at the instruction after the interrupted instruction.
Exception	Application or kernel unexpected behaviors	Synchronous	Exceptions on the other hand are divided into three kinds. These are Faults, Traps and Aborts. Faults are detected and serviced by the processor before the faulting instructions. Traps are serviced after the instruction causing the trap. Aborts are used only to signal severe system problems, when operation is no longer possible.
System call	Application requests	Asynchronous or Synchronous	A way for programs to interact with the operating system. A computer program makes a system call when it makes a request to the operating system's kernel. System call provides the services of the operating system to the user programs via API.

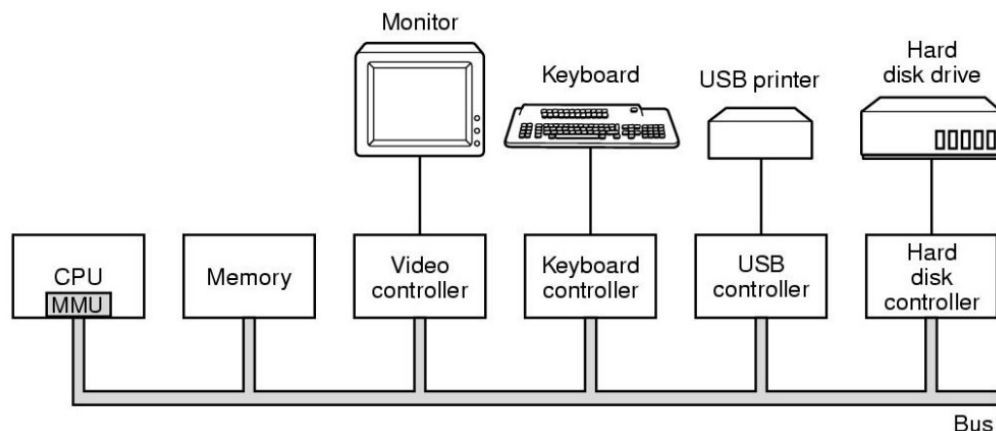
Outline

- What is interrupt?
 - Interrupt vs Polling
 - Advanced programmable interrupt controller
 - Interrupt processing
- Interrupt types and affinity
 - Hardware and software interrupt
 - Interrupt affinity



Hardware interrupt

- Hardware interrupts are used to signal that a particular device needs attention:
 - a disk read completed, or
 - a network was connected, or
 - a timer is ready



Hardware interrupt: cat /proc/interrupts

```
1. fish /home/pi/Downloads (ssh)
fish /home/pi/Dow... 81

pi@raspberrypi ~/Downloads> cat /proc/interrupts
```

	CPU0	CPU1	CPU2	CPU3			
16:	0	0	0	0	bcm2836-timer	0 Edge	arch_timer
17:	21666	63962	23766	9395	bcm2836-timer	1 Edge	arch_timer
21:	0	0	0	0	bcm2836-pmu	9 Edge	arm-pmu
23:	3600	0	0	0	ARMCTRL-level	1 Edge	3f00b880.mailbox
24:	29	0	0	0	ARMCTRL-level	2 Edge	VCHIQ doorbell
46:	0	0	0	0	ARMCTRL-level	48 Edge	bcm2708_fb dma
48:	0	0	0	0	ARMCTRL-level	50 Edge	DMA IRQ
50:	0	0	0	0	ARMCTRL-level	52 Edge	DMA IRQ
51:	360	0	0	0	ARMCTRL-level	53 Edge	DMA IRQ
54:	4908	0	0	0	ARMCTRL-level	56 Edge	DMA IRQ
59:	0	0	0	0	ARMCTRL-level	61 Edge	bcm2835-auxirq
62:	139344	0	0	0	ARMCTRL-level	64 Edge	dwc_otg, dwc_otg_pcd, dwc_otg_hcd:usb1
86:	870	0	0	0	ARMCTRL-level	88 Edge	mmc0
87:	7237	0	0	0	ARMCTRL-level	89 Edge	uart-pl011
92:	101080	0	0	0	ARMCTRL-level	94 Edge	mmc1
169:	0	0	0	0	lan78xx-irqs	17 Edge	usb-001:004:01
FIQ:	usb_fiq						
IPI0:	0	0	0	0	CPU wakeup interrupts		
IPI1:	0	0	0	0	Timer broadcast interrupts		
IPI2:	7855	81594	13799	6452	Rescheduling interrupts		
IPI3:	10	6	9	8	Function call interrupts		
IPI4:	0	0	0	0	CPU stop interrupts		
IPI5:	7342	26013	5157	2072	IRQ work interrupts		
IPI6:	0	0	0	0	completion interrupts		
Err:	0						



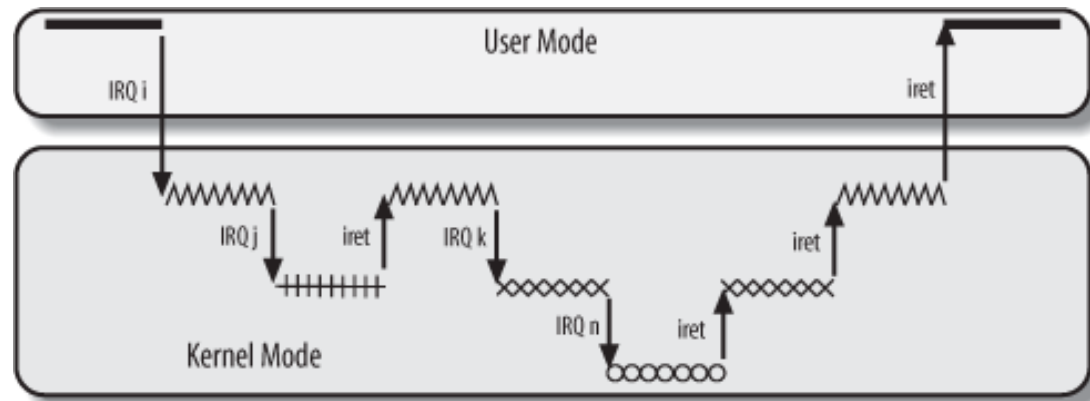
\$ watch -n 1 -d cat /proc/interrupts

```
watch /home/ksuo
Every 1.0s: cat /proc/interrupts
LinuxKernel2: Sun Feb 16 23:21:05 2020
```

	CPU0	CPU1	CPU2	CPU3			
0:	33	0	0	0	IO-APIC	2-edge	timer
1:	0	9	0	0	IO-APIC	1-edge	i8042
6:	0	0	3	0	IO-APIC	6-edge	floppy
8:	0	0	1	0	IO-APIC	8-edge	rtc0
9:	0	0	0	0	IO-APIC	9-fasteoi	acpi
10:	0	0	0	0	IO-APIC	10-fasteoi	virtio0
11:	0	0	0	0	IO-APIC	11-fasteoi	uhci_hcd:usb1
12:	15	0	0	0	IO-APIC	12-edge	i8042
14:	0	0	0	0	IO-APIC	14-edge	ata_piix
15:	0	0	0	0	IO-APIC	15-edge	ata_piix
24:	0	0	0	0	PCI-MSI	81920-edge	virtio2-config
25:	31	0	0	0	PCI-MSI	81921-edge	virtio2-virtqueues
26:	0	0	0	0	PCI-MSI	98304-edge	virtio3-config
27:	0	0	877586	0	PCI-MSI	98305-edge	virtio3-req.0
28:	0	0	0	0	PCI-MSI	65536-edge	virtio1-config
29:	568	0	0	10272921	PCI-MSI	65537-edge	virtio1-input.0
30:	0	3	0	0	PCI-MSI	65538-edge	virtio1-output.0
NMI:	0	0	0	0			Non-maskable interrupts
LOC:	32591583	22611290	143121561	26274990			Local timer interrupts
SPU:	0	0	0	0			Spurious interrupts
PMI:	0	0	0	0			Performance monitoring interrupts
IWI:	0	0	1	0			IRQ work interrupts
RTR:	0	0	0	0			APIC ICR read retries
RES:	7691195	7665686	1340645	9627071			Rescheduling interrupts
CAL:	269029	290892	63368	263701			Function call interrupts
TLB:	11652	12950	12247	13551			TLB shootdowns
TRM:	0	0	0	0			Thermal event interrupts
THR:	0	0	0	0			Threshold APIC interrupts
DFR:	0	0	0	0			Deferred Error APIC interrupts
MCE:	0	0	0	0			Machine check exceptions
MCP:	42203	42203	42203	42203			Machine check polls
HYP:	0	0	0	0			Hypervisor callback interrupts
ERR:	0						
MIS:	0						
PIN:	0	0	0	0			Posted-interrupt notification event
NPI:	0	0	0	0			Nested posted-interrupt event
PIW:	0	0	0	0			Posted-interrupt wakeup event

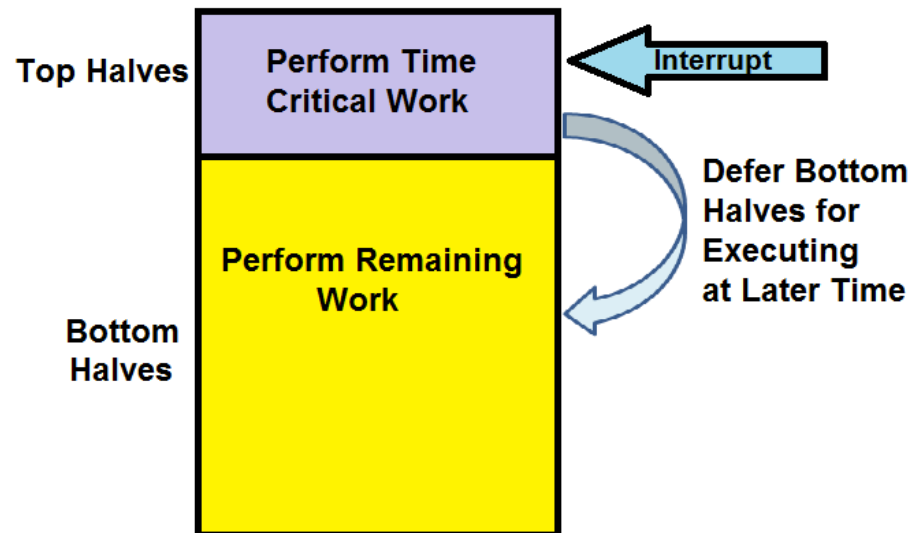
Hardware interrupt feature

- **Definition:** raised by hardware
- **Nest:** hardirq in Linux can be nested
- **Interrupt controller:** IRQ number is provided by APIC



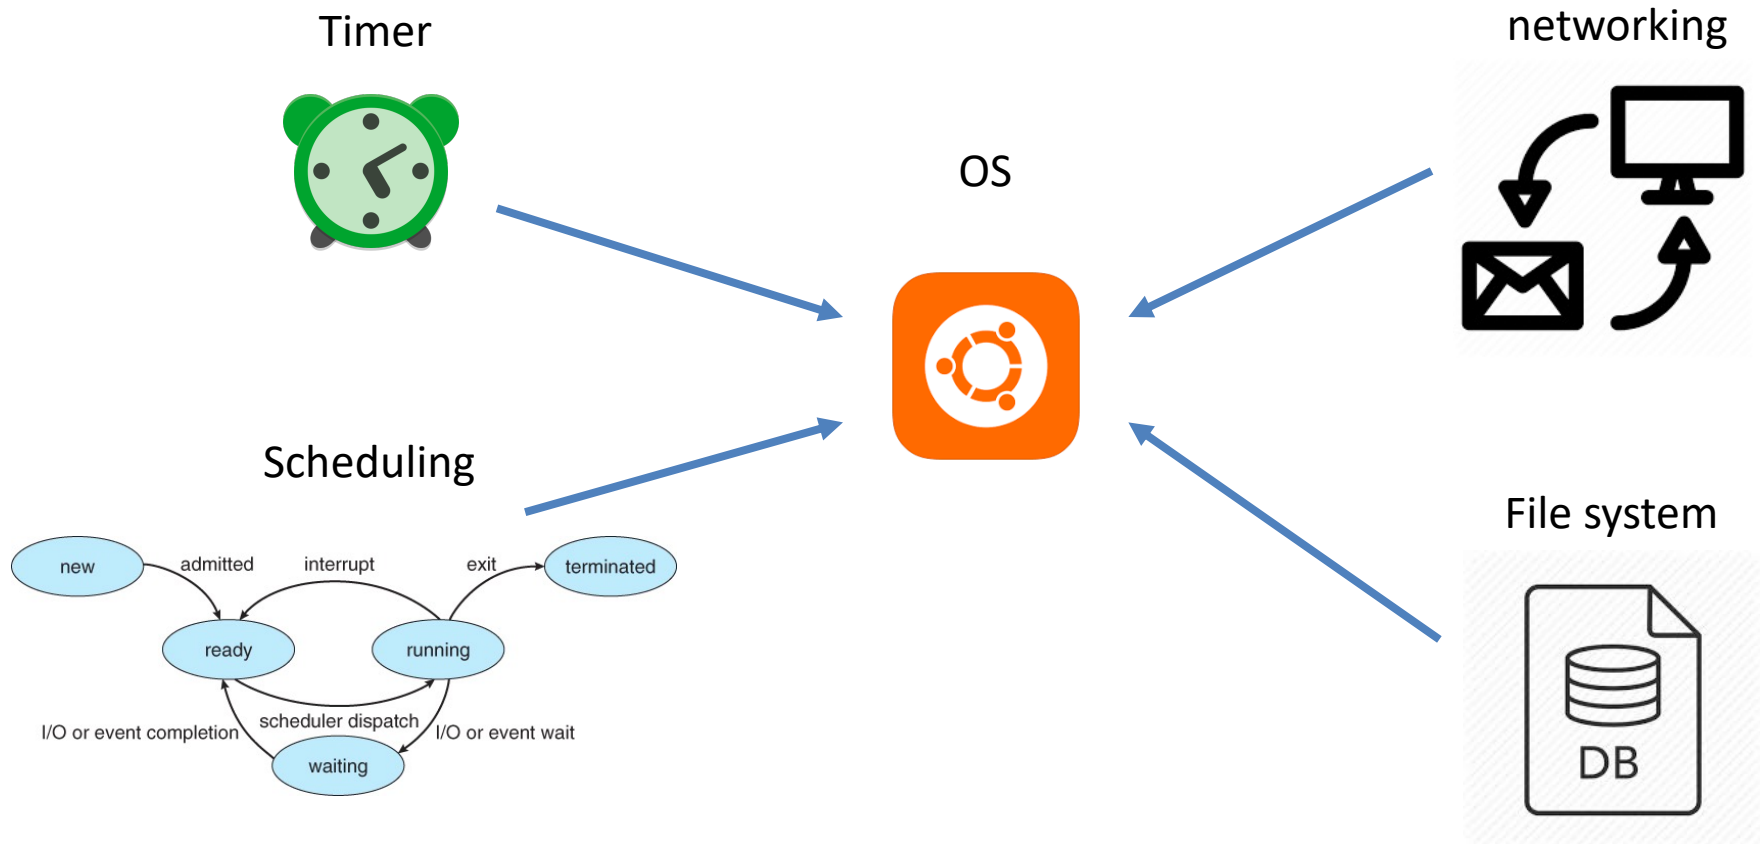
Hardware interrupt feature

- **Mask:** hardirq can be masked
- **Top/Bottom:** hardirq handler (small) ensures that it completes the task quickly
- **Preemption:** hardirq has priority and can preempt over softirq



Software interrupt feature

- **Definition:** raised by software or execution

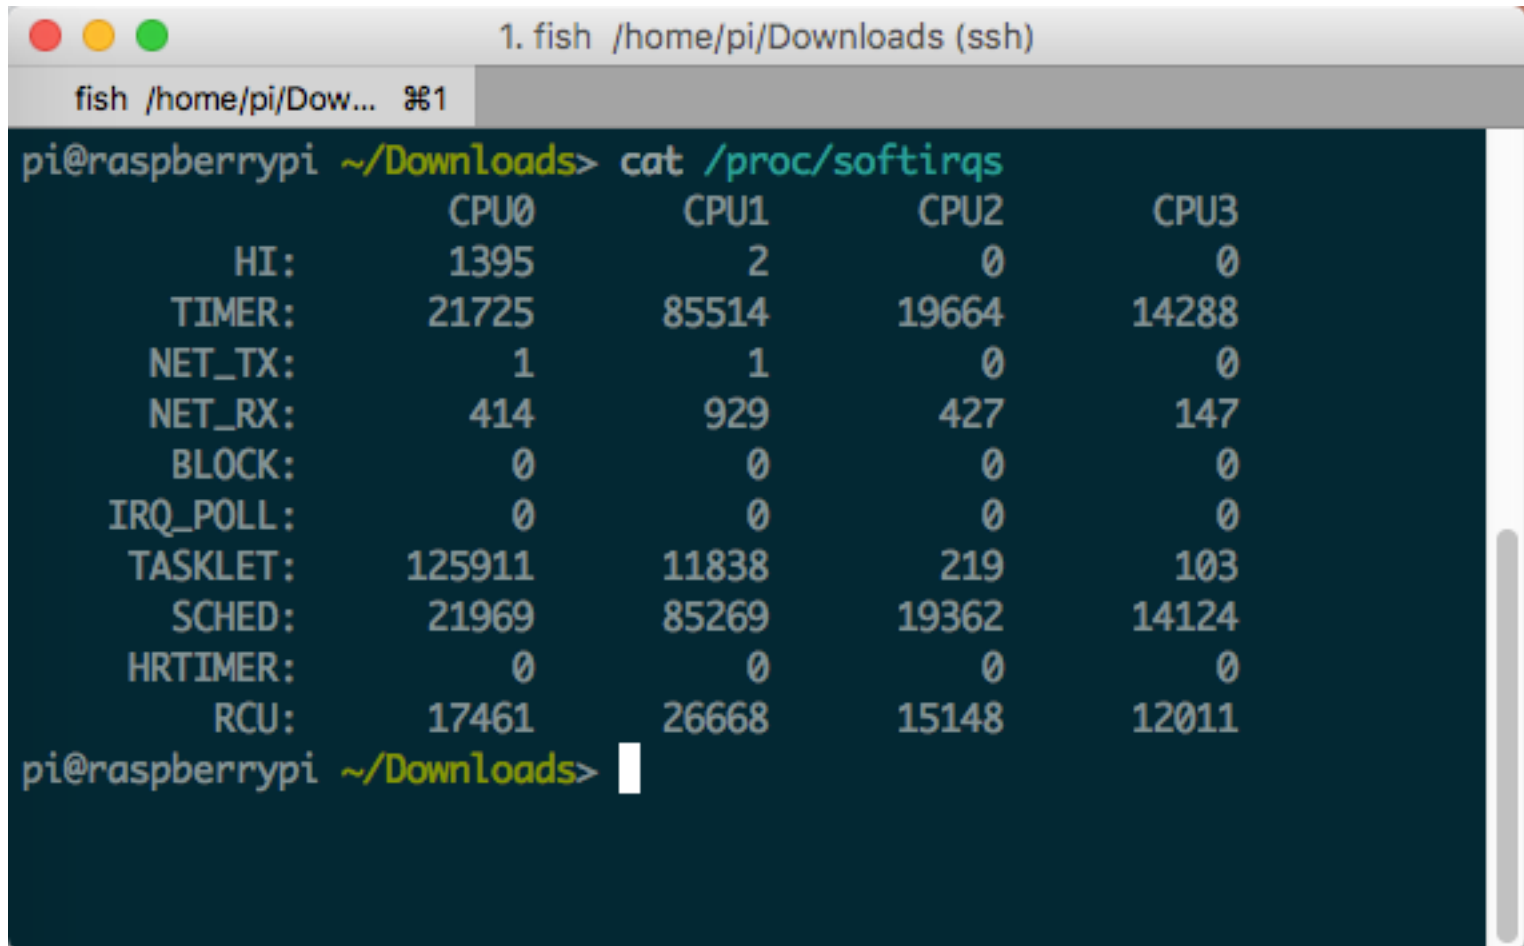


Software interrupt feature

- **Definition:** raised by software or execution
- **Nest:** softirq in Linux cannot be nested
- **Interrupt controller:** softirq does not have IRQ number
- **Mask:** softirq cannot be masked
- **Top/Bottom:** softirq is responsible the bottom-half work of the interrupt
- **Preemption:** softirq cannot preempt with each other



Software interrupt: cat /proc/softirqs



```
1. fish /home/pi/Downloads (ssh)
fish /home/pi/Dow... %1
pi@raspberrypi ~/Downloads> cat /proc/softirqs
```

	CPU0	CPU1	CPU2	CPU3
HI:	1395	2	0	0
TIMER:	21725	85514	19664	14288
NET_TX:	1	1	0	0
NET_RX:	414	929	427	147
BLOCK:	0	0	0	0
IRQ_POLL:	0	0	0	0
TASKLET:	125911	11838	219	103
SCHED:	21969	85269	19362	14124
HRTIMER:	0	0	0	0
RCU:	17461	26668	15148	12011

```
pi@raspberrypi ~/Downloads>
```

\$ watch -n 1 -d cat /proc/softirqs

```
watch /home/ksuo
Every 1.0s: cat /proc/softirqs      LinuxKernel2: Sun Feb 16 23:26:54 2020
```

	CPU0	CPU1	CPU2	CPU3
HI:	0	1	0	0
TIMER:	11005104	7993553	75625648	11401153
NET_TX:	0	1	8	1
NET_RX:	1211	703	778	10282130
BLOCK:	0	0	0	0
IRQ_POLL:	0	0	0	0
TASKLET:	14	2	0	0
SCHED:	9317734	7430993	73278709	10394289
HRTIMER:	0	0	0	0
RCU:	7450009	2604225	33874861	4333671




Hardware IRQ vs. Software IRQ

	Hardware IRQ	Software IRQ
Definition	raised by hardware	raised by software or execution
Nest	can be nested	cannot be nested
Interrupt controller	IRQ number is provided by APIC	softirq does not have IRQ number
Mask	can be masked	cannot be masked
Top/Bottom	hardirq handler ensures that it completes the task quickly	softirq is responsible the bottom-half work of the interrupt
Preemption	hardirq has priority and can preempt over softirq	softirq cannot preempt with each other

Interrupt Affinity

- IRQ affinity determines the CPU cores that are allowed to execute the processing for that IRQ
- IRQ affinity can be used to improve application performance (e.g., for multi CPU system like NUMA)

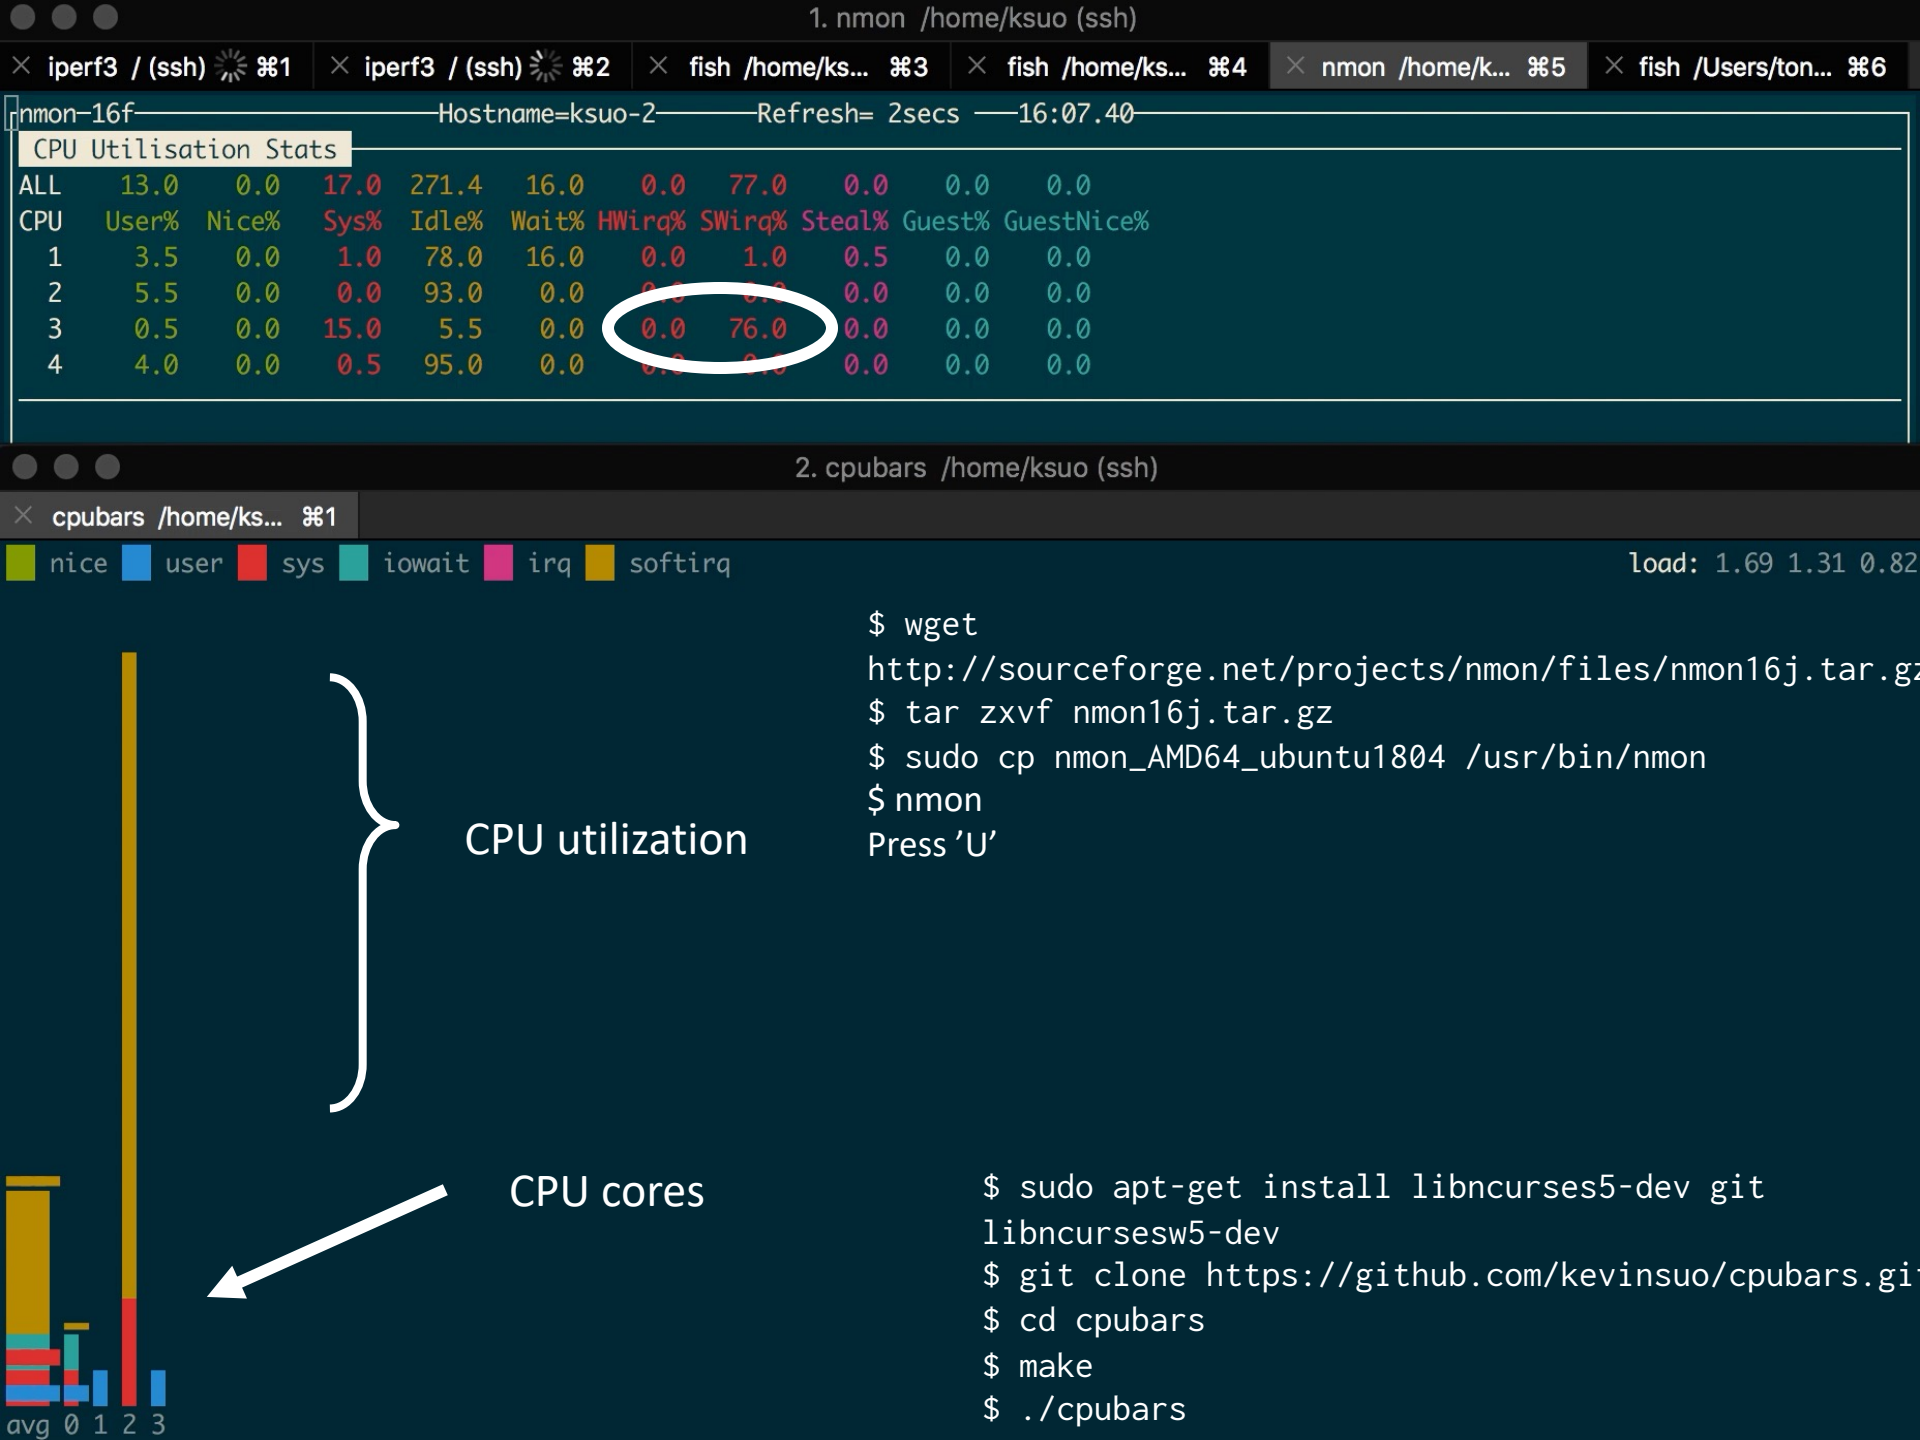


Interrupt Affinity

`watch -n 1 -d cat /proc/interrupts`

```
1. fish /home/pi/Downloads (ssh)
fish /home/pi/Dow... %1
pi@raspberrypi: /Downloads> cat /proc/interrupts
```

	CPU0	CPU1	CPU2	CPU3			
16:	0	0	0	0	bcm2836-timer	0 Edge	arch_timer
17:	21666	63962	23766	9395	bcm2836-timer	1 Edge	arch_timer
21:	0	0	0	0	bcm2836-pmu	9 Edge	arm-pmu
23:	3600	0	0	0	ARMCTRL-level	1 Edge	3f00b880.mailbox
24:	29	0	0	0	ARMCTRL-level	2 Edge	VCHIQ doorbell
46:	0	0	0	0	ARMCTRL-level	48 Edge	bcm2708_fb dma
48:	0	0	0	0	ARMCTRL-level	50 Edge	DMA IRQ
50:	0	0	0	0	ARMCTRL-level	52 Edge	DMA IRQ
51:	360	0	0	0	ARMCTRL-level	53 Edge	DMA IRQ
54:	4908	0	0	0	ARMCTRL-level	56 Edge	DMA IRQ
59:	0	0	0	0	ARMCTRL-level	61 Edge	bcm2835-auxirq
62:	139344	0	0	0	ARMCTRL-level	64 Edge	dwc_otg, dwc_otg_pcd, dwc_otg_hcd:usb1
86:	870	0	0	0	ARMCTRL-level	88 Edge	mmc0
87:	7237	0	0	0	ARMCTRL-level	89 Edge	uart-pl011
92:	101080	0	0	0	ARMCTRL-level	94 Edge	mmc1
169:	0	0	0	0	lan78xx-irqs	17 Edge	usb-001:004:01
FIQ:		usb_fiq					
IPI0:	0	0	0	0	CPU wakeup interrupts		
IPI1:	0	0	0	0	Timer broadcast interrupts		
IPI2:	7855	81594	13799	6452	Rescheduling interrupts		
IPI3:	10	6	9	8	Function call interrupts		
IPI4:	0	0	0	0	CPU stop interrupts		
IPI5:	7342	26013	5157	2072	IRQ work interrupts		
IPI6:	0	0	0	0	completion interrupts		
Err:	0						



Interrupt Affinity

- Read the IRQ affinity of specific hardIRQ

f = 1111, means all 4 CPU cores will receive the interrupt

```
# cat /proc/irq/32/smp_affinity  
f
```

- Set the IRQ affinity to specific hardIRQ

Interrupt from IRQ32 will only go to CPU core 1

```
# echo 1 >/proc/irq/32/smp_affinity  
# cat /proc/irq/32/smp_affinity  
1
```

- SoftIRQ will execute on the core which it is raised and cannot be set with affinity
 - SoftIRQ can be scaled with RPS/RFS in the kernel settings (irq migration)



Conclusion

- What is interrupt?
 - Interrupt vs Polling
 - Advanced programmable interrupt controller
 - Interrupt processing
- Interrupt types and affinity
 - Hardware and software interrupt
 - Interrupt affinity

