

# CS 3502

# Operating Systems

## Scheduling

**Kun Suo**

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

# Outline

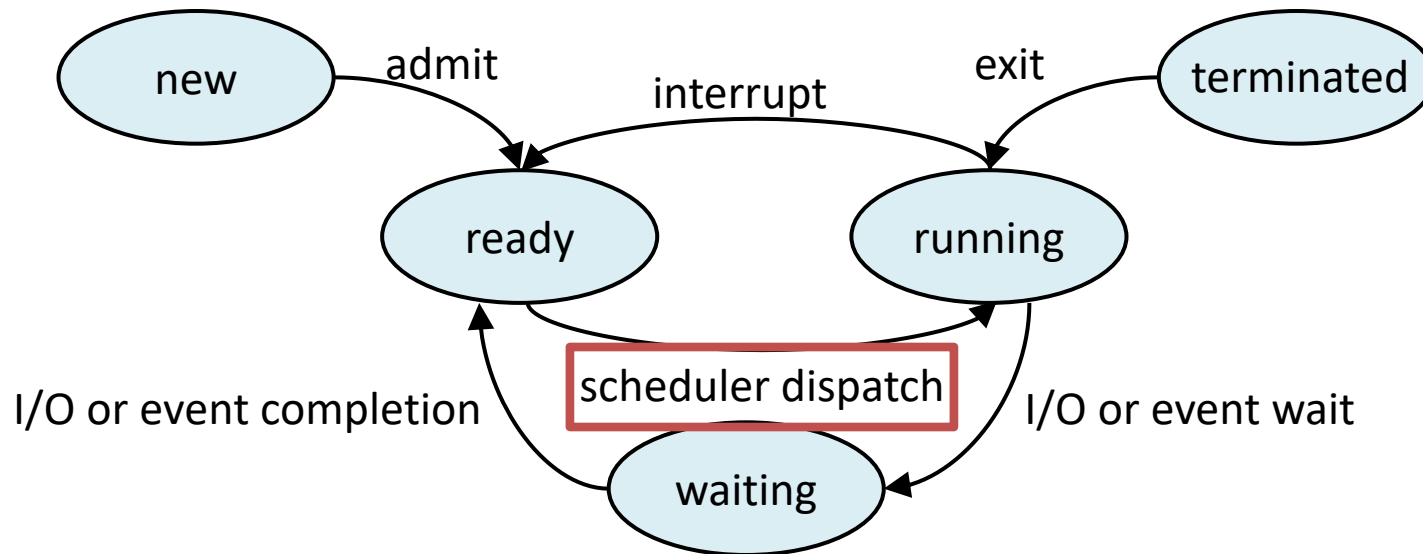
---

- Introduction to CPU scheduling
  - What is CPU scheduling
  - Why we need CPU scheduling
  - When scheduling happens
- Scheduling policies
  - FCFS, SJF, RR, Priority
  - Scheduling on multiple CPUs



# What is CPU scheduling?

- The five-state process model

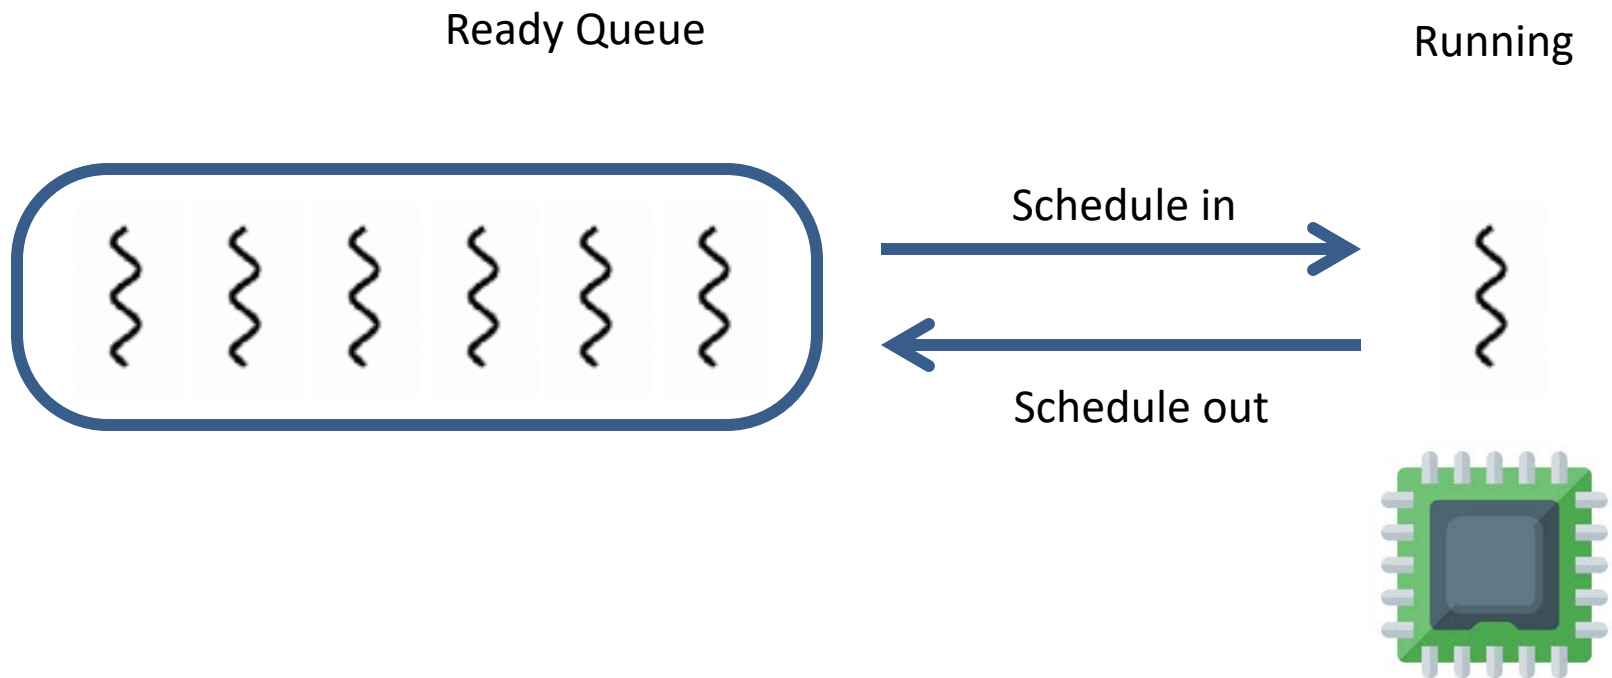


## CPU scheduling

Selects from among the processes/threads that are ready to execute, and allocates the CPU to it


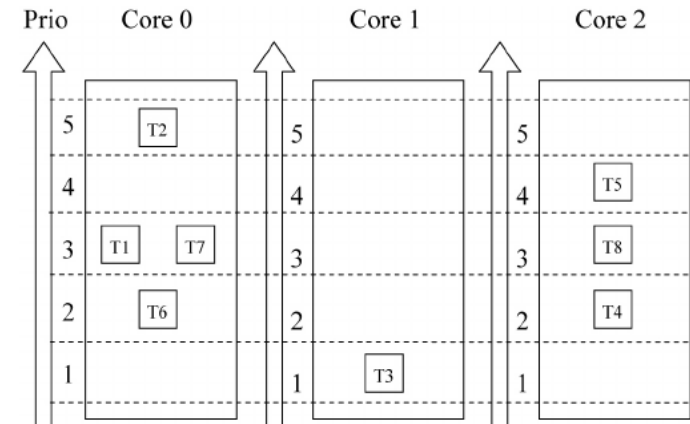

# What is CPU scheduling?

---



# Why CPU scheduling?

- In support of **multiprogramming**
  - uniprocessor systems
    - ▶ Time-sharing processor
  - multiprocessor systems
    - ▶ Efficiently distributing tasks
  - Real-time systems
    - ▶ Reliably guaranteeing deadlines



# Why CPU scheduling? – Different Goals


---

- All systems
  - Fairness - giving each process a fair share of the CPU
  - Policy enforcement - seeing that stated policy is carried out
  - Balance - keeping all parts of the system busy



# Why CPU scheduling? – Different Goals

- Batch systems
  - Throughput - maximize jobs per hour
  - Turnaround time - minimize time between submission and termination
  - CPU utilization - keep the CPU busy all the time




National lab:  
Scientific research  
Energy, nuclear, climate, ...



# Why CPU scheduling? – Different Goals

---

- Interactive systems
  - Response time - respond to requests quickly
  - Proportionality - meet users' expectations

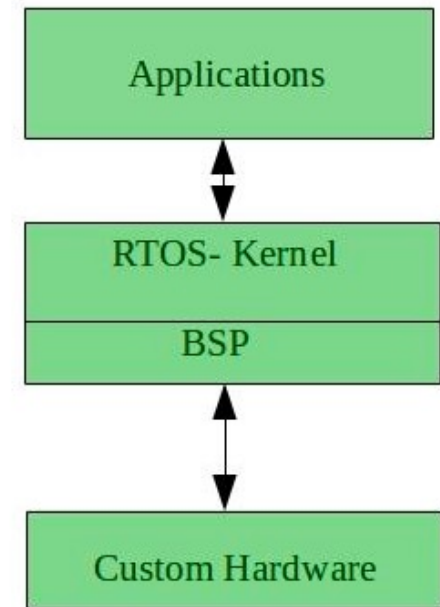




# Why CPU scheduling? – Different Goals

---

- Real-time systems
  - Meeting deadlines - avoid losing data
  - Predictability - avoid quality degradation in multimedia systems



# Why CPU scheduling? – Different Goals

---

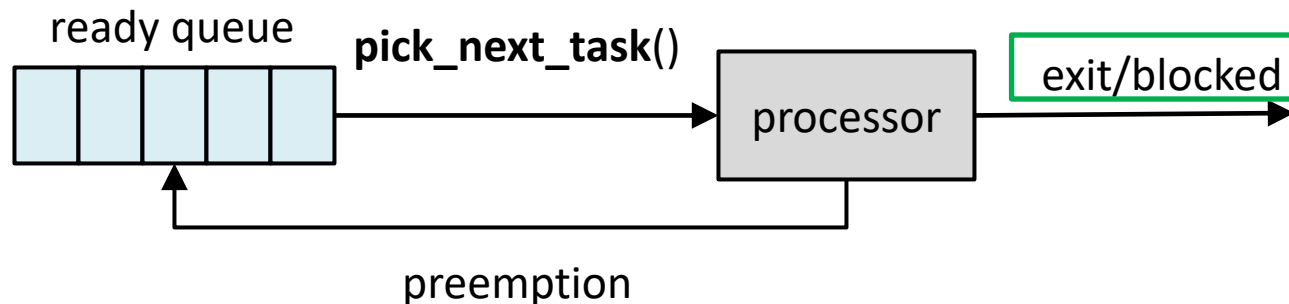
- All systems
  - Fairness - giving each process a fair share of the CPU
  - Policy enforcement - seeing that stated policy is carried out
  - Balance - keeping all parts of the system busy
- Batch systems
  - Throughput - maximize jobs per hour
  - Turnaround time - minimize time between submission and termination
  - CPU utilization - keep the CPU busy all the time
- Interactive systems
  - Response time - respond to requests quickly
  - Proportionality - meet users' expectations
- Real-time systems
  - Meeting deadlines - avoid losing data
  - Predictability - avoid quality degradation in multimedia systems



# When scheduling happens?

---

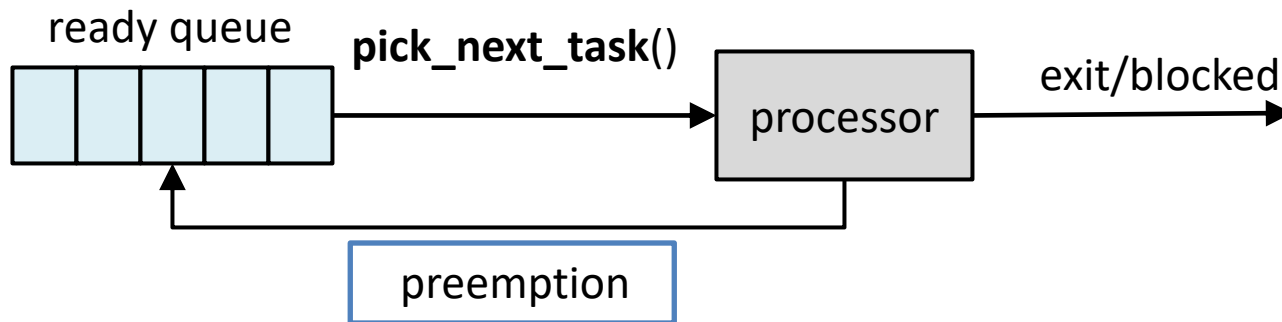
- Non-preemptive
  - Scheduling only when current process **terminates** or **gives up** control



# When scheduling happens?

---

- Preemptive
  - Processes can be **forced** to give up control

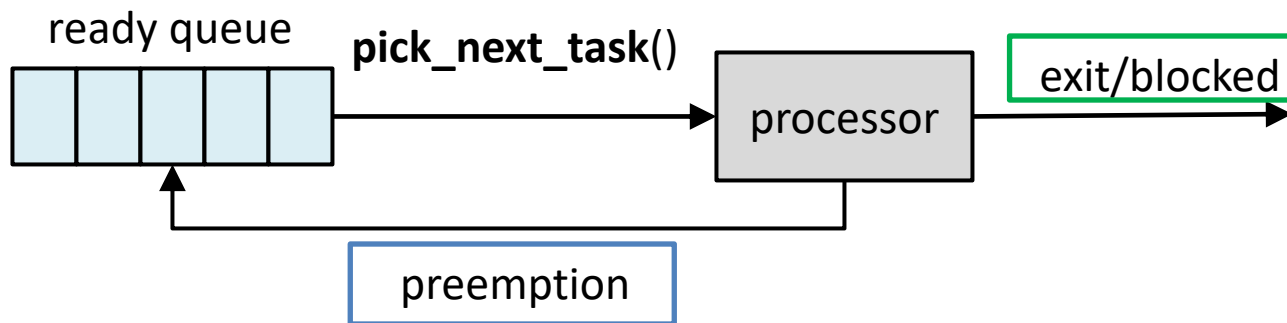


# When scheduling happens?

- CPU scheduling may take place at
  - Clock interrupts
  - I/O interrupts
  - I/O completion
  - Termination

} preemptive

} non-preemptive



# Outline

---

- Introduction to CPU scheduling
  - What is CPU scheduling
  - Why we need CPU scheduling
  - When scheduling happens
- Scheduling policies
  - FCFS, SJF, RR, Priority
  - Scheduling on multiple CPUs



# Scheduling Policies

Not exist best scheduling.  
It depends on your goals.

- Batch Systems
  - First-Come First-Serve (FCFS)
  - Shortest Job First
  - Shortest Remaining Time Next
- Interactive Systems
  - Round-Robin
  - Priority Scheduling
  - Multiple Queues
  - Shortest Process Next
  - Guaranteed Scheduling
  - Lottery Scheduling
- Real-time Systems
  - Rate Monotonic Scheduling
  - Earliest Deadline First Scheduling


Determine the next  
ready task to run



# Turnaround time = End time – Arrival time

---

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5





# Response time = Start time – Arrival time

---

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5



0

1  $\longleftrightarrow$  8

2  $\longleftrightarrow$  12

3  $\longleftrightarrow$  21



# Calculate the Turnaround time and Response time for scheduling

---

- Step 1: draw the timeline based on the scheduling policy
- Step 2: calculate turnaround time and response time for each task



# First-Come, First-Serve (FCFS)

Turnaround time = End time – Arrival time  
Response time = Start time – Arrival time

- CPU schedules the task that arrived earliest, non-preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5



Average turnaround time = ?



# First-Come, First-Serve (FCFS)

---

- CPU schedules the task that arrived earliest, non-preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5



Average turnaround time =  $((8-0)+(12-1)+(21-2)+(26-3)) / 4 = 15.25$



# First-Come, First-Serve (FCFS)

Turnaround time = End time – Arrival time  
Response time = Start time – Arrival time

- CPU schedules the task that arrived earliest, non-preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5



Average response time = ?



# First-Come, First-Serve (FCFS)

---

- CPU schedules the task that arrived earliest, non-preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5



$$\text{Average response time} = (0 + (8 - 1) + (12 - 2) + (21 - 3)) / 4 = 8.75$$




# Shortest Job First (SJF)

---

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive




# Shortest Job First (SJF)

- CPU schedules the task with the shortest remaining time


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive



$P_2$  has to wait  
 $P_1$  due to non-  
preemptive

pool






# Shortest Job First (SJF)

- CPU schedules the task with the shortest remaining time


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive



$P_3$  has to wait  
 $P_1$  due to non-  
preemptive

pool




# Shortest Job First (SJF)

- CPU schedules the task with the shortest remaining time


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive



$P_4$  has to wait  
 $P_1$  due to non-  
preemptive

pool




# Shortest Job First (SJF)

- CPU schedules the task with the shortest remaining time


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive



$P_2$  is selected  
due to shortest  
job

pool




# Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive



pool




# Shortest Job First (SJF)

- CPU schedules the task with the shortest remaining time


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive



$P_4$  is selected  
due to shortest  
job

pool



# Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive



pool



# Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive



pool



# Shortest Job First (SJF)


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive



pool





# Shortest Job First (SJF)

Turnaround time = End time – Arrival time

Response time = Start time – Arrival time

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive



Average turnaround time = ?



# Shortest Job First (SJF)

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive



$$\text{Average turnaround time} = ((8-0)+(12-1)+(26-2)+(17-3)) / 4 = 14.25$$



# Shortest Job First (SJF)

Turnaround time = End time – Arrival time

Response time = Start time – Arrival time

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive



Average response time = ?



# Shortest Job First (SJF)

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

non-preemptive



$$\text{Average response time} = (0 + (8 - 1) + (17 - 2) + (12 - 3)) / 4 = 7.75$$




# Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

preemptive



pool




# Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

preemptive



If allowing preemption,  
 $P_2$  is selected due to  
shortest job

pool




# Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

preemptive



pool

$P_1(7)$

# Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

preemptive



P2 remain =3

pool






# Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

preemptive



pool



# Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

preemptive



$P_4$  is selected  
due to shortest  
job

pool

$P_1(7)$

$P_3(9)$

$P_4(5)$

# Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

preemptive



pool



# Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

preemptive



pool




# Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

preemptive



pool



# Shortest Job First (SJF) with preemption


- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

preemptive



pool



# Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

Turnaround time = End time – Arrival time  
Response time = Start time – Arrival time

preemptive



Average turnaround time = ?



# Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

Turnaround time = End time – Arrival time  
Response time = Start time – Arrival time

preemptive



$$\text{Average turnaround time} = ((17-0)+(5-1)+(10-3)+(26-2)) / 4 = 13$$





# Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

Turnaround time = End time – Arrival time  
Response time = Start time – Arrival time

preemptive



Average response time = ?



# Shortest Job First (SJF) with preemption

- CPU schedules the task with the shortest remaining time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

Turnaround time = End time – Arrival time

Response time = Start time – Arrival time

preemptive



$$\text{Average response time} = (0 + (1 - 1) + (5 - 3) + (17 - 2)) / 4 = 4.25$$





# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool





# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool




# Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool



$P_1$  is put at the end of queue after scheduled out





# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool



$q=4$



0 4 8

# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool





# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool







# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool





# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool





# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool





# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool





# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool





# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool




# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool

$P_3 (1)$


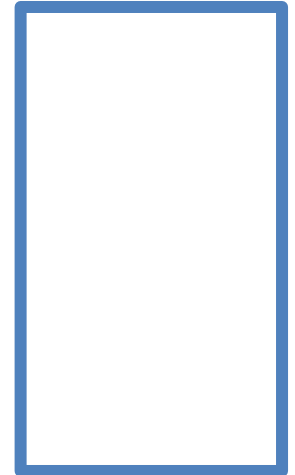


# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool






# Round Robin (RR)

Turnaround time = End time – Arrival time

Response time = Start time – Arrival time

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5




Average turnaround time = ?



# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5




$$\text{Average turnaround time} = ((20-0)+(8-1)+(26-2)+(25-3)) / 4 = 18.25$$

# Round Robin (RR)

Turnaround time = End time – Arrival time  
Response time = Start time – Arrival time

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5




Average response time = ?



# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5



$$\text{Average response time} = (0 + (4 - 1) + (8 - 2) + (12 - 3)) / 4 = 4.5$$

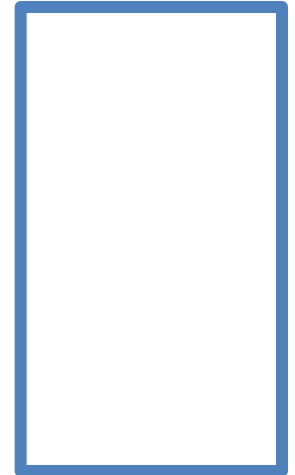


# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool



$q=5$



# Round Robin (RR)

Select the task at the beginning

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5


pool

$P_2 (4)$

$P_3 (9)$

$P_4 (5)$

$q=5$



# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool

$P_3 (9)$

$P_4 (5)$

$P_1 (3)$

$q=5$



P1 is put at the end of queue after scheduled out




# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool



$q=5$






# Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool



$q=5$




# Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool



$q=5$




# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool



$q=5$



# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool

$P_1 (3)$

$P_3 (4)$

$q=5$




# Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool



$q=5$




# Round Robin (RR)


- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool



$q=5$




# Round Robin (RR)

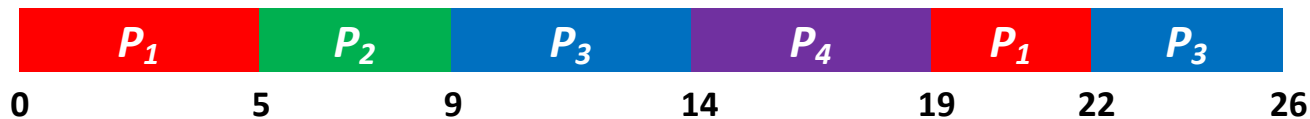
- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

pool



$q=5$



# Round Robin (RR)


Turnaround time = End time – Arrival time

Response time = Start time – Arrival time

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

$q=5$



Average turnaround time = ?






# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

$q=5$



$$\text{Average turnaround time} = ((22-0)+(9-1)+(26-2)+(19-2)) / 4 = 17.5$$



# Round Robin (RR)


Turnaround time = End time – Arrival time

Response time = Start time – Arrival time

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

$q=5$



Average response time = ?




# Round Robin (RR)

- Like FCFS, but with limited time slices, preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

$q=5$




$$\text{Average response time} = (0 + (5 - 1) + (9 - 2) + (14 - 3)) / 4 = 5.5$$




# Priority Scheduling

- CPU schedules the highest priority (smaller value) first, FCFS within the same priority



# Priority Scheduling

- CPU schedules the highest priority (smaller value) first, FCFS within the same priority



# Comparison

---

	Turnaround time	Response time
FCFS	15.25	8.75
SJF-preemptive	<b>13</b>	<b>4.25</b>
RR (q=5)	17.5	5.5
Priority scheduling	N/A	N/A


	Throughput	Response time	Starvation
FCFS	TBD	TBD	No
SJF-preemptive	High	Good	Yes
RR	Can be low	Good	No
Priority scheduling	Can be high	Can be good	Can remove



# Challenges on Emerging Hardware and Applications

---

- Multi-processor → Single queue



Multiprocessor = more powerful processor

**pick\_next\_task()** will be the bottleneck

Linux 2.4.x

## Pros:

1. Easy to implement
2. Perfect load balancing

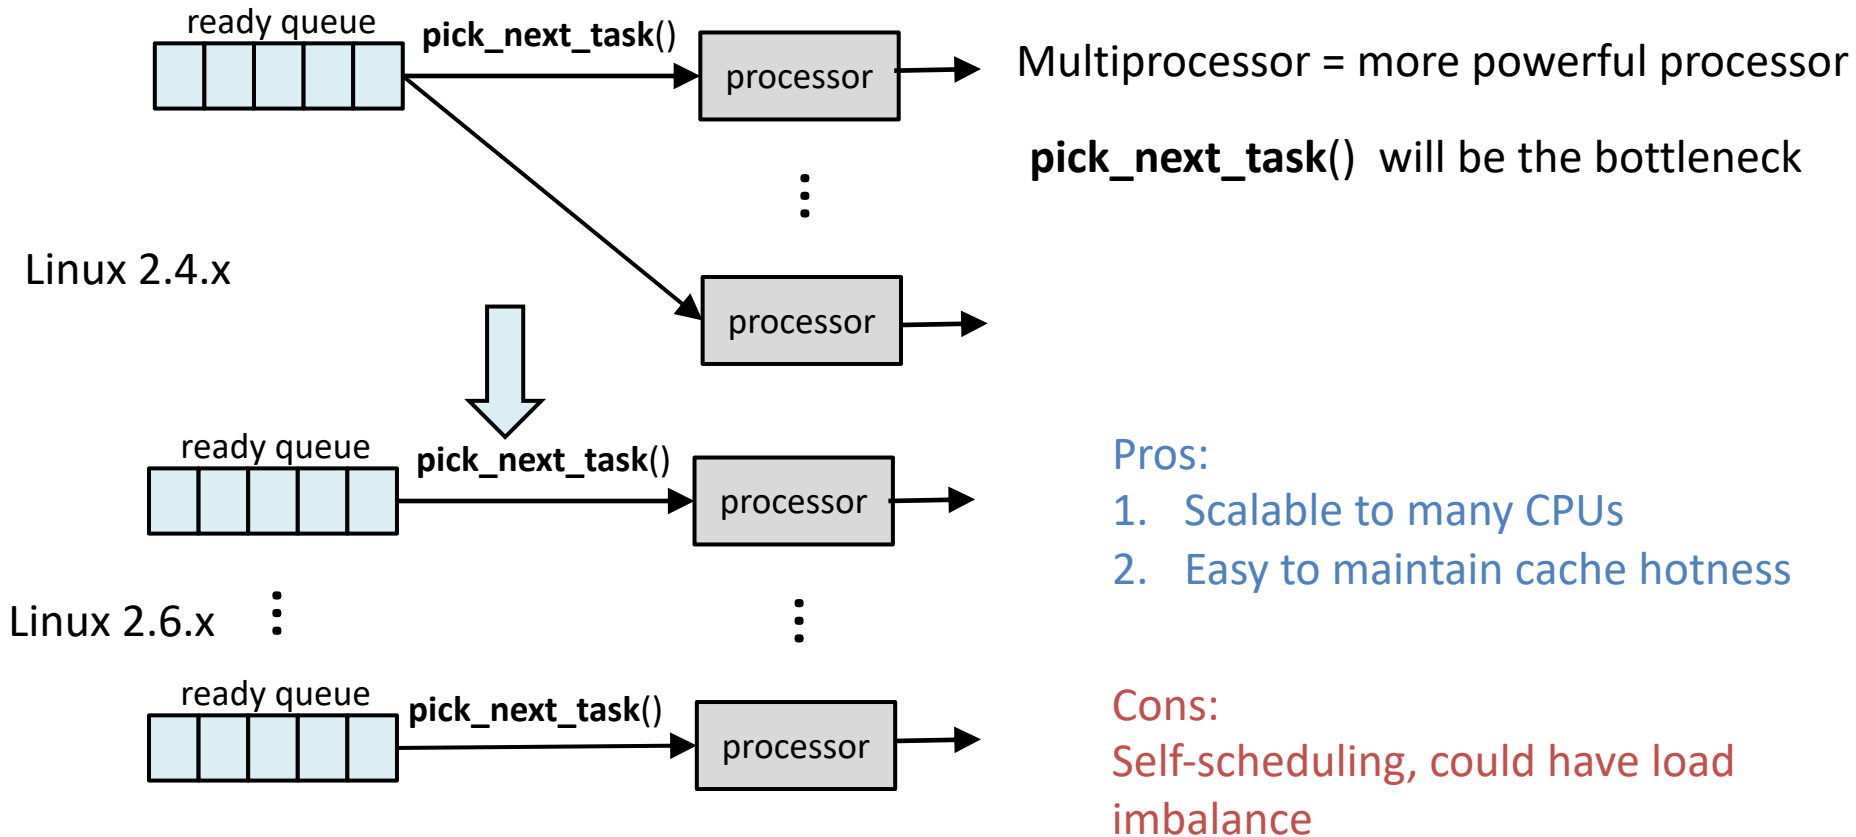
## Cons:

1. Scalability issues due to centralized synchronization
2. High overhead and low efficiency
3. Hard to maintain cache hotness due to global scheduling



# Challenges on Emerging Hardware and Applications

- Multi-processor → Many queues






# Overcome Load Imbalance

- Push model

Every a while, a kernel thread checks load imbalance and move threads


Made by OS




- Pull model

Whenever a queue becomes empty, steal a thread from non-empty queues


Made by local queue. Both are widely used



# Load balance on SMP vs. NUMA



symmetric multiprocessing (SMP):  
The distance to memory is the same



Non-uniform memory access (NUMA):  
The distance to memory is different

## Multi processor/core scheduling

The scheduling policy not only considers the fairness, throughput, etc., but also needs to consider the **hardware architecture** (e.g., locality)

# Conclusion

---

- Introduction to CPU scheduling
  - What is CPU scheduling
  - Why we need CPU scheduling
  - When scheduling happens
- Scheduling policies
  - FCFS, SJF, RR, Priority
  - Scheduling on multiple CPUs



# Why CPU scheduling?

---

- It is (maybe) the most important part in a OS
  - Why some OS seems to be faster than others?
  - Why I do not see performance improvement when upgrading to a 16-core computer?




# Scheduling Goals: A Different Point of View

---

- User oriented → minimize
  - Response time (wait time): the time that the first response is received (interactivity)
  - Turnaround time: the time that the task finishes
  - Predictability: variations in different runs
- System oriented → maximize
  - Throughput: # of tasks that finish per time unit
  - Utilization: the percentage of time the CPU is busy
  - Fairness: avoid starvation



# CPU-bound tasks vs. I/O-bound tasks



```
#include <stdio.h>

int main() {
    int a = 1;

    while ( a < 1000000 )
    {
        a++;
    }

    return 0;
}
```

a CPU-bound/CPU-intensive process

```
#include <stdio.h>

int main()
{
    char name[20];
    printf("Enter name: ");

    scanf("%s", name);
    printf("Your name is %s.", name);

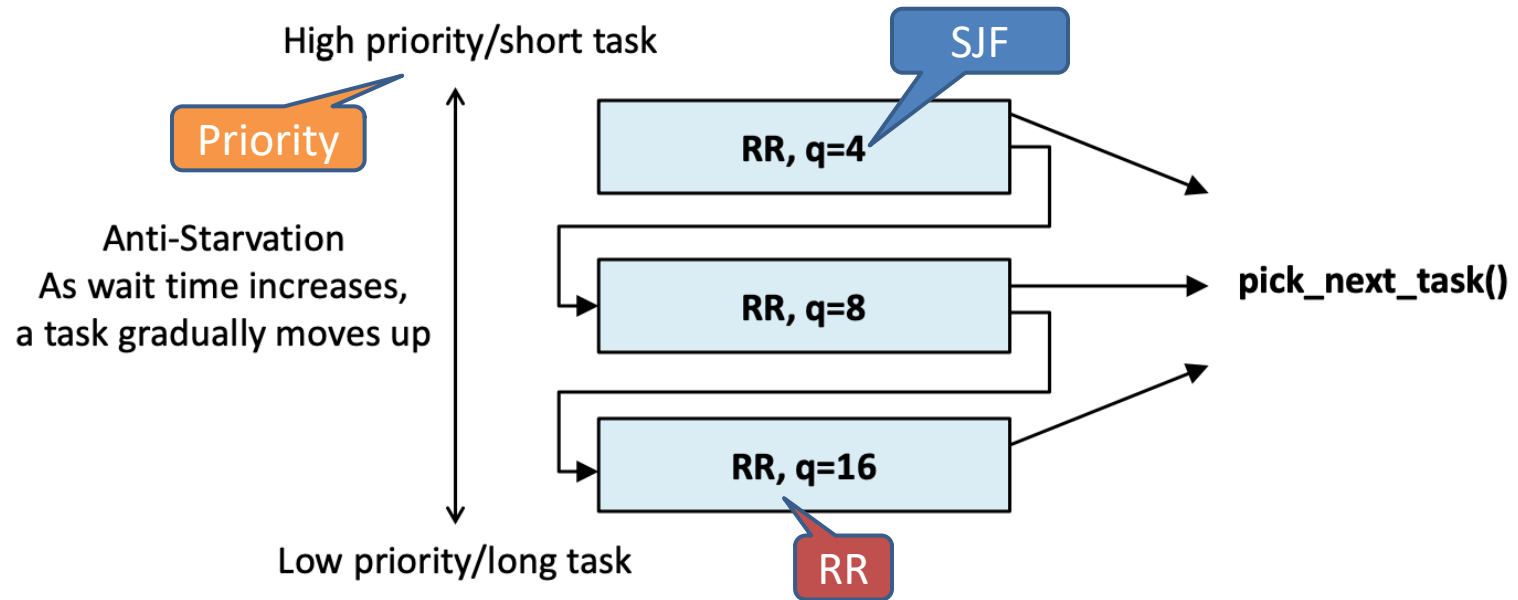
    return 0;
}
```

a I/O-bound/I/O-intensive process

I/O is when a process enters the blocked state waiting for an external device to complete its work

# Multilevel Feedback Queue

	Throughput	Response time	Starvation
FCFS	TBD	TBD	No
SJF-preemptive	High	Good	Yes
RR	Can be low	Good	No
Priority scheduling	Can be high	Can be good	Can remove



**Windows XP, Mac OS X, Linux 2.6.22 and before**

# Scheduling in Linux

- Linux Completely Fair Scheduler (CFS)
  - CFS = RR + SJF + Priority + smart data structure + hardware consideration
  - One ready queue for one processor
  - Red-black tree based ready queue

Print it  
in project 1

**vruntime:** how much time the  
task consumes the CPU cycles

**pick\_next\_task()**  
Select the leftmost task  
which has the least vruntime

