

Kennesaw State University

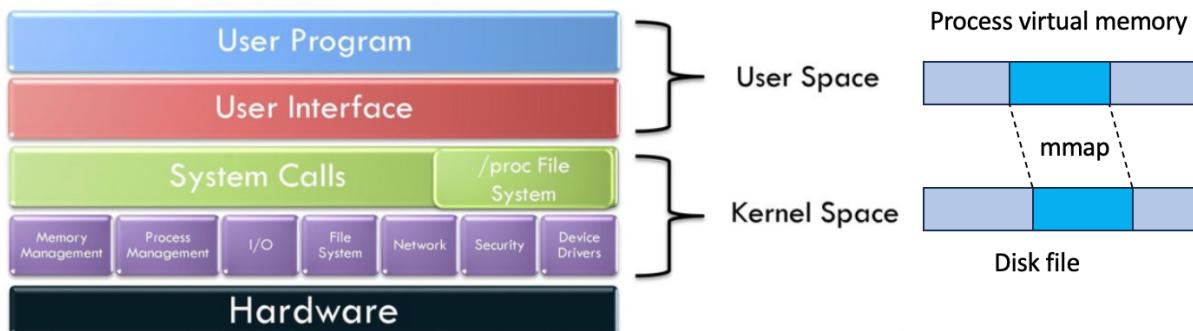
CSE 3502 Operating Systems

Project 3 - The /Proc File Systems and mmap

Instructor: Kun Suo

Points Possible: 100

Difficulty: ★★★★★



Part 1: Create a helloworld kernel module (20 pts)

The following code is a complete helloworld module. Name it as *new_module.c*

<https://github.com/kevinsuo/CS3502/blob/master/project-4-1.c>

```
#include <linux/module.h>
#include <linux/kernel.h>

int init_new_module(void)
{
    printk(KERN_INFO "Hello, world!\n");
    return 0;
}

void exit_new_module(void) {
    printk(KERN_INFO "Goodbye, world!\n");
}

module_init(init_new_module);
module_exit(exit_new_module);
```

The module defines two functions. `init_module` is invoked when the module is loaded into the kernel and `exit_module` is called when the module is removed from the kernel.

`module_init` and `module_exit` are special kernel macros to indicate the role of these two functions.

Use the following makefile to compile the module. Name it as ***Makefile***
<https://github.com/kevinsuo/CS3502/blob/master/project-4-1-Makefile>

Note that here ***new_module.o*** is the output after compiling.

```
-----  
obj-m += new_module.o  
all:  
    sudo make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules  
clean:  
    sudo make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean  
-----
```

Compile the `new_module.c` file using make command.

`$ make`

To insert the module into the Linux kernel:

`$ sudo insmod new_module.ko`

Use the following command to verify the module has been loaded:

`$ lsmod`

Module	Size	Used by
new_module	16384	0
btrfs	1179648	0
xor	24576	1 btrfs
zstd_compress	163840	1 btrfs
raid6_pq	114688	1 btrfs
ufs	81920	0
qnx4	16384	0
hfsplus	110592	0
hfs	61440	0
minix	36864	0
ntfs	106496	0
msdos	20480	0
jfs	188416	0
xfs	1245184	0
libcrc32c	16384	2 btrfs,xfs
crc10dif_pclmul	16384	1

To remove the module from the kernel:

`$ sudo rmmod new_module`

When you insert or remove the module, corresponding information will be printed out under the dmesg.

```
ksuo@ksuo-VirtualBox ~/hw4> dmesg
[79806.620385] Hello, world!
[79808.949265] Goodbye, world!
```

Part 2: Create an entry in the /proc file system for user level read and write (30 pts)

Write a new kernel module following steps in Part 1. This module creates an entry in the /proc file system. Use the following code skeleton to write the module:

<https://github.com/kevinsuo/CS3502/blob/master/project-4-2.c>

```
#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/proc_fs.h>
#include <linux/string.h>
#include <linux/vmalloc.h>
#include <linux/slab.h>
#include <linux/uaccess.h>

#define MAX_LEN          4096
static struct proc_dir_entry *proc_entry;

ssize_t read_proc(struct file *f, char *user_buf, size_t count, loff_t *off)
{
    //output the content of info to user's buffer pointed by page
    return count;
}

ssize_t write_proc(struct file *f, const char *user_buf, size_t count, loff_t *off)
{
    //copy the written data from user space and save it in info
    return count;
}

struct file_operations proc_fops = {
    read: read_proc,
    write: write_proc
};

int init_module( void )
{
    int ret = 0;
    //create the entry named myproc and allocated memory space for the proc entry
    printk(KERN_INFO "test_proc created.\n");

    return ret;
}

void cleanup_module( void )
{
    //remove the entry named myproc and free info space
```

}

Step 1: create an entry in proc file system named *myproc* when the kernel module is loaded; this entry *myproc* will be deleted when the kernel mode is deleted. You can use [*\\$ ls /proc/*](#) to check whether it is existed. (Hint: *proc_create()* and *remove_proc_entry()* are needed.)

Step 2: implement *read_proc* and *write_proc* function to read/write the proc file entry in Step 1. You need to add codes for allocating memory in *init_module* and releasing the memory in *cleanup_module* for the proc file entry. (Hint: *copy_to_user()* is needed for the read and *copy_from_user()* is needed for write.)

To test your results, load the kernel module and there should be a new entry created under */proc*. Use *cat* and *echo* to verify and change the content of the new entry.

```
ksuo@ksuo-VirtualBox ~/hw4-2> sudo insmod my_proc.ko
[sudo] password for ksuo:
ksuo@ksuo-VirtualBox ~/hw4-2> ls /proc/
1/    1283/  1471/  23/   39/   497/  683/      diskstats  pagetypeinfo
10/   1284/  15/    24/   4/    50/   686/      dma        partitions
11/   1285/  1502/  249/  40/   500/   7/       driver/   pressure/
1114/  1287/  1504/  250/  41/   502/  702/      execdomains sched_debug
1119/  1288/  1522/  251/  42/   503/  748/      fb        schedstat
1124/  1295/  154/   254/  423/  506/   8/       filesystems scsi/
1137/  13/    155/   26/   43/   509/  803/      fs/       self@
1142/  1303/  156/   27/   438/  52/   804/      interrupts slabinfo
1144/  1304/  157/   275/  44/   523/  817/      iomem     softirqs
1168/  1314/  158/   276/  440/  53/   821/      ioports   stat
1174/  1315/  159/   28/   449/  532/  823/      irq/     swaps
1189/  1321/  1598/  282/  45/   533/  891/      kallsyms sys/
1190/  1323/  16/    29/   453/  534/   9/       kcore     sysrq-trigger
12/    1325/  161/   292/  454/  537/  905/      keys      sysvipc/
1205/  1327/  162/   3/    455/  54/   912/      key-users thread-self@
1209/  1331/  1684/  30/   457/  56/   931/      kmsg      timer_list
1210/  1332/  1685/  32/   46/   571/  950/      kpagecgroup tty/
1212/  1337/  1695/  321/  47/   572/  954/      kpagecount uptime
1218/  1338/  17/    33/   48/   59/   975/      kpageflags version
1229/  1372/  173/   331/  482/  6/    acpi/      loadavg   vmallocinfo
1233/  1383/  1763/  335/  484/  60/   asound/    locks     vmstat
1241/  1384/  18/    336/  485/  606/  buddyinfo mdstat   zoneinfo
1245/  14/    19/   34/   489/  607/  bus/       meminfo
1252/  1412/  192/  35/   49/   608/  cgroups   misc
1261/  1415/  193/  36/   490/  61/   cmdline   modules
1266/  1439/  2/    360/  491/  614/  consoles  mounts@
1271/  1442/  20/   37/   493/  634/  cpuinfo   mtrr
1275/  1446/  21/   38/   494/  659/  crypto    myproc
1279/  1460/  22/   384/  496/  677/  devices   net@
```

You can use the following to test the read or write on the entry of proc file system. Here the root user is needed.

Expected output:

Write “your name” into /proc/myproc. For instance, a student named “Sisi”

```
root@ksuo-VirtualBox /h/k/hw4-2# echo Sisi > /proc/myproc
```

Read /proc/myproc and printout its content:

```
root@ksuo-VirtualBox /h/k/hw4-2# cat /proc/myproc
Sisi
```

Part 3: Exchange data between the user and kernel space via mmap (50 pts)

Write a kernel module that create an entry in the /proc file system. The new entry cannot be directly read or written using cat and echo commands. Instead, map the new entry to a user space memory area so that user-level processes can read from and write to the kernel space via mmap. The skeleton of the kernel module is given below:

<https://github.com/kevinsuo/CS3502/blob/master/project-4-3-1.c>

```
#include <linux/module.h>
#include <linux/list.h>
#include <linux/init.h>
#include <linux/kernel.h>
#include <linux/types.h>
#include <linux/kthread.h>
#include <linux/proc_fs.h>
#include <linux/sched.h>
#include <linux/mm.h>
#include <linux/fs.h>
#include <linux/slab.h>

static struct proc_dir_entry *tempdir, *tempinfo;
static unsigned char *buffer;
static unsigned char array[12]={0,1,2,3,4,5,6,7,8,9,10,11};

static void allocate_memory(void);
static void clear_memory(void);
static int my_map(struct file *filp, struct vm_area_struct *vma);

static const struct file_operations myproc_fops = {
    .mmap    = my_map,
};

static int my_map(struct file *filp, struct vm_area_struct *vma)
{
    // map vma of user space to a continuous physical space
    return 0;
}

static int init_myproc_module(void)
{
```

```

tempdir=proc_mkdir("mydir", NULL);
if(tempdir == NULL) {
    printk("mydir is NULL\n");
    return -ENOMEM;
}

tempinfo = proc_create("myinfo", 0, tempdir, &myproc_fops);
if(tempinfo == NULL) {
    printk("myinfo is NULL\n");
    remove_proc_entry("mydir",NULL);
    return -ENOMEM;
}
printk("init myproc module successfully\n");

allocate_memory();

return 0;
}

static void allocate_memory(void)
{
    /* allocation memory */
    buffer = (unsigned char *)kmalloc(PAGE_SIZE,GFP_KERNEL);
    /* set the memory as reserved */
    SetPageReserved(virt_to_page(buffer));
}

static void clear_memory(void)
{
    /* clear reserved memory */
    ClearPageReserved(virt_to_page(buffer));
    /* free memory */
    kfree(buffer);
}

static void exit_myproc_module(void)
{
    clear_memory();
    remove_proc_entry("myinfo", tempdir);
    remove_proc_entry("mydir", NULL);
    printk("remove myproc module successfully\n");
}

module_init(init_myproc_module);
module_exit(exit_myproc_module);
MODULE_LICENSE("GPL");

```

The above code will create an entry **/proc/mydir/myinfo** under the proc file system. You are required to implement the **my_map** function to map one piece of memory (**char array[12]**) into user space. Then write a user space program using mmap to visit the memory space of the proc file and print the data in that memory area. You can use the following skeleton:

<https://github.com/kevinsuo/CS3502/blob/master/project-4-3-2.c>

```

#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <linux/fb.h>
#include <sys/mman.h>
#include <sys/ioctl.h>

```

```

#define PAGE_SIZE 4096

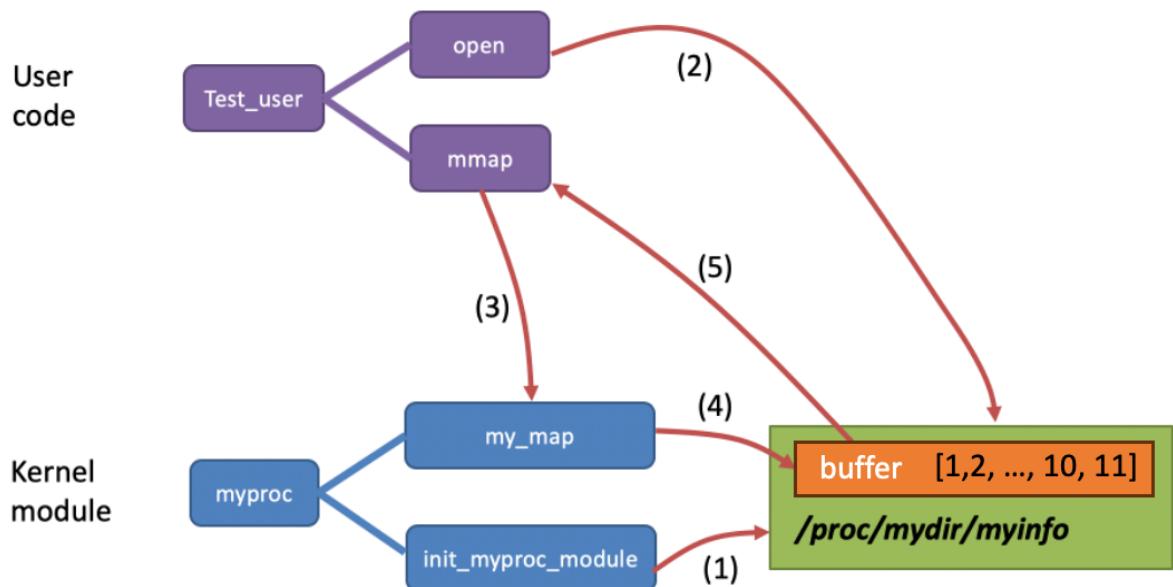
int main(int argc , char *argv[])
{
    int fd;
    int i;
    unsigned char *p_map;

    /* open proc file */
    fd = open("/proc/mydir/myinfo", O_RDWR);
    if(fd < 0) {
        printf("open fail\n");
        exit(1);
    }else {
        printf("open successfully\n");
    }

    // map p_map to the proc file and grant read & write privilege
    // read data from p_map
    // unmap p_map from the proc file

    return 0;
}

```



The above figure shows the entire workflow:

- (1) Kernel module create a proc file: **/proc/mydir/myinfo**
- (2) User process open the created proc file
- (3) User process calls mmap function, which further executed my_map defined in the kernel
- (4) my_map() then maps one piece of memory into user space (e.g., buffer) and puts some data inside
- (5) User process visits this piece of memory and prints the data out.

Expected output:

For instance, a student named Sisi should upload a screenshot like:

```
ksuo@ksuo-VirtualBox ~/hw4-3> sudo ./test_user.o
open successfully by Sisi
0
1
2
3
4
5
6
7
8
9
10
11
Printed by Sisi
```

Submission requirements:

Submit your assignment file through D2L using the appropriate link.

The submission must include the source code, and a report describe your code logic (including part 1/2/3). Output screenshot of your code should be included in the report.