

Kennesaw State University


CSE 3502 Operating Systems

Project 3 - The /Proc File Systems and mmap

Instructor: Kun Suo

Points Possible: 100

Difficulty: ★★★★★



Part 1: Create a helloworld kernel module (20 pts)

The following code is a complete helloworld module. Name it as *new_module.c*

<https://github.com/kevinsuo/CS3502/blob/master/project-4-1.c>

```
#include <linux/module.h>
#include <linux/kernel.h>

int init_new_module(void)
{
    printk(KERN_INFO "Hello, world!\n");
    return 0;
}

void exit_new_module(void)
{
    printk(KERN_INFO "Goodbye, world!\n");
}

module_init(init_new_module);
module_exit(exit_new_module);
```

The module defines two functions. `init_module` is invoked when the module is loaded into the kernel and `exit_module` is called when the module is removed from the kernel.

module_init and module_exit are special kernel macros to indicate the role of these two functions.

Use the following makefile to compile the module. Name it as **Makefile**
<https://github.com/kevinsuo/CS3502/blob/master/project-4-1-Makefile>

Note that here **new_module.o** is the output after compiling.

```
-----  
obj-m += new_module.o  
all:  
    sudo make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules  
clean:  
    sudo make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean  
-----
```

Compile the new_module.c file using make command.


\$ make

To insert the module into the Linux kernel:

\$ sudo insmod new_module.ko

Use the following command to verify the module has been loaded:

\$ lsmod



Module	Size	Used by
new_module	16384	0
btrfs	1179648	0
xor	24576	1 btrfs
zstd_compress	163840	1 btrfs
raid6_pq	114688	1 btrfs
ufs	81920	0
qnx4	16384	0
hfsplus	110592	0
hfs	61440	0
minix	36864	0
ntfs	106496	0
msdos	20480	0
jfs	188416	0
xfs	1245184	0
libcrc32c	16384	2 btrfs, xfs
crc32c_pclmul	16384	1

To remove the module from the kernel:

\$ sudo rmmod new_module

When you insert or remove the module, corresponding information will be printed out under the dmesg.

```
ksuo@ksuo-VirtualBox ~/hw4> dmesg
[79806.620385] Hello, world!
[79808.949265] Goodbye, world!
```

Part 2: Create an entry in the /proc file system for user level read and write (30 pts)

Write a new kernel module following steps in Part 1. This module creates an entry in the /proc file system. Use the following code skeleton to write the module:

<https://github.com/kevinsuo/CS3502/blob/master/project-4-2.c>

```
#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/proc_fs.h>
#include <linux/string.h>
#include <linux/vmalloc.h>
#include <linux/slab.h>
#include <linux/uaccess.h>

#define MAX_LEN          4096
static struct proc_dir_entry *proc_entry;

ssize_t read_proc(struct file *f, char *user_buf, size_t count, loff_t *off)
{
    //output the content of info to user's buffer pointed by page
    return count;
}

ssize_t write_proc(struct file *f, const char *user_buf, size_t count, loff_t *off)
{
    //copy the written data from user space and save it in info
    return count;
}

struct file_operations proc_fops = {
    read: read_proc,
    write: write_proc
};

int init_module( void )
{
    int ret = 0;
    //create the entry named myproc and allocated memory space for the proc entry
    printk(KERN_INFO "test_proc created.\n");
    return ret;
}

void cleanup_module( void )
{
    //remove the entry named myproc and free info space
```

}

Step 1: create an entry in proc file system named *myproc* when the kernel module is loaded; this entry *myproc* will be deleted when the kernel mode is deleted. You can use [*\\$ ls /proc/*](#) to check whether it is existed. (Hint: `proc_create()` and `remove_proc_entry()` are needed.)

Step 2: implement `read_proc` and `write_proc` function to read/write the proc file entry in Step 1. You need to add codes for allocating memory in `init_module` and releasing the memory in `cleanup_module` for the proc file entry. (Hint: `copy_to_user()` is needed for the read and `copy_from_user()` is needed for write.)

To test your results, load the kernel module and there should be a new entry created under `/proc`. Use `cat` and `echo` to verify and change the content of the new entry.

```
ksuo@ksuo-VirtualBox ~/hw4-2> sudo insmod my_proc.ko
[sudo] password for ksuo:
ksuo@ksuo-VirtualBox ~/hw4-2> ls /proc/
1/ 1283/ 1471/ 23/ 39/ 497/ 683/ diskstats pagetypeinfo
10/ 1284/ 15/ 24/ 4/ 50/ 686/ dma partitions
11/ 1285/ 1502/ 249/ 40/ 500/ 7/ driver/ pressure/
1114/ 1287/ 1504/ 250/ 41/ 502/ 702/ execdomains sched_debug
1119/ 1288/ 1522/ 251/ 42/ 503/ 748/ fb schedstat
1124/ 1295/ 154/ 254/ 423/ 506/ 8/ filesystems scsi/
1137/ 13/ 155/ 26/ 43/ 509/ 803/ fs/ self@
1142/ 1303/ 156/ 27/ 438/ 52/ 804/ interrupts slabinfo
1144/ 1304/ 157/ 275/ 44/ 523/ 817/ iomem softirqs
1168/ 1314/ 158/ 276/ 440/ 53/ 821/ ioports stat
1174/ 1315/ 159/ 28/ 449/ 532/ 823/ irq/ swaps
1189/ 1321/ 1598/ 282/ 45/ 533/ 891/ kallsyms sys/
1190/ 1323/ 16/ 29/ 453/ 534/ 9/ kcore sysrq-trigger
12/ 1325/ 161/ 292/ 454/ 537/ 905/ keys sysvipc/
1205/ 1327/ 162/ 3/ 455/ 54/ 912/ key-users thread-self@
1209/ 1331/ 1684/ 30/ 457/ 56/ 931/ kmsg timer_list
1210/ 1332/ 1685/ 32/ 46/ 571/ 950/ kpagecgroup tty/
1212/ 1337/ 1695/ 321/ 47/ 572/ 954/ kpagecount uptime
1218/ 1338/ 17/ 33/ 48/ 59/ 975/ kpageflags version
1229/ 1372/ 173/ 331/ 482/ 6/ acpi/ loadavg vmallocinfo
1233/ 1383/ 1763/ 335/ 484/ 60/ asound/ locks vmstat
1241/ 1384/ 18/ 336/ 485/ 606/ buddyinfo mdstat zoneinfo
1245/ 14/ 19/ 34/ 489/ 607/ bus/ meminfo
1252/ 1412/ 192/ 35/ 49/ 608/ cgroups misc
1261/ 1415/ 193/ 36/ 490/ 61/ cmdline modules
1266/ 1439/ 2/ 360/ 491/ 614/ consoles mounts@
1271/ 1442/ 20/ 37/ 493/ 634/ cpufreq mtrr
1275/ 1446/ 21/ 38/ 494/ 659/ crypto myproc
1279/ 1460/ 22/ 384/ 496/ 677/ devices net@
```

You can use the following to test the read or write on the entry of proc file system. Here the root user is needed.

Expected output:

Write “your name” into /proc/myproc. For instance, a student named “Sisi”

```
root@ksuo-VirtualBox /h/k/hw4-2# echo Sisi > /proc/myproc
```

Read /proc/myproc and printout its content:

```
root@ksuo-VirtualBox /h/k/hw4-2# cat /proc/myproc
Sisi
```

Part 3: Exchange data between the user and kernel space via mmap (50 pts)

Write a kernel module that create an entry in the /proc file system. The new entry cannot be directly read or written using cat and echo commands. Instead, map the new entry to a user space memory area so that user-level processes can read from and write to the kernel space via mmap. The skeleton of the kernel module is given below:

<https://github.com/kevinsuo/CS3502/blob/master/project-4-3-1.c>

```
#include <linux/module.h>
#include <linux/list.h>
#include <linux/init.h>
#include <linux/kernel.h>
#include <linux/types.h>
#include <linux/kthread.h>
#include <linux/proc_fs.h>
#include <linux/sched.h>
#include <linux/mm.h>
#include <linux/fs.h>
#include <linux/slab.h>

static struct proc_dir_entry *tempdir, *tempinfo;
static unsigned char *buffer;
static unsigned char array[12]={0,1,2,3,4,5,6,7,8,9,10,11};

static void allocate_memory(void);
static void clear_memory(void);
static int my_map(struct file *filp, struct vm_area_struct *vma);

static const struct file_operations myproc_fops = {
    .mmap    = my_map,
};

static int my_map(struct file *filp, struct vm_area_struct *vma)
{
    // map vma of user space to a continuous physical space
    return 0;
}

static int init_myproc_module(void)
{
```

```

tempdir=proc_mkdir("mydir", NULL);
if(tempdir == NULL) {
    printk("mydir is NULL\n");
    return -ENOMEM;
}

tempinfo = proc_create("myinfo", 0, tempdir, &myproc_fops);
if(tempinfo == NULL) {
    printk("myinfo is NULL\n");
    remove_proc_entry("mydir",NULL);
    return -ENOMEM;
}
printk("init myproc module successfully\n");

allocate_memory();

return 0;
}

static void allocate_memory(void)
{
    /* allocation memory */
    buffer = (unsigned char *)kmalloc(PAGE_SIZE,GFP_KERNEL);
    /* set the memory as reserved */
    SetPageReserved(virt_to_page(buffer));
}

static void clear_memory(void)
{
    /* clear reserved memory */
    ClearPageReserved(virt_to_page(buffer));
    /* free memory */
    kfree(buffer);
}

static void exit_myproc_module(void)
{
    clear_memory();
    remove_proc_entry("myinfo", tempdir);
    remove_proc_entry("mydir", NULL);
    printk("remove myproc module successfully\n");
}

module_init(init_myproc_module);
module_exit(exit_myproc_module);
MODULE_LICENSE("GPL");

```

The above code will create an entry **/proc/mydir/myinfo** under the proc file system. You are required to implement the **my_map** function to map one piece of memory (**char array[12]**) into user space. Then write a user space program using mmap to visit the memory space of the proc file and print the data in that memory area. You can use the following skeleton:

<https://github.com/kevinsuo/CS3502/blob/master/project-4-3-2.c>

```

#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <linux/fb.h>
#include <sys/mman.h>
#include <sys/ioctl.h>

```

```

#define PAGE_SIZE 4096


int main(int argc , char *argv[])
{
    int fd;
    int i;
    unsigned char *p_map;

    /* open proc file */
    fd = open("/proc/mydir/myinfo", O_RDWR);
    if(fd < 0) {
        printf("open fail\n");
        exit(1);
    }else {
        printf("open successfully\n");
    }

    // map p_map to the proc file and grant read & write privilege
    // read data from p_map
    // unmap p_map from the proc file

    return 0;
}

```



The above figure shows the entire workflow:

- (1) Kernel module create a proc file: **/proc/mydir/myinfo**
- (2) User process open the created proc file
- (3) User process calls mmap function, which further executed my_map defined in the kernel
- (4) my_map() then maps one piece of memory into user space (e.g., buffer) and puts some data inside
- (5) User process visits this piece of memory and prints the data out.

Expected output:

For instance, a student named Sisi should upload a screenshot like:

```
ksuo@ksuo-VirtualBox ~/hw4-3> sudo ./test_user.o
open successfully by Sisi
0
1
2
3
4
5
6
7
8
9
10
11
Printed by Sisi
```

Submission requirements:

Submit your assignment file through D2L using the appropriate link.

The submission must include the source code, and a report describe your code logic (including part 1/2/3). Output screenshot (including part 1/2/3) of your code should be included in the report.