

# CS 3502

# Operating Systems

## Project Lab

**Kun Suo**

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

# Outline

---

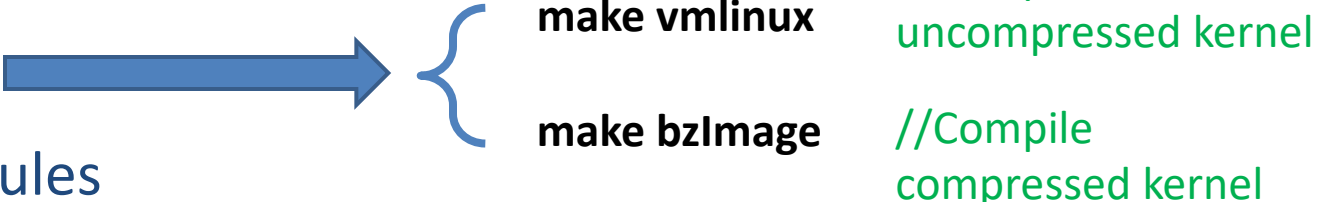
- Part 1: Create a helloworld kernel module (20')
- Part 2: Create an entry in the /proc file system for user-level read and write (30')
- Part 3: Exchange data between the user and kernel space via mmap (50')



# Part 1: Kernel module

---

- Compile and install the Linux kernel

- make
  - make modules
  - make modules\_install
  - make install
- 
- make vmlinux** //Compile uncompressed kernel
- make bzImage** //Compile compressed kernel

# Part 1: Kernel module

- Compile and install the Linux kernel

- make

- make modules



make vmlinux

make bzImage

//Compile  
uncompressed kernel

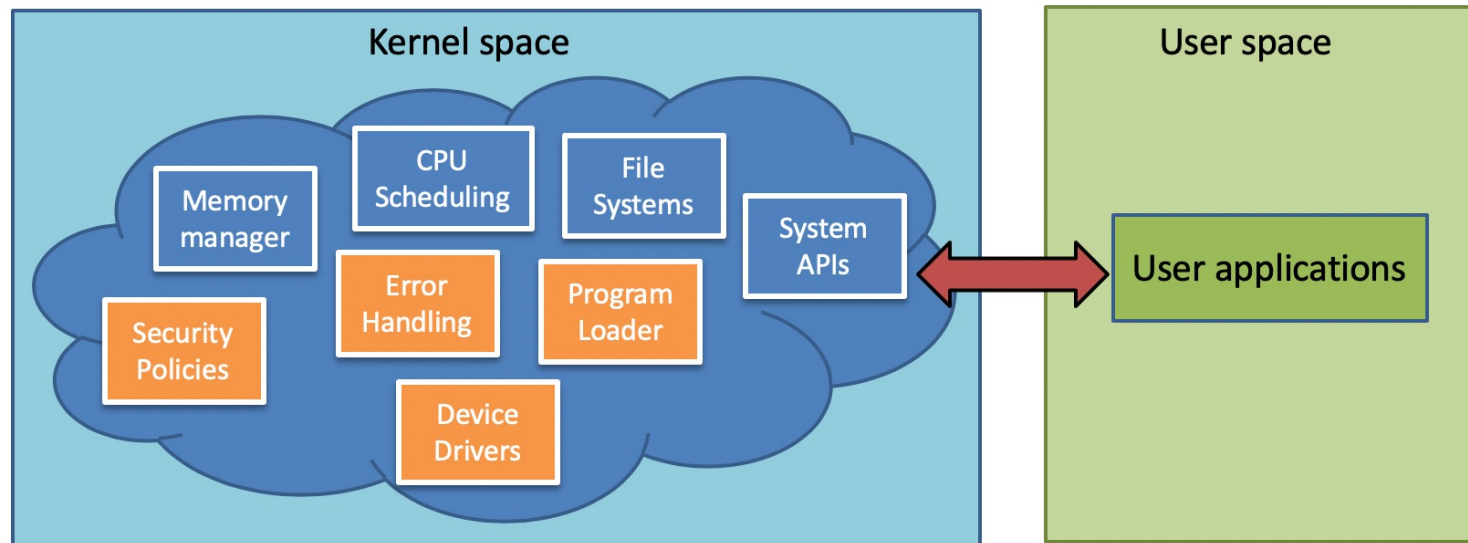
//Compile  
compressed kernel

```
vm [Running]
Thu 11:38
fish /home/ksuo/Downloads/linux-5.1
File Edit View Search Terminal Help
ksuo@ksuo-VirtualBox ~/D/linux-5.1> ls /boot/
config-5.0.0-23-generic      memtest86+.elf
config-5.0.0-25-generic      memtest86+_multiboot.bin
config-5.1.0                System.map-5.0.0-23-generic
grub/                       System.map-5.0.0-25-generic
initrd.img-5.0.0-23-generic  System.map-5.1.0
initrd.img-5.0.0-25-generic  vmlinuz-5.0.0-23-generic
initrd.img-5.1.0            vmlinuz-5.0.0-25-generic
memtest86+.bin              vmlinuz-5.1.0
ksuo@ksuo-VirtualBox ~/D/linux-5.1>
```

# Part 1: Kernel module

---

- Compile and install the Linux kernel
  - `make`
  - `make modules`
  - `make modules_install`
  - `make install`



# List all modules in the kernel

```
fish /home/ksuo/hw4 (ssh)
ksuo@ksuo-VirtualBox ~/hw4> lsmod
Module                Size  Used by
btrfs                 1179648  0
xor                   24576  1 btrfs
zstd_compress         163840  1 btrfs
raid6_pq              114688  1 btrfs
ufs                   81920  0
qnx4                  16384  0
hfsplus              110592  0
hfs                   61440  0
minix                 36864  0
ntfs                  106496  0
msdos                 20480  0
jfs                   188416  0
xfs                   1245184  0
libcrc32c             16384  2 btrfs,xfs
crct10dif_pclmul      16384  1
crc32_pclmul          16384  0
ghash_clmulni_intel   16384  0
vmwgfx                290816  2
ttm                   102400  1 vmwgfx
drm_kms_helper        180224  1 vmwgfx
aesni_intel           372736  0
snd_intel8x0           45056  2
snd_ac97_codec        135168  1 snd_intel8x0
aes_x86_64            20480  1 aesni_intel
crypto_simd           16384  1 aesni_intel
cryptd                24576  3 crypto_simd,ghash_clmulni_intel,aesni_intel
ac97_bus              16384  1 snd_ac97_codec
glue_helper           16384  1 aesni_intel
snd_pcm               102400  2 snd_intel8x0,snd_ac97_codec
```

# Build your module

<https://github.com/kevinsuo/CS3502/blob/master/project-4-1.c>

## *new\_module.c*

```
#include <linux/module.h>
#include <linux/kernel.h>

int init_new_module(void)
{
    printk(KERN_INFO "Hello, world!\n");
    return 0;
}

void exit_new_module(void) {
    printk(KERN_INFO "Goodbye, world!\n");
}

module_init(init_new_module);
module_exit(exit_new_module);
```

init\_module is invoked when the module is loaded into the kernel

exit\_module is called when the module is removed from the kernel



# Compile your module

<https://github.com/kevinsuo/CS3502/blob/master/project-4-1-Makefile>

## *Makefile*

new\_module.o is the output file

```
obj-m += new_module.o
all:
    sudo make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
clean:
    sudo make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
```

```
ksuo@ksuo-VirtualBox ~/hw4> make
sudo make -C /lib/modules/5.1.0/build M=/home/ksuo/hw4 modules
[sudo] password for ksuo:
make: Entering directory '/home/ksuo/linux-5.1-modified'
Building modules, stage 2.
MODPOST 1 modules
WARNING: modpost: missing MODULE_LICENSE() in /home/ksuo/hw4/new_module.o
see include/linux/module.h for more information
make: Leaving directory '/home/ksuo/linux-5.1-modified'
ksuo@ksuo-VirtualBox ~/hw4> ls
Makefile      Module.symvers  new_module.ko    new_module.mod.o
modules.order new_module.c     new_module.mod.c new_module.o
```



# Insert the module into the Linux kernel

---

- `sudo insmod new_module.ko`

```
fish /home/ksuo/hw4 (ssh)
ksuo@ksuo-VirtualBox ~/hw4> lsmod
Module                Size  Used by
new_module             16384  0
btrfs                 1179648  0
xor                   24576  1 btrfs
zstd_compress         163840  1 btrfs
raid6_pq              114688  1 btrfs
ufs                   81920  0
qnx4                  16384  0
hfsplus               110592  0
hfs                   61440  0
minix                 36864  0
ntfs                  106496  0
msdos                 20480  0
jfs                   188416  0
xfs                  1245184  0
libcrc32c             16384  2 btrfs,xfs
crct10dif_pclmul      16384  1
```

# Remove the module from the kernel

- `sudo rmmod new_module`

```
fish /home/ksuo/hw4 (ssh)
ksuo@ksuo-VirtualBox ~/hw4> lsmod
Module                Size  Used by
btrfs                 1179648  0
xor                   24576  1 btrfs
zstd_compress         163840  1 btrfs
raid6_pq              114688  1 btrfs
ufs                   81920  0
qnx4                  16384  0
hfsplus              110592  0
hfs                   61440  0
minix                 36864  0
ntfs                  106496  0
msdos                 20480  0
jfs                   188416  0
xfs                  1245184  0
libcrc32c             16384  2 btrfs,xfs
crct10dif_pclmul      16384  1
crc32_pclmul          16384  0
ghash_clmulni_intel   16384  0
vmwgfx                290816  2
ttm                   102400  1 vmwgfx
drm_kms_helper        180224  1 vmwgfx
aesni_intel           372736  0
```



# Related codes

---

- ***new\_module.c:***

<https://github.com/kevinsuo/CS3502/blob/master/project-4-1.c>

- ***Makefile:***

<https://github.com/kevinsuo/CS3502/blob/master/project-4-1-Makefile>



# Outline

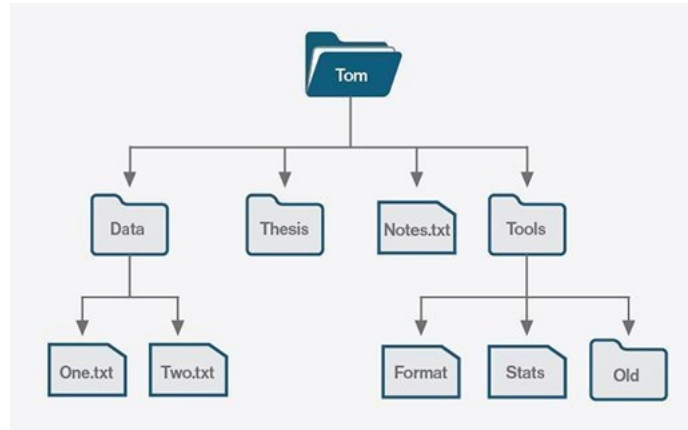
---

- Part 1: Create a helloworld kernel module (30')
- Part 2: Create an entry in the /proc file system for user level read and write (30')
- Part 3: Exchange data between the user and kernel space via mmap (50')

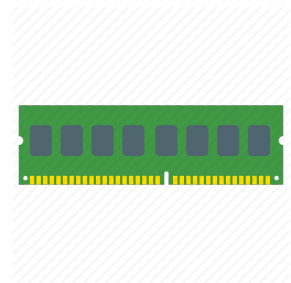


# File System

OS { File system  
Proc File system



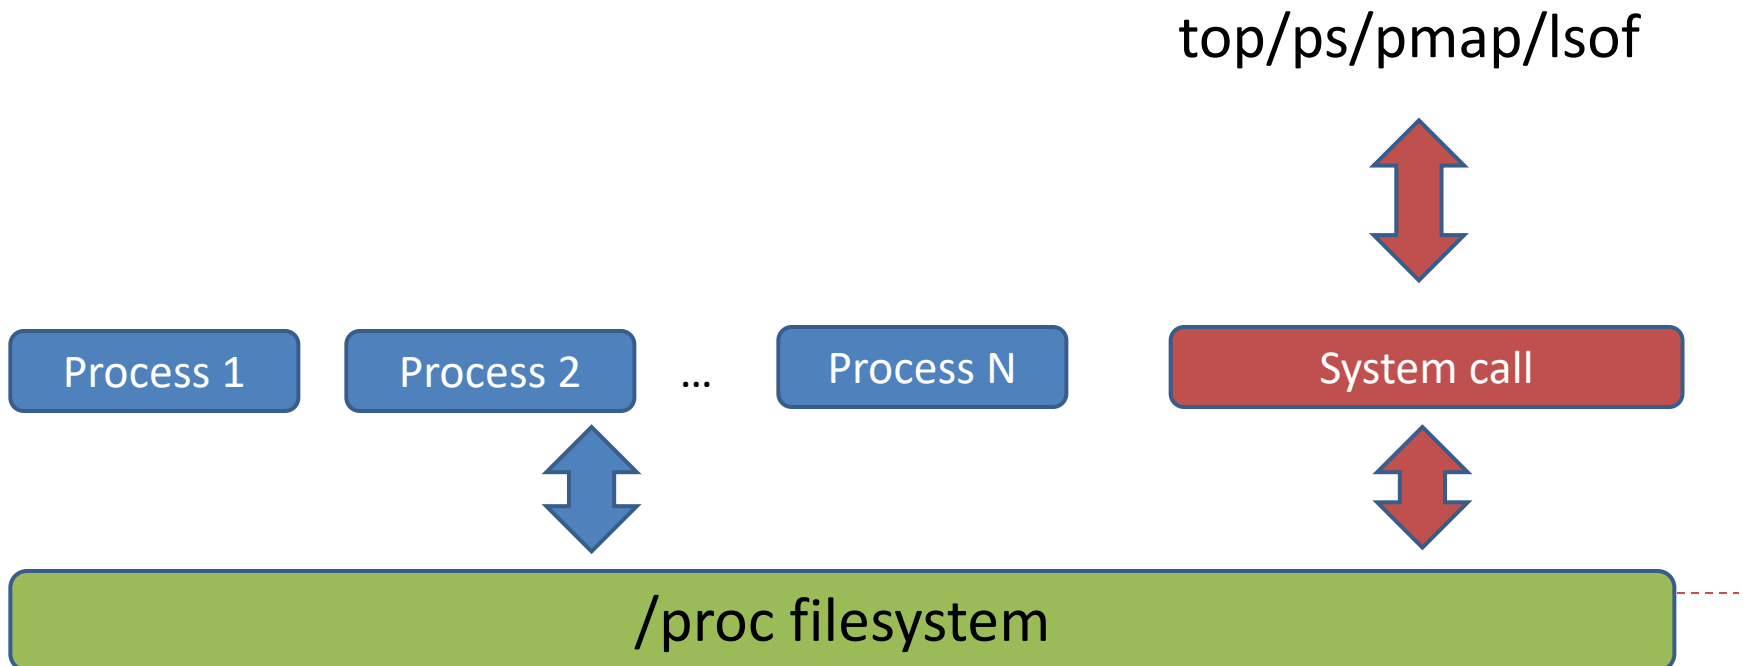
Virtual  
FS



# Proc file system

---

- The proc filesystem (procfs) is a special filesystem in Unix-like OS that presents information about processes and other system information in a hierarchical file-like structure, providing a convenient and standardized method for dynamically accessing process data held in the kernel



# Proc file system

```
administrator@ubuntuvm-1604: ~
nmon-14g Hostname=ubuntuvm-1604Refresh= 2secs 11:44.54
CPU Utilisation
-----+-----+-----+-----+-----+
CPU  User%  Sys%  Wait%  Idle|0      |25      |50      |75      |100|
  1   7.1   3.5   0.0   89.4|UUUS  >
  2   7.0   4.5   0.0   88.5|UUUS  >
-----+-----+-----+-----+
Avg   7.1   3.8   0.0   89.2|UUUS  >
-----+-----+-----+-----+
Memory Stats
-----+-----+-----+-----+
          RAM      High      Low      Swap      Page Size=4 KB
Total MB      7977.3      -0.0      -0.0      975.0
Free MB        285.7      -0.0      -0.0      975.0
Free Percent    3.6%    100.0%    100.0%    100.0%
          MB
          Cached= 4988.9      Active= 3972.7
Buffers=  588.8 Swapcached=  0.0 Inactive = 2755.3
Dirty  =   0.0 Writeback=  0.0 Mapped  =  326.6
Slab   =  849.7 Commit_AS = 4989.8 PageTables=  29.3
-----+-----+-----+-----+
Network I/O
-----+-----+-----+-----+
I/F Name Recv=KB/s Trans=KB/s packin packout insize outsize Peak->Recv Trans
ens160    1.6     4.1     18.0     16.0     92.7    263.6     17.7    192.6
lo         0.0     0.0     0.0     0.0     0.0     0.0      0.0     0.0
-----+-----+-----+-----+
Disk I/O ---/proc/diskstats---mostly in KB/s---Warning:contains duplicates
DiskName Busy  Read WriteKB|0      |25      |50      |75      |100|
loop0     0%    0.0   0.0|>disk busy not available
sr0        0%    0.0   0.0|>
sda        0%    0.0  10.0|>
sda1       0%    0.0  10.0|>
sda2       0%    0.0   0.0|>
sda5       0%    0.0   0.0|>
Totals Read-MB/s=0.0      Writes-MB/s=0.0      Transfers/sec=2.0
```

# Proc file system

- ls /proc/

```
fish /home/ksuo/hw4 (ssh)
ksuo@ksuo-VirtualBox ~/hw4> ls /proc/
1/      1332/  159/   234/   34/    478/   6/      crypto  net@
10/     1347/  16/    235/   35/    479/   60/     devices pagetypeinfo
1014/   1368/  160/   238/   354/   48/    602/    diskstats partitions
1017/   1377/  162/   24/    36/    484/   61/     dma      pressure/
1022/   14/    1644/  240/   37/    488/   641/    driver/  sched_debug
1035/   1411/  1650/  2470/  372/   49/    651/    execdomains schedstat
1040/   1432/  1651/  2472/  373/   494/   652/    fb        scsi/
1042/   1433/  1676/  2490/  38/    496/   660/    filesystems self@
1066/   1437/  1678/  2516/  39/    499/   696/    fs/        slabinfo
1072/   1446/  1698/  26/    4/     50/    706/    interrupts softirqs
1087/   1452/  17/    260/   40/    503/   709/    iomem      stat
1088/   1458/  1701/  261/   401/   508/   734/    ioports    swaps
11/     1463/  1704/  27/    41/    51/    735/    irq/       sys/
1115/   1483/  1722/  2717/  414/   513/   750/    kallsyms   sysrq-trigger
1130/   1486/  173/   2718/  418/   514/   767/    kcore      sysvipc/
1140/   1497/  1738/  275/   42/    52/    776/    keys       thread-self@
```



# Proc file system

---

- `/proc/cmdline` – Kernel command line information.
- `/proc/console` – Information about current consoles including `tty`.
- `/proc/devices` – Device drivers currently configured for the running kernel.
- `/proc/dma` – Info about current DMA channels.
- `/proc/fb` – Framebuffer devices.
- `/proc/filesystems` – Current filesystems supported by the kernel.
- `/proc/iomem` – Current system memory map for devices.
- `/proc/ioports` – Registered port regions for input output communication with device.
- `/proc/loadavg` – System load average.
- `/proc/locks` – Files currently locked by kernel.
- `/proc/meminfo` – Info about system memory (see above example).
- `/proc/misc` – Miscellaneous drivers registered for miscellaneous major device.
- `/proc/modules` – Currently loaded kernel modules.
- `/proc/mounts` – List of all mounts in use by system.
- `/proc/partitions` – Detailed info about partitions available to the system.
- `/proc/pci` – Information about every PCI device.
- `/proc/stat` – Record of various statistics kept from last reboot.
- `/proc/swap` – Information about swap space.
- `/proc/uptime` – Uptime information (in seconds).
- `/proc/version` – Kernel version, gcc version, and Linux distribution installed.

# Part 2: Create an entry in the /proc file system

---

- <https://github.com/kevinsuo/CS3502/blob/master/project-4-2.c>

```
int init_module( void )
{
    int ret = 0;
    //create the entry and allocated memory space for the proc entry

    printk(KERN_INFO "test_proc created.\n");

    return ret;
}

void cleanup_module( void )
{
    //remove the proc entry and free info space

    printk(KERN_INFO "test_proc deleted.\n");
}
```



# Part 2: Create an entry in the /proc file system

---

- <https://github.com/kevinsuo/CS3502/blob/master/project-4-2.c>

```
int init_module( void )
{
    int ret = 0;
    //create the entry and allocated memory space for the proc entry

    printk(KERN_INFO "test_proc created.\n");

    return ret;
}
```

Google search “proc\_create”

```
void cleanup_module( void )
{
    //remove the proc entry and free info space

    printk(KERN_INFO "test_proc deleted.\n");
}
```

Google search “remove\_proc\_entry”



## Part 2: after you insert your module, check whether it exists under /proc

```
ksuo@ksuo-VirtualBox ~/hw4-2> ls /proc/
1/      1283/  1471/  23/    39/    497/  683/    diskstats  pagetypeinfo
10/     1284/  15/    24/    4/     50/   686/    dma        partitions
11/     1285/  1502/  249/   40/    500/  7/      driver/    pressure/
1114/   1287/  1504/  250/   41/    502/  702/    execdomains sched_debug
1119/   1288/  1522/  251/   42/    503/  748/    fb         schedstat
1124/   1295/  154/   254/   423/   506/  8/      filesystems scsi/
1137/   13/   155/   26/    43/    509/  803/    fs/        self@
1142/   1303/  156/   27/    438/   52/   804/    interrupts slabinfo
1144/   1304/  157/   275/   44/    523/  817/    iomem      softirqs
1168/   1314/  158/   276/   440/   53/   821/    ioports    stat
1174/   1315/  159/   28/    449/   532/  823/    irq/       swaps
1189/   1321/  1598/  282/   45/    533/  891/    kallsyms   sys/
1190/   1323/  16/    29/    453/   534/  9/      kcore      sysrq-trigger
12/     1325/  161/   292/   454/   537/  905/    keys       sysvipc/
1205/   1327/  162/   3/     455/   54/   912/    key-users  thread-self@
1209/   1331/  1684/  30/    457/   56/   931/    kmsg       timer_list
1210/   1332/  1685/  32/    46/    571/  950/    kpagecgrou tty/
1212/   1337/  1695/  321/   47/    572/  954/    kpagecount uptime
1218/   1338/  17/    33/    48/    59/   975/    kpageflags version
1229/   1372/  173/   331/   482/   6/    acpi/    loadavg    vmallocinfo
1233/   1383/  1763/  335/   484/   60/    asound/   locks      vmstat
1241/   1384/  18/    336/   485/   606/    buddyinfo mdstat      zoneinfo
1245/   14/    19/    34/    489/   607/    bus/      meminfo
1252/   1412/  192/   35/    49/    608/    cgroups   misc
1261/   1415/  193/   36/    490/   61/    cmdline  modules
1266/   1439/  2/     360/   491/   614/    consoles mounts@
1271/   1442/  20/    37/    493/   634/    cpuinfo   mnt
1275/   1446/  21/    38/    494/   659/    crypto    myproc
1279/   1460/  22/    384/   496/   677/    devices  net@
```

Here my module is named as  
“myproc”

## Part 2: after you remove your module, check whether it exists under /proc

- When your module is removed, it should disappear from the /proc

```
ksuo@ksuo-VirtualBox ~/hw4-2> sudo rmmod my_proc
fish: "sudo rmmod my_proc" terminated by signal SIGKILL (Forced quit)
ksuo@ksuo-VirtualBox ~/hw4-2> ls /proc/
```

1/	1229/	1314/	1460/	17/	275/	38/	47/	52/	634/	950/	fs/	mounts@	tty/
10/	1233/	1315/	1471/	173/	276/	384/	48/	523/	659/	954/	interrupts	mtrr	uptime
11/	1241/	1321/	15/	1791/	28/	39/	482/	53/	677/	975/	iomem	net@	version
1114/	1245/	1323/	1502/	18/	282/	4/	484/	532/	683/	acpi/	ioports	pagetypeinfo	vmallocinfo
1119/	1252/	1325/	1504/	19/	29/	40/	485/	533/	686/	asound/	irq/	partitions	vmstat
1124/	1261/	1327/	1522/	192/	292/	41/	489/	534/	7/	buddyinfo	kallsyms	pressure/	zoneinfo
1137/	1266/	1331/	154/	193/	3/	42/	49/	537/	702/	bus/	kcore	sched_debug	
1142/	1271/	1332/	155/	2/	30/	423/	490/	54/	748/	cgroups	keys	schedstat	
1144/	1275/	1337/	156/	20/	32/	43/	491/	56/	8/	cmdline	key-users	scsi/	
1168/	1279/	1338/	157/	21/	321/	438/	493/	571/	803/	consoles	kmsg	self@	
1174/	1283/	1372/	158/	22/	33/	44/	494/	572/	804/	cpuinfo	kpagecgroup	slabinfo	
1189/	1284/	1383/	159/	23/	331/	440/	496/	59/	817/	crypto	kpagecount	softirqs	
1190/	1285/	1384/	1598/	24/	335/	449/	497/	6/	821/	devices	kpageflags	stat	
12/	1287/	14/	16/	249/	336/	45/	50/	60/	823/	diskstats	loadavg	swaps	
1205/	1288/	1412/	161/	250/	34/	453/	500/	606/	891/	dma	locks	sys/	
1209/	1295/	1415/	162/	251/	35/	454/	502/	607/	9/	driver/	mdstat	sysrq-trigger	
1210/	13/	1439/	1684/	254/	36/	455/	503/	608/	905/	execdomains	meminfo	sysvipc/	
1212/	1303/	1442/	1685/	26/	360/	457/	506/	61/	912/	fb	misc	thread-self@	
1218/	1304/	1446/	1695/	27/	37/	46/	509/	614/	931/	filesystems	modules	timer_list	

# Part 2: read/write the proc entry your created in your module

```
ssize_t read_proc(struct file *f, char *user_buf, size_t count, loff_t *off )
{
    //output the content of info to user's buffer pointed by page
    printk(KERN_INFO "procfs_read: read %lu bytes\n", count);
    return count;
}
```

Write data from user space to  
your proc file memory  
Google search "copy\_to\_user"

```
ssize_t write_proc(struct file *f, const char *user_buf, size_t count, loff_t *off)
{
    //copy the written data from user space and save it in info
    printk(KERN_INFO "procfs_write: write %lu bytes\n", count);
    return count;
}
```

Write data from user space to  
your proc file memory  
Google search "copy\_from\_user"

```
int init_module( void )
{
    int ret = 0;
    //create the entry and allocated memory space for the proc entry

    printk(KERN_INFO "test_proc created.\n");

    return ret;
}
```

Allocate memory space  
for your proc file

```
void cleanup_module( void )
{
    //remove the proc entry and free info space

    printk(KERN_INFO "test_proc deleted.\n");
}
```

Release memory space  
from your proc file

## Part 2: read/write the proc entry your created in your module

---

- Use the following to test the read or write on the entry of proc file system
  - # echo to write data into your proc entry
  - # cat command to read data from your proc entry

```
root@ksuo-VirtualBox /h/k/hw4-2# echo 12345 > /proc/myproc
root@ksuo-VirtualBox /h/k/hw4-2# cat /proc/myproc
12345
```



# Outline

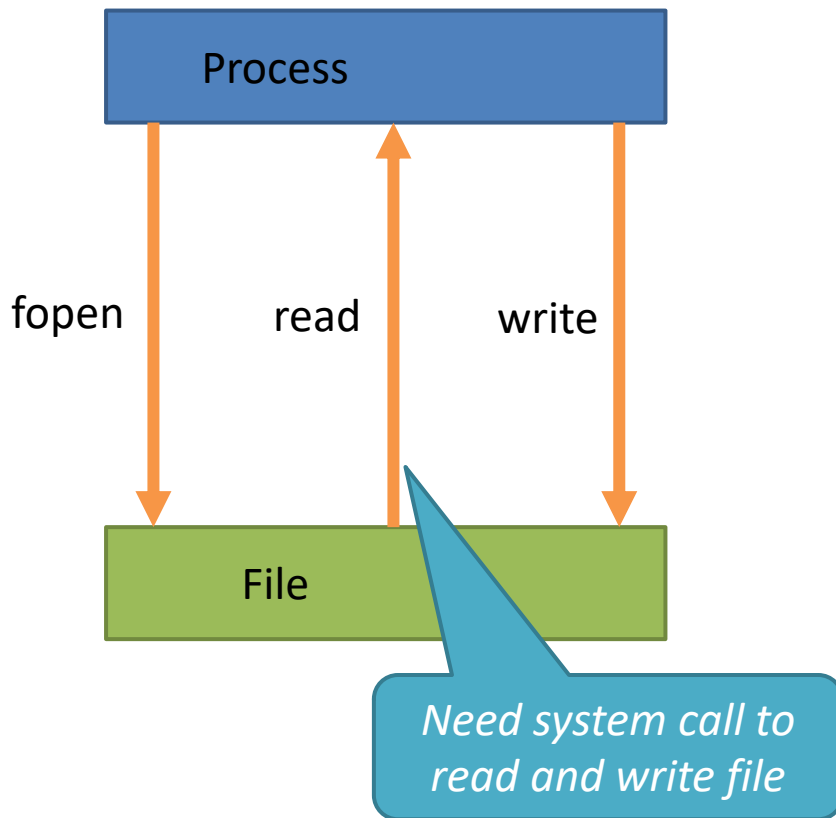
---

- Part 1: Create a helloworld kernel module (20')
- Part 2: Create an entry in the /proc file system for user-level read and write (30')
- Part 3: Exchange data between the user and kernel space via mmap (50')





# Memory-mapped Files

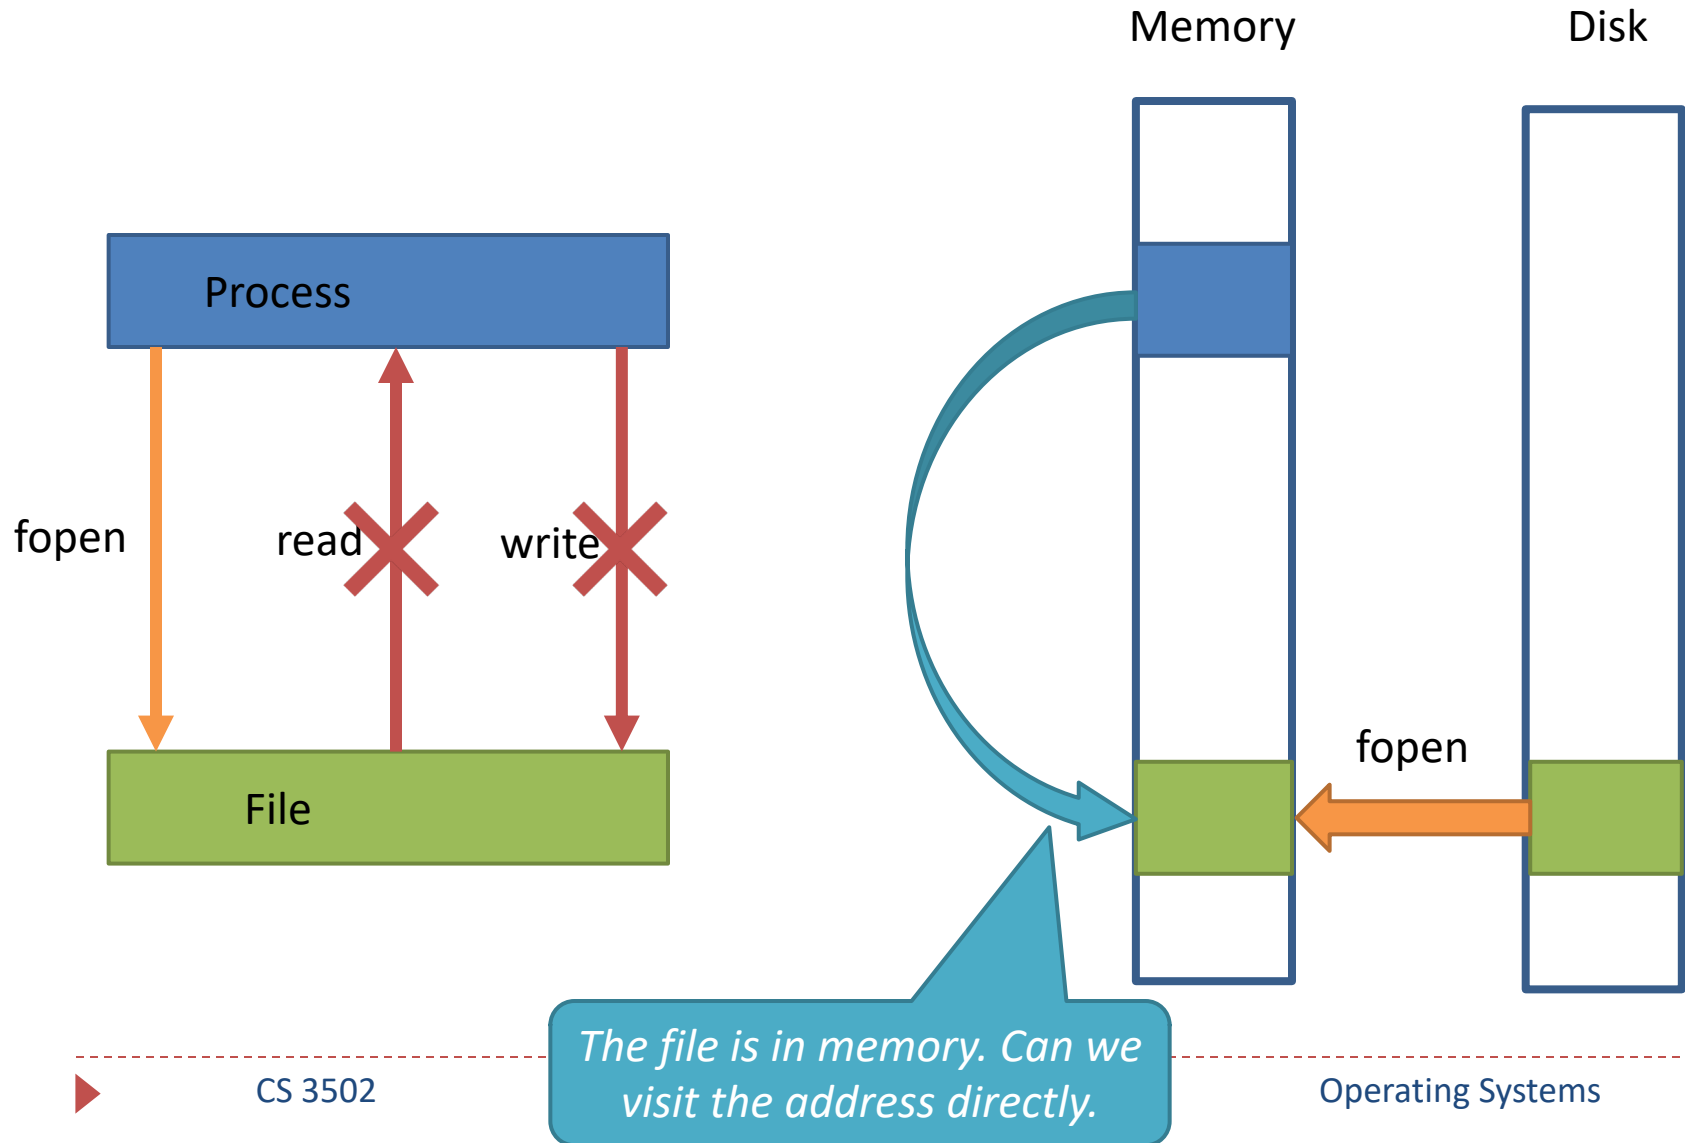


```
/* Open the input file and create the output file */
in_fd = open(argv[1], O_RDONLY); /* open the source file */
if (in_fd < 0) exit(2);           /* if it cannot be opened, exit */
out_fd = creat(argv[2], OUTPUT_MODE); /* create the destination file */
if (out_fd < 0) exit(3);          /* if it cannot be created, exit */

/* Copy loop */
while (TRUE) {
    rd_count = read(in_fd, buffer, BUF_SIZE); /* read a block of data */
    if (rd_count <= 0) break;                 /* if end of file or error, exit loop */
    wt_count = write(out_fd, buffer, rd_count); /* write data */
    if (wt_count <= 0) exit(4);               /* wt_count <= 0 is an error */
}

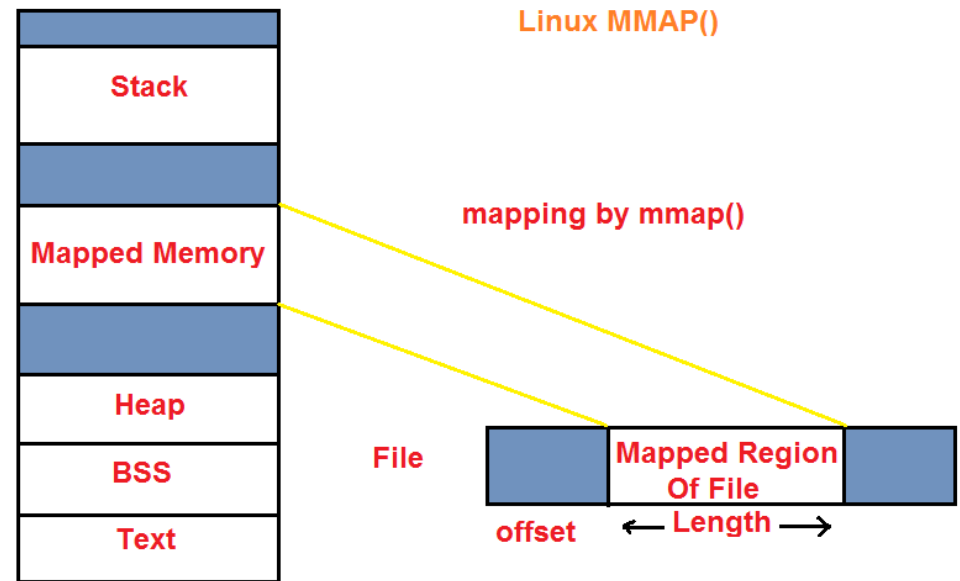
/* Close the files */
close(in_fd);
close(out_fd);
if (rd_count == 0) /* no error on last read */
    exit(0);
else /* error on last read */
    exit(5);
```

# Memory-mapped Files



# Memory-mapped Files

- OS provide a way (map and unmap) to map files into the address space of a running process
  - No read or write system calls are needed thereafter
- Advantages
  - Improved I/O performance and avoidance of kernel to user data copying

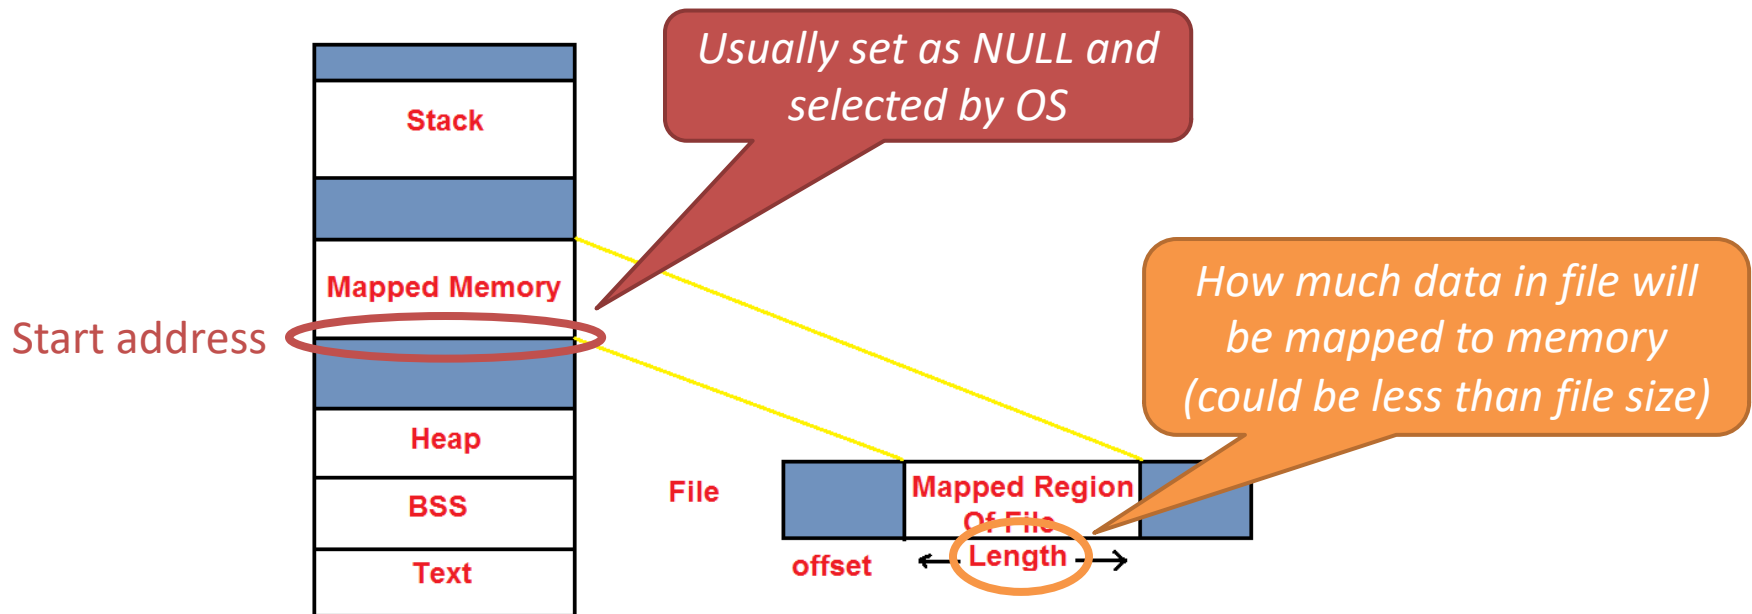


# Memory-mapped Files

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t length, int prot, int flags, int fd, off_t offset);
```

<https://pubs.opengroup.org/onlinepubs/009695399/functions/mmap.html>

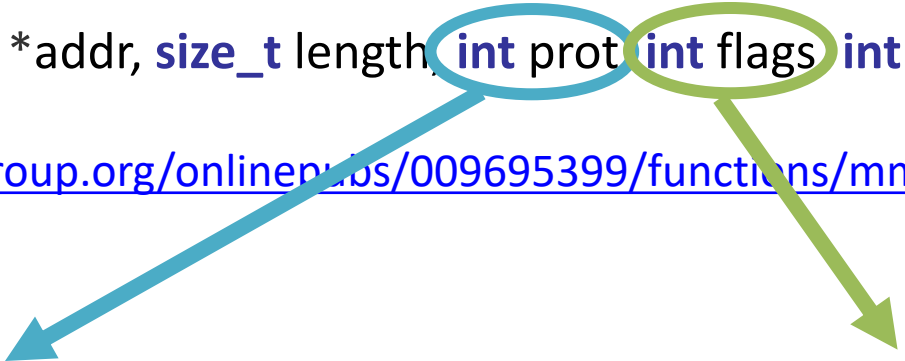


# Memory-mapped Files

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t length, int prot, int flags, int fd, off_t offset);
```

<https://pubs.opengroup.org/onlinepubs/009695399/functions/mmap.html>



Symbolic Constant	Description
PROT_READ	Data can be read.
PROT_WRITE	Data can be written.
PROT_EXEC	Data can be executed.
PROT_NONE	Data cannot be accessed.

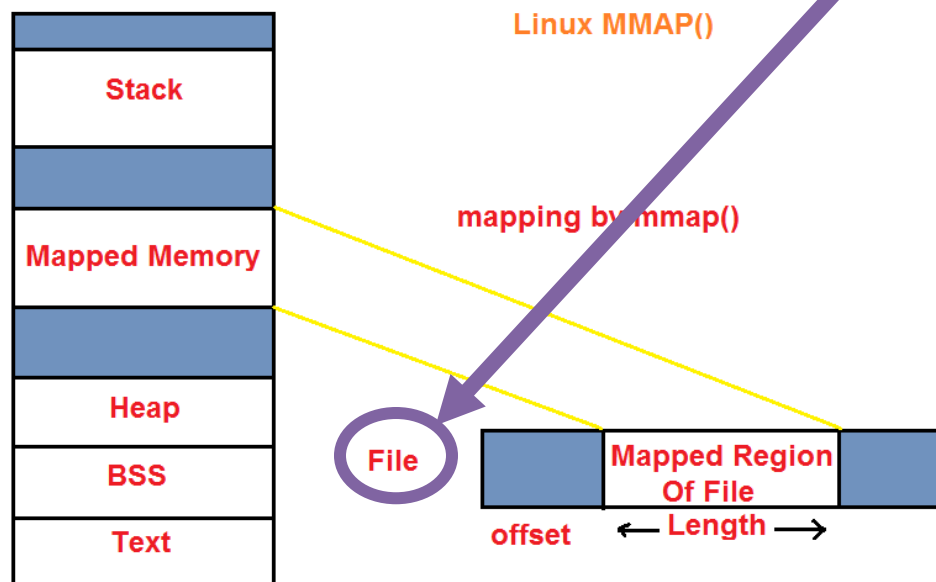
Symbolic Constant	Description
MAP_SHARED	Changes are shared.
MAP_PRIVATE	Changes are private.
MAP_FIXED	Interpret <i>addr</i> exactly.

# Memory-mapped Files

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t length, int prot, int flags, int fd, off_t offset);
```

<https://pubs.opengroup.org/onlinepubs/009695399/functions/mmap.html>



# Memory-mapped File Example

```
#include <sys/mman.h> /* for mmap and munmap */
#include <sys/types.h> /* for open */
#include <sys/stat.h> /* for open */
#include <fcntl.h> /* for open */
#include <unistd.h> /* for lseek and write */
#include <stdio.h>
```

```
int main(int argc, char **argv)
{
```

```
    int fd;
    char *mapped_mem, *p;
    int flength = 1024;
    void *start_addr = 0;
```

*Return the  
mapped  
memory address*

*Print out the  
data in the  
memory*

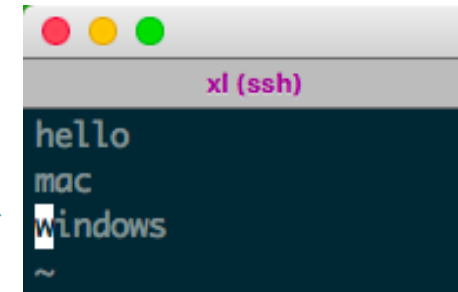
```
    fd = open(argv[1], O_RDWR | O_CREAT, S_IRUSR | S_IWUSR);
    flength = lseek(fd, 1, SEEK_END);
    lseek(fd, 0, SEEK_SET);
```

```
    mapped_mem = mmap(start_addr, flength, PROT_READ, MAP_PRIVATE, fd, 0);
```

```
    printf("%s\n", mapped_mem);
    close(fd);
    munmap(mapped_mem, flength);
    return 0;
```

```
}
```

text.txt



*Allow read*

*Set private, do not allow  
other process to read*

*0: beginning of the file  
1024: mapped memory size*

# Memory-mapped File Example

---

```
#include <sys/mman.h> /* for mmap and munmap */
#include <sys/types.h> /* for open */
#include <sys/stat.h> /* for open */
#include <fcntl.h> /* for open */
#include <unistd.h> /* for lseek and write */
#include <stdio.h>

int main(int argc, char **argv)
{
    int fd;
    char *mapped_mem, *p;
    int flength = 1024;
    void *start_addr = 0;

    fd = open(argv[1], O_RDWR | O_CREAT, S_IRUSR | S_IWUSR);
    flength = lseek(fd, 1, SEEK_END);
    lseek(fd, 0, SEEK_SET);

    mapped_mem = mmap(start_addr, flength, PROT_READ, MAP_PRIVATE, fd, 0);

    printf("%s\n", mapped_mem);
    close(fd);
    munmap(mapped_mem, flength);
    return 0;
}
```

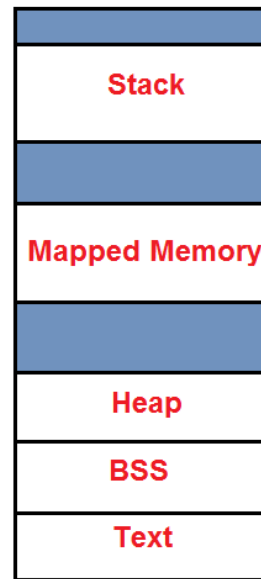




# Memory-mapped File Example

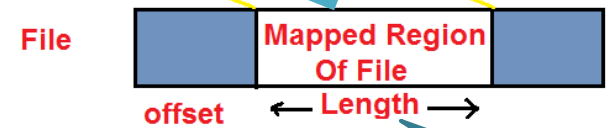
*Print out the data in this memory area*

```
ksuo@centos65-pv-3 mymmap$ ./a.out text.txt
hello
mac
windows
ksuo@centos65-pv-3 mymmap$
```

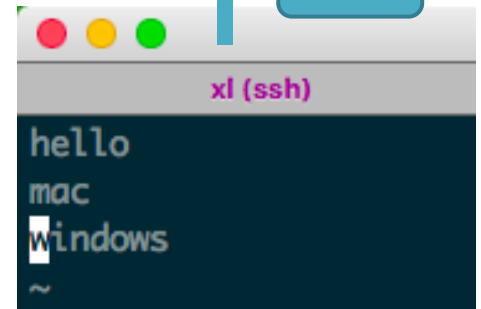


Linux MMAP()

mapping by mmap()



1024



text.txt

# Comparison of regular file and memory-mapped file

```
/* Open the input file and create the output file */
in_fd = open(argv[1], O_RDONLY); /* open the source file */
if (in_fd < 0) exit(2);           /* if it cannot be opened, exit */
out_fd = creat(argv[2], OUTPUT_MODE); /* create the destination file */
if (out_fd < 0) exit(3);         /* if it cannot be created, exit */

/* Copy loop */
while (TRUE) {
    rd_count = read(in_fd, buffer, BUF_SIZE); /* read a block of data */
    if (rd_count <= 0) break;                  /* if end of file or error, exit loop */
    wt_count = write(out_fd, buffer, rd_count); /* write data */
    if (wt_count <= 0) exit(4);                /* wt_count <= 0 is an error */
}

/* Close the files */
close(in_fd);
close(out_fd);
if (rd_count == 0) /* no error on last read */
    exit(0);
else /* error on last read */
    exit(5);
```

```
#include <sys/mman.h> /* for mmap and munmap */
#include <sys/types.h> /* for open */
#include <sys/stat.h> /* for open */
#include <fcntl.h> /* for open */
#include <unistd.h> /* for lseek and write */
#include <stdio.h>
```

```
int main(int argc, char **argv)
{
```

```
    int fd;
    char *mapped_mem, *p;
    int flength = 1024;
    void *start_addr = 0;
```

```
    fd = open(argv[1], O_RDWR | O_CREAT, S_IRUSR | S_IWUSR);
    flength = lseek(fd, 1, SEEK_END);
    lseek(fd, 0, SEEK_SET);
```

```
    mapped_mem = mmap(start_addr, flength, PROT_READ, MAP_PRIVATE, fd, 0);
```

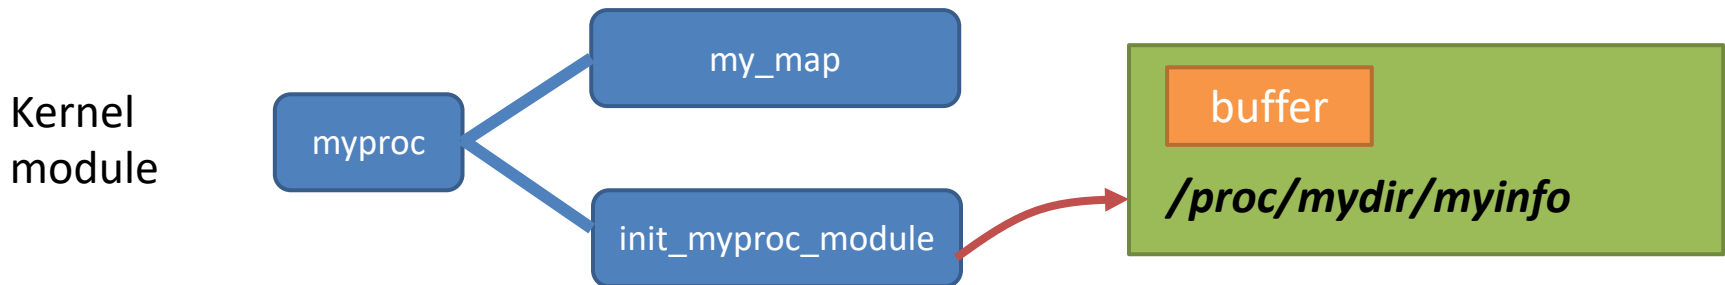
```
    printf("%s\n", mapped_mem);
    close(fd);
    munmap(mapped_mem, flength);
    return 0;
```

```
}
```



# Part 3: Exchange data between the user and kernel space via mmap

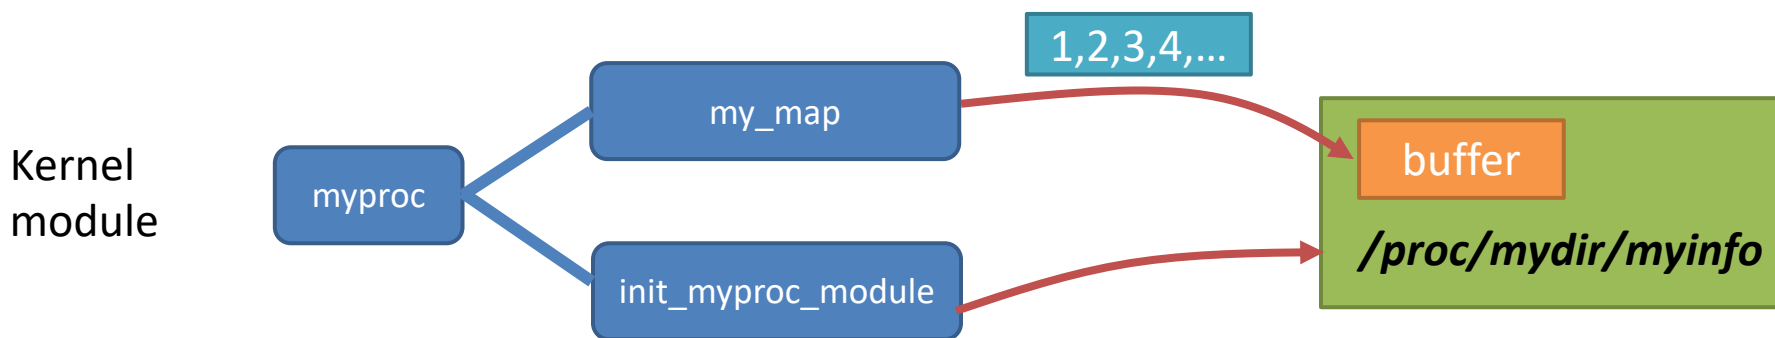
- <https://github.com/kevinsuo/CS3502/blob/master/project-4-3-1.c>
- The above code will create an entry ***/proc/mydir/myinfo*** under the proc file system and allocate a buffer under this entry



# Part 3: Exchange data between the user and kernel space via mmap

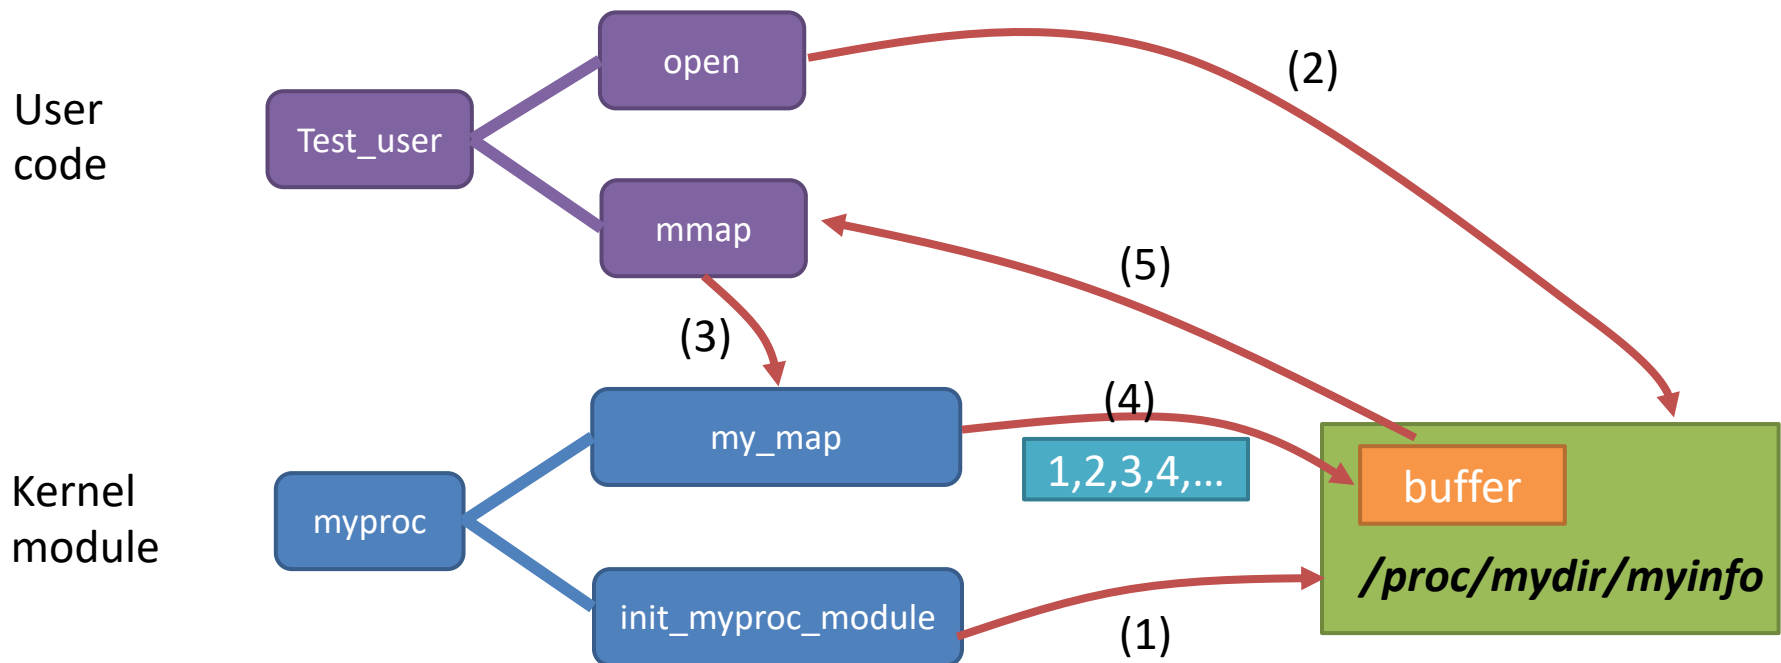
- You are required to implement the *my\_map* function to map one piece of memory (*char array[12]*) into user space.

```
static unsigned char array[12]={0,1,2,3,4,5,6,7,8,9,10,11};
```



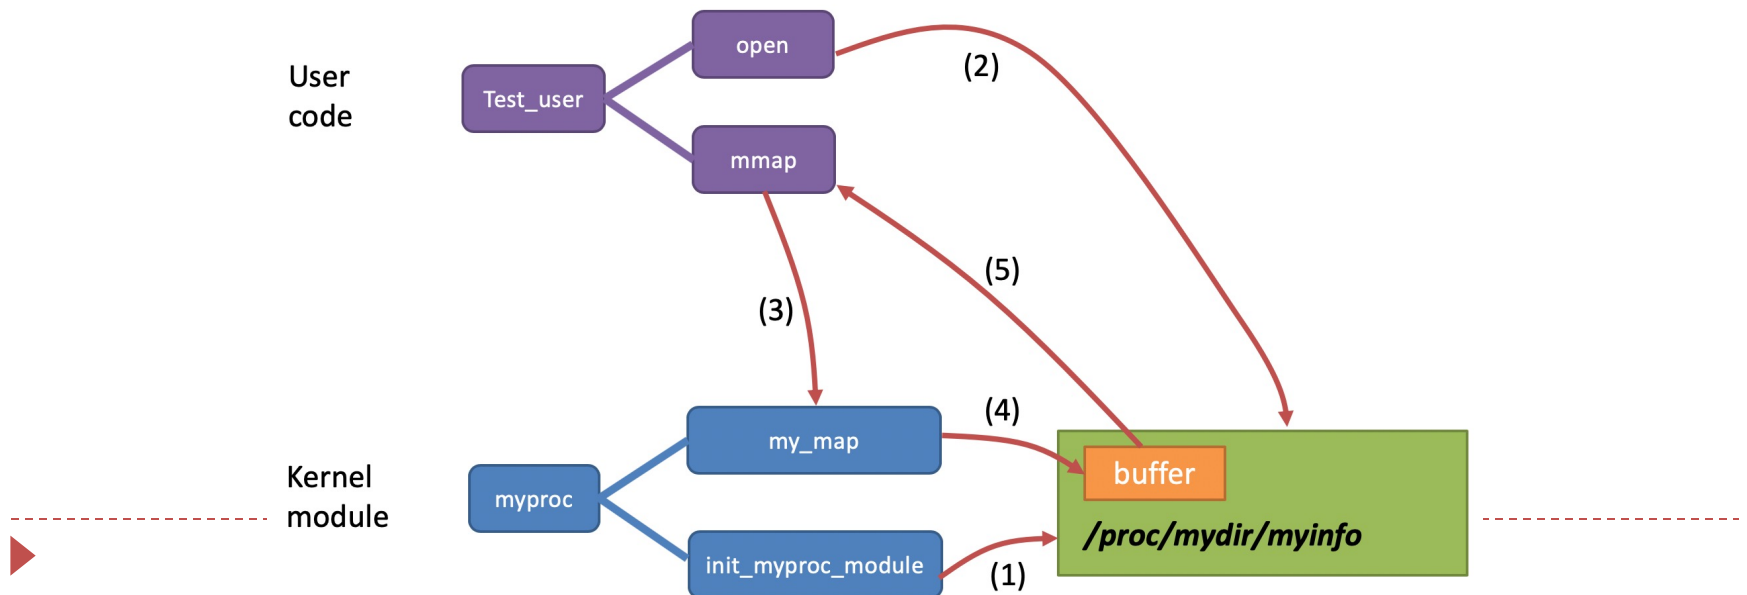
# Part 3: Exchange data between the user and kernel space via mmap

- You are required to write a user space program using mmap to visit the memory space of the proc file and print the data in that memory area.
- <https://github.com/kevinsuo/CS3502/blob/master/project-4-3-2.c>



# Part 3: Exchange data between the user and kernel space via mmap

1. Kernel module create a proc file: ***/proc/mydir/myinfo***
2. User process open the created proc file
3. User process calls mmap function, which further executed my\_map defined in the kernel
4. my\_map() then maps one piece of memory into user space (e.g., buffer) and puts some data inside
5. User process visits this piece of memory and prints the data out.



# Conclusion

---

- Part 1: Create a helloworld kernel module (20')
- Part 2: Create an entry in the /proc file system for user level read and write (30')
- Part 3: Exchange data between the user and kernel space via mmap (50')

