

HPC & Parallel Programming

Pthread Lab


Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

High Performance Computing: Pthread Project

Parallel Array Summation & Synchronization Analysis




The Objective

- **Task:** Sum array elements where $a[i] = 2*i$
- **Scale:** 10,000 elements
- **Threads:** Variable (1 to 10)
- **Output:** Execution Time + Global Sum


Array a[10000]	
0	0
1	2
2	4
3	6
4	8
...	
9999	19998

The Toolkit: Pthread Lifecycle



The Shared Memory Model

Process vs. Thread



THE PHANTOM BUG: WHEN 50,000 DOES NOT EQUAL 50,000

THE EXPECTATION

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

int counter = 0;

void *compute()
{
    int i = 0;
    while (i < 10000) {
        counter = counter + 1;
        i++;
    }
    printf("Counter value: %d\n", counter);
}

int main()
{
    pthread_t thread1, thread2, thread3, thread4, thread5;

    pthread_create(&thread1, NULL, compute, (void *)&thread1);
    pthread_create(&thread2, NULL, compute, (void *)&thread2);
    pthread_create(&thread3, NULL, compute, (void *)&thread3);
    pthread_create(&thread4, NULL, compute, (void *)&thread4);
    pthread_create(&thread5, NULL, compute, (void *)&thread5);

    pthread_exit(NULL);
    exit(0);
}
```



Expected Result: **50,000**

THE REALITY

```
pi@raspberrypi ~/Downloads> ./race_condition.o
Counter value: 14467
Counter value: 10410
Counter value: 12080
Counter value: 22745
Counter value: 32725
Counter value: 32,725
```

A race condition occurs when concurrent threads access shared data simultaneously. The order of access creates unpredictable results.

Diagnosing the Root Cause: The Critical Section


```
while (i < 100) {  
    counter = counter + 1;  
    i++;  
}
```




The Danger Zone

Critical Section: A section of code where a thread accesses a shared resource. If not protected, this is where the race condition lives.

Task 1: The Naive Architecture



The Danger: Race Conditions



Result: One increment is lost.
Final Value: 6 (Expected: 7)


```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
```

```
int counter = 0;

void *compute()
{
    int i = 0;
    while (i < 100) {
        counter = counter + 1;
        i++;
    }
    printf("Counter value: %d\n", counter);
}
```

Unprotected Shared Access

The Solution: Mutex Locks




Task 1 Implementation Flow

```
while (global_index < 10000) {  
    pthread_mutex_lock(&lock); ← Serialization Point  
    idx = global_index;  
    global_index++;  
    pthread_mutex_unlock(&lock); ← Serialization Point  
    val = 2 * idx;  
    // Add val to global_sum (requires lock) } } Parallel Execution  
} }
```


The diagram illustrates the execution flow of the provided C code. It features two red horizontal arrows pointing left from the `pthread_mutex_lock` and `pthread_mutex_unlock` calls, labeled "Serialization Point". Between these points is a red curly brace labeled "CriticalSection (Very Short)". After the `pthread_mutex_unlock` call, there is a blue curly brace labeled "Parallel Execution", which covers the assignment of `val` and the comment about adding it to `global_sum`.

The Bottleneck: High Contention




Threads spend more time fighting for the lock than computing.
The ratio of synchronization overhead to actual work is too high.

Task 2 Strategy: Chunking




Lock Frequency Reduced by 10x

Task 2 Implementation Logic




Task 3 Analysis: The Cost of Concurrency



For very small workloads, the overhead of managing threads exceeds the computation time saved.




Restoring Order: Mutual Exclusion (Locks)



Strategy 1: The Spinlock (Busy-Waiting)

The Spin: Constantly polling CPU




```
pthread_spin_lock(&slock);  
  
// Critical Section  
counter = counter + 1;  
pthread_spin_unlock(&slock);
```

Mechanism: The thread never sleeps. It remains active on the CPU, repeatedly asking 'Are you ready?' until the lock opens.


Deep Dive: Spinlocks

```
do {  
    // Entry section: Check and acquire lock  
    → while (atomic_flag_test_and_set(&lock_flag)) {  
        // Spin-wait: Busy-waiting  
    }  
  
    // Critical Section  
    critical_section_code();  
  
    // Exit section: Release lock  
    atomic_flag_clear(&lock_flag);  
} while (TRUE);
```




Efficiency Paradox: efficient only if the wait time is shorter than the time it takes to **context switch**. Ideal for the increment operations in Task 1.

The Cost of Spinning: Burning CPU Cycles




Strategy 2: The Mutex (Sleep & Wakeup)



```
Lock(mutex);  
// Critical Section  
Unlock(mutex);
```

A Mutex (Mutual Exclusion) allows the thread to sleep if the lock is **unavailable**. The OS wakes it up **only when** the resource is free.

Deep Dive: Mutex




The 'Sleep-and-Wakeup' cycle introduces significant OS overhead. If the Critical Section is just `i++`, the thread spends more time sleeping and waking than working.

Locking Mechanisms: Spinlock vs. Mutex

 Spinlock (Busy-Wait)	 Mutex (Sleep-Wake)
Behavior <pre>do { check lock } while (locked) } while (locked)</pre>	Behavior If locked -> Sleep -> Scheduler -> Wake
Pros No context switch, fast for short waits.	Pros Yields CPU to other tasks.
Cons Burns CPU cycles.	Cons High latency overhead.

High Contention Analysis



Because the integer addition is instant, putting a thread to sleep (Mutex) is wasteful. Spinning is preferred here.

Summary

Threads	Global Sum	Run Time (s)
1	10000000000	1.25
2	10000000000	0.85
3	10000000000	0.65
4	10000000000	0.55
5	10000000000	0.50
6	10000000000	0.48
7	10000000000	0.46
8	10000000000	0.45
9	10000000000	0.44
10	10000000000	0.44

- Correct Global Sum (Formula Validation)
- Task 1 (Naive) Performance Data
- Task 2 (Chunking) Performance Data
- Written Analysis (Single vs Multi, Spin vs Mutex)