

HPC & Parallel Programming

Introduction

Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

Outline

- Why study HPC & parallel programming?
- What to learn?
- Course structure
- Course policy
- An example of HPC & parallel programming



Self Introduction

- Kun Suo, Ph.D.
 - Homepage, <https://kevinsuo.github.io/>
 - Research interests:
 - Cloud computing and virtualization;
 - Parallel and Distributed Computation, containers and kubernetes;
 - Software defined network (SDN) and network function virtualization (NFV)
 - Big data systems and machine learning systems
 - Projects you may be interested in:
 - Several projects in Cloud & Data & Edge
 - <https://kevinsuo.github.io/code-lab.html>



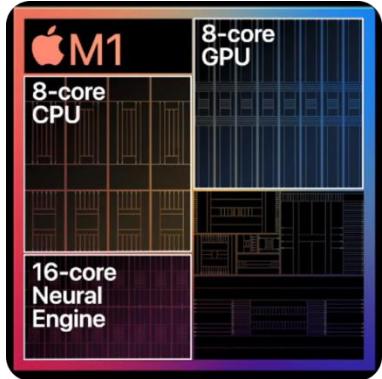
Now it's your turn

- Name, program/year, where from
 - Your interests in Computer Science
 - Have you ever used or heard of parallel and distributed system? Can you name some of them? What do you expect from this course?
- <https://www2.eecs.berkeley.edu/Research/Areas/CS/>

If you are in the online course, introduce yourself in D2L,
Discussions → Self-Introduction



Example of HPC & Parallel Programming



personal computer



internet



cloud



Social system



An Example of Parallel Computing

- CPU v.s. GPU

Single core

[NVIDIA: Adam and Jamie explain parallel processing on GPU's \(youtube.com\)](#)



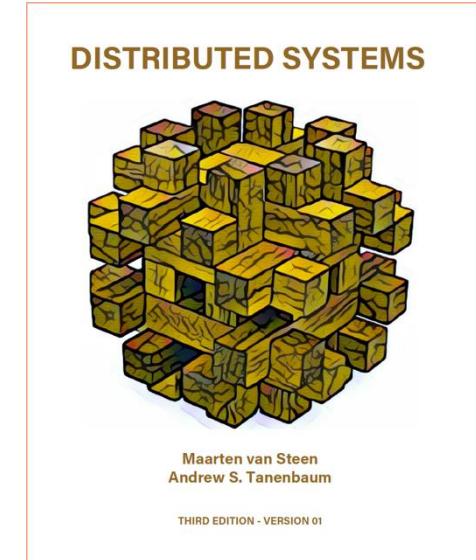
Course Information

- Instructor: Dr. Kun Suo
- Office: J-3230
- Email: ksuo@kennesaw.edu
 - Only reply to e-mails that are sent from KSU student email accounts and title the course number [CS4522]
- Office Hours:
 - Email or Microsoft Teams
 - By appointment
- Course Materials
 - Homework assignments, lecture slides, and other materials will be posted in the webpage (<https://kevinsuo.github.io/teaching.html>) and D2L.



Reference Book

- “Distributed Systems 3rd edition (2017)” by M. van Steen and A.S. Tanenbaum:
 - ISBN-13: 978-1543057386
 - You can get a digital copy of this book for free: <https://www.distributed-systems.net/index.php/books/ds3/>



Prerequisites

- Computer basics that are supposed to be covered in *(CS 3502) Operating Systems, (CS 3503) Computer Organization and Architecture course, (CS 4504) Parallel System course.*
- C programming (code reading, kernel development and debugging). ([Famous projects in C](#))
- Linux command line environment (compiling, Makefile, debugging, simple shell programming).



For C and Linux beginners

- C tutorial
 - <https://www.tutorialspoint.com/cprogramming/>
 - <https://www.learn-c.org>
 - <https://www.cprogramming.com/tutorial/c-tutorial.html>
- Linux tutorial
 - <https://ryanstutorials.net/linuxtutorial/>
 - <http://www.ee.surrey.ac.uk/Teaching/Unix/>
 - <https://www.tutorialspoint.com/unix/>



Project Environment

- Recommend project environment (local)

- VirtualBox + Ubuntu + Linux

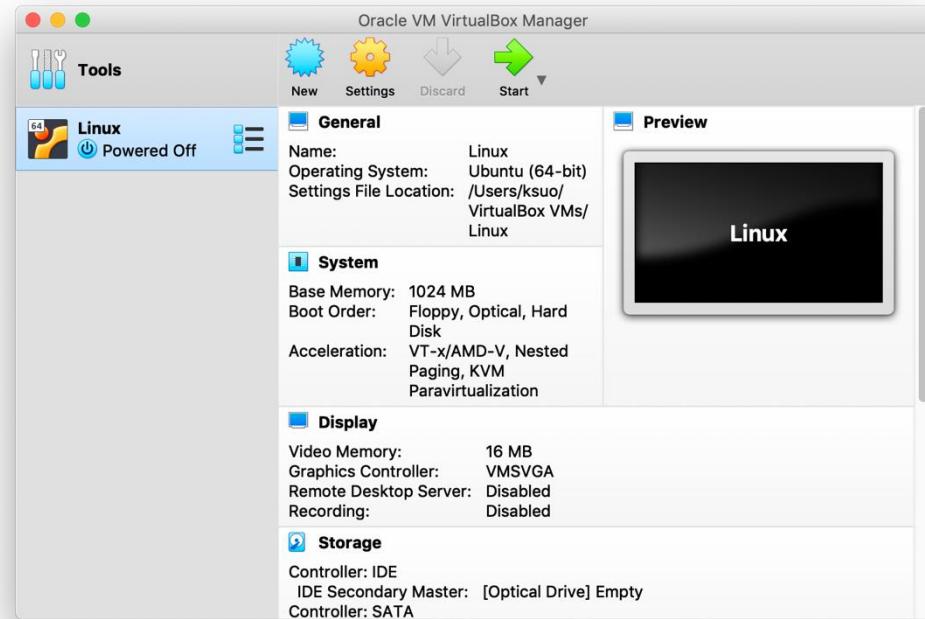
Virtual machine

VM OS

VM OS Kernel

<https://www.virtualbox.org/>

<https://ubuntu.com/download/desktop>



Project Environment

- Recommend project environment (local)
 - VirtualBox + Ubuntu + Linux
- New to VirtualBox?
 - <https://oracle-base.com/articles/vm/virtualbox-creating-a-new-vm>
 - https://www.youtube.com/watch?v=sB_5fqysi4
 - <https://youtu.be/GDoCrfPma2k> (MacOS)
- You can access to VMs in KSU data centers (cloud) through <https://cseview.kennesaw.edu/>,
 - username: administration; password: linuxadmin



Why study HPC & Parallel Programming?

- Most computer systems today are a certain form of HPC/parallel/distributed systems
 - Internet, datacenters, super computers, mobile devices
 - Most of the applications are parallel or even distributed apps
(example: debug decompress file app, [link](#), starts at 2:33)
- To learn useful techniques to build large systems
 - A system with 10,000 nodes is different from one with 100 nodes
- How to deal with imperfections
 - Machines can fail; network is slow; topology is not flat



What to learn

- HPC & Parallel Programming:
 - Parallel hardware
 - Matrix multiplication optimization
 - Pthread programming
 - MPI programming
 - OpenMP programming
 - GPU programming



Expected Outcomes

- Familiar with popular parallel programming libraries (Pthread, OpenMP, MPI, GPU)
- Familiar with fundamentals of program optimization
- The ability to
 - Evaluate the performance of parallel and HPC systems
 - Write simple parallel and HPC programs
 - Understand the tradeoffs in program design



Course Structure

- Lectures
 - Time/Location
 - D2L/Course website
- Projects
 - 5 programming assignments
 - 1 paper/project presentation
- Exams (open books)
 - Midterm: online D2L, TBA.
 - Final: online D2L, TBA



Course Policy

- Grading scale

Percentage	Grade
90 - 100	A
80 - 89	B
70 - 79	C
60 - 69	D
Below 60	F



Grading Policy (cont.)

- Grading percentage
 - Projects (x5): 50%
 - Presentation: 10%, including project or paper presentation
 - Midterm: 20%
 - Final exam: 20%

Late submission policy: late submission will **not be accepted** and **no credits**



Academic Integrity

- Academic dishonesty

- Cheating
 - Plagiarism
 - Collusion
 - The submission for credit of any work or materials that are attributable in whole or in part to another person
 - Taking an examination for another person
 - Any act designed to give unfair advantage to a student or the attempt to commit
- [https://scai.kennesaw.edu/KSU Codes
of Conduct 2019-2020.pdf](https://scai.kennesaw.edu/KSU_Codes_of_Conduct_2019-2020.pdf)
- Receiving, attempting to receive, knowingly giving or attempting to give unauthorized assistance...
- Do not upload course documents to 3rd party website without author's permission



Where to go for help ?

- Ask questions in class
- Ask questions outside class
 - Classmates and friends
- Attend office hours
 - Send Dr. Kun Suo emails or leave message on teams
- Search on the web
 - Stand on the shoulder of giants



HPC & Parallel Programming

Start from An Example

Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

An example of HPC & Parallel Programming: Matrix Multiplication

$$\begin{array}{c} \vec{b}_1 \quad \vec{b}_2 \\ \downarrow \quad \downarrow \\ \vec{a}_1 \rightarrow \begin{bmatrix} 1 & 7 \\ 2 & 4 \end{bmatrix} \cdot \begin{bmatrix} 3 & 3 \\ 5 & 2 \end{bmatrix} = \begin{bmatrix} \vec{a}_1 \cdot \vec{b}_1 & \vec{a}_1 \cdot \vec{b}_2 \\ \vec{a}_2 \cdot \vec{b}_1 & \vec{a}_2 \cdot \vec{b}_2 \end{bmatrix} \\ \vec{a}_2 \rightarrow \end{array}$$

A

B

C



Matrix multiply

<https://github.com/kevinsuo/CS7172/blob/master/matrix.c>

```
int main()
{
    initMatrix();

    double time_spent = 0.0;
    clock_t begin = clock();

    matrixMultiply();

    clock_t end = clock();
    time_spent += (double)(end - begin) / CLOCKS_PER_SEC;
    printf("Time elapsed is %f seconds", time_spent);

    return 0;
}
```



Matrix multiply

<https://github.com/kevinsuo/CS7172/blob/master/matrix.c>

```
#include <stdio.h>
#include <time.h>
#include <stdlib.h>
#define N 1000

double A[N][N], B[N][N], C[N][N];

void initMatrix()
{
    int i, j = 0;
    for (i = 0; i < N; i++) {
        for (j = 0; j < N; j++) {
            A[i][j] = rand() % 100 + 1; //generate a number between [1, 100]
            B[i][j] = rand() % 100 + 1; //generate a number between [1, 100]
        }
    }
}
```

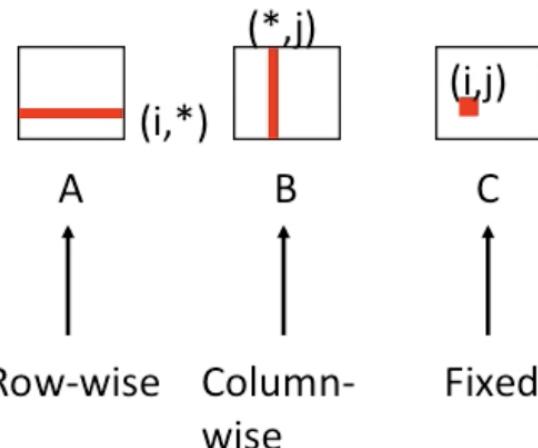


Matrix multiply

<https://github.com/kevinsuo/CS7172/blob/master/matrix.c>

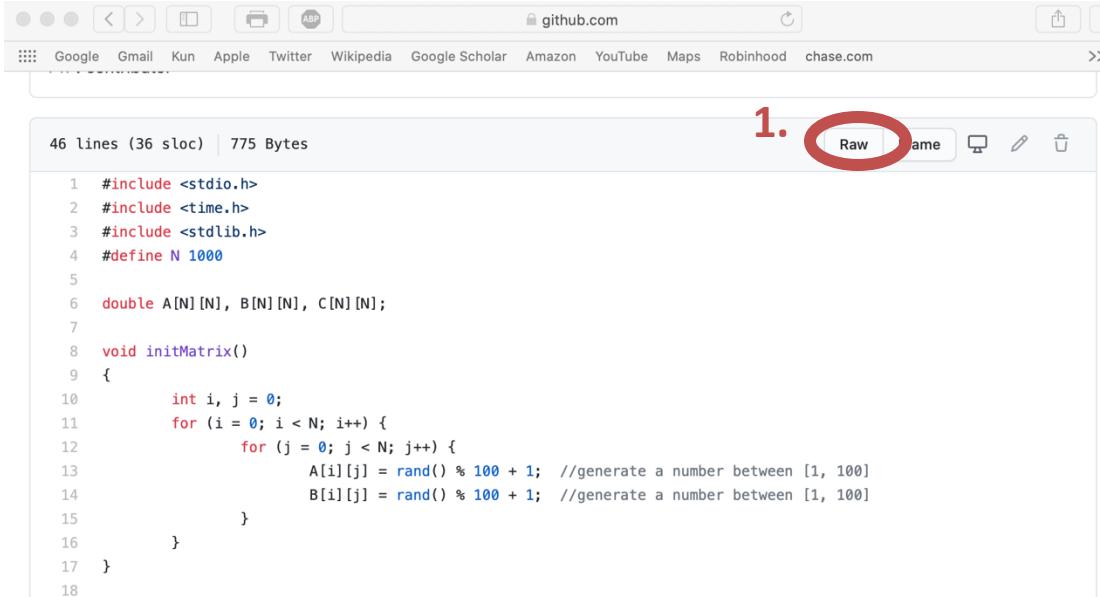
```
void matrixMultiply() {  
    int i, j, k = 0;  
    for (i = 0; i < N; i++) {  
        for (j = 0; j < N; j++) {  
            for (k = 0; k < N; k++) {  
                C[i][j] += A[i][k] * B[k][j];  
            }  
        }  
    }  
}
```

Inner loop:



Matrix multiply

<https://github.com/kevinsuo/CS7172/blob/master/matrix.c>

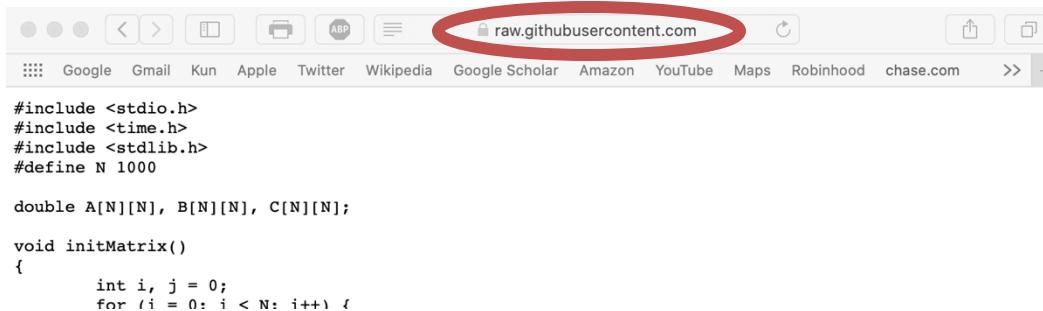


1. **Raw** (highlighted with a red circle)

```
46 lines (36 sloc) | 775 Bytes

1 #include <stdio.h>
2 #include <time.h>
3 #include <stdlib.h>
4 #define N 1000
5
6 double A[N][N], B[N][N], C[N][N];
7
8 void initMatrix()
9 {
10     int i, j = 0;
11     for (i = 0; i < N; i++) {
12         for (j = 0; j < N; j++) {
13             A[i][j] = rand() % 100 + 1; //generate a number between [1, 100]
14             B[i][j] = rand() % 100 + 1; //generate a number between [1, 100]
15         }
16     }
17 }
```

2. Copy the URL



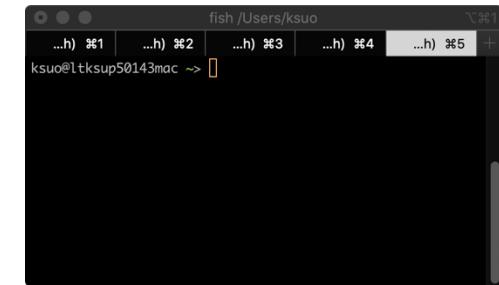
2. **Copy the URL**

raw.githubusercontent.com (highlighted with a red circle)

```
#include <stdio.h>
#include <time.h>
#include <stdlib.h>
#define N 1000

double A[N][N], B[N][N], C[N][N];

void initMatrix()
{
    int i, j = 0;
    for (i = 0; i < N; i++) {
```



3.

\$ wget URL

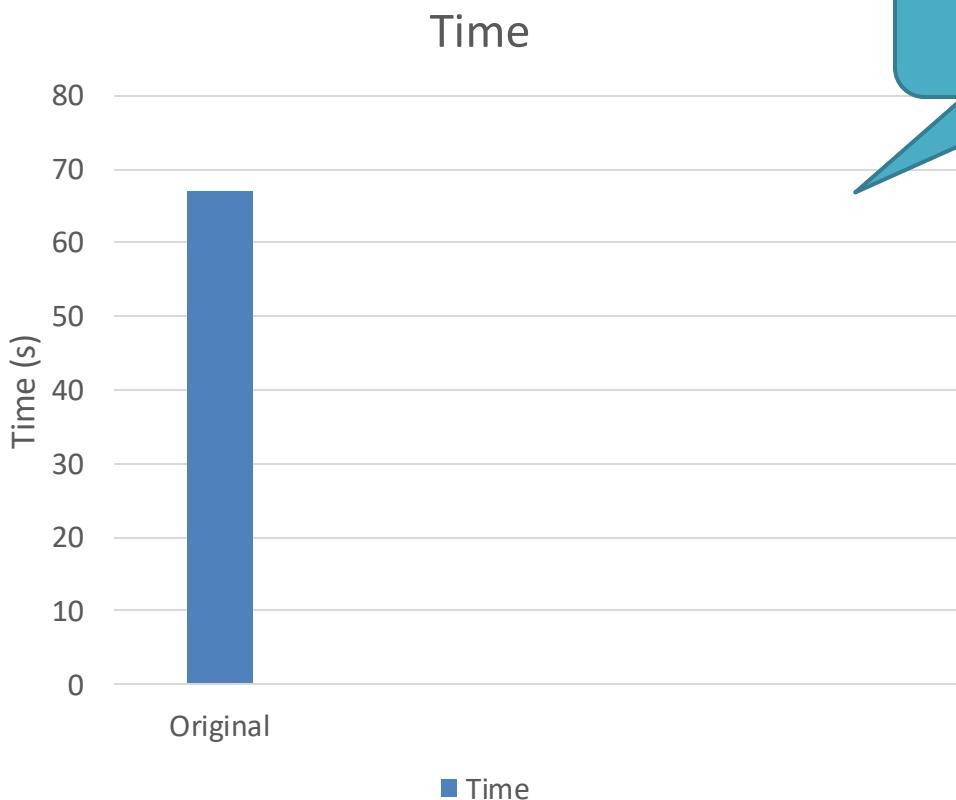
\$ gcc filename.c –o filename.o

\$./filename.o

(if no wget/gcc,
\$ sudo apt install wget, gcc)



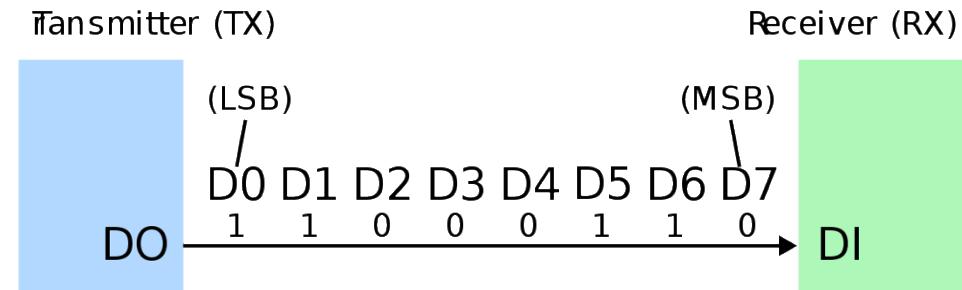
Matrix multiply



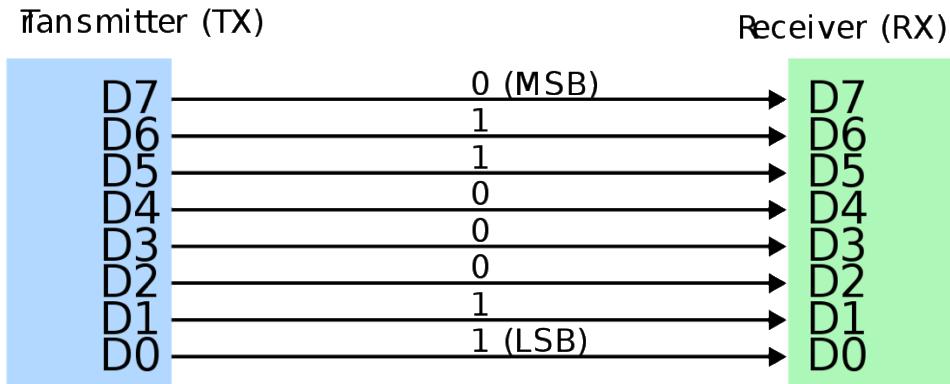
1. Accelerate serial execution
2. Accelerate in parallel



How to run it faster?



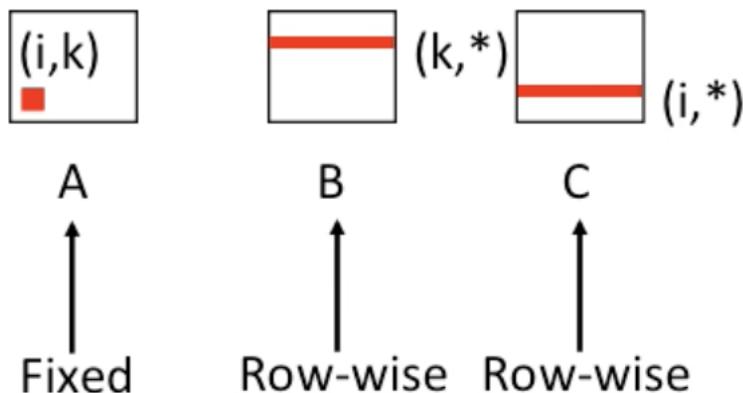
1. Accelerate serial execution
Reduce unnecessary steps



Option 1: Optimization using locality

```
void matrixMultiply() {  
    int i, j, k = 0;  
    for (k = 0; k < N; k++) {  
        for (i = 0; i < N; i++) {  
            for (j = 0; j < N; j++) {  
                C[i][j] += A[i][k] * B[k][j];  
            }  
        }  
    }  
}
```

Inner loop:



[https://github.com/kevinsuo/CS
7172/blob/master/matrix-opt.c](https://github.com/kevinsuo/CS7172/blob/master/matrix-opt.c)

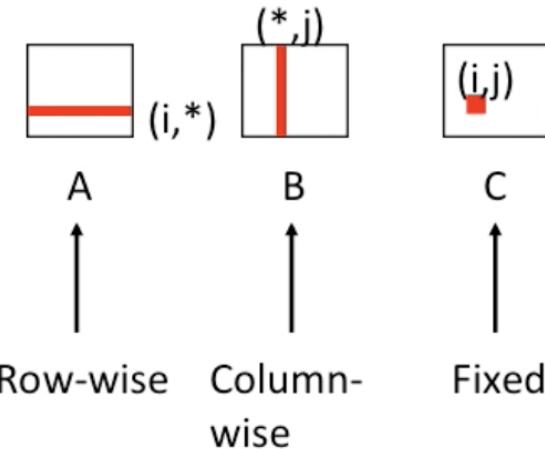
Option 1: Optimization using locality

```
ksuo@ksuo-VirtualBox ~/cs7172> ./a.o  
Time elapsed is 67.452589 seconds  
ksuo@ksuo-VirtualBox ~/cs7172>  
ksuo@ksuo-VirtualBox ~/cs7172>  
ksuo@ksuo-VirtualBox ~/cs7172>  
ksuo@ksuo-VirtualBox ~/cs7172>  
ksuo@ksuo-VirtualBox ~/cs7172>  
ksuo@ksuo-VirtualBox ~/cs7172> ./a2.o  
Time elapsed is 18.149353 seconds
```

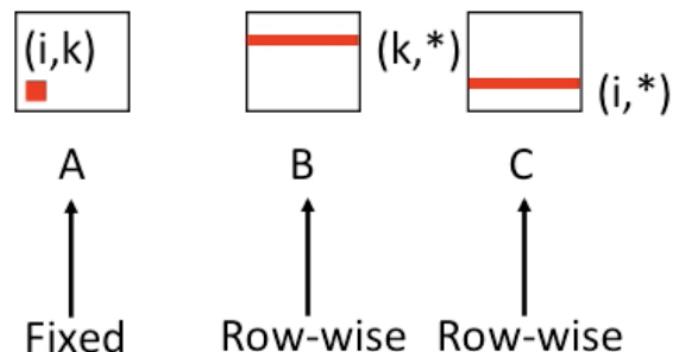
N=2000

3.7x

Inner loop:

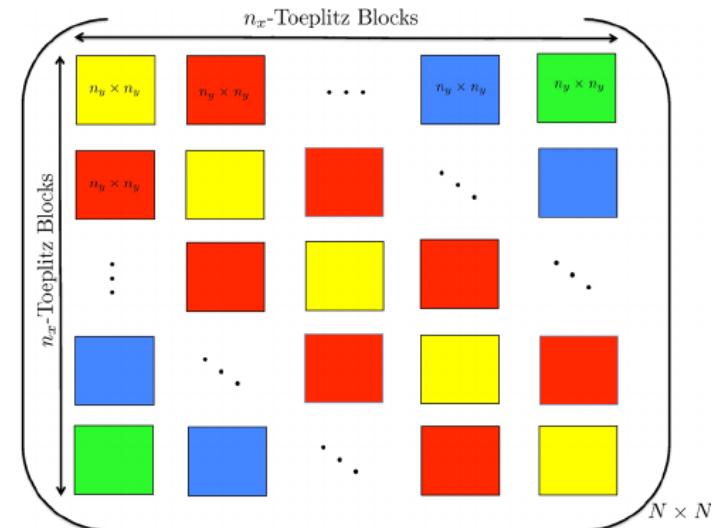
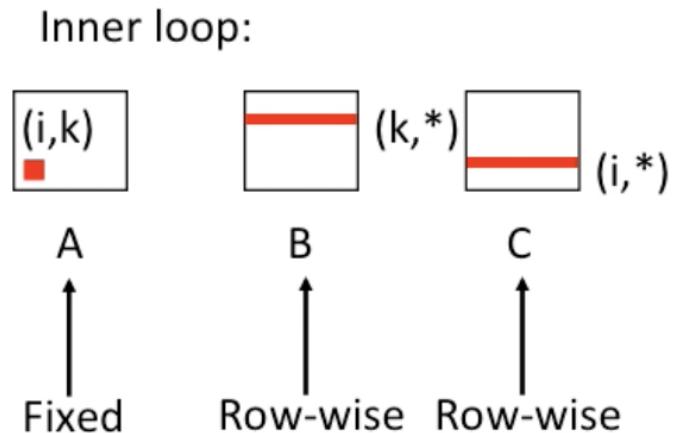


Inner loop:



Option 1: Optimization using locality

- Temporal locality
 - Every inner loop reuse the value of $A[i, k]$
- Spatial locality
 - Divide the large matrix into smaller ones and put it inside the cache during calculation



Option 1: Optimization using locality

```
void matrixMultiply() {  
    int i, j, k = 0;  
    int i2, j2, k2 = 0;  
  
    for (k2 = 0; k2 < N; k2+=BLOCK_SIZE) {  
        for (i2 = 0; i2 < N; i2+=BLOCK_SIZE) {  
            for (j2 = 0; j2 < N; j2+=BLOCK_SIZE) {  
                //inside each block  
                for (k = k2; k < k2+BLOCK_SIZE; k++) {  
                    for (i = i2; i < i2+BLOCK_SIZE; i++) {  
                        for (j = j2; j < j2+BLOCK_SIZE; j++) {  
                            C[i][j] += A[i][k] * B[k][j];  
                        }  
                    }  
                }  
            }  
        }  
    }  
}
```

<https://github.com/kevinsuo/CS7172/blob/master/matrix-opt2.c>

N = 2000

$$\left(\begin{array}{c|cc|cc} J_1 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & J_2 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & J_3 \\ 0 & 0 & 0 & 0 & 0 \end{array} \right)$$

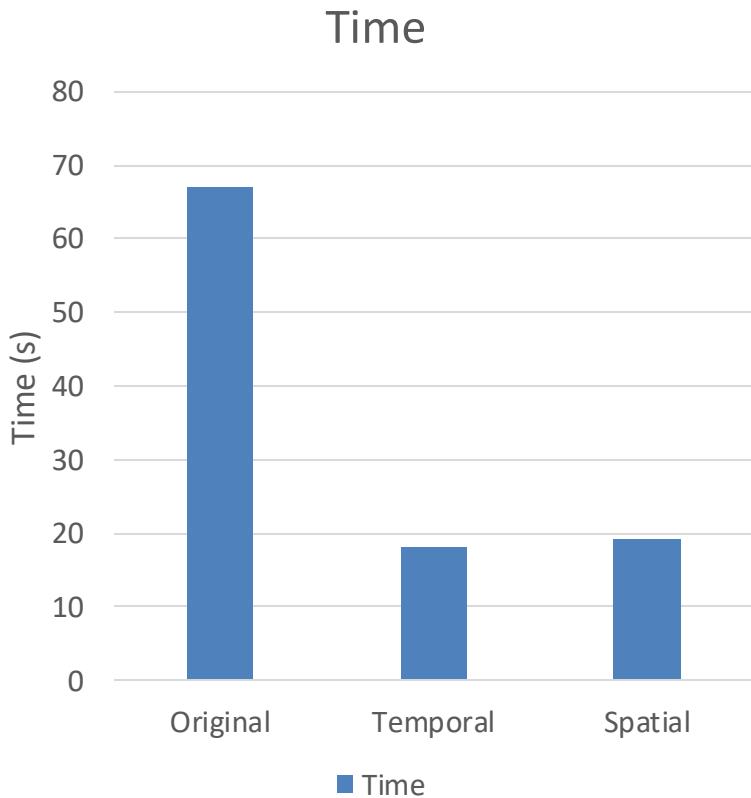


Option 1: Optimization using locality

$$A = \left(\begin{array}{cc|cc} a_{11} & a_{12} & a_{13} & a_{1n} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ \hline a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{array} \right) \implies \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix}$$
$$A_{11} = \begin{pmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{pmatrix}, A_{12} = \begin{pmatrix} a_{13} & a_{14} \\ a_{23} & a_{24} \end{pmatrix}$$
$$A_{21} = \begin{pmatrix} a_{31} & a_{32} \\ a_{41} & a_{42} \end{pmatrix}, A_{22} = \begin{pmatrix} a_{33} & a_{34} \\ a_{43} & a_{44} \end{pmatrix}$$



Option 1: Optimization using locality



```
ksuo@ksuo-VirtualBox ~/cs7172> ./a.o
Time elapsed is 67.845517 seconds ↵
ksuo@ksuo-VirtualBox ~/cs7172>
ksuo@ksuo-VirtualBox ~/cs7172>
ksuo@ksuo-VirtualBox ~/cs7172>
ksuo@ksuo-VirtualBox ~/cs7172>
ksuo@ksuo-VirtualBox ~/cs7172> ./a3.o
Time elapsed is 19.115410 seconds ↵
```



Optimal 2: Optimization using parallel

$$\begin{pmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{pmatrix} \cdot \begin{pmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{pmatrix} \rightarrow \begin{pmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{pmatrix}$$

(a)

$$\text{Task 1: } C_{1,1} = A_{1,1}B_{1,1} + A_{1,2}B_{2,1}$$

$$\text{Task 2: } C_{1,2} = A_{1,1}B_{1,2} + A_{1,2}B_{2,2}$$

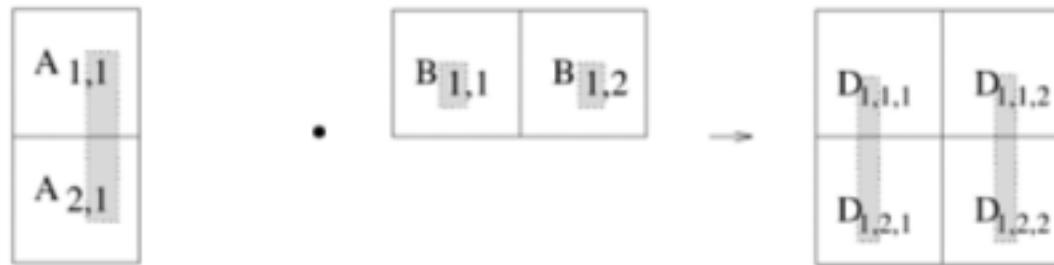
$$\text{Task 3: } C_{2,1} = A_{2,1}B_{1,1} + A_{2,2}B_{2,1}$$

$$\text{Task 4: } C_{2,2} = A_{2,1}B_{1,2} + A_{2,2}B_{2,2}$$

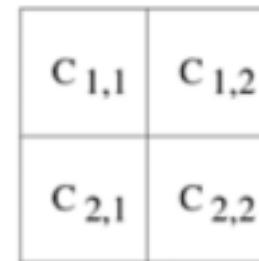
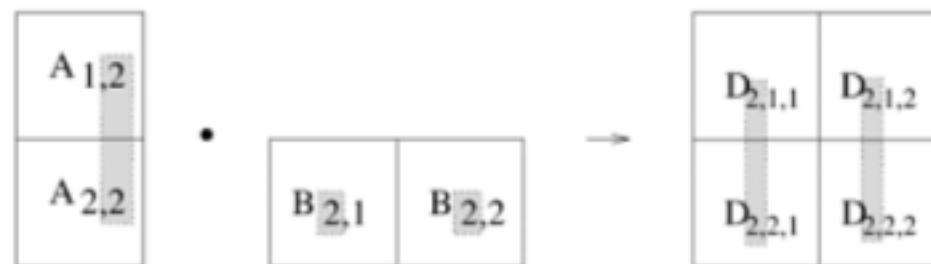


Optimal 2: Optimization using parallel

Thread 1:



Thread 2:



Reading Materials

- Strassen algorithm
 - In 1969, Volker Strassen proposed a matrix multiplication algorithm with a complexity of $O(n^{2.807})$: [Link](#). This was the first time in history that computational complexity of matrix multiplication was reduced below $O(n^3)$.
- Coppersmith–Winograd algorithm
 - In 1990, Don Coppersmith and Shmuel Winograd made a groundbreaking achievement by reducing the complexity to $O(n^{2.3727})$. Paper: [Link](#)



Optimization and Speedup

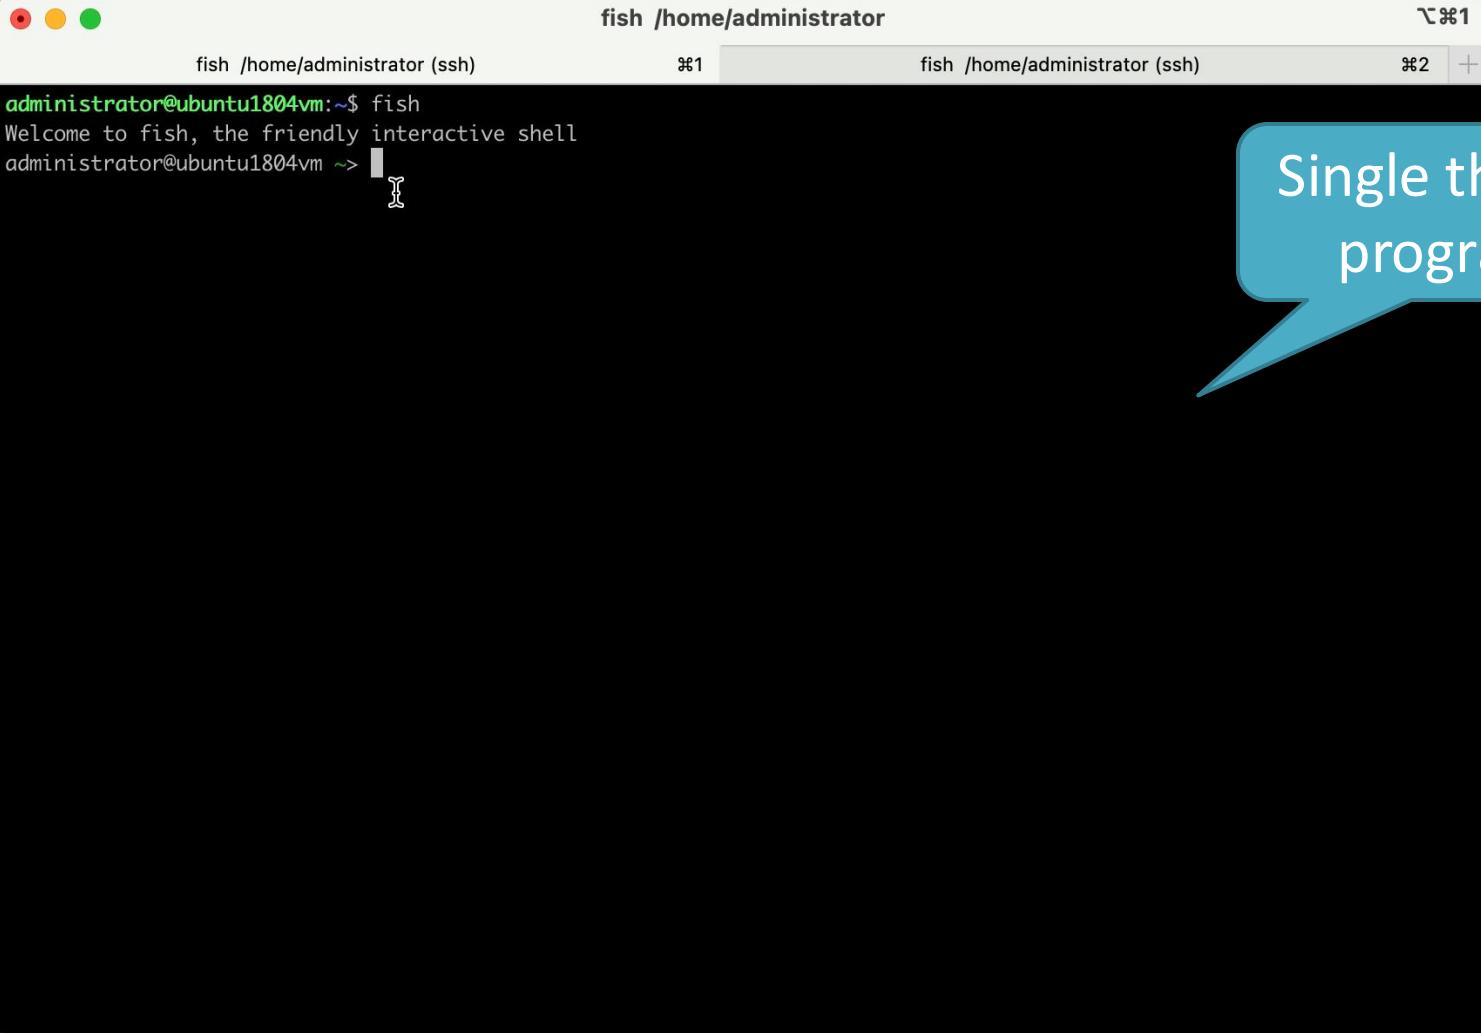
	N=200	N=400	N=800	N=1600
matrix				
matrix-opt1				
matrix-opt2				

	N=200	N=400	N=800	N=1600
matrix				
matrix-opt1				
matrix-opt2				



Single thread app demo

<https://youtu.be/dlsBhvQ9mA>



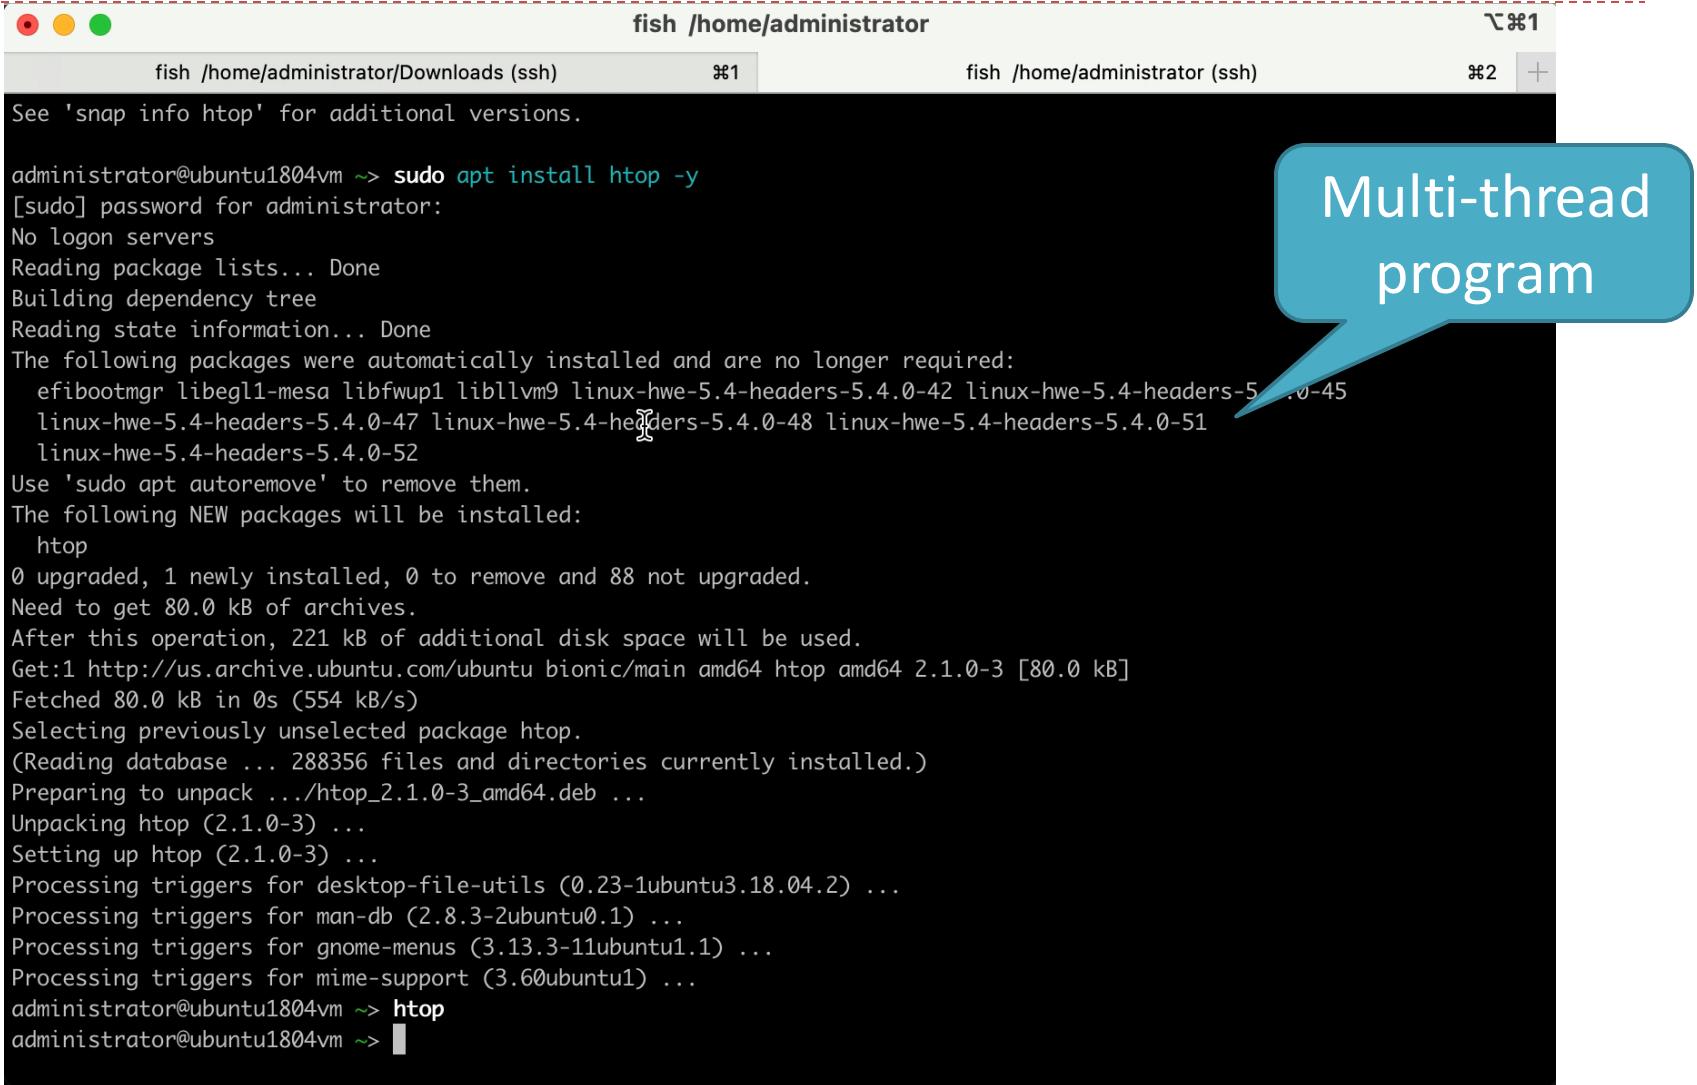
```
fish /home/administrator
fish /home/administrator (ssh)
administrator@ubuntu1804vm:~$ fish
Welcome to fish, the friendly interactive shell
administrator@ubuntu1804vm ~> |
```

Single thread
program



Multi-thread app demo

<https://youtu.be/ubLB2fb8cdc>

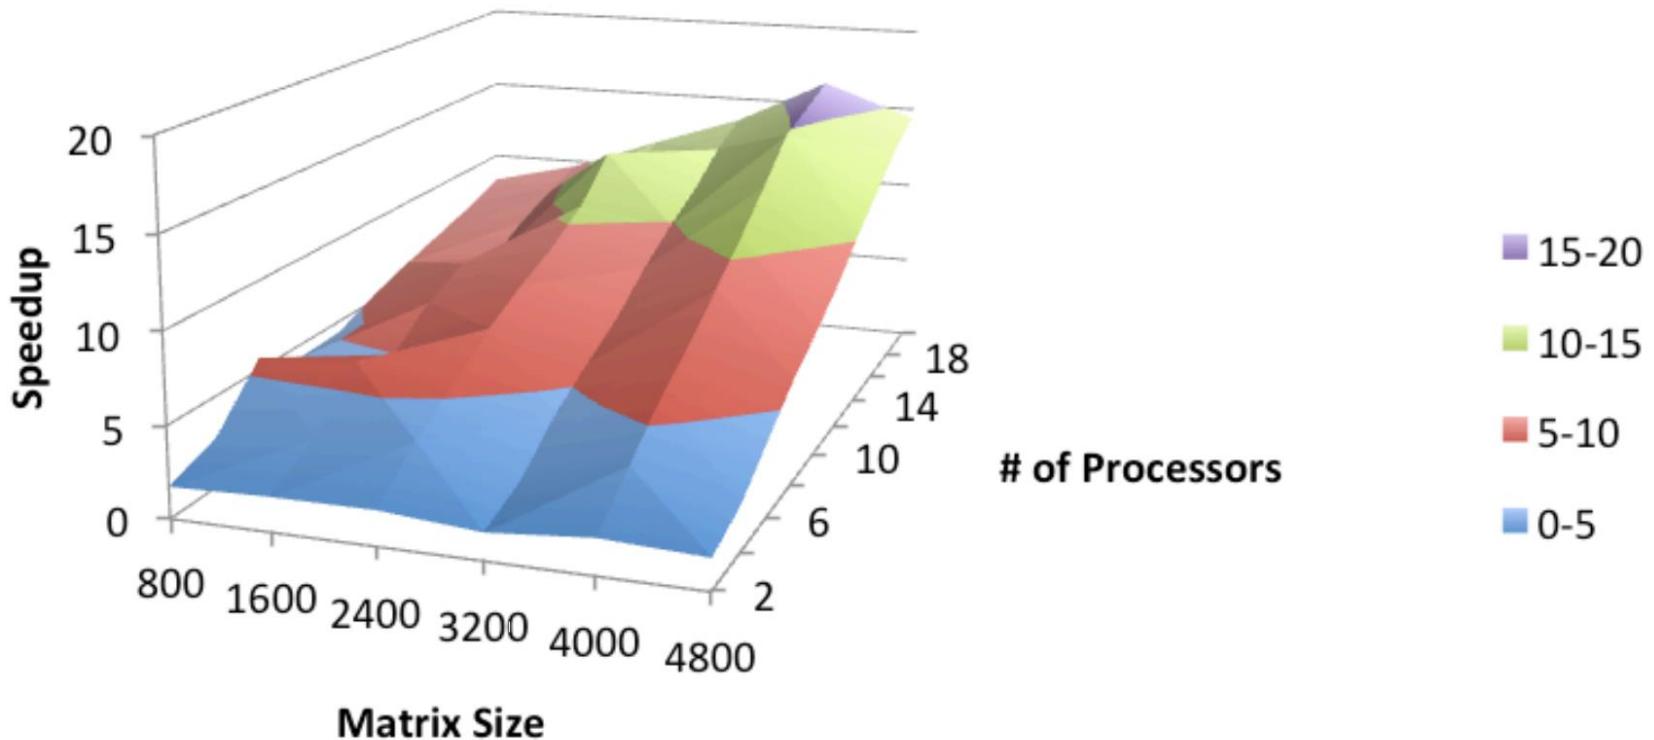


```
fish /home/administrator
fish /home/administrator/Downloads (ssh) #1
See 'snap info htop' for additional versions.

administrator@ubuntu1804vm ~> sudo apt install htop -y
[sudo] password for administrator:
No logon servers
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following packages were automatically installed and are no longer required:
  efibootmgr libegl1-mesa libfwup1 libllvm9 linux-hwe-5.4-headers-5.4.0-42 linux-hwe-5.4-headers-5.4.0-45
    linux-hwe-5.4-headers-5.4.0-47 linux-hwe-5.4-headers-5.4.0-48 linux-hwe-5.4-headers-5.4.0-51
    linux-hwe-5.4-headers-5.4.0-52
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
  htop
0 upgraded, 1 newly installed, 0 to remove and 88 not upgraded.
Need to get 80.0 kB of archives.
After this operation, 221 kB of additional disk space will be used.
Get:1 http://us.archive.ubuntu.com/ubuntu bionic/main amd64 htop amd64 2.1.0-3 [80.0 kB]
Fetched 80.0 kB in 0s (554 kB/s)
Selecting previously unselected package htop.
(Reading database ... 288356 files and directories currently installed.)
Preparing to unpack .../htop_2.1.0-3_amd64.deb ...
Unpacking htop (2.1.0-3) ...
Setting up htop (2.1.0-3) ...
Processing triggers for desktop-file-utils (0.23-1ubuntu3.18.04.2) ...
Processing triggers for man-db (2.8.3-2ubuntu0.1) ...
Processing triggers for gnome-menus (3.13.3-11ubuntu1.1) ...
Processing triggers for mime-support (3.60ubuntu1) ...
administrator@ubuntu1804vm ~> htop
administrator@ubuntu1804vm ~>
```



Optimal 2: Optimization using parallel



https://www.cse.unr.edu/~fredh/class/415/Nolan/matrix_multiplication/writeup.pdf



Example of distributed system: sorting

- Sorting on a single machine, e.g., Database

```
select field_a from table_b order by field_a limit 100, 10;
```

```
db.collection_b  
.find()  
.sort({"field_a":1})  
.skip(100)  
.limit(10);
```

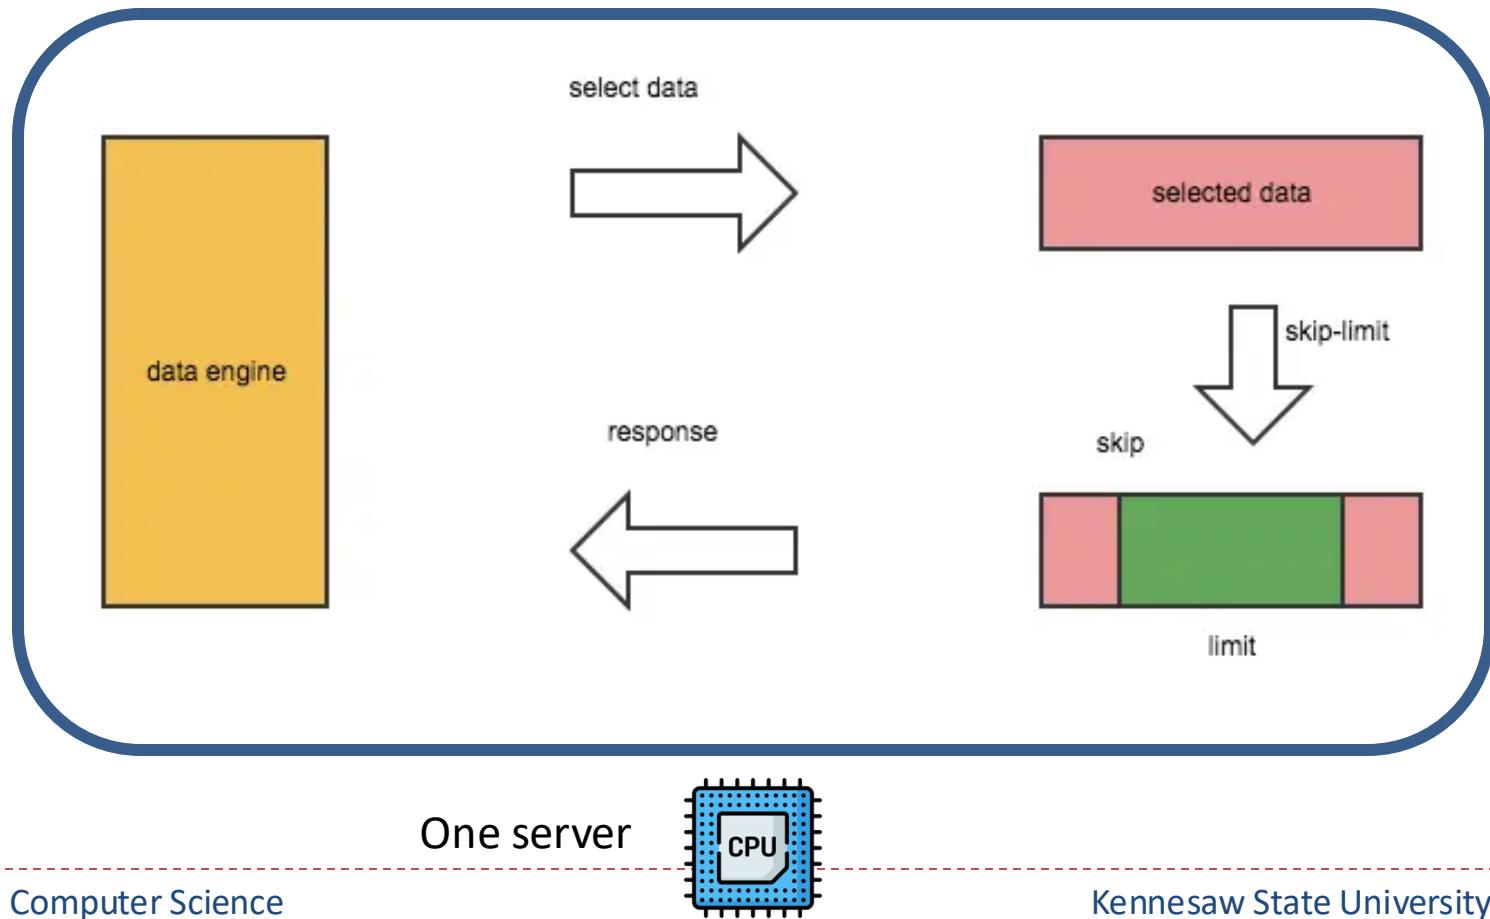
	field_a	field_b	field_c
100			
...			
...			
110			

From line 100
to the next 10
lines of data



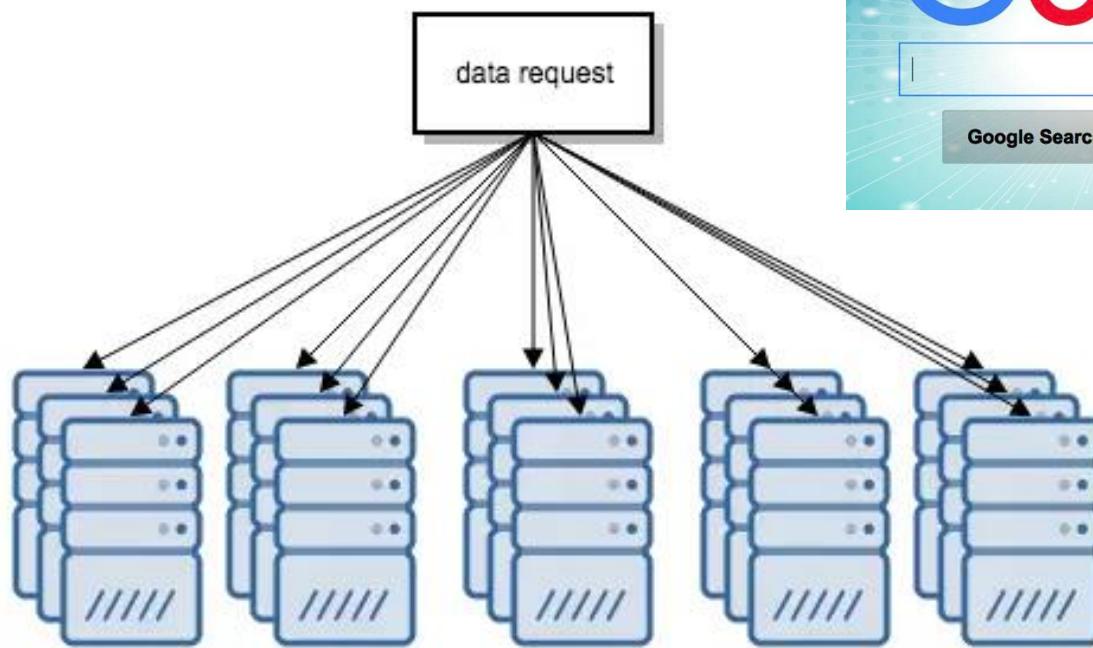
Example of distributed system: sorting

- Workflow on a single node



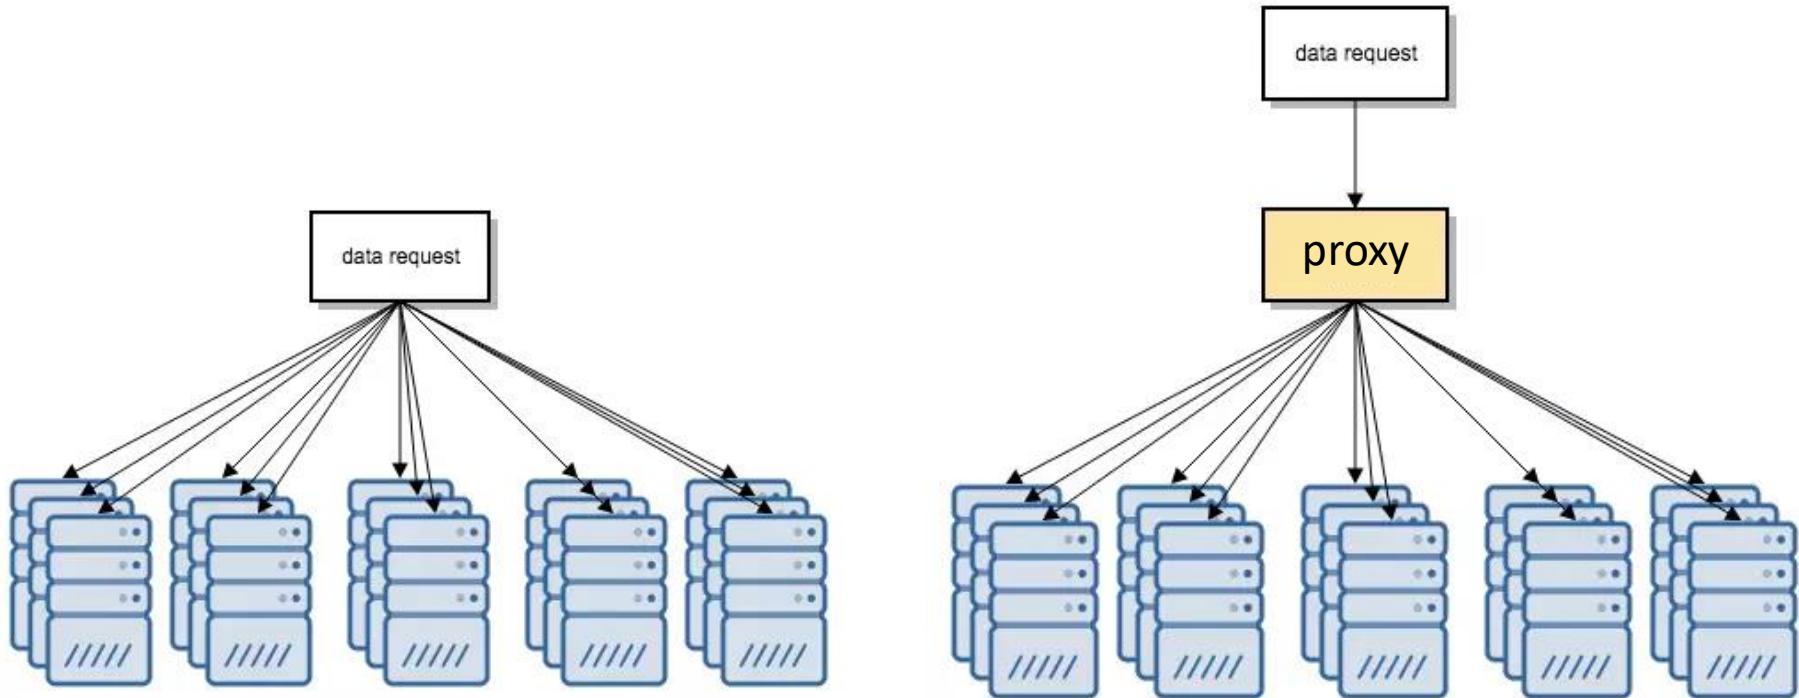
Example of distributed system: sorting

- If the data is too much and single node cannot hold



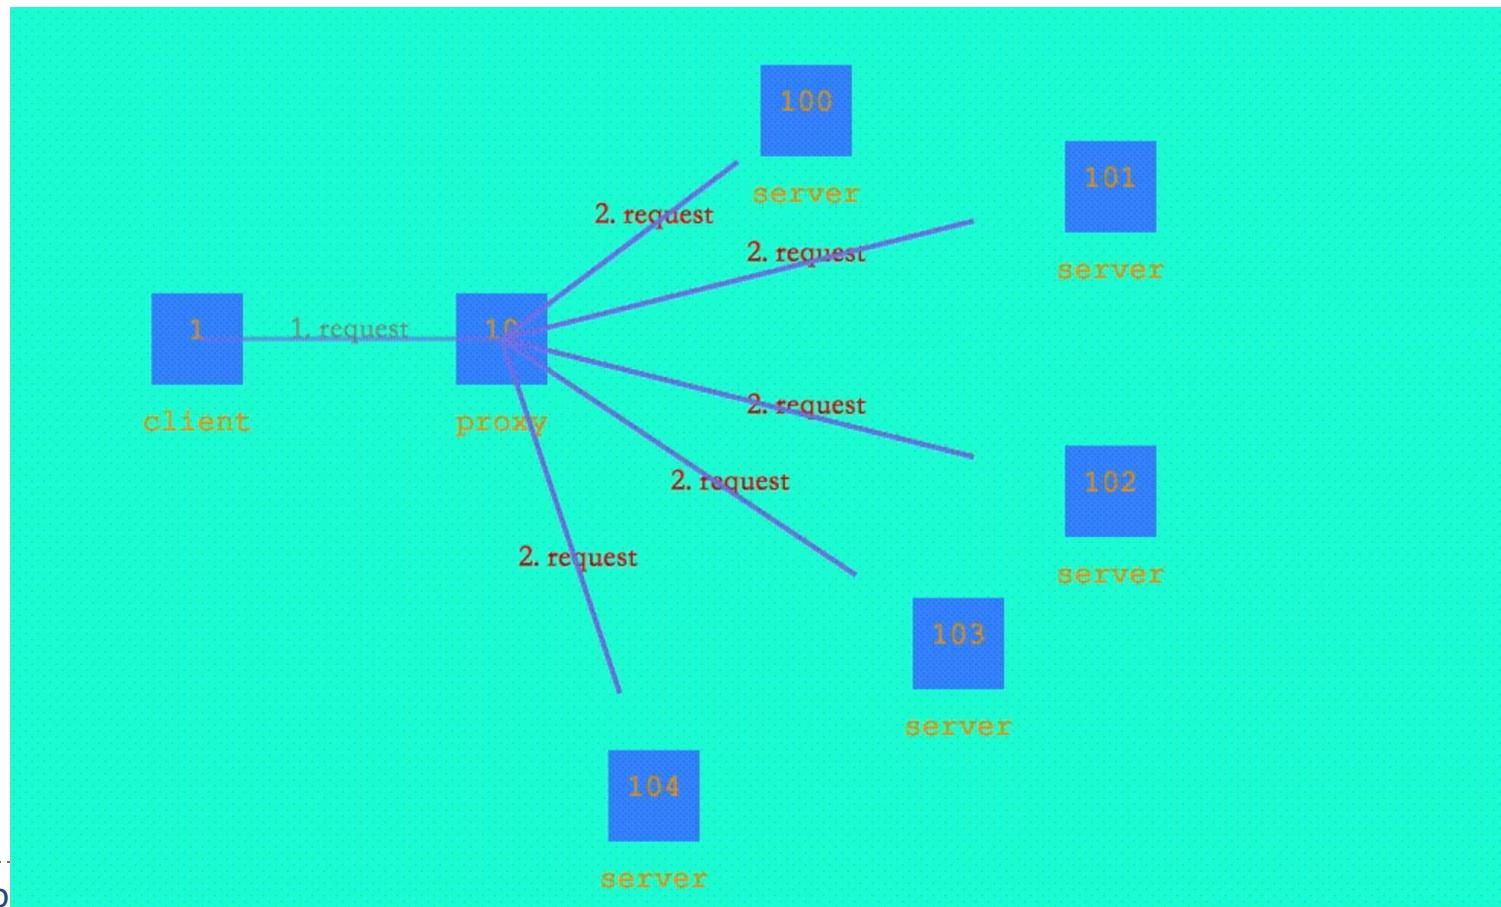
Example of distributed system: sorting

- Choose a node for merge processing



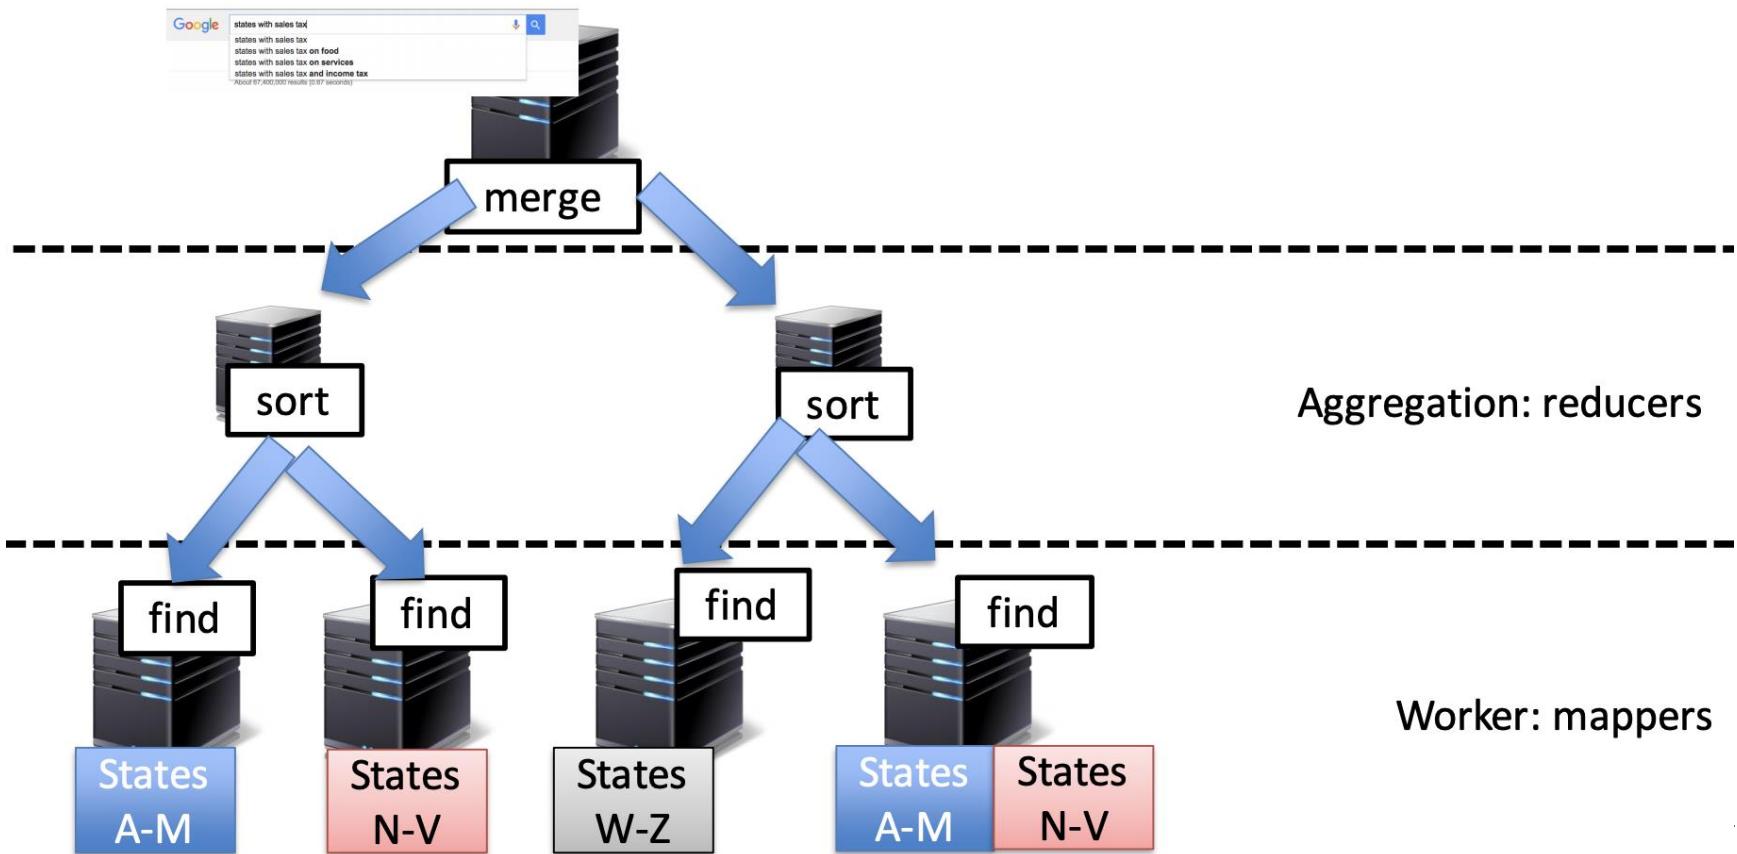
Example of distributed system: sorting

- Workflow



Example of distributed system: sorting

- How Do Requests Get Processed in a Data Center



Example of distributed system: sorting

- How Google Search Works
- <https://www.youtube.com/watch?v=0eKVizvYSUQ>



Conclusion

- Why study HPC & parallel programming?
- What to learn?
- Course structure
- Course policy
- An example of HPC & parallel programming

