

HPC & Parallel Programming

Thread

Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

Outline

- What is thread?
 - Multiple thread application
 - Thread vs Process
 - Advantage and disadvantage of thread
- Thread in Linux
- Thread design
 - Kernel space vs User space
 - Local thread vs Global thread scheduling



Process review

- Definition
 - An instance of a *program* running on a computer
 - An *abstraction* that supports running programs - -> cpu virtualization
 - An *execution stream* in the context of a particular *process state* - -> dynamic unit
 - A *sequential* stream of execution in its *own address space* - -> execution code line by line




What is a thread?

- Thread
 - A finer-granularity entity for execution and parallelism
 - Lightweight process
 - A program in execution without dedicated address space
- Multithreading
 - Running multiple threads within a single process



Finer-granularity entity



Process : thread = 1: N





What is a thread?

- Thread
 - A finer-granularity entity for execution and parallelism
 - Lightweight process
 - A program in execution without dedicated address space
- Multithreading
 - Running multiple threads within a single process




Process review

- Two parts of a process
 - Sequential execution of instructions
 - Process state
 - ▶ registers: PC (program counter), SP (stack pointer), ...
 - ▶ Memory: address space, code, data, stack, heap ...
 - ▶ I/O status: open files ...




Lightweight process

Occupy more memory,
complex switching (e.g., save old data, switch to new data),
low CPU utilization (e.g., slow context switch)



Three processes



Three threads




What is a thread?

- Thread
 - A finer-granularity entity for execution and parallelism
 - Lightweight process
 - A program in execution without dedicated address space
- Multithreading
 - Running multiple threads within a single process




Dedicated address space


0x111111




0x222222



0x333333



0x666666



Three processes

Three threads




What is a thread?

- Thread
 - A finer-granularity entity for execution and parallelism
 - Lightweight process
 - A program in execution without dedicated address space
- Multithreading
 - Running multiple threads within a single process



What is a thread? Multithreading apps

Programs, Apps, Processes & Threads



* this image may not reflect the reality for the show-cased apps

What is a thread? Multithreading apps


Activity Monitor (All Processes)

Process Name	% CPU	CPU Time	Threads	Idle Wake Ups	PID	User
Activity Monitor	7.0	1:21.83	7	17	691	ksuo
hidd	3.8	13:18.76	5	1	99	_hidd
coreaudiod	1.9	1:44.86	6	55	151	_coreaudiod
Google Chrome	1.3	37:38.24	32	4	683	ksuo
sysmond	0.9	9:24.21	3	1	260	root
launchd	0.8	2:04.06	2	1	1	root
Dock	0.7	30.94	3	0	696	ksuo
iconservicesagent	0.6	12.97	5	0	351	ksuo
Google Chrome Helper	0.6	30.03	8	19	1951	ksuo
Microsoft Outlook	0.4	12:19.10	23	16	686	ksuo
Backup and Sync from Google	0.4	7:55.60	46	17	830	ksuo
Google Chrome Helper	0.2	5:56.29	8	5	739	ksuo
欧路词典	0.2	26.01	8	3	740	ksuo
Microsoft PowerPoint	0.2	7:24.65	19	5	689	ksuo
powerd	0.1	40.42	2	1	57	root
Google Chrome Helper	0.1	6.01	14	1	9082	ksuo
Finder	0.1	1:57.51	7	1	698	ksuo
mds	0.1	1:58.49	8	3	67	root
fseventsds	0.0	43.58	10	3	47	root
launchservicesd	0.0	29.57	4	0	79	root
mds_stores	0.0	2:51.27	5	0	238	root
smcFanControl	0.0	54.43	7	1	831	ksuo
...

System: 8.14% CPU LOAD User: 12.30% Threads: 2105 Idle: 79.56% Processes: 519

Multithreading apps

<https://youtu.be/Me0mL44TxW0>



Example 1: A word process with three threads




Four score and seven years ago, our fathers brought forth upon this continent a new nation, conceived in liberty, and dedicated to the proposition that all men are created equal. Now we are engaged in a great civil war, testing whether that nation, or any nation so conceived and so dedicated, can long endure. We are met on a great battlefield of that war.

We have come to dedicate a portion of that field as a final resting place for those who here gave their lives that this nation might live. It is for us to be dedicated to the unfinished work which they who fought here have thus far so nobly advanced. It is rather for us to be here dedicated to the great task remaining before us, that from these honored dead we take increased devotion to that cause for which they gave the last full measure of devotion, that we here highly resolve that these dead shall not have died in vain that this nation, under God, shall have a new birth of freedom and that government of the people by the people, for the people.


Second thread handles screen display

Why not multiple processes ?

First thread handles keyboard input



Keyboard




Third thread handles saving the document to disk

A word process with three threads.



Example 2: a multi-threaded web server



Thread Type 1: A dispatcher thread waits for and accepts network connections




A multithreaded Web server.



Example 2: a multi-threaded web server




Example 2: a multi-threaded web server





Why not multiple processes ?

Occupy more memory,
complex switching
low CPU utilization
Difficult to communicate
Expensive data sharing



Three processes



Three threads



Why multiple threads

- Good example from Wikipedia: multiple threads within a single process are like multiple cooks trying to prepare the same meal together.



- Each one is doing one thing.
- They are probably doing different things.
- They all share the same recipe but may be looking at different parts of it.
- They have private state but can communicate easily.
- They must coordinate!



Outline

- What is thread?
 - Multiple thread application
 - Thread vs Process
 - Advantage and disadvantage of thread
- Thread in Linux
- Thread design
 - Kernel space vs User space
 - Local thread vs Global thread scheduling




Process vs. Thread comparison

	Process	Thread
Concurrency and protection	?	?
Data structure	?	?
Performance	?	?



Processes vs. Threads (Concurrency and protection)




Processes vs. Threads (Concurrency and protection)

	Process	Threads
Concurrency	Parallel execution stream of instructions 	Maintain parallel execution stream of instructions 
Protection		



Processes vs. Threads (Concurrency and protection)




Processes vs. Threads (Concurrency and protection)

	Process	Threads
Concurrency	Parallel execution stream of instructions 	Maintain parallel execution stream of instructions 
Protection	A dedicated address space 	Share address space with other threads 

Separate concurrency from protection


Processes vs. Threads (Data structure)



Process	Thread
Have data/code/heap	
Include at least one thread	
Have own address space, isolated from other processes	




Processes vs. Threads (Data structure)



Context	Process	Thread
File pointer	*	
Stack	*	*
Memory	*	
State	*	*
Priority	*	*
I/O state	*	
Authority	*	




Processes vs. Threads (Data structure)



Context	Process	Thread
Scheduling	*	
Statistics	*	
File description	*	
Read/Write pointer	*	
Event/Signal	*	
Registers	*	*



Processes vs. Threads (Performance)




Process	Thread
Expensive to create	Inexpensive to create
Expensive context switching	Inexpensive context switching
IPC can be expensive	Efficient communication



Paper: Performance of Multi-Process and Multi-Thread Processing on Multi-core SMT Processors

- Our results showed that both models (multi-process vs. multi-thread) achieved almost comparable performance, whereas the multi-thread model achieved much **better SMT scalability** and higher performance.


<https://pdfs.semanticscholar.org/d54/a215131c5eacd997708c5c4612345fe989c7.pdf>



Demo

- <https://github.com/tsuna/contextswitch>
 - timectxsw: Benchmarks the overhead of context switching between 2 processes.
 - timetctxsw: Benchmarks the overhead of context switching between 2 threads.

```
$ gcc xxx.c -o xxx.o -pthread  
$ ./xxx.o
```



A terminal window showing the results of benchmarks for context switching. The window title is "fish /home/NewUsersDir/ksuo/ksuo/contextswitch". The terminal shows two commands being run:

```
ksuo@ccsegpu4 ~/contextswitch (master)> ./timectxsw.o  
2000000 process context switches in 5001406293ns (2500.7ns/ctxsw)  
ksuo@ccsegpu4 ~/contextswitch (master)> ./timetctxsw.o  
2000000 thread context switches in 3958269002ns (1979.1ns/ctxsw)
```

Outline

- What is thread?
 - Multiple thread application
 - Thread vs Process
 - Advantage and disadvantage of thread
- Thread in Linux
- Thread design
 - Kernel space vs User space
 - Local thread vs Global thread scheduling



Advantages of Threads

- Efficient creation
 - Only create the thread context
- Express concurrency
 - Lightweight, better performance, higher scalability
- Efficient communication
 - Communication can be carried out via shared data objects within the shared address space



Disadvantages of Threads

- Shared data - -> Security
 - Global variables are shared between threads.
 - Accidental data changes can cause errors.
- Lack of robustness
 - Crash in one thread will crash the entire process.
- Some library functions may not be thread-safe
 - Library Functions that return pointers to static internal memory. E.g. `gethostbyname()`



Test

- The following statements about threads, which ones are correct (ABCDE).
 - A. The thread is introduced to improve the execution efficiency of the system and reduce the idle time of the processor and the scheduling switching time
 - B. Thread is the basic unit independently scheduled by the system
 - C. The thread itself basically does not own system resources, but it can share all the resources owned by the process with other threads belonging to the same process
 - D. Thread is also called lightweight process
 - E. Multiple threads in the same process can execute concurrently



Test

- A process can contain multiple threads, each thread (A).
 - A. Shares process virtual address space
 - B. The address space of each thread is completely independent
 - C. is the unit of resource allocation
 - D. Shares stack



Test

- In the following description, (D) is not a feature of multi-threaded systems.
 - A. Use threads to perform matrix multiplication operations in parallel
 - B. Web server uses thread to request http service
 - C. GUI-based debugger uses different threads to process user input, calculation, tracking and other operations
 - D. The keyboard driver creates a thread for each running process in response to the corresponding keyboard input



Test

- The following description of the comparison between processes and threads (D) is wrong:
 - A) The time overhead of concurrent execution of threads is less than that of processes
 - B) Inter-thread communication is simpler than process
 - C) Threads have fewer resources than processes
 - D) Thread is more stable than process



Outline

- What is thread?
 - Multiple thread application
 - Thread vs Process
 - Advantage and disadvantage of thread
- Thread in Linux
- Thread design
 - Kernel space vs User space
 - Local thread vs Global thread scheduling



Processes vs. Threads in Linux

- On Mon, 5 Aug 1996, Peter P. Eiserloh wrote:
*We need to keep a clear the concept of threads.
Too many people seem to confuse a thread with
a process. The following discussion does not
reflect the current state of linux, but rather is an
attempt to stay at a high level discussion.*

- <http://lkml.iu.edu/hypermail/linux/kernel/9608/0191.html>
- The way Linux thinks about this (and the way I want things to work) is that there _is_ no such thing as a “process” or a “thread”. There is only the totality of the COE (called "task" by Linux).



Thread creation: clone(), not fork()

- Create new threads in Linux
 - Use `clone()` to create threads instead of using `fork()`
 - `clone()` is usually not called directly but from some threading libraries, such as `pthread`.
 - <http://man7.org/linux/man-pages/man2/clone.2.html>

```
int main(int argc, char *argv[])
{
    char *stack;                                /* Start of stack buffer */
    char *stackTop;                             /* End of stack buffer */
    pid_t pid;
    struct utsname uts;

    if (argc < 2) {
        fprintf(stderr, "Usage: %s <child-hostname>\n", argv[0]);
        exit(EXIT_SUCCESS);
    }

    /* Allocate stack for child */

    stack = malloc(STACK_SIZE);
    if (stack == NULL)
        errExit("malloc");
    stackTop = stack + STACK_SIZE; /* Assume stack grows downward */

    /* Create child that has its own UTS namespace;
       child commences execution in childFunc() */


    pid = clone(childFunc, stackTop, CLONE_NEWUTS | SIGCHLD, argv[1]);
    if (pid == -1)
        errExit("clone");
    printf("clone() returned %ld\n", (long) pid);

    /* Parent falls through to here */
}
```



Clone system call

- No. 56, <https://filippo.io/linux-syscall-table/>



A screenshot of a web browser window titled "filippo.io". The page displays a table of Linux system calls, their corresponding kernel functions, and the source files where they are implemented. The table has four columns: number, user-space name, kernel-space name, and source file. The "clone" row is highlighted with a blue background.


51	getsockname	sys_getsockname	net/socket.c
52	getpeername	sys_getpeername	net/socket.c
53	socketpair	sys_socketpair	net/socket.c
54	setsockopt	sys_setsockopt	net/socket.c
55	getsockopt	sys_getsockopt	net/socket.c
56	clone	stub_clone	kernel/fork.c
57	fork	stub_fork	kernel/fork.c
58	vfork	stub_vfork	kernel/fork.c
59	execve	stub_execve	fs/exec.c
60	exit	sys_exit	kernel/exit.c
61	wait4	sys_wait4	kernel/exit.c
62	kill	sys_kill	kernel/signal.c



Clone system call

Pthread_create()
in the user space

Clone()
in the kernel



Pthread Creation Example

```
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

void *myThread1(void)
{
    int i;
    for (i=0; i<3; i++)
    {
        printf("This is the 1st pthread.\n");
        sleep(1);
    }
}

int main()
{
    int ret=0;
    pthread_t id1;

    printf("This is main thread!\n");

    ret = pthread_create(&id1, NULL, (void*)myThread1, NULL);
    if (ret)
    {
        printf("Create pthread error!\n");
        return 1;
    }

    pthread_exit(NULL);

    return 0;
}
```


[https://github.com/kevinsuo/CS7172
/blob/master/pthread_create.c](https://github.com/kevinsuo/CS7172/blob/master/pthread_create.c)

```
pi@raspberrypi ~/Downloads> ./test.o
This is main thread!
This is the 1st pthread.
This is the 1st pthread.
This is the 1st pthread.
```

Clone() in the kernel

Fork vs clone

- Fork:
 - All resources in PCB are **copied** from parent process to child process
- Clone:
 - The resources in PCB are **partly copied** from one thread to another thread (the copied context is different in vfork)




Outline

- What is thread?
 - Multiple thread application
 - Thread vs Process
 - Advantage and disadvantage of thread
- Thread in Linux
- Thread design
 - Kernel space vs User space
 - Local thread vs Global thread scheduling




Thread design

- When you have multiple threads, where to put them?




A comparison

- Services in or out of campus




User-level Thread

- When you have multiple threads, where to put them?




User-level Thread

- User-level threads: the kernel knows nothing about them and managed by applications




Discussion

- Advantages?




User-level Thread - Discussions

- Advantages
 - No OS thread-support need
 - Lightweight: thread switching vs. process switching
 - Less memory
 - Faster context switch
 - Easier to communicate and data sharing
 - Management in the user space



Discussion

- Disadvantages?



User-level Thread - Discussions

- Disadvantages (less kernel support, need to implement all logic by application itself)

- Scheduling:


- How blocking system calls implemented? Called by a thread?
 - How to change blocking system calls to non-blocking?

- Memory management:

- How to deal with page faults?


- Interrupt

- How to stop a thread from running forever? No clock interrupts




Kernel-level Thread

- When you have multiple threads, where to put them?




Discussion

- Advantages?




Implementing Threads in the Kernel

- Kernel-level threads: managed by the kernel
 - Threads known to OS, use OS service directly
 - Scheduling, memory management, storage, I/O, etc.




Discussion

- Disadvantages?




Implementing Threads in the Kernel

- Kernel-level threads: managed by the kernel
 - Threads known to OS, use OS service directly
 - Scheduling, memory management, storage, I/O, etc.
 - Slow
 - Trap into the kernel mode
 - Expensive to create and switch
 - Create memory inside the kernel
 - Context switch between the kernel and user space




Threads in Hybrid-Space

- When you have multiple threads, where to put them?




Hybrid Implementations

- Use kernel-level threads and then *multiplex* user-level threads onto some or all of the kernel-level threads
- Multiplexing user-level threads onto kernel-level threads
- Enjoy the benefits of user and kernel level threads
- Too complex !




Thread in user level/kernel level/hybrid



Thread design

- When you have multiple threads, by what order should we execute them? - -> Scheduling




Possible: A1, A2, A3, A1, A2, A3
Also possible: A1, B1, A2, B2, A3, B3



Local Thread Scheduling

- Next thread is picked from among the threads belonging to the **current process**
- Each process gets a timeslice from kernel.
- Then the timeslice is divided up among the threads within the current process
- Scheduling decision requires only **local knowledge** of threads within the current process. It can be implemented by using **user or kernel level threads**



Possible: A1, A2, A3, A1, A2, A3


Not possible: A1, B1, A2, B2, A3, B3

For example, say process timeslice may be N ms, and each thread within the process runs for $N/3$ ms per cycle



Global Thread scheduling

- Next thread to be scheduled is picked up from **ANY process** in the system.
 - Not just the current process
- Timeslice is allocated at the granularity of threads
 - No notion of per-process timeslice
- Global scheduling can be implemented **only with kernel-level threads**
 - Picking the next thread requires global knowledge of threads in all processes.




Possible: A1, A2, A3, A1, A2, A3
Also possible: A1, B1, A2, B2, A3, B3

For example, say process timeslice may be N ms, For example each thread runs for $2N/7$ ms per cycle




Local Thread Scheduling vs Global Thread scheduling



Possible: A1, A2, A3, A1, A2, A3

Not possible: A1, B1, A2, B2, A3, B3



Possible: A1, A2, A3, A1, A2, A3

Also possible: A1, B1, A2, B2, A3, B3

1, next thread

2, time slice for each thread

3, implementation: user/kernel level vs. only kernel level thread




Multiple threads → Pandora's Concurrency Box

- Multiple threads → Concurrency
- The illusion of concurrency is both **powerful** and **useful**:
 - It helps us think about how to structure our applications - -> multiple threads apps
 - It hides latencies caused by hardware devices - -> accelerate execution
- Unfortunately, concurrency also creates **problems**:
 - **Coordination**: how do we enable efficient communication between the multiple threads involved in performing a single task?
 - **Correctness**: how do we ensure that shared information remains consistent when being accessed by multiple threads concurrently?



Concurrency and multi-threads

- Unless precisely synchronized, threads may:
 - Be run in **any order**,
 - Be stopped and restarted at **any time**,
 - Remain stopped for **arbitrary lengths of time**.



- Generally these are **good things** -- the operating system is responsible for how to allocate resources.

Talk about CPU scheduling, memory management, lock, synchronization in future talks

Conclusion

- What is thread?
 - Multiple thread application
 - Thread vs Process
 - Advantage and disadvantage of thread
- Thread in Linux
- Thread design
 - Kernel space vs User space
 - Local thread vs Global thread scheduling

