

# HPC & Parallel Programming

## Message Passing Interface (MPI)

Kun Suo

Computer Science, Kennesaw State University

<https://kevinsuo.github.io/>

# Outline

---

- MPI introduction
  - Helloworld of MPI
- Performance evaluation
- Example: how to solve problems in MPI
  - Trapezoidal problem



# What is MPI?

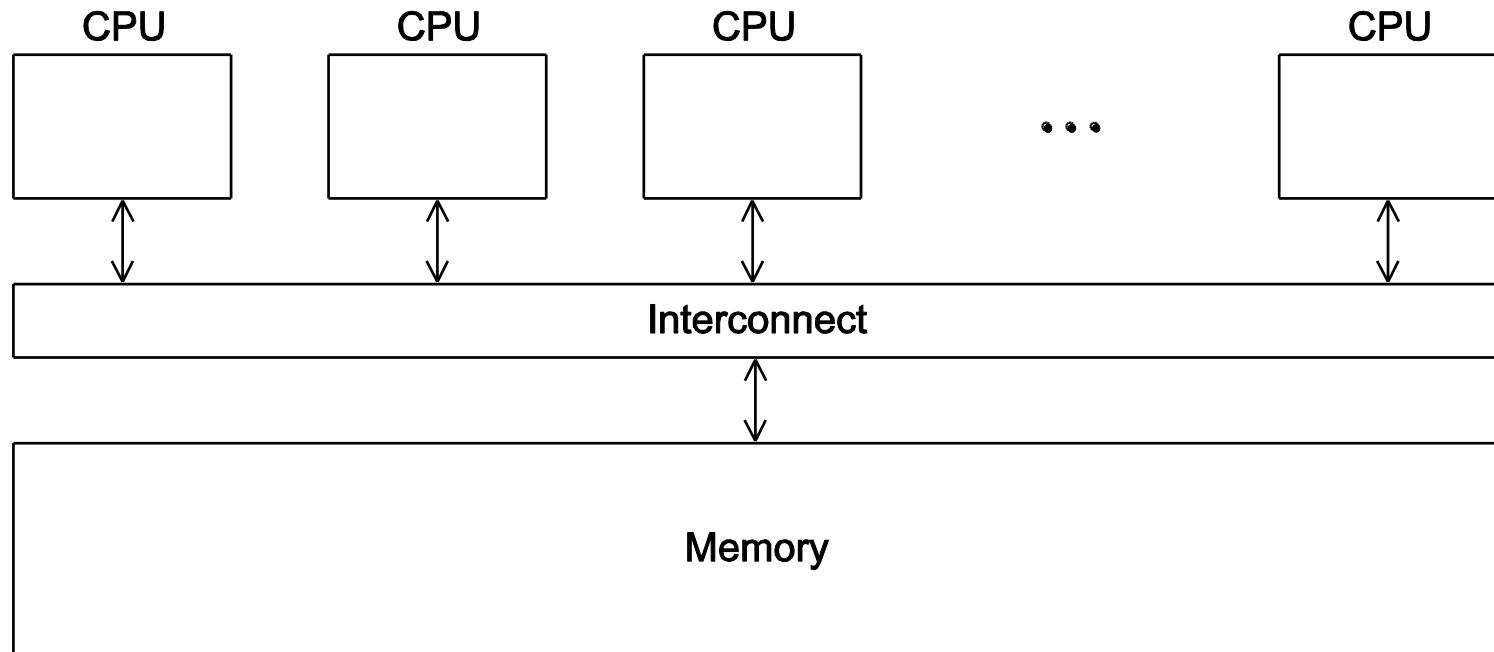
---



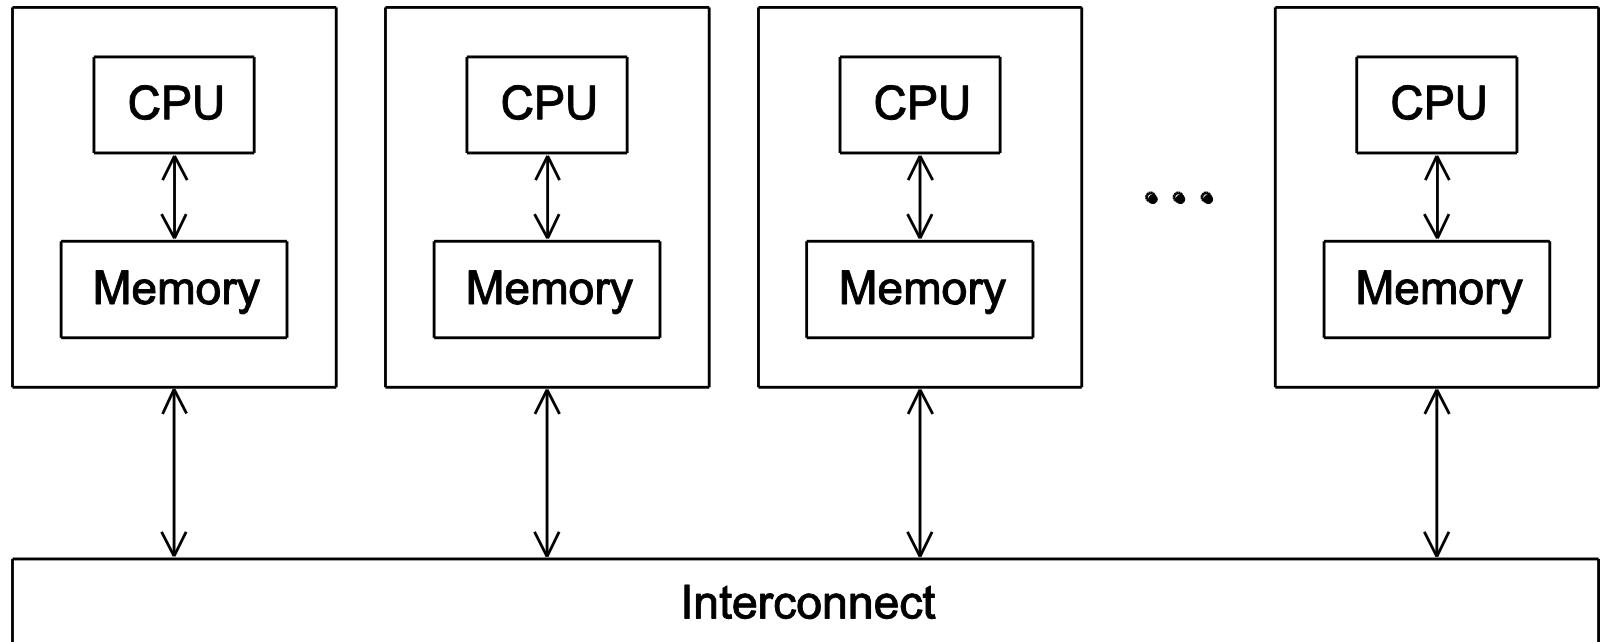
- Message Passing Interface (MPI) is a standardized library for parallel computing



# A shared memory system



# A distributed memory system



# Hello World!

---

```
#include <stdio.h>

int main(void) {
    printf("hello, world\n");

    return 0;
}
```



# helloworld-mpi.c

[https://github.com/kevinsuo/CS4504  
/blob/master/helloworld-mpi.c](https://github.com/kevinsuo/CS4504/blob/master/helloworld-mpi.c)

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    // Initialize the MPI environment
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    // Print off a hello world message
    printf("Hello world from processor %s, rank %d out of %d processors\n",
           processor_name, world_rank, world_size);

    // Finalize the MPI environment.
    MPI_Finalize();
}
```

# MPI Programs

- Written in C.
  - Has main.
  - Uses stdio.h, string.h, etc.
- Need to add **mpi.h** header file.
- Identifiers defined by MPI start with “MPI\_”.
- First letter following underscore is uppercase (e.g., MPI\_Init).
  - For function names and MPI-defined types.
  - Helps to avoid confusion.

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    // Initialize the MPI environment
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    // Print off a hello world message
    printf("Hello world from processor %s, rank %d out of %d processors\n",
           processor_name, world_rank, world_size);

    // Finalize the MPI environment.
    MPI_Finalize();
}
```



# MPI Components

- **MPI\_Init**

- Tells MPI to do all the necessary setup.

```
int MPI_Init(  
    int*      argc_p /* in/out */,  
    char*** argv_p /* in/out */);
```

- **MPI\_Finalize**

- Tells MPI we're done, so clean up anything allocated for this program.

```
int MPI_Finalize(void);
```



# Basic Outline

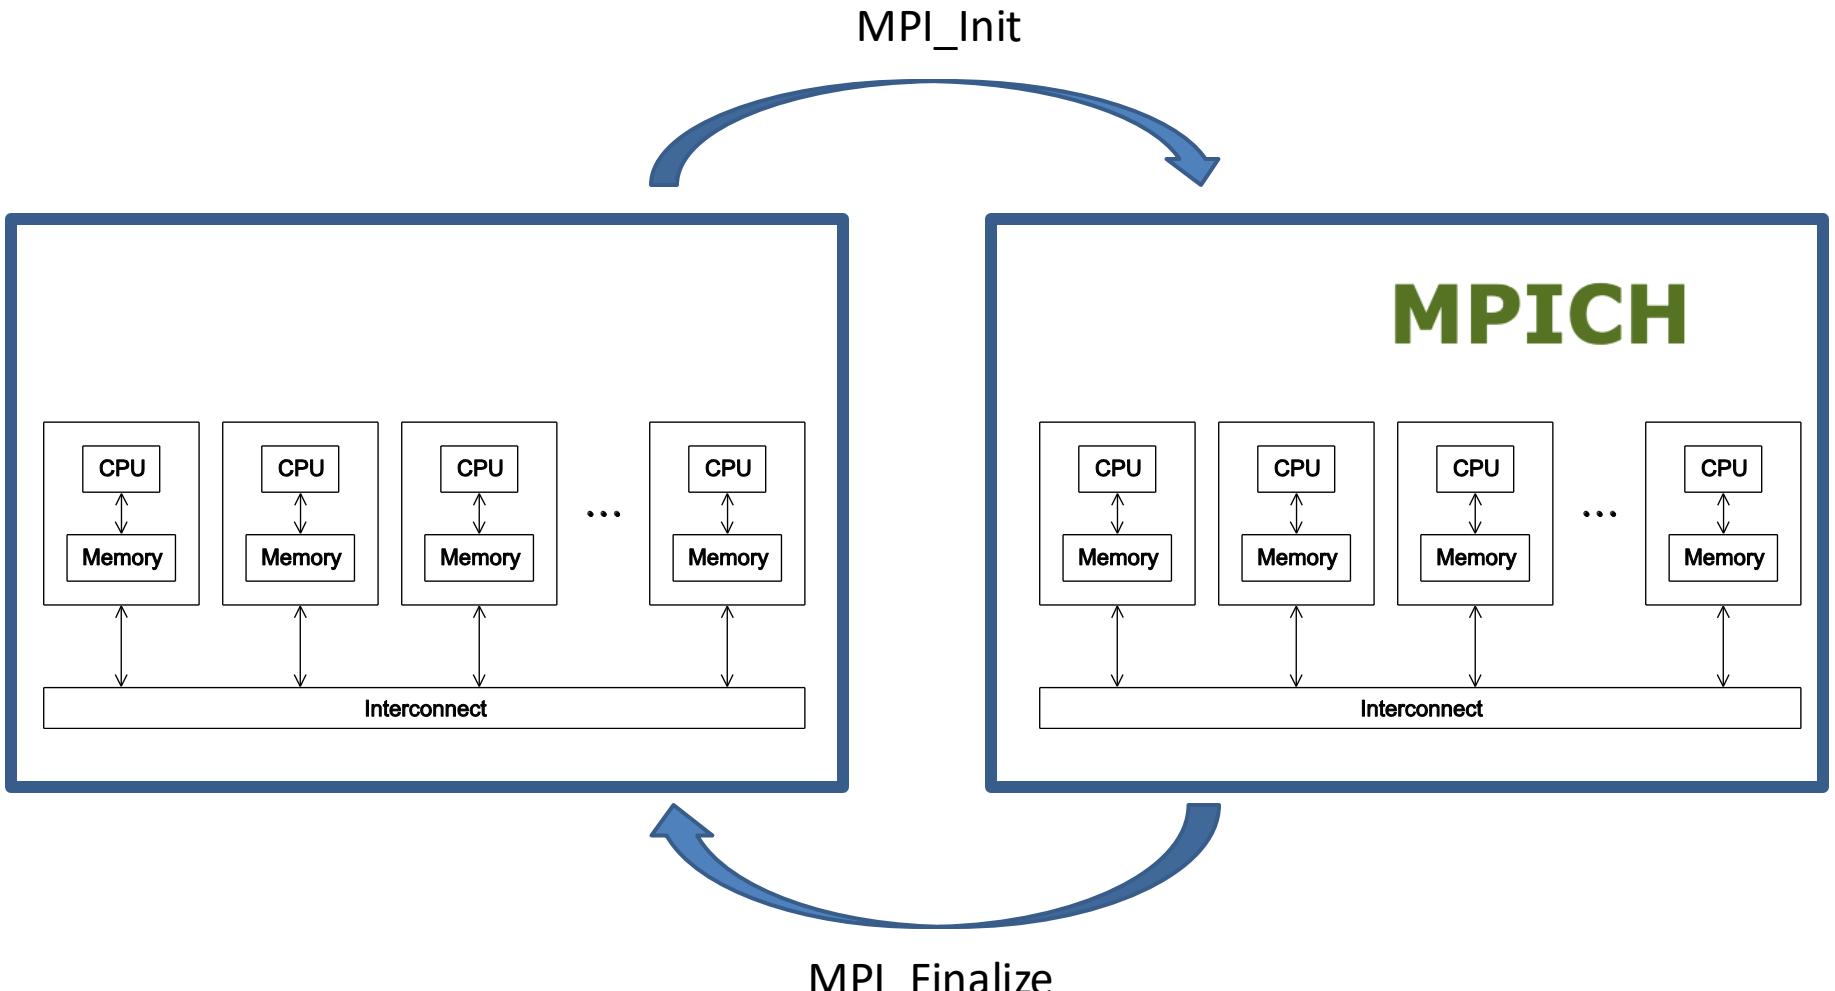
---

```
#include <mpi.h>

int main(int argc, char* argv[]) {
    . . .
    /* No MPI calls before this */
    MPI_Init(&argc, &argv);
    . . .
    MPI_Finalize();
    /* No MPI calls after this */
    . . .
    return 0;
}
```



# MPI Components



# helloworld-mpi.c

[https://github.com/kevinsuo/CS4504  
/blob/main/helloworld-mpi.c](https://github.com/kevinsuo/CS4504/blob/main/helloworld-mpi.c)

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    // Initialize the MPI environment
    MPI_Init(NULL, NULL);   
  

    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);  
  

    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);  
  

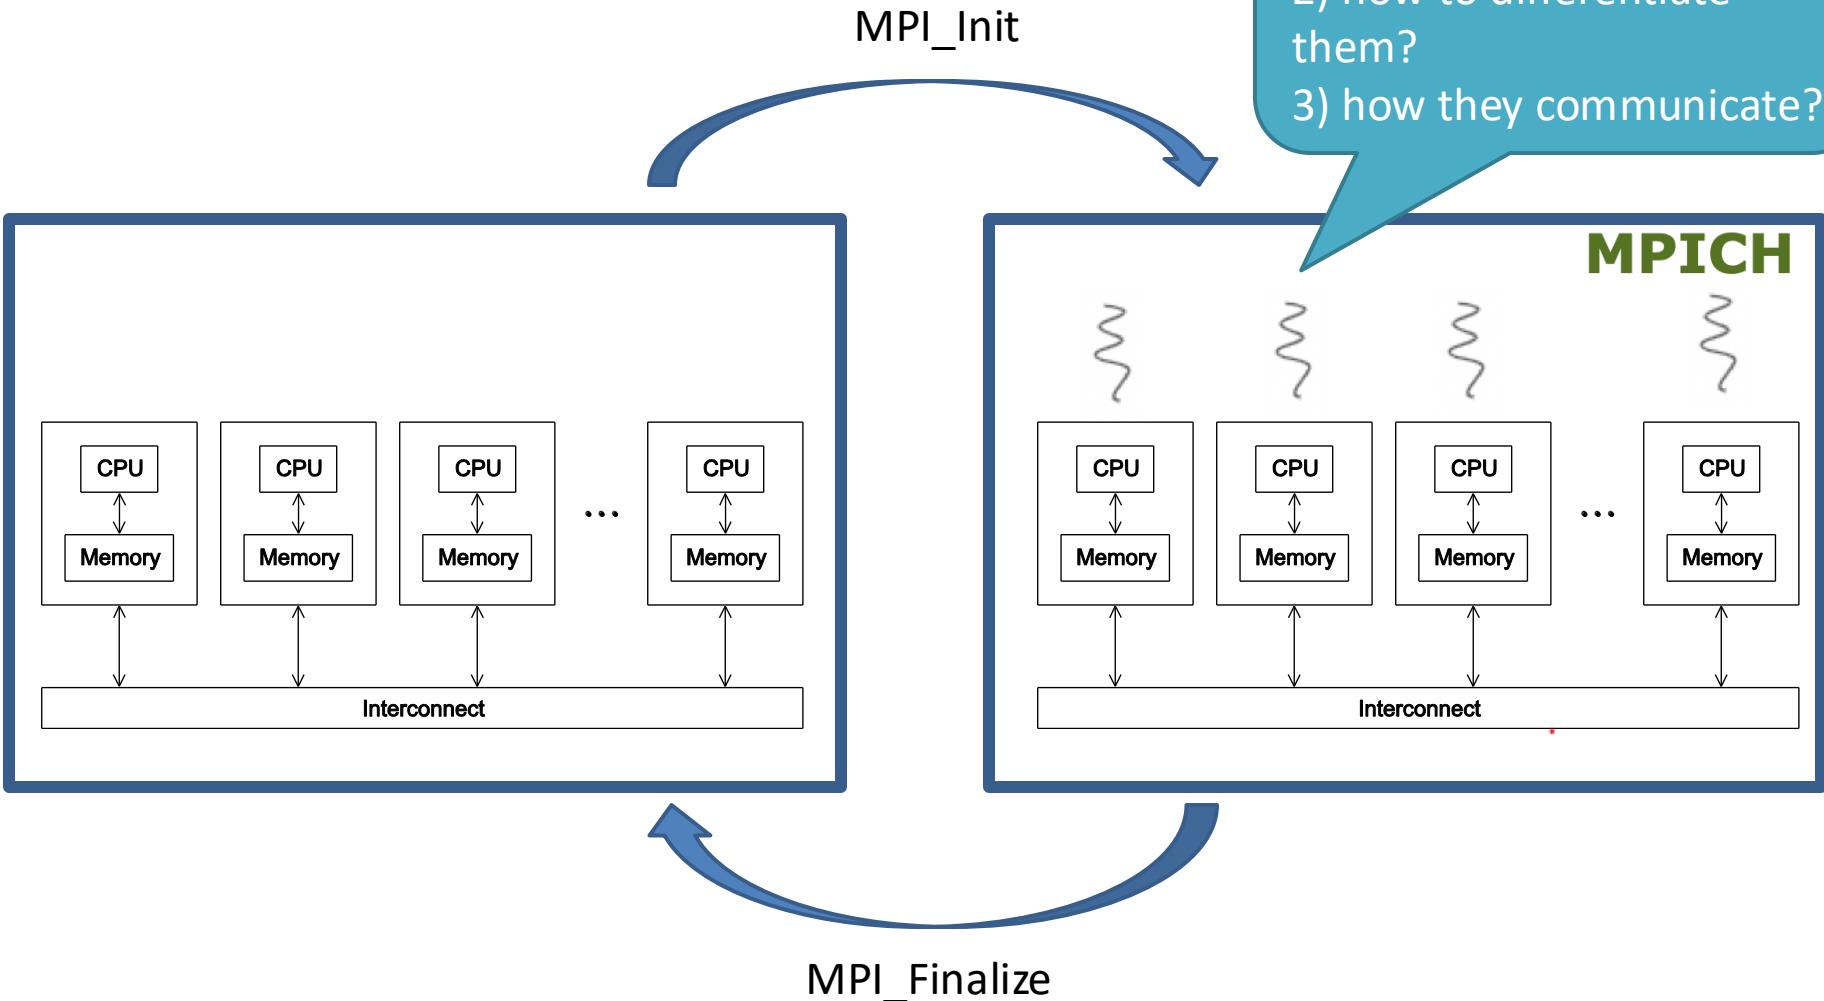
    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);  
  

    // Print off a hello world message
    printf("Hello world from processor %s, rank %d out of %d processors\n",
           processor_name, world_rank, world_size);  
  

    // Finalize the MPI environment. 
    MPI_Finalize();
}
```

# MPI Components

- 1) how many processes do I have?
- 2) how to differentiate them?
- 3) how they communicate?



# Communicators

---

- A collection of processes that can send messages to each other.
- `MPI_Init` defines a communicator that consists of all the processes created when the program is started (defined by user).
- Called `MPI_COMM_WORLD`. → Number of process



# helloworld-mpi.c

[https://github.com/kevinsuo/CS4504  
/blob/master/helloworld-mpi.c](https://github.com/kevinsuo/CS4504/blob/master/helloworld-mpi.c)

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    // Initialize the MPI environment
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
        
```

```
    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    // Print off a hello world message
    printf("Hello world from processor %s, rank %d out of %d processors\n",
           processor_name, world_rank, world_size);

    // Finalize the MPI environment.
    MPI_Finalize();
}
```

# Communicators

```
int MPI_Comm_size(  
    MPI_Comm    comm          /* in */,  
    int*        comm_sz_p     /* out */);
```



*number of processes in the communicator*

```
int MPI_Comm_rank(  
    MPI_Comm    comm          /* in */,  
    int*        my_rank_p     /* out */);
```



*my rank*

*(the process making this call)*



# helloworld-mpi.c

[https://github.com/kevinsuo/CS4504  
/blob/main/helloworld-mpi.c](https://github.com/kevinsuo/CS4504/blob/main/helloworld-mpi.c)

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    // Initialize the MPI environment
    MPI_Init(NULL, NULL);

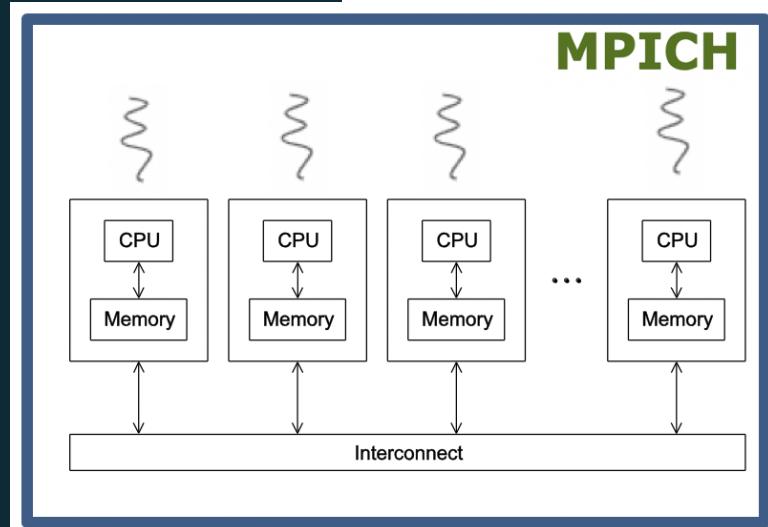
    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size); ✓

    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank); ✓

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

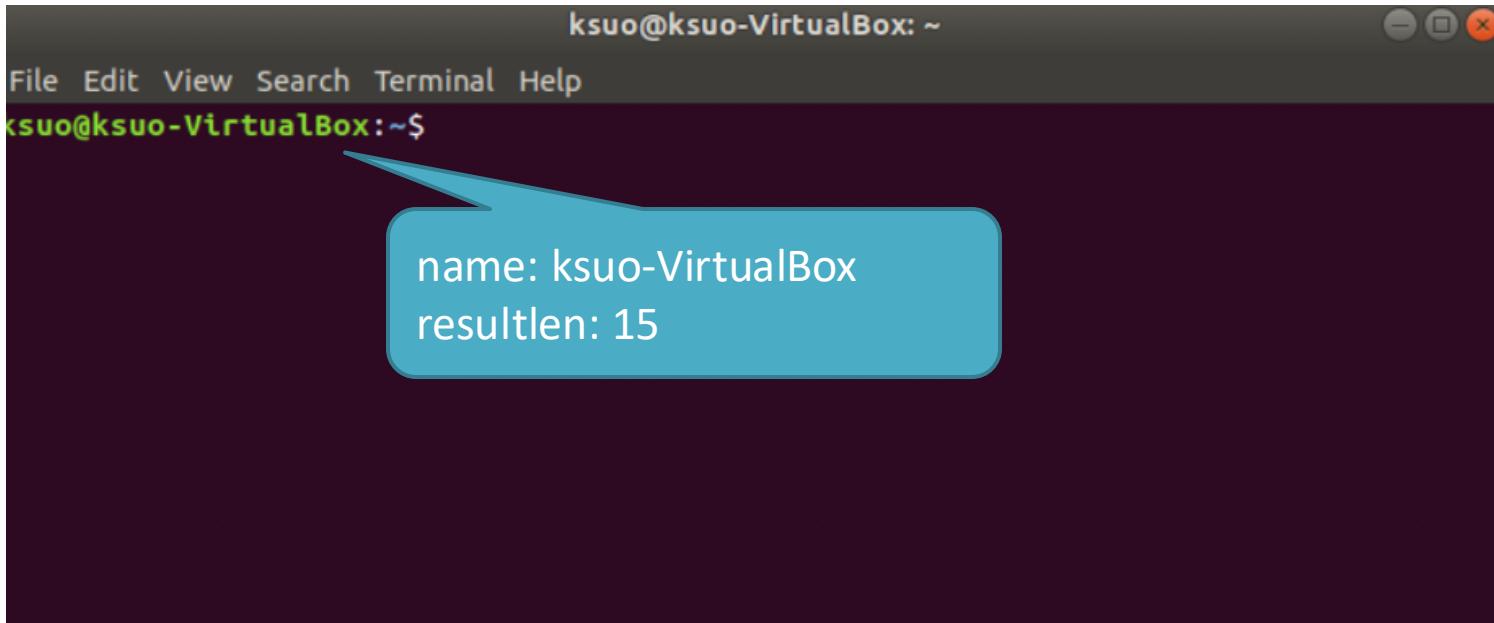
    // Print off a hello world message
    printf("Hello world from processor %s, rank %d out of %d processors\n",
           processor_name, world_rank, world_size);

    // Finalize the MPI environment.
    MPI_Finalize(); ✓
}
```



# **MPI\_Get\_processor\_name(char \*name,int \*resultlen)**

- name: your machine name string
- resultlen: your machine name string length



A screenshot of a terminal window titled "ksuo@ksuo-VirtualBox: ~". The window has a dark background and a light gray header bar. The terminal menu bar includes "File", "Edit", "View", "Search", "Terminal", and "Help". The prompt "ksuo@ksuo-VirtualBox:~\$" is visible. A blue callout bubble points from the bottom left towards the terminal window. Inside the bubble, the text "name: ksuo-VirtualBox" and "resultlen: 15" is displayed in white.

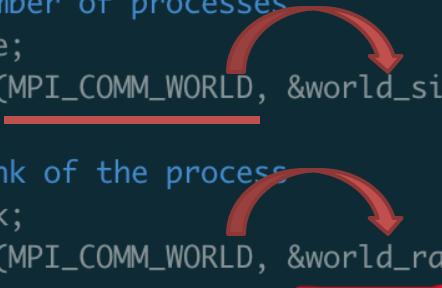
```
name: ksuo-VirtualBox
resultlen: 15
```

# helloworld-mpi.c

[https://github.com/kevinsuo/CS4504  
/blob/main/helloworld-mpi.c](https://github.com/kevinsuo/CS4504/blob/main/helloworld-mpi.c)

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    // Initialize the MPI environment
    MPI_Init(NULL, NULL); ✓

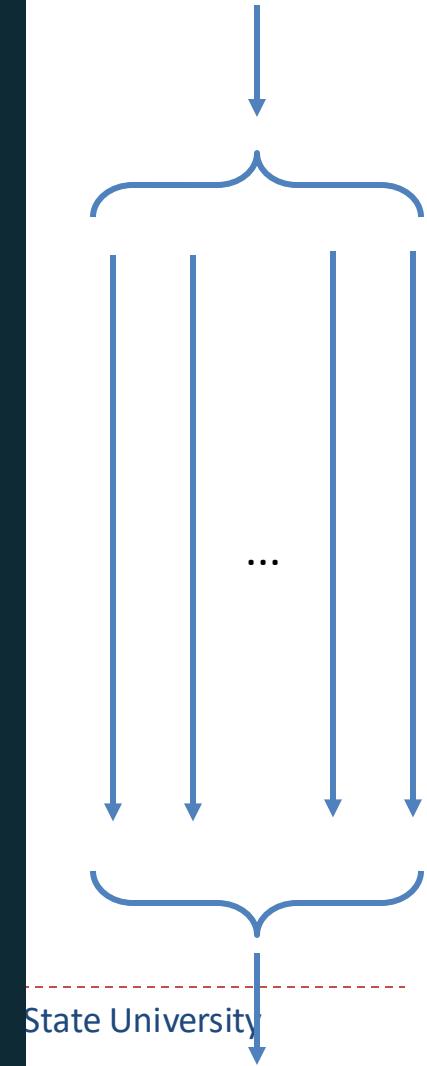
    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size); ✓
     ↗

    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank); ✓
     ↗

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len); ✓
     ↗

    // Print off a hello world message
    printf("Hello world from processor %s, rank %d out of %d processors\n",
           processor_name, world_rank, world_size);

    // Finalize the MPI environment.
    MPI_Finalize(); ✓
}
```



# Compile

*wrapper script to compile*

*source file*

**mpicc -g -Wall -o mpi\_hello mpi\_hello.c**

*produce debugging information*

*create this executable file name  
(as opposed to default a.out)*

*turns on all warnings*

```
fish /home/administrator
administrator@ubuntu1804vm ~>
mpicc -g -Wall helloworld-mpi.c -o helloworld-mpi.o

Command 'mpicc' not found, but can be installed with:

sudo apt install lam4-dev
sudo apt install libmpich-dev
sudo apt install libopenmpi-dev
```



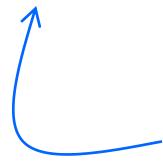
# Execution

---

```
mpiexec -n <number of processes> <executable>
```

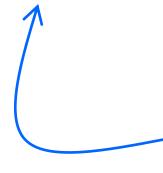
---

```
mpiexec -n 1 ./mpi_hello
```



*run with 1 process*

```
mpiexec -n 4 ./mpi_hello
```



*run with 4 processes*



# Execution

```
mpiexec -n 1 ./mpi_hello
```

Greetings from process 0 of 1 !

```
mpiexec -n 4 ./mpi_hello
```

Greetings from process 0 of 4 !

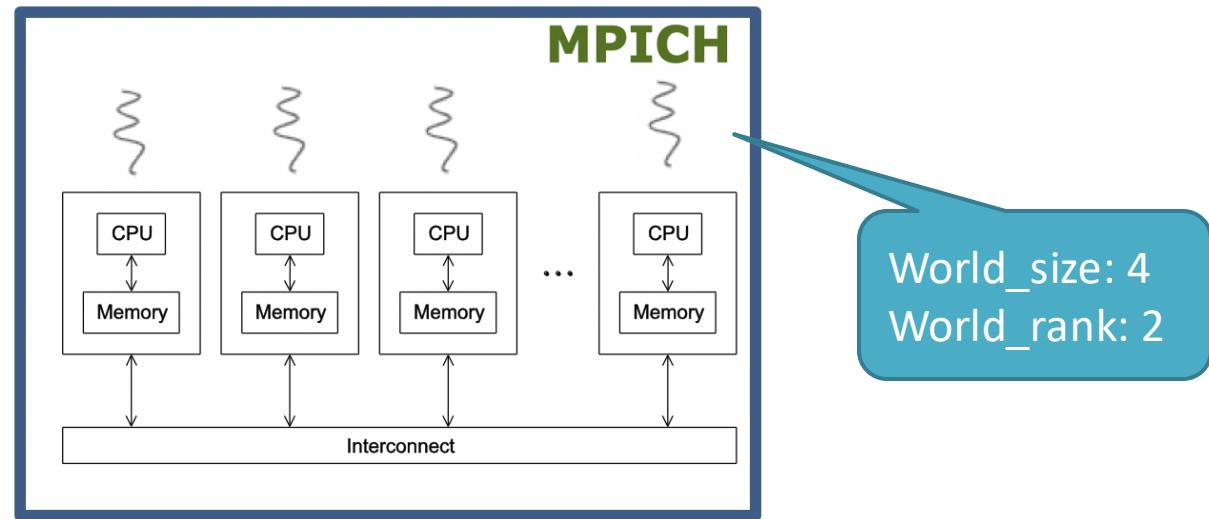
Greetings from process 1 of 4 !

Greetings from process 2 of 4 !

Greetings from process 3 of 4 !



```
ksuo@ksuo-VirtualBox ~/cs7172> mpiexec -n 4 ./helloworld-mpi.o
Hello world from processor ksuo-VirtualBox, rank 0 out of 4 processors
Hello world from processor ksuo-VirtualBox, rank 1 out of 4 processors
Hello world from processor ksuo-VirtualBox, rank 2 out of 4 processors
Hello world from processor ksuo-VirtualBox, rank 3 out of 4 processors
ksuo@ksuo-VirtualBox ~/cs7172> mpiexec -n 2 ./helloworld-mpi.o
Hello world from processor ksuo-VirtualBox, rank 1 out of 2 processors
Hello world from processor ksuo-VirtualBox, rank 0 out of 2 processors
ksuo@ksuo-VirtualBox ~/cs7172> mpiexec -n 1 ./helloworld-mpi.o
Hello world from processor ksuo-VirtualBox, rank 0 out of 1 processors
```



# helloworld-mpi.c

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    // Initialize the MPI environment
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    // Print off a hello world message
    printf("Hello world from processor %s, rank %d out of %d processors\n",
           processor_name, world_rank, world_size);

    // Finalize the MPI environment.
    MPI_Finalize();
}
```

[https://github.com/kevinsuo/CS4504  
/blob/main/helloworld-mpi.c](https://github.com/kevinsuo/CS4504/blob/main/helloworld-mpi.c)

```
ksuo@ksuo-VirtualBox ~/cs7172> mpiexec -n 4 ./helloworld-mpi.o
Hello world from processor ksuo-VirtualBox, rank 0 out of 4 processors
Hello world from processor ksuo-VirtualBox, rank 1 out of 4 processors
Hello world from processor ksuo-VirtualBox, rank 2 out of 4 processors
Hello world from processor ksuo-VirtualBox, rank 3 out of 4 processors
ksuo@ksuo-VirtualBox ~/cs7172> mpiexec -n 2 ./helloworld-mpi.o
Hello world from processor ksuo-VirtualBox, rank 1 out of 2 processors
Hello world from processor ksuo-VirtualBox, rank 0 out of 2 processors
ksuo@ksuo-VirtualBox ~/cs7172> mpiexec -n 1 ./helloworld-mpi.o
Hello world from processor ksuo-VirtualBox, rank 0 out of 1 processors
```

# Performance evaluation

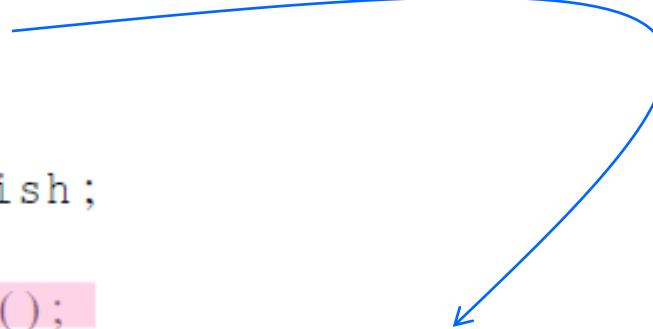
---



# Elapsed parallel time

- Returns the number of seconds that have elapsed since some time in the past.

```
double MPI_Wtime(void);  
  
double start, finish;  
.  
.  
.  
start = MPI_Wtime();  
/* Code to be timed */  
.  
.  
finish = MPI_Wtime();  
printf("Proc %d > Elapsed time = %e seconds\n"  
      my_rank, finish-start);
```



# Example

[https://raw.githubusercontent.com/kevinsuo/CS4504/master/time\\_sample.c](https://raw.githubusercontent.com/kevinsuo/CS4504/master/time_sample.c)

```
#include "mpi.h"
#include <stdio.h>

int main( int argc, char *argv[] )
{
    double t1, t2 = 0;

    MPI_Init( 0, 0 );
    sleep(10);
    printf("MPI_Wtime measured a 1 second sleep to be: %1.2f\n", t2-t1);fflush(stdout);
    MPI_Finalize( );
    return 0;
}
```



# Example

[https://raw.githubusercontent.com/kevinsuo/CS4504/master/time\\_sample.c](https://raw.githubusercontent.com/kevinsuo/CS4504/master/time_sample.c)

```
#include "mpi.h"
#include <stdio.h>

int main( int argc, char *argv[] )
{
    double t1, t2;

    MPI_Init( 0, 0 );
    t1 = MPI_Wtime(); ←
    sleep(10);
    t2 = MPI_Wtime(); ←
    printf("MPI_Wtime measured a 1 second sleep to be: %1.2f\n", t2-t1);fflush(stdout);
    MPI_Finalize( );
    return 0;
}
```



# Elapsed serial time

- In this case, you don't need to link in the MPI libraries.
- Returns time in microseconds elapsed from some point in the past.

```
#include "timer.h"  
.  
.  
double now;  
.  
.  
GET_TIME(now);
```



# Elapsed serial time

---

```
#include "timer.h"
.
.
double start, finish;
.
.
GET_TIME( start );
/* Code to be timed */
.
.
GET_TIME( finish );
printf("Elapsed time = %e seconds\n", finish-start);
```



# Example

[https://raw.githubusercontent.com/kevinsuo/CS4504/master/time\\_sample2.c](https://raw.githubusercontent.com/kevinsuo/CS4504/master/time_sample2.c)

```
#include "mpi.h"
#include <stdio.h>
#include <sys/time.h>

#define GET_TIME(now) { \
    struct timeval t; \
    gettimeofday(&t, NULL); \
    now = t.tv_sec + t.tv_usec/1000000.0; \
}

int main( int argc, char *argv[] )
{
    double t1, t2 = 0;

    MPI_Init( 0, 0 );
    sleep(10);
    printf("Elapsed time = %e\n", t2-t1);fflush(stdout);
    MPI_Finalize( );
    return 0;
}
```



# Example

[https://raw.githubusercontent.com/kevinsuo/CS4504/master/time\\_sample2.c](https://raw.githubusercontent.com/kevinsuo/CS4504/master/time_sample2.c)

```
#include "mpi.h"
#include <stdio.h>
#include <sys/time.h>

#define GET_TIME(now) { \
    struct timeval t; \
    gettimeofday(&t, NULL); \
    now = t.tv_sec + t.tv_usec/1000000.0; \
}

int main( int argc, char *argv[] )
{
    double t1, t2;

    MPI_Init( 0, 0 );
    GET_TIME(t1); ←
    sleep(10);
    GET_TIME(t2); ←
    printf("Elapsed time = %e\n", t2-t1);fflush(stdout);
    MPI_Finalize( );
    return 0;
}
```



# Elapsed serial time in nanoseconds

```
#include <time.h>

{
    struct timespec start, end;

    clock_gettime(CLOCK_MONOTONIC, &start);

    //... do something

    clock_gettime(CLOCK_MONOTONIC, &end);

    u_int64_t diff = 1000000000L * (end.tv_sec - start.tv_sec) + end.tv_nsec - start.tv_nsec;

    printf("elapsed time = %llu nanoseconds\n", (long long unsigned int) diff);
}
```



# Example

[https://github.com/kevinsuo/CS4504/blob/master/  
time sample3.c](https://github.com/kevinsuo/CS4504/blob/master/time%20sample3.c)

```
#include "mpi.h"
#include <stdio.h>
#include <time.h>

int main( int argc, char *argv[] )
{
    struct timespec t1, t2;

    MPI_Init( 0, 0 );
    sleep(10);

    double diff = 0;

    printf("Time = %llu nanoseconds\n", (long long unsigned int)diff) ;
    fflush(stdout);
    MPI_Finalize( );
    return 0;
}
```



# Example

[https://github.com/kevinsuo/CS4504/blob/master/time\\_sample3.c](https://github.com/kevinsuo/CS4504/blob/master/time_sample3.c)

```
#include "mpi.h"
#include <stdio.h>
#include <time.h>

int main( int argc, char *argv[] )
{
    struct timespec t1, t2;

    MPI_Init( 0, 0 );
    clock_gettime(CLOCK_MONOTONIC, &t1); ←

    sleep(10);

    clock_gettime(CLOCK_MONOTONIC, &t2); ←

    double diff = 1000000000L * (t2.tv_sec - t1.tv_sec) + t2.tv_nsec - t1.tv_nsec;

    printf("Time = %llu nanoseconds\n", (long long unsigned int)diff) ; fflush(stdout);
    MPI_Finalize( );
    return 0;
}
```



# Run-times of serial and parallel matrix-vector multiplication

comm_sz	Order of Matrix				
	1024	2048	4096	8192	16,384
1	4.1	16.0	64.0	270	1100
2	2.3	8.5	33.0	140	560
4	2.0	5.1	18.0	70	280
8	1.7	3.3	9.8	36	140
16	1.7	2.6	5.9	19	71

(Seconds)



# Speedup

---

$$S(n, p) = \frac{T_{\text{serial}}(n)}{T_{\text{parallel}}(n, p)}$$



# Speedups of Parallel Matrix-Vector Multiplication

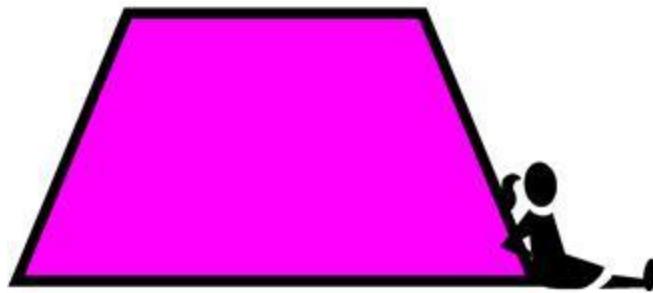
---

comm_sz	Order of Matrix				
	1024	2048	4096	8192	16,384
1	1.0	1.0	1.0	1.0	1.0
2	1.8	1.9	1.9	1.9	2.0
4	2.1	3.1	3.6	3.9	3.9
8	2.4	4.8	6.5	7.5	7.9
16	2.4	6.2	10.8	14.2	15.5

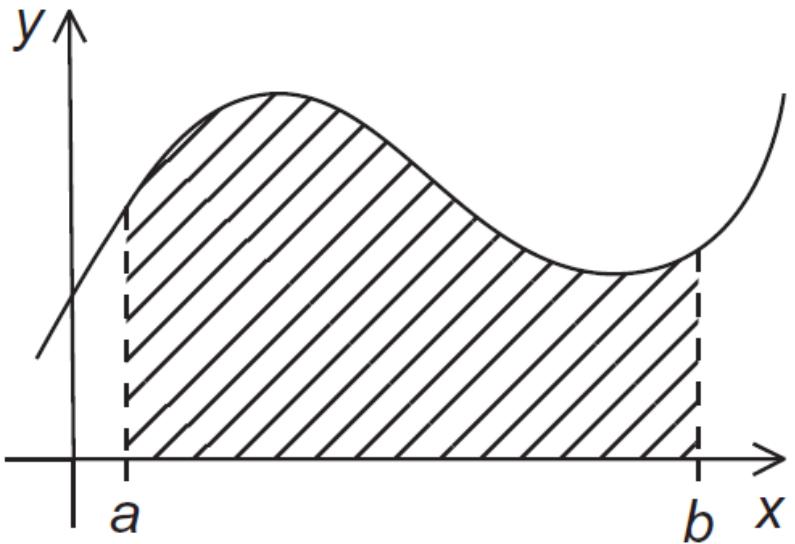


# Trapezoidal rule in mpi

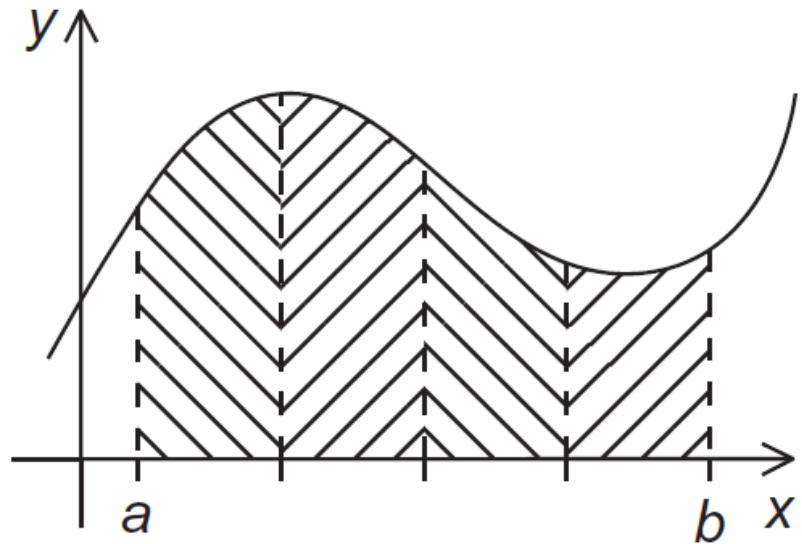
---



# The Trapezoidal Rule



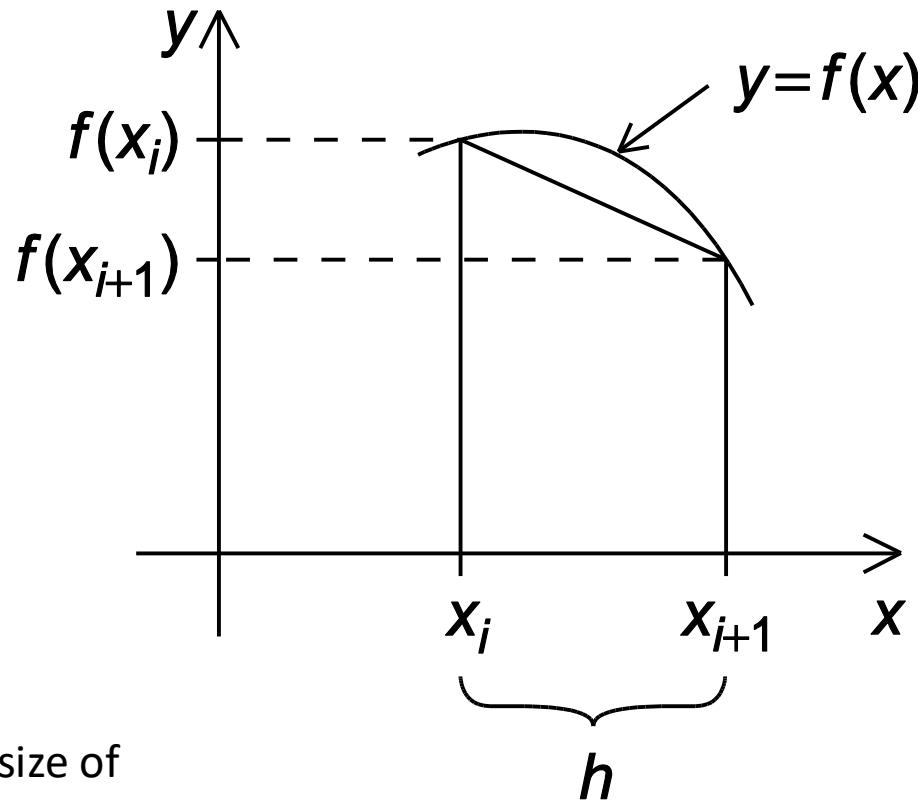
(a)



(b)



# One trapezoid



How to get the size of  
one trapezoid?

$$\text{Area of one trapezoid} = \frac{h}{2}[f(x_i) + f(x_{i+1})]$$



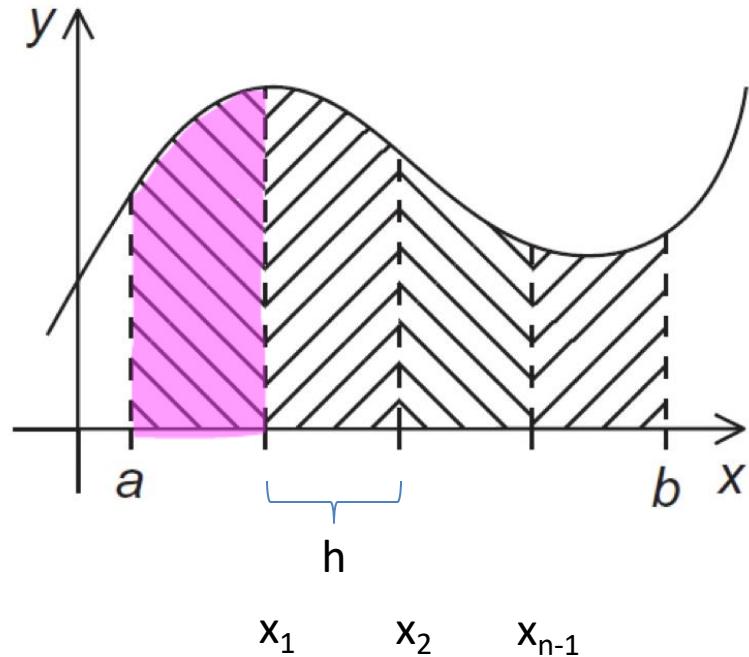
# The Trapezoidal Rule

$$\text{Area of one trapezoid} = \frac{h}{2}[f(x_i) + f(x_{i+1})]$$

$$h = \frac{b-a}{n}$$

$$x_0 = a, x_1 = a+h, x_2 = a+2h, \dots, x_{n-1} = a+(n-1)h, x_n = b$$

$$\frac{h}{2} [f(a) + f(x_1)] = h \left[ \frac{f(a)}{2} + \frac{f(x_1)}{2} \right]$$



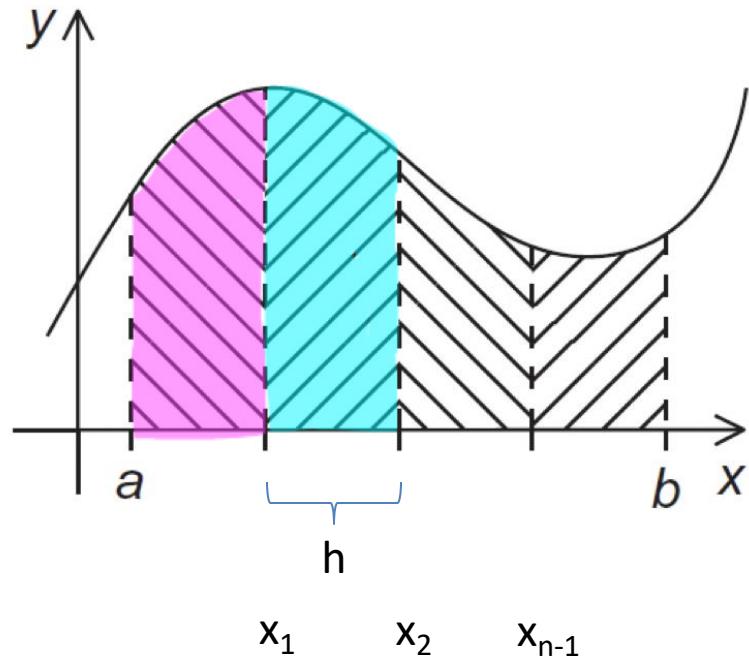
# The Trapezoidal Rule

$$\text{Area of one trapezoid} = \frac{h}{2}[f(x_i) + f(x_{i+1})]$$

$$h = \frac{b-a}{n}$$

$$x_0 = a, x_1 = a+h, x_2 = a+2h, \dots, x_{n-1} = a+(n-1)h, x_n = b$$

$$S = h \left[ \frac{f(a)}{2} + \frac{f(x_1)}{2} + \frac{f(x_2)}{2} + \frac{f(x_3)}{2} + \dots + \frac{f(x_{n-1})}{2} + \frac{f(b)}{2} \right]$$



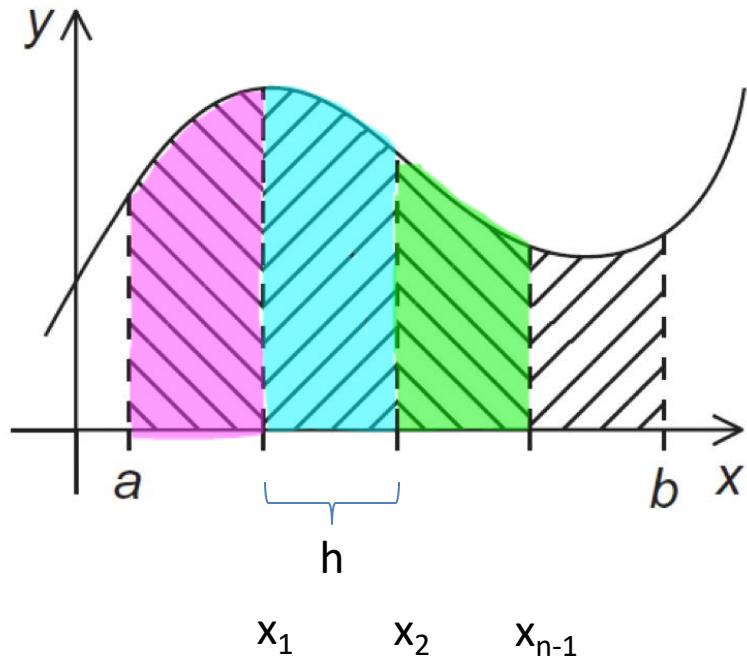
# The Trapezoidal Rule

$$\text{Area of one trapezoid} = \frac{h}{2}[f(x_i) + f(x_{i+1})]$$

$$h = \frac{b-a}{n}$$

$$x_0 = a, x_1 = a+h, x_2 = a+2h, \dots, x_{n-1} = a+(n-1)h, x_n = b$$

$$S = h \left[ \frac{f(a)}{2} + \frac{f(x_1)}{2} + \frac{f(x_1)}{2} + \frac{f(x_2)}{2} + \frac{f(x_2)}{2} + \frac{f(x_3)}{2} \right]$$

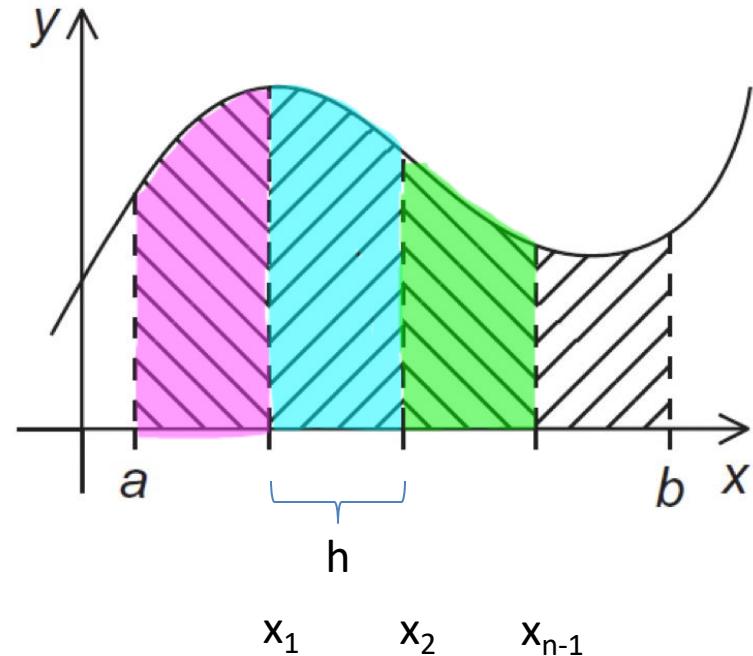


# The Trapezoidal Rule

$$\text{Area of one trapezoid} = \frac{h}{2}[f(x_i) + f(x_{i+1})]$$

$$h = \frac{b-a}{n}$$

$$x_0 = a, x_1 = a+h, x_2 = a+2h, \dots, x_{n-1} = a+(n-1)h, x_n = b$$



$$\text{Sum of trapezoid areas} = h[f(x_0)/2 + f(x_1) + f(x_2) + \dots + f(x_{n-1}) + f(x_n)/2]$$

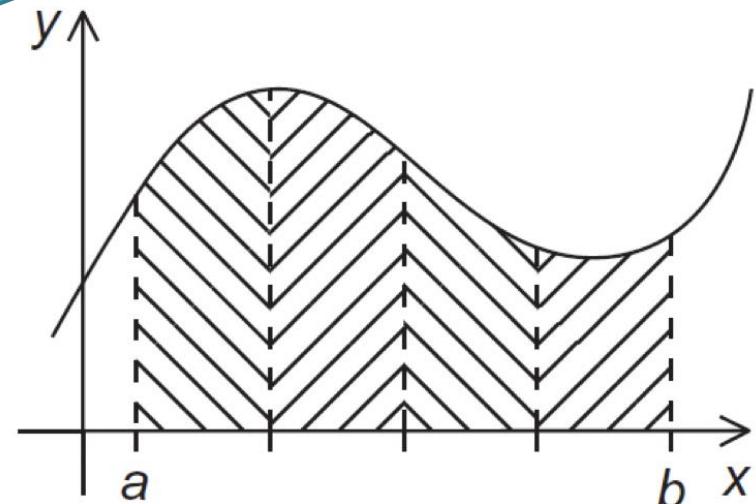


# Pseudo-code for a serial program

Sum of trapezoid areas =  $h[f(x_0)/2 + f(x_1) + f(x_2) + \cdots + f(x_{n-1}) + f(x_n)/2]$

$$\begin{aligned} & (f(a) + f(b))/2 \\ & = (f(x_0) + f(x_n))/2 \end{aligned}$$

```
/* Input:  a, b, n */
h = (b-a)/n;
approx = (f(a) + f(b))/2.0;
for (i = 1; i <= n-1; i++) {
    x_i = a + i*h;
    approx += f(x_i);
}
approx = h*approx;
```



# Parallelizing the Trapezoidal Rule

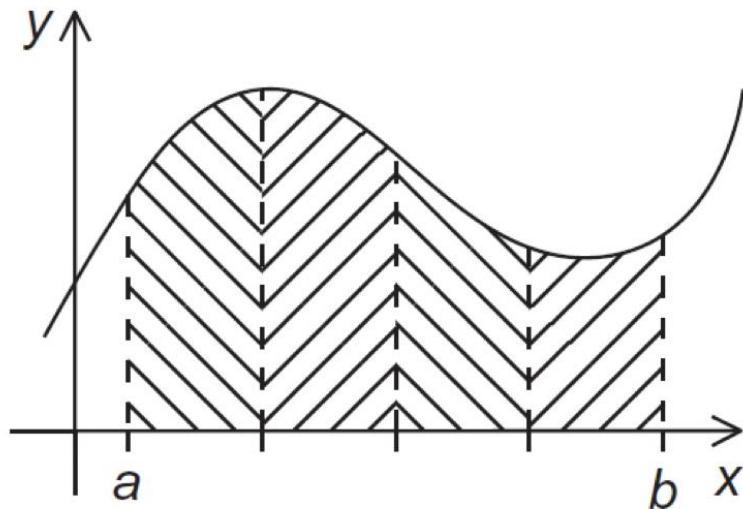
---

1. Partition problem solution into tasks.
2. Identify communication channels between threads.
3. Aggregate tasks into composite tasks.
4. Map composite tasks to threads/cores.



# Parallelizing the Trapezoidal Rule

1. Partition problem solution into tasks.



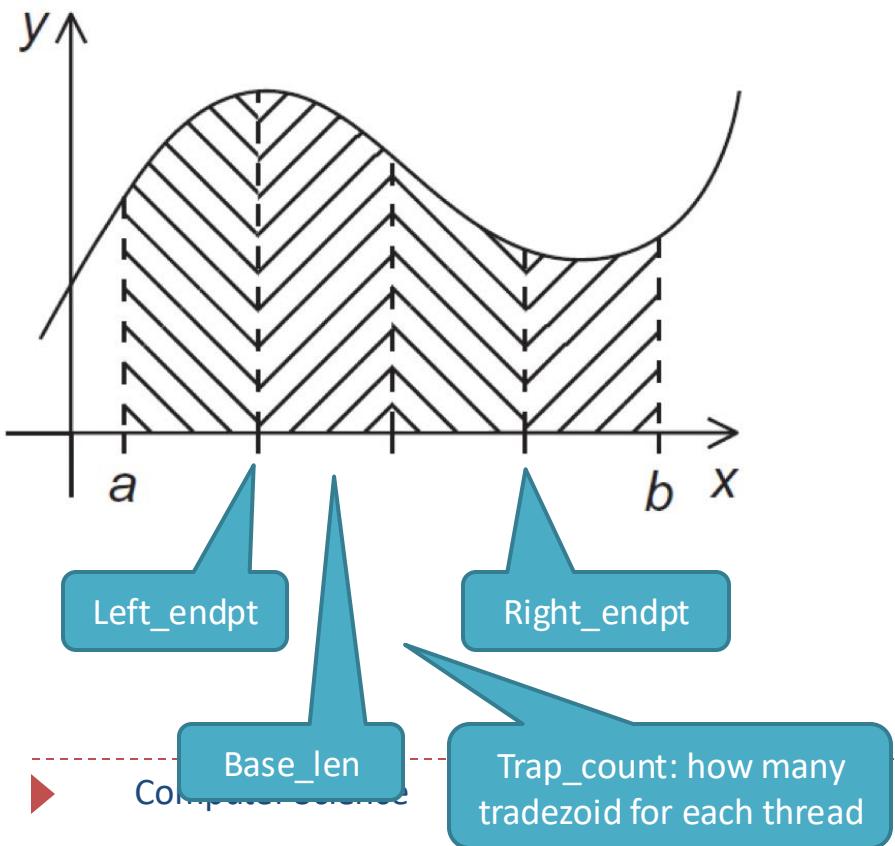
What is the problem?

What is the task?



# Parallelizing the Trapezoidal Rule

## 1. Partition problem solution into tasks.



```
double Trap(  
    double left_endpt /* in */,  
    double right_endpt /* in */,  
    int trap_count /* in */,  
    double base_len /* in */) {  
  
    // function logic  
  
    return estimate;  
} /* Trap */
```

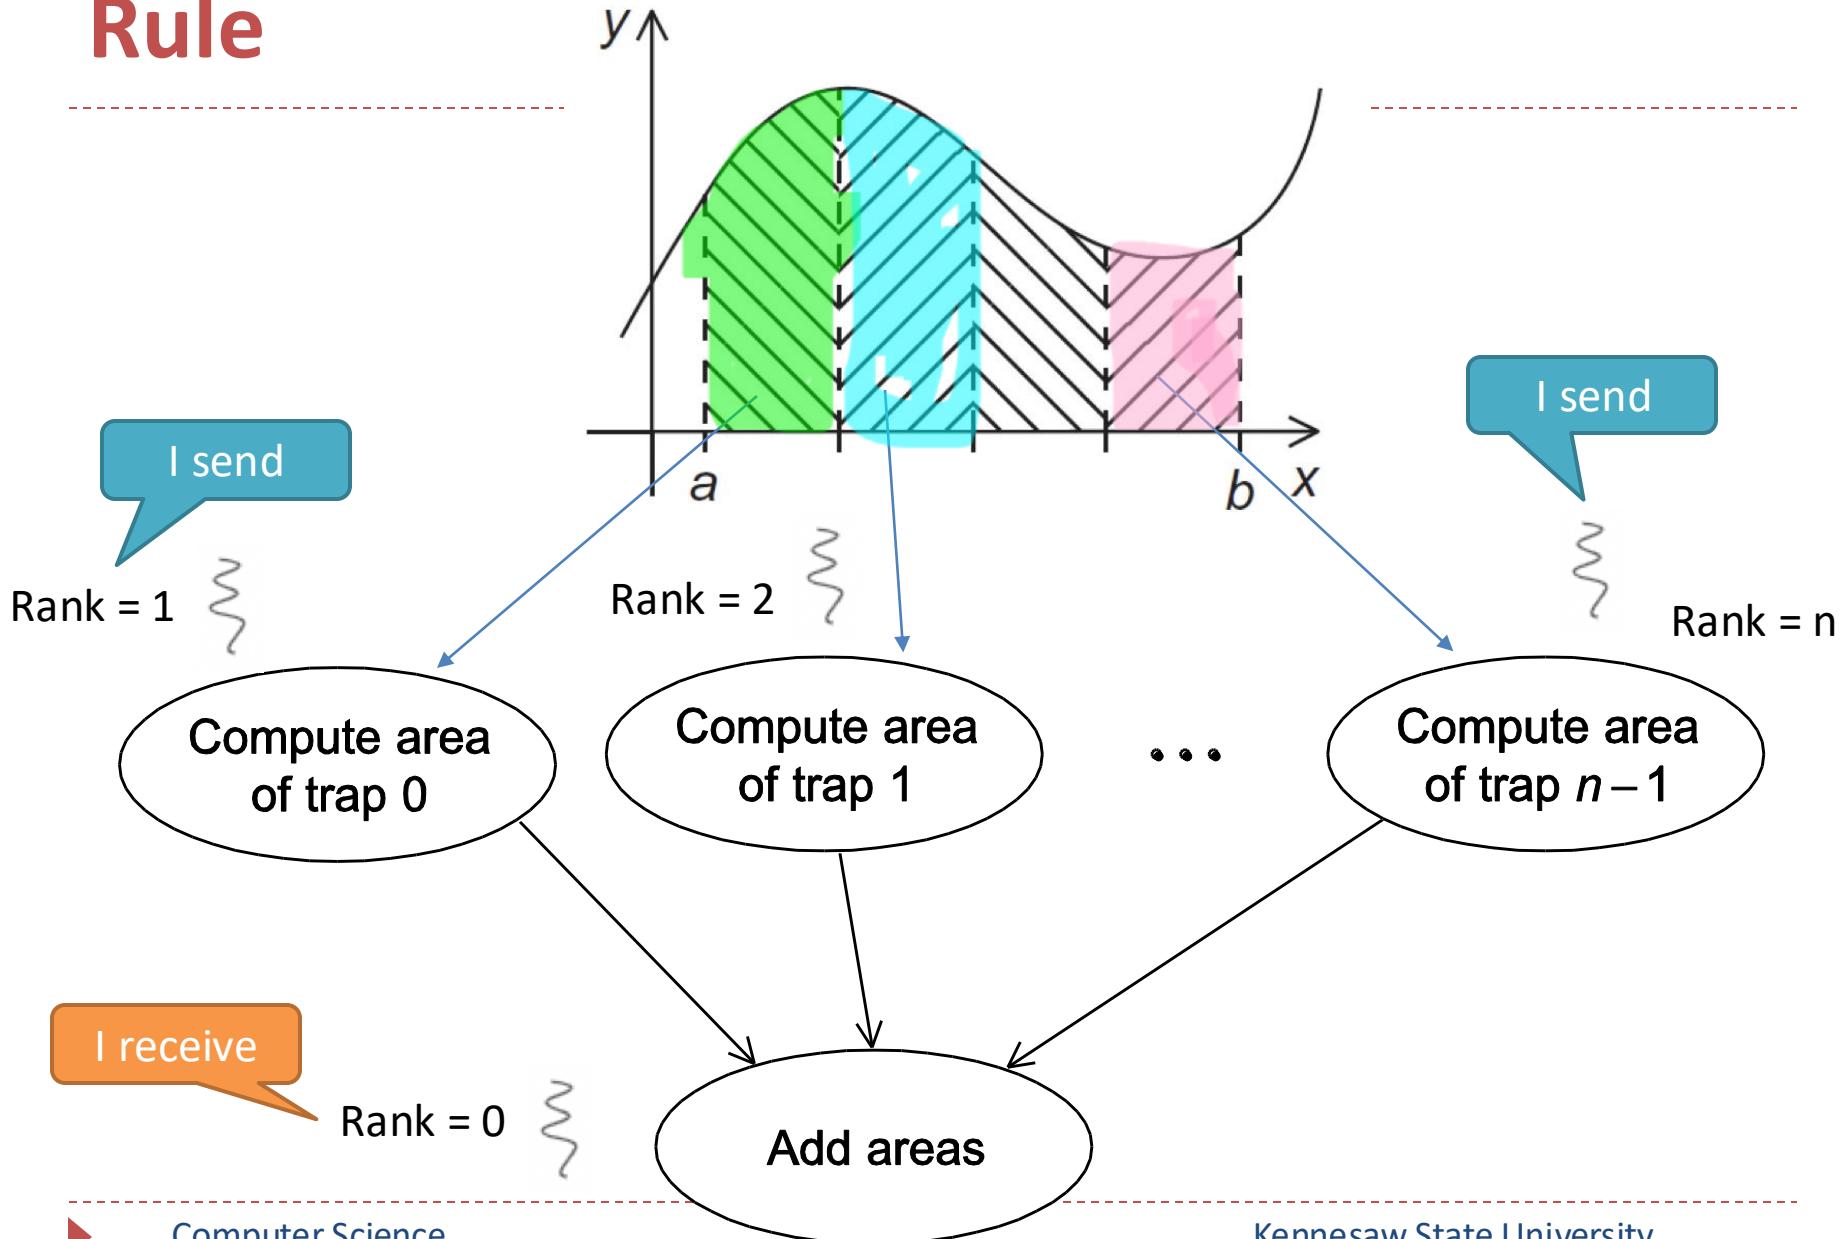
# Parallelizing the Trapezoidal Rule

---

1. ~~Partition problem solution into tasks.~~
2. Identify communication channels between threads.



# Tasks and communications for Trapezoidal Rule



# Communication

```
int MPI_Send(
```

Send what?

How many?

```
void*
```

```
msg_buf_p
```

```
/* in */ ,
```

```
int
```

```
msg_size
```

```
/* in */ ,
```

```
MPI_Datatype
```

```
msg_type
```

```
/* in */ ,
```

```
int
```

```
dest
```

```
/* in */ ,
```

```
int
```

```
tag
```

```
/* in */ ,
```

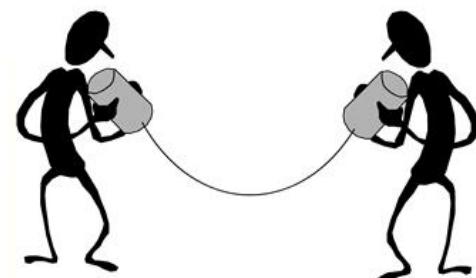
```
MPI_Comm
```

```
communicator
```

```
/* in */ );
```

To where?

Type?



# Data types

MPI datatype	C datatype
MPI_CHAR	<b>signed char</b>
MPI_SHORT	<b>signed short int</b>
MPI_INT	<b>signed int</b>
MPI_LONG	<b>signed long int</b>
MPI_LONG_LONG	<b>signed long long int</b>
MPI_UNSIGNED_CHAR	<b>unsigned char</b>
MPI_UNSIGNED_SHORT	<b>unsigned short int</b>
MPI_UNSIGNED	<b>unsigned int</b>
MPI_UNSIGNED_LONG	<b>unsigned long int</b>
MPI_FLOAT	<b>float</b>
MPI_DOUBLE	<b>double</b>
MPI_LONG_DOUBLE	<b>long double</b>
MPI_BYTE	
MPI_PACKED	



# Communication

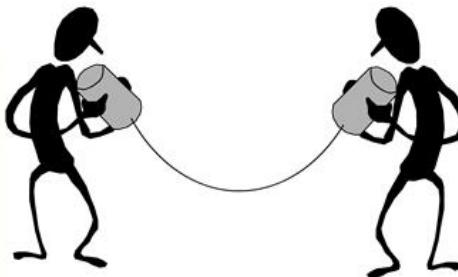
```
int MPI_Recv(  
    void*           msg_buf_p /* out */,  
    int             buf_size  /* in  */,  
    MPI_Datatype   buf_type  /* in  */,  
    int             source    /* in  */,  
    int             tag       /* in  */,  
    MPI_Comm        communicator /* in  */,  
    MPI_Status*     status_p  /* out */);
```

From where?

Recv what?

How many?

Type?



# Message Sent From q to r

```
MPI_Send(send_buf_p, send_buf_sz, send_type, dest, send_tag,  
         send_comm);
```

dest  
*r*

*MPI\_Send*

*src = q*



```
MPI_Recv(recv_buf_p, recv_buf_sz, recv_type, src, recv_tag,  
        recv_comm, &status);
```

src  
*q*



# Parallelizing the Trapezoidal Rule

---

- ~~1. Partition problem solution into tasks.~~
- ~~2. Identify communication channels between threads.~~
3. Aggregate tasks into composite tasks.
4. Map composite tasks to threads/cores.



```

#include <mpi.h>
#include <stdio.h>
#include <string.h>

const int MAX_STRING = 100;

int main(int argc, char** argv) {
    char greeting[MAX_STRING];
    int comm_sz; //number of processes
    int my_rank; //my process rank

    // Initialize the MPI environment
    MPI_Init(NULL, NULL);
    MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    if (my_rank != 0) {
        sprintf(greeting, "Greetings from processor %s, rank %d out of %d processors\n",
                processor_name, my_rank, comm_sz);
        MPI_Send(greeting, strlen(greeting)+1, MPI_CHAR, 0, 0, MPI_COMM_WORLD);
    } else {
        printf("This is process: %d, Greetings from processor %s\n",
               my_rank, processor_name);
        for (int q = 1; q < comm_sz; q++) {
            MPI_Recv(greeting, MAX_STRING, MPI_CHAR, q, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
            printf("%s", greeting);
        }
    }
    // Finalize the MPI environment.
    MPI_Finalize();
}

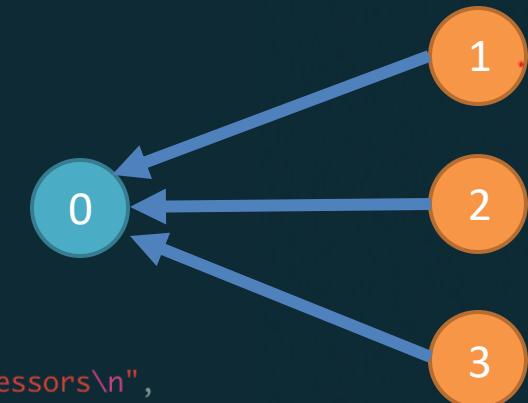
```

# Example: 1/2/3 send to 0 and 0 prints out messages

```

ksuo@ksuo-VirtualBox ~/cs7172> mpiexec -n 4 ./helloworld-mpi.o
This is process: 0, Greetings from processor ksuo-VirtualBox
Greetings from processor ksuo-VirtualBox, rank 1 out of 4 processors
Greetings from processor ksuo-VirtualBox, rank 2 out of 4 processors
Greetings from processor ksuo-VirtualBox, rank 3 out of 4 processors

```



Send to 0 from 1,...,comm\_z-1  
Kennesaw State University

# Sample code:

```
ksuo@ksuo-VirtualBox ~/cs7172> mpiexec -n 4 ./helloworld-mpi.o
This is process: 0, Greetings from processor ksuo-VirtualBox
Greetings from processor ksuo-VirtualBox, rank 1 out of 4 processors
Greetings from processor ksuo-VirtualBox, rank 2 out of 4 processors
Greetings from processor ksuo-VirtualBox, rank 3 out of 4 processors
```

<https://github.com/kevinsuo/CS4504/blob/master/helloworld-mpi-sample.c>

```
#include <mpi.h>
#include <stdio.h>
#include <string.h>

const int MAX_STRING = 100;

int main(int argc, char** argv) {
    char greeting[MAX_STRING];
    int comm_sz; //number of processes
    int my_rank; //my process rank

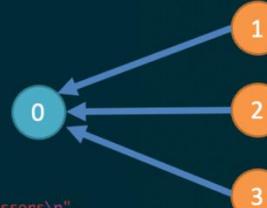
    // Initialize the MPI environment
    MPI_Init(NULL, NULL);
    MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    if (my_rank != 0) {
        sprintf(greeting, "Greetings from processor %s, rank %d out of %d processors\n",
                processor_name, my_rank, comm_sz);
        MPI_Send(greeting, strlen(greeting)+1, MPI_CHAR, 0, 0, MPI_COMM_WORLD);
    } else {
        printf("This is process: %d, Greetings from processor %s\n",
               my_rank, processor_name);
        for (int q = 1; q < comm_sz; q++) {
            MPI_Recv(greeting, MAX_STRING, MPI_CHAR, q, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
            printf("%s", greeting);
        }
    }
    // Finalize the MPI environment.
    MPI_Finalize();
}
```

**Example: 1/2/3 send to 0 and  
0 prints out messages**

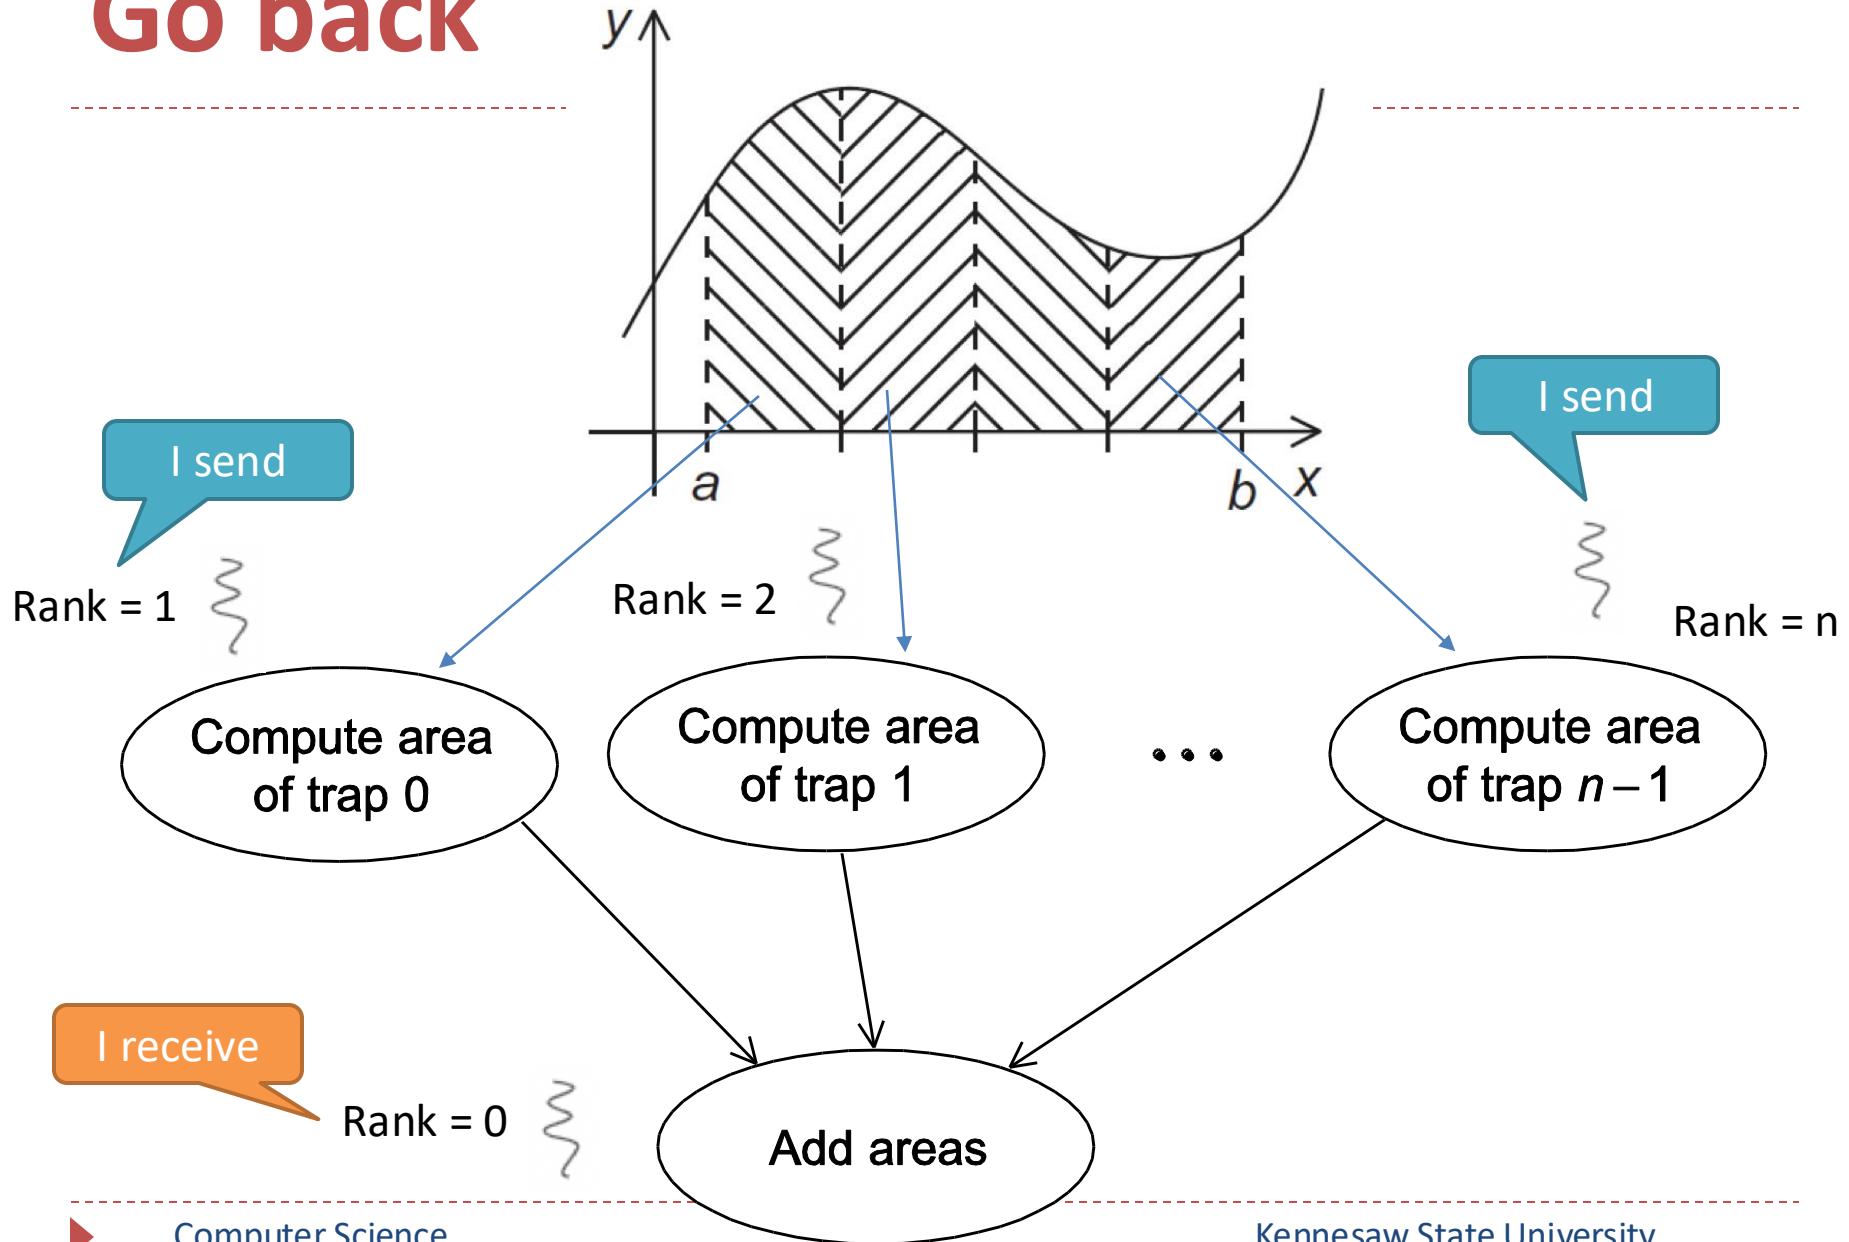
```
ksuo@ksuo-VirtualBox ~/cs7172> mpiexec -n 4 ./helloworld-mpi.o
This is process: 0, Greetings from processor ksuo-VirtualBox
Greetings from processor ksuo-VirtualBox, rank 1 out of 4 processors
Greetings from processor ksuo-VirtualBox, rank 2 out of 4 processors
Greetings from processor ksuo-VirtualBox, rank 3 out of 4 processors
```



Send to 0 from 1,...,comm\_z-1  
Kennesaw State University



# Go back



# Parallel pseudo-code

```
1      Get a, b, n;
2      h = (b-a)/n;
3      local_n = n/comm_sz;
4      local_a = a + my_rank*local_n*h;
5      local_b = local_a + local_n*h; Task
6      local_integral = Trap(local_a, local_b, local_n, h);
7      if (my_rank != 0)
8          Send local_integral to process 0;
9      else /* my_rank == 0 */
10         total_integral = local_integral;
11         for (proc = 1; proc < comm_sz; proc++) {
12             Receive local_integral from proc;
13             total_integral += local_integral;
14         }
15     }
16     if (my_rank == 0)
17         print result;
```



# First version (1)

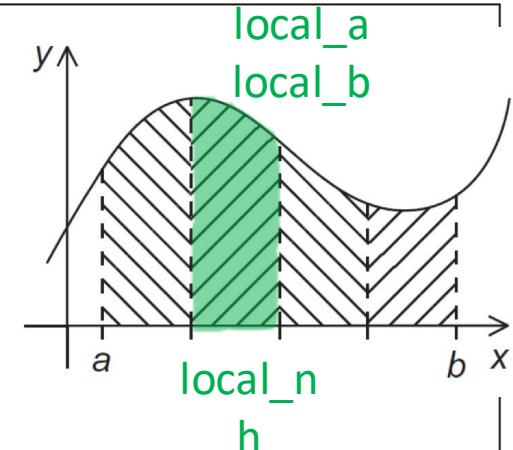
```
1 double Trap(
2     double left_endpt /* in */,
3     double right_endpt /* in */,
4     int trap_count /* in */,
5     double base_len /* in */) {
6     double estimate, x;
7     int i;
8
9     estimate = (f(left_endpt) + f(right_endpt))/2.0;
10    for (i = 1; i <= trap_count -1; i++) {
11        x = left_endpt + i*base_len;
12        estimate += f(x);
13    }
14    estimate = estimate*base_len;
15
16    return estimate;
17 } /* Trap */
```



Please read the code by yourself  
based on the pseudo-code

# First version (2)

```
1 int main(void) {
2     int my_rank, comm_sz, n = 1024, local_n;
3     double a = 0.0, b = 3.0, h, local_a, local_b;
4     double local_int, total_int;
5     int source;
6
7     MPI_Init(NULL, NULL);
8     MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
9     MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);
10
11    h = (b-a)/n;           /* h is the same for all processes */
12    local_n = n/comm_sz;   /* So is the number of trapezoids */
13
14    local_a = a + my_rank*local_n*h;
15    local_b = local_a + local_n*h;
16    local_int = Trap(local_a, local_b, local_n, h);
17
18    if (my_rank != 0) {
19        MPI_Send(&local_int, 1, MPI_DOUBLE, 0, 0,
20                  MPI_COMM_WORLD);
```



# First version (3)

```
21 } else {
22     total_int = local_int;
23     for (source = 1; source < comm_sz; source++) {
24         MPI_Recv(&local_int, 1, MPI_DOUBLE, source, 0,
25                  MPI_COMM_WORLD, MPI_STATUS_IGNORE);
26         total_int += local_int;
27     }
28 }
29
30 if (my_rank == 0) {
31     printf("With n = %d trapezoids, our estimate\n", n);
32     printf("of the integral from %f to %f = %.15e\n",
33            a, b, total_int);
34 }
35 MPI_Finalize();
36 return 0;
37 } /* main */
```



# Conclusion

---

- MPI introduction
  - Helloworld of MPI
- Performance evaluation
- Example: how to solve problems in MPI
  - Trapezoidal problem



# Practice 1

[https://github.com/kevinsuo/CS4504  
/blob/master/helloworld-mpi.c](https://github.com/kevinsuo/CS4504/blob/master/helloworld-mpi.c)

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    // Initialize the MPI environment
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    // Print off a hello world message
    printf("Hello world from processor %s, rank %d out of %d processors\n",
           processor_name, world_rank, world_size);

    // Finalize the MPI environment.
    MPI_Finalize();
}
```

Thread 0:

print out “000000”

Thread 1 & 2:

print out “121212”

Thread 3:

print out “333333”



# Practice 2

<https://github.com/kevinsuo/CS4504/blob/master/helloworld-mpi-sample.c>

```
#include <mpi.h>
#include <stdio.h>
#include <string.h>

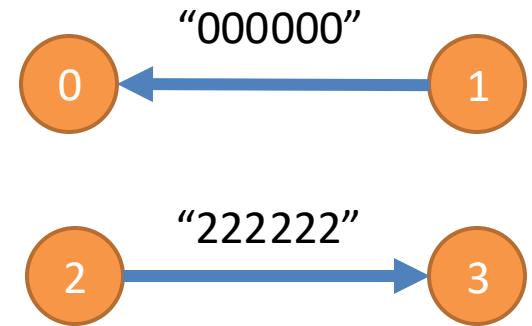
const int MAX_STRING = 100;

int main(int argc, char** argv) {
    char greeting[MAX_STRING];
    int comm_sz; //number of processes
    int my_rank; //my process rank

    // Initialize the MPI environment
    MPI_Init(NULL, NULL);
    MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    if (my_rank != 0) {
        sprintf(greeting, "Greetings from processor %s, rank %d out of %d processors\n",
                processor_name, my_rank, comm_sz);
        MPI_Send(greeting, strlen(greeting)+1, MPI_CHAR, 0, 0, MPI_COMM_WORLD);
    } else {
        printf("This is process: %d, Greetings from processor %s\n",
               my_rank, processor_name);
        for (int q = 1; q < comm_sz; q++) {
            MPI_Recv(greeting, MAX_STRING, MPI_CHAR, q, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
            printf("%s", greeting);
        }
    }
    // Finalize the MPI environment.
    MPI_Finalize();
}
```



# Practice 3

<https://github.com/kevinsuo/CS4504/blob/master/helloworld-mpi-sample.c>

```
#include <mpi.h>
#include <stdio.h>
#include <string.h>

const int MAX_STRING = 100;

int main(int argc, char** argv) {
    char greeting[MAX_STRING];
    int comm_sz; //number of processes
    int my_rank; //my process rank

    // Initialize the MPI environment
    MPI_Init(NULL, NULL);
    MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    if (my_rank != 0) {
        sprintf(greeting, "Greetings from processor %s, rank %d out of %d processors\n",
                processor_name, my_rank, comm_sz);
        MPI_Send(greeting, strlen(greeting)+1, MPI_CHAR, 0, 0, MPI_COMM_WORLD);
    } else {
        printf("This is process: %d, Greetings from processor %s\n",
               my_rank, processor_name);
        for (int q = 1; q < comm_sz; q++) {
            MPI_Recv(greeting, MAX_STRING, MPI_CHAR, q, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
            printf("%s", greeting);
        }
    }
    // Finalize the MPI environment.
    MPI_Finalize();
}
```

